Tools to Design a Biomimetic Humanoid Robot Actuated by Pneumatic Artificial Muscles and a Synthetic Neural Network

by

Ben Bolen

A dissertation submitted in partial fulfillment of the
requirements for the degree of

Doctor of Philosophy
in
Mechanical Engineering

Dissertation Committee:
Alexander Hunt, Chair
Bin Jiang
Radhika Reddy
David Turcic

Portland State University
2026

©Copyright by Ben Bolen
July, 2026
All Rights Reserved

Abstract

Biomimetic robots offer a platform for studying neuromechanical control through experiments that would be impractical or unethical with living subjects. This dissertation advances a long-running effort at Portland State University to build a humanoid robot, from the trunk down, actuated by pneumatic artificial muscles and controlled by a synthetic nervous system. The isometric force of Festo braided pneumatic actuators (BPAs) was characterized over a wide range of resting lengths, contractions, and pressures using a custom test jig. The analysis produced novel equations for the maximum isometric force as a function of resting length and a three-coefficient nondimensional force surface that predicts isometric force from relative strain and relative pressure. The actuator model was carried to the joint level through isometric knee torque experiments on two test stands: a simplified pinned knee and a biomimetic knee with a migrating instantaneous center of rotation. This approach isolated the sources of prediction error, revealing constant length offsets, series compliance in brackets and artificial tendons, and a loss of usable actuator length when BPAs wrap around the joint. The biomimetic knee driven by optimized BPA placements was shown to meet or exceed human isometric knee torque over most of the range of motion. Tools were created for curating relevant neuromechanical research, as well as simulating neuromechanical models based on human and robot biomechanical models.

Dedication

*

Darian

azīzam, jānam,

nūr-e cheshmam — the light of my eyes

qorbānat beram

Bābā kheyli dūset dāre

↘ ↘ ↘

"Los corazones obreros que laten

por nuestra causa, felices serán.

Si entusiasmados y unidos combaten,

de la victoria, la palma obtendrán."

Acknowledgments

This work was funded by the Department of Mechanical and Materials Engineering at Portland State University, the National Science Foundation (NSF) grant for NeuroNex: Communication, Coordination, and Control in Nervous Systems (C3NS) #2015317, NSF grant #1943483, and the NIH BUILD EXITO program (grant numbers 5RL5GM118963-06, 5TL4GM118965-06, 5RL5GM118963-07, and 5TL4GM118965-07).

I would like to thank my advisor, Dr. Alexander Hunt, for his guidance, patience, and the freedom to pursue this project, and my committee members — Dr. Bin Jiang, Dr. Radhika Reddy, and Dr. David Turcic — for their time and thoughtful feedback.

I am grateful to my collaborators and labmates: Mohammad Elzein and Lawrence Pang for their work on the actuator characterization experiments; Lindie Burgess and Connor Morrow for their work on the knee torque experiments; Cody Scharzenberger for building the Muscle Mutt platform and for his assistance proofreading; Alex Steele, who designed the initial biomimetic four-bar knee linkage used in the torque tests and generously shared solid models and prototype experience; and Jasmine Bradley, whose work reworking the figures improved their aesthetic quality and accessibility for readers with color vision deficiencies.

Finally, I thank my family for their support throughout this long process.

Contents

Abstract **ii**

Dedication **iii**

Acknowledgments **iv**

1 Introduction **1**

1.1 Motivation 1

1.2 The Biomimetic Humanoid Robot Project 4

1.3 Research Objectives 5

1.4 Dissertation Organization 6

2 Background **9**

2.1 Biomimetic and Musculoskeletal Robots 9

2.2 Robot Actuation Technologies 10

2.2.1 Overview 10

2.2.2 Braided Pneumatic Actuators 14

2.2.3 BPA Force–Length–Pressure Models 16

2.3 Human Lower-Limb Biomechanics 18

2.4 Biomimetic Knee and Bipedal Robot Model 19
2.5 Neural Control of Locomotion: Central Pattern Generators 21
2.6 Sensory Feedback in Locomotion 24
2.7 Computational Neuroscience 25
2.8 Synthetic Nervous Systems 25
2.9 Neuromechanical Simulation 28
2.9.1 Simulation Tools 29
2.10 Testing Neuromechanical Hypotheses on Robots: The Ethical Case 31

3 Isometric Force Characterization of Braided Pneumatic Actuators 35
3.1 Methods 36
3.1.1 Actuator Force Model 36
3.1.2 BPA Isometric Force Characterization Experiment 37
3.2 Results 39
3.2.1 BPA Force Characterization 39
3.3 Discussion 45
3.3.1 Effects of Resting Length 45
3.3.2 Nondimensionalized Force Prediction 50
3.3.3 Comparison with Other Published BPA Models 51
3.3.4 Implications for Real-Time Control and Simulation 55
3.3.5 Biomimetic Artificial Muscle Analogy 56

4 Isometric Knee Torque 60
4.1 Methods 61

4.1.1 Robot Leg Model and Knee Test Stands 61
4.1.2 Knee Torque Data Collection and Equipment 65
4.1.3 Hybrid Torque Calculation Method 66
4.1.4 Improved Torque Model 67
4.1.5 Optimization Algorithms and Cost Functions 83
4.2 Results 87
4.2.1 Flexor on Pin Joint 88
4.2.2 Extensor on Pin Joint 93
4.2.3 Extensor on Biomimetic Knee 96
4.2.4 Flexor on Biomimetic Knee 97
4.3 Discussion 101
4.3.1 Sources of Error in Isometric Torque Generation 101
4.3.2 Follow-Up Identification and Route Redesign 103
4.3.3 Mono- and Bi-Pulley Tendon Routing 108
4.3.4 Physical Tests of the Redesigned Routes 108

5 The Sensory Afferent Database 112
5.1 The Resource 112
5.2 Curation Methods 113
5.2.1 AI-Assisted Triage 115
5.3 Coverage and Use 116
5.4 Discussion 116
5.4.1 Use in Neural-Controller Design 116

5.4.2 Further improvement and work 117

6 Synthetic Nervous System Simulation 119

6.1 The AnimatLab Bipedal Walker 119

6.2 Knowledge Gaps and the New Tools 120

6.3 Software Tools 123

6.3.1 Model Conversion and Actuator Coupling 123

6.3.2 Support, Stimulation, and Parameter Search 125

6.3.3 Graphical Circuit Construction 126

6.4 Verification of the Implemented Interfaces 129

6.4.1 Reflex Response 129

6.4.2 Published Neural Behaviors 129

6.4.3 Matched Neural Implementations 131

6.4.4 Closed-Loop Model Coupling 134

6.4.5 Perturbation Robustness and Pathway Deletions 135

6.4.6 Feedback and Reduced Hand Drive 138

6.5 Capabilities and Remaining Limits 138

7 Discussion 140

7.1 The Work as a Whole 140

7.2 Future Work 143

7.2.1 Expanding the BPA Force Characterization 143

7.2.2 Expanding Attachment and Routing Design 146

7.2.3 Expanding the Sensory Afferent Database 148

7.2.4 Closing the Loop in Simulation . . . 148
7.2.5 Extending the Sensorium . . . 150
7.2.6 Real-Time Hardware Realization . . . 151
7.2.7 Toward Dynamic Whole-Body Maneuvers . . . 154

8 Conclusion **156**

Appendices **158**
A Appendix A: Muscle Force and Joint Torque Calculations . . . 159
A.1 Human muscle force calculations . . . 159
A.1.1 OpenSim Thelen2003 and Millard2012 Formulations . . . 159
A.1.2 OpenSim Moment Arms and Human Torque Targets . . . 160
A.2 Muscle Model Formulations and Cross-Platform Force Comparison . . . 161
A.2.1 MuJoCo Muscle Actuator (as converted) . . . 161
A.2.2 Cross-Platform Force Comparison . . . 161
A.3 BPA Force Model Implementation . . . 163
A.3.1 Goodness-of-Fit Metrics . . . 163
A.3.2 Strain Measures . . . 165
A.3.3 Normalized Force Surface . . . 165
A.3.4 Maximum Force . . . 166
A.4 Muscle Path Length and Moment Arm . . . 166
A.5 Compliance Model Implementation . . . 167
A.6 Wrapping-Loss Implementation . . . 169
B Appendix B: Neuron Model Equations . . . 170

B.1 Spiking Neuron Model . . . 170
B.2 Non-Spiking (Graded) Chemical Synapses . . . 171
B.3 Relation to the SNS-Toolbox Implementation . . . 172
B.4 Spiking-Mirror Implementation, Calibration, and Verification . . . 173
B.4.1 Adapting Integrate-and-Fire Cells and Hybrid Synapses . . . 173
B.4.2 Calibration of Record . . . 175
B.4.3 Verification Gates of the MuJoCo Mirror . . . 176
B.4.4 AnimatLab Conversion Variants . . . 176
B.5 Connectomes of the Spinal-Network Models . . . 181
C Appendix C: Model Geometry, Bracket Frames, and Compliance Parameters . 185
C.1 Static BPA Pathway Data . . . 185
C.2 Joint Transformation Matrices . . . 187
C.3 Bracket Reference Points and Frames . . . 190
C.4 Stiffness Arrays and Compliance . . . 192
C.5 Adopted Correction-Term Values . . . 194
C.6 Redesigned Extensor Route Geometry . . . 196
C.7 Supporting Route Searches and Transmission Sweeps . . . 197
C.8 Supporting Torque-Prediction Fit Statistics . . . 208
D Appendix D: Supplementary Neuromechanical Verification . . . 210
D.1 Model Construction and Recording . . . 210
D.2 Controller Ports and Reference Checks . . . 210
D.2.1 Archived AnimatLab Models . . . 214
D.3 Tuning and Supported Evaluations . . . 214

D.3.1 Reference-Library Comparison 216
D.3.2 Pathway-Removal Tests 216
D.4 Spiking-Neuron Comparisons 217
D.4.1 MuJoCo Network and Ground Evaluation 218
D.4.2 Simulink Demonstrations 221
D.4.3 AnimatLab Conversions 221
D.5 Supported BPA-Robot Evaluations 223
E Appendix E: Supporting SADb Records 225
E.1 Record Sets and Coverage 225
E.2 Afferent Classes and Linked Pathways 226

List of Tables

2.1 Coefficients of the 10 mm Festo BPA pressure model of Equation (2.1), corrected from **?**. 17

3.1 Maximum force equation coefficient values with confidence intervals (CI) and goodness-of-fit measures. Equation (3.5) is compared to data taken at 620 kPa. Equations (3.7) and (3.8) are compared against maximum force from using the Festo tool. It was not necessary to fit an adjusted R^2 in these cases. 44

3.2 Normalized isometric BPA force equation (Equation (3.4)) coefficient values and goodness-of-fit measures. 44

3.3 BPA force-model comparison. Diameter and resting length are in mm; RMSE and maximum absolute error are in N. FVU is dimensionless. 59

4.1 Pre-optimization pinned-knee flexor torque-prediction goodness of fit per test, comparing the three pre-optimization torque predictions of Figure 4.6: the Hunt-equation rigid-body prediction, the unoptimized prediction of the improved BPA characterization, and the back-calculated (hybrid) torque. RMSE and maximum absolute residual (Max.) are in N m, and FVU denotes the fraction of variance unexplained. 90

4.2 Pinned-knee flexor torque-prediction goodness of fit: the improved BPA characterization alone (the zero-correction prediction of the text) and the corrected model at the adopted correction terms (Table 4.4); training and validation tests follow the adopted cross-validation fold. RMSE and maximum absolute residual (Max.) are in N m, and FVU denotes the fraction of variance unexplained. 92

4.3 Pinned-knee extensor torque-prediction goodness of fit, at the adopted correction terms with the effective-stiffness pair locked to the flexor-identified values (Table 4.4): the improved BPA characterization alone (the zero-correction prediction of the text) and the corrected model. RMSE and maximum absolute residual (Max.) are in N m, and FVU denotes the fraction of variance unexplained. 96

4.4 Adopted transmission corrections. The stiffness pair is shared. The wrapping factor was identified on the extensor and is also used in the flexor redesign; it was not fitted to the pinned-flexor tests. 103

List of Figures

1.1 Dissertation roadmap 7

2.1 Kengoro human-mimetic musculoskeletal humanoid 11
2.2 Representative robotic actuator technologies 13
2.3 Festo BPAs with and without inflation 15
2.4 Steele biomimetic knee mechanism 20
2.5 Robot skeleton and robot/human muscle paths 21
2.6 Equivalent geometric constructions of joint torque 22
2.7 Biceps femoris long-head torque comparison from the master's study 23
2.8 Rybak two-level locomotor CPG concept 26
2.9 Shevtsova extended CPG model with two-joint-muscle circuits 26
2.10 Decerebrate-cat treadmill experiment 32

3.1 A picture of a BPA in the isometric force test stand with components labeled. 40

3.2 Results for finding the relationship between F_{620} and l_0. (**a**) $F_{620_{10}}$ as a function of l_0, at P_{620}. Dashed line is the $F_{620_{10}}$ data from Festo. Solid line is the fit from Equation (3.7). (**b**) $F_{620_{20}}$ as a function of l_0 at P_{620}. Dashed line is $F_{620_{20}}$ data from Festo. Solid line is the fit from Equation (3.8). 42

3.3 Results for finding the relationship between F_{620}, l_0, P_{620}, and $\epsilon_{620_{10}}$. (**a**) Isoclines of the surface fit for $F_{620_{10}}(l_0, P_{620})$. (**b**) $\epsilon_{620_{10}}$ versus l_0 at P_{620}. Although there is a general trend of longer resting lengths producing more contraction, no conclusive relationship between ϵ_{620} and l_0 could be deduced from this experiment. The error bars show the effect of being ± 1 mm in measuring l_0 and l_{620}. 46

3.4 Surface fit for $F^*(\epsilon^*, P^*)$. Solid lines are from our model. Dashed lines from Festo supplied data. In total, 537 data points were collected and used to fit the 3D surfaces (321 data points for ∅10 mm BPAs and 216 data points for ∅20 mm BPAs). In total, 80% of these data were used for training, and 20% were used for validation. For figure clarity, not all data is included and plotted circles represent collected data at ± 10 kPa of the stated pressure. (**a**) Fit data for ∅10 mm Festo BPAs. (**b**) Fit data for ∅20 mm Festo BPAs. 47

3.5 Comparison between measured and predicted BPA force at varying amounts of contraction and pressure. A $y = x$ perfect fit line is added for reference. The BPAs compared are: (**a**) ∅10 mm, $l_0 = 257$ mm, (**b**) ∅10 mm, $l_0 = 233$ mm, (**c**) ∅20 mm, $l_0 = 300$ mm, and (**d**) ∅20 mm, $l_0 = 450$ mm. 53

4.1 **(A)** CAD solid model assembly of the revolute joint knee in the test stand. **(B)** 1-DoF pinned-joint robot knee in the test apparatus with important components labeled. Leg is shown with 30 degrees flexion, i.e. $\theta_k = -30^\circ$. **(C)** Picture from isometric torque testing. Setup shows the revolute knee configured for a test of the extensor BPA. The knee is flexed with $\theta_k = -120^\circ$. Also note the compressed shape of the inflated BPA on the knee during this high degree of flexion. 62

4.2 ICR for biomimetic robot knee as a function of knee angle. Measured from θ_1 on the tibia. Plot is overlayed on model of the knee at $\theta_k = 0^\circ$ (solid line), and -60° and -120° (dotted dash lines). White circles represent measured points in -15° increments from -150° to -15°, -5° increments from -15° to 0°, and then at 2°, 5°, and 10°. 63

4.3 Two-rotation construction of the bracket frame 71

4.4 Bracket frame orientations for: **(A)** flexor configuration on the pinned knee, **(B)** extensor configuration on the pinned knee, **(C)** flexor and extensor configurations on the biomimetic knee with ∅10 mm BPAs, and **(D)** flexor configuration on the biomimetic knee with the ∅20 mm BPA. 72

4.5 Compliance projections and BPA wrapping geometry 79

4.6 Pre-optimization pinned-knee torque versus knee angle for the four cross-validation flexor tests. Knee angle is negative in flexion, and flexion torque is negative. Gold circles are the hybrid calculation, which combines the measured muscle length and moment arm with the modeled force; indigo filled circles are experimentally measured torque; the gray dashed line is the rigid-body prediction; and the light orange dash-dot line is the unoptimized prediction, the improved BPA force characterization applied to the same rigid kinematics. Panel titles give the nominal resting length; the actual resting lengths are $l_0 = 48.5\,\text{cm}$ **(A)**, $l_0 = 41.0\,\text{cm}$ **(B)**, $l_0 = 45.7\,\text{cm}$ **(C)**, and $l_0 = 40.6\,\text{cm}$ with tendon **(D)**. Panels **(A)** and **(B)** are the training tests and panels **(C)** and **(D)** the validation tests of the adopted cross-validation fold. 89

4.7 Optimized pinned-knee torque versus knee angle for the four cross-validation flexor tests, at the adopted correction terms (Table 4.4). Knee angle is negative in flexion, and flexion torque is negative. Blue filled circles are experimentally measured torque; the dashed light orange line is the unoptimized prediction of the improved BPA characterization; the magenta line is the optimized prediction of the corrected model. Panel titles give the nominal resting length; the actual resting lengths are $l_0 = 48.5\,\text{cm}$ **(A)**, $l_0 = 41.0\,\text{cm}$ **(B)**, $l_0 = 45.7\,\text{cm}$ **(C)**, and $l_0 = 40.6\,\text{cm}$ with tendon **(D)**. Panels **(A)** and **(B)** are the training tests and panels **(C)** and **(D)** the validation tests of the adopted cross-validation fold; the goodness-of-fit values are given in Table 4.2. 91

4.8 Pinned-knee torque versus knee angle for the extensor BPAs: the six identification-pool tests at the adopted correction terms, with the extensor-specific terms (χ_0, χ_3) on the shared stiffness pair (Table 4.4). Knee angle is negative in flexion, and extension torque is positive. Gold circles are the hybrid calculation; indigo filled circles are experimentally measured torque; the dashed light orange line is the unoptimized prediction of the improved BPA characterization; the magenta line is the optimized prediction of the corrected model. Panel titles give the nominal resting length; the actual resting lengths are $l_0 = 43.6\,\text{cm}$ **(A)**, $l_0 = 43.2\,\text{cm}$ with tendon **(B)**, $l_0 = 45.7\,\text{cm}$ **(C)**, $l_0 = 40.0\,\text{cm}$ **(D)**, $l_0 = 40.0\,\text{cm}$ with tendon **(E)**, and $l_0 = 46.5\,\text{cm}$ **(F)**. Panels **(A)–(C)** are the training tests and panels **(D)–(F)** the held-out validation tests of the adopted cross-validation fold; the legend is given in panel **(A)**. 94

4.9 Measured versus expected isometric torque results using a $\varnothing 10\,\text{mm}$, $l_0 = 51.8\,\text{cm}$ extensor BPA on the biomimetic knee. Measured results are blue dots. Original prediction is a dashed black line. Optimized prediction is a solid magenta line. Hybrid calculation shown in yellow dots. 97

4.10 A configuration that is particularly hard to model. The BPA passes between the femoral condyles close to ICR. 98

4.11 Biomimetic-knee torque versus knee angle for the flexor BPAs. Knee angle is negative in flexion, and flexion torque is negative. **(A)** ∅10 mm, $l_0 = 41.5$ cm actuator with a 21 mm artificial tendon at 604 kPa. **(B)** ∅20 mm, $l_0 = 41.5$ cm actuator with a 14 mm artificial tendon at 620 kPa. Indigo filled circles are measured torque; the magenta line is the optimized prediction of the corrected model using the flexor terms identified on the pinned knee; the dashed light orange line is the unoptimized prediction of the improved BPA characterization alone. In **(B)**, the black dotted line is the maximum human isometric torque, and the pink circles are measured points at the reduced pressures 560 kPa, 421 kPa, 325 kPa and 281 kPa (from left to right, respectively), recorded because full pressure would load the printed femur beyond its ultimate strength. **(C)** The redesigned two-BPA flexor route . . . 100

4.12 Cam-extensor route geometry and hardware . . . 107

4.13 Physical tests of the cam-extensor route . . . 110

4.14 Physical tests of the redesigned flexor route . . . 111

5.1 SADb citation-neighborhood view . . . 114

5.2 The SADb Explorer pivot view . . . 117

6.1 AnimatLab walker and controller hierarchy . . . 121

6.2 Simulation-to-robot pipeline . . . 127

6.3 Simulink SNS library blocks and neuron implementation . . . 128

6.4 Reduced-order knee-reflex circuit topology . . . 130

6.5 W2L rhythm-core traces across engines . . . 132

6.6 W2L perturbations with the stored conductances applied 137

7.1 The first-pass design workflow the enablers make possible 144

7.2 Construction of the braided pneumatic actuator 147

7.3 Control stack with the planned supraspinal additions 152

1 Introduction

1.1 Motivation

Biomimetic robots provide a platform for investigating how mechanical structure and control interact to generate movement, offering complementary insights to simulation-based studies [**???**]. Wheeled robots remain the platform of choice in structured environments, where their speed, efficiency, and payload capacity are difficult to match; in the unstructured terrain of the natural and built world, legged machines hold the advantage, and biomimetic legs that reproduce human force capabilities have the potential to excel where wheels cannot [**??**]. Musculoskeletal humanoids and limb systems driven by tendon-like actuators have been shown to support biologically meaningful experiments on motion generation, internal force transmission, and neural control hypotheses [**?????**]. The dual perspective of using robots to learn about biology, and using biology to design more adaptive robots, has been emphasized in both biomimetic robotics and biomechanics [**??**]. With such robotic systems, experiments can be performed that would not be practical or ethical with human test subjects, such as removing a specific muscle from the system or blocking a sensory feedback pathway. Experiments of this kind allow direct testing of how the components of the neuromechanical system interact, and they offer new opportunities for testing neural control theories that are faster and cheaper than human observation studies [**???**].

Artificial muscles are a key enabling technology for this class of robot. Among the

available actuation technologies, McKibben-style braided pneumatic actuators (BPAs) combine low mass, high force-to-weight ratio, and intrinsic compliance with a force–length curve that qualitatively resembles that of skeletal muscle [**??**]. BPAs have been used in quadrupedal robots [**???**], musculoskeletal humanoids [**??**], and dexterous manipulation systems [**?**]. However, BPAs also differ from biological muscle in important ways: their maximum tensile force occurs at resting length rather than at an optimal fiber length, their maximum contraction ratio is substantially lower than that of human muscle, and their force output is a highly nonlinear function of internal pressure, contraction, and loading state [**???**]. My experience with 10 mm actuators made this gap concrete: manufacturer specifications overstate the force and contraction available at the short resting lengths a leg design requires, and designs based on those specifications delivered less range of motion and joint torque than intended. Nor could the published empirical models fill in the gaps, because each is fitted to measurements at a single resting length [**???**]; none accurately predicts how force changes when a designer chooses a different actuator length. Because muscle resting length is one of the primary degrees of freedom available to the robot designer, the absence of a force model valid at arbitrary resting length blocked accurate actuator selection. Designing a biomimetic robot around BPAs therefore requires quantitative, predictive models of the actuator and of the joint torques it can produce.

Actuator force alone, however, does not determine the torque a joint produces. The robot knee transmits BPA force to the skeleton through mounting brackets, artificial tendons, and paths that wrap over the joint, and each of these elements departs from the rigid-body idealization: brackets and fixtures deflect under load, tendons stretch, and a wrapping actuator loses usable length as the knee flexes. At the start of this work it

was unknown how much these transmission effects reduce the torque delivered at the knee relative to the rigid-body prediction, and no method existed for correcting that prediction in a form that would transfer between joint configurations. The biomimetic knee's optimized muscle routing [**?**] compounded the problem: it had never been validated experimentally or compared against human isometric torque data, so even the nominal torque envelope of the robot leg rested on unverified calculations.

Attempts to implement neural control models exposed two further gaps. First, the project lacked a structured, searchable resource connecting measured sensory-afferent properties to the species, preparations, and pathway roles on which controller design depends. Those findings were distributed across primary papers and reviews [**???**], and were difficult for researchers entering the field to locate and compare. Second, the project lacked a tested framework for connecting neural controllers to its human musculoskeletal reference and BPA-actuated models, with control over mechanical support, sensory inputs, and disturbances. Published CPG architectures and simulation tools supplied components [**????**], but this project still needed to combine and test them in the human and BPA-routed robot models. The SADb and simulation chapters address these evidence and implementation gaps, respectively.

The dissertation therefore addresses four connected needs: an actuator-force model that accounts for resting length, corrections for transmission losses at the knee, a curated sensory-afferent resource, and tools for developing and checking neuromechanical simulations. Together, they support the design and evaluation of the robot; achieving unsupported walking remains a separate control and integration problem.

1.2 The Biomimetic Humanoid Robot Project

This dissertation is part of a long-running effort at Portland State University to build a biomimetic humanoid robot from the trunk down, actuated by pneumatic artificial muscles and controlled by a synthetic nervous system (SNS). The project is deliberately iterative. Knowledge from biomechanics and neuroscience guides the design of robots intended to perform complex tasks in unstructured environments with improved efficiency, balance, adaptability, and resilience. Those robots then become physical models with which biological hypotheses can be tested, validated, and refined through controlled changes that would be impractical or unethical in living subjects. The resulting biological insight informs the next generation of robot design, and lessons from both directions can be translated into improved prosthetic and orthotic devices.

The present work builds directly on a series of studies by me and my collaborators. In my master's project, the kinematics and dynamics of a biomimetic robot leg were matched against the Gait2392 OpenSim model, artificial muscle attachment and wrapping locations were determined for the robot legs, and MATLAB software was developed for evaluating muscle lengths, moment arms, and isometric joint torques over the full range of motion [**??**]. That work showed that simply transplanting human muscle attachment locations onto a BPA-actuated robot does not reproduce human torque profiles; the placement optimization that closes this gap was completed during this dissertation and is summarized in Section 7.1. The rest of the platform advanced in parallel, and this dissertation builds on its hardware and demonstrations: the biomimetic four-bar knee linkage with a migrating instantaneous center of rotation, designed and validated by Steele [**???**], is the knee of

the robot leg whose torque transmission Chapter 4 measures, and the canine-inspired quadruped platform Muscle Mutt, formerly called DoggyDeux, demonstrated antagonistic BPA pairs under synthetic nervous system control on a physical robot [**?**], extending an earlier AARL demonstration of central pattern generator control of a pneumatically actuated robot [**?**]. The humanoid leg of this dissertation carries that line forward.

1.3 Research Objectives

This dissertation investigates actuator force and knee torque experimentally, consolidates sensory-afferent evidence in the Sensory Afferent Database [**??**], and develops neural-model and simulation interfaces with representative verification checks. The objectives are:

1. **Characterize the isometric force of Festo BPAs** with uninflated diameters of 10 mm and 20 mm over a wide range of resting lengths, contractions, and pressures; derive improved maximum-force and normalized force–strain–pressure models; and compare them against existing models (Chapter 3).

2. **Validate and improve the prediction of isometric knee torque** produced by BPAs on pinned-joint and biomimetic four-bar knees; identify and model losses caused by constant length offsets, series compliance, and actuator wrapping; and compare the resulting torque envelopes with human capabilities (Chapter 4).

3. **Consolidate the sensory afferent literature into a structured afferent resource**, the Sensory Afferent Database, by auditing and curating the published record of afferent

properties across animal species and their role in motor control, and by evaluating AI-assisted literature triage for expanding the corpus (Chapter 5).

4. **Develop tools and a framework for neuromechanical control** of the humanoid robot by converting its OpenSim musculoskeletal model to MuJoCo, implementing an SNS-Toolbox two-level spinal CPG with proprioceptive and heel/toe contact feedback, and establishing a staged tuning and verification workflow (Chapters 6 and 7).

1.4 Dissertation Organization

Chapter 2 introduces the robot, BPA actuation, musculoskeletal design, locomotor control, and the scientific rationale for robotic models. Chapter 3 presents the isometric force characterization and actuator models, adapted from my publication [**?**]. Chapter 4 tests the joint-torque predictions, identifies compliance and wrapping corrections, and uses them to redesign actuator routes.

Chapter 5 presents the Sensory Afferent Database, its curation, and its use limits. Chapter 6 begins with the AnimatLab attempt to combine the Deng CPG and Li walker, identifies the development requirements exposed by that attempt, and presents the tools and representative checks built to address them. Chapter 7 discusses the contributions together and develops the future work; Chapter 8 states the conclusions. Figure 1.1 summarizes these connections.

Appendices A–C provide the muscle calculations, neural equations, model geometry, and supporting mechanical results. Appendix D collects the larger simulation evaluations,

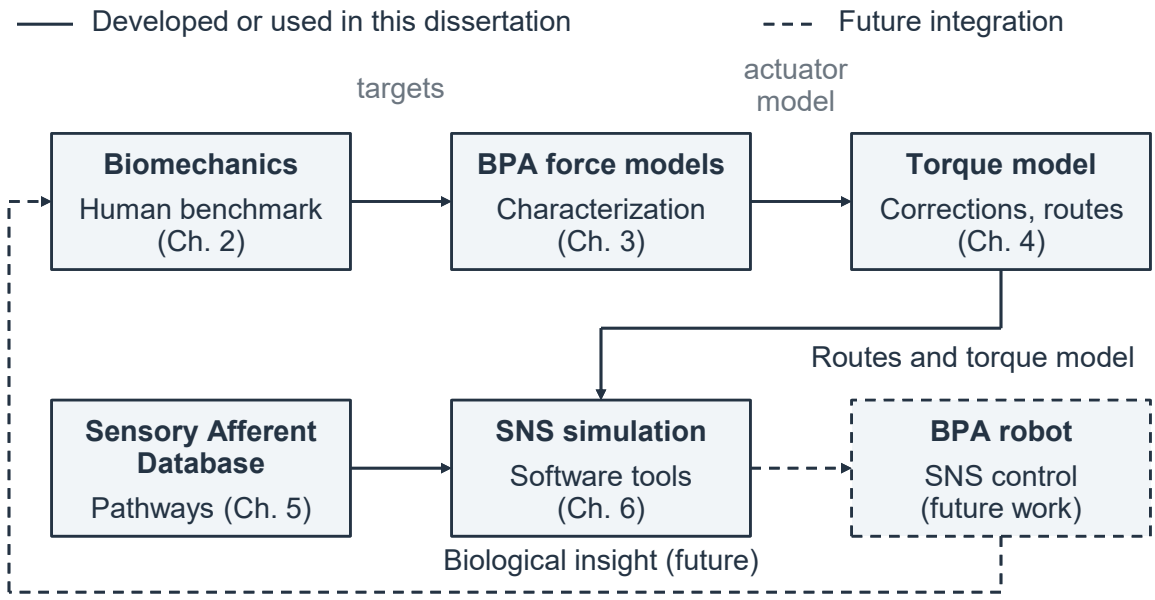


Figure 1.1: Dissertation roadmap. Biomechanical targets guide actuator characterization and route design; the sensory-afferent resource informs feedback selection; and simulation tools support controller evaluation. Dashed connections indicate future integration and the intended return of biological insight.

and Appendix E provides the database coverage and pathway tables.

2 Background

This chapter reviews the established knowledge against which the dissertation's contributions are defined: biomimetic and musculoskeletal robotics, the actuation technologies available to them, the human lower-limb biomechanics they imitate, the neuroscience of locomotor control they implement, and the simulation and ethical context in which they are tested. Work from my own program that goes beyond this background, including the prior AnimatLab bipedal walker, the Sensory Afferent Database, and the neuromechanical simulation pipeline, is presented as dissertation work in Chapters 5 and 6 rather than repeated here.

2.1 Biomimetic and Musculoskeletal Robots

Biomimetic robots offer many potential applications, ranging from assisting in factories and environmental exploration to serving as physical models of human movement. By developing high-fidelity biomimetic robots, we can perform tests that help improve our understanding of both biomechanics and the underlying neuromechanical systems that control movement [**??????**]. A few such biomimetic robots exist, such as the impressive Kengoro and Kenshiro robots of **?** and **?**, which have proportionally human bodies, anatomically correct musculotendon arrangements, rolling-sliding joints, and joint ranges and output power comparable with human-measured values; their artificial muscles are implemented with brushless DC motors, wires as tendons, gears, and integrated tension measurement. Kengoro, the more recent of the two, stands 1670 mm tall, weighs 56.5 kg,

and carries 174 degrees of freedom when the end effectors are included [**?**]. A dog-like robot developed by **?** and actuated with Festo BPAs was designed from a planar kinematic model scaled to an adult greyhound, with three sagittal degrees of freedom per leg driven by antagonistic BPA pairs; its front-leg pair alone lifted a 13.5 kg payload, almost double the complete robot's predicted mass. That platform was used to show how central pattern generators provide robust locomotion patterns that are adaptive to perturbations in ground speed [**?**].

Robots built with artificial muscles thus provide opportunities to investigate how muscles are controlled to produce desired motions; however, artificial muscles do not exactly replicate muscle properties, and existing robots do not capture all of the degrees of freedom and joint torques of the animal limbs they mimic. The biomimetic leg developed and tested in this dissertation sits squarely in this lineage: a human-proportioned skeleton, muscle-like actuators, and the goal of human-matched joint torques.

Figure 2.1 illustrates the human-scale skeletal structure and tendon routing of this approach.

2.2 Robot Actuation Technologies

2.2.1 Overview

A wide range of actuation technologies have been explored in robotics, including electric motors, hydraulic systems, smart-material actuators, and artificial muscles [**??**]. Electric motors are widely used due to their controllability and maturity, but they are limited by thermal constraints, gearing requirements, and reduced compliance compared to muscle-

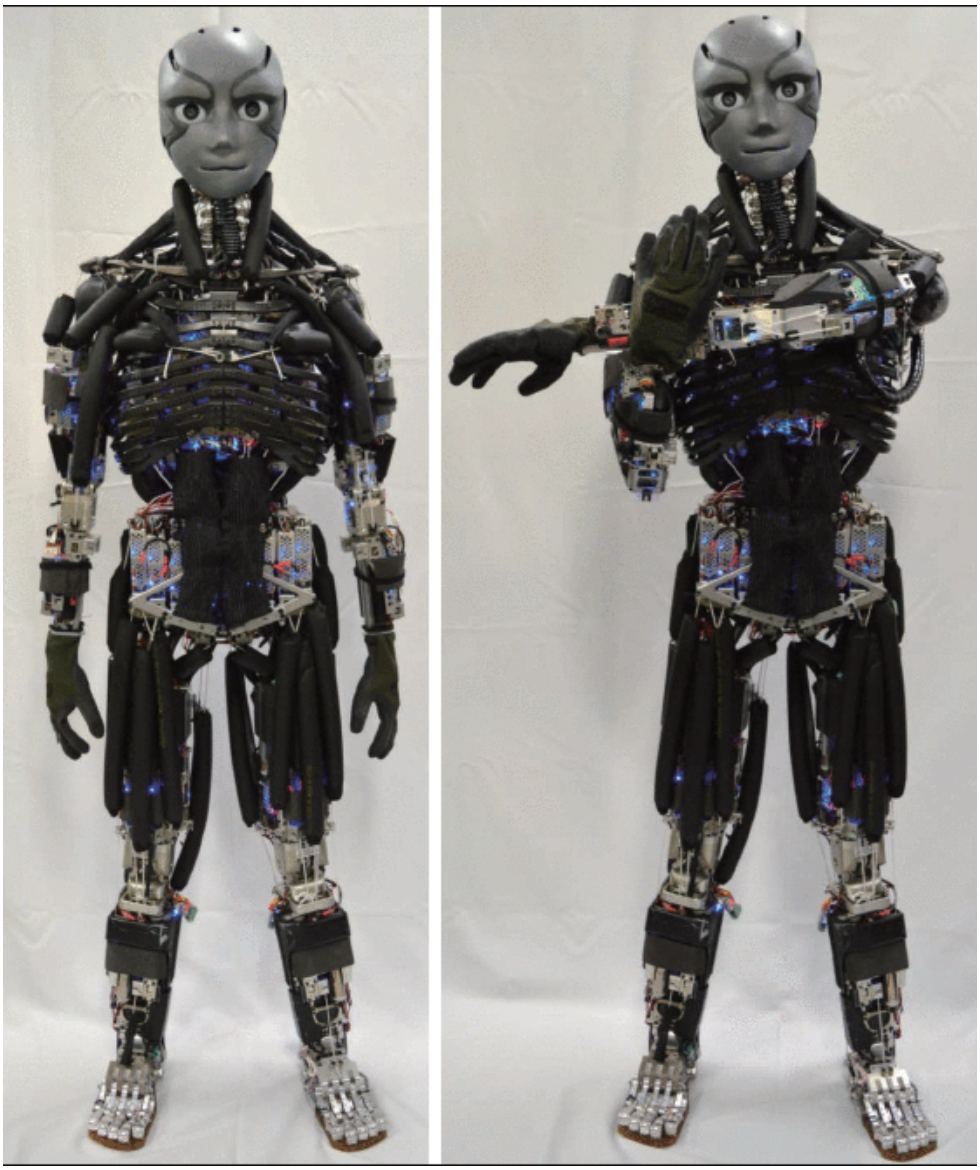


Figure 2.1: Kengoro human-mimetic musculoskeletal humanoid in neutral and upper-limb poses, illustrating its human-scale skeletal structure and densely routed tendon-driven musculature [**?**].

like actuators [**??**]. Hydraulic actuators offer exceptionally high power density and have enabled highly dynamic legged systems, but they remain relatively heavy, less efficient than biological muscle, and require a dedicated fluid infrastructure [**?**].

Among smart-material actuators, dielectric elastomer actuators (DEAs), a class of electronic electroactive polymers, contract by a different mechanism: an applied voltage squeezes a soft polymer film between compliant electrodes through electrostatic stress, and the thinning and radial spreading of the film produces strains that can exceed biological muscle [**?**]. Comparative studies report high specific work and good bandwidth for DEAs [**?**], but kilovolt-range driving voltages, electromechanical instability, and the compliant-electrode requirement have kept them out of legged robots. Thermally activated artificial muscles reach the most extreme performance figures of any technology surveyed by **?**: twisted-fiber muscles support loads more than one hundred times those of a biological muscle of the same length and weight, with output stresses of 22 MPa. Their activation route, however, is shared with shape-memory alloys, whose heating is fast but whose cooling is limited by heat transfer away from the element, and that asymmetry sets the cycle time.

Closer to muscle, series-elastic and pneumatic actuators trade peak force and bandwidth for compliance, a trade that wearable and rehabilitation robotics exploit because force control and backdrivability matter most when a machine interacts with a human [**???**]. A recent review of pneumatic artificial muscle technology documents continued progress in actuator design and in assistive applications [**?**]. None of these technologies yet combines muscle-like force, strain, and compliance in one actuator, which is the combination a biomimetic humanoid leg demands and that braided pneumatic actuators approximate

most closely.

Figure 2.2 compares the operating mechanisms of these actuator technologies.

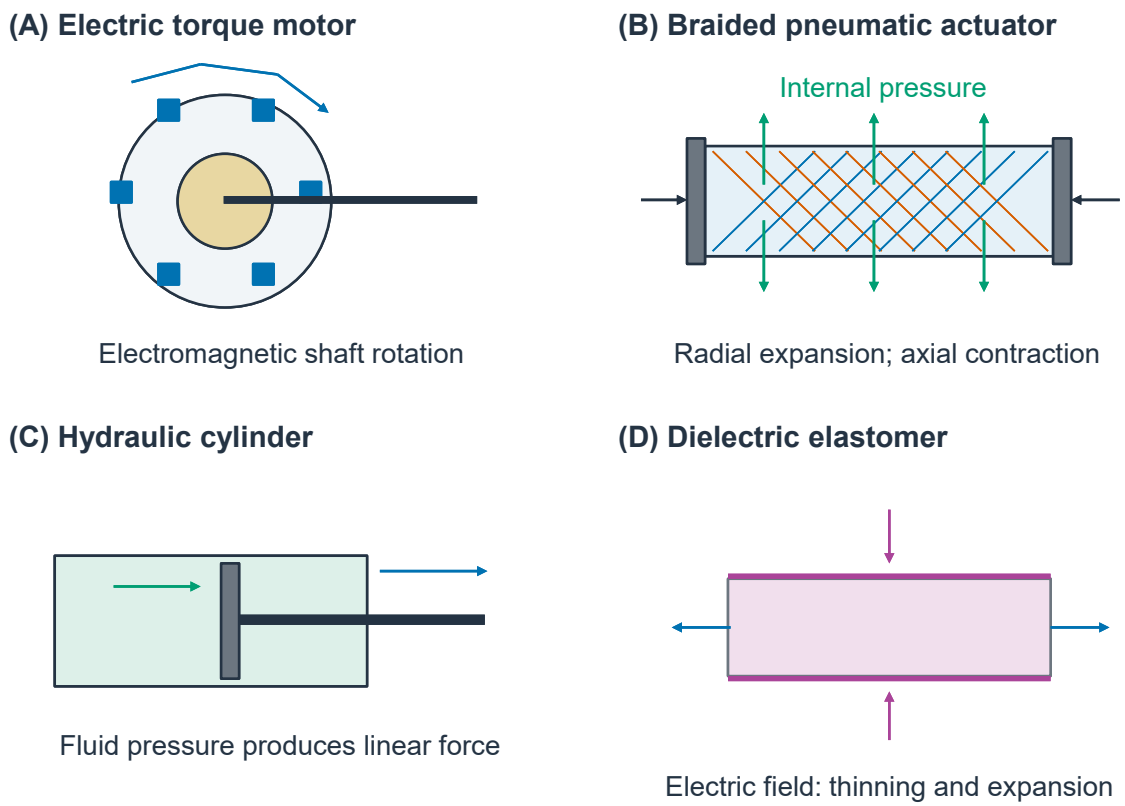


Figure 2.2: Representative operating mechanisms for robotic actuators. **(A)** An electric torque motor produces controlled shaft rotation. **(B)** A braided pneumatic actuator shortens as internal pressure expands its sleeve radially. **(C)** A hydraulic cylinder converts fluid pressure into linear piston force. **(D)** A dielectric elastomer thins and expands laterally under an applied electric field. The panels use a consistent schematic style to emphasize the distinct energy-conversion mechanisms and design tradeoffs summarized in the text [**??**].

2.2.2 Braided Pneumatic Actuators

In contrast to the technologies above, McKibben-style braided pneumatic actuators (BPAs) combine low mass, high force-to-weight ratio, and intrinsic compliance with a force–length curve that qualitatively resembles that of skeletal muscle [**??**]. A BPA consists of an inflatable bladder inside a braided sleeve; pressurizing the bladder expands the bladder radially and contracts the sleeve axially, producing tensile force. As shown by Chou and Hannaford, the force–length relationship of McKibben actuators is governed by geometric constraints imposed by the braided shell, leading to a nonlinear dependence of force on actuator length through the braid angle [**?**]: actuator force is proportional to pressure and to a nonlinear geometric term $(3\cos^2\theta - 1)$, which governs both the maximum contraction limit and the curvature of the force–length relationship. Predicting their force as a function of length and pressure is therefore the first improvement of the torque tools developed in this dissertation, and the state of that prediction is the subject of the next subsection.

Figure 2.3 shows the change in actuator shape with inflation.

However, BPAs exhibit several important differences from biological muscle. First, BPA maximum tensile force occurs at resting (i.e., uninflated) length (l_0) rather than at an optimal fiber length [**?**]. Second, their maximum contraction ratio is substantially lower than that of human muscle, presenting challenges for achieving biomimetic joint ranges of motion [**???**]. Third, BPAs have a highly nonlinear dependence on internal pressure (P), contraction (ϵ), and loading state [**??**], which complicates control and motivates high-fidelity force–length–pressure modeling. Despite these differences, their muscle-like properties have made them attractive for biomimetic quadrupeds [**??**], humanoids [**??**],



Figure 2.3: Festo BPAs in the test fixture: the inflated actuator is wider and shorter than the deflated actuator beside it.

and musculoskeletal arms [**??**].

2.2.3 BPA Force–Length–Pressure Models

Previous studies have generated empirical models of the force–length–pressure relationship of BPAs. **?** discuss several high-fidelity BPA force models. They present a static force model with a 21-coefficient polynomial function for a Festo MAS-20-200N (i.e., ∅20 mm, $l_0 = 200$ mm) with an impressive $R^2 = 0.9994$, as well as Sárosi's static force model for a Festo DMSP-20-400N-RM-RM (i.e., ∅20 mm, $l_0 = 400$ mm), with six coefficients and a $R^2 = 0.9995$ [**??**]. Martens and Boblan present an even more accurate model, in terms of absolute error, using five coefficients for a DMSP-10-250 (i.e., ∅10 mm, $l_0 = 250$ mm) and a DMSP-20-300 (i.e., ∅20 mm, $l_0 = 300$ mm) [**?**]. However, as demonstrated in this work, different resting lengths produce different force–length curves, preventing the use of these single-length models for designing systems with different resting lengths.

? looked at six resting lengths of ∅10 mm Festo BPAs, accounting for differences in maximum contractile percentages, and elucidated the force-length-pressure relationship. In particular, for a given contraction percent and force F (in Newtons), the scalar pressure P (in kPa) required to produce this force for a ∅10 mm Festo artificial muscle can be determined by solving the equation:

$$P = P_0 + k_F F + k_S S + P_t \tan\left(c_a \left[\frac{\epsilon}{\epsilon_{\max} - k_E F} - c_b\right]\right) \tag{2.1}$$

where F is the required force, ϵ is the contraction, S encodes hysteresis ($S = 1$ shortening, $S = -1$ lengthening, $S = 0$ static), and the seven coefficients are defined in Table 2.1. The

coefficients have been updated with corrected values, as those reported in **?** contained typographical errors. The contraction ϵ and the maximum free-load contraction at 620 kPa, which appears as $\epsilon_{\max}$ in Equation (2.1), are defined in Section 3.1.1, together with the relative contraction $\epsilon^* = \epsilon/\epsilon_{620}$ on which the force surface developed in this dissertation is built. Because Equation (2.1) expresses pressure in closed form as a function of force and contraction, it can also be solved onto a grid of pressures and contractions to build a force–pressure–relative-contraction lookup table for design studies.

Table 2.1: Coefficients of the 10 mm Festo BPA pressure model of Equation (2.1), corrected from **?**.

Symbol	Value	Role
P_0	254 kPa	pressure offset
k_F	1.23 kPa N^{-1}	linear force-to-pressure gain
k_S	15.6 kPa	hysteresis offset (sign of S)
P_t	192 kPa	tangent-term amplitude
c_a	2.03	contraction-shape coefficient
c_b	0.46	contraction-shift coefficient
k_E	0.331×10^{-3} N^{-1}	force-dependent reduction of usable contraction

Festo provides a manufacturer tool, MuscleSIM, for predicting force as a function of pressure and contraction [**?**], but experimental comparisons indicate that it does not adequately account for variation in l_0 or actuator-specific contraction limits. Comparisons with experimental data indicate that none of the existing models effectively take into account differences in maximum force due to initial actuator length, and the Hunt model was taken with low external forces ($\leqslant$106 N), so it is unclear how well it captures actuator behavior at higher forces. Chapter 3 compares these models with data collected for this dissertation and develops the improved characterization on which the torque and

controller work of the later chapters depends.

2.3 Human Lower-Limb Biomechanics

Designing a biomimetic leg requires a quantitative benchmark of human muscle–joint behavior. This dissertation uses OpenSim [**???**] and the Gait2392 model, named for its 23 degrees of freedom and 92 musculotendon actuators. Gait2392 descends from the interactive lower-extremity model of **?** and includes musculoskeletal geometry from **?** and planar knee kinematics from **?**; related lower-limb models of the same family are characterized and validated by **?**. Muscle moment arms, force-generating properties, and functional roles in such models have been characterized extensively [**??????**], and standard reporting conventions for kinematics and joint mechanics are established by the International Society of Biomechanics [**????**].

Moment arms relate muscle force to joint torque. For a fixed-axis joint, the signed moment arm is $r_{\mathrm{m}} = \hat{\mathbf{a}} \cdot (\mathbf{r} \times \hat{\mathbf{u}})$, where $\hat{\mathbf{a}}$ is the joint-axis direction, $\mathbf{r}$ locates a point on the muscle's line of action relative to the joint, and $\hat{\mathbf{u}}$ is the force direction [**?**]. The tendon-excursion method instead evaluates the sensitivity of musculotendon length to the joint coordinate, with sign determined by the coordinate and force conventions [**?**]. For multiarticular muscles, the result also depends on adjacent-joint configurations [**??**]. OpenSim uses length sensitivities, so the reference moment arms in this dissertation are properties of the Gait2392-derived geometry [**??**]. Chapter 4 compares the robot's resulting torque predictions with test-stand measurements.

Knee modeling requires particular care because the instantaneous center of rotation

changes with flexion. Four-bar descriptions reproduce this behavior through linkage geometry [**??**]. Gait2392 instead prescribes planar translations as functions of knee flexion [**?**]. These different kinematics produce different moment arms for the same nominal attachment locations. Joint-coordinate conventions must therefore be stated consistently when comparing models and measurements [**?**]. Gait2392 provides the human reference in this dissertation, while Steele's four-bar mechanism supplies the physical robot knee.

2.4 Biomimetic Knee and Bipedal Robot Model

Steele's biomimetic knee joint supplies the physical knee used in this dissertation [**???**]. It is a one-degree-of-freedom four-bar linkage whose crossed links reproduce the posterior migration of the human knee's instantaneous center of rotation (ICR) during flexion. The rolling-and-sliding behavior of the human joint therefore emerges from the link geometry itself rather than from cams or gear profiles, and the design was calculated to reproduce an average sliding-to-rolling ratio of 0.079 over a 151° range of motion [**?**]. Because the ICR migrates, muscle moment arms about the knee change with flexion angle in a human-like way, and a torque prediction tool for this knee must capture that variation rather than assume a constant arm. The biomimetic knee tests of Chapter 4 therefore exercise the corrected torque model against this geometry, and Figure 2.4 summarizes the mechanism.

In my master's work, muscle lengths, moment arms, and isometric joint torques were computed over the range of motion for both the Gait2392 model and a PAM-actuated robot leg model, demonstrating that several robot muscle torque profiles could closely match the human model while many fell far short for part of the RoM [**??**]. Hand placement of

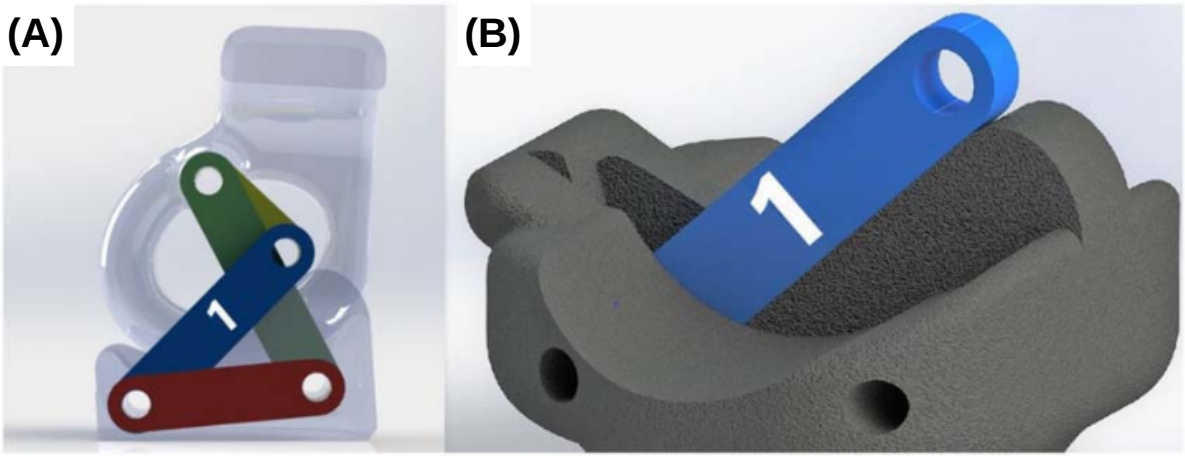


Figure 2.4: Biomimetic knee mechanism. **(A)** Sagittal cutaway at full extension showing the crossed four-bar linkage. **(B)** Exploded view of the mechanical stop. Reproduced from **?**, Fig. 6 under CC BY 4.0.

muscle origins and insertions was a painstaking and iterative task. This resulted in the creation of a placement optimization algorithm [**?**], discussed in Chapter 4. Figure 2.5 places the laboratory skeleton beside the robot PAM routes and human muscle paths used in that comparison.

The comparison evaluated muscle-path lengths, moment arms, and isometric torques across the human and robot models [**?**]. I adjusted BPA origins, insertions, and via-point locations to respect usable contraction and avoid interference. The biceps femoris long head illustrates the remaining mismatch: the robot path and its knee- and hip-torque surfaces differ substantially from the human reference. These discrepancies motivated placement optimization [**?**]. Figures 2.6 and 2.7 show the moment-arm construction and the torque comparison.

(A) **(B)** **(C)**

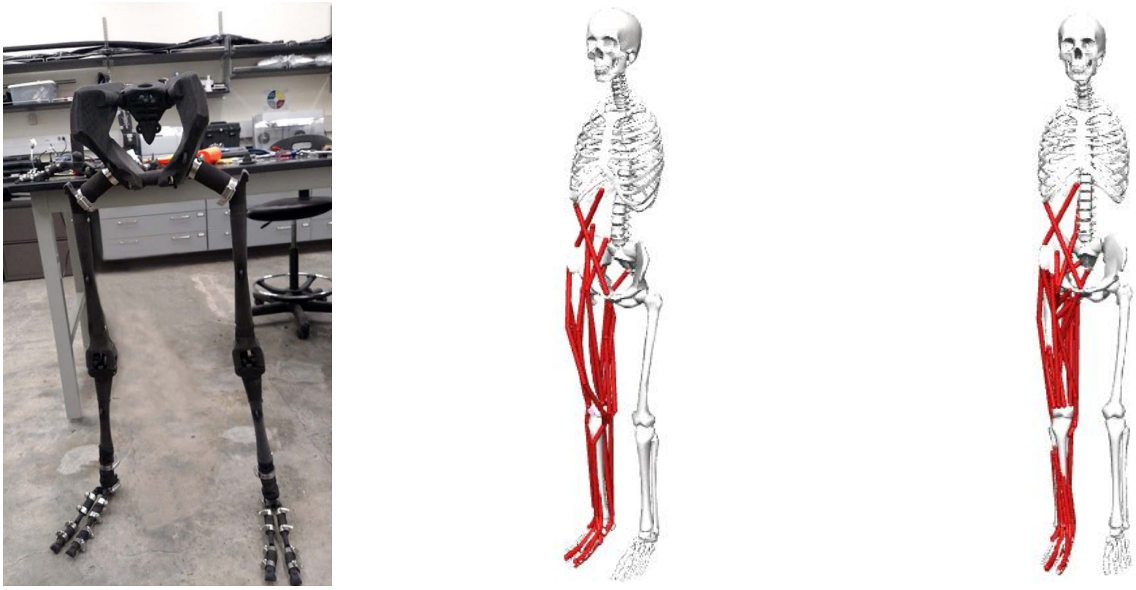


Figure 2.5: **(A)** Laboratory robot skeleton. **(B)** Robot PAM routes. **(C)** Human Gait2392 muscle paths used as design references [**??**].

2.5 Neural Control of Locomotion: Central Pattern Generators

Locomotion in mammals is driven by spinal central pattern generators (CPGs), neural circuits that can produce alternating rhythmic activity of flexor and extensor motoneurons in the absence of rhythmic input and proprioceptive feedback [**??**]. The founding hypothesis is Brown's half-center: alternating flexor and extensor activity arises from two reciprocally inhibitory central populations with opposing bias [**?**]. Modern instantiations of the half-center endow each neuron with persistent sodium channels, whose slow inactivation terminates each burst so that reciprocal inhibition switches the populations without any phasic input [**?**]; rhythmogenesis, in this tradition, is a property of the intrinsic membrane dynamics, not of a clock elsewhere in the network. The modern architectural

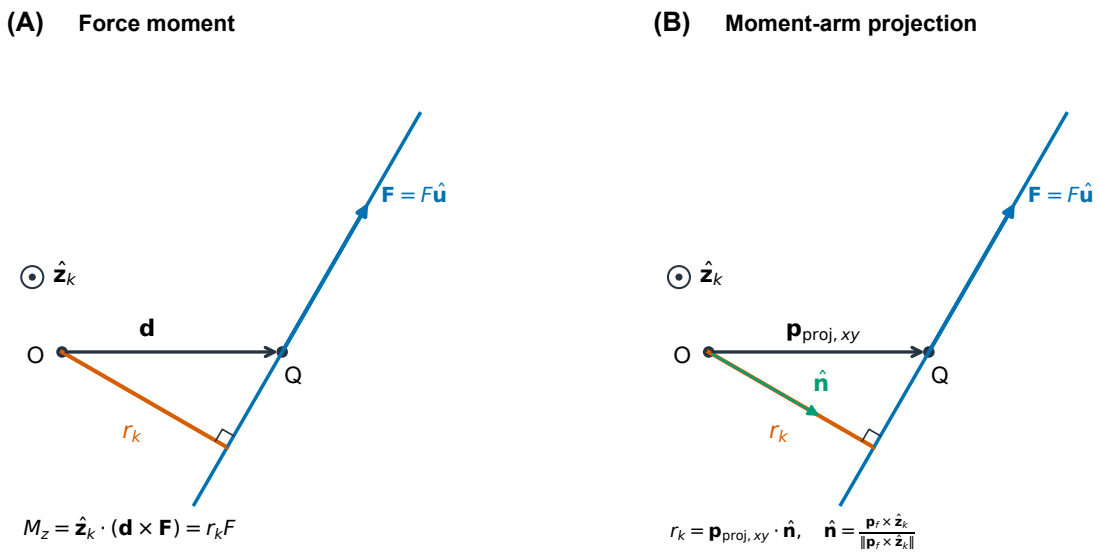


Figure 2.6: Equivalent fixed-axis constructions. **(A)** The cross product gives the moment of a force about the joint. **(B)** Projecting a joint-to-path vector onto the force-path normal gives the signed moment arm used in Equation (4.2). The joint axis points out of the page.

understanding of mammalian locomotor CPGs comes largely from the modeling work of the Rybak group and collaborators. **?** proposed that the locomotor CPG consists of two levels: a half-center rhythm generator (RG) and a pattern formation (PF) network, with reciprocal inhibitory interactions between antagonist neural populations at each level. This two-level organization explains, among other phenomena, spontaneous deletions of motoneuron activity during fictive locomotion in which the locomotor period and phase are maintained, and it permits independent control of the level of motoneuron activity and of step-cycle timing. **?** reviewed the electrophysiological support for this organization, and **?** related the model architecture to genetically identified spinal interneuron classes. Left–right coordination is provided by commissural pathways between the two half-centers [**??**], and the rhythmogenic core of the CPG continues to be refined experimentally [**?**].

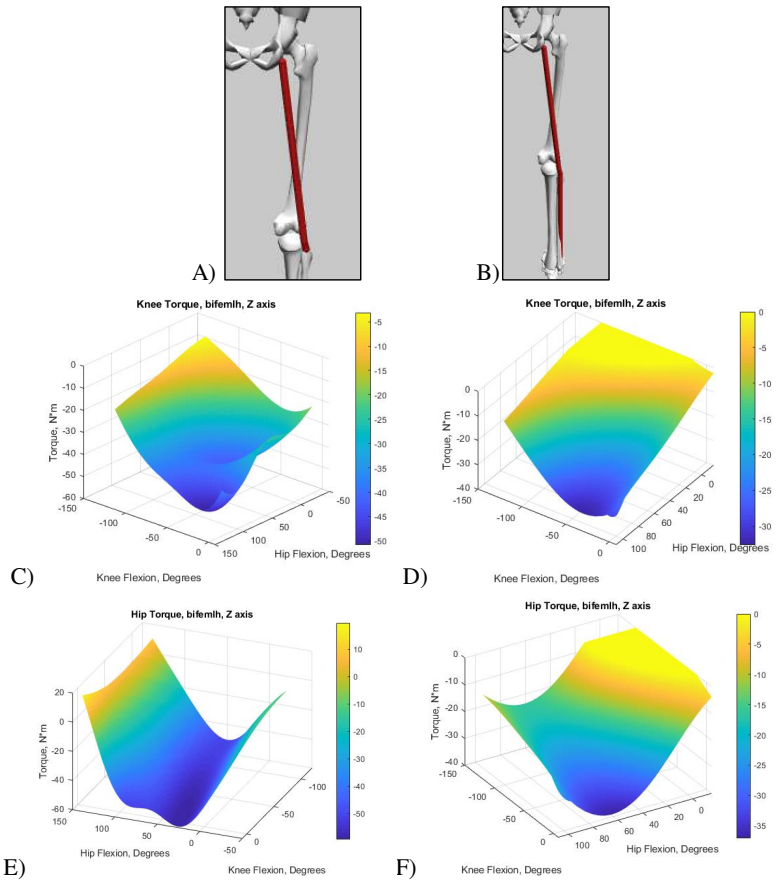


Figure 2.7: Biceps femoris long head in the master's study: muscle paths in the OpenSim model (A, C, E) and the robot model (B, D, F), with isometric torque surfaces about the knee (C, D) and the hip axis (E, F) over the hip–knee range [**?**, Fig. 8].

Two-level models separate the timing of the locomotor rhythm from the distribution of activity to muscles. This distinction has been extended to genetically identified interneuron classes [**?**] and circuits controlling muscles that cross two joints [**?**]. The latter problem also arises in a robot: activating a two-joint muscle can affect motion at both joints, so its control cannot be assigned to either joint in isolation.

Coupling the circuit to a musculoskeletal model adds the effects of mechanics and sensory feedback. Such models show how afferent feedback shapes CPG output [**?**] and how body mechanics and feedback influence the behavior of different CPG architectures [**?**]. Human locomotion models combine locomotor primitives with Ia, II, and Ib pathways and Renshaw cells to examine their contributions to control [**?**]. Other models examine interlimb responses to sudden loss of ground support [**?**]. Section 2.7 describes these circuits, and Chapter 6 develops tools for implementing and testing related controllers.

2.6 Sensory Feedback in Locomotion

Sensory feedback is not a passive correction layer but an integral part of locomotor pattern generation. Reviews of afferent control of the CPG [**?**] and of muscle proprioceptive feedback and spinal networks [**??**] summarize the classical picture: group Ia afferents from muscle spindles signal muscle length and velocity, group Ib afferents from Golgi tendon organs signal force, and group II afferents signal length, with load-related feedback playing a role distinct from length-related feedback during locomotion [**?**]. Removing proprioceptive feedback degrades locomotor patterns in mice [**?**]; proprioceptive input regulates and adapts locomotor activity [**?**]; and the relative contributions of proprioceptive

and vestibular systems to locomotion are the subject of ongoing study [**??**]. Tonic central and sensory stimuli have been shown to facilitate involuntary air-stepping in humans [**?**], demonstrating that simple graded stimulation of motoneuronal and rhythmic circuits can elicit stepping-like output, an experiment readily reproduced in simulation by tonic stimulation of pattern-formation and rhythm-generator populations. The proprioceptive and contact feedback channels designed in Chapter 6 follow this afferent organization.

2.7 Computational Neuroscience

Computational neuroscience expresses hypotheses about neural circuitry as executable models. In conductance-based models, membrane voltages evolve through capacitance, leak, synaptic, and applied currents, and connections represent identifiable pathways [**?**]. Model predictions can be compared with recordings or with responses to stimulation and pathway removal. This provides a way to test a proposed mechanism rather than infer it from a wiring diagram alone. Synthetic nervous systems use this approach to control physical and simulated bodies.

2.8 Synthetic Nervous Systems

For robots, these biological architectures can be implemented as synthetic nervous systems (SNSs): networks of conductance-based spiking and non-spiking neurons that are small enough to run in real time yet biologically grounded enough to be compared against animal data. The approach has been pursued most persistently by the insect-biorobotics lineage of the Quinn laboratory at Case Western Reserve University and by laboratories closely

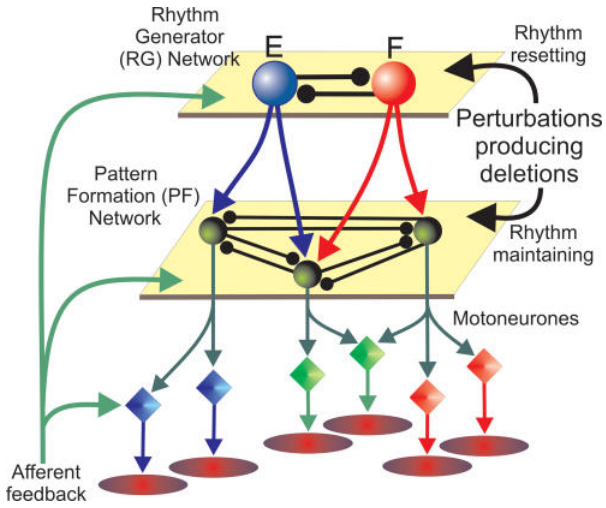


Figure 2.8: Two-level CPG: rhythm generation (RG), pattern formation (PF), motor output, and afferent feedback. Reproduced from **?**, Fig. 1.

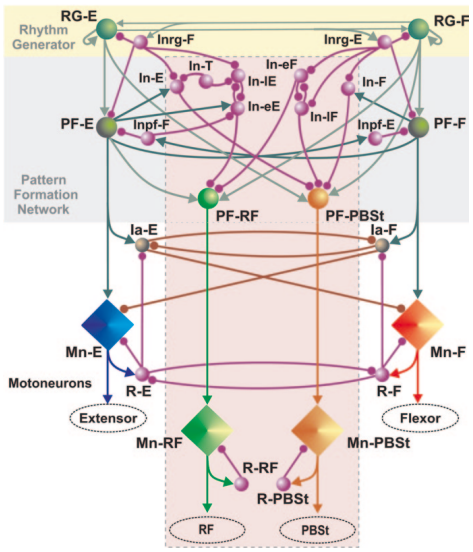


Figure 2.9: Extended CPG with pattern-formation and motor populations for the two-joint posterior biceps–semitendinosus (PBSt) and rectus femoris (RF). Reproduced from **?**, Fig. 5.8.

connected to it, my own included. Works in this lineage span several platforms: CPG-based neural control of a pneumatically actuated dog robot [**?**], analytically designed dynamic neural networks for human balance control [**?**], the full-muscle rat hindlimb model and its moment-arm analysis [**?**], the two-layer CPG neuromechanical walker built on it [**?**], which supplies this dissertation's controller architecture, and the Muscle Mutt quadruped, a canine-inspired BPA robot designed at my laboratory, the Agile and Adaptive Robotics Laboratory (AARL) at Portland State University, for synthetic nervous system control [**?**]. Related controllers from the wider neuromechanical community include sensory-entrained CPG control of bipedal walking [**?**], inter-leg coordination analyses in stick-insect-inspired controllers [**?**], reinforcement-reflex hybrids in insects [**?**], and adaptive hindlimb walking controllers [**?**].

The SNS-Toolbox [**?**] makes such networks practical: it is an open-source Python package that simulates hundreds to thousands of spiking and non-spiking neurons in real time or faster, supports multiple software and hardware backends including GPUs and embedded platforms, and couples directly to physics simulators such as MuJoCo for musculoskeletal control. The field it serves is active and growing: the perspective review of **?** traces neuromorphic legged-locomotion control from its insect-robot beginnings through current platforms to open problems, and physical testbeds in this lineage now extend from insect-scale robots to a robotic cat hindlimb [**?**].

2.9 Neuromechanical Simulation

Neuromechanical simulation, the coupling of a musculoskeletal model to a model of the neural controller that drives it, has advanced along three partially overlapping lines. The first line embeds reflex physiology directly into the controller. The muscle-reflex walking model of **?** demonstrated that positive force, length, and velocity feedback around each muscle, organized with simple interlimb coordination, is sufficient to produce self-stabilizing human-like walking dynamics and realistic muscle activations without any central pattern generator. That model became the reference point for a generation of reflex-based neuromechanical controllers, including the cat hindlimb models of Section 2.7 [**??**] and their extensions. Hybrid work makes the complementarity of the two strategies explicit: **?** combined feed-forward muscle activations with neuronal reflex feedback in simulated human walking and found that the robustness benefit of each depends on which muscle carries it.

The second line replaces hand-designed circuitry with machine learning. **?** trained deep reinforcement learning policies on a physiologically realistic musculoskeletal model in OpenSim, reproducing human kinematics and muscle activations across locomotion tasks, with the learned policy re-usable for experiments that would be difficult in the laboratory. This approach inherits the fidelity of the musculoskeletal model but encodes control in opaque function approximations: the trained policy can reproduce behavior, but the individual "neurons" of the controller do not correspond to identifiable physiological circuits, so the class of scientific questions it can answer, such as lesions, deletions, and pathway stimulation, is narrower than its behavioral repertoire. Recent perspectives

on predictive musculoskeletal simulation [**?**] place both learned and optimal-control approaches in the same frame: the models are mature, and the open problems are the cost of solving them and the biological interpretability of the control they find.

The third line uses biologically grounded synthetic nervous systems with explicit neuron and synapse models [**?**]. Their circuit structure allows a researcher to change a specified pathway and examine its effect on behavior. Related work includes motor-module bipedal walkers [**?**], two-layer RG/PF control of a musculoskeletal macaque model [**?**], and two-level CPG control of muscle-driven robotic hind legs [**?**]. This dissertation develops the interfaces needed to study such controllers with the project's human reference and artificial-muscle models. The actuator and knee-torque experiments provide mechanical evidence, while the simulation framework provides a way to examine controller behavior under specified support and perturbation conditions.

2.9.1 Simulation Tools

All three lines depend on a software substrate that is itself part of the method: elaborate musculoskeletal geometry and complex neural models. The tools used in this dissertation are the following.

OpenSim [**??**], distributed through the SimTK (Simulation Toolkit) repository, provides the biomechanical modeling side: musculoskeletal models with detailed bone geometry, muscle routing, and force-generating properties, along with tools for analyzing moment arms and isometric torque. These models are approximations with known limits rather than exact anatomies: muscle paths are interpolations over discrete attachment and wrap-

ping data [**?**]. Their geometry is accurate for a purpose, namely moment arms and torque predictions over the functional range of motion, and that is the standard against which the robot model of this program is compared.

The same SimTK infrastructure hosts SCONE, an open-source package for predictive simulation and optimization of neuromuscular control of biological motion [**?**].

AnimatLab [**?**] provides a graphical environment in which physical bodies and synthetic neural systems are constructed and co-simulated; the AARL's prior walker (Section 6.1) was built in it.

MuJoCo [**?**] provides the physics-engine side: fast, accurate simulation of rigid-body dynamics with contacts, along with muscle-like actuators that include activation dynamics.

MyoConverter bridges the biomechanical and physics-engine worlds by converting OpenSim musculoskeletal models to MuJoCo format with optimized muscle kinematics and kinetics [**??**].

SNS-Toolbox supplies the neuron and synapse models used to simulate the controller [**?**].

I used these tools to develop a robotic OpenSim leg model, convert it to MuJoCo, and couple it to a synthetic nervous system. Chapter 6 describes the assembled interfaces and their evaluation. The coupled simulation supports repeated parameter, stimulation, and pathway-removal tests under controlled conditions. Whether those tests support a biological hypothesis depends on the model's validation and the experimental comparison, as well as simulation speed.

2.10 Testing Neuromechanical Hypotheses on Robots: The Ethical Case

The experiments this research program is designed to enable, including deleting a muscle from a limb, blocking a sensory pathway, lesioning part of a rhythm generator, and applying tonic stimulation to spinal circuits, cannot be performed on healthy humans. Invasive intervention on the human neuromuscular system is limited to clinical populations, is tightly constrained by ethics review, and permits neither the within-subject control nor the reversibility that neuromechanical hypotheses demand. The same experiments in animals also carry real moral cost, and their justification is governed by the Three Rs framework: replacement of animal subjects with non-animal methods, reduction of the numbers required, and refinement of the remaining procedures [**??**]. Robotic and computational models allow proposed neural interventions to be tested under controlled conditions before corresponding animal experiments are considered. The same model and initial conditions can be reused across interventions to isolate the effect of a changed pathway. Repeating a model does not create independent biological subjects, and its relevance to animal experiments depends on validation against biological observations.

Figure 2.10 documents the animal preparation behind many foundational spinal-locomotion experiments [**??**]. In this preparation, the forebrain is disconnected by a brainstem transection while spinal and brainstem circuits continue to produce motor activity. The image is included to make the experimental context visible alongside the circuit schematics. Robotic and computational models can support preliminary intervention tests, but their relevance to these biological preparations depends on reproducing the observations being studied.

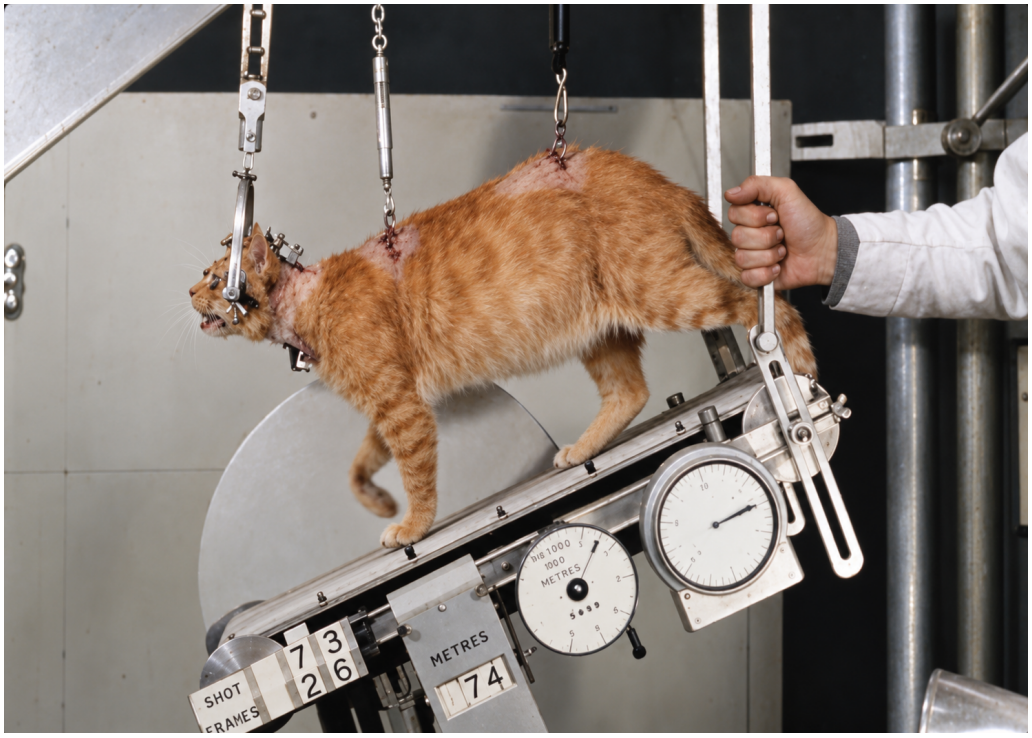


Figure 2.10: Restored and colorized frame from archival footage of a decerebrate-cat treadmill experiment [**?**].

The procedures behind that record are documented in films that remain publicly available. A 50 min classroom recording posted by the UTHealth Houston School of Dentistry first shows an intact cat playing with its handler, righting itself when lowered and placing its paws accurately on the edge of a table, and then documents the decerebration performed for the class [**?**]: the animal is anesthetized with open-drop ether, more than the horse in the same demonstration series received, the carotid arteries are ligated and a tracheal cannula installed, and the skull is opened with a trephine before the brainstem is transected between the superior and inferior colliculi. The treadmill preparations that established locomotor neurophysiology are likewise preserved in the historical literature: **?** reproduce the arrangement of the high-decerebrate cat treadmill experiments, with electromyographic recording from the hindlimb muscles, from the original Russian-language reports [**?**].

The honest accounting is murkier than the Three Rs framing alone suggests, because the models a robot implements are themselves products of animal experimentation. The half-center organization, the two-level rhythm and pattern architecture, and the length, force, and load signaling of the afferents were each established in preparations like the ones described above [**???**], and the databases of afferent properties that parameterize sensory feedback are compiled from animal recordings [**?**]. A synthetic nervous system therefore replaces future experiments; it does not replace, and cannot repay, the animal work on which its architecture is built. Recognizing this inheritance sharpens rather than undermines the case for robotic subjects: the knowledge purchased at that cost is precisely what makes a robotic replacement scientifically meaningful, and the robot is the instrument that keeps that knowledge productive without adding to the ledger.

For a robot to serve as a replacement subject, however, it must first earn standing as a scientific model. Webb's influential analyses of bio-robotic methodology [**??**] argue that a robot is a good model of biological behavior when the mechanism it implements is tested at the same level of organization as the target phenomenon, matching dynamics, control structure, and sensory ecology rather than surface appearance. The validation chain of this dissertation is constructed to that standard: the actuator force model is validated against bench data (Section 3.2.1), the joint torque model against isometric measurements and human torque envelopes (Section 4.2), and selected controller responses are compared with published tonic stimulation and deletion phenomena [**??**], the verification step that precedes locomotion tuning in this program. Only then can a robot experiment be read as evidence about the hypothesized biological circuit rather than about the robot.

The argument is ultimately bidirectional, completing the loop described in Section 1.2. Experiments that are unethical, impractical, or statistically unreachable in humans and animals become routine on the robot; the resulting insights into neuromechanical control feed back into biology, into the design of the next robot iteration, and directly into prosthetic and orthotic devices that face the same problem of controlling muscle-like actuators for human users, whose central pattern generator controllers are an active topic in lower-limb exoskeleton research [**??**]. The ethical case and the scientific case are the same case: the robot is the instrument that makes these experiments possible.

3 Isometric Force Characterization of Braided Pneumatic Actuators

Publication note: Sections 3.1.1–3.1.2, Section 3.2.1, and Sections 3.3.1–3.3.5, with their associated figures and tables, appeared in **?** (Bolen, Elzein, Pang, and Hunt, *Actuators*, 2026). For that work, B.B. performed the conceptualization, methodology, software, validation, formal analysis, investigation, data curation, visualization, and original draft preparation; M.E. and L.P. contributed methodology, software, investigation, and data curation; and A.H. supervised the work, administered the project, and acquired funding. All co-authors contributed to writing, review, and editing.

This chapter develops and validates the actuator force models on which the rest of the dissertation builds, and it reads as one self-contained study, from methods through results to discussion. Sections 3.1.1 and 3.1.2 present the methods: the force equations used throughout this work, and the test jig and measurement procedure used to characterize Festo BPAs with 10 mm and 20 mm uninflated diameters over a wide range of resting lengths, contractions, and pressures. Section 3.2.1 reports the resulting maximum-force and normalized force–strain–pressure models and compares them with the manufacturer's tool. The remaining sections discuss the effects of resting length, the nondimensionalized force prediction, the comparison with other published BPA models, the implications for real-time control and simulation, and the analogy between these artificial muscles and biological muscle. The validated force model is the one used for the knee-torque predictions of Chapter 4 and for the BPA actuation of the neuromechanical simulations in

Chapter 6.

3.1 Methods

3.1.1 Actuator Force Model

The actuator force model used throughout this dissertation is described here. First, the amount of contraction ε in a BPA is calculated as

$$\varepsilon = \frac{(l_0 - l_{\mathrm{m}})}{l_0} \tag{3.1}$$

where l_0 is uninflated resting length of the BPA measured from hose clamp to hose clamp. Current muscle length is l_{m}.

Similarly, ε_{620} is the amount of contraction in a BPA without external load when inflated at 620 kPa, calculated as

$$\varepsilon_{620} = \frac{(l_0 - l_{620})}{l_0} \tag{3.2}$$

where l_{620} is the muscle length measured at 620 kPa and free from external loads.

The amount of relative contraction of a BPA ε^* can then be calculated by dividing equation (3.1) by (3.2), such that

$$\varepsilon^* = \frac{\varepsilon}{\varepsilon_{620}} \tag{3.3}$$

describes the BPA's current contraction as a normalized percentage of its total.

The amount of normalized isometric force F^* in a Festo BPA is then calculated as

$$F^*(\varepsilon^*, P^*) = \begin{cases} c_0 \cdot \left(\exp(-c_1 \cdot \varepsilon^*) - 1\right) + P^* \cdot \exp\left(-c_2(\varepsilon^*)^2\right) & \text{for } F^* > 0 \\ 0 & \text{for } F^* \leqslant 0 \end{cases} \tag{3.4}$$

where P^* is pressure P normalized by 620 kPa. BPA diameter dependent coefficients are determined by the actuator characterization experiments reported in Section 3.2.1 (Table 3.2).

The maximum force F_{max} in a BPA is calculated as

$$F_{max}(l_0, P_{\max}) = a_1 \cdot P \cdot \arctan\left(a_2 \cdot P \cdot (l_0 - 0.0075)\right) \tag{3.5}$$

with coefficients determined by the actuator characterization experiments reported in Section 3.2.1 (Table 3.1). When maximum system pressure is set to $P = 620\,\text{kPa}$, $F_{\max}(l_0, 620) = F_{620}(l_0)$. Multiplying (3.4) by (3.5) yields

$$F(\varepsilon^*, P^*, l_0) = F^*(\varepsilon^*, P^*) \cdot F_{620}(l_0) \tag{3.6}$$

as the equation for isometric force F in a Festo BPA.

3.1.2 BPA Isometric Force Characterization Experiment

To measure artificial muscle force as a function of length and pressure, we built a test jig from extruded aluminum and 3D printed parts. Festo BPAs with 10 mm and 20 mm diameter were tested with various resting lengths and different amounts of contraction.

These data are used to develop an improved model that more accurately predicts maximum muscle force based on l_0, current muscle length l_m, and pressure P. The results are compared with existing models and the new data are used to create an updated model.

A test jig was made of 80/20 brand 1515 series extruded aluminum (Figure 3.1). Artificial muscles were placed vertically in the jig one at a time. The upper end was attached to a load cell. Force data was collected using a CALT DYLY-103 100 kg S-shaped load cell for the ∅10 mm BPAs and a Loadstar RAS1-2HKS-S for the ∅20 mm BPAs. The load cell was used in conjunction with a HX711 Load Cell Amplifier. Pressure data was obtained from a Freescale MPX5700 GP 5 V pressure sensor. Building air supply pressure was controlled with two pressure regulators in series. The first was a Parker model 20R113GC $0 - 120$ psi pressure regulator. The second was a Husky 3/8 in. High Performance Air Regulator HDA72200. A Festo VTUG-10-MRCR-S1T-26V20-T516LA-UL-532S-8K valve manifold VTUG-G was used to deliver air from the pressure regulator to the actuator. This manifold is comprised of eight two-in-one bidirectional normally closed Festo VUVG-S10-T32C-AZT-M5-1T1L valves.

Pressure and load cell amplifier data were sent to MATLAB 2021b via an Arduino Uno style Sparkfun BlackBoard C microcontroller. The computers that used MATLAB were running either Windows 10 or Windows 11. During phases when the Arduino was collecting force and pressure data to send to MATLAB, the Arduino would also trigger (via MATLAB) an Onsemi 2N4401 NPN transistor to make the valve manifold open or close the valve. For other data collection, the valve was manually opened and closed.

The inner distance between the hose clamps on each BPA was measured with a FANUC tape measure to determine the muscle's length l and resting length (l_0). This is how **?**

defines l_0, although in **?** it was measured to also include end cap length. We then inflated each BPA to $P_{620} = 620\,\text{kPa}$, with one end allowed to move freely in the axial Degree of Freedom (DoF), and measured the length l_{620} to calculate maximum contraction at 620 kPa using Equation (3.2). The distance between the crossmembers was then controlled to obtain different amounts of contraction, the muscles were inflated to various pressures, and the contractile force was recorded at the pressure–contraction pairs. This was done for ∅10 mm BPA resting lengths (l_0) of 112 mm–518 mm. For ∅20 mm BPAs, resting lengths (l_0) of 300 mm–509 mm were used.

3.2 Results

3.2.1 BPA Force Characterization

This section reports the results of the isometric force characterization experiments described in Section 3.1.2. First, maximum isometric force at 620 kPa is modeled as a function of resting length. The remaining data are then normalized to build a predictive force surface, and the resulting model is compared with the manufacturer's tool and with other published BPA models.

Maximum Force at 620 kPa

Data of the maximum force from BPA characterization tests of the 10 mm and 20 mm diameters at 620 kPa show a dependency on resting length (Figure 3.2). This is a previously unreported characteristic of these artificial muscles. Detailed analysis of the Festo tool [**?**]

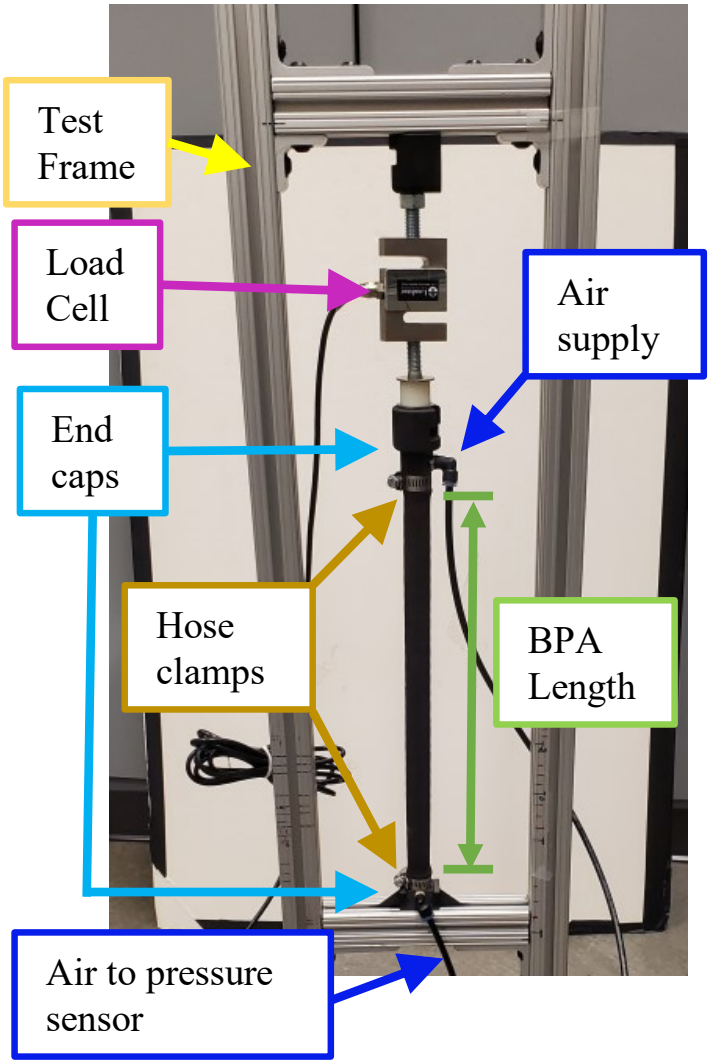


Figure 3.1: A picture of a BPA in the isometric force test stand with components labeled.

does in fact predict a change in maximum force with the resting length; however, the Festo tool predicts increased force with shorter lengths, while our collected data indicate decreasing force with shorter lengths. The data show a force response resembling an arctan curve along the l_0 dimension. Using the Nonlinear Least Squares method and a Least Absolute Residual robustness, we fit an arctan curve to the data to get the maximum force at 620 kPa given a resting length as:

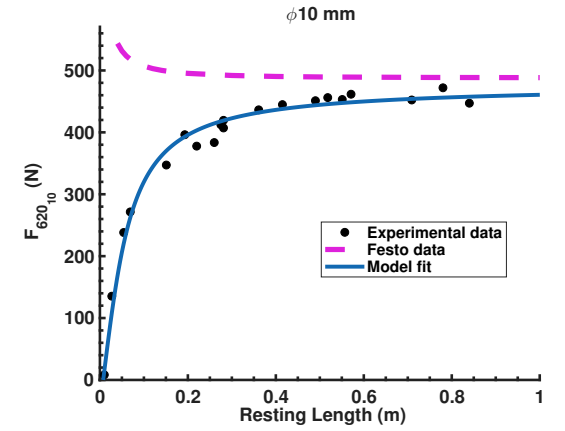
$$F_{620_{10}} = 303.5\,\mathrm{N} \cdot \arctan\left(19.03\,\mathrm{m}^{-1} \cdot (l_0 - 0.0075)\right) \tag{3.7}$$

$$F_{620_{20}} = 922.4\,\mathrm{N} \cdot \arctan\left(15.37\,\mathrm{m}^{-1} \cdot (l_0 - 0.013)\right) \tag{3.8}$$

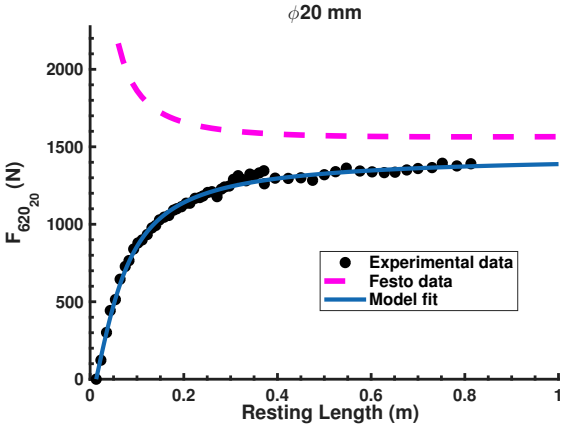
The length is offset by 0.0075 m and 0.013 m because solid modeling showed that the end caps contact each other at these lengths. At these lengths, the actuator would not be able to contract to produce force.

Force as a Function of Pressure and Resting Length

Data from BPA characterization tests of the 10 mm muscles resulted in force–pressure pairings for different muscle resting lengths (l_0) (Figure 3.3a). Similar to the maximum force data, these data show a force response resembling an arctan curve along the l_0 dimension with a more linear response to changes in pressure. Using the Nonlinear Least Squares method and a Least Absolute Residual robustness, we fit an arctan curve to the data to get the maximum force given a resting length and pressure as described by the model form given in Section 3.1.1, where $a_1 = 0.4895\,\mathrm{N\,kPa}^{-1}$ and $a_2 = 0.030\,68\,\mathrm{kPa}^{-1}\,\mathrm{m}^{-1}$



(**a**)



(**b**)

Figure 3.2: Results for finding the relationship between F_{620} and l_0. (**a**) $F_{620_{10}}$ as a function of l_0, at P_{620}. Dashed line is the $F_{620_{10}}$ data from Festo. Solid line is the fit from Equation (3.7). (**b**) $F_{620_{20}}$ as a function of l_0 at P_{620}. Dashed line is $F_{620_{20}}$ data from Festo. Solid line is the fit from Equation (3.8).

for the 10 mm actuator, and $a_1 = 1.49\,\mathrm{N\,kPa^{-1}}$, $a_2 = 0.0248\,\mathrm{kPa^{-1}\,m^{-1}}$ for the 20 mm actuator. Goodness-of-fit measures are given in Table 3.1; throughout this dissertation, goodness of fit is reported as the root-mean-square error (RMSE), the adjusted fraction of variance unexplained (FVU, for which $R^2 = 1 - \mathrm{FVU}$), and the maximum absolute residual, as defined alongside the reference implementation in Appendix A.

Maximum Contraction at 620 kPa

Figure 3.3(b) shows an attempt at a linear fit for maximum contraction at 620 kPa (ϵ_{620}) as a function of resting length (l_0). Error bars have been included to show the effect of ± 1 mm accuracy with measurements of l_0 and l_{620}. There was a large amount of variance in the data, with the linear fit giving an adjusted $\mathrm{R}^2 = 0.4124$ and an RMSE $= 0.0083$. Since there is no direct, predicable relationship between maximum contraction and l_0, this value should be recorded in each muscle used on a robot to best predict the force it may produce at different pressures and contraction.

Force as a Function of Pressure and Contraction

In addition to being a function of the pressure and resting length, the force produced by the actuator is also a function of the amount of actuator contraction, with less force being applied as the actuator contracts. To build this model, the collected data of force, pressure, and contraction are normalized by dividing by F_{620}, P_{620}, and ϵ_{620}, respectively. This had the effect of compressing the data into a 3D surface ranging from 0 to 1 on all axes. Normalized data are compared with pressure isoclines of force predicted by the

Table 3.1: Maximum force equation coefficient values with confidence intervals (CI) and goodness-of-fit measures. Equation (3.5) is compared to data taken at 620 kPa. Equations (3.7) and (3.8) are compared against maximum force from using the Festo tool. It was not necessary to fit an adjusted R^2 in these cases.

Equation	Coefficient	CI (95%)	Our Model			Festo Tool		
			Adj. R^2	RMSE	Max. Error	Adj. R^2	RMSE	Max. Error
(3.5)	$a_1 = 0.4895\,\mathrm{N\,kPa^{-1}}$ $a_2 = 0.030\,68\,\mathrm{kPa^{-1}\,m^{-1}}$	(0.4822, 0.4968) (0.0282, 0.03317)	0.9945	11.61 N	55.3 N	0.9854	14.7 N	30.9 N
(3.7)	$b_1 = 303.5\,\mathrm{N}$ $b_2 = 19.03\,\mathrm{m^{-1}}$	(300, 308) (17.48, 20.57)	0.9854	14.72 N	30.9 N	–	189.9 N	375.6 N
(3.8)	$b_1 = 922.4\,\mathrm{N}$ $b_2 = 15.37\,\mathrm{m^{-1}}$	(914.2, 930.7) (14.75, 15.98)	0.9945	23.83 N	62.1 N	–	668.4 N	1590.1 N

Table 3.2: Normalized isometric BPA force equation (Equation (3.4)) coefficient values and goodness-of-fit measures.

BPA	Coefficient	CI (95%)	Model			Validation		
			Adj. R^2	RMSE	Max. Error	Adj. R^2	RMSE	Max. Error
∅10 mm	$c_0 = 0.5682$	(0.5584, 0.578)	0.9998	0.005118	10.3%	0.9994	0.0245409	10.6%
	$c_1 = 4.254$	(4.126, 4.383)						
	$c_2 = 0.5597$	(0.5429, 0.5766)						
∅20 mm	$c_0 = 0.2579$	(0.2401, 0.2756)	0.992	0.02294	6.8%	0.9943	0.0231303	5.7%
	$c_1 = 6.477$	(5.558, 7.396)						
	$c_2 = 1.321$	(1.239, 1.403)						

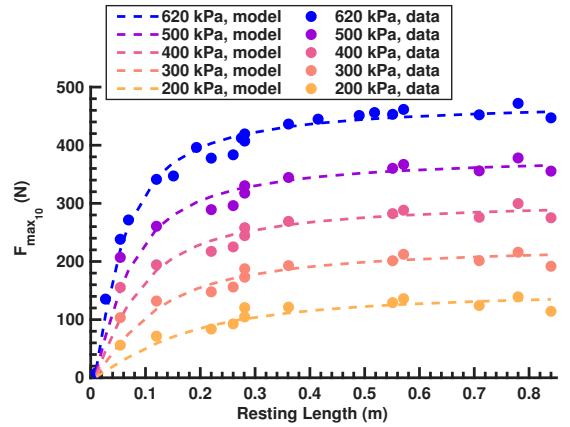
Festo tool divided by the maximum force equation described in the previous section in Figure 3.4. With this comparison, it is clear that the Festo tool over-predicts the expected force, especially at lower pressures, contraction, and resting lengths.

We therefore derived our own equation for isometric force in the BPA as a function of pressure and contraction. Visual analysis of the experimental data shows an exponential relationship between ϵ^* and F^*, and a linear relationship between P^* and F. We fit a surface to the data using Nonlinear Least Squares and a Least Absolute Residual robustness such that it is described by Equation (3.4). For the ∅10 mm BPAs, the result of the improved fit can be seen in Figure 3.4a. The adjusted $R^2 = 0.9998$, a RMSE $= 0.004537$, and a maximum absolute residual of 10.6%. Solving (3.4) for the ∅20 mm BPA, yields different coefficients, and results are seen in Figure 3.4b. Coefficient values and goodness-of-fit statistics are found in Table 3.2.

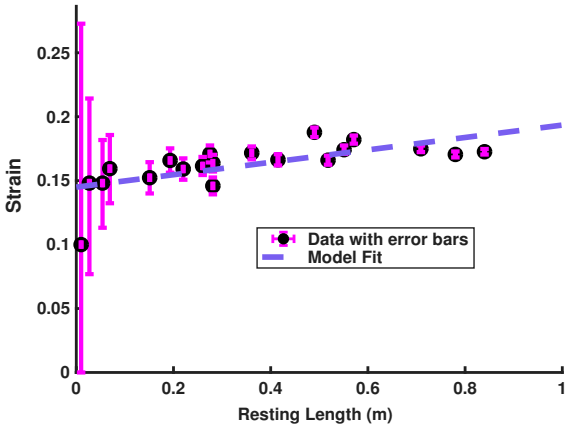
3.3 Discussion

3.3.1 Effects of Resting Length

We have developed new F_{620} equations for 10 mm and 20 mm diameter Festo BPAs that capture the change in maximum force produced by the actuators at 620 kPa as a function of their resting length. When examining Equation (3.7) and measured data in Figure 3.2b, it can be seen that as l_0 goes to infinity, $F_{620_{10}}$ goes to 470 N. When the Festo tool [**?**] was queried at $l_0 = 1$ m as an approximation of infinity, it predicts the ∅10 mm BPA can produce $F = 490$ N at P_{620}. This is within the 10% variability that Festo states may occur in manufacturing tolerances. Similarly, the Festo data predict a ∅20 mm BPA will

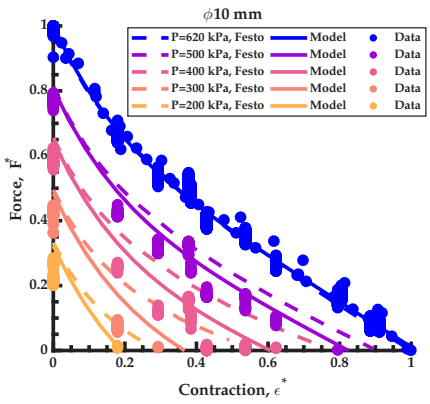


(**a**)

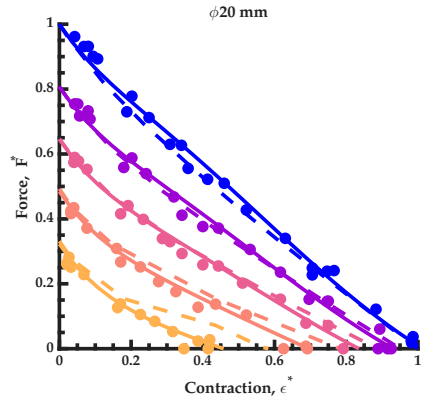


(**b**)

Figure 3.3: Results for finding the relationship between F_{620}, l_0, P_{620}, and $\epsilon_{620_{10}}$. (**a**) Isoclines of the surface fit for $F_{620_{10}}(l_0, P_{620})$. (**b**) $\epsilon_{620_{10}}$ versus l_0 at P_{620}. Although there is a general trend of longer resting lengths producing more contraction, no conclusive relationship between ϵ_{620} and l_0 could be deduced from this experiment. The error bars show the effect of being ± 1 mm in measuring l_0 and l_{620}.



(**a**)



(**b**)

Figure 3.4: Surface fit for $F^*(\epsilon^*, P^*)$. Solid lines are from our model. Dashed lines from Festo supplied data. In total, 537 data points were collected and used to fit the 3D surfaces (321 data points for ∅10 mm BPAs and 216 data points for ∅20 mm BPAs). In total, 80% of these data were used for training, and 20% were used for validation. For figure clarity, not all data is included and plotted circles represent collected data at ± 10 kPa of the stated pressure. (**a**) Fit data for ∅10 mm Festo BPAs. (**b**) Fit data for ∅20 mm Festo BPAs.

produce $F = 1570\,\text{N}$. Our results for Equation (3.8) show that $F_{620_{20}}$ goes to 1460 N as l_0 approaches infinity. Although the difference is much larger, it is also within the 10% manufacturing tolerance.

Our measured data differ significantly from the Festo model as l_0 goes to zero. For the measured data, as the length gets smaller, the force does as well, whereas the Festo tool predicts exponential force increase (see Figure 3.2). As a specific example, when specifying $l_0 = 112\,\text{mm}$ for the 10 mm diameter BPAs, the Festo predicted F_{620} is 498.6 N. However, using Equation (3.7), F_{620} is calculated to be 335.3 N. Actual force in the $l_0 = 112\,\text{mm}$ was measured at 350.9 N. Therefore, the error in predicted F_{620} for the ∅10 mm BPA is 4.5% for our model and 42.1% using the Festo tool. Researchers using BPA resting lengths under 300 mm should take note of these results.

? (p. 12) states that force will be reduced by approximately 10% for a minimum nominal length. In **?** (p. 14) they say the force can deviate by 10% based on factors such as production variation, material variation, and deviation from nominal length (l_0). **?** (p. 13) notes that testing was done with BPAs that are ten times the uninflated diameter, and that theoretical force can increase by up to 10%. In each datasheet, the characteristic curves appear to be the same. As demonstrated above, the 100 mm ∅10 mm BPA should have much less force than what is shown in the curve. This is our experience and the reason for the testing. It is unclear why the MuscleSIM tool [**?**] then predicts higher force as resting length decreases. It could be a reversed sign coefficient in their internal curve-fit model. We have shown here that as the nominal length goes to zero, so does force.

BPA models commonly predict higher forces than are measured experimentally because they do not account for end effects in regions of the pressure vessel where those effects

dominate. Nominally, the clamped ends fix the inner and outer diameter while allowing axial and radial displacement. In the isometric case, there is only radial displacement. The boundary conditions force steep gradients in stress and strain near the ends. Near the boundary conditions, axial, radial, and circumferential strains deviate from analytic solutions for a cylindrical model. For a constant pressure, the end effect region appears to have the same length. Therefore, in shorter BPAs, the end effects dominate. The axial length scale of the boundary section is comparable to l_0, and the active force differs substantially from that of an assumed perfect cylinder.

Manufacturing variations seem to play an important role in the behavior of these actuators. It remains to be seen how ϵ_{620} can be known a priori. We still suspect that it is a function of l_0, but it might also be a function of other factors such as product batch number or UV light exposure [**?**] (p. 15); [**?**] (p. 6); [**?**] (p. 13). We found there to be larger variations in the 10 mm BPAs than in the 20 mm BPAs. Modeling work on ideal McKibben actuators describes how the relationship between contraction and initial length is a function of the total number of twists and the angle of twist in the usable muscle area [**?**]. It is unclear if there are variations in the angle and number of twists in the Festo BPAs due to manufacturing. Uncovering the relationship between ϵ_{620} and l_0 for the Festo DMSP artificial muscles, if it exists in a meaningful way, will require additional controlled tests.

Similarly, it is not clear how the maximum force might change with manufacturing variations. This was reasonably corrected for by normalizing the equations to each individual's maximum contraction, but more nuanced details might emerge with more data collection. Not knowing these relationships a priori means there is still a need for

robots and other systems to go through an iterative design stage if it is discovered that the system does not produce the torque or have the RoM that the design team expects. However, the data and models provided here should enable faster redesigns with fewer additional measurements.

3.3.2 Nondimensionalized Force Prediction

We have also introduced the concept of nondimensionalized isometric force that is a function of relative strain and relative pressure (Equation (3.4)). This elegant equation has only three coefficients with low error (see Table 3.2). In the work presented here, we introduce the relative pressure term, $P^* = P/P_{620}$. Our data show that by normalizing force, contraction, and pressure, we are able to create a simplified force equation as a function of contraction and pressure that scales well with initial actuator length, as expressed by Equation (3.6) in Section 3.1.1.

For example, given $l_0 = 54\,\text{mm}$ and $P = 300\,\text{kPa}$, the measured force was $F_{actual} = 103\,\text{N}$. Force prediction using only the Festo tool predicts the force to be $F_{predict} = 270\,\text{N}$. This is 161% greater than the measured force. Using Equations (3.7) and (3.4) with Equation (3.6) gives $F_{predict} = F^*(0, 0.48) \cdot F_{620}(0.054) = 0.49 \cdot 230\,\text{N} = 112\,\text{N}$. This is an error of 8%, which is much more accurate than using the Festo tool only.

The presented work is capable of predicting force output for 10 mm and 20 mm muscles with different initial muscle lengths, internal pressure, and contraction amount with minimal initial data collection. However, it remains to be seen if these relationships are maintained over the full lifecycle of the muscles. There are not yet any peer reviewed

studies investigating the long-term behavior of these artificial muscles, though the Festo datasheet provides some indication of what can be expected [**?**] (p.19). They report that the 40 mm BPA can complete 1 million cycles at 6000 N (maximum rated load), or 10 million cycles at 4000 N. The service life of the other size BPAs is not reported; however, it is mentioned that thermal stressing may reduce the lifespan and they should be inspected every 500,000 cycles for signs of aging [**?**] (pp. 10–11). Future research should look at the behavior of these muscles over time.

In our work, $P_{620} = 620\,\text{kPa}$ is the maximum supply pressure for our system, and other users of this actuator may use a different supply pressure. However, this supply pressure is not required, and normalizing by this value was chosen as a method for improving the optimization techniques by reducing the sensitivity of the term associated with pressure. This equation will work for any pressure supplies below 620 kPa. We also anticipate the equation will work for pressures up to 800 kPa (the maximum rated pressure for Festo ∅10 mm BPAs); however, we do not have the data to confirm this.

3.3.3 Comparison with Other Published BPA Models

The proposed model was compared with the formulations by **?**, **?**, and **?**, to evaluate their predictive performance across actuator sizes and operating conditions (see Table 3.3 and Figure 3.5). For the ∅20 mm BPA, parameters from [**?**] were used, corresponding to actuators with resting lengths of 200 mm and 400 mm. These were compared against the experimental data at $l_0 = 300\,\text{mm}$ and $l_0 = 450\,\text{mm}$, respectively. For the ∅10 mm BPA with resting lengths of 257 mm and 233 mm, parameters from [**?**] were applied.

In each case, the maximum contraction parameter (k_{max}) was determined numerically by solving $F(6.2\,\text{bar}, k_{max}) = 0$, ensuring consistency with the experimental maximum pressure condition.

While the **?** and **?** models capture general trends in force–contraction behavior, discrepancies were observed when applying parameters derived from different actuator lengths. In particular, the model tends to over-predict force at higher pressures and low amounts of contraction, suggesting sensitivity to geometric scaling effects not explicitly accounted for in the formulation.

The **?** model is derived from a geometric and material-based framework and provides a more physically grounded representation of actuator behavior. However, its accuracy depends strongly on precise geometric parameters, including braid angle, wall thickness, and fiber length. Small deviations in these parameters can lead to significant differences in predicted force, particularly at higher pressures.

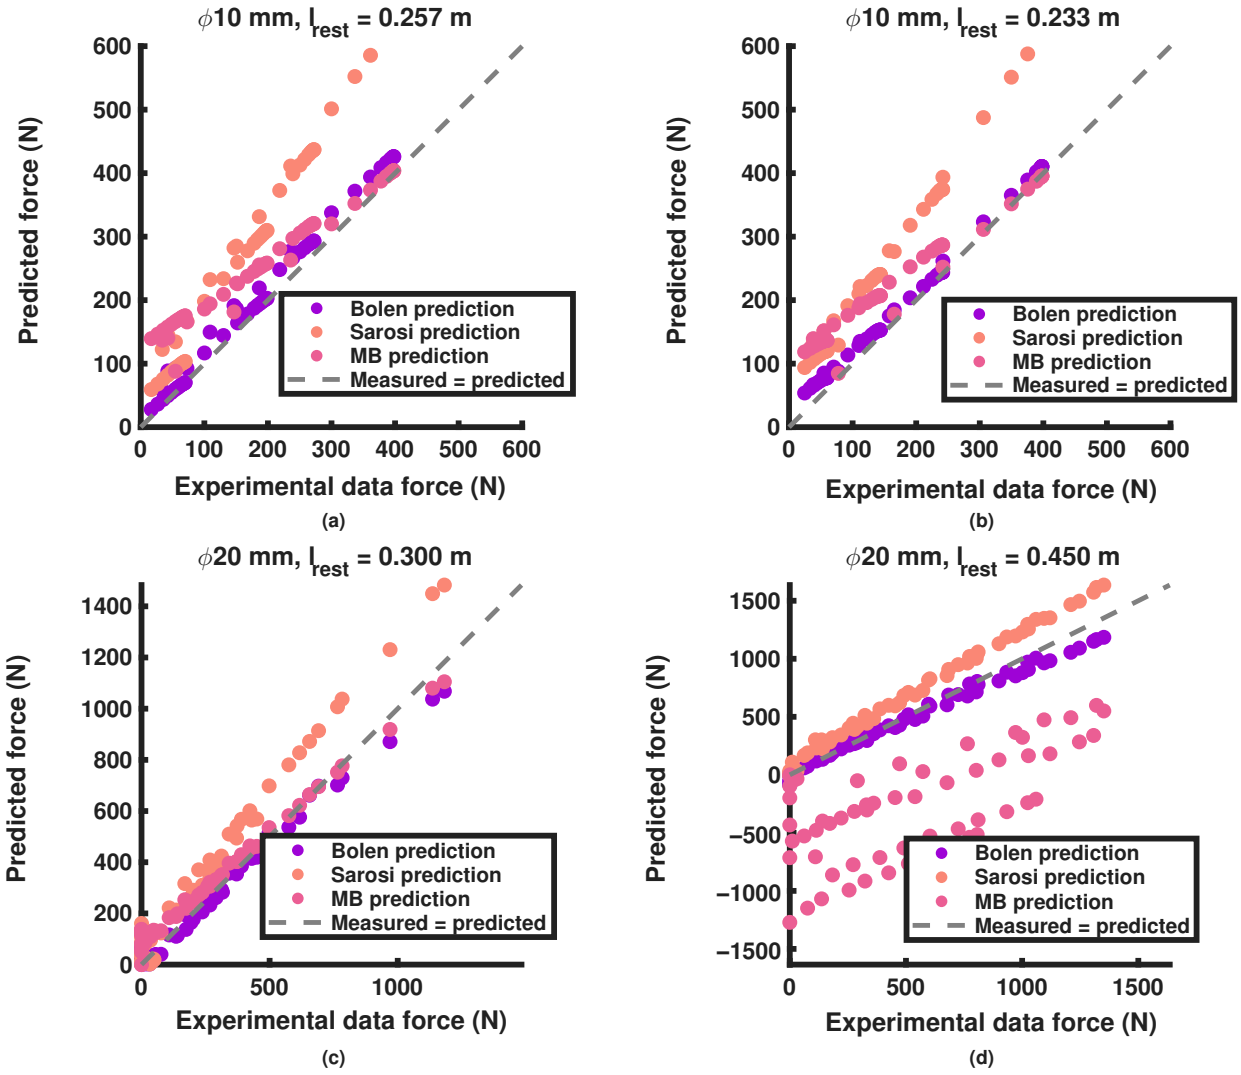


Figure 3.5: Comparison between measured and predicted BPA force at varying amounts of contraction and pressure. A $y = x$ perfect fit line is added for reference. The BPAs compared are: (**a**) ∅10 mm, $l_0 = 257$ mm, (**b**) ∅10 mm, $l_0 = 233$ mm, (**c**) ∅20 mm, $l_0 = 300$ mm, and (**d**) ∅20 mm, $l_0 = 450$ mm.

The Martens and Boblan model was first evaluated at the same pressure–contraction points as in their experimental data [**?**] (Tables A1 and A2), indicating that it reproduces

the general published force map shape but is sensitive to the assumed initial geometry. Examining their ∅20 mm BPA, for example, using the explicitly stated $l_0 = 300\,\text{mm}$ and initial diameter $D_0 = 20\,\text{mm}$ significantly under-predicted the results compared to using the implied $l_0 = 296\,\text{mm}$ ([**?**] (Table A1)) and $D_0 + H_0 = 21.8\,\text{mm}$, where $H_0 = 1.8\,\text{mm}$ is the material thickness. The diameter adjustment can be inferred from [**?**]. This adjusted diameter was also included for the ∅10 mm BPA to give a better fit. As can be seen in Figure 3.5 and Table 3.3, the highest accuracy is observed when ∅20 mm and $l_0 = 300\,\text{mm}$, which is at or very near their experimental measurements. For the ∅10 mm BPAs the accuracy decreases even though their experimentally measured $l_0 = 250\,\text{mm}$ is close to our $l_0 = 257\,\text{mm}$ and $l_0 = 233\,\text{mm}$. The largest error occurs at ∅20 mm and $l_0 = 450\,\text{mm}$ and the model starts to predict negative force. We have calculated that the $F_{PE} \cdot \frac{dD}{DL}$ term becomes very large compared to the others [**?**] (p. 5). This suggests that the model is sensitive to geometric parameter selection and interpretation, particularly with respect to initial length and diameter definitions.

In comparison, the present normalized model demonstrates improved consistency across actuator lengths by incorporating experimentally derived scaling through F_{620} and ϵ_{620}. This normalization reduces sensitivity to geometric variation and enables a unified representation of actuator behavior, while maintaining predictive accuracy over the tested range of pressures and contractions.

3.3.4 Implications for Real-Time Control and Simulation

Although pneumatic muscle actuators exhibit hysteresis, thermodynamic effects, and other dynamic behavior, accurate and high-performance controller designs can still be achieved using a static force map approach when the underlying static force model is accurate. Martens and Boblan explicitly note that BPA dynamics are often dominated by the air dynamics inside the actuator and the mass flow behavior of the valve, such that precise modeling of the static force characteristic remains crucial for torque control accuracy [**?**]. This view is consistent with more recent BPA control work, including a sensor-less torque control interface for BPA-driven joints [**?**], stable force/position/stiffness control of antagonistic pneumatic-muscle joints [**??**], and sensor-less angle stiffness control of antagonistic BPA systems [**?**].

In this context, the present model is well suited for real-time use because it reduces the static isometric force prediction to a low-order normalized surface, $F^*(\epsilon^*, P^*)$, scaled by the resting-length-dependent term $F_{620}(l_0)$. This structure is computationally inexpensive to evaluate online and is also suitable for lookup-table implementation. While the present work does not model the full actuator–valve dynamics, it provides the static force relationship needed as a feedforward or inner-model component in real-time torque control, simulation, and preliminary design workflows. It is anticipated that improved real-time control for high-speed and high-precision systems will benefit more from improved modeling of airflow dynamics than from modeling actuator hysteresis, damping, and other dynamic behavior. Characterizing the actuators under dynamic operating conditions, such as isokinetic (constant-velocity), isobaric, and isotonic (constant-load) contractions,

is a natural next step beyond the isometric characterization reported here, for which the pressure–force dynamics model of **?** provides the modeling foundation.

3.3.5 Biomimetic Artificial Muscle Analogy

BPA artificial muscles are often said to be analogous to biological muscles because they have force–length curves and can produce force only in tension. This analogy is worth closer inspection. For example, the passive force–length curve has been described as an exponential function arising from connective tissue elements, while the active force–length relationship is typically modeled as a bell-shaped curve reflecting actin–myosin overlap [**?**]. These formulations emphasize that muscle force generation is governed by nonlinear geometric and material interactions rather than simple linear elasticity. McKibben-type pneumatic artificial muscles (of which Festo is an example) exhibit analogous nonlinear behavior, though arising from fundamentally different physical mechanisms. As shown by Chou and Hannaford, the force–length relationship of McKibben actuators is derived from geometric constraints imposed by the braided shell, leading to a nonlinear dependence of force on actuator length through the braid angle [**?**]. In their formulation, actuator force is proportional to pressure and a nonlinear geometric term $(3\cos^2\theta - 1)$, which governs both the maximum contraction limit and the curvature of the force–length relationship. The work presented here demonstrates that this nonlinear curve can be well approximated by an exponential as well.

However, the behavior of these artificial muscles differs distinctly from that of biological muscles. For example, the optimal fiber length of a biological muscle, l_{OFL}, is not equivalent

to l_0 in an artificial muscle. Artificial muscles produce their maximum active force at their resting length, which is also their maximum length. In contrast, the maximum active force that a biological muscle can produce is somewhere in the middle of the total contractile length of the muscle.

Furthermore, biological muscles can adapt the optimal operating length through structural remodeling (e.g., sarcomere addition or tendon compliance) while BPAs have a fixed geometric configuration. This allows biological systems to self-optimize for the particular tasks the animal is subjected to. In order to re-optimize a robot for new tasks, new actuators would have to be swapped into the system. This distinction highlights a key limitation of BPAs in biomimetic applications, while also reinforcing the importance of accurate nonlinear modeling for control and design.

The analogy also suggests two design observations. First, the force–length penalty is shared by both families: force falls as the actuator operates away from its optimum, so a limb design must route each actuator such that its highest-force operating range coincides with the joint angles where torque is needed most, because the force lost at the extremes of the range of motion cannot be recovered by pressure alone. Second, the two families gain moment arm in different ways. Skeletal muscle attaches to the skeleton itself, through tendons or directly over broad areas of bone, so the moment arms of large animals follow from skeletal geometry. The robot's BPAs instead act through brackets that protrude from the limb segments; that hardware is mechanically exposed, and its series compliance is one of the loss mechanisms quantified in Chapter 4.

The analogy finally raises a question that this dissertation can pose but not settle. Every muscle-like actuator loses force as it operates away from its optimal length, and a soft

actuator working against soft structure has no rigid skeleton from which to gain a moment arm. Given those two facts, it may be no accident that nearly all large land animals are built on rigid skeletons and that large soft-bodied land animals are essentially absent. Whether the force–length penalty contributes to that absence, or whether it reflects other constraints such as support and protection, is a question for comparative biomechanics rather than for a single actuator study. For the designer, the lesson holds regardless of the answer: some moment arm must come from somewhere, and on this robot it comes from protruding brackets whose behavior must be modeled rather than ignored.

Table 3.3: BPA force-model comparison. Diameter and resting length are in mm; RMSE and maximum absolute error are in N. FVU is dimensionless.

Diameter	Length	Model	RMSE	FVU	Max. error
10 mm	257	Bolen	22.9	0.0353	46.0
10 mm	257	S'arosi	143.8	1.3903	235.8
10 mm	257	Martens and Boblan	73.8	0.3668	122.9
10 mm	233	Bolen	16.8	0.0187	30.3
10 mm	233	S'arosi	132.7	1.1655	222.6
10 mm	233	Martens and Boblan	62.6	0.2590	101.2
20 mm	300	Bolen	35.6	0.0144	114.0
20 mm	300	S'arosi	132.1	0.1992	313.4
20 mm	300	Martens and Boblan	72.8	0.0605	137.8
20 mm	450	Bolen	67.6	0.0258	166.6
20 mm	450	S'arosi	177.7	0.1785	294.8
20 mm	450	Martens and Boblan	862.1	4.2022	1321.1

4 Isometric Knee Torque

Publication note: This chapter is the subject of an unpublished manuscript in preparation by Bolen, Burgess, Morrow, Elzein, Pang, Scharzenberger, and Hunt [**?**]. For that work, B.B. performed the conceptualization, methodology, software, validation, formal analysis, investigation, data curation, visualization, and original draft preparation; M.E., L.P., L.B., and C.M. contributed methodology, software, investigation, and data curation; and A.H. supervised the work, administered the project, and acquired funding. All co-authors contributed to writing, review, and editing.

This chapter connects the actuator force characterization of Chapter 3 to torque about the robot knee. The rigid-body torque model of the earlier robot work assumed inextensible muscle paths on rigid links, and the isometric tests reported here show where that assumption fails and by how much. The chapter proceeds from the simple pinned knee to the biomimetic joint, identifies the correction terms that remove the rigid-body idealization from the torque prediction, validates them on held-out tests, and closes with the attachment-route redesign that brings the actuator configurations to human-scaled torque targets. The methods, results, and discussion follow the manuscript noted above; the physical-validation paragraph of Section 4.3.3 additionally reports the newer nine-point sweep results, which postdate the manuscript draft.

4.1 Methods

4.1.1 Robot Leg Model and Knee Test Stands

The primary components of the robot leg are the knee joint, femur, tibia, BPA assemblies, and artificial tendons (Fig. 4.1**A–C**). The artificial bone components are 3D printed using a combination of Onyx, carbon fiber, and high-strength high-temperature fiberglass on Markforged Onyx One and Mark Two printers, as well as machined links for the knee mechanism. Onyx is a proprietary Markforged material that consists of chopped carbon fiber in nylon. Artificial tendons are made with Shimano bicycle brake cable.

The leg assembly geometry and dimensions are based on bones scans of a 6 ft (1.83 m) tall person [**????**]. Two knee joints were tested, one pure rotational joint, and one biomimetic joint with rotation and translation. Each joint is driven by two antagonistic Festo actuators. Both robot joints have the same center of rotation in the femur frame at $\theta_k = 0\,\text{deg}$.

Details on the biomimetic robot knee used in this study can be found in [**???**]. In brief, this 1-DoF joint (Fig. 4.2) uses a four-bar linkage to change the effective moment arm during rotation. The intersection point of the internal crossed links define the instantaneous center of rotation (ICR). This causes the joint ICR location in the femur frame to translate posteriorly in the X direction and distally in the Y direction except for a brief proximal move near full flexion. Reference coordinate frames are given by **??**. This is similar to how the human knee behaves [**?????**]. The correct location of the ICR is important for both theoretical and experimental torque calculations.

Uniarticular knee extensor and knee flexor are used to actuate the joint. Each knee joint type was tested with a single ∅10 mm Festo BPA per muscle. The biomimetic knee was

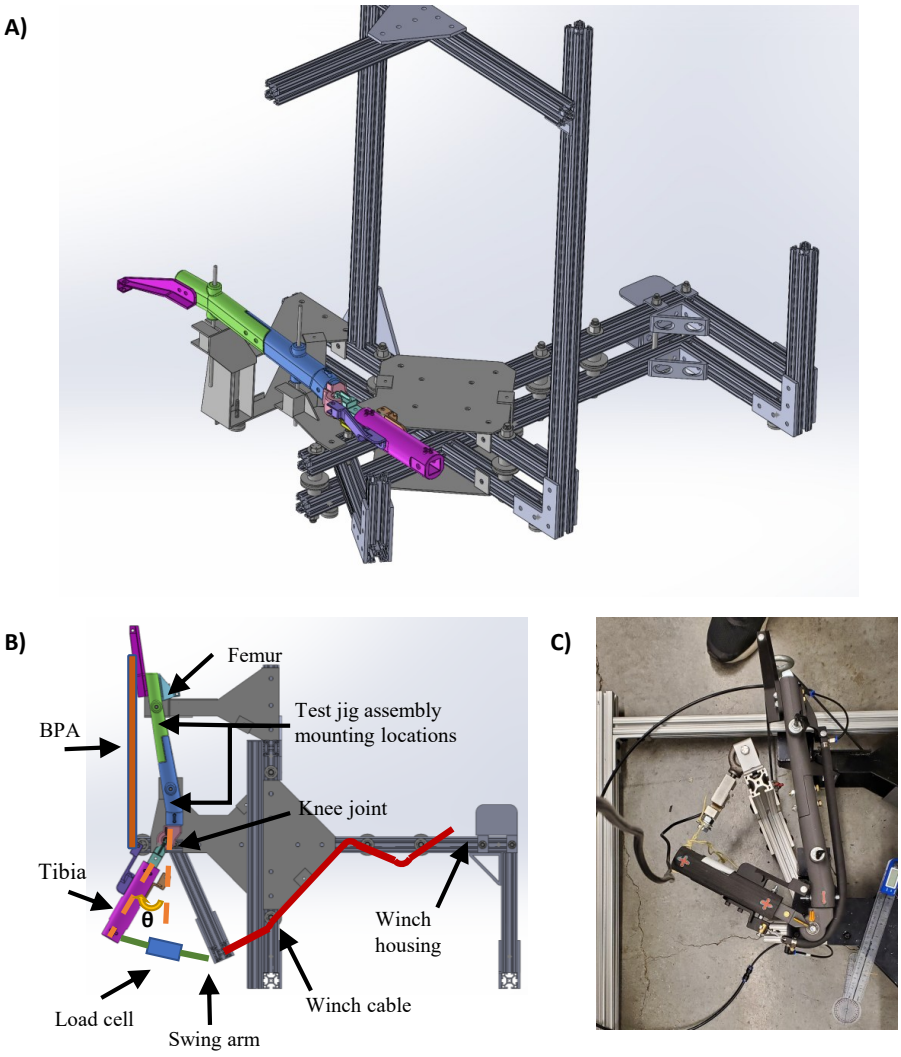


Figure 4.1: **(A)** CAD solid model assembly of the revolute joint knee in the test stand. **(B)** 1-DoF pinned-joint robot knee in the test apparatus with important components labeled. Leg is shown with 30 degrees flexion, i.e. $\theta_k = -30^\circ$. **(C)** Picture from isometric torque testing. Setup shows the revolute knee configured for a test of the extensor BPA. The knee is flexed with $\theta_k = -120^\circ$. Also note the compressed shape of the inflated BPA on the knee during this high degree of flexion.

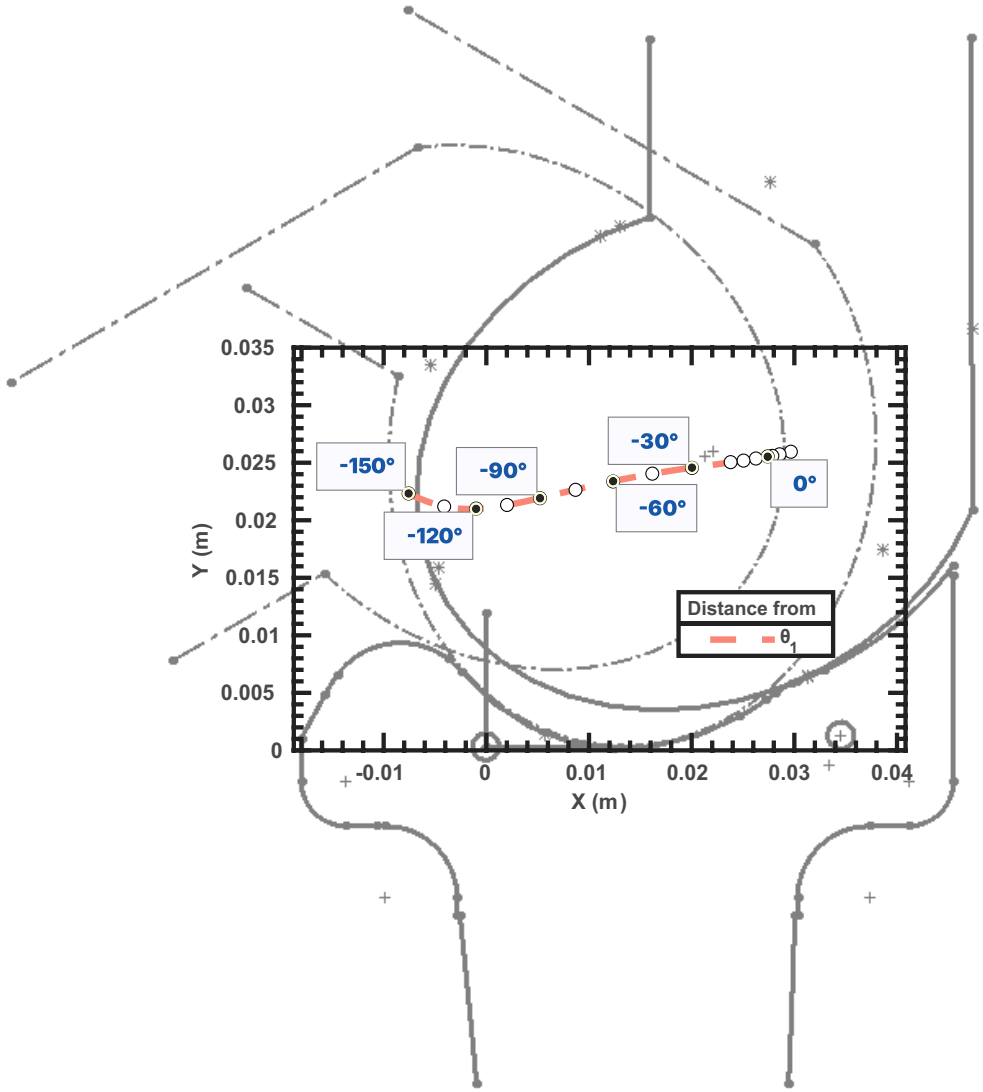


Figure 4.2: ICR for biomimetic robot knee as a function of knee angle. Measured from θ_1 on the tibia. Plot is overlayed on model of the knee at $\theta_k = 0^\circ$ (solid line), and -60° and -120° (dotted dash lines). White circles represent measured points in -15° increments from -150° to -15°, -5° increments from -15° to 0°, and then at 2°, 5°, and 10°.

additionally tested with ∅20 mm Festo BPAs in a paired configuration: the extensor with two BPAs in parallel, and the flexor with two BPAs whose pathways are mirrored about the Z axis. Each artificial muscle is pinned at the muscle origin location on the femur while the other end is either pinned or attached to the muscle insertion location on the tibia via an artificial tendon.

Torque about the knee is generated by the moment of the BPA force about the joint's instantaneous center of rotation (ICR). Two vectors enter that moment: the position vector $\mathbf{d}$, drawn from the ICR to the point where the actuator force is applied, and the force vector $\mathbf{F} = \|\mathbf{F}\|\hat{\mathbf{u}}$, whose direction $\hat{\mathbf{u}}$ is the line of action of the actuator. The moment of the force about the joint is $\mathbf{m} = \mathbf{d} \times \mathbf{F}$, and its magnitude is the force magnitude multiplied by the moment arm r, the perpendicular distance from the joint axis to the line of action,

$$r = \|\mathbf{d} \times \hat{\mathbf{u}}\| \tag{4.1}$$

The moment arm is therefore a scalar, and it is the same quantity whether it is computed from the position vector with Eq. (4.1) or measured directly with a tape measure. In this study it is calculated using the method specified by **?**. The scalar moment arm r_k about the joint axis $\hat{\mathbf{z}}_k$ (i.e., the Z axis in the knee frame) is calculated as

$$r_k = \mathbf{p}_{proj,xy} \cdot \frac{\mathbf{p}_f \times \hat{\mathbf{z}}_k}{\|\mathbf{p}_f \times \hat{\mathbf{z}}_k\|} \tag{4.2}$$

A muscle path may be modeled as a single segment running from origin to insertion, or as a chain of segments joined at via points. The segment that crosses from one reference frame into the other, here from the femur frame to the tibia frame, is the free muscle

segment: because its endpoints ride on different bodies, its length and its direction respond directly to joint rotation, whereas a segment contained within one body's frame rotates rigidly with that body and changes neither. Here $\mathbf{p}_f$ is the force-path segment vector, and $\mathbf{p}_{proj,xy}$ locates a point on that line relative to the joint, projected onto the plane normal to the joint axis. The normalized cross product in Equation (4.2) is the in-plane normal to the force path. Its dot product with the projected position gives the signed moment arm (Figure 2.6). Moment arm is calculated in the tibial body frame.

Details of the existing model and the torque calculation methodology in the modeled robot are given by **??**.

4.1.2 Knee Torque Data Collection and Equipment

A test jig was built to collect knee torque measurements at different knee angles over its range of motion. The test stand frame is made predominantly out of 80/20 1515 series extruded aluminum components. The knee joint is free to rotate the tibia while the femur is fixed to the frame. A force sensor has one end connected to the tibia and the other is connected to a swing arm. The swing arm is tied to a winch with 3/16 inch Kevlar rope from Quality Nylon Rope routed over intermediate idler pulleys (see Fig. 4.1**B** and 4.1**C**). At each measured position of θ_k the winch would be locked and the BPA would be inflated and then deflated. This test procedure kept the measured results as isometric as possible. The knee joint is co-axial with gravity to eliminate gravitational torque (see Fig. 4.1**A** and Fig. 4.1**C**).

Force data was collected using one of two different sensors. The first is a MODERN

STEP 300 kg digital crane scale. The second force sensor is a CALT DYLY-103 100 kg S shaped load cell. The load cell is used in conjunction with a HX711 Load Cell Amplifier. Pressure is measured with a Freescale MPX5700 GP 5 V pressure sensor. Building air supply pressure is controlled with two pressure regulators in series. The first is a Parker model 20R113GC $0 - 120$ psi pressure regulator. The second is a Husky $3/8$ inch High Performance Air Regulator HDA72200. A Festo VTUG-10-MRCR-S1T-26V20-T516LA-UL-T532S-8K valve manifold was used to supply air to the actuator. Pressure and load cell amplifier data are sent to MATLAB via an Arduino Uno style Sparkfun BlackBoard C microcontroller.

The length measurements were made using a tape measure. When there was sufficient curvature in the BPA during knee extensor torque tests, a flexible piece of string was used to mark the axial length and then transferred to the tape measure. Angle measurements were taken with a Medigauge digital electronic goniometer or with MALENOO analog goniometers of 6, 8, or 12 inch lengths.

4.1.3 Hybrid Torque Calculation Method

A Hybrid Torque calculation method was developed to help uncover inconsistencies between the model and recorded torque values. This method measures the actual muscle length and angles of the robot to calculate force and joint torque. With this method, it is easier to determine which parts of the torque equation did not achieve their predicted value and provide a correction to that part of the equation. The classical mechanics way to calculate the torque about a joint is to take the cross product of the position vector **d** and

the force vector **F**, both defined in Section 4.1.1,

$$\mathbf{m} = \mathbf{d} \times \mathbf{F} \tag{4.3}$$

Knee torque is the scalar component of this moment about the joint axis, written M_z to make its axis explicit. Because the moment arm is the perpendicular distance from that axis to the line of action, the vector cross product of Eq. (4.3) reduces to an ordinary scalar product,

$$M_z = r\,F \tag{4.4}$$

where r is the scalar moment arm of Eq. (4.1), here measured directly with a tape measure as the perpendicular distance from the knee axis to the actuator's line of action, and F is the scalar force magnitude, calculated by measuring the BPA's P, l_0, l_{620} and l_{m} values and using the model fit equations determined in the following section. The hybrid method therefore replaces the model-predicted moment arm r_k of Eq. (4.2) and the model-predicted muscle length with their measured counterparts, leaving the force model as the only predicted quantity.

4.1.4 Improved Torque Model

The classical kinematic torque model assumed rigid links and inextensible muscle paths. Data collected during the first flexor experiments (Section 4.2) showed that this assumption does not hold for a pneumatic system: measured muscle lengths were consistently shorter

than predicted, and measured torque deviated from prediction in a load-dependent manner. This section describes the corrections that were developed to account for a constant length offset, series compliance in the attachment hardware and artificial tendons, and the loss of usable muscle length when a BPA wraps around the joint.

These corrections are deliberately formulated to be identified from experiment rather than derived analytically, because the physical effects they represent are distributed and not separately measurable. The lumped compliance absorbs the cantilevered three-dimensional-printed attachment brackets, whose geometry is a printed CAD part rather than a prismatic beam with a tractable second moment of area, together with the bone fixtures, the cable winch of the test stand, and the artificial tendon. Predicting that aggregate stiffness from first principles would require finite-element models of each component with material properties (anisotropic printed polymer, braided sleeve, construction-dependent tendon stiffness) that are themselves unmeasured. The correction terms therefore parameterize only what the torque model needs: χ_0 absorbs the constant length discrepancy observed at zero load, χ_1 and χ_2 are effective stiffnesses in the bracket frame's principal directions, and χ_3 scales the wrapping-length loss. Each is identified by global optimization against the measured torque rather than solved in closed form, for two reasons. First, the terms enter the prediction through an implicit force balance, because the compliance displacement depends on the force that it itself modifies, so the inverse problem has no analytic solution. Second, the parameters are only jointly identifiable: preliminary fits showed that widely different stiffness pairs can reproduce the training tests nearly equally well, so the adopted solution must be selected by cross-validated prediction on held-out tests rather than by residual minimization alone, which is what the multiobjective

Pareto-front procedure of Section 4.1.5 provides.

Length Offset and System Stiffness

Analysis of measured muscle length data shows that it was consistently shorter than the predicted muscle length at all joint angles. This shorter length results in less muscle force, a shifted RoM, and less torque in the experimental setup. That this shorter length occurs when $\varepsilon = 1$ and $F = 0$ indicates that a constant length offset should be used. Therefore, to produce an improved torque prediction model, l_{m} was recalculated as

$$l_{\mathrm{m}} = l_{LMT} - 2\, l_{fitting} - \chi_0 \tag{4.5}$$

where l_{LMT} is the path length of the musculotendon from origin to insertion point, $l_{fitting}$ is the length of the end caps, and χ_0 is a constant length offset.

Additionally, as the muscles produce more force, compliant components in the limb segments and attachment hardware deform. This deformation changes the effective braided pneumatic actuator (BPA) length, which alters contraction and therefore changes the force predicted by the BPA model. This effect is most apparent at high forces and low contractions. The purpose of the compliance model developed next is predictive, not structural: it exists so that the torque calculation can correct the actuator length for load-dependent deformation, and it makes no claim about the stiffness of any one component in isolation. Accordingly, the stiffness values that populate the model, χ_1 and χ_2, are not measured on the hardware directly; they are identified by the optimization of Section 4.1.5 from the measured torque, and every stiffness value quoted later in this chapter is an

optimizer output. The combined hardware compliance is approximated using Hooke's law in the bracket's principal directions.

A local bracket frame was established at the attachment bracket. The $\hat{\mathbf{x}}_{br}$ axis was aligned with the vector from the bracket reference point to the BPA attachment point and represents the stiffest direction of the bracket. The $\hat{\mathbf{y}}_{br}$ axis was chosen to point in the posterior direction with respect to the tibia frame, corresponding to the primary cantilever bending direction (see Fig. 4.4), while the $\hat{\mathbf{z}}_{br}$ axis completes the right-handed frame. In this frame the bracket is modeled as three independent linear springs along the principal axes, with stiffnesses k_x, k_y, and k_z in the x_{br}, y_{br}, and z_{br} directions.

For the two-rotation evaluator, the bracket frame was constructed from the vector joining the bracket origin to attachment point $\mathbf{p}_A$ or $\mathbf{p}_B$. The first rotation, through θ_z about the fixed space-frame axis $\hat{\mathbf{z}}_s$, points the intermediate axis $\hat{\mathbf{x}}_1$ toward the sagittal-plane projection, $\mathbf{p}_A^{\text{sag}}$ for the insertion point $\mathbf{p}_A$ and $\mathbf{p}_B^{\text{sag}}$ for the origin point $\mathbf{p}_B$. The second rotation, through $-\theta_y$ about the current axis $\hat{\mathbf{y}}_1 = \hat{\mathbf{y}}_{br}$, brings the final bracket axis $\hat{\mathbf{x}}_{br}$ into alignment with the full three-dimensional point vector. Figure 4.3 shows these geometric targets. The resulting orientation and homogeneous transform are

$$\mathbf{R}_{sb} = \mathbf{R}_z(\theta_z)\mathbf{R}_y(\theta_y)^T, \qquad \mathbf{T}_{sb} = \begin{bmatrix} \mathbf{R}_{sb} & \mathbf{p}_{br} \\ \mathbf{0}^T & 1 \end{bmatrix}. \tag{4.6}$$

$$\mathbf{K}_{br} = \operatorname{diag}\big(k_x,\ k_y,\ k_z\big) \tag{4.7}$$

These three principal stiffnesses are not independent in this model, and the reason follows from the frame construction itself. The first rotation of Eq. (4.6) is fixed by the

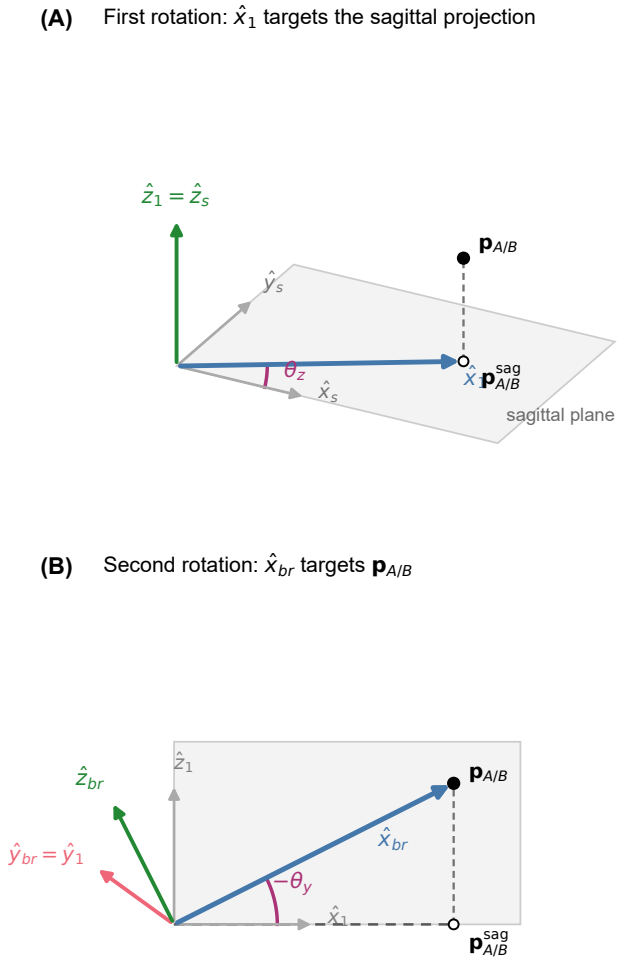


Figure 4.3: Two-rotation construction of the bracket frame from attachment point $\mathbf{p}_A$ or $\mathbf{p}_B$. **(A)** Rotation through θ_z about fixed $\hat{\mathbf{z}}_s$ aligns $\hat{\mathbf{x}}_1$ with the point's sagittal-plane projection. **(B)** Rotation through $-\theta_y$ about current $\hat{\mathbf{y}}_1 = \hat{\mathbf{y}}_{br}$ aligns $\hat{\mathbf{x}}_{br}$ with the full point vector.

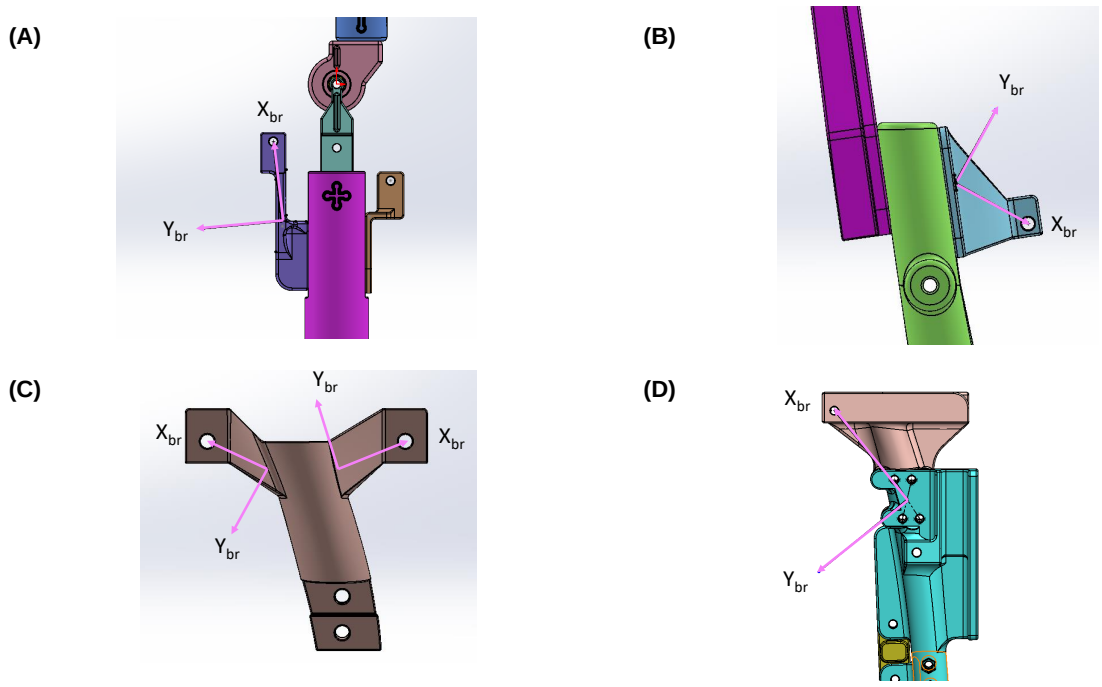


Figure 4.4: Bracket frame orientations for: **(A)** flexor configuration on the pinned knee, **(B)** extensor configuration on the pinned knee, **(C)** flexor and extensor configurations on the biomimetic knee with ∅10 mm BPAs, and **(D)** flexor configuration on the biomimetic knee with the ∅20 mm BPA.

sagittal-plane projection of the attachment point, but the orientation of the bracket frame about its own $\hat{\mathbf{y}}_{br}$ axis is a modeling choice: the pitch-only convention (no second rotation) and the two-rotation convention of Fig. 4.3 describe the same physical bracket and differ by exactly a rotation about $\hat{\mathbf{y}}_{br}$. A diagonal stiffness array is invariant under such a rotation, $\mathbf{R}_y(\theta)\,\mathbf{K}_{br}\,\mathbf{R}_y(\theta)^T = \mathbf{K}_{br}$ for all θ, if and only if the two stiffnesses in the plane of the rotation are equal, $k_x = k_z$. Adopting this y-symmetric form therefore makes the compliance prediction independent of the frame convention, and the redundant third stiffness is eliminated: with this array the two frame conventions produce identical torque predictions to numerical precision.

The change of notation that follows is deliberate. The symbols k_x, k_y, and k_z name the stiffness components of the general bracket-frame array, Eq. (4.7); χ_1 and χ_2 are the two identified stiffness parameters, the quantities the optimization of Section 4.1.5 solves for, and each enters the diagonal array in a fixed axis assignment. Two stiffness values remain, χ_1 in the x_{br} and z_{br} directions and χ_2 in the cantilever bending direction y_{br}:

$$\mathbf{K}_{br} = \begin{bmatrix} \chi_1 & 0 & 0 \\ 0 & \chi_2 & 0 \\ 0 & 0 & \chi_1 \end{bmatrix} \tag{4.8}$$

The corresponding local compliance matrix is

$$\mathbf{C}_{br} = \mathbf{K}_{br}^{-1} \tag{4.9}$$

The force applied to the bracket was expressed in the bracket frame as $\mathbf{F}_{br}$. This force vector can be written in terms of its magnitude and unit direction as

$$\mathbf{F}_{br} = F\,\hat{\mathbf{u}} \tag{4.10}$$

where $F = \|\mathbf{F}\|$ is the scalar force magnitude of Eq. (4.4) and $\hat{\mathbf{u}} = [u_x, u_y, u_z]^T$ is the BPA force direction expressed in the bracket frame. The bracket deflection vector, $\boldsymbol{\epsilon}_{br}$, was calculated from Hooke's law as

$$\mathbf{F}_{br} = \mathbf{K}_{br}\,\boldsymbol{\epsilon}_{br} \tag{4.11}$$

Solving Eq. 4.11 for the bracket deflection gives

$$\boldsymbol{\epsilon}_{br} = \mathbf{K}_{br}^{-1}\,\mathbf{F}_{br} = \mathbf{C}_{br}\,\mathbf{F}_{br} \tag{4.12}$$

Substituting Eq. 4.10 into Eq. 4.12 gives

$$\boldsymbol{\epsilon}_{br} = F\,\mathbf{K}_{br}^{-1}\,\hat{\mathbf{u}} \tag{4.13}$$

For the diagonal stiffness matrix in Eq. 4.8, Eq. 4.12 is equivalent to the component-wise expression

$$\boldsymbol{\epsilon}_{br} = \begin{bmatrix} F\,u_x/\chi_1 \\ F\,u_y/\chi_2 \\ F\,u_z/\chi_1 \end{bmatrix} \tag{4.14}$$

Although $\boldsymbol{\epsilon}_{br}$ is a three-dimensional deflection vector, the BPA force model depends on the scalar change in effective BPA length (δ). Therefore, only the component of $\boldsymbol{\epsilon}_{br}$

along the BPA line of action directly changes the effective BPA length. This scalar bracket deflection is the projection of $\boldsymbol{\epsilon}_{br}$ onto $\hat{\mathbf{u}}$:

$$\delta_{br} = \hat{\mathbf{u}}^T \boldsymbol{\epsilon}_{br} \tag{4.15}$$

Substituting Eq. 4.12 into Eq. 4.15 gives

$$\delta_{br} = \hat{\mathbf{u}}^T \mathbf{K}_{br}^{-1} \mathbf{F}_{br} \tag{4.16}$$

Using Eq. 4.10, Eq. 4.16 becomes

$$\delta_{br} = F \, \hat{\mathbf{u}}^T \mathbf{K}_{br}^{-1} \hat{\mathbf{u}} \tag{4.17}$$

The scalar stiffness along the BPA line of action due to the bracket compliance is therefore

$$k_{\hat{u},br} = \frac{F}{\delta_{br}} = \left(\hat{\mathbf{u}}^T \mathbf{K}_{br}^{-1} \hat{\mathbf{u}} \right)^{-1} \tag{4.18}$$

This projected-compliance form is used because the force is applied along $\hat{\mathbf{u}}$, while the bracket is allowed to deflect independently along its local $\hat{\mathbf{x}}_{br}$, $\hat{\mathbf{y}}_{br}$, and $\hat{\mathbf{z}}_{br}$ directions. The reduced-order Hooke's law expression for the bracket deflection along the BPA line of action is therefore

$$F = k_{eq} \, \delta \tag{4.19}$$

Here k_{eq} is the series sum of every compliant element in the actuator path, and the

bracket-only case is the special case with no tendon. When the BPA attaches directly to the bracket, k_{eq} excludes the tendon and equals $k_{\hat{u},br}$ of Eq. (4.18). When an artificial tendon was used, it was modeled as an axial spring with stiffness k_{ten}: bracket deflection and tendon stretch both shorten the effective BPA length along the actuator line of action, the two elements carry the same force, and their scalar compliances add in series,

$$k_{eq} = \left(\frac{1}{k_{\hat{u},br}} + \frac{1}{k_{ten}} \right)^{-1} \tag{4.20}$$

This series reduction is exact only under a specific set of conditions, each of which corresponds to a way the physical hardware could violate it. The bracket and the tendon must carry the same force, which holds for two massless elements coupled in series under quasi-static loading; a compliant coupling between them would break this. Their displacements must be kinematically additive along the line of action, with no displacement lost to off-axis geometry. Within a single iteration of the force balance the line of action $\hat{\mathbf{u}}$ is treated as fixed, which is an approximation, and one the results justify: once the balance returns the compliance displacement, the calculation updates the attachment locations with the deflected bracket positions and re-evaluates $\hat{\mathbf{u}}$ and the moment arm at the deflected geometry, as described below. The elements are assumed to carry no moment, so rotational compliance and force–moment coupling at the attachments are neglected. Both elements are linear, with stiffness independent of deflection, and the tendon is modeled without preload, so its tangent stiffness does not stiffen with load state. Finally, compliance deflection shortens the actuator, as Eq. (4.21) assumes, because a BPA produces force only in tension: under its own pulling force the compliant attachments give way in the direction

that relieves that tension. A deflection that would lengthen the actuator can be produced by an antagonistic pull, but that loading lies outside this treatment.

Since δ is the scalar compliance displacement along the BPA line of action, then the compliance-corrected active BPA length used in the force balance is

$$l_m(\delta) = l_m(0) - \delta \tag{4.21}$$

where $l_m(0)$ is the initial muscle length calculated from (4.5).

For a specified set of correction parameters, the following calculation solves the attachment-compliance displacement at each knee angle; it is distinct from the optimization that identifies those parameters from experimental torque data (Section 4.1.5). The corresponding relative contraction is $\varepsilon^*(\delta)$. Since the BPA force depends on contraction, and contraction depends on the compliance displacement, δ is the root of the nonlinear force-balance equation

$$F_{\max} \cdot F_{\mathrm{BPA}}(\varepsilon^*(\delta), P^*) - k_{eq}\,\delta = 0 \tag{4.22}$$

The root is found with a bracketed scalar solver (`fzero` in MATLAB) on the physically admissible interval between zero deflection and the deflection that would bring the actuator to its maximum usable contraction. Termination is controlled by the bracket, not by a residual tolerance: the iteration closes in on the root and stops when the bracket has collapsed to adjacent floating-point values, so the returned δ is converged to machine precision. Two safeguards act as secondary stop conditions. If the actuator already sits at or beyond its maximum usable contraction at zero deflection, no root is sought and δ is set

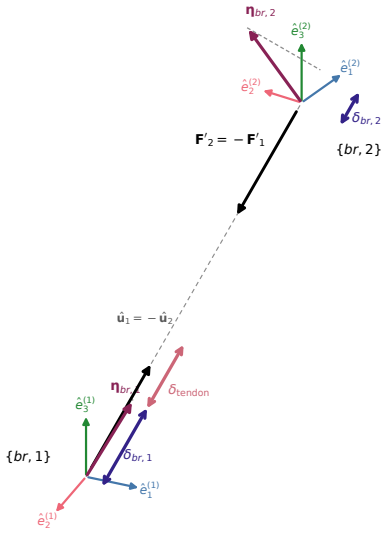
to zero; if the residual shows no sign change across the interval, the solver likewise returns $\delta = 0$ rather than a spurious root.

After solving Eq. 4.22, the corresponding force magnitude was used in Eq. 4.12 to compute the vector bracket deflection. This deflection was transformed back into the limb reference frame and used to update the BPA attachment location. The routed BPA length, force direction, moment arm, and torque were then recalculated using the updated attachment location. In the present implementation, the optimized stiffness terms χ_1 and χ_2 were treated as effective compliance parameters for the attachment system rather than direct measurements of an isolated bracket. This allowed unmodeled compliance in the surrounding hardware to be represented by the reduced-order stiffness model.

The actuator path can carry compliant hardware at both ends: the insertion-side bracket of the preceding development together with a second, origin-side bracket. Figure 4.5 distinguishes the three-dimensional displacement of each bracket from the scalar displacement along the BPA force path. For bracket j, the displacement resolves into components $\eta_{1,j}$, $\eta_{2,j}$, and $\eta_{3,j}$ along the local bracket-frame basis, whereas only its projection $\delta_{br,j}$ along the corresponding force direction $\hat{\mathbf{u}}_j$ changes the actuator length. The two bracket projections act in series with the tendon elongation δ_{tendon}:

$$\begin{aligned} \boldsymbol{\eta}_{br,j} &= \sum_{i=1}^{3} \eta_{i,j} \hat{\mathbf{e}}_i^{(j)}, \\ \delta_{br,j} &= \hat{\mathbf{u}}_j^T \boldsymbol{\eta}_{br,j}, \qquad j \in \{1, 2\}, \\ \delta_{\text{series}} &= \delta_{br,1} + \delta_{\text{tendon}} + \delta_{br,2}. \end{aligned} \tag{4.23}$$

(A)



(B)

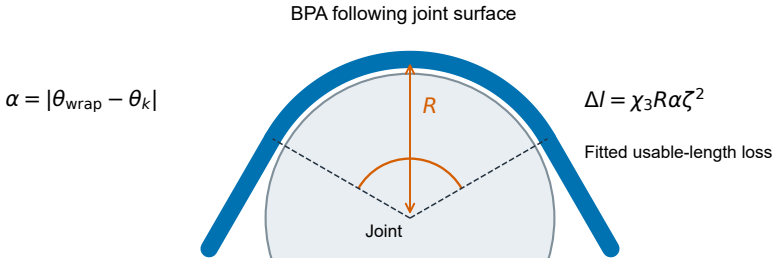


Figure 4.5: **(A)** Bracket displacement and its projection along the force path; tendon elongation acts in series. **(B)** Wrapping radius R and bend angle $|\theta_{\text{wrap}} - \theta_{\text{k}}|$ in the length-loss correction of Equation (4.24). The lost usable length is a fitted correction, not the difference between arc and chord lengths.

BPA Fictive Contraction

The extensor configuration required one further correction, for a loss with a different physical origin than the compliance terms. As the knee flexes, the extensor BPA bends over the front of the joint and makes contact with it, so part of the actuator follows the joint surface instead of a straight line. Bent over the joint this way, the BPA behaves like a beam: the pressurized wall on the inside of the bend carries less tension than the wall on the outside, which both unbalances the actuator's force and consumes part of the wrapped arc length that would otherwise shorten the path. The wrapped section therefore stops contributing usable contraction length, and the actuator behaves as if it were more contracted than it physically is. This excess is termed fictive strain: the additional contraction the force model must see to reproduce the observed force, beyond the actuator's physical contraction. Figure 4.5**B** shows the wrapping geometry. A wrapping-loss term χ_3 was introduced to scale the length that this fictive strain removes, such that

$$\Delta l = \begin{cases} \chi_3 \; R \; |\theta_{\text{wrap}} - \theta_{\text{k}}| \; \zeta^2 & \theta_{\text{k}} \leqslant \theta_{\text{wrap}} \\ 0 & \theta_{\text{k}} \geqslant \theta_{\text{wrap}} \end{cases} \tag{4.24}$$

where Δl is the change in usable length of the BPA, R is the wrapping radius, $|\theta_{\text{wrap}} - \theta_{\text{k}}|$ is the angle through which the BPA is bent once contact begins, and θ_{wrap} is the knee flexion angle at which the BPA starts to bend over the knee joint, determined from solid modeling. ζ is the additive complement to relative strain, such that

$$\zeta = 1 - \varepsilon^* \tag{4.25}$$

so the wrapping loss grows quadratically as the actuator approaches its resting length and vanishes at full contraction. Combining (4.25) into (4.24), and subtracting Δl from (4.5) yields

$$l_{\mathrm{m}} = l_{LMT} - 2\, l_{fitting} - \chi_0 - \Delta l \tag{4.26}$$

Multiobjective Identification Algorithm

Multiobjective optimization using `gamultiobj` in MATLAB was used to determine the length offset (χ_0) and the stiffness terms (χ_1 and χ_2). The error between the model and the experimental torque was minimized across the pinned-knee flexor tests, with the $l_0 = 45.7\,\mathrm{cm}$ test held out for validation.

The optimization algorithm was again `gamultiobj` in MATLAB. This time, the muscle origins bracket was chosen. The bracket frame is shown in Fig. 4.4**B**. Since the bracket forms an arch shape and extends for 75 mm in the Z direction, for this configuration, we set χ_2 in the X_{br} direction, and χ_1 in the Y_{br} and Z_{br} directions. The optimization algorithm for this set solves for the best average goodness-of-fit (GoF) measures RMSE, FVU, and maximum absolute residual for two BPAs, while the third one was held out. Three independent rounds of optimization were performed, so that one of each of the BPAs could be held out as validation. The solution was then selected by the shortest distance between the average normalized GoF and the validation data, subject to improved GoF

values for all the individual BPAs, with the exception of maximum absolute residual on the $l_0 = 48.0\,\text{cm}$ BPA.

This initial identification was subsequently extended to the full test sets in two stages. First, the flexor model was given a second compliant bracket at the origin-side attachment (a hip-frame cantilever carrying the BPA origin), so that the compliance of both ends of the actuator path is represented; the stiffness pair (χ_1, χ_2) is shared by the two brackets in fixed axis assignments (Appendix C, Table C6), leaving a three-term fit (χ_0, χ_1, χ_2) cross-validated over the four remaining flexor tests, two of them held out per fold; the encoder-corrected $l_0 = 47.9\,\text{cm}$ test was not used in the two-bracket identification. Second, the extensor refit locked (χ_1, χ_2) at the flexor-identified values, enforcing a single stiffness pair across the flexor and extensor configurations, and solved only the extensor-specific terms (χ_0, χ_3) on a three-test training set with three tests held out for validation, with the extensor length offset allowed to take negative values. The bounds and the adopted values of both stages are recorded in Appendix C (Tables C7 and C8); the fit outcomes are summarized in Tables 4.2 and 4.3.

The goodness-of-fit metrics used to judge the optimization results are the root mean square error (RMSE), the fraction of variance unexplained (FVU), and the maximum absolute residual, each computed between the measured and predicted isometric torque over the tested range of motion.

Each metric isolates a different aspect of disagreement between the two curves. RMSE reports the overall magnitude of the error and, when the curve shapes already agree, is dominated by the remaining level offset. FVU normalizes the error variance by the variance of the measured torque, making it scale-free: it asks whether the prediction

reproduces the variation of the measurement rather than whether it is close in absolute units, is the complement of the Nash–Sutcliffe model efficiency [**?**], and takes values below one when the prediction outperforms the mean of the measurement itself. The maximum absolute residual is a worst-case local measure. A small shift of the predicted curve along the knee-angle axis produces large pointwise residuals where the torque gradient is steep and small residuals elsewhere, so this metric is sensitive to phase misalignment while responding only weakly to uniform level offsets.

This interpretation follows the standard decomposition of curve-comparison error into offset, amplitude, and pattern components [**?**]: the mean square error between two curves separates into the squared mean offset, an amplitude mismatch, and a residual shape term. An equivalent three-term formulation, splitting waveform similarity into offset, amplitude, and shape, underlies the linear fit method developed for clinical gait waveform comparison [**?**]. The metrics used here map onto this decomposition, with RMSE carrying the level disagreement, FVU the shape and amplitude agreement, and the maximum absolute residual serving as the phase-sensitive worst case. FVU is treated as the primary criterion because it is the only one of the three that is directly comparable across tests with different torque scales.

4.1.5 Optimization Algorithms and Cost Functions

Optimization enters this chapter in several places, and in every one of them it replaces a calculation that cannot be carried out any other way. The correction terms of the improved torque model are quantities that no instrument measures directly: they are hidden

inside an implicit force balance, and their best values are the ones that make the model's predicted torque match the measured torque. The route redesign is the mirror problem: the attachment locations are free design variables, and their best values are the ones that make the redesigned routes meet or exceed human torque capability. The cam-profile variant of the extensor route extends the redesign once more, searching over the shape of the tibial-side routing itself. In each case a search algorithm proposes candidate values, the physics model evaluates them, and the search proceeds until a stopping rule is met. This section states what is searched for, which solvers were chosen and why, and when each search stops.

Two families of optimization problem underpin the results of this dissertation and the design framework built on them. The first is the identification of the correction terms of Section 4.1.4 (χ_0, χ_1, χ_2, and χ_3) from the pinned-knee test data. The second is the redesign of the biomimetic knee route geometry, meaning the BPA origin $\mathbf{p}_1$, insertion $\mathbf{p}_{\text{end}}$, resting length, and tendon length (Appendix C), so that the flexor and extensor configurations meet or exceed human knee torque capability over the range of motion, together with its cam-profile extension, in which the tibial-side route points are free to trace a cam profile (Section 4.3.2). Both are carried out in MATLAB [**?**].

The correction-term identification is a genuine multiobjective problem: candidate parameter sets are judged simultaneously on RMSE, FVU, and maximum absolute residual (Section 4.1.4). It is solved with the genetic multiobjective solver `gamultiobj`, an implementation of the elitist nondominated-sorting genetic algorithm NSGA-II [**??**], which returns a Pareto front of parameter sets from which the adopted solution is selected by the holdout-validation procedure described above. The production runs use adaptive-

feasible mutation and scattered crossover at a fraction of 0.8, with MATLAB's default Pareto fraction of 0.35; the population settings are stage-specific. The flexor two-bracket fit evolves a population of 150 for up to 600 generations and terminates early once the geometric average relative change in the spread of the Pareto front falls below 4×10^{-3}; the extensor refit uses a population of 50 for up to 150 generations with a spread tolerance of 1×10^{-3}. In both stages the parameters are searched in a normalized space, with length offsets in centimeters and stiffnesses on a base-ten logarithmic scale, so that the genetic search samples the stiffness bounds, which span an order of magnitude, uniformly by decade, and the initial population is seeded around the expected solution region. The search bounds and the adopted values are recorded in Appendix C.

The route redesign replaces the multiobjective formulation with a single scalar cost assembled from weighted penalty terms, because the route prediction (contact activation, wrapping, and compliance correction) is expensive to evaluate and only moderately smooth. The cost has the form

$$J = w_w J_{\text{worst}} + w_t J_{\text{torque}} + \textstyle\sum_a J_a + J_{\text{geom}} + J_{\text{len}}, \tag{4.27}$$

where J_{worst} is the largest fraction by which robot torque falls short of the human target at any angle across the range of motion, weighted heavily ($w_w = 10^5$); J_{torque} is the mean-squared shortfall fraction ($w_t = 10^3$), weighted two orders of magnitude below the worst-angle term so that no shortfall at a single angle can be traded away for a better average; the J_a are configuration-specific auxiliary penalties described below; and J_{geom} and J_{len} are regularizers that keep the endpoints and actuator lengths near their seed

values. Human torque targets are computed from the Gait2392 lower-extremity model [**??**]: the flexor is graded against the biceps femoris short head and the extensor against vastus medialis. Each target is a maximum-isometric torque curve: OpenSim moment arms, obtained as the partial derivative of muscle-tendon length with respect to the knee coordinate, $r(\theta_k) = \partial l_{\mathrm{MT}} / \partial \theta_k$, are evaluated across the range of motion and multiplied by each muscle's maximum isometric force (provenance and exported tables in Appendix A). Hard constraints (bone clearance, wrap feasibility, actuator strain and contraction limits, tendon geometry) are enforced alongside the cost, and any candidate whose prediction fails is rejected with a prohibitive penalty.

The solution pipeline is common to both configurations. Each run begins with a full sanity evaluation of the seed design and a small smoke test of the solver setup, then executes `surrogateopt`, MATLAB's surrogate-assisted global solver for expensive black-box objectives, which builds a spatial interpolation surrogate over sampled candidate points and searches it adaptively [**?**], in parallel under per-configuration evaluation budgets of 3,000 for the extensor route and 1,000 for the flexor route (the later cam variant of Section 4.3.2 uses 1,000 surrogate and 1,500 refinement evaluations), and finally refines the best feasible point with `patternsearch`, a derivative-free mesh-adaptive direct search method [**??**], for 15,000 further evaluations. The surrogate budgets sit well past the observed plateau of a few hundred evaluations, so they cap run time rather than define convergence. Each stage stops on its evaluation budget; the pattern search also stops early when its mesh collapses below the mesh tolerance with the constraints satisfied, which is the usual exit once the design has converged. The entry scripts and their cost functions are archived with the code.

Although both configurations share this pipeline, their cost functions differ in ways that reflect the mechanics. The flexor BPA routes around a single moving wrap point, so clearance from the tibia and femoral condyles is enforced by continuous signed-distance constraints, and its cost adds an off-axis torque penalty (torques about the joint axes other than the knee axis, relative to the original single-BPA route) and a penalty if the wrap fails to release cleanly as the knee extends. The extensor BPA wraps over the front of the joint through as many as seven sequential contacts, so clearance is enforced instead by the contact-elimination rule that governs the route, and its cost penalizes torque overshoot above the human curve and excessive actuator contraction, which risks bottoming out at high flexion. The flexor and extensor targets alike carry a required torque margin of 5 % above the human curve, so that a route that merely touches the target at any angle is rejected. With two BPAs per joint, the flexor route is mirrored about the mid-sagittal plane and the two actuator forces are summed, while the extensor model scales the force of a single shared route by the actuator count.

4.2 Results

With the actuator force model validated at the actuator level, the same model was used to predict isometric torque about the knee joint for the test stands described in Sections 4.1.1 and 4.1.3. Measured torque is compared against three predictions: the original rigid-body torque model of Section 2.4, a prediction using the improved BPA force characterization alone (Chapter 3), and the improved torque model of Section 4.1.4 after optimization. The hybrid torque method (Section 4.1.3) is used to isolate the sources of error.

4.2.1 Flexor on Pin Joint

Figure 4.6 shows the measured flexor torque on the pinned knee next to the pre-optimization prediction series for the four cross-validation tests. Both pre-optimization force models overpredicted the magnitude of the measured flexor torque, and on the angles where the two predictions are both defined the improved BPA characterization fit the measurements better than the original rigid-body prediction on every test, although it degraded near full extension, where the rigid-body prediction is undefined altogether (Table 4.1). The back-calculated (hybrid) torque reproduced the measurements closely, which localizes the misprediction to the modeled kinematics, the muscle length l_{m} and the moment arm r, rather than to the force model.

For Table 4.1, the rigid-body and improved-characterization predictions are scored on the same points where the rigid-body strain is defined: 75 of 90, 17 of 20, 14 of 19, and 15 of 22 points in panels A–D, respectively. The hybrid calculation is scored on all measured points. These pre-optimization scores therefore use a different point set from the full-test corrected-model scores in Table 4.2.

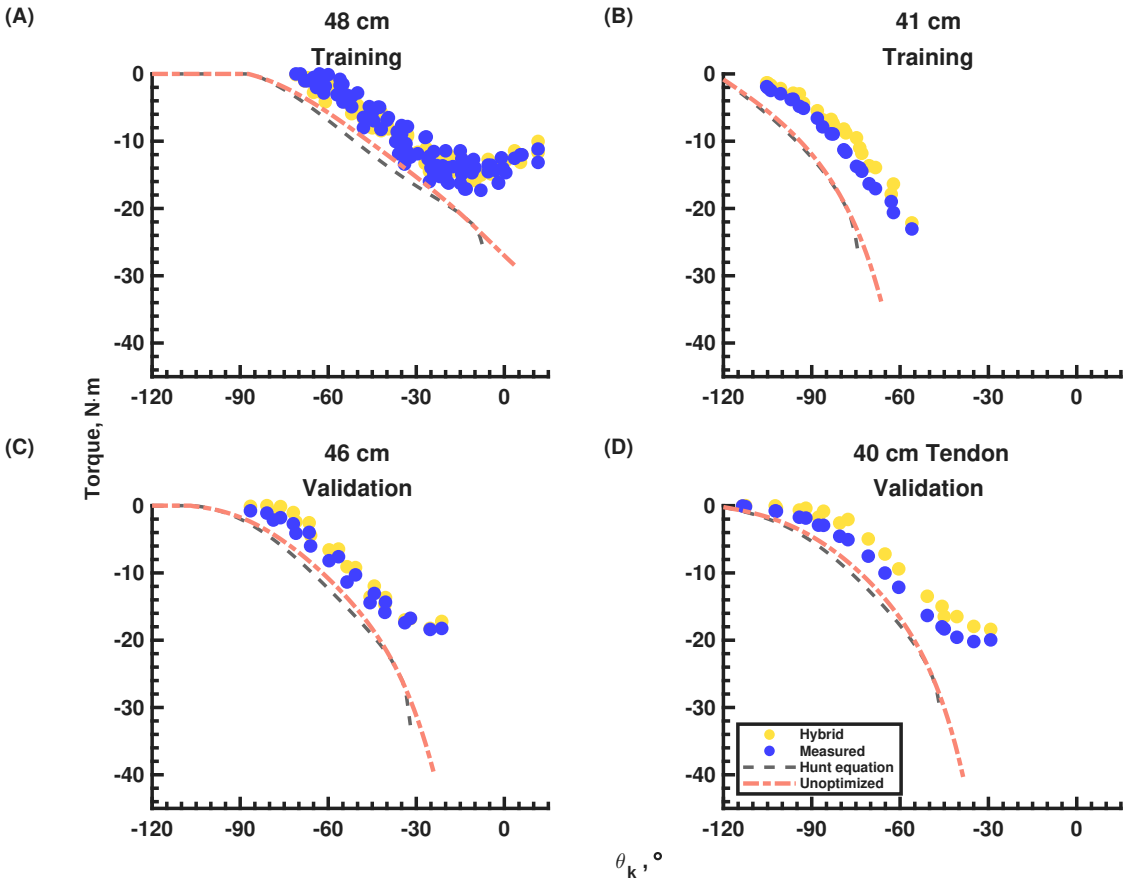


Figure 4.6: Pre-optimization pinned-knee torque versus knee angle for the four cross-validation flexor tests. Knee angle is negative in flexion, and flexion torque is negative. Gold circles are the hybrid calculation, which combines the measured muscle length and moment arm with the modeled force; indigo filled circles are experimentally measured torque; the gray dashed line is the rigid-body prediction; and the light orange dash-dot line is the unoptimized prediction, the improved BPA force characterization applied to the same rigid kinematics. Panel titles give the nominal resting length; the actual resting lengths are $l_0 = 48.5\,\text{cm}$ **(A)**, $l_0 = 41.0\,\text{cm}$ **(B)**, $l_0 = 45.7\,\text{cm}$ **(C)**, and $l_0 = 40.6\,\text{cm}$ with tendon **(D)**. Panels **(A)** and **(B)** are the training tests and panels **(C)** and **(D)** the validation tests of the adopted cross-validation fold.

Table 4.1: Pre-optimization pinned-knee flexor torque-prediction goodness of fit per test, comparing the three pre-optimization torque predictions of Figure 4.6: the Hunt-equation rigid-body prediction, the unoptimized prediction of the improved BPA characterization, and the back-calculated (hybrid) torque. RMSE and maximum absolute residual (Max.) are in N m, and FVU denotes the fraction of variance unexplained.

		Hunt equation			**Improved char. alone**			**Back-calculated (hybrid)**		
l_0 (cm)	l_{ten} (mm)	**RMSE**	**FVU**	**Max.**	**RMSE**	**FVU**	**Max.**	**RMSE**	**FVU**	**Max.**
48.5	–	5.779	1.1761	9.798	4.871	0.8354	9.865	1.635	0.0991	3.842
41.0	–	6.633	3.2313	11.409	6.074	2.7090	9.262	2.187	0.1184	4.243
45.7	–	5.315	0.9408	9.991	4.465	0.6640	9.536	1.246	0.0409	2.275
40.6	35	3.945	0.6750	7.527	3.226	0.4512	7.376	2.085	0.0755	3.073

The initial identification, fitted on the $l_0 = 48.5\,\text{cm}$ test, yielded a length offset $\chi_0 = 11.9\,\text{mm}$ and effective stiffnesses $\chi_1 = 102\,\text{kN m}^{-1}$ and $\chi_2 = 10.6\,\text{kN m}^{-1}$. The corrected model at these terms followed the measured torque across the range of motion, including the loss of torque near full extension that the rigid-body prediction did not produce, and Figure 4.7 shows this corrected-family behavior at the adopted two-stage terms. The $l_0 = 45.7\,\text{cm}$ actuator, configured to operate over a different joint range of motion, served as the validation test of the initial fit; the corrected prediction reproduced the measured torque there without refitting. At these initial terms, the corrected prediction for the training test had a root-mean-square error (RMSE) of 1.616 N m, a fraction of variance unexplained (FVU) of 0.0958, and a maximum absolute residual of 3.2527 N m; validation on the held-out actuator gave 2.1214 N m, 0.0878, and 5.1690 N m, respectively (Table C11).

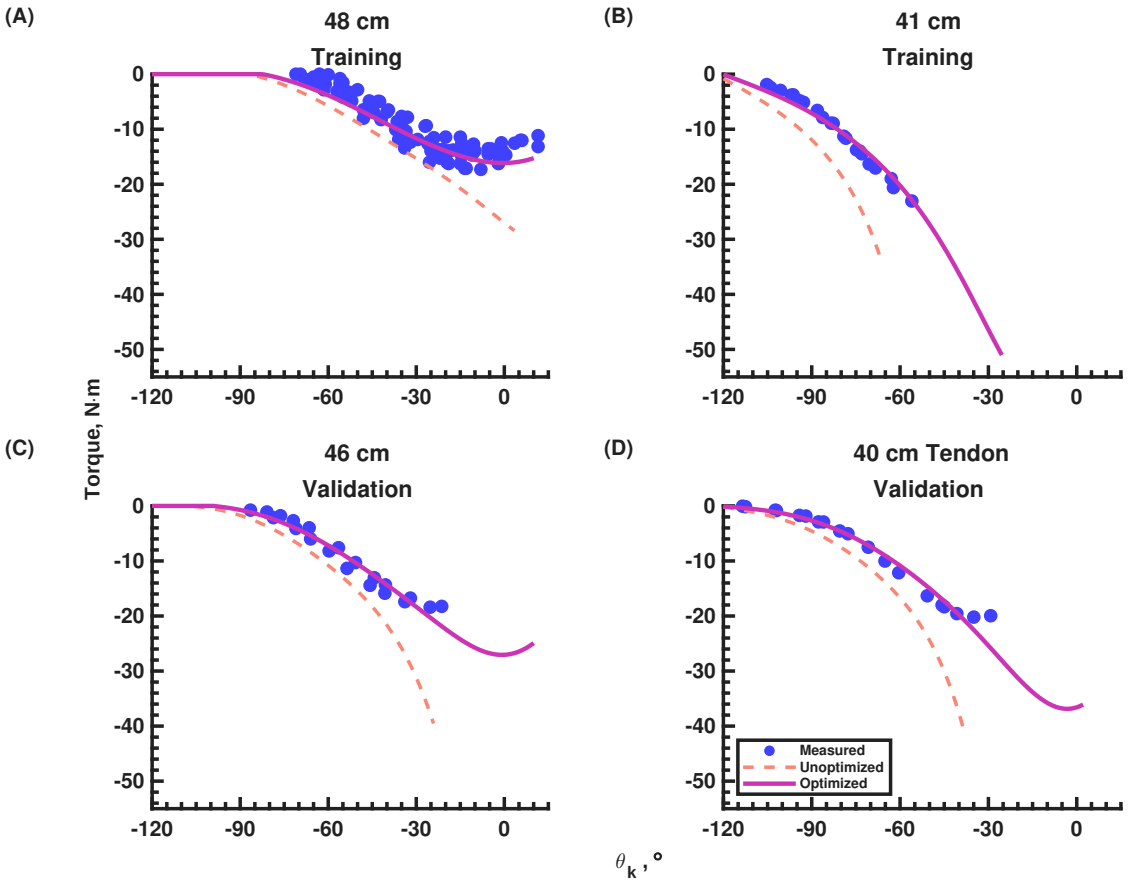


Figure 4.7: Optimized pinned-knee torque versus knee angle for the four cross-validation flexor tests, at the adopted correction terms (Table 4.4). Knee angle is negative in flexion, and flexion torque is negative. Blue filled circles are experimentally measured torque; the dashed light orange line is the unoptimized prediction of the improved BPA characterization; the magenta line is the optimized prediction of the corrected model. Panel titles give the nominal resting length; the actual resting lengths are $l_0 = 48.5\,\text{cm}$ **(A)**, $l_0 = 41.0\,\text{cm}$ **(B)**, $l_0 = 45.7\,\text{cm}$ **(C)**, and $l_0 = 40.6\,\text{cm}$ with tendon **(D)**. Panels **(A)** and **(B)** are the training tests and panels **(C)** and **(D)** the validation tests of the adopted cross-validation fold; the goodness-of-fit values are given in Table 4.2.

Table 4.2: Pinned-knee flexor torque-prediction goodness of fit: the improved BPA characterization alone (the zero-correction prediction of the text) and the corrected model at the adopted correction terms (Table 4.4); training and validation tests follow the adopted cross-validation fold. RMSE and maximum absolute residual (Max.) are in N m, and FVU denotes the fraction of variance unexplained.

		Improved BPA char. alone			**Corrected model**		
l_0 (cm)	l_{ten} (mm)	**RMSE**	**FVU**	**Max.**	**RMSE**	**FVU**	**Max.**
Training							
48.5	–	7.617	2.1278	25.531	2.165	0.1718	5.467
41.0	–	9.694	2.1149	21.076	1.084	0.0277	1.766
Mean		8.656	2.1213	23.303	1.624	0.0998	3.617
Validation							
45.7	–	8.803	1.9380	26.264	1.368	0.0468	3.614
40.6	35	6.373	0.5935	16.859	1.611	0.0427	5.891
Mean		7.588	1.2657	21.562	1.490	0.0447	4.752

The 47.9 cm flexor test was not used in identification and is evaluated separately in Appendix C. Table 4.2 completes this comparison at the adopted correction terms of the full re-identification (Section 4.1.4). Across the four tests of the two-bracket identification, the corrected model reduces the training-mean RMSE, over the $l_0 = 48.5$ cm and 41.0 cm training tests, from $8.66\,\text{N}\cdot\text{m}$ under the zero-correction prediction, the same kinematics and force model with every correction term disabled, to $1.62\,\text{N}\cdot\text{m}$, and the validation

mean over the held-out $l_0 = 45.7\,\text{cm}$ and $40.6\,\text{cm}$ tendon tests improves from $7.59\,\text{N}\cdot\text{m}$ to $1.49\,\text{N}\cdot\text{m}$. The behavior of the original rigid-body torque model on these tests is shown graphically in Fig. 4.6.

4.2.2 Extensor on Pin Joint

When the new stiffness and length offset model was used to predict torque for the knee extensor, measured torque was slightly higher than the original rigid-body model predicts for the 48.0 cm resting length BPA, slightly lower than expected for the 45.7 cm BPA, and much lower for the 41.5 cm BPA. Only the χ_1 and χ_2 terms from the previous optimization results were used with the extensor configuration. By inspection, the zero point of the measured torque was shifted to the right of the original model, indicating an offset opposite in sign to the χ_0 term in the previous section. The BPA with $l_0 = 45.7\,\text{cm}$ illustrates the pattern: the hybrid-method calculated torque correlates well with the originally predicted torque, while the actual measured torque diverges.

Results are shown in Fig. 4.8. The initial identification was fitted in cross-validated rounds over the 48.0 cm, 45.7 cm, and 41.5 cm actuators, holding each actuator out in turn, and the quoted round is the one holding out $l_0 = 41.5\,\text{cm}$; it resulted in $\chi_0 = -10.9\,\text{mm}$ and $\chi_3 = 0.123$, the extensor length offset χ_0 being nearly equal in magnitude to the flexor offset but opposite in sign. The goodness-of-fit values of the initial round are found in Table C11. The two-stage extensor refit of Section 4.1.4 used a different six-test pool (Table 4.3) drawn from the same test campaign, in which the 41.5 cm and 48.0 cm tests of the initial set were excluded from the pool and the 45.7 cm test served as a training test of

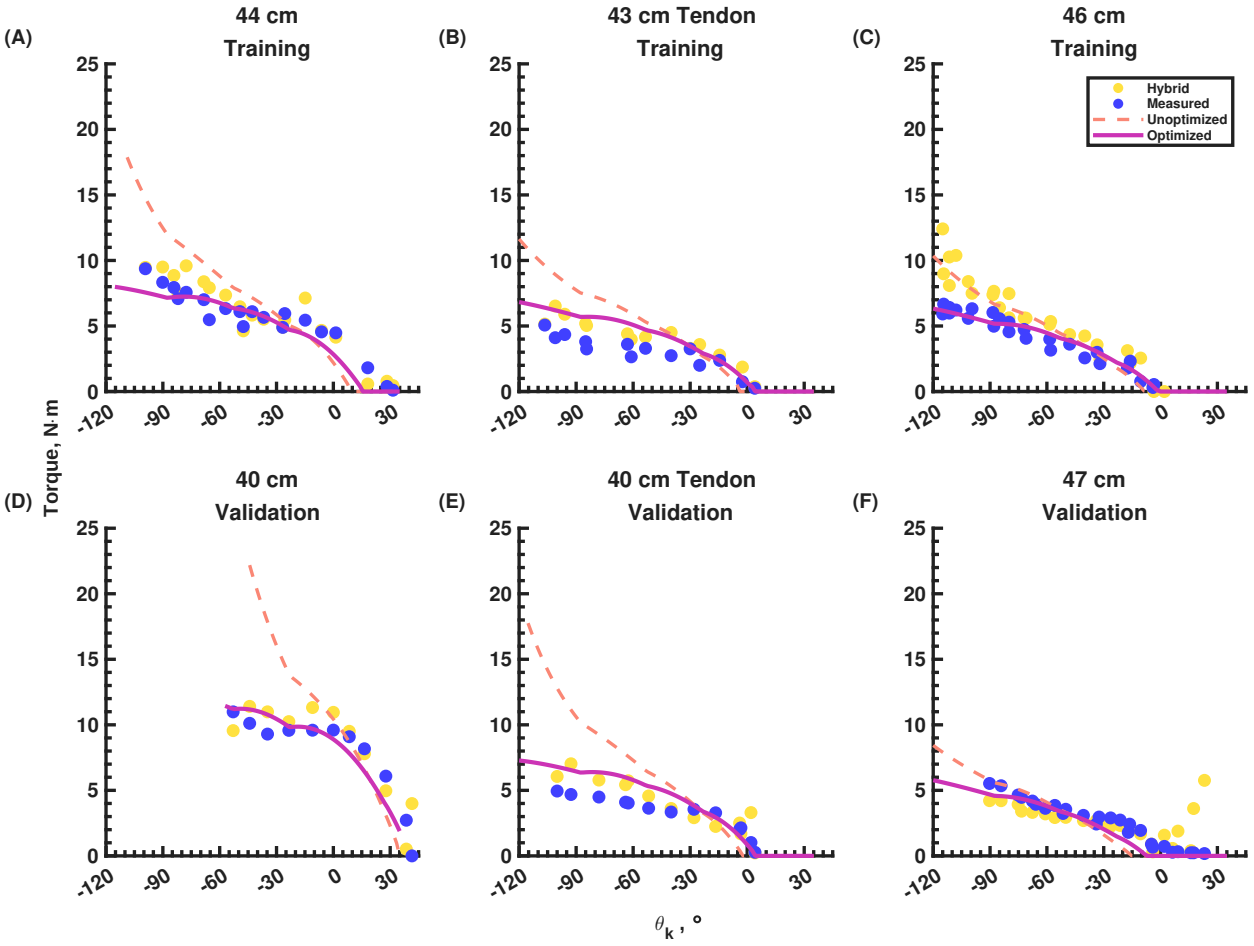


Figure 4.8: Pinned-knee torque versus knee angle for the extensor BPAs: the six identification-pool tests at the adopted correction terms, with the extensor-specific terms (χ_0, χ_3) on the shared stiffness pair (Table 4.4). Knee angle is negative in flexion, and extension torque is positive. Gold circles are the hybrid calculation; indigo filled circles are experimentally measured torque; the dashed light orange line is the unoptimized prediction of the improved BPA characterization; the magenta line is the optimized prediction of the corrected model. Panel titles give the nominal resting length; the actual resting lengths are $l_0 = 43.6\,\text{cm}$ **(A)**, $l_0 = 43.2\,\text{cm}$ with tendon **(B)**, $l_0 = 45.7\,\text{cm}$ **(C)**, $l_0 = 40.0\,\text{cm}$ **(D)**, $l_0 = 40.0\,\text{cm}$ with tendon **(E)**, and $l_0 = 46.5\,\text{cm}$ **(F)**. Panels **(A)–(C)** are the training tests and panels **(D)–(F)** the held-out validation tests of the adopted cross-validation fold; the legend is given in panel **(A)**.

the adopted fold; Fig. 4.8 shows the six pool tests at the adopted terms.

At the adopted correction terms of the full re-identification (Section 4.1.4), Table 4.3 reports the per-test goodness of fit on the six extensor tests used by the identification, three for training and three for validation: the corrected model improves the RMSE of every used test over the zero-correction prediction and holds FVU below unity on all but the $l_0 = 43.2\,\mathrm{cm}$ tendon test. The original rigid-body model on these tests is shown graphically in Fig. 4.8.

Table 4.3 gives the per-test extensor fit, including training and validation means. The earlier cross-configuration fit is retained in Appendix C.8.

Table 4.3: Pinned-knee extensor torque-prediction goodness of fit, at the adopted correction terms with the effective-stiffness pair locked to the flexor-identified values (Table 4.4): the improved BPA characterization alone (the zero-correction prediction of the text) and the corrected model. RMSE and maximum absolute residual (Max.) are in N m, and FVU denotes the fraction of variance unexplained.

		Improved BPA char. alone			**Corrected model**		
l_0 (cm)	l_{ten} (mm)	**RMSE**	**FVU**	**Max.**	**RMSE**	**FVU**	**Max.**
Training							
43.6	–	2.437	0.9618	4.284	0.978	0.1634	1.818
43.2	22	2.878	4.7304	4.817	1.525	1.3270	2.481
45.7	–	2.178	1.0586	4.974	0.490	0.0535	1.206
Mean		2.498	2.2503	4.692	0.997	0.5146	1.835
Validation							
40.0	–	3.092	0.5638	4.842	1.244	0.1303	2.196
40.0	40	3.622	6.3919	7.845	1.350	0.8878	1.890
46.5	–	0.949	0.3152	2.272	0.699	0.1714	1.559
Mean		2.554	2.4236	4.986	1.098	0.3965	1.882

4.2.3 Extensor on Biomimetic Knee

The biomimetic knee was tested with a ∅10 mm BPA in the extensor configuration. The bracket frame origin is the centroid of the cantilevered segment that holds the muscle origin bracket to the proximal femur (see Fig. 4.4**C**, right side).

The torque results for a $l_0 = 51.8$ cm BPA are shown in Fig. 4.9. The torque values measured are significantly smaller than both the original prediction and back-calculated

torque using the hybrid method. The results from the previous extensor configuration lead to a better torque prediction. GoF values can be found in table C11. Fig.4.10 shows the difficulty of modeling a BPA path that passes through the femoral condyles. Using the optimization results for extensor BPAs on the pinned knee joint, the $l_0 = 51.8\,\text{cm}$ BPA has an RMSE $= 0.8920$, FVU $= 0.3396$, and a maximum absolute residual of 1.4188.

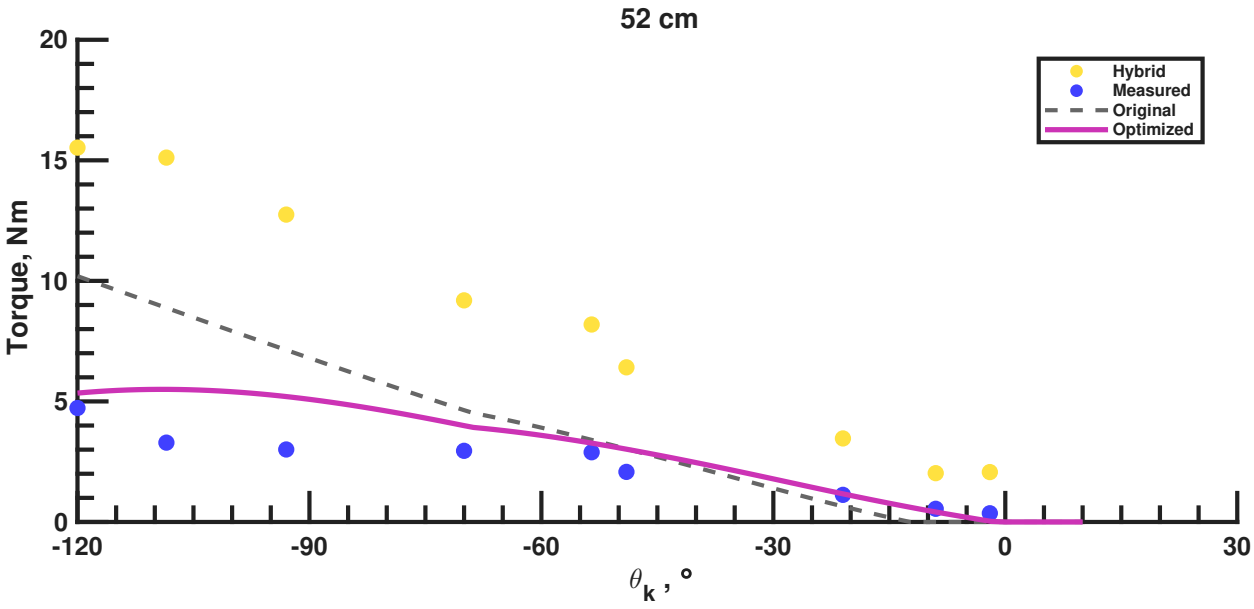


Figure 4.9: Measured versus expected isometric torque results using a ∅10 mm, $l_0 = 51.8\,\text{cm}$ extensor BPA on the biomimetic knee. Measured results are blue dots. Original prediction is a dashed black line. Optimized prediction is a solid magenta line. Hybrid calculation shown in yellow dots.

4.2.4 Flexor on Biomimetic Knee

Fig. 4.11 shows the biomimetic-knee flexor measurements with the flexor terms carried over unchanged from the pinned-knee identification; the bracket frame for the ∅10 mm

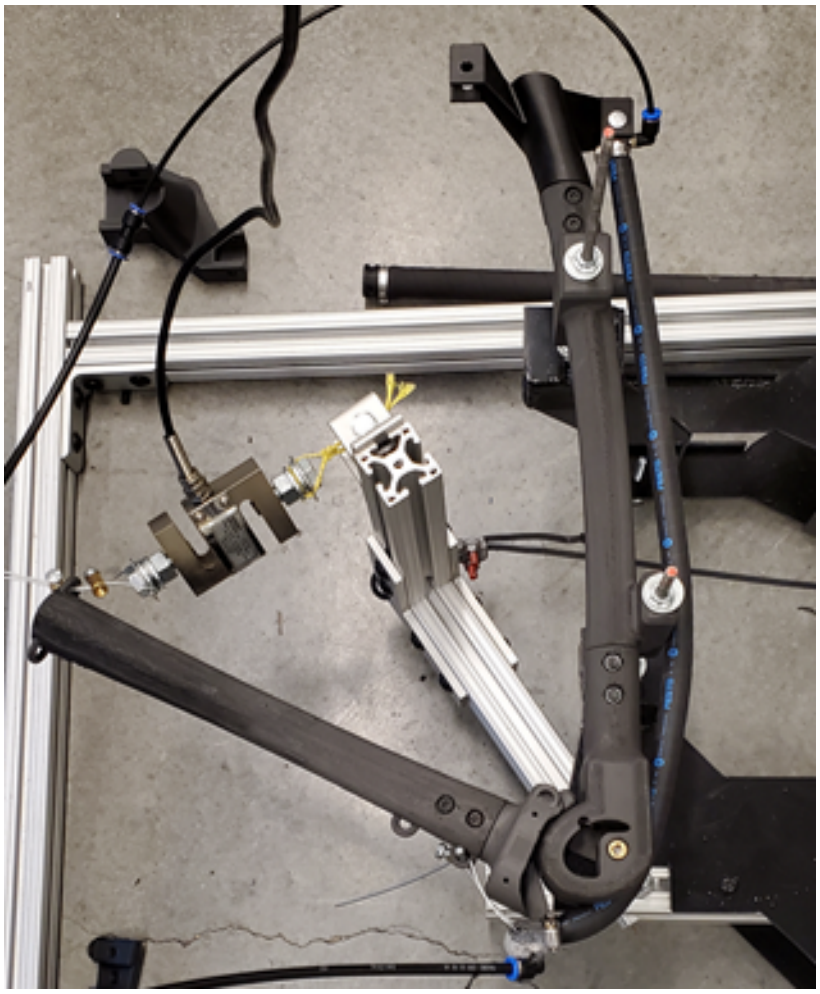


Figure 4.10: A configuration that is particularly hard to model. The BPA passes between the femoral condyles close to ICR.

BPA is shown in Fig. 4.4**C**, and for the ∅20 mm BPA in Fig. 4.4**D**. Panel **A** presents the ∅10 mm actuator at $l_0 = 41.5$ cm with a 21 mm artificial tendon (∅1.5 mm Shimano bicycle brake cable wrapped four times), pressurized to 604 kPa; the corrected model reached an RMSE of 0.3007 N m, an FVU of 0.0102, and a maximum absolute residual of 0.4647 N m. Panel **B** presents the ∅20 mm actuator at $l_0 = 41.5$ cm with a 14 mm artificial tendon made out of the same brake cable, wrapped three times so that six cable segments share the load, each wrap contributing two tendon lengths and cross-sectional areas. Panel **B** was tested at 620 kPa, for which the corrected model reached an RMSE of 1.8024 N m, an FVU of 0.0499, and a maximum absolute residual of 3.5678 N m.

The torque of the single ∅20 mm flexor actuator of panel **B** did not exceed the magnitude of the maximum human isometric flexion torque between 60° and 100° of knee flexion, and within a ±20% band about the human mean it still fell short of human-level torque between 78° and 103° of flexion. Because operating at 620 kPa across the full desired range of motion would load the printed femur beyond its ultimate strength, single points were also recorded at the reduced pressures 560 kPa, 421 kPa, 325 kPa and 281 kPa, which show that the desired range of motion is achievable at reduced pressure. This shortfall in flexion torque defines the target that the attachment-route redesign of Section 4.3.2 ...addresses. Panel **C** shows the outcome for the redesigned two-BPA flexor route of Section 4.3.4: across its fifteen tested measurements from −119° to 5°, the measured torque magnitude met or exceeded the human biceps-femoris short-head maximum at every point, with margins from 0.42 N m at −105° to 17.0 N m at −14°.

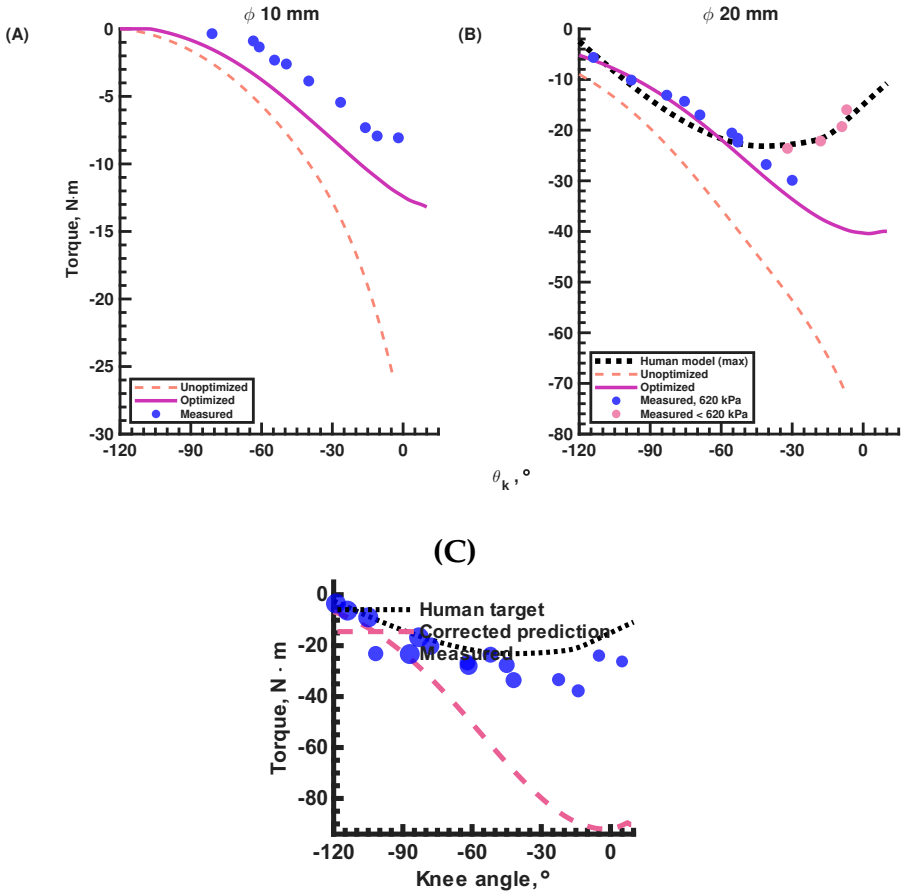


Figure 4.11: Biomimetic-knee torque versus knee angle for the flexor BPAs. Knee angle is negative in flexion, and flexion torque is negative. **(A)** ∅10 mm, $l_0 = 41.5$ cm actuator with a 21 mm artificial tendon at 604 kPa. **(B)** ∅20 mm, $l_0 = 41.5$ cm actuator with a 14 mm artificial tendon at 620 kPa. Indigo filled circles are measured torque; the magenta line is the optimized prediction of the corrected model using the flexor terms identified on the pinned knee; the dashed light orange line is the unoptimized prediction of the improved BPA characterization alone. In **(B)**, the black dotted line is the maximum human isometric torque, and the pink circles are measured points at the reduced pressures 560 kPa, 421 kPa, 325 kPa and 281 kPa (from left to right, respectively), recorded because full pressure would load the printed femur beyond its ultimate strength. **(C)** The redesigned two-BPA flexor route

4.3 Discussion

4.3.1 Sources of Error in Isometric Torque Generation

Our improved BPA characterization predicts torque more accurately than the original BPA model (see Fig. 4.6). This only accounts for part of the error. Many factors affect the torque results. Compliance in the testing system, positioning tolerances, flexibility of brackets, imperfect path modeling, and variation in individual BPAs are other factors.

Simplifying and testing the model revealed deficiencies in our previous analysis. The isometric system is not rigid. There is compliance in the winch cable, artificial tendons, and brackets that bend and stretch. We introduced a constant length offset term χ_0, axial stiffness χ_1, and bending stiffness χ_2. The bracket frame of reference was placed at the onset of the cantilevered section and rotated about Z. The optimization identified parameter values for a $\varnothing$10 mm BPA with $l_0 = 48.5$ cm in the simple pinned knee configuration. These results were validated with a second $\varnothing$10 mm BPA with $l_0 = 45.7$ cm.

The modeling of extensor BPA paths that wrap around the knee joint is an approximation with inherent error (Fig. 4.10). When the moment arm is short, small changes to it can have a large effect on the resultant torque [**?**]. It was assumed that wrapping the BPA around the joint would be easier to account for than adding a patella, deterministically designing its location as a function of knee angle, and accounting for friction. At greater knee flexion angles, shorter l_0 BPAs were observed to stretch and compress (see Fig. 4.1**C**). The same loading sometimes caused the BPA to slip off the anterior head of the bolt that holds the femur to the knee. This produced a $\pm$20 mm displacement in the Z direction, changing both the force vector and the torque. Fig. 4.8 shows a deviation between mea-

sured and originally predicted results that widens as expected torque increases. This indicates system compliance. However, the hybrid-calculated torque matches or exceeds the predicted torque, except at greater flexion angles for the $l_0 = 41.5\,\text{cm}$ BPA, indicating a loss of functional l_m; this is why the χ_3 term was created in (4.24). Fig. 4.8 shows the improved fit from running the optimization that includes this term.

The biomimetic knee joint with an extensor BPA was used to validate the extensor optimization results. The corrected prediction more closely matches the measured torque values, although it remains higher than the measurements.

The nonlinear biomimetic robotic system presented here replaces rigid-body simplifications commonly used in other robotic and biomechanical systems. Instead, progressing from simple to complex systems elucidated compliance mechanisms that can lead to more accurate first-iteration designs. Technologies have been developed to predict system bending and axial compliance, a constant length offset, and a loss of usable length that occurs as artificial muscle bends around a joint.

The quasi-static force characterization also fixes the coefficients that a dynamic actuator model needs. Writing the actuator's contraction as a second-order system, $m\ddot{\varepsilon}^* + c(\dot{\varepsilon}^*)\dot{\varepsilon}^* + k(\varepsilon^*)\varepsilon^* = F$, the stiffness term is the tangent slope of the normalized force surface of Equation (3.4), $k(\varepsilon^*, P^*) = \frac{\partial F}{\partial \varepsilon^*}\Big|_{P^*} = -c_0 c_1 e^{-c_1\varepsilon^*} - 2c_2\varepsilon^* P^* e^{-c_2(\varepsilon^*)^2}$. Setting the mass and the external force to zero then yields the equivalent viscous damping coefficient, $c(\dot{\varepsilon}^*) = (k(\varepsilon^*)\varepsilon^*)/\dot{\varepsilon}^*$, the term that dynamic state-space actuator models, such as the pressure–force formulation of **?**, pair with this stiffness.

4.3.2 Follow-Up Identification and Route Redesign

The coefficient values reported with the results figures belong to the initial single-bracket identification. Re-identifying the correction terms over the full test sets with the two-stage procedure of Section 4.1.4 yields the adopted values $\chi_0 = 3.94\,\mathrm{mm}$, $\chi_1 = 3.998 \times 10^4\,\mathrm{N\,m^{-1}}$, and $\chi_2 = 1.473 \times 10^4\,\mathrm{N\,m^{-1}}$ for the flexor, and $\chi_0 = -6.38\,\mathrm{mm}$ with $\chi_3 = 0.158$ for the extensor on that same locked stiffness pair (Table 4.4); the search bounds and the adopted values are recorded in Appendix C (Tables C7 and C8).

Table 4.4: Adopted transmission corrections. The stiffness pair is shared. The wrapping factor was identified on the extensor and is also used in the flexor redesign; it was not fitted to the pinned-flexor tests.

Configuration	χ_0 (m)	χ_1 (N/m)	χ_2 (N/m)	χ_3
Flexor	0.00394	3.998×10^4	1.473×10^4	0.158
Extensor	-0.00638	3.998×10^4	1.473×10^4	0.158

What the correction terms do to the prediction is quantified in Tables 4.2 and 4.3. The corrected flexor model holds FVU at or below 0.18 on the four tests of the identification, the worst case being the 48.5 cm training test at 0.172, and the post hoc evaluation of the unused 47.9 cm test improves as well (Appendix C). On the extensor identification the corrected model improves the RMSE on every used test, cutting the training-triple mean from $2.50\,\mathrm{N \cdot m}$ to $1.00\,\mathrm{N \cdot m}$ and the validation-triple mean from $2.55\,\mathrm{N \cdot m}$ to $1.10\,\mathrm{N \cdot m}$, and holds FVU below unity on every used test except the 43.2 cm tendon test. On the biomimetic extensor test ($l_0 = 51.8\,\mathrm{cm}$) the corrected prediction improves RMSE from $3.03\,\mathrm{N \cdot m}$ to $1.17\,\mathrm{N \cdot m}$ and FVU from 4.49 to 0.67, so the shared-pair model explains the bulk of that configuration's measured torque variance as well as its level. The genetic

multiobjective solver is retained for the identification, where the Pareto front is required for holdout validation, and the surrogate-plus-pattern-search pipeline serves the route redesign (Section 4.1.5).

The initial route searches bounded the BPA origins and insertions near the corresponding human muscle attachments to preserve biomimetic routing. Those anatomical bounds had to be relaxed where the available BPA contraction could not cover the knee range of motion. The resulting routes then required a second tradeoff: reducing the actuator forces carried by the attachment hardware after parts failed during testing. The cam redesign changes the moment arm and contact geometry to retain the torque requirement at lower actuator force.

The route redesign built on these values addresses the single-muscle shortfall at high flexion angles documented in Section 4.2: each configuration now carries two 20 mm BPAs per route. The flexor solution meets the flexor torque target with a minimum margin of $+5.01\,\%$ at -107°, with a separately solved route for each actuator. The second insertion mirrors the first across the sagittal plane, while its origin preserves the first route's lateral endpoint offset; its wrapping and clearance are evaluated independently. The direct extensor route meets or exceeds its human-scaled target with the required $5\,\%$ overshoot as the thinnest margin at every angle, $+5.02\,\%$ at worst, while respecting the wrap-release schedule and contraction limits of the route-elimination rule (Appendix C); it was later superseded by the cam redesign of Section 4.3.2. These are mechanical design predictions. Their complete incorporation into the converted neuromechanical models remains future integration work.

Redesigned-Route Results on the Gait2392 Robot-Body Model

This subsection records the quantitative outcomes of the redesigned routes on the Gait2392-derived robot-body model, whose muscle force and torque targets define the redesign objectives. All three result families were produced by the surrogate-plus-pattern-search pipeline of Section 4.1.5 under the adopted correction terms (Appendix C, Table C8) with one shared actuator plan: two 20 mm BPAs per route, a 620 kPa pressure ceiling, and a maximum usable contraction of 0.255 in the force model. Each family minimizes its own objective function, a unitless weighted sum of normalized penalty terms in which zero corresponds to no violated penalty, led by the worst-angle and mean-square torque shortfall against the required 5 %-overshoot human envelope, followed by overshoot, rest-length, strain, contraction, and mild design-change penalties. The three functions differ in penalty families and normalizations (the flexor and pulley objectives add off-axis-excess, wrap-release, and second-BPA anchor terms, and the pulley objective evaluates its torque requirement on the insertion-side torque after the tackle's force loss), so their numerical values are not commensurable with one another. The human-scaled requirements are the OpenSim isometric torque profiles of the reference human muscles over the same design-grid angles, the vastus medialis for the extensor and the biceps femoris short head for the flexor.

The direct extensor route exceeds the human envelope in magnitude at every angle, with zero shortfall points: across the design range the predicted two-BPA extensor torque spans 35.45 N · m to 109.4 N · m against the human-scaled requirement of 20.25 N · m to 50.32 N · m, and the required 5 % overshoot is consumed as the thinnest margin exactly as

designed, at a minimum margin of $+5.02\,\%$ at the worst-case angle of -20.2°.

Allowing the tibial-side route points to trace a cam profile produced the current extensor design (Figure 4.12). The final cam geometry includes the optimized route and the verified adjustment of the distal route point and tendon length. Its minimum torque margin is $+30.83\,\%$ with respect to the required envelope, the human target plus the 5 % overshoot, and $+39.77\,\%$ with respect to the plain human target. The pressure required to meet the envelope spans 269 kPa–507 kPa across the range of motion, comfortably above the 150 kPa regulated floor of the VTUG valve manifold that supplies the hardware. The geometry confirms that the headroom is not bought with routing risk: the tibial wrap follows a bend radius of 96.9 mm, the guiding route point clears the femoral-condyle envelope by 43.1 mm at 10° of knee extension, and the moment arm spans 59.4 mm–102.8 mm (mean 72.4 mm) over the range of motion.

The flexor route, built from two separately evaluated BPA paths, meets the flexor torque target with a minimum margin fraction of 0.0501 against the required 0.05. Its predicted knee-axis torque spans $-81.23\,\mathrm{N\cdot m}$ to $-3.864\,\mathrm{N\cdot m}$ against the human-scaled requirement of $-23.16\,\mathrm{N\cdot m}$ to $-2.461\,\mathrm{N\cdot m}$ (peak magnitude $23.16\,\mathrm{N\cdot m}$), again exceeding the human envelope at every angle. The flexor redesign applies the adopted correction terms of Table C8, including the wrapping factor identified on the extensor configuration.

The Mono/Bi-pulley algorithms extend the route search space by allowing the artificial tendon to be reeved over pulleys before insertion. On the knee-flexor route, the pulley variant converged to a collision-feasible layout with zero bound violations and a peak predicted flexion torque of $85.56\,\mathrm{N\cdot m}$ on the insertion-side measure, the torque applied at the insertion after the tackle's force loss; Section 4.3.3 quantifies what the redirect adds on

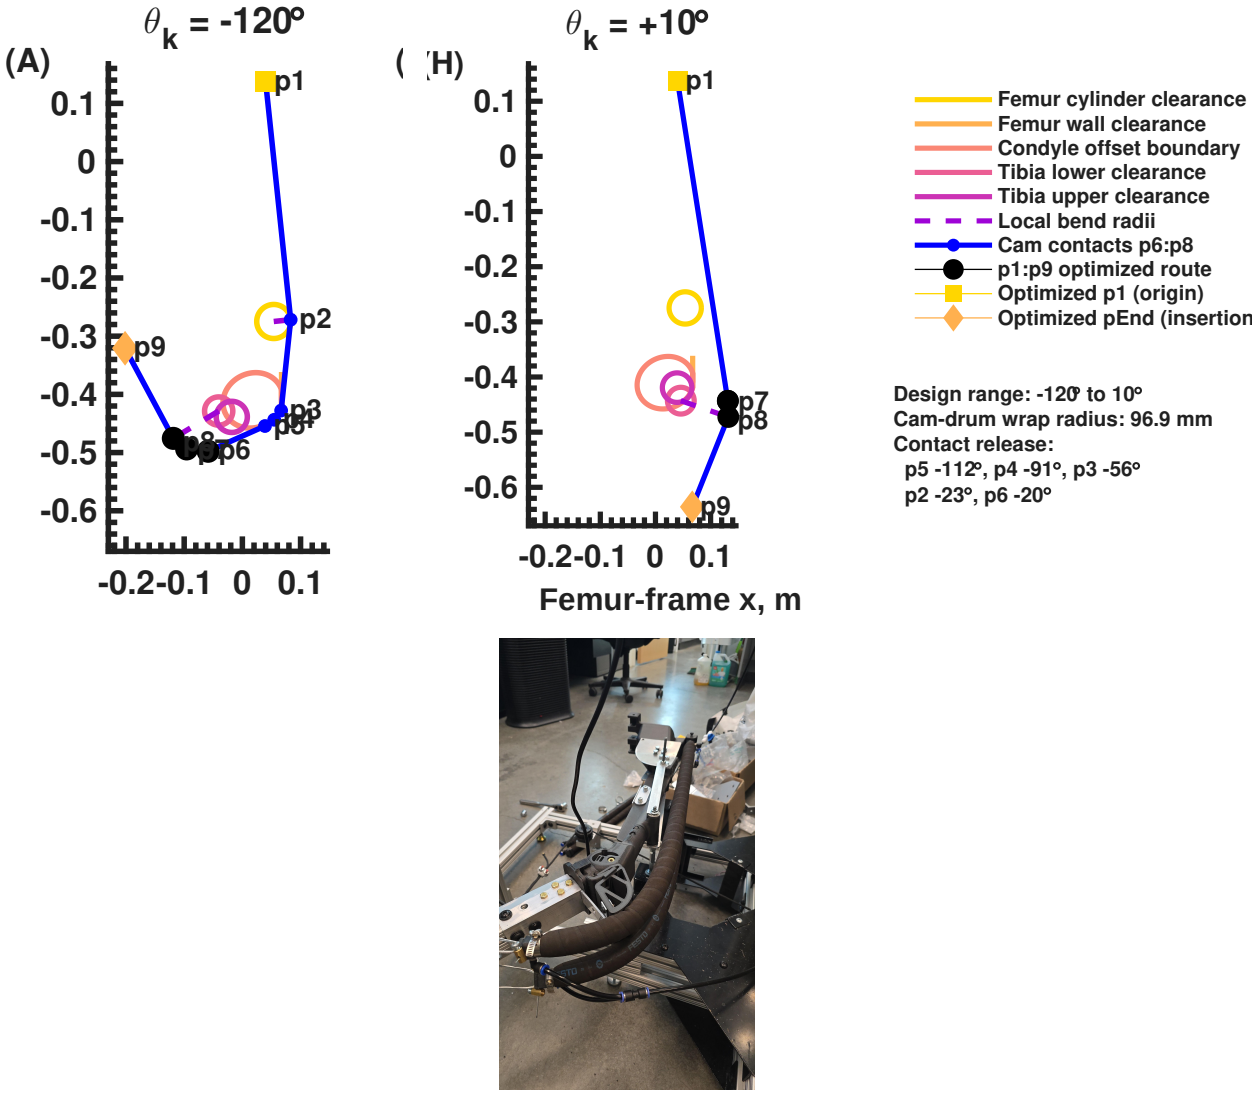


Figure 4.12: Cam-extensor route (top) at -120° flexion (A) and 10° extension (H), with the original clearance and contact legend. Squares mark the femoral origin, diamonds the tibial insertion, and blue lines the actuator route. The photograph (bottom) shows the assembled two-BPA hardware. Panel letters follow the full sequence in Figures C5 and C6.

the knee-axis measure and the off-axis torque it bounds.

4.3.3 Mono- and Bi-Pulley Tendon Routing

At unity transmission, the knee-flexor pulley route increased peak predicted knee-axis torque from 81.23 N m to 85.74 N m, while limiting off-axis torque to 0.803 N m. This comparison measures the effect of the two optimized routing configurations; their different objective functions do not isolate the pulley alone. The pulley serves as a redirect in this configuration rather than a force multiplier. Appendix C.7 records the optimizer diagnostics and the broader route searches.

A subsequent transmission study re-solved actuator sizes for each tackle configuration at five anatomical muscle sites. A two-BPA pair through a two-part tackle preserves net actuator force while doubling insertion travel; increasing travel can extend the angle range over which a route meets its torque requirement. The corresponding torque and margin surfaces are shown in Figures C3 and C4. These calculations use the moving-force-point routing, with the adjacent joint fixed for the biarticular evaluation planes; they do not establish simultaneous performance across both joints. Physical verification of the pulley configurations remains future work.

4.3.4 Physical Tests of the Redesigned Routes

The redesigned cam-extensor route was tested at nine isometric holds, and the two-BPA flexor route at fifteen holds. Both test sets span knee angles from -119° to 5°. The test-specific analysis scripts transform the measured load-cell wrench to the moving knee

center and evaluate model predictions at each measured angle and pressure. Figures 4.13 and 4.14 compare these reductions with the human-reference torques and the optimized 620 kPa curves.

The measured torque magnitudes reach or exceed the human-reference magnitudes over most of the tested range. Agreement with the pressure-specific prediction is uneven, particularly for the flexor at intermediate angles. The tests therefore support the redesigned routes' torque capacity while identifying prediction errors that remain to be explained. Since pressure varied between holds, the measured points do not constitute a full-pressure envelope at every angle.

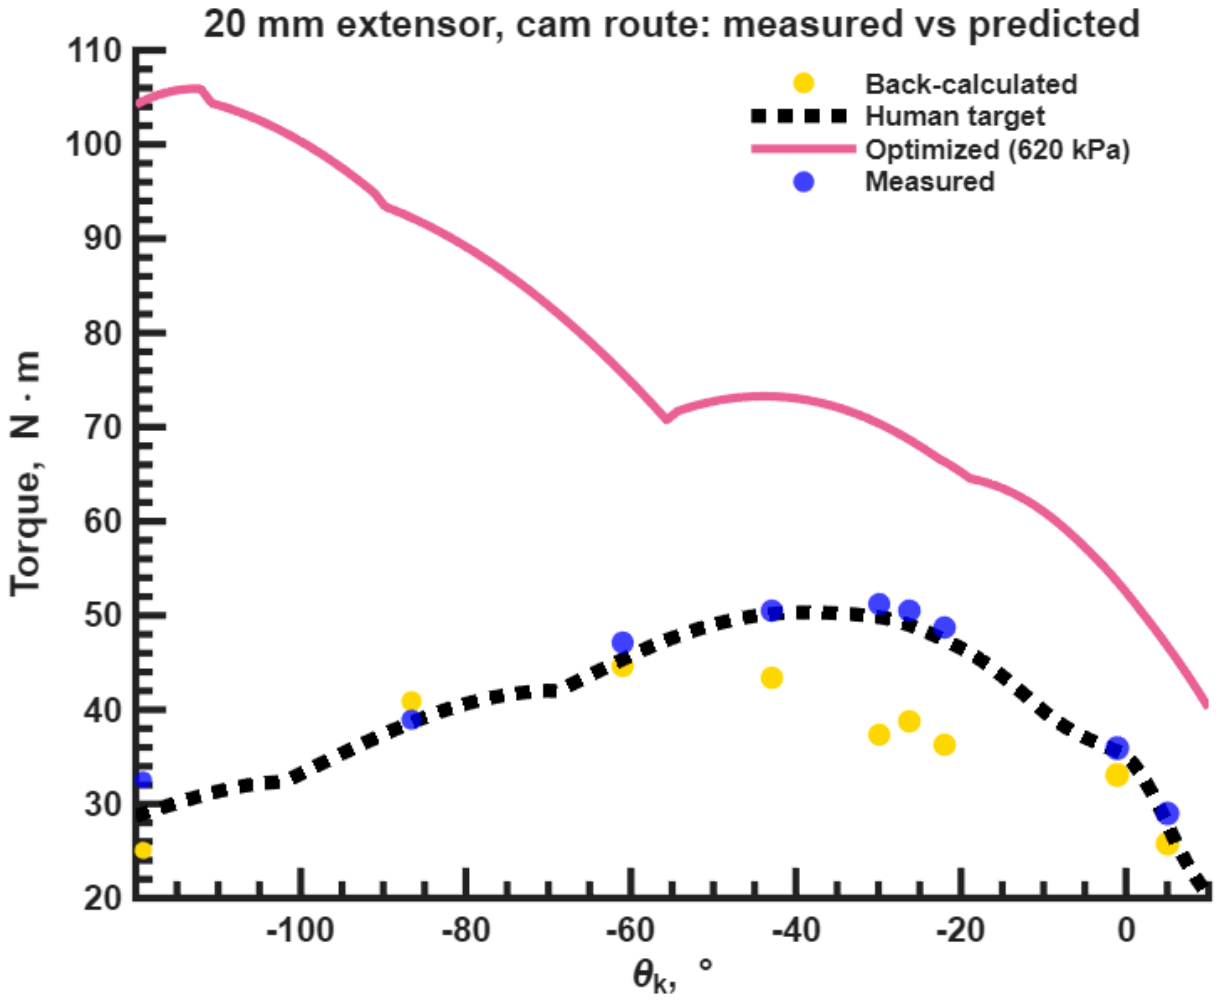


Figure 4.13: Cam-extensor tests. Blue points are measured torques; yellow points are model predictions at the measured angles and pressures. The pink curve is the optimized 620 kPa prediction, and the black dotted curve is the human target. Point area scales with pressure.

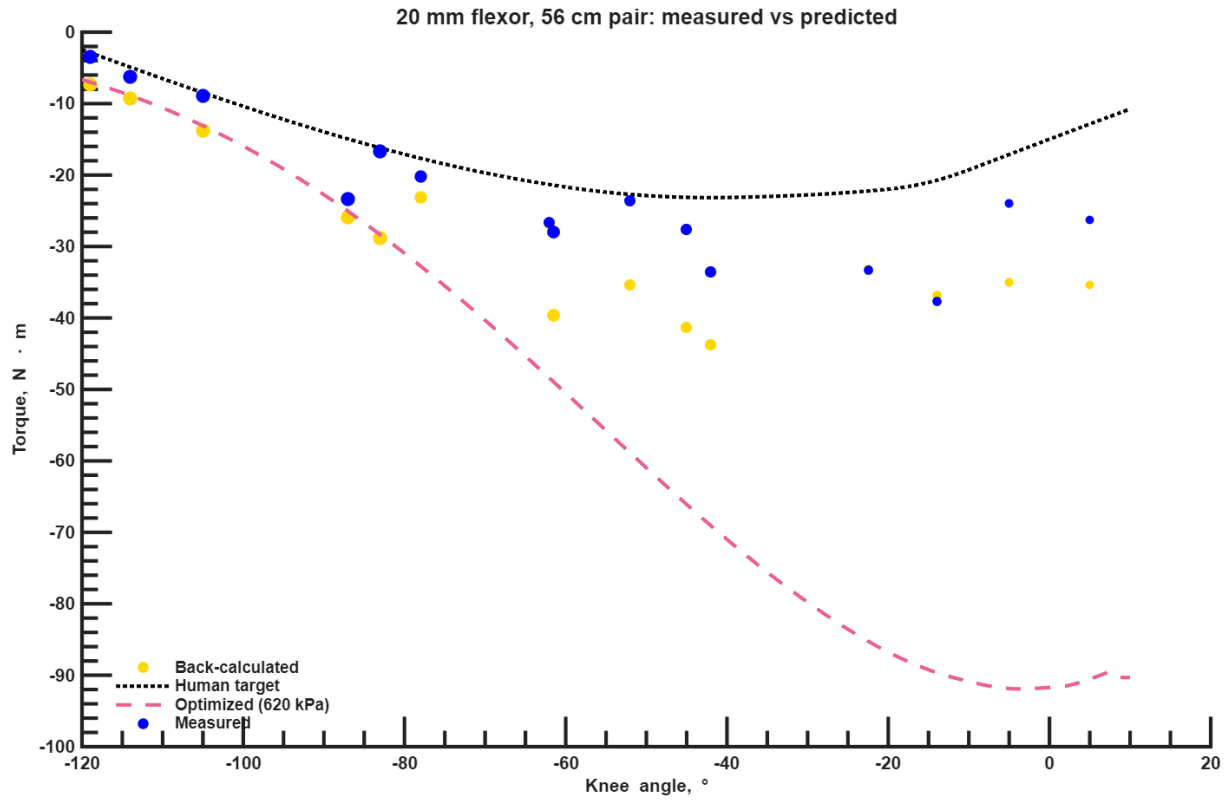


Figure 4.14: Two-BPA flexor tests with 56 cm actuators. Blue points are measured torques; yellow points are predictions at the measured angles and pressures. The dashed pink curve is the optimized 620 kPa prediction, and the black dotted curve is the human target. Point area scales with pressure.

5 The Sensory Afferent Database

5.1 The Resource

The Sensory Afferent Database (SADb) is a curated literature resource for studying sensory feedback in locomotion and posture. It connects papers to afferent classes, experimental preparations, feedback pathways, and computational models. I developed it within the AARL's NeuroNex Communication, Coordination, and Control in Nervous Systems (C3NS) effort and presented the resource at Neuroscience 2023 [**?**].

When I began this work, the project lacked a structured resource linking measured afferent properties and feedback pathways to the evidence needed for robot-controller design. Relevant findings were dispersed across species, preparations, tasks, and primary papers. Reviews organized parts of this evidence [**???**], but selecting a feedback rule still required recovering its experimental conditions and separating measurements from modeling assumptions. I built the SADb to connect those findings to their sources, afferent classes, and pathway roles so that controller choices could be traced to the evidence and remaining unknowns could be identified.

The primary deliverable is a browser application for searching paper records, counting them by selected classifications, and exploring citation neighborhoods (Figure 5.1). Its embedded data support offline use; its separate online-search tab requires a network connection. The distributed snapshot contains 1,393 records: 943 papers and 450 linked model, review, feedback-pathway, and recorder records. Records removed from the working

Airtable base remain available in the application as archived records. This separates the collection used for ongoing curation from the broader set of records considered during its assembly. The pivot and citation views operate on papers; the table view searches across record types. For a selected paper, references are the works it cites, whereas citing papers are works that cite it. The citation view distinguishes these directions.

The working base is contained in Airtable, an online database program. Its linked tables organize Papers, Feedback Pathways, Models, and Review Papers. A paper record contains bibliographic information and curation notes: short summaries of its reported mechanism, experimental conditions, and relevance to feedback design. A feedback-pathway record describes a sensory influence on motor control, such as Ib feedback influencing the stance-to-swing transition, and links that description to supporting papers. A model record connects a computational implementation to its source literature; a review record summarizes literature relevant to the collection. Thus, a pathway identifies a mechanism, the linked papers supply its experimental support, and model records identify implementations to examine. These links organize evidence; they do not establish that every relationship has been entered or verified. Paper matching uses DOI where available and normalized titles otherwise. The paper-table view is shown in Appendix E (Figure E1).

5.2 Curation Methods

This was a curated research collection, rather than a registered systematic review or an exhaustive search of every relevant publication. The collection began with the libraries used to provide background and understanding for the previous publications on locomotion

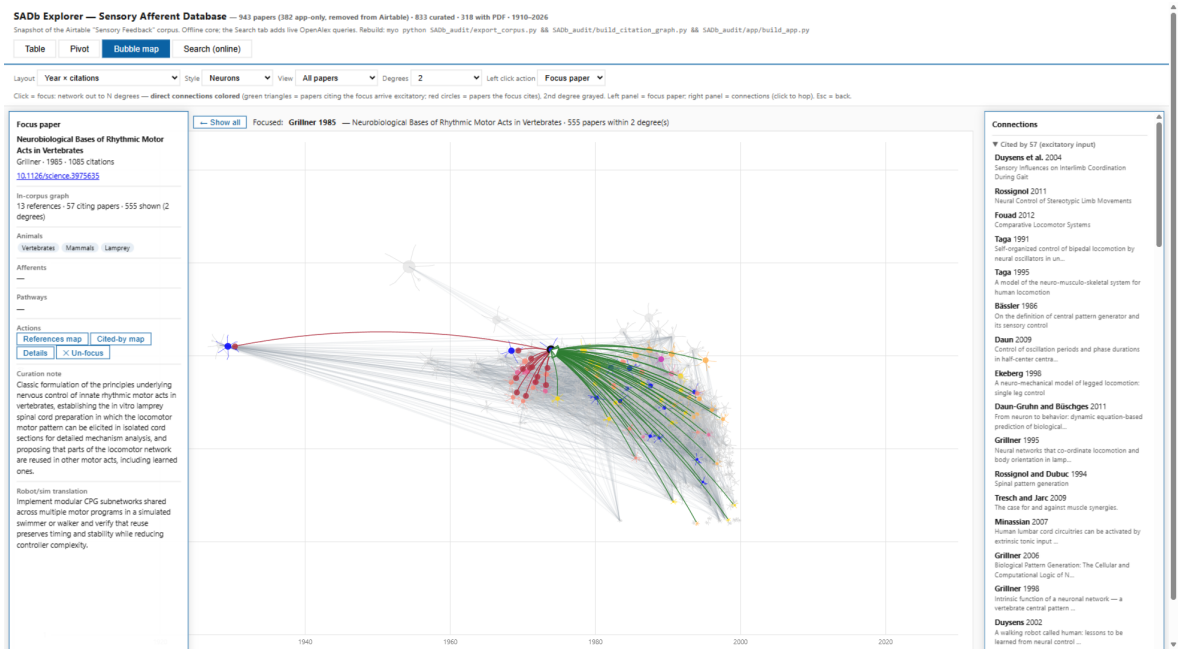


Figure 5.1: Paper-only citation view focused on Grillner (1985). Soma size represents citation count; axes show year and citation count. Triangles mark 57 papers in the corpus that cite Grillner (1985); circles mark 13 papers cited by Grillner (1985). Gray papers connect through one of those neighbors rather than directly to the selected paper. The panels distinguish incoming citations from the selected paper's references.

and postural balance control in the AARL lab. Records were reconciled by DOI and title, then expanded through citation searches. Reviews of spinal and proprioceptive feedback supplied an initial organization of the literature [**???**].

Records were retained when they reported measurements from sensory afferents, investigated feedback in locomotion or posture, or modeled a sensory pathway relevant to those tasks. Here, a sensory pathway means the route by which a sensory signal influences neural activity or motor output, such as force feedback from a muscle influencing the transition from stance to swing. Reviews helped organize the subject and locate primary evidence. Duplicate records and material outside this scope were archived. Classification used a consistent set of tags for the experimental preparation and afferent class, allowing papers that used different terminology to be retrieved together. Each mechanism summary was based on the available source text, using the abstract when full text was unavailable.

5.2.1 AI-Assisted Triage

I evaluated AI-assisted literature identification, summarization, and categorization in work presented at Neuroscience 2024 [**?**]. The comparison included ResearchRabbit, Litmaps, txyz, SciSpace, and MirrorThink. Some tools helped identify related papers or summarize straightforward material, but categorization remained inadequate: none received a good rating, and only txyz was rated adequate. Complex models and nuanced physiological findings required substantial human correction. The tools therefore offered limited help with the central curation task of accurately assigning papers to database categories. I retained responsibility for inclusion, classification, and deletion decisions after checking

the source.

5.3 Coverage and Use

The audited paper set contains 943 records; 561 were retained in the working base. Of those retained papers, 523 have curation notes, 426 have preparation tags, 86 have an afferent classification, and 258 have notes relating their findings to robot or simulation use. These counts describe curation coverage, not the number of independent physiological findings or a measure of their reliability. The more detailed record counts and linked pathway examples are given in Appendix E.

A user of the database can begin with a preparation or afferent class, retrieve the matching papers, and inspect their notes and source links. The pivot view counts papers along two selected dimensions (Figure 5.2); the citation view provides a second route from an anchor paper to related literature (Figure 5.1). These views help locate evidence, while the original experiment determines the conditions under which a finding applies.

5.4 Discussion

5.4.1 Use in Neural-Controller Design

The resource helped organize the sensory channels used in Chapter 6: spindle-related length and velocity signals, force feedback, and contact events. The linked records provide sources for selecting a pathway and examining its preparation, sign, and phase dependence. They are not ready-made controller parameters. Translating an experimental result into a

Afferent types ↓ / Animal →	(untagged)	Cat	Human	Mice	Mammals	Insects	Vertebrates	Rat
(untagged)	327	162	202	64	51	40	31	29
Cutaneous	5	28	6	1	2	·	1	·
Ia	4	29	5	2	·	·	1	1
Ib	5	21	7	·	3	·	·	1
II	2	16	6	·	1	·	1	1
Mechanosensory	·	·	·	·	·	2	·	·
Flexor reflex afferents	1	4	·	·	·	·	·	·
III/IV	·	1	·	·	·	·	1	·
Nociceptive	·	1	·	·	·	·	·	·

Figure 5.2: Excerpt of the application pivot matrix, with afferent classifications as rows and selected experimental preparations as columns. Cells count tagged papers; selecting a cell retrieves its records. Multiple tags allow a paper to contribute to more than one cell.

robot circuit still requires decisions about sensor encoding, species and task differences, and gain selection, followed by verification of the implemented behavior.

5.4.2 Further improvement and work

The present corpus is uneven. Its afferent classifications cover only part of the retained literature, and the depth of the mammalian locomotor material should not be taken as equal coverage of vestibular, visual, or other sensory systems. The application makes the curated papers easier to inspect and compare; additional curation and experimental interpretation remain necessary for extending the robot's sensory system. Further work should check and complete the links between feedback pathways and their supporting papers, and between reviews, papers, and pathways. Each link needs to identify the evidence for the relationship, so that a reader can distinguish a directly measured sensory

effect from a model assumption or a review summary. The priority is to verify these relationships and fill gaps in the existing collection before treating it as a complete basis for controller design.

6 Synthetic Nervous System Simulation

The actuator and torque models of Chapters 3 and 4 predict how the robot's muscles produce joint motion. Studying neural control requires connecting those mechanics to a synthetic nervous system (SNS). This chapter develops a simulation framework in which neural controllers can drive the Gait2392 musculoskeletal model and separate BPA-actuated models, receive sensory feedback, and be tested under specified support and disturbance conditions. Its purpose is to make controller behavior measurable as neural organization, actuator mechanics, and mechanical support are changed. The contribution is the implemented framework and its demonstrated tests; unsupported walking remains an open control problem.

6.1 The AnimatLab Bipedal Walker

For a course project in ME 610: Neuromechanical Modeling, Morrow and I combined the bipedal mechanics of **?** with the two-level central pattern generator (CPG) of **?** in AnimatLab [**?**]. The mechanical model had a pelvis and bilateral thigh, shank, foot, and toe segments, with sagittal-plane joints and antagonist Hill-type muscles at the hip, knee, and ankle. I added a passive toe joint. The initial controller used one rhythm generator (RG) to control both sides, with separate pattern-formation (PF) groups for each leg: one for the hip and one for the knee and ankle. The shared RG layer contained four neurons. We added afferent feedback at the motoneuron (MN) level. Figure 6.1 shows the model and controller organization.

We adjusted applied currents, synaptic weights, and muscle parameters by hand. The suspended model developed alternating muscle activity and stepping-like motion, presented at the 2021 NIH BRAIN Initiative Investigators Meeting [**?**]. It did not achieve weight-bearing walking. The later bilateral-RG extension was a separate development; the suspended result does not establish its weight-bearing performance.

I subsequently expanded the neural model into left- and right-side circuits with interconnected rhythm generators (RGs). I constructed the right-side subsystems from the left-side organization and linked the two RGs through commissural pathways. This later model is referred to as the Walker Two-Layer (W2L) model. Further circuit schematics and implementation records are provided in Appendix D.

The suspended motion left an important question unresolved: could the controller maintain coordinated activity when the model contacted the ground and had to support load? Answering it required control over the amount and direction of mechanical support, access to sensory and motor signals, and repeatable changes to the controller or body. Those requirements motivated the framework developed here. The Li-derived W2L walker, the Gait2392 human model, and the BPA-routed robot model are distinct mechanical models; a result obtained with one is identified accordingly.

6.2 Knowledge Gaps and the New Tools

The project lacked a tested way to connect its neural controllers to the human reference and artificial-muscle models and evaluate them under comparable conditions. The missing capability was more specific than generating rhythmic joint motion. I needed to choose

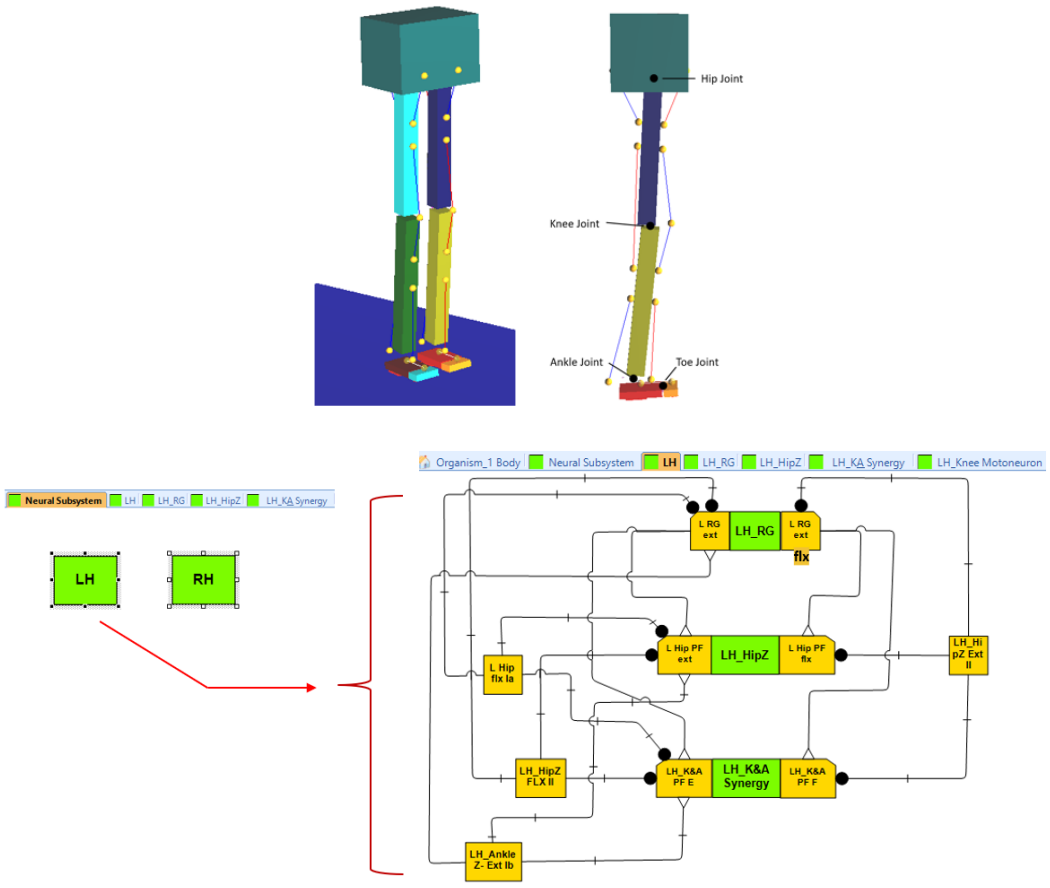


Figure 6.1: AnimatLab walker (top) and bilateral controller hierarchy (bottom). The model combines the Li walker mechanics and Deng CPG organization. These diagrams show construction; the suspended stepping result is described in the text.

a controller, define the mechanical model and sensory inputs, specify how the body was supported, and measure what changed after a controlled intervention. Without those controls, an improvement in motion could not be assigned to neural feedback, muscle mechanics, or external support.

I addressed this gap by connecting OpenSim-derived models to MuJoCo and SNS-Toolbox, implementing a BPA-force calculation, and adding procedures for support, stimulation, perturbation, and parameter search. On the human reference model, alternative neural organizations can be evaluated against the same mechanical description. Separate W2L and BPA-actuated models extend the workflow to different bodies and actuators. The SNS-Simscape library provides graphical circuit construction and component demonstrations in MATLAB. These environments serve different parts of the work; importing a mechanism into Simscape is a separate achievement from controlling it with the demonstrated circuits.

The framework is evaluated at three levels. Circuit tests ask whether a controller produces the intended neural response. Interface tests ask whether activation, muscle force, and sensory measurements pass correctly between the neural and mechanical models. Coupled tests ask how the resulting motion changes with support, sensory pathways, or disturbances. The results below demonstrate each level within a stated scope. They establish a working basis for controlled neuromechanical experiments while identifying the balance, model-transfer, and controller-tuning problems that remain.

Existing tools supply the components of this framework. OpenSim provides musculoskeletal geometry, MyoConverter translates that geometry to MuJoCo, and SNS-Toolbox supplies neuron and synapse models [**????**]. MuJoCo supports scripted dynamics, contact,

and repeated evaluations; the version used here models rigid tendons. Simulink supplies graphical circuit assembly, and Simscape Multibody supplies mechanical imports, including CAD-derived mechanisms. AnimatLab contains the original walker and reference circuits. My contribution connects these components through the model conversions, actuator calculations, circuit implementations, and tests described below.

6.3 Software Tools

The tools connect musculoskeletal geometry, actuator forces, and neural circuits. Figure 6.2 shows these connections and their relationship to the proposed hardware system.

6.3.1 Model Conversion and Actuator Coupling

My first conversion attempt used the robot model, which combined BPA routing on the Gait2392 skeleton with unchanged human muscle parameters. MyoConverter's muscle-equilibrium checks rejected infeasible combinations of routed length and muscle parameters. I therefore separated conversion of the human reference from reconstruction of the robot routing. This distinction matters because an error caused by an artificial-actuator route should not be interpreted as a property of the original human model.

The human-model conversion also required specific repairs. Three small hip rotators per side, the quadratus femoris, gemelli, and piriformis, were disabled at model load. The converted model retains 92 actuator entries, but these six contribute no force. The model retains the OpenSim and International Society of Biomechanics coordinate convention; frame transformations and individual-actuator tests check that joint motion and torque

directions have the intended signs.

The source represents patella motion through muscle insertion locations that vary with knee angle. In the converted model, limits on the joints representing these moving path points conflicted with their coupling constraints. Removing those conflicting limits allowed the coupled path points to follow their prescribed trajectories. The rectus femoris path also required a patella-tracking via point to restore its knee-extension action. Explicit collision pairs introduced unwanted contacts, so contact was restricted to the foot–ground pairs. Passive joint torques supplied joint-stop behavior absent from the conversion. These repairs explain why conversion needed checks beyond successful loading: a model could load while producing an incorrect torque direction or unintended contact force.

Comparisons with the OpenSim attachment paths and single-actuator tests checked the transferred geometry and force directions. They do not establish that every muscle reproduces the source force throughout its range or that the controller produces walking. The neural and mechanical models exchange signals at a 2 ms solver time step.

For the BPA-routed robot model, an actuator-force calculation connects the neural and mechanical simulations. It takes the simulated muscle length and pressure as inputs and evaluates the quasi-static force model from Chapter 3; the resulting force acts along the simulated actuator path. This calculation does not include valve or pressure dynamics. Length, velocity, force, and foot-contact signals return to the controller through sensory channels. The correction terms of Chapter 4 do not by themselves establish that transferred routing and compliance are correct in a dynamic model; those properties require separate checks.

6.3.2 Support, Stimulation, and Parameter Search

The MuJoCo simulations allow mechanical support to be specified independently of the neural controller. A virtual harness applies forces and torques to the body: it can hold the model above the ground, lower it into contact, retain vertical support, or constrain selected directions of translation and rotation. These settings represent different experimental conditions. Suspended motion tests whether the controller and actuators can move the limbs, whereas supported ground contact adds loading and sensory feedback. Progression with retained support is reported separately from unsupported walking.

The training sequence, or curriculum, progresses from simpler conditions toward greater mechanical demand. Its stages include air stepping, supported standing, ground-contact stepping, and reduced support. Parameters can be held fixed to test transfer between stages or re-tuned at a new stage to seek a viable configuration. These procedures answer different questions: successful re-tuning shows that a configuration can be found, while a fixed-controller test shows whether an existing controller tolerates the change. A stage is demonstrated only when its stated motion and stability criteria are met.

Neural interventions include tonic currents, brief stimuli to selected circuit populations, and removal of specified feedback or coupling pathways. Mechanical interventions include timed forces at the pelvis and changes to segment mass and inertia. The recorded outputs include membrane voltage, actuator force, joint motion, foot contact, body position and orientation, and harness force. Together, these controls allow the neural response and the resulting motion to be examined in the same experiment.

Optuna automates parameter searches within specified bounds. Each candidate is

evaluated using measures such as bilateral contact cycles, inter-leg phase, forward displacement, joint trajectories, posture, and survival. The score determines which candidate is selected; a separate replay checks that candidate under the stated conditions. Optuna is used as a search tool, without a claim that it finds a global optimum or outperforms other optimizers. Interpretation depends on the objective and validity checks, including confirmation that every searched parameter reaches the implemented model.

6.3.3 Graphical Circuit Construction

The SNS-Simscape library implements neuron, synapse, afferent-encoder, activation, and muscle-force elements as editable Simulink blocks. The library retains the conductance-based neuron and synapse formulation supplied by SNS-Toolbox; it does not introduce a new neuron model. Verification uses selected circuits under matched inputs. Figure 6.3 shows the core blocks and graded-neuron signal flow.

Transcribing the source circuits also required distinguishing synapse-type parameters from per-instance editor fields. The Deng AnimatLab project contains 90 neurons, 196 synapses, and 58 named synapse types. Many instance fields retain a common editor default, while the named types specify different pathway conductances. The extraction records both fields so that a uniform default does not replace the source's pathway-specific strengths. A matching neuron count alone would not detect that error.

OpenSim-derived and SolidWorks-derived mechanisms were imported into Simscape Multibody. These mechanical imports are separate from the standalone reflex demonstrations.

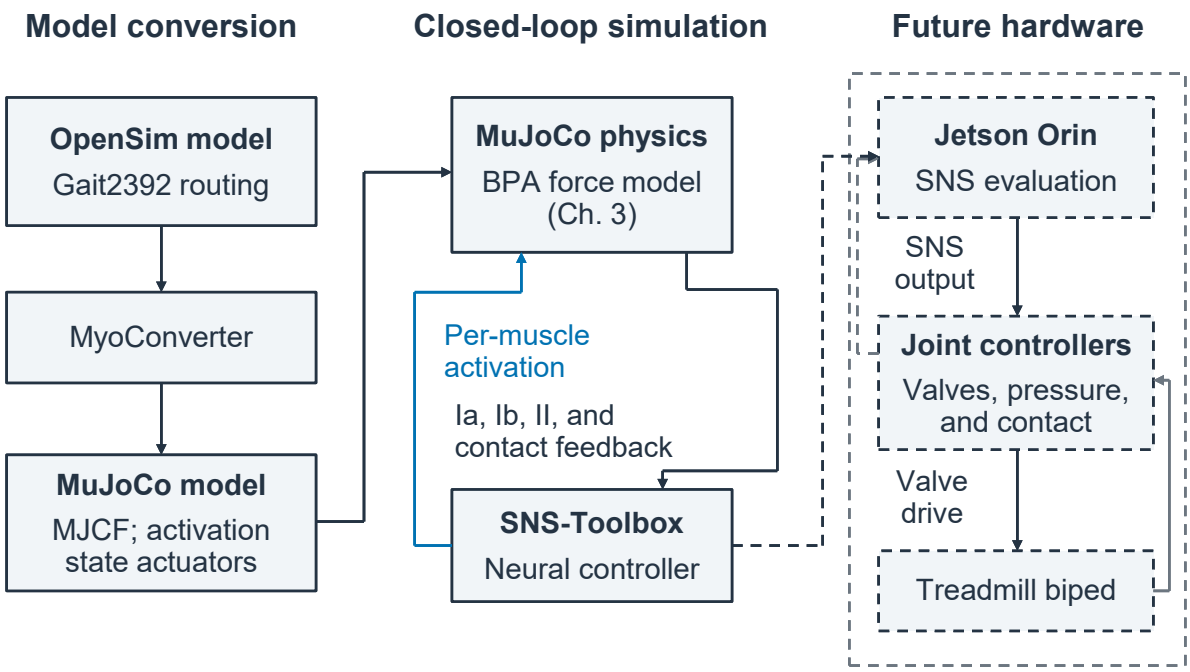


Figure 6.2: Model conversion, closed-loop simulation, and proposed hardware deployment. Solid simulation interfaces connect neural activation to BPA force and return afferent signals. The embedded valve-control and network-evaluation system is future work (Chapter 7).

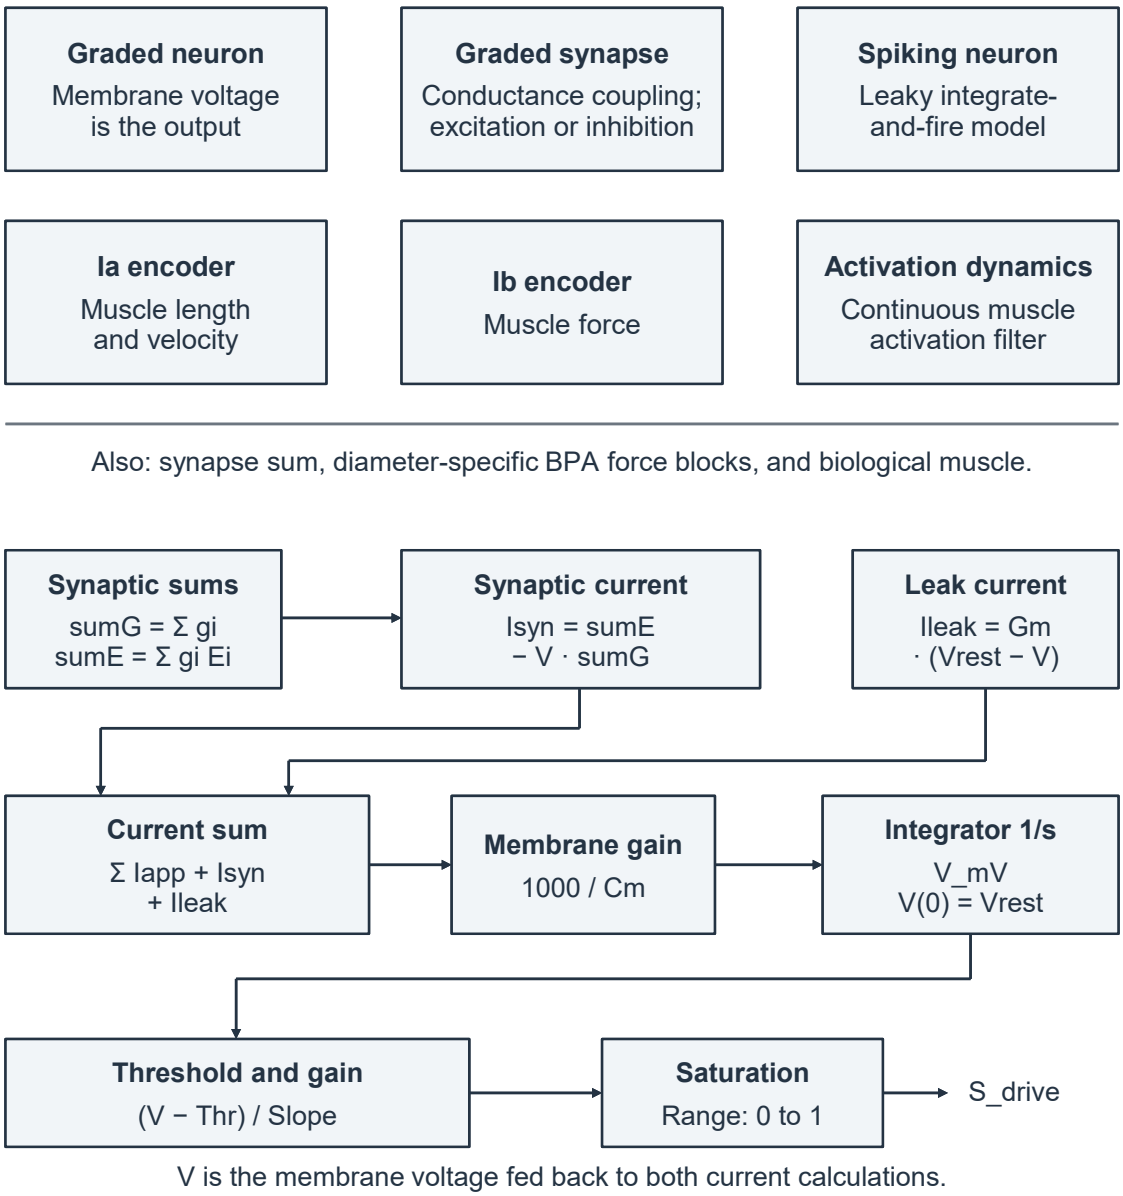


Figure 6.3: Core SNS-Simscape blocks (top) and graded-neuron implementation (bottom). Summed currents drive membrane integration; threshold, inverse-slope gain, and saturation bound the output. Voltage feeds both current calculations. The generic BPA demonstration remains provisional.

6.4 Verification of the Implemented Interfaces

The tests address three distinct questions: whether a circuit produces its expected response, whether implementations agree under matched inputs, and whether the neural and mechanical models exchange signals in a closed loop. The larger evaluation sets and their limitations are collected in Appendix D; neuron equations and calibration details remain in Appendix B.

6.4.1 Reflex Response

The library assembles Ia stretch excitation, reciprocal inhibition, and Ib autogenic inhibition in a reduced-order antagonist knee model (Figure 6.4). A second demonstration tests the stretch reflex in an arm holding a beverage container as it is loaded. With the reflex enabled, arm sag was 2.00°, compared with 7.85° when it was disabled. The comparison checks the implemented feedback response; experimental calibration of the mechanical parameters remains necessary.

6.4.2 Published Neural Behaviors

An activity-based reduction of the two-level circuit of **?** was tested against selected published neural behaviors. At the published drive settings, increasing extensor drive produced a longer extensor phase, while increasing flexor drive reversed that relationship. The deletion tests distinguished the two circuit levels: a PF-level pulse suppressed flexor MN bursts without resetting the RG schedule, whereas an RG-level pulse shifted the

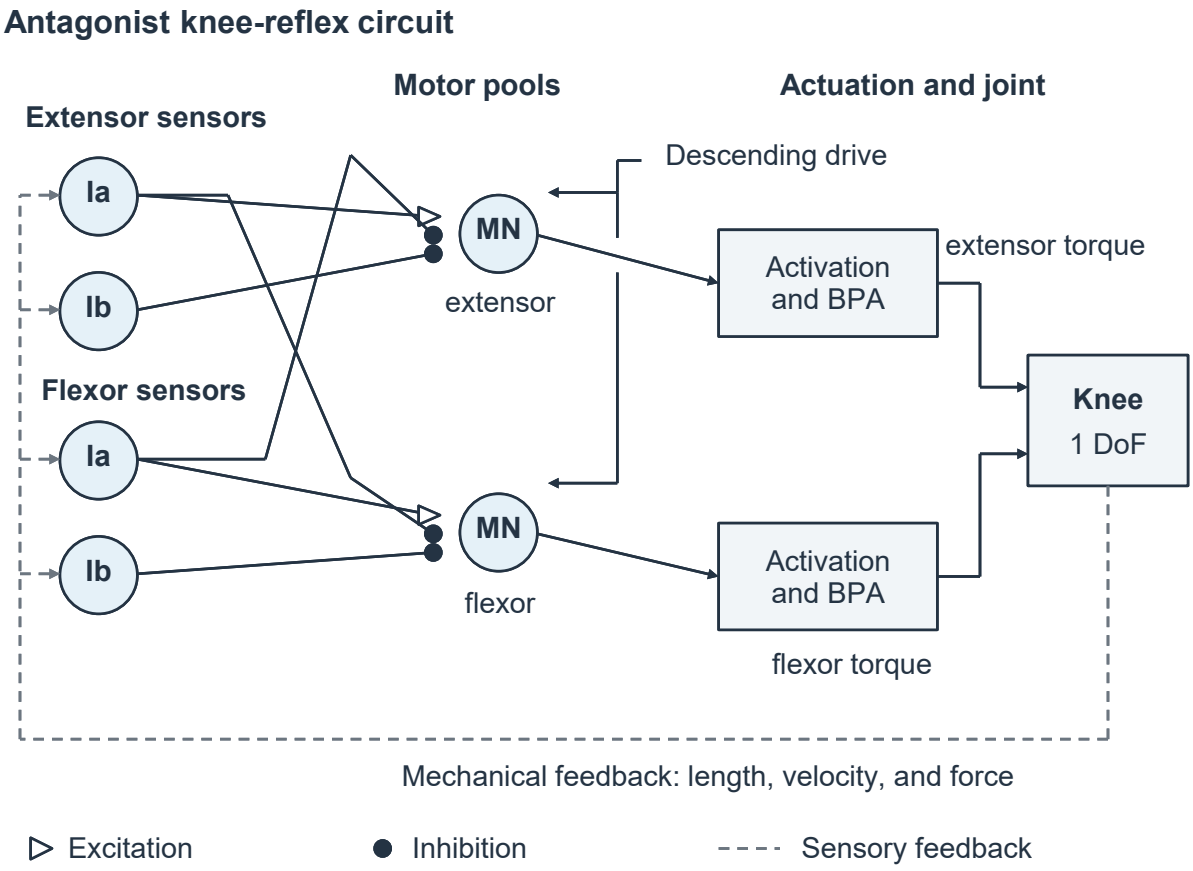


Figure 6.4: Reduced-order antagonist knee-reflex circuit. Triangles mark excitation, and filled circles mark inhibition. Ia input excites the homonymous motor pool and inhibits its antagonist; Ib input inhibits the homonymous pool. Motor output drives activation and BPA-force blocks. Mechanical parameters and the generic BPA block remain provisional.

subsequent rhythm. The first returning flexor burst differed from the extrapolated schedule by 0.001 periods after the PF perturbation and 0.416 periods after the RG perturbation. All 13 specified checks passed, including deterministic repeat runs.

This result reproduces selected qualitative behaviors of the published neural model. The implementation uses activity-based populations in place of the original spiking formulation, calibrated global conductances, and adapted pulse amplitudes and integration settings. It therefore tests the two-level explanation of resetting and non-resetting deletions without claiming an exact numerical reproduction of the original model or any walking mechanics. The Deng-derived and Di Russo-derived tests below have different scopes and should not inherit this result.

6.4.3 Matched Neural Implementations

The contact-driven W2L rhythm core was implemented in NumPy and native Simulink blocks from a common connectivity specification. Both implementations received the same input protocol. Burst periods matched at 1.000 s, all ten expected bursts were reproduced, and voltage-trace correlation was at least 0.99996 (Figure 6.5). The agreement is specific to this circuit and input protocol. Together with the contact-removal test, it checks the contact-driven design principle; it is not a quantitative comparison with the original Li experiment.

Removing the phasic heel-contact inputs caused the circuit to latch in both implementations. The alternating activity therefore depends on the specified sensory drive.

A separate Simulink implementation of a Deng-derived RG/PF circuit, with a fixed

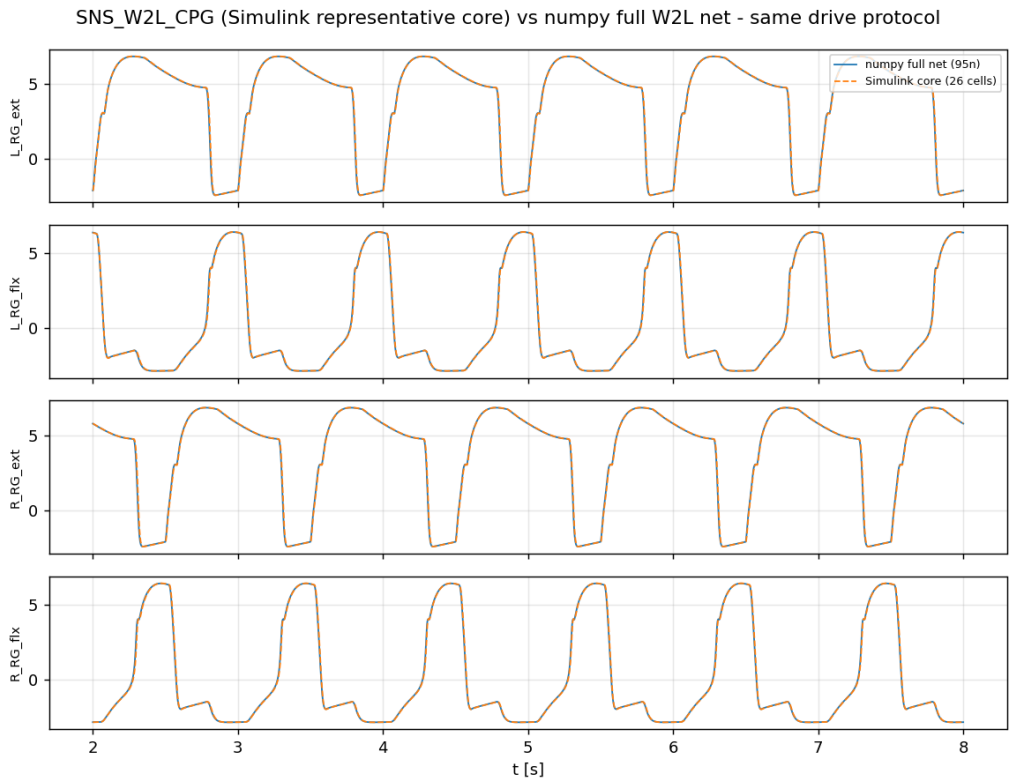


Figure 6.5: W2L rhythm-generator voltages under matched inputs in NumPy and Simulink. Burst periods and counts agree; full-trace correlation is at least 0.99996. The residual 2 ms lag is one discrete solver time step. The test concerns the driven rhythm core, not walking mechanics.

sodium-inactivation time constant of 350 ms, produced ten bursts over 20 s following a single 10 nA, 1 ms pulse. Its period was 1.938 s, compared with approximately 1.94 s in our NumPy reference, which was itself translated from the Nourse/Deng parameter tables. This test demonstrates self-sustained activity after a brief input, unlike the contact-driven comparison above. The original Deng AnimatLab project was extracted for connectivity and parameters but was not run for this comparison. The published Deng perturbation experiment therefore remains untested in translation. This comparison also does not demonstrate pointwise voltage agreement: the logged maximum voltage differences were approximately 63 mV.

Additional tests checked the graphical implementation at different scales. In the five-neuron SNS-Toolbox tutorial circuit, final membrane voltages agreed with the NumPy implementation to within 1.1×10^{-5} mV. The persistent-sodium tutorial pair reproduced its winner and loser steady states, but retained transient differences. For the 410-neuron export, the maximum difference across all final membrane voltages at 0.3 s was 4.201×10^{-6} mV with a 2 ms Euler time step. This short-window check tests translation of the network equations and wiring; it does not establish long-duration agreement or transfer of a tuned walking controller.

A Li-derived contact-driven circuit was also assembled with 58 neurons and 98 connections. This implementation substitutes graded stance cells with slow adaptation for the original spiking cells. Under scripted alternating heel inputs, it produced rhythmic activity with an inter-leg correlation of -0.271. It is an adaptation tested with imposed contacts, rather than a reproduction of the original circuit's closed-loop walking experiment.

6.4.4 Closed-Loop Model Coupling

The MuJoCo–Simulink bridge was tested against an independent Python integration using single-muscle activation, different control rates, and sensor readback. Readback was bit-exact over the scored window. The exported neural network also drove all 92 actuators during a 5 s open-loop run without afferent return. A separate reflex experiment exercised commands and sensory measurements through the interface. These tests establish selected force and signal paths across the interface, rather than a complete transferred locomotor controller.

A separate W2L BPA-actuation test replaced all 12 muscle actuators with the Festo force model. Both knees completed at least three air-stepping excursions, with periods of approximately 1.54 s, and the model subsequently maintained height and foot loading with retained support. Peak pressure was 372 kPa, below the 620 kPa ceiling. This is direct evidence that the controller, BPA-force calculation, and mechanical model operated together. The retained harness supplied substantial support, and no matched Hill-muscle versus BPA comparison was performed. This W2L test is distinct from the converted 92-actuator human model used in the bridge tests.

The Simulink air-stepping demonstration connects eight BPA-force elements to two planar legs with a fixed pelvis and no ground contact. It establishes that neural commands can drive the actuator and joint calculations together. However, the approximately 0.31 s swing period differed substantially from the RG period of 1.938 s. The observed limb oscillation therefore does not establish that the mechanics follow the neural rhythm. Mechanical imports from OpenSim and SolidWorks remain another stage of integration;

the standalone reflex and air-stepping demonstrations do not establish control of those imported mechanisms.

6.4.5 Perturbation Robustness and Pathway Deletions

The W2L tests examine whether a controller retains bilateral contact cycling when the body or selected neural pathways change. An implementation check showed that the initial search applied only six of the 61 saved parameter fields; the other 55 fields did not reach the model. Those records describe a controller with the remaining pathways at their default strengths. The results reported here use a later replay with the stored individual fields applied. Because those values had not been optimized under the corrected mapping, this replay tests a candidate controller rather than a newly optimized solution.

With the stored individual fields applied and settings held fixed, the control completed 17 left and 16 right heel-contact cycles during the 16 s record (Figure 6.6). Removing the crossed heel pathway produced 16/16 cycles, the Ia chain 13/13, the Renshaw pathways 16/14, the heel cascade 17/15, and commissural coupling 13/12. Removing Ib load feedback produced no heel-contact cycles on either side. The specific result is that Ib feedback was necessary for contact cycling in this configuration, while commissural removal reduced the count without abolishing it. Contact counts alone do not show that phase, posture, or gait quality remained unchanged.

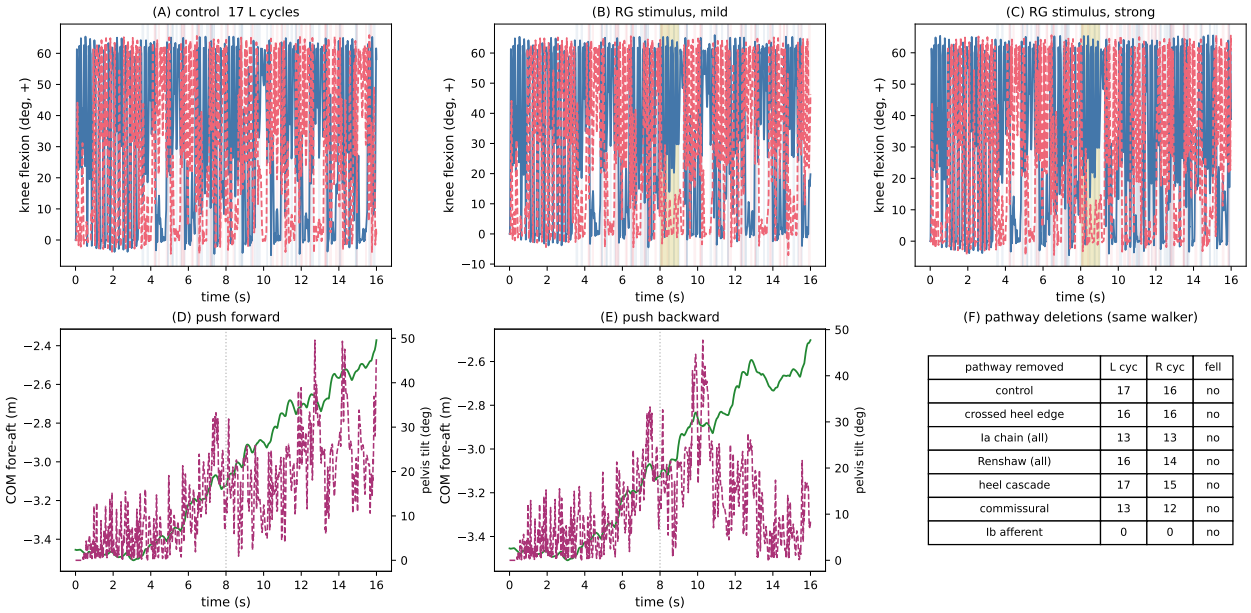
One-second RG stimuli and 40 N, 0.15 s fore-aft pushes also left bilateral contact cycling present. Nevertheless, the stronger stimulus reached approximately 62.8° pelvis tilt, and the forward and backward pushes reached approximately 49.6° and 47.7°. These runs do

not satisfy a 45° posture criterion merely because cycling persists. Recovery of the original gait requires a separately defined phase, posture, and timing measure.

The corrected candidate was also tested with thigh or distal-leg masses and inertias changed by ±10 and ±20 percent, and whole-leg properties changed by ±10 percent. All ten conditions retained bilateral contact cycling, but nine exceeded the 45° tilt gate. Only one condition passed the recorded survival criterion. Minimum pelvis height across these records was approximately 0.898 m. This is evidence of persistent cycling with inadequate posture robustness, rather than ten successful robust-walking tests.

All these evaluations retained a virtual harness with vertical height feedback and constraints on lateral motion, roll, and yaw. The nominal weight-support setting was 25 %; the realized force also depended on the feedback terms. Earlier searches reduced the nominal setting from 70 to 10 percent with re-tuning between settings, but those records precede full application of the individual fields. Their failed local conductance probes, 9/12 at nominal 25 percent and 7/12 at nominal 10 percent, also describe the earlier effective controller. They cannot certify the robustness of the corrected candidate. These are W2L results, not demonstrations of the BPA-routed Gait2392 robot model or harness-free walking.

A separate controller derived from **?** was coupled to the Gait2392 human model. It uses per-muscle reflex parameters and a staged optimization procedure, with changes to the source model and optimization method. The recorded second-stage stability search completed 120 trials without a configuration passing its survival criterion. This is a negative result for that implementation and search, rather than a reproduction of the published walking result. It also shows why completing a connectome transcription or



(F) pathway deletions (same walker)

pathway removed	L cyc	R cyc	fell
control	17	16	no
crossed heel edge	16	16	no
Ia chain (all)	13	13	no
Renshaw (all)	16	14	no
heel cascade	17	15	no
commissural	13	12	no
Ib afferent	0	0	no

Figure 6.6: W2L records after applying the stored individual conductances, with retained harness support. (A–C) Knee motion in control and one-second RG-stimulus conditions. (D, E) Center-of-mass position and pelvis tilt after forward and backward pushes at 8 s. (F) Independently counted left and right heel-contact cycles after pathway removals. The control has 17/16 cycles; commissural removal retains 13/12, while Ib removal gives 0/0. Continued cycling does not establish posture recovery: both push conditions exceed the model's 45° tilt threshold for a fall. The "fell" column summarizes the deletion conditions only.

obtaining a fitted neural output cannot substitute for testing the coupled mechanics.

6.4.6 Feedback and Reduced Hand Drive

The human-model feedback-tuning campaign provides a complementary test of sustained activity. With the selected feedback parameters fixed, bilateral stepping-like motion persisted as the added hand-drive currents were reduced to zero, with 8 and 11 contact cycles at zero scale. The original feedback configuration lost bilateral activity below 0.85 of those added currents. However, the descending drive remained active, the mechanics retained support, and the motion was in place. This comparison tests the effect of feedback re-tuning on dependence on added drive; it is not a fixed-controller feedback-intact versus feedback-disabled experiment. A matched feedback-ablation pair on the walking BPA model remains absent.

6.5 Capabilities and Remaining Limits

The framework addresses the project's model-integration gap by making neural control, muscle mechanics, sensory return, and support conditions explicit and adjustable. It connects the Gait2392 reference model to neural controllers, provides force calculations for BPA-actuated models, and supplies procedures for staged tuning and controlled interventions. Graphical circuit construction in Simulink complements the scripted evaluations. These capabilities allow a controller to be examined at the circuit, interface, and coupled-model levels before a whole-model result is interpreted.

The tests show what that framework can establish. The Rybak-derived neural model

reproduced selected phase-resetting behaviors, the contact-driven rhythm core agreed across implementations, and the BPA gate demonstrated air stepping and supported standing with the actuator force model inside the simulation. W2L pathway tests linked Ib feedback to bilateral contact cycling at one operating point. They also showed why cycling and walking must be distinguished: the corrected mass-change candidate exceeded the posture criterion in nine of ten conditions. The Li transcription changes the original neuron dynamics, and the Simulink air-stepping mechanics do not yet follow the RG period. These limits identify specific model and control problems for subsequent work.

The present contribution is a usable framework for developing and testing neuromechanical controllers, with demonstrated components and explicit limits. Establishing a general biological mechanism will require matched interventions and comparisons with the corresponding biological or published-model results. Achieving unsupported walking will require further balance and weight-bearing control. The framework makes those questions accessible to controlled simulation; it does not resolve them by the existence of a tuned trajectory. Chapter 7 discusses the remaining integration and hardware work.

7 Discussion

This chapter considers how the actuator models, torque corrections, sensory-afferent resource, and simulation tools support biomimetic robot design, and identifies the experiments and integration work still required.

7.1 The Work as a Whole

The outcome that organizes this discussion is that improved first-pass design of a BPA-actuated biomimetic leg is now achievable: a designer can predict actuator force, joint torque, and range of motion accurately enough to commit to an actuator set, a resting length, and a route before the hardware exists (Figure 7.1). At the center of that outcome are the identified design factors that remove the rigid-body assumptions from the torque model (Chapter 4): a constant length offset χ_0, effective stiffnesses χ_1 and χ_2 that lump the series compliance of the attachment brackets, fixtures, and artificial tendons into the bracket frame's principal directions, and a wrapping-loss factor χ_3 that removes usable actuator length as an extensor bends over the joint. The terms were identified on the pinned knee, and they hold up beyond it: carried unchanged to the biomimetic knee with its reinforced frame, the corrected model explains the bulk of the measured torque variance there as well, cutting the biomimetic extensor's fraction of variance unexplained from 4.49 to 0.67 with the shared terms (Section 4.3.2). Whether the set is complete is itself now an experimental question: a fourth term χ_4, a stiffness acting along the force path independently of force orientation and presently absorbed by the stiffness pair (χ_1 and

χ_2), is arguably the next candidate, and the identification machinery delivered here is the instrument for settling it.

Improved BPA characterization is the first enabler (Chapter 3). The maximum isometric force of Festo BPAs was shown to depend strongly on resting length, a previously unreported characteristic that corrects the manufacturer's tool at exactly the short lengths a leg design requires, and a nondimensional force surface predicts isometric force from relative strain and relative pressure with only three diameter-dependent coefficients. Because resting length is a primary degree of freedom in actuator selection, this model is what makes that selection rational. The quasi-static surface is complemented at the dynamic end by the state-space, optimization-based model of pressure and force dynamics in pulse-modulated BPAs [**?**], which supplies the actuator behavior that the isometric maps cannot.

Attachment and routing optimization is the second enabler, and it is tiered. The placement optimization of Morrow, Bolen, and Hunt found artificial muscle attachment locations that reproduce human torque profiles and ranges of motion about the trunk and leg joints [**?**], converting the earlier finding that transplanted human attachment locations do not transfer directly to a BPA-actuated robot [**??**] into a usable design tool; the route redesigns of Section 4.3.2 rerun that machinery under the corrected torque model. At the route-selection tier, with two 20 mm BPAs per route and routes chosen with the actuators' bending in mind, the redesigned flexor and extensor each meet or exceed the magnitude of their human-scaled torque targets across the full knee range of motion, and the cam-profile extensor converts the direct route's thin margin into genuine design headroom (Section 4.3.2). Stated as a design program, the generalization is one or more BPAs replacing one or more human muscles across one or more joints, serving the two degrees of freedom

that matter most at each joint, with routes grown outward from positions close to the bone and held to the spatial constraints of the candidate actuators and of the BPAs already mounted; the broader route-search diagnostics are collected in Appendix C.7. Where the attachments are fixed and only the route and the tendon length can change, the cam-shape solver applies: the tibial-side route points trace a designed cam profile, trading biomimicry for torque-profile optimization and reduced component forces.

The Sensory Afferent Database (SADb) is the third enabler (Chapter 5). It organizes the sensory-afferent literature into a structured afferent resource [**?**]: of the 561 papers retained after the audit and the prune, 523 carry structured curation notes, 86 carry an afferent classification, and 258 carry notes translating each paper toward robot or simulation use, and the full audited corpus of 943 papers, including the 382 pruned records preserved under an archived flag, remains reconstructible in the packaged tool. The toolset matters as much as the corpus. Its detailed table drills from any filter down to a paper's full notes; its pivot view counts the corpus along any two dimensions at once; its bubble map draws every paper as a neuron, soma sized by citations, with citation links rendered as excitatory and inhibitory synapses in the laboratory's convention and a focus mode that walks the citation neighborhood. The monthly literature-scraping extension is under development; it is separate from the completed curation and application described here. The evaluation of AI-assisted literature triage established the division of labor that makes expansion affordable [**?**]: the machine drafts, and the human rules.

The fourth enabler is the neuromechanical simulation framework (Chapter 6). It connects OpenSim-derived mechanics to neural controllers, evaluates BPA forces within coupled models, and provides graphical circuit construction and matched implementation

tests. Adjustable support, staged parameter tuning, neural stimuli, pathway removal, and mechanical perturbations make controller behavior testable under specified conditions. The demonstrated results include agreement of the contact-driven rhythm core across implementations, BPA-actuated air stepping and supported standing, and W2L pathway tests. The latter identify Ib dependence at one operating point while also exposing posture failures under mechanical perturbation. No harness-free walking has been demonstrated. In separate hardware work, a BPA-actuated limb maintained balance using pseudo-vestibular inertial feedback [**???**]. Extending these components to unsupported bilateral walking, wider model transfer, and embedded execution remains future work.

7.2 Future Work

The subsections that follow are options for the researchers who take this program up, ordered from the actuator outward to the whole body. Each draws directly on the delivered models, the database, and the simulation tools.

7.2.1 Expanding the BPA Force Characterization

The isometric characterization of Chapter 3 leaves one actuator curve unidentified and one regime unmeasured, and both are tractable next experiments. The unidentified curve is the maximum-contraction function: ε_{620} could not be pinned to a predictive relationship with resting length, so each actuator's free contraction is recorded individually before use. Controlled tests designed for that relationship, across multiple product batches of the 10 mm, 20 mm and 40 mm diameters, would either produce general fit surfaces or establish

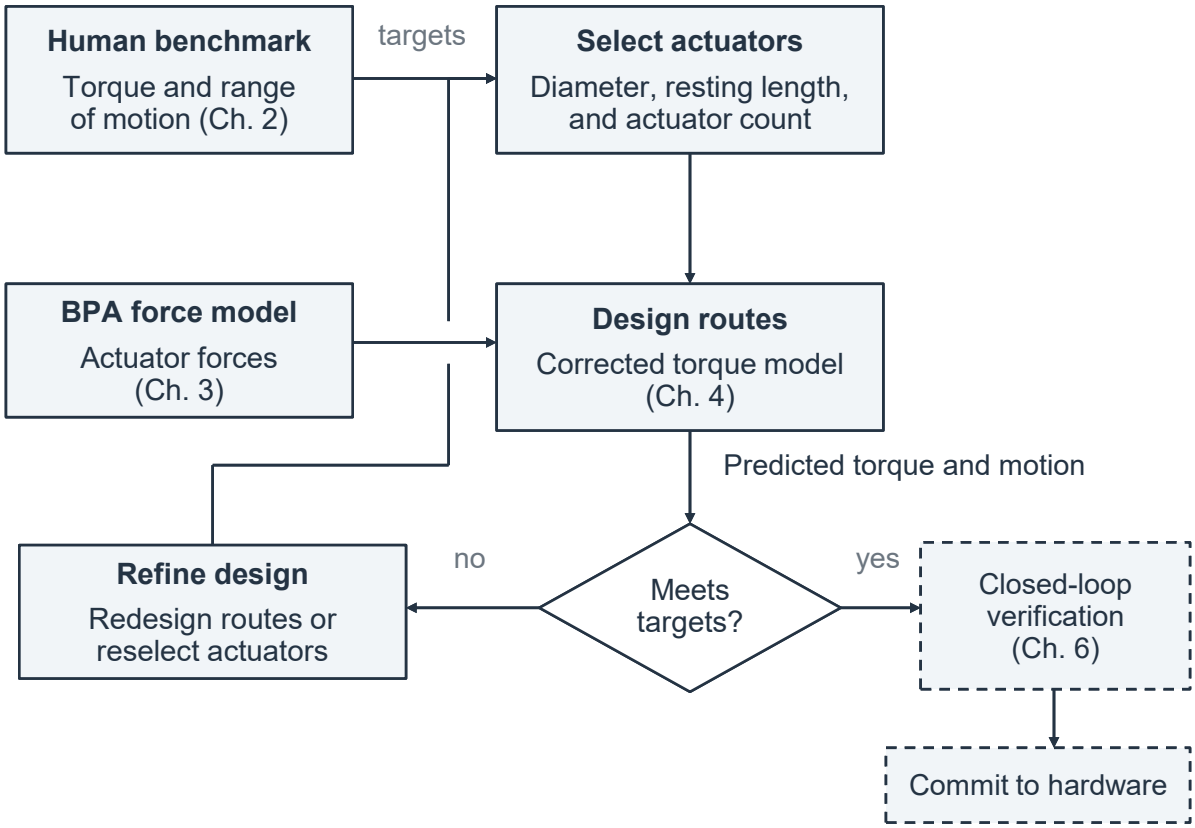


Figure 7.1: The first-pass design workflow the delivered enablers make possible. Human torque and range-of-motion targets from the OpenSim benchmark and the validated BPA force model of Chapter 3 drive actuator selection; route design runs under the corrected torque model of Chapter 4; the predicted torque and range of motion are checked against the targets, with route redesign or actuator reselection as the loop when they fall short; and the candidate configuration is checked in closed-loop simulation with the tools of Chapter 6 before hardware is cut.

batch variation as the binding uncertainty. The 20 mm actuators characterized here came from a single product batch, so batch-to-batch variation is unquantified and those fits in particular carry that limitation.

The unmeasured regime is everything off the isometric plateau. A biological muscle model separates active force from passive force and couples both to activation and contraction dynamics [**??**], and Appendix A, which collects the Thelen and Millard formulations used as torque targets in this work, is the natural home for a BPA counterpart: extending it with the remaining equations of those models, the Gaussian active-force shape factors, the activation and deactivation time constants, and the tendon force curves, would let the BPA model be decomposed into active pressurized-braid force, passive resistance of the unpressurized actuator and sleeve, and contraction dynamics. Isokinetic and isotonic or isobaric characterizations could then be assimilated into the same framework rather than fitted separately, with the pressure–force dynamics model of **?** supplying the actuator-side dynamics.

Two experimental programs would complete the mechanical picture. Instron tensile testing of actuators with their Onyx and aluminum end caps, unpressurized and pressurized, would measure the end-effect region that the characterization identifies as the mechanism behind the resting-length dependence of maximum force. A finite-element program sketched from the construction disclosed in Festo's patent literature would give the actuator a first-principles model to check the empirical surfaces against [**?**]: a symmetry half-model of the actuator, a half circle fixed at the outer and inner diameter at one end with the diameter free to expand at the other and an optional lateral force or fixity, assembled from thin-shell or membrane layers in inter-layer contact, with the two handed aramid

weaves of the braided sleeve represented as beam or truss elements. Figure 7.2 shows the construction such a model would represent.

7.2.2 Expanding Attachment and Routing Design

Each routing tier of Section 7.1 leaves a named next step. The knee's flexor redesign assumes both actuators share one attachment route; making the mirrored pair explicit is the next refinement. The bi-pulley sweep showed that reeved routes do not yet fit the anatomical corridors, where geometry rather than torque was the binding constraint, so clearance-relaxed attachment sites are where that freedom should pay. The whole-body pass needs its route-builder failure fixed and either larger actuators or relaxed corridors on the two routes that satisfy geometry but fall short of torque.

The design element this program has so far declined is a patella. The torque model wraps the extensor BPA over the joint surface and accounts for the loss with χ_3, which was easier than adding a patella, deterministically designing its location as a function of knee angle, and accounting for friction (Chapter 4). A designed sesamoid is the biomimetic answer: an artificial patella, its location optimized over the range of motion the way route points now are, would buy back moment arm at high flexion and shield the actuator from the joint, at the cost of a friction model and a contact schedule the optimizer must respect. The cam-profile solver is the proof that a moving routing surface can be a design variable; a patella is its biomimetic limit case.

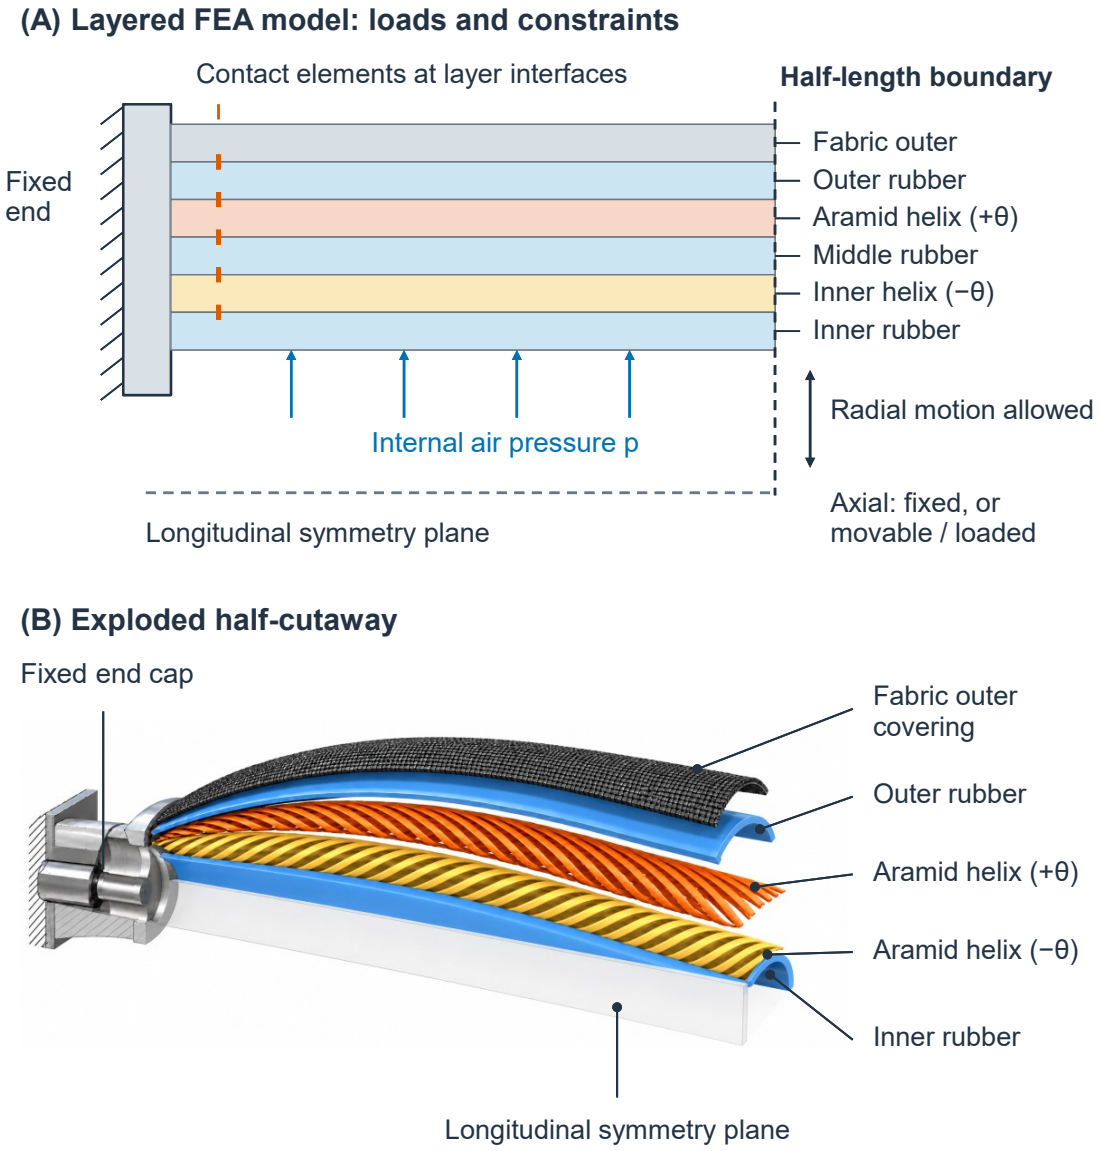


Figure 7.2: Proposed layered BPA finite-element model, not to scale. (A) Contact interfaces, internal pressure, and end constraints. Longitudinal symmetry permits in-plane expansion; at half length, fix axial displacement or allow axial motion with an optional load. (B) Illustrative cutaway identifying materials and longitudinal symmetry. Middle rubber is separated only in (A); cap details and layer spacing are schematic.

7.2.3 Expanding the Sensory Afferent Database

A monthly literature-scraping extension is being implemented to identify candidate additions to the database. Its retrieval and classification must be evaluated against human decisions before it can support routine updates. Further curation can also target missing afferent classifications, pathway links, and preparation metadata in the existing records. The AI-assisted triage study establishes the relevant division of responsibility: automated tools draft and retrieve, while I retain the inclusion and interpretation decisions [**?**].

7.2.4 Closing the Loop in Simulation

The simulation laboratory of Chapter 6 is assembled, audited, and in active use. The converted robot model carries all of its muscle-tendon actuators with audited geometry, and the BPA force interface computes actuator forces from the validated model inside the loop; the implemented interfaces and their limits are described in Section 6.3. The 410-neuron SNS-Toolbox instantiation of the spinal network, with per-leg rhythm-generator half-centers, phase-shifted pattern-formation groups, motoneuron pools for every actuator, and Ia, Ib, and II afferent populations, couples to the converted model at a 2 ms neural step, and supported stepping configurations have been tuned against a walking reference and evaluated on the ground (Section D.3.1). A drive-weaning study has since produced a feedback-sustained bilateral configuration whose stepping persists with the hand-drive withdrawn entirely, still in place under the support rig, and the robot-side model progresses under its harness at weaned drive (Section D.5). What those same evaluations localize as missing defines the work remaining here: the balance and weight-bearing layers that

unsupported, bilateral walking requires; transfer from a single tuning reference to the diversity of the locomotion reference library; and the systematic verification series that follows.

The verification series comes first because it is cheap and decisive. The circuits' signs, gains, and phase structure are checked against the experimental literature with the stimulus protocols the physiology itself uses: a tonic stimulus applied at the motoneuron, pattern-formation, and rhythm-generator layers, with the resulting outputs compared against the tonic central and sensory stimulation that facilitates involuntary air-stepping in humans [**?**] and the deletion and stimulation phenomena reproduced by the two-level CPG models and their neuromechanical implementations [**???**]. Isolated rhythm-generator tests checked antiphase responses under tonic drive; the motoneuron-level and pattern-formation-level series remain to be run, together with the motoneuron skipping-deletion phenomenon of the reference models.

The laboratory's hypothesis-testing program is stimulus, deletion, and addition. The stimulus and deletion families have their first sweeps (Section D.3.2). Pathway addition is their complement and the one the tools have not yet exercised: adding a candidate circuit to a tuned configuration, a commissural class, a sensory gain, or a modulatory layer, and re-tuning, which turns connectome hypotheses into falsifiable experiments. The closed-loop comparison of spiking and non-spiking controllers of the same topology belongs to the same program, with the spiking mirror's untuned bilateral ground gait as the standing lead (Section D.4).

The tools also compose across programs, and the composition is now cheap enough to be a standard workflow. A full-body OpenSim model with arms and trunk, of the

kind MyoSuite distributes [**?**], can be brought into the SNS-Simscape library for graphical assembly and inspection of the controller, and the same network imported back into SNS-Toolbox on MuJoCo [**?**] for the speed that tuning and long evaluations require. A controller built in that loop inherits the conversion audit, the cross-engine checks, and the deletion apparatus of Chapter 6 without rebuilding any of them.

7.2.5 Extending the Sensorium

Once stepping is stable, three supraspinal contributions will be added, and a fourth is in line behind them (Figure 7.3). First, vestibular feedback will be simulated with IMU-like sensing of body orientation and angular velocity, extending the balance-platform work, in which an inverted pendulum with pseudo-vestibular IMU feedback maintained balance on an oscillating platform [**???**] and analytic neural networks reproduced human balance dynamics [**?**]. Second, ocular feedback will be modeled to investigate how visual stabilization cues can help reset IMU drift, a well-known practical problem for inertial sensing on robots, by treating visual measurements of the horizon and surroundings as a low-frequency correction to the vestibular estimate. Third, a cerebellar layer will be explored as an adaptive gateway between higher-level locomotion commands and the CPG: a structure that can gate between gaits or movement patterns (e.g., walking, running, and turning) and tune the phase relationships among them, in line with the cerebellum's established role in adaptation and coordination. Multilevel spinal-plus-cerebellar synthetic nervous systems have already been demonstrated for insect-like legs [**?**], and hierarchical cortico-cerebellar-spinal models have been proposed for adaptive step-length control in

bipedal locomotion [**?**]. The fourth addition, basal-ganglia circuits realizing the inhibitory and excitatory gating that selects among actions and locomotor modes, would give the controller a substrate for choosing between behaviors rather than only executing one. The database already holds the entry points for this program: the human balance and posture cluster and the multisensory recalibration cluster are curated and connected to the locomotor core through the citation graph (Chapter 5), and the proprioceptive–vestibular interaction literature will guide how the signals are weighted [**?**].

7.2.6 Real-Time Hardware Realization

The final stage returns to hardware, and its near version is a step-by-step integration problem for which the laboratory's existing equipment already stocks the parts. The stack is fixed: Teensy-class microcontrollers handle the low-latency per-joint loops (valve control, pressure sensing, and contact sensing), and an NVIDIA Jetson Orin serves as the companion computer for SNS network evaluation and activation computation. The integration steps then run in order. Execute the SNS-Toolbox network in real time on the treadmill biped. Close feedback from the sensors the lab already uses: the joint encoders, the actuator pressure sensors, the liquid-wire strain gauges, the pressure manifolds, and the IMUs. Compute the afferent signals from those streams exactly as the simulation tools compute them, Ia length and velocity and II length from the encoders, and Ib force from pressure-based force estimates checked against the strain gauges. Tune computer-in-the-loop, with the embedded stack executing the network against the simulated body before it executes against the physical one, so that gain changes are tested at simulation

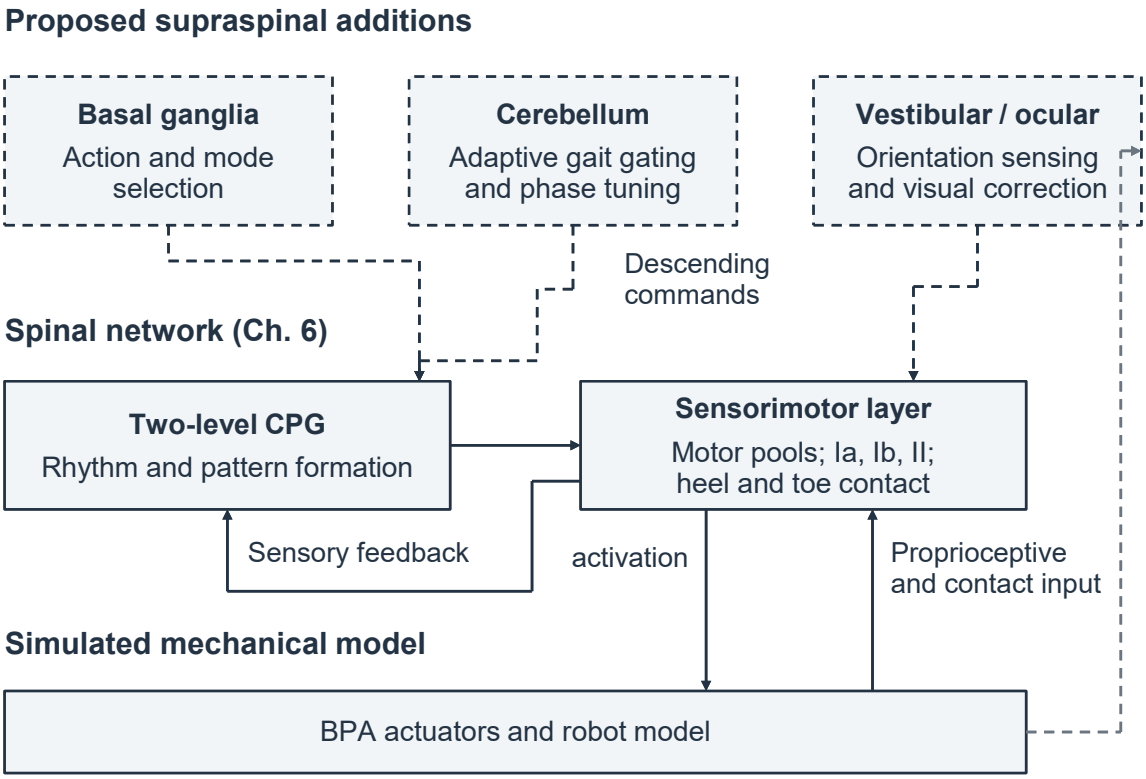


Figure 7.3: The control stack with the planned supraspinal additions. The spinal network tier, the two-level central pattern generator with motor-neuron pools and Ia, Ib, and II afferent plus heel and toe contact feedback, and the BPA-actuated body tier are implemented and demonstrated (Chapter 6), with the balance-platform studies as the demonstrated precursor of the vestibular channel. The dashed supraspinal tier is the extension program of this section: basal-ganglia circuits for action and mode selection, a cerebellar adaptive gateway between high-level commands and the pattern generator, and vestibular and ocular sensing closing body-orientation and visual feedback onto the same loop.

speed and validated on hardware only when they survive. Keep the serial connection to the host computer for data logging and high-level commands only (e.g., run, walk, or stimulate a named neural pathway), so the control loop stays fully embedded. The wiring schematics, the neural-pathway schematics, and the step-by-step commissioning procedure that this integration produces belong in a hardware appendix rather than in this chapter; the connectome figures of Appendix B already serve as the pathway schematics for the spinal network.

The precedent and the margins both support feasibility. Walking control of a physical musculoskeletal biped with McKibben muscles has been demonstrated with spring-loaded inverted-pendulum controllers [**?**], and the AARL's quadruped has already traveled part of this road: a MuJoCo model of the BPA quadruped was coupled to SNS-Toolbox, produced closed-loop coordinated stepping, and was transferred to hardware with potentiometer and pressure feedback on a treadmill [**?**]; the present program extends that approach from the quadruped to a leg built on human anatomy. An actuation-level design lever also remains open: the BPA's limited contraction can force long musculotendon routings, a burden visible in the extensor redesign bounds of Appendix C and encountered in my master's routing work [**?**]. A block-and-tackle "reverse pulley" arrangement, demonstrated on the AARL's self-balancing inverted-pendulum platform, addresses this by trading transmitted force for travel: paired BPAs pulled through compound pulleys on a tendon coupled to the limb, increasing range of motion at the cost of reduced transmitted force and increased reaction force on the structure [**???**]. Because joint torque is computed in the body frame, configurations that turn that reaction force to advantage may exist, making the arrangement a candidate for joints whose optimized routes demand impractically long

actuators.

Because the static force and torque models of Chapters 3 and 4 are computationally inexpensive lookup-table candidates, they will serve as the feedforward model inside the real-time controller, while the dynamic pressure–force behavior under pulse-modulated control will draw on the state-space model of **?**. Embedded execution of the complete CPG-plus-reflex network on this stack is the intended demonstration of the dissertation project: a robot that walks with human-like biomechanics under a biologically grounded neural controller, and whose nervous system, unlike a living animal's, can be instrumented, stimulated, and lesioned at will.

7.2.7 Toward Dynamic Whole-Body Maneuvers

Beyond steady walking, the long-term goal of this line of research is dynamic, whole-body maneuvering: transitions between gaits, reactive balancing, and ultimately ballistic maneuvers that require coordinating the full actuator set through brief, precisely timed activation bursts. My aspiration for this program is a robot that can perform a kickflip, a maneuver that, like its skateboard namesake, demands explosive coordinated extension, precise timing, aerial body orientation, and controlled landing. It is exactly the kind of behavior that exercises every layer of the architecture described above, from contact receptors and spinal pattern generation to vestibular orientation sensing and cerebellar coordination, and it would demonstrate that biomimetic actuation and neuromechanical control scale from steady locomotion to dynamic whole-body movement.

These subsections are options, not prerequisites; each can be pursued on the delivered

framework independently of the others. What they all draw on, the actuator and torque models, the routing machinery, the afferent database, and the simulation laboratory, is consolidated in Chapter 8.

8 Conclusion

This dissertation moved the biomimetic humanoid robot project from an unvalidated design concept to a quantitatively characterized actuator set and experimentally tested knee-torque models, and it defined the path from those experimental foundations to neuromechanical control.

At the actuator level, the analysis in Chapter 3 produced novel equations for calculating isometric force in Festo 10 mm and 20 mm-diameter BPAs. The work elucidated the relationship between maximum BPA force at 620 kPa and resting length, $F_{620}(l_0)$, a previously unreported characteristic of these actuators that the manufacturer's tool predicts with the opposite trend, and it produced a nondimensional force surface, $F^*(\epsilon^*, P^*)$, with only three diameter-dependent coefficients. Taken together, these fits predict isometric BPA force across resting lengths from a small number of measurements, correct the large overpredictions of the manufacturer's tool at short resting lengths, and outperform published single-length models. In a design workflow, these equations make it possible to select actuator diameter and resting length and to evaluate whether a proposed muscle routing can meet force and range-of-motion requirements, before committing to a prototype.

At the joint level (Chapter 4), the BPA model was combined with transmission corrections for effective length, series compliance, and wrapping. Identification on the pinned knee improved torque prediction, and the shared correction terms also improved predictions on the biomimetic four-bar knee. The corrected model supports route redesign: the

optimized two-BPA flexor and cam-extensor configurations are predicted to meet their human-scaled torque targets across the design range. The isometric tests demonstrate that the redesigned configurations meet or exceed the human reference torques over most of the tested range of motion.

The sensory-afferent resource of Chapter 5 makes relevant physiological reports easier to locate and compare while retaining their experimental context. Chapter 6 provides a framework for connecting neural controllers to mechanical models and testing them under specified support, stimulation, and perturbation conditions. Independent implementations reproduced the contact-driven rhythm core with exact period agreement and trace correlations above 0.9999. BPA-actuated models demonstrated air stepping and supported standing. The coupled pathway and perturbation tests identified model-specific feedback dependence and posture limitations. These results establish implemented experimental capabilities; harness-free walking remains unachieved.

The next steps are to complete model integration, extend the mechanical characterization to dynamic conditions, improve load-bearing and balance control, and test the resulting system under progressively reduced support (Chapter 7). These steps are required before using whole-body robot behavior as evidence about a biological control hypothesis.

The dissertation provides experimentally grounded actuator and joint models, a curated afferent resource, and simulation tools for that work. Their scientific value lies in making assumptions explicit and testable as the project moves toward controlled neuromechanical experiments on a physical robot.

Bibliography

Turgay Akay. Sensory Feedback Control of Locomotor Pattern Generation in Cats and Mice. *Neuroscience*, 450:161–167, 2020. doi: 10.1016/j.neuroscience.2020.05.008.

Turgay Akay and Andrew J. Murray. Relative Contribution of Proprioceptive and Vestibular Sensory Systems to Locomotion: An Overview. *International Journal of Molecular Sciences*, 22(3):1467, 2021. doi: 10.3390/ijms22031467.

Turgay Akay, William G. Tourtellotte, Silvia Arber, and Thomas M. Jessell. Degradation of mouse locomotor pattern in the absence of proprioceptive sensory feedback. *Proceedings of the National Academy of Sciences*, 111(47):16877–16882, 2014. doi: 10.1073/pnas.1419045111.

Efe Aksöz, Marco Laubacher, Robert Riener, and Kenneth Hunt. Design of an isokinetic knee dynamometer for evaluation of functional electrical stimulation strategies. *Medical Engineering & Physics*, 73, August 2019. doi: 10.1016/j.medengphy.2019.07.010.

J. G. Andrews. A General Method for Determining the Functional Role of a Muscle. *Journal of Biomechanical Engineering*, 107(4):348–353, November 1985. ISSN 0148-0731. doi: 10.1115/1.3138568. URL `https://doi.org/10.1115/1.3138568`.

Bryson Armstead, Christian Jackson, William Nourse, and Roger D. Quinn. Design and validation of a robotic cat hindlimb. In *Biomimetic and Biohybrid Systems*, Lecture Notes in Computer Science. Springer, 2025. doi: 10.1007/978-3-032-07448-5_6.

Edith M. Arnold, Samuel R. Ward, Richard L. Lieber, and Scott L. Delp. A Model of the Lower Limb for Analysis of Human Movement. *Annals of Biomedical Engineering*, 38(2): 269–279, February 2010. ISSN 0090-6964, 1573-9686. doi: 10.1007/s10439-009-9852-5. URL `http://link.springer.com/10.1007/s10439-009-9852-5`.

Yuki Asano, Kei Okada, and Masayuki Inaba. Design principles of a human mimetic humanoid: Humanoid platform to study human intelligence and internal body system. *Science Robotics*, 2(13), December 2017. ISSN 2470-9476. doi: 10.1126/scirobotics.aaq0899. URL `https://robotics.sciencemag.org/content/2/13/eaaq0899`.

Yuki Asano, Kei Okada, and Masayuki Inaba. Musculoskeletal design, control, and application of human mimetic humanoid Kenshiro. *Bioinspiration & Biomimetics*, 14 (3):036011, April 2019. ISSN 1748-3190. doi: 10.1088/1748-3190/ab03fc. URL https://iopscience.iop.org/article/10.1088/1748-3190/ab03fc.

K.S. Aschenbeck, N.I. Kern, R.J. Bachmann, and R.D. Quinn. Design of a Quadruped Robot Driven by Air Muscles. In *The First IEEE/RAS-EMBS International Conference on Biomedical Robotics and Biomechatronics, 2006. BioRob 2006.*, pages 875–880, February 2006. doi: 10.1109/BIOROB.2006.1639201. URL https://ieeexplore.ieee.org/abstract/document/1639201.

Charles Audet and J. E. Dennis, Jr. Mesh adaptive direct search algorithms for constrained optimization. *SIAM Journal on Optimization*, 16(3):988–1015, 2006. doi: 10.1137/040603371.

B. P. Bolen and A. J. Hunt. Torque Values About Artificial Knee Using Braided Pneumatic Actuators, June 2021. URL https://www.braininitiative.org/event/the-7th-annual-brain-initiative-investigators-meeting/. Poster 4008. Work funded by NSF NeuroNex C3NS #2015317. Includes bipedal walker CPG results (Bipedal Walker and CPG Architecture sections).

B. P. Bolen and A. J. Hunt. Using AI Tools to Enhance the Sensory Afferent Database: A Comparative Study, October 2024. URL https://www.sfn.org/Meetings/Neuroscience-2024. Poster presentation, Neuroscience 2024, Society for Neuroscience, Chicago, IL.

B. P. Bolen, K. Deng, and A. J. Hunt. Sensory Afferent Database: Advancing towards a comprehensive understanding of sensory afferents and motor control, November 2023. URL https://neuronex5.io/. Poster presentation, Neuroscience 2023, Society for Neuroscience, Washington, D.C.

Ben Bolen. Bipedal Robot Legs Project. MSME project paper, Portland State University, Portland, OR, 2019.

Ben Bolen, J. Stu McNeal, Sunil Williams, and Alexander J. Hunt. Self-balancing robot leg with vestibular feedback, April 2024a. URL https://nwrsymposium.github.io/2024/. Poster presentation, Northwest Robotics Symposium (NWRS), April 26, Corvallis, OR, United States.

Ben Bolen, Lindie Burgess, Connor Morrow, Mohammad Elzein, Lawrence Pang, Cody Scharzenberger, and Alexander Hunt. Isometric Torque Values About Robotic Knee

Using Braided Pneumatic Actuators. Unpublished manuscript in preparation; not submitted for publication, 2025.

Ben Bolen, Mohammad Elzein, Lawrence Pang, and Alexander Hunt. Isometric Force Characterization of Braided Pneumatic Actuators. *Actuators*, 15(5):230, 2026. doi: 10.3390/act15050230.

Ben P. Bolen and Alexander J. Hunt. Determination of Artificial Muscle Placement for Biomimetic Humanoid Robot Legs. In Uriel Martinez-Hernandez, Vasiliki Vouloutsi, Anna Mura, Michael Mangan, Minoru Asada, Tony J. Prescott, and Paul F.M.J. Verschure, editors, *Biomimetic and Biohybrid Systems*, volume 11556, pages 15–26. Springer International Publishing, Cham, 2019. ISBN 978-3-030-24740-9 978-3-030-24741-6. doi: 10.1007/978-3-030-24741-6_2. URL http://link.springer.com/10.1007/978-3-030-24741-6_2.

Ben P. Bolen, J. Stu McNeal, Sunil Williams, and Alexander J. Hunt. Balance Platform and Inverted Pendulum for Robotic Vestibular and Proprioceptive Control, May 2024b. URL https://nwbs2024.wixsite.com/uoregon. Abstract and poster presentation, Northwest Biomechanics Symposium 2024, May 17–18, Eugene, OR.

Ben P. Bolen, J. Stu McNeal, Sunil Williams, and Alexander J. Hunt. Self-balancing Robot Leg, May 2024c. URL https://pdxscholar.library.pdx.edu/studentsymposium/2024/posters/60/. Poster presentation, Portland State University Student Research Symposium, May 8, Portland, OR.

R. A. Brand, R. D. Crowninshield, C. E. Wittstock, D. R. Pedersen, C. R. Clark, and F. M. van Krieken. A Model of Lower Extremity Muscular Anatomy. *Journal of Biomechanical Engineering*, 104(4):304–310, November 1982. ISSN 0148-0731, 1528-8951. doi: 10.1115/1.3138363. URL https://asmedigitalcollection.asme.org/biomechanical/article/104/4/304/453784/A-Model-of-Lower-Extremity-Muscular-Anatomy.

T. Graham Brown. The intrinsic factors in the act of progression in the mammal. *Proceedings of the Royal Society B: Biological Sciences*, 84(572):308–319, 1911. doi: 10.1098/rspb.1911.0077.

Vittorio Caggiano, Huawei Wang, Guillaume Durandau, Massimo Sartori, and Vikash Kumar. MyoSuite: A Contact-rich Simulation Suite for Musculoskeletal Motor Control. In *Proceedings of The 4th Annual Learning for Dynamics and Control Conference*, volume 168 of *Proceedings of Machine Learning Research*, pages 492–507. PMLR, 2022.

C. Capaday and R. B. Stein. Amplitude modulation of the soleus H-reflex in the human during walking and standing. *The Journal of Neuroscience*, 6(5):1308–1313, 1986. doi: 10.1523/JNEUROSCI.06-05-01308.1986.

Jiahao Chen, Shanlin Zhong, Erlong Kang, and Hong Qiao. Realizing human-like manipulation with a musculoskeletal system and biologically inspired control scheme. *Neurocomputing*, 339:116–129, April 2019. ISSN 0925-2312. doi: 10.1016/j.neucom.2018.12.069. URL https://www.sciencedirect.com/science/article/pii/S0925231218315315.

Z. Chen, Y. Yu, F. Han, and K. Deng. A cortico-cerebellar–spinal hierarchical model for adaptive step-length control in bipedal locomotions. *The European Physical Journal Special Topics*, 234:9129–9145, 2025. doi: 10.1140/epjs/s11734-025-02028-6.

Ching-Ping Chou and B. Hannaford. Measurement and modeling of McKibben pneumatic artificial muscles. *IEEE Transactions on Robotics and Automation*, 12(1):90–102, February 1996. ISSN 2374-958X. doi: 10.1109/70.481753. URL https://ieeexplore.ieee.org/document/481753/.

David Cofer, Gennady Cymbalyuk, James Reid, Ying Zhu, W. J. Heitler, and Donald H. Edwards. AnimatLab: A 3D graphics environment for neuromechanical simulations. *Journal of Neuroscience Methods*, 187(2):280–288, 2010. doi: 10.1016/j.jneumeth.2010.02.005.

Simon M. Danner, S. D. Wilshin, Natalia A. Shevtsova, and Ilya A. Rybak. Central control of interlimb coordination and speed-dependent gait expression in quadrupeds. *The Journal of Physiology*, 594(23):6947–6967, 2016. doi: 10.1113/JP272787.

Friedl De Groote and Antoine Falisse. Perspective on musculoskeletal modelling and predictive simulations of human movement to assess the neuromechanics of gait. *Proceedings of the Royal Society B: Biological Sciences*, 288(1946):20202432, 2021. doi: 10.1098/rspb.2020.2432.

Kalyanmoy Deb, Amrit Pratap, Sameer Agarwal, and T. Meyarivan. A fast and elitist multiobjective genetic algorithm: NSGA-II. *IEEE Transactions on Evolutionary Computation*, 6 (2):182–197, 2002. doi: 10.1109/4235.996017.

Scott L. Delp, Frank C. Anderson, Allison S. Arnold, Peter Loan, Ayman Habib, Chand T. John, Eran Guendelman, and Darryl G. Thelen. OpenSim: Open-Source Software to Create and Analyze Dynamic Simulations of Movement. *IEEE Transactions on Biomedical Engineering*, 54(11):1940–1950, November 2007. ISSN 0018-9294. doi: 10.1109/TBME.2007.901024. URL http://ieeexplore.ieee.org/document/4352056/.

S.L. Delp, J.P. Loan, M.G. Hoy, F.E. Zajac, E.L. Topp, and J.M. Rosen. An interactive graphics-based model of the lower extremity to study orthopaedic surgical procedures. *IEEE Transactions on Biomedical Engineering*, 37(8):757–767, August 1990. ISSN 00189294. doi: 10.1109/10.102791. URL http://ieeexplore.ieee.org/document/102791/.

Kaiyu Deng, Nicholas S. Szczecinski, David Arnold, Emanuel Andrada, Martin S. Fischer, Roger D. Quinn, and Alexander J. Hunt. Neuromechanical model of rat hindlimb walking with two-layer CPGs. *Biomimetics*, 4(1):21, 2019. doi: 10.3390/biomimetics4010021.

Kaiyu Deng et al. Biomechanical and sensory feedback regularize the behavior of different locomotor central pattern generators. *Biomimetics*, 7(4):226, 2022. doi: 10.3390/biomimetics7040226.

Timothy R. Derrick, Antonie J. van den Bogert, Andrea Cereatti, Raphael Dumas, Silvia Fantozzi, and Alberto Leardini. ISB recommendations on the reporting of intersegmental forces and moments during human motion analysis. *Journal of Biomechanics*, 99:109533, January 2020. ISSN 0021-9290. doi: 10.1016/j.jbiomech.2019.109533. URL https://www.sciencedirect.com/science/article/pii/S0021929019307870.

A. Di Russo et al. Investigating the roles of reflexes and central pattern generators in the control and modulation of human locomotion using a physiologically plausible neuromechanical model. *Journal of Neural Engineering*, 20(6):066006, 2023. doi: 10.1088/1741-2552/acfdcc.

Volker Dietz and Jacques Duysens. Significance of load receptor input during locomotion: a review. *Gait & Posture*, 11(2):102–110, 2000. doi: 10.1016/S0966-6362(99)00052-1.

L. E. Domínguez-Rodríguez, K. Stecina, D. L. García-Ramírez, E. Mena-Avila, J. J. Milla-Cruz, L. Martínez-Silva, M. Zhang, H. Hultborn, and J. N. Quevedo. Candidate interneurons mediating the resetting of the locomotor rhythm by extensor group I afferents in the cat. *Neuroscience*, 450:96–112, 2020. doi: 10.1016/j.neuroscience.2020.09.017.

W. Duan, W. Chen, J. Wang, J. Liu, Z. Pei, and J. Chen. Gait-generation strategy for lower limb exoskeleton based on central pattern generator. *IEEE/ASME Transactions on Mechatronics*, 29(6):4191–4200, 2024. URL https://ieeexplore.ieee.org/document/10460752.

Rosamond M. Eccles and A. Lundberg. Integrative pattern of Ia synaptic actions on motoneurones of hip and knee muscles. *The Journal of Physiology*, 144(2):271–298, 1958. doi: 10.1113/jphysiol.1958.sp006101.

S. A. Edgley and E. Jankowska. An interneuronal relay for group I and II muscle afferents in the midlumbar segments of the cat spinal cord. *The Journal of Physiology*, 389(1): 647–674, 1987. doi: 10.1113/jphysiol.1987.sp016676.

Mohammad Elzein, Jack Lutz, Ben Bolen, and Alexander Hunt. Pressure and Force Dynamics in Artificial Muscle Actuators: A State-Space and Optimization-Based Approach. In *Biomimetic and Biohybrid Systems*, pages 592–605. Springer Nature Switzerland, Cham, 2025. ISBN 978-3-032-07448-5. doi: 10.1007/978-3-032-07448-5_49.

Festo AG & Co. KG. Fluidic Muscle DMSP/MAS. Product Datasheet, Festo AG & Co. KG, April 2013. URL https://ftp.festo.com/Public/PNEUMATIC/SOFTWARE_SERVICE/Documentation/2015/EN/DMSP-MAS_EN.PDF.

Festo AG & Co. KG. Fluidic Muscle DMSP/MAS. Product Datasheet, Festo AG & Co. KG, January 2018. URL https://www.festo.com/net/en_us/SupportPortal/Files/732773/DMSP_MAS_2018-01f_8081496g1.pdf.

Festo AG & Co. KG. Fluidic Muscle DMSP. Product Datasheet, Festo AG & Co. KG, January 2026. URL https://www.festo.com/media/catalog/202851_documentation.pdf.

H. Forssberg. Stumbling corrective reaction: a phase-dependent compensatory reaction during locomotion. *Journal of Neurophysiology*, 42(4):936–953, 1979. doi: 10.1152/jn.1979.42.4.936.

Jonas Frey, Roi Mattamala, Hae-Won Park, Mayank Mittal, Georg Martius, Michael A. Osborne, Robert Sparrow, and Marco Hutter. Advances, challenges, and opportunities for legged robots. *Science Robotics*, 11:aee0787, 2026. doi: 10.1126/scirobotics.aee0787.

Yasuhiro Fukuoka, Yasuyuki Habu, and Ning Inoue. Autonomous speed adaptation by a muscle-driven hind leg robot modeled on a cat without intervention from brain. *International Journal of Advanced Robotic Systems*, 18, 2021. doi: 10.1177/17298814211044936.

Thomas Geijtenbeek. SCONE: Open Source Software for Predictive Simulation of Biological Motion. *Journal of Open Source Software*, 4(38):1421, 2019. doi: 10.21105/joss.01421.

Hartmut Geyer and Hugh Herr. A muscle-reflex model that encodes principles of legged mechanics produces human walking dynamics and muscle activities. *IEEE Transactions on Neural Systems and Rehabilitation Engineering*, 18(3):263–273, 2010. doi: 10.1109/TNSRE.2010.2047592.

C. A. Goldsmith, R. D. Quinn, and N. S. Szczecinski. Investigating the role of low level reinforcement reflex loops in insect locomotion. *Bioinspiration & Biomimetics*, 16(6): 065008, October 2021. ISSN 1748-3190. doi: 10.1088/1748-3190/ac28ea. URL https://dx.doi.org/10.1088/1748-3190/ac28ea.

Ryan Grummer, Christian B. Jackson, and Roger D. Quinn. Multilevel synthetic nervous system control for legged locomotion. In *Biomimetic and Biohybrid Systems*, volume 14930 of *Lecture Notes in Computer Science*, pages 382–396. Springer, 2024. doi: 10.1007/978-3-031-72597-5_26.

Sehoon Ha et al. Learning-based legged locomotion: State of the art and future challenges. *The International Journal of Robotics Research*, 2025. doi: 10.1177/02783649241312698.

Daniel F. B. Haeufle, Benjamin Schmortte, Hartmut Geyer, and Andre Seyfarth. The benefit of combining neuronal feedback and feed-forward control for robustness in step down perturbations of simulated human walking depends on the muscle function. *Frontiers in Computational Neuroscience*, 12:80, 2018. doi: 10.3389/fncom.2018.00080.

Jun He and Feng Gao. Mechanism, Actuation, Perception, and Control of Highly Dynamic Multilegged Robots: A Review. *Chinese Journal of Mechanical Engineering*, 33(1):79, November 2020. ISSN 2192-8258. doi: 10.1186/s10033-020-00485-9. URL https://doi.org/10.1186/s10033-020-00485-9.

G. W. Hiebert, P. J. Whelan, A. Prochazka, and K. G. Pearson. Contribution of hind limb flexor muscle afferents to the timing of phase transitions in the cat step cycle. *Journal of Neurophysiology*, 75(3):1126–1137, 1996. doi: 10.1152/jn.1996.75.3.1126.

Wade Hilts, Nicholas S. Szczecinski, Roger D. Quinn, and Alexander J. Hunt. A Dynamic Neural Network Designed Using Analytical Methods Produces Dynamic Control Properties Similar to Human Test Subjects. *IEEE Control Systems Letters*, 3(2):320–325, 2019. doi: 10.1109/LCSYS.2018.2871126.

Arne Hitzmann, Hiroaki Masuda, Shuhei Ikemoto, and Koh Hosoda. Anthropomorphic musculoskeletal 10 degrees-of-freedom robot arm driven by pneumatic artificial muscles. *Advanced Robotics*, 32(15):865–878, August 2018. ISSN 01691864. doi: 10.1080/01691864.2018.1494040. URL http://stats.lib.pdx.edu/proxy.php?url=http://search.ebscohost.com/login.aspx?direct=true&db=aph&AN=131926820&site=ehost-live.

Melissa Hoy, Felix Zajac, and Michael Gordon. A musculoskeletal model of the human lower extremity: the effect of muscle, tendon, and moment arm on the moment-angle re-

lationship of musculotendon actuators at the hip, knee, and ankle. *Journal of Biomechanics*, 23(2):157–169, 1990.

Alexander Hunt, Alexander Graber-Tilton, and Roger Quinn. Modeling length effects of braided pneumatic actuators. In *Volume 5A: 41st Mechanisms and Robotics Conference*, volume 58172, page V05AT08A008, Cleveland, OH, August 2017a. ASME. ISBN 978-0-7918-5817-2. doi: 10.1115/detc2017-67458. URL https://asme.pinetec.com/detc2017/data/pdfs/trk-8/DETC2017-67458.pdf.

Alexander Hunt, Nicholas Szczecinski, and Roger Quinn. Development and Training of a Neural Controller for Hind Leg Walking in a Dog Robot. *Frontiers in Neurorobotics*, Volume 11 - 2017, April 2017b. ISSN 1662-5218. doi: 10.3389/fnbot.2017.00018. URL http://journal.frontiersin.org/article/10.3389/fnbot.2017.00018/full.

Daisuke Ichimura and Tetsuro Yamazaki. A pathological condition affects motor modules in a bipedal locomotion model. *Frontiers in Neurorobotics*, 13:79, 2019. doi: 10.3389/fnbot.2019.00079.

Auke J. Ijspeert. Amphibious and Sprawling Locomotion: From Biology to Robotics and Back. *Annual Review of Control, Robotics, and Autonomous Systems*, 3(1):173–193, 2020. doi: 10.1146/annurev-control-091919-095731. URL https://doi.org/10.1146/annurev-control-091919-095731. _eprint: https://doi.org/10.1146/annurev-control-091919-095731.

Marco Iosa, Andrea Cereatti, Antonio Merlo, Ilaria Campanini, Stefano Paolucci, and Aurelio Cappozzo. Assessment of waveform similarity in clinical gait data: the linear fit method. *BioMed Research International*, 2014:214156, 2014. doi: 10.1155/2014/214156.

E. Jankowska, M. G. M. Jukes, S. Lund, and A. Lundberg. The effect of DOPA on the spinal cord 5. Reciprocal organization of pathways transmitting excitatory action to alpha motoneurones of flexors and extensors. *Acta Physiologica Scandinavica*, 70(3-4): 369–388, 1967. doi: 10.1111/j.1748-1716.1967.tb03636.x.

E. Jankowska, T. Johannisson, and J. Lipski. Common interneurones in reflex pathways from group 1a and 1b afferents of ankle extensors in the cat. *The Journal of Physiology*, 310(1):381–402, 1981. doi: 10.1113/jphysiol.1981.sp013556.

Shunichi Kurumaya, Koichi Suzumori, Hiroyuki Nabae, and Shuichi Wakimoto. Musculoskeletal lower-limb robot driven by multifilament muscles. *ROBOMECH Journal*, 3 (1):18, December 2016. ISSN 2197-4225. doi: 10.1186/s40648-016-0061-3. URL https://robomechjournal.springeropen.com/articles/10.1186/s40648-016-0061-3.

Tania Lam and Keir G. Pearson. The role of proprioceptive feedback in the regulation and adaptation of locomotor activity. *Advances in Experimental Medicine and Biology*, 508: 343–355, 2002. doi: 10.1007/978-1-4615-0713-0_40.

Lang, Robert. MuscleSIM, 2005. Dept. TK-ES / Robert Lang.

Ronald Leibach. *Development of a Tunable Compliance Energy Return Actuator*. PhD thesis, Case Western Reserve University, 2020. URL https://etd.ohiolink.edu/apexprod/rws_olink/r/1501/10?p10_etd_subid=184456&clear=10.

Jianfeng Li, Shiping Zuo, Chenghui Xu, Leiyu Zhang, Mingjie Dong, Chunjing Tao, and Run Ji. Influence of a Compatible Design on Physical Human-Robot Interaction Force: a Case Study of a Self-Adapting Lower-Limb Exoskeleton Mechanism. *Journal of Intelligent & Robotic Systems*, 98(2):525–538, May 2020. ISSN 0921-0296, 1573-0409. doi: 10.1007/s10846-019-01063-5. URL http://link.springer.com/10.1007/s10846-019-01063-5.

Wansheng Li, Nicholas S. Szczecinski, and Roger D. Quinn. A neural network with central pattern generators entrained by sensory feedback controls walking of a bipedal model. *Bioinspiration & Biomimetics*, 12(6):065002, 2017. doi: 10.1088/1748-3190/aa8290.

Yinlong Li, Yinlai Jiang, and Kou Hosoda. Controller design and experimental validation of walking for a musculoskeletal bipedal lower limb robot based on the spring-loaded inverted pendulum model. *Frontiers in Robotics and AI*, 11:1449721, 2024. doi: 10.3389/frobt.2024.1449721.

Wei Liang, Hao Liu, Kunyang Wang, Zhihui Qian, Luquan Ren, and Lei Ren. Comparative study of robotic artificial actuators and biological muscle. *Advances in Mechanical Engineering*, 12(6):1687814020933409, June 2020. ISSN 1687-8140. doi: 10.1177/1687814020933409. URL https://doi.org/10.1177/1687814020933409.

Xiangxiao Liu, Andre Rosendo, Shuhei Ikemoto, Masahiro Shimizu, and Koh Hosoda. Robotic investigation on effect of stretch reflex and crossed inhibitory response on bipedal hopping. *Journal of the Royal Society Interface*, 15(140):20180024, March 2018. ISSN 1742-5689. doi: 10.1098/rsif.2018.0024. Place: London WOS:000428576200018.

M. J. Lutz. Mechanical improvements and control implementation on a biomimetic quadruped model. Master's thesis, Portland State University, 2025.

J.D.W. Madden, N.A. Vandesteeg, P.A. Anquetil, P.G.A. Madden, A. Takshi, R.Z. Pytel, S.R. Lafontaine, P.A. Wieringa, and I.W. Hunter. Artificial muscle technology: physical principles and naval prospects. *IEEE Journal of Oceanic Engineering*, 29(3):706–728, July

2004. ISSN 1558-1691. doi: 10.1109/JOE.2004.833135. Conference Name: IEEE Journal of Oceanic Engineering.

Sergey N. Markin, Alexander N. Klishko, Natalia A. Shevtsova, Michel A. Lemay, Boris I. Prilutsky, and Ilya A. Rybak. Afferent control of locomotor CPG: insights from a simple neuromechanical model. *Annals of the New York Academy of Sciences*, 1198(1):21–34, 2010. doi: 10.1111/j.1749-6632.2010.05435.x.

Martin Martens, Alexander Passon, and Ingo Boblan. A Sensor-Less Torque Control Interface for Pneumatic Muscle Actuators. *Actuators*, 7(4):82, 2018. doi: 10.3390/act7040082.

Mirco Martens and Ivo Boblan. Modeling the Static Force of a Festo Pneumatic Muscle Actuator: A New Approach and a Comparison to Existing Models. *Actuators*, 6(4):33, November 2017. doi: 10.3390/act6040033. URL http://www.mdpi.com/2076-0825/6/4/33.

N. Mazzaro, M. J. Grey, O. Feix do Nascimento, and T. Sinkjær. Afferent-mediated modulation of the soleus muscle activity during the stance phase of human walking. *Experimental Brain Research*, 173(4):713–723, 2006. doi: 10.1007/s00221-006-0451-5.

D. A. McCrea, S. J. Shefchyk, M. J. Stephens, and K. G. Pearson. Disynaptic group I excitation of synergist ankle extensor motoneurones during fictive locomotion in the cat. *The Journal of Physiology*, 487(2):527–539, 1995. doi: 10.1113/jphysiol.1995.sp020897.

David A. McCrea and Ilya A. Rybak. Organization of mammalian locomotor rhythm and pattern generation. *Brain Research Reviews*, 57(1):134–146, 2008. doi: 10.1016/j.brainresrev.2007.08.006.

Luciano Luporini Menegaldo, Agenor de Toledo Fleury, and Hans Ingo Weber. Moment arms and musculotendon lengths estimation for a three-dimensional lower-limb model. *Journal of Biomechanics*, 37(9):1447–1453, September 2004. ISSN 0021-9290. doi: 10.1016/j.jbiomech.2003.12.017. URL https://www.sciencedirect.com/science/article/pii/S0021929003004809.

Matthew Millard, Thomas Uchida, Ajay Seth, and Scott L. Delp. Flexing Computational Muscle: Modeling and Simulation of Musculotendon Dynamics. *Journal of Biomechanical Engineering*, 135(2):021005, February 2013. ISSN 0148-0731, 1528-8951. doi: 10.1115/1.4023390. URL https://asmedigitalcollection.asme.org/biomechanical/article/doi/10.1115/1.4023390/371394/Flexing-Computational-Muscle-Modeling-and.

Yaroslav I. Molkov, Bartholomew J. Bacak, Adolfo E. Talpalar, and Ilya A. Rybak. Mechanisms of Left-Right Coordination in Mammalian Locomotor Pattern Generation Circuits: A Mathematical Modeling View. *PLOS Computational Biology*, 11(6):e1004270, 2015. doi: 10.1371/journal.pcbi.1004270.

Connor Morrow, Benjamin Bolen, and Alexander Hunt. Optimization of Artificial Muscle Placements for a Humanoid Bipedal Robot. In *Biomimetic and Biohybrid Systems*, volume 12413 of *Lecture Notes in Computer Science*, pages 257–269. Springer, Cham, July 2020. ISBN 978-3-030-64312-6. doi: 10.1007/978-3-030-64313-3_25. URL https://doi.org/10.1007/978-3-030-64313-3_25.

J. E. Nash and J. V. Sutcliffe. River flow forecasting through conceptual models part I — A discussion of principles. *Journal of Hydrology*, 10(3):282–290, 1970. doi: 10.1016/0022-1694(70)90255-6.

K. Nishizaki, R. Adachi, Y. Ambe, S. Aoi, et al. Gait transition mechanism from quadrupedal to bipedal locomotion in the Japanese macaque based on inverted pendulum. *eLife*, 2026. doi: 10.7554/eLife.109826.1. Reviewed preprint.

William R. P. Nourse, Nicholas S. Szczecinski, Alexander J. Hunt, and Roger D. Quinn. Analyzing the Interplay Between Local CPG Activity and Sensory Signals for Inter-leg Coordination in Stick Insect. In *Biomimetic and Biohybrid Systems*, pages 342–345. Springer International Publishing, 2019. doi: 10.1007/978-3-030-24741-6_34.

William R. P. Nourse, Axel Jackson, Nicholas S. Szczecinski, and Roger D. Quinn. SNS-Toolbox: An Open Source Tool for Designing Synthetic Nervous Systems and Interfacing Them with Cyber-Physical Systems. *Biomimetics*, 8(2):247, 2023. doi: 10.3390/biomimetics8020247.

K. G. Pearson, J. E. Misiaszek, and K. Fouad. Enhancement and resetting of locomotor activity by muscle afferents. *Annals of the New York Academy of Sciences*, 860(1):203–215, 1998. doi: 10.1111/j.1749-6632.1998.tb09050.x.

H. Pery, T. Ben Eli, and E. Holdengreber. Benchmarking open-loop central pattern generator controllers for lower-limb exoskeleton walking trajectories. *IEEE Access*, 2026. URL https://ieeexplore.ieee.org/document/11523455.

C. A. Pratt and L. M. Jordan. Ia inhibitory interneurons and Renshaw cells as contributors to the spinal mechanisms of fictive locomotion. *Journal of Neurophysiology*, 57(1):56–71, 1987. doi: 10.1152/jn.1987.57.1.56.

A. Prochazka and M. Gorassini. Ensemble firing of muscle afferents recorded during normal locomotion in cats. *The Journal of Physiology*, 507(1):293–304, 1998. doi: 10.1111/j.1469-7793.1998.293bu.x.

Vladimir Rancic and Simon Gosgnach. Recent Insights into the Rhythmogenic Core of the Locomotor CPG. *International Journal of Molecular Sciences*, 22(3):1394, 2021. doi: 10.3390/ijms22031394.

Neider Romero, Edwin Rua Ramirez, and Gonzalo Moreno. FOUR-BAR MECHANISM'S ROTATION CENTER DETERMINATION USING NATURAL COORDINATES. *Journal of Engineering and Applied Sciences*, 15, November 2020.

William M. S. Russell and Rex L. Burch. *The Principles of Humane Experimental Technique*. Methuen & Co. Ltd, London, 1959. Reissued 1992 by the Universities Federation for Animal Welfare.

Ilya A. Rybak, Natalia A. Shevtsova, Myriam Lafreniere-Roula, and David A. McCrea. Modelling spinal circuitry involved in locomotor pattern generation: insights from deletions during fictive locomotion. *The Journal of Physiology*, 577(Pt 2):627–639, December 2006. ISSN 0022-3751. doi: 10.1113/jphysiol.2006.118703.

Ilya A. Rybak, Kevin J. Dougherty, and Natalia A. Shevtsova. Organization of the Mammalian Locomotor CPG: Review of Computational Model and Circuit Architectures Based on Genetically Identified Spinal Interneurons. *eNeuro*, 2(5):ENEURO.0069–15.2015, 2015. doi: 10.1523/ENEURO.0069-15.2015.

József Sárosi and Zoltán Fabulya. Mathematical analysis of the function approximation for the force generated by pneumatic artificial muscle. *Scientific Bulletin of the Politehnica University of Timişoara, Transactions on Mechanics*, pages 59–64, 2012.

Cody Warren Scharzenberger. Design of a Canine Inspired Quadruped Robot as a Platform for Synthetic Neural Network Control. Master's thesis, Portland State University, Digital, May 2019. URL `https://www.proquest.com/docview/2289587100?pq-origsite=gscholar&fromopenview=true`.

Malte Schilling, Jan Paskarbeit, Helge Ritter, Axel Schneider, and Holk Cruse. From Adaptive Locomotion to Predictive Action Selection – Cognitive Control for a Six-Legged Walker. *IEEE Transactions on Robotics*, 38(2):666–682, April 2022. ISSN 1941-0468. doi: 10.1109/TRO.2021.3106832. Conference Name: IEEE Transactions on Robotics.

E. D. Schomburg, N. Petersen, I. Barajon, and H. Hultborn. Flexor reflex afferents reset the step cycle during fictive locomotion in the cat. *Experimental Brain Research*, 122(3): 339–350, 1998. doi: 10.1007/s002210050522.

V. A. Selionov, Yuri P. Ivanenko, Irina A. Solopova, and Victor S. Gurfinkel. Tonic Central and Sensory Stimuli Facilitate Involuntary Air-Stepping in Humans. *Journal of Neurophysiology*, 101(6):2847–2858, April 2009. doi: 10.1152/jn.90895.2008.

Ajay Seth, Michael Sherman, Jeffrey A. Reinbolt, and Scott L. Delp. OpenSim: a musculoskeletal modeling and simulation framework for in silico investigations and exchange., 2011.

Ajay Seth, Jennifer L. Hicks, Thomas K. Uchida, Ayman Habib, Christopher L. Dembia, James J. Dunne, Carmichael F. Ong, Matthew S. DeMers, Apoorva Rajagopal, Matthew Millard, Samuel R. Hamner, Edith M. Arnold, Jennifer R. Yong, Shrinidhi K. Lakshmikanth, Michael A. Sherman, Joy P. Ku, and Scott L. Delp. OpenSim: Simulating musculoskeletal dynamics and neuromuscular control to study human and animal movement. *PLOS Computational Biology*, 14(7):e1006223, July 2018. ISSN 1553-7358. doi: 10.1371/journal.pcbi.1006223. URL `https://dx.plos.org/10.1371/journal.pcbi.1006223`.

F. V. Severin, G. N. Orlovsky, and M. L. Shik. [work of the muscles during controlled locomotion]. *Biofizika*, 1967. In Russian; English translation in Biophysics. Verify exact bibliographic details against Stuart (2007).

Michael A. Sherman, Ajay Seth, and Scott L. Delp. What is a Moment Arm? Calculating Muscle Effectiveness in Biomechanical Models Using Generalized Coordinates. In *Volume 7B: 9th International Conference on Multibody Systems, Nonlinear Dynamics, and Control*, page V07BT10A052, Portland, Oregon, USA, August 2013. American Society of Mechanical Engineers. ISBN 978-0-7918-5597-3. doi: 10.1115/DETC2013-13633. URL `https://asmedigitalcollection.asme.org/IDETC-CIE/proceedings/IDETC-CIE2013/55973/Portland,%20Oregon,%20USA/256002`.

C. S. Sherrington. Flexion-reflex of the limb, crossed extension-reflex, and reflex stepping and standing. *The Journal of Physiology*, 40(1-2):28–121, 1910. doi: 10.1113/jphysiol.1910.sp001362.

Natalia A. Shevtsova, Khaldoun Hamade, Samit Chakrabarty, Sergey N. Markin, Boris I. Prilutsky, and Ilya A. Rybak. Modeling the Organization of Spinal Cord Neural Circuits Controlling Two-Joint Muscles. In Boris I. Prilutsky and Donald H. Edwards, editors,

Neuromechanical Modeling of Posture and Locomotion, Springer Series in Computational Neuroscience, pages 121–162. Springer, New York, NY, 2016. ISBN 978-1-4939-3267-2. doi: 10.1007/978-1-4939-3267-2_5. URL https://doi.org/10.1007/978-1-4939-3267-2_5.

Hirofumi Shin, Shuhei Ikemoto, and Koh Hosoda. Constructive understanding and reproduction of functions of gluteus medius by using a musculoskeletal walking robot. *Advanced Robotics*, 32(4):202–214, February 2018. ISSN 0169-1864. doi: 10.1080/01691864.2018.1434015. URL https://www.tandfonline.com/doi/citedby/10.1080/01691864.2018.1434015.

Takaya Shin and Kiminao Kogiso. Sensorless angle and stiffness control of antagonistic PAM actuator using reference set. *Advanced Robotics*, 36(9):423–437, 2022. doi: 10.1080/01691864.2022.2046502.

Kota Shinohara, Yuichi Ambe, Yongi Kim, Fu Mano, Shravan Tata Ramalingasetty, Andrew B. Lockhart, Sergey N. Markin, Jessica Ausborn, Ilya A. Rybak, Simon M. Danner, and Shinya Aoi. Mechanisms of adaptive interlimb coordination to sudden ground loss: a neuromusculoskeletal modeling study. *bioRxiv*, page 2025.11.11.687930, 2025. doi: 10.1101/2025.11.11.687930.

Seungmoon Song, Łukasz Kidziński, Xue Bin Peng, Carmichael F. Ong, Jennifer Hicks, Sergey Levine, Christopher G. Atkeson, and Scott L. Delp. Deep reinforcement learning for modeling human locomotion control in neuromechanical simulation. *Journal of NeuroEngineering and Rehabilitation*, 18:126, 2021. doi: 10.1186/s12984-021-00919-y.

Alexander G Steele, Alexander Hunt, and Appolinaire C Etoundi. Development of a Bio-inspired Knee Joint Mechanism for a Bipedal Robot. In *Biomimetic and Biohybrid Systems*, Lecture Notes in Computer Science, pages 418–427. Springer, Cham, July 2017. ISBN 978-3-319-63536-1 978-3-319-63537-8. doi: 10.1007/978-3-319-63537-8_35. URL https://link.springer.com/chapter/10.1007/978-3-319-63537-8_35.

Alexander G. Steele, Alexander Hunt, and Appolinaire C. Etoundi. Biomimetic Knee Design to Improve Joint Torque and Life for Bipedal Robotics. In Manuel Giuliani, Tareq Assaf, and Maria Elena Giannaccini, editors, *Towards Autonomous Robotic Systems*, pages 91–102, Cham, 2018. Springer International Publishing. ISBN 978-3-319-96728-8.

Alexander Gabriel Steele. *Biomimetic Design and Construction of a Bipedal Walking Robot*. Thesis, Portland State University, Portland, Oregon, USA, 2018.

M. J. Stephens and J. F. Yang. Short latency, non-reciprocal group I inhibition is reduced during the stance phase of walking in humans. *Brain Research*, 743:24–31, 1996. doi: 10.1016/S0006-8993(96)00977-8.

Douglas G. Stuart. Reflections on integrative and comparative movement neuroscience. *Integrative and Comparative Biology*, 47(4):482–504, 2007. doi: 10.1093/icb/icm037.

Bethany L. Suderman and Anita N. Vasavada. Moving muscle points provide accurate curved muscle paths in a model of the cervical spine. *Journal of Biomechanics*, 45(2): 400–404, January 2012. ISSN 0021-9290. doi: 10.1016/j.jbiomech.2011.10.014. URL https://www.sciencedirect.com/science/article/pii/S0021929011006373.

Koichi Suzumori and Ahmad Athif Faudzi. Trends in hydraulic actuators and components in legged and tough robots: a review. *Advanced Robotics*, 32(9):458–476, May 2018. ISSN 0169-1864. doi: 10.1080/01691864.2018.1455606. URL https://doi.org/10.1080/01691864.2018.1455606. _eprint: https://doi.org/10.1080/01691864.2018.1455606.

Nicholas S. Szczecinski, Clarissa A. Goldsmith, William R. P. Nourse, and Roger D. Quinn. A perspective on the neuromorphic control of legged locomotion in past, present, and future insect-like robots. *Neuromorphic Computing and Engineering*, 3(2):023001, 2023. doi: 10.1088/2634-4386/acc04f.

József Sárosi. New approximation algorithm for the force of fluidic muscles. In *Applied Computational Intelligence and Informatics (SACI), 2012 7th IEEE International Symposium on*, pages 229–233. IEEE, 2012. URL http://ieeexplore.ieee.org/xpls/abs_all.jsp?arnumber=6250007.

József Sárosi, Zoltán Fabulya, János Gyeviki, Gábor Keszthelyi-Szabó, and Péter Szendr Ho. Investigation of Accuracy of the Newest Function Approximation for the Force Generated by Pneumatic Artificial Muscle. *Analecta Technica Szegedinensia, Review of Faculty of Engineering*, 2013:1–2, 2013. URL http://www.analecta.hu/index.php/issues/201312/60-investigation-of-accuracy-of-the-newest-function-approximation-for-the-force-generfile.

József Sárosi, Ján Piteľ, Mária Tóthová, Alexander Hošovský, and István Bíró. COMPARATIVE SURVEY OF VARIOUS STATIC AND DYNAMIC MODELS OF PNEUMATIC ARTIFICIAL MUSCLES. *Transactions of the Canadian Society for Mechanical Engineering*, 41(5):825–844, December 2017. ISSN 0315-8977. doi: 10.1139/tcsme-2017-514. URL http://www.nrcresearchpress.com/doi/10.1139/tcsme-2017-514.

Jerrold Tannenbaum and B. Taylor Bennett. Russell and Burch's 3Rs then and now: The need for clarity in definition and purpose. *Journal of the American Association for Laboratory Animal Science*, 54(2):120–132, 2015.

Karl E. Taylor. Summarizing multiple aspects of model performance in a single diagram. *Journal of Geophysical Research: Atmospheres*, 106(D7):7183–7192, 2001. doi: 10.1029/2000JD900719.

The MathWorks, Inc. Optimization Toolbox, release r2025b, 2025. URL `https://www.mathworks.com/products/optimization.html`. Software documentation for `gamultiobj`, `surrogateopt`, and `patternsearch`.

Darryl G. Thelen. Adjustment of Muscle Mechanics Model Parameters to Simulate Dynamic Contractions in Older Adults. *Journal of Biomechanical Engineering*, 125(1):70–77, February 2003. ISSN 0148-0731, 1528-8951. doi: 10.1115/1.1531112. URL `https://asmedigitalcollection.asme.org/biomechanical/article/125/1/70/463793/Adjustment-of-Muscle-Mechanics-Model-Parameters-to`.

Emanuel Todorov, Tom Erez, and Yuval Tassa. MuJoCo: A physics engine for model-based control. In *2012 IEEE/RSJ International Conference on Intelligent Robots and Systems*, pages 5026–5033, October 2012. doi: 10.1109/IROS.2012.6386109.

Barkan Ugurlu, Paolo E. Forni, Hadi Dallali, Darwin G. Caldwell, and Nikos G. Tsagarakis. Stable Control of Force, Position, and Stiffness for a Pneumatic Muscle Actuator Joint. *IEEE Transactions on Industrial Informatics*, 15(2):1240–1251, 2019. doi: 10.1109/TII.2018.2849330.

UTHealth Houston School of Dentistry. 1155 Explicit Decerebrate Cat [video]. `https://www.youtube.com/watch?v=1nkczBM5YiM`, 2025. Classroom demonstration film. Accessed October 2026.

Henry W. A. A. Van de Crommert, Theo Mulder, and Jacques Duysens. Neural control of locomotion: sensory control of the central pattern generator and its relation to treadmill training. *Gait & Posture*, 7(3):251–263, 1998. doi: 10.1016/S0966-6362(98)00010-1.

Allan Joshua Veale and Shane Quan Xie. Towards compliant and wearable robotic orthoses: A review of current and emerging actuator technologies. *Medical Engineering & Physics*, 38(4):317–325, April 2016. ISSN 1350-4533. doi: 10.1016/j.medengphy.2016.01.010. URL `https://www.sciencedirect.com/science/article/pii/S135045331600031X`.

Huawei Wang, Vittorio Caggiano, Guillaume Durandau, Massimo Sartori, and Vikash Kumar. MyoSim: Fast and physiologically realistic MuJoCo models for musculoskeletal and exoskeletal studies. In *2022 IEEE International Conference on Robotics and Automation (ICRA)*, pages 10410–10417, 2022. doi: 10.1109/ICRA46639.2022.9811684.

Barbara Webb. Can robots make good models of biological behaviour? *Behavioral and Brain Sciences*, 24(6):1033–1050, 2001. doi: 10.1017/S0140525X01000127.

Barbara Webb. Robots in invertebrate neuroscience. *Nature*, 417(6886):359–363, 2002. doi: 10.1038/417359a.

WhiteCoatWaste. Decerebrated cat treadmill experiments [video]. `https://www.youtube.com/watch?v=hAteWL78NMc`, n.d. Archival footage of a decerebrate-cat locomotion experiment; restored and colorized copy archived with the dissertation materials. Accessed October 2026.

Uwe Windhorst. Muscle proprioceptive feedback and spinal networks. *Brain Research Bulletin*, 73(4):155–202, 2007. doi: 10.1016/j.brainresbull.2007.03.010.

Uwe Windhorst. Spinal Cord Circuits: Models and Reality. *Neurophysiology*, 2022. doi: 10.1007/s11062-022-09927-7.

Jack M. Winters. Hill-Based Muscle Models: A Systems Engineering Perspective. In Jack M. Winters and Savio L.-Y. Woo, editors, *Multiple Muscle Systems: Biomechanics and Movement Organization*, pages 63–93. Springer, New York, NY, 1990. doi: 10.1007/978-1-4613-9030-5_5.

Ge Wu and PeterR Cavanagh. FOR STANDARDIZATION IN THE REPORTING OF KINEMATIC DATA. *Journal Biomechanics*, 28(10):1257–1261, 1995. doi: 0021-9290(95)00017-8.

Ge Wu, Sorin Siegler, Paul Allard, Chris Kirtley, Alberto Leardini, Dieter Rosenbaum, Mike Whittle, Darryl D D'Lima, Luca Cristofolini, Hartmut Witte, Oskar Schmid, and Ian Stokes. ISB recommendation on definitions of joint coordinate system of various joints for the reporting of human joint motion—part I: ankle, hip, and spine. *Journal of Biomechanics*, 35(4):543–548, April 2002. ISSN 0021-9290. doi: 10.1016/S0021-9290(01)00222-6. URL `https://www.sciencedirect.com/science/article/pii/S0021929001002226`.

Ge Wu, Frans C. T. van der Helm, H. E. J. (DirkJan) Veeger, Mohsen Makhsous, Peter Van Roy, Carolyn Anglin, Jochem Nagels, Andrew R. Karduna, Kevin McQuade, Xuguang Wang, Frederick W. Werner, and Bryan Buchholz. ISB recommendation on

definitions of joint coordinate systems of various joints for the reporting of human joint motion—Part II: shoulder, elbow, wrist and hand. *Journal of Biomechanics*, 38 (5):981–992, May 2005. ISSN 0021-9290. doi: 10.1016/j.jbiomech.2004.05.042. URL https://www.sciencedirect.com/science/article/pii/S002192900400301X.

Gary T. Yamaguchi and Felix E. Zajac. A planar model of the knee joint to characterize the knee extensor mechanism. *Journal of Biomechanics*, 22(1):1–10, January 1989. ISSN 0021-9290. doi: 10.1016/0021-9290(89)90179-6. URL http://www.sciencedirect.com/science/article/pii/0021929089901796.

Fletcher Young, Christian Rode, Alex Hunt, and Roger Quinn. Analyzing Moment Arm Profiles in a Full-Muscle Rat Hindlimb Model. *Biomimetics*, 4(1):10, March 2019. doi: 10.3390/biomimetics4010010. URL https://www.mdpi.com/2313-7673/4/1/10.

Felix E. Zajac and Michael E. Gordon. Determining Muscle's Force and Action in Multi-Articular Movement. *Exercise and Sport Sciences Reviews*, 17(1):187–230, January 1989. ISSN 0091-6331. URL https://journals.lww.com/acsm-essr/Citation/1989/00170/Determining_Muscle_s_Force_and_Action_in.9.aspx.

Altynai Zhagiparova, Viktor Golubev, and Daekyu Kim. Recent developments in pneumatic artificial muscle actuators. *Actuators*, 14(12):582, 2025. doi: 10.3390/act14120582.

APPENDICES

Appendix A: Muscle Force and Joint Torque Calculations

This appendix documents the reference implementation of the force, torque, and compliance equations used in Chapters 3 and 4. When present in MATLAB code, file names are given so that any number in this dissertation can be traced to the code that produced it.

A.1 Human muscle force calculations

A.1.1 OpenSim Thelen2003 and Millard2012 Formulations

OpenSim's Thelen2003Muscle comprises a contractile element (first-order activation dynamics with distinct rise and fall time constants; active force–length curve; Hill-type force–velocity curve) in series with a compliant tendon, plus a parallel passive fiber element, with fiber pennation. Unlike the MuJoCo actuator, the fiber length is *not* the musculotendon length: at every instant the model solves the force equilibrium $F_{\text{tendon}}(L_{\text{mt}} - L_{\text{fiber}} \cos\alpha) = F_{\text{fiber}}(L_{\text{fiber}}, \dot{L}_{\text{fiber}}, a)$, so tendon compliance filters and delays force transmission and stores energy. The tendon is a quadratic-then-linear "strip" curve anchored at the slack length L_s (force zero until L_{mt} stretches the tendon beyond L_s). Millard2012EquilibriumModel, OpenSim's newer default, retains the same equilibrium structure but replaces the curve shapes (an asymmetric Gaussian active-force–length curve with a plateau, an exponentially stiffening tendon) and adds fiber damping; it is numerically better behaved at extreme lengths. Both OpenSim models therefore differ from the

converted MuJoCo model in the same single respect: series-elastic compliance.

A.1.2 OpenSim Moment Arms and Human Torque Targets

The human torque targets of Section 4.1.5 are OpenSim tendon-excursion moment arms: for each muscle, the moment arm about the knee is the partial derivative of muscle-tendon length with respect to the knee coordinate, $r(\theta_k) = \partial l_{\mathrm{MT}} / \partial \theta_k$, evaluated across the range of motion. The exported analysis tables archived with the test data (`OpenSim_Vasti_MomentArm.txt`, `OpenSim_Vasti_Results.txt`, `OpenSim_Bifem_MomentArm.txt`, `OpenSim_Bifem_Results.txt`) carry the per-muscle moment arms and the corresponding maximum-isometric torque curves (moment arm multiplied by each muscle's maximum isometric force) for vastus intermedius, lateralis, and medialis, and for the biceps femoris short head; the optimization loads the vastus medialis and biceps femoris short-head columns as its targets (`buildKneeExtContext20mm.m` lines 119–140; `Knee_Flexor_Data_20mm.m` line 546). This derivative definition and the geometric perpendicular-foot construction of Equation (A4) below are equivalent for a straight muscle path crossing a planar revolute joint; the derivative form remains correct for the wrapping and moving-path-point geometry of the biological knee, which is why the human targets are computed in OpenSim rather than by the robot-side construction.

A.2 Muscle Model Formulations and Cross-Platform Force Comparison

A.2.1 MuJoCo Muscle Actuator (as converted)

The conversion from OpenSim to MuJoCo re-expresses each Thelen2003 muscle as a MuJoCo `general` actuator with the built-in muscle dynamics. Activation is a first-order process with distinct rise and fall time constants (0.01 s and 0.04 s in the converted model). Force is computed from the actuator's normalized length $\bar{L} = (L - L_{\min})/(L_{\max} - L_{\min})$, where the length range is taken from the converted geometry, as

$$F = F_{\max} \left[a\, f_L(\bar{L})\, f_V(\dot{L}) - f_P(\bar{L}) \right], \tag{A1}$$

with $a \in [0, 1]$ the activation, f_L a force–length curve that peaks inside the normalized range, f_V a force–velocity curve, and f_P a passive element. The peak force scale $F_{\max}$ (actuator parameter `gainprm[2]`) is set by the converter such that the peak force over the model's range of motion matches the OpenSim `max_isometric_force`; audited across all 84 active muscles, the deviation is $\leqslant 15\%$ (Sec. A.2.2). Two properties matter for interpretation: (i) the tendon is *rigid*, so there is no series-elastic element and no tendon slack length, and force at the joint equals the contractile-element force immediately; and (ii) all curve parameters are shared defaults of the converted model rather than per-muscle fits.

A.2.2 Cross-Platform Force Comparison

Both models were driven with identical kinematics and identical activations along the subject01 walking trajectory (the same IK frames used for the static-optimization back-solve).

The OpenSim side is the StaticOptimization solution's own muscle forces (full Thelen dynamics with compliant tendon); the MuJoCo side replays the identical joint trajectory with those activations. Table A1 reports, per muscle, the coefficient of determination R^2 between the two force time series, the shape correlation, and the best-aligning lag [Provenance: `Code/MuJoCo_SNS/spinal/_muscle_force_compare.py`; per-muscle values for all twelve analyzed muscles are archived in `muscle_force_compare.csv` beside Figure A1].

Table A1: MuJoCo vs. OpenSim muscle force along the subject01 walk at matched activations. Lag in analysis frames (≈ 17 ms each).

Muscle	R^2	shape corr^2	lag [frames]
soleus	0.742	0.800	1
med. gastrocnemius	0.888	0.894	2
lat. gastrocnemius	0.884	0.887	2
tibialis anterior	0.945	0.945	0
vastus lateralis	0.941	0.945	1
rectus femoris	-1.15	0.384	0
semimembranosus	0.565	0.621	0
biceps femoris short	-0.26	0.868	0
psoas	0.927	0.934	1
gluteus maximus (2)	0.926	0.943	1
tensor fasciae latae	0.933	0.935	0
peroneus brevis	0.977	0.989	0

For eight of the twelve representative muscles the two formulations agree with $R^2 =$ 0.87–0.98 and mean-force ratios within 10%, with *no* systematic phase shift (best-aligning lag 0–2 frames, $\leqslant 34$ ms); soleus and semimembranosus agree moderately ($R^2 = 0.74$ and 0.57, respectively). The absence of a phase lead in the rigid-tendon model indicates that, along physiological trajectories, missing tendon slack compliance does not desynchronize force transmission; its signature is instead visible in waveform *shape*: for the soleus (the

largest series-elastic energy store in human gait), the OpenSim compliant tendon carries substantial force through mid-stance while the MuJoCo rigid-tendon force follows the activation dips. This is the expected smoothing of the series-elastic element, visible in the soleus panel of Figure A1. The two remaining outliers are anatomically explained: rectus femoris was rerouted in the conversion over the vastii's patella-tracking via point, and its moment-arm polyfit deviates from OpenSim's path (the same coupler artifact noted for the vastii); biceps femoris short head agrees in shape ($r^2 = 0.87$) but carries a scale offset. We conclude that the converted model reproduces OpenSim muscle force generation along physiological trajectories to a degree sufficient for neural-control studies, and that if tendon-elasticity phenomena (energy storage, force smoothing) become material to a claim, the appropriate upgrade is a compliant-tendon formulation rather than parameter retuning of the rigid-tendon model.

A.3 BPA Force Model Implementation

A.3.1 Goodness-of-Fit Metrics

All fitted quantities in this dissertation are scored by the shared routine `Functions/Go_OfF.m`: the root-mean-square error RMSE $= \sqrt{\frac{1}{n}\sum_i (y_i - \hat{y}_i)^2}$, the adjusted fraction of variance unexplained FVU $= \frac{n-1}{n}\ \sum_i (y_i - \hat{y}_i)^2 \ /\ \sum_i (y_i - \bar{y})^2$ (so that $R^2 = 1 - \text{FVU}$), and the maximum absolute residual $\max_i |y_i - \hat{y}_i|$ (`Go_OfF.m` lines 45–52). Lower is better for all three, which makes the triplet directly usable inside optimization objectives.

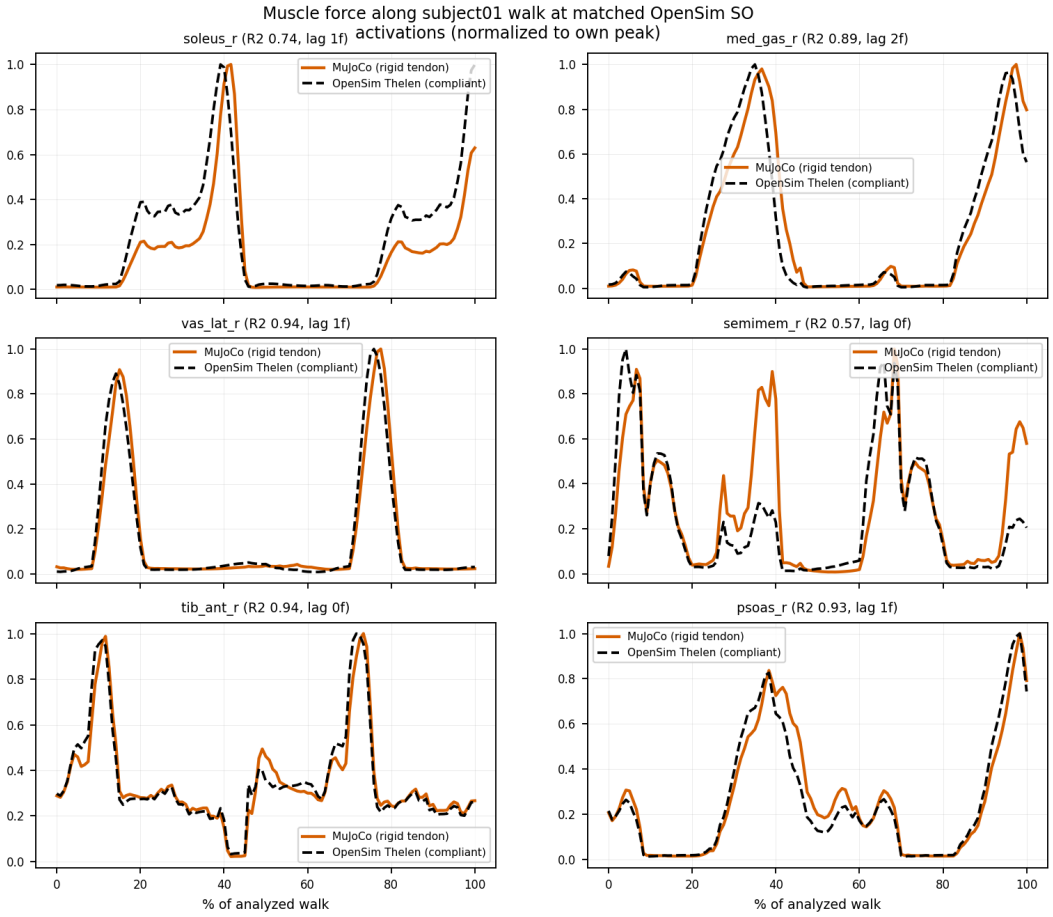


Figure A1: Normalized muscle force along the analyzed walk at matched OpenSim static-optimization activations (MuJoCo: rigid tendon; OpenSim: Thelen compliant tendon). Soleus illustrates the signature of the missing series compliance: OpenSim force is carried through mid-stance by tendon stretch while MuJoCo force follows the activation.

A.3.2 Strain Measures

The musculotendon path length l_{LMT} is the sum of straight segments between route via points, transformed into a common frame at each frame crossing. The contractile strain used by the force model subtracts the artificial tendon and both end-cap fittings from the path length before normalization by the resting length,

$$\varepsilon \;=\; \frac{l_0 - \left(l_{LMT} - l_{ten} - 2\,l_{fitting} - \chi_0\right)}{l_0}, \tag{A2}$$

where χ_0 is the constant length offset of Section 4.1.4 (zero in the uncorrected model). [Provenance: `MonoPamDataExplicit.m` get.Contraction, lines 210–217; `minimizeFlxPin.m` lines 415–426. End-cap fitting length $l_{fitting} = 25.4$ mm is a constructor default, `MonoPamDataExplicit.m` line 69.]

The maximum free-load contraction at 620 kPa is $\varepsilon_{620} = (l_0 - l_{620})/l_0$, where l_{620} is the BPA length at 620 kPa with no external load (code variable `KMAX`; `MonoPamDataExplicit.m` line 278). The relative strain is $\varepsilon^* = \varepsilon/\varepsilon_{620}$ (line 281).

A.3.3 Normalized Force Surface

The normalized isometric force surface of Equation (3.4) is implemented in closed form in `Functions/f_festo.m` (lines 1–26),

$$F^* = c_0\left(e^{-c_1\varepsilon^*} - 1\right) + P^*\,e^{-c_2(\varepsilon^*)^2}, \qquad F^* \leqslant 0 \;\Rightarrow\; 0, \tag{A3}$$

with $P^* = P/620\,\mathrm{kPa}$. The zero-clamp $F^* \leqslant 0 \Rightarrow 0$ is applied by the callers (e.g. `Functions/festo4.m` line 35), not inside `f_festo.m` itself. The coefficients were fitted with `fittype` nonlinear least squares and least-absolute-residual robustness (`Testing_Data/bpaFits.m`, lines 45–60 and 115–140, lower bounds $[0, 0, 0]$) and are reported in Table 3.2. The code additionally holds a 40 mm fit not used in this dissertation ($c_0 = 0.1224$, $c_1 = 10.47$, $c_2 = 2.023$; `f_festo.m` lines 17–20).

At runtime the torque model may alternatively evaluate a spline interpolation of the manufacturer's datasheet curves (`Functions/festo4.m`, backed by `FestoLookup.mat`), with force clamped to zero for $\varepsilon^* > 1$ or $F^* < 0$. Equation (A3) replaced this lookup in the corrected model.

A.3.4 Maximum Force

The resting-length-dependent maximum force of Equation (3.5) is implemented in `Functions/maxBPAforc` (10 mm: line 46; 20 mm: line 57; duplicate for 20 mm in `Functions/Fmax20.m`) with the fixed 7.5 mm end-contact offset. The fitted 20 mm coefficients carry confidence intervals $a_1 \in [1.4745, 1.5011]\,\mathrm{N\,kPa^{-1}}$ and $a_2 \in [0.0238, 0.0258]\,\mathrm{kPa^{-1}\,m^{-1}}$ (`Fmax20.m` lines 26–27). Total force is the product $F = F^*(\varepsilon^*, P^*)\, F_{620}(l_0)$, resolved into a vector along the force direction (`MonoPamDataExplicit.m` lines 283–287).

A.4 Muscle Path Length and Moment Arm

Path length accumulates the Euclidean norm of each segment between via points, applying the body-frame transformation whenever the path crosses between the femur and tibia

frames (`MonoPamDataExplicit.m` lines 78–95, total at lines 122–133).

The vector moment arm is implemented as the perpendicular foot of the joint origin onto the line of action,

$$\mathbf{r} \;=\; \mathbf{p} \;-\; \hat{\mathbf{u}}\,(\hat{\mathbf{u}}\cdot\mathbf{p})\,, \tag{A4}$$

where $\mathbf{p}$ is the vector from the joint origin to the muscle-path crossing point and $\hat{\mathbf{u}}$ is the unit force direction (`MonoPamDataExplicit.m` lines 156–168; identical implementations in `MonoMuscleData.m` lines 111–123, the `Mom` function of `minimizeFlxPin.m` at line 539, and `MonoPam_mult.m` lines 186–201). This is the same quantity as the cross-product projection of Equation (4.2); the planar scalar moment arm about the joint axis is the magnitude of the in-plane components, computed as `hypot(mA(:,1), mA(:,2))` (`predictKneeFlexor20mm.m` line 217). Torque about the joint z-axis is $\mathbf{m} = \mathbf{r} \times \mathbf{f}$, evaluated only where the BPA is in tension (values are discarded for $\varepsilon < -0.02$, the excessive stretching region; `MonoPamDataExplicit.m` line 308).

A.5 Compliance Model Implementation

The compliance-corrected muscle length entering Equation (4.5) is computed as $l_m = l_{LMT} - \chi_0 - l_{ten} - 2\,l_{fitting}$ (`minimizeFlxPin.m` lines 415–426; identically `minimizeExtX3.m` line 345). In the two-bracket fit of record, the insertion-bracket stiffness matrix and the projected scalar stiffness along the force direction are

$$\mathbf{K}_{br} = \mathrm{diag}(\chi_1,\, \chi_2,\, \chi_1), \qquad k_{\hat{u},br} = \left(\hat{u}^{\,T}\mathbf{K}_{br}^{-1}\hat{u}\right)^{-1}, \tag{A5}$$

matching Equation (4.8), and the origin-side second bracket contributes an additional scalar compliance c_{b2} in series (Appendix C, Table C6). As executed by the evaluator of record, both bracket frames are assembled pitch-only, at `minimizeFlxPin.m` line 289 (insertion) and line 306 (origin), with the two-rotation assemblies carried as commented alternatives; the insertion array is $[X_1, X_2, X_1]$ (line 387) and the origin array $[X_1, X_1, X_2]$ (line 400). The two-rotation-capable twin of the same block is `minimizeFlxPin2brk.m` lines 387–418. An earlier arm-1 array $[X_1, X_2, X_2]$ is superseded and survives only in the provenance comment at line 387, and the legacy single-bracket evaluator uses $[X_1, X_2, X_1]$ (`minimizeFlxPinX3.m` line 267).

The artificial tendon is modeled as an axial spring with stiffness

$$k_{ten} \;=\; \frac{n\,A_c\,E}{L_{ten}}, \tag{A6}$$

with $n = 2$ parallel cable strands ($\times 6$ for the 20 mm assembly), $A_c = 1.51 \times 10^{-6}\,\mathrm{m}^2$ (19-strand cable), and $E = 193 \times 10^{9}$ Pa (`MonoPam_mult.m`, function `Spr()`, lines 627–637; the evaluators obtain the spring rate through `Spr(bpa_i)`, `minimizeFlxPin.m` line 156). Series combination with the projected bracket compliance gives $k_{eq} = \left(k_{\hat{u},br}^{-1} + k_{ten}^{-1}\right)^{-1}$ (`minimizeFlxPin.m` lines 394–409).

The scalar compliance displacement δ of Equation (4.22) is found numerically with `fzero` on the bracket $[0,\, l_m - l_{620}]$,

$$F_{620}(l_0)\,F^*(\varepsilon^*(\delta), P^*) - k_{eq}\,\delta = 0, \tag{A7}$$

with $\delta \equiv 0$ when the uncorrected relative strain already exceeds unity (`minimizeFlxPin.m`

lines 434–441). The vector bracket deflection is then $\boldsymbol{\epsilon}_{br} = \mathbf{K}_{br}^{-1}\,(\|\mathbf{f}\|\,\hat{\mathbf{u}})$ (line 461), the cable stretch is $\|\mathbf{f}\|/k_{ten}$ (line 478), and the attachment point, path, moment arm, and torque are recomputed. For the two-BPA flexor grouping, the coupled balance shares one common force between both muscles and is solved analogously on $[0, F_{620}]$ (`MonoPam_mult.m` lines 556–612).

A.6 Wrapping-Loss Implementation

For the pinned-knee extensor, Equation (4.24) is implemented segment-wise over the wrap arc with radii $R_1 = 12\,\text{mm}$ and $R_2 = 40\,\text{mm}$ and segment boundaries $\theta_1 = 27^\circ$ and $\theta_2 = -23^\circ$ (`minimizeExtX3.m`, nested `Contraction`, lines 327–370). For the biomimetic-knee 20 mm cell, the arc-length factor is the routed bend measure $b(\theta_k)$ (arc length swept by the wrap, built with the `wRap` bend radius in `buildKneeFlexorRoute20mm.m` lines 265–268), giving $\Delta l = \chi_3\, b(\theta_k)\,(1-\varepsilon^*)^2$ (`MonoPam_mult.m` lines 398–441), with a hardcoded-geometry fallback for degenerate routes (lines 428–434). The loss enters only the force-producing strain; the kinematic (prediction) strain excludes it (`minimizeExtX3.m` lines 236–240).

Appendix B: Neuron Model Equations

This appendix documents the neuron and synapse models used by the synthetic nervous systems in this project, both in the AnimatLab walker model and in the SNS-Toolbox implementation of the MuJoCo pipeline (Chapter 6). The authoritative model state is the working AnimatLab project file (`Biped_2xCPG_wSubs/Biped_2xCPG_wSubs.aproj`); after the network reorganization described in Section 6.3, a fresh standalone export (`Biped_2xCPG_wSubs_Standalone.asim`) runs headless and carries the rhythm-generator charts, and the earlier phase-1 instrumented export (`walk new new tester added 2 axis_phase1.asim`) records the run reported in Section 6.4.

B.1 Spiking Neuron Model

Each spiking neuron is a conductance-based leaky integrator with threshold and adaptation,

$$C_m \dot{V} \;=\; -\,g_{\text{leak}}\,(V - E_{\text{rest}}) \;+\; I_{\text{syn}} \;+\; I_{\text{app}}, \tag{B1}$$

where V is the membrane potential, C_m the membrane capacitance ($g_{\text{leak}} = C_m/\tau_m$ with membrane time constant τ_m), E_{rest} the resting potential, I_{syn} the summed synaptic current, and I_{app} any applied (tonic stimulation) current. A spike is emitted when V exceeds the firing threshold, after which V resets and an after-hyperpolarization (AHP) of amplitude A_{AHP} with time constant τ_{AHP} is applied; the threshold itself accommodates with the firing history (relative accommodation r_{acc}, time constant τ_{acc}). Optional calcium channels are

defined but disabled in the current model ($G_{\text{max,Ca}} = 0$).

The parameter values shared by the neurons in the walker model are given in Table B1.

Table B1: Spiking neuron parameters in the AnimatLab bipedal walker model (`Biped_2xCPG_wSubs`). All neurons share these values except the motor neuron threshold noted below.

Parameter	**Symbol**	**Value**
Resting potential	E_{rest}	$-60\,\text{mV}$
Membrane time constant	τ_m	$5\,\text{ms}$
Initial firing threshold	V_{th}	$-55\,\text{mV}$ (motor neurons: $-59\,\text{mV}$)
Relative accommodation	r_{acc}	0.3
Accommodation time const	τ_{acc}	$10\,\text{ms}$
AHP amplitude	A_{AHP}	$1\,\text{mV}$
AHP time constant	τ_{AHP}	$3\,\text{ms}$

B.2 Non-Spiking (Graded) Chemical Synapses

Pattern-formation and motoneuron-drive connections in the walker use graded chemical synapses, whose conductance is a clipped linear ramp of the presynaptic potential,

$$g_{\text{syn}}(V_{\text{pre}}) \;=\; g_{\max}\;\text{clamp}\!\left(\frac{V_{\text{pre}} - V_{\text{th}}}{V_{\text{sat}} - V_{\text{th}}},\,0,\,1\right), \qquad I_{\text{syn}} = g_{\text{syn}}\left(V_{\text{post}} - E_{\text{syn}}\right), \tag{B2}$$

where V_{th} and V_{sat} are the synaptic threshold and saturation potentials and E_{syn} the reversal potential. The walker model instantiates three non-spiking synapse types:

- **Nicotinic ACh (excitatory drive):** $E_{\text{syn}} = -10\,\text{mV}$, $g_{\max} = 0.5\,\mu\text{S}$, $V_{\text{th}} \to V_{\text{sat}} = -65\,\text{mV} \to -20\,\text{mV}$.

- **Hyperpolarizing IPSP (inhibition):** $E_{\text{syn}} = -70\,\text{mV}$, $V_{\text{th}} \rightarrow V_{\text{sat}} = -55\,\text{mV} \rightarrow -40\,\text{mV}$.
- **Depolarizing IPSP:** $E_{\text{syn}} = -55\,\text{mV}$, $V_{\text{th}} \rightarrow V_{\text{sat}} = -60\,\text{mV} \rightarrow -40\,\text{mV}$.

Spiking chemical synapses and electrical synapses (rectifying coupling weight 0.1, non-rectifying 0.2) are also defined in the model and are available for the sensory afferent pathways of the MuJoCo/SNS implementation.

B.3 Relation to the SNS-Toolbox Implementation

These are the standard AnimatLab firing-rate/conductance models. SNS-Toolbox [**?**] implements corresponding classes of non-spiking and spiking neurons with conductance-based synapses, and the AnimatLab parameters above served as the physically motivated initialization for the first SNS-Toolbox instantiation of the spinal network: a 410-neuron network with per-leg rhythm-generator half-centers, phase-shifted pattern-formation groups, motoneuron pools for all 92 muscle-tendon actuators of the converted model, and Ia, Ib, and II afferent populations. In that implementation the non-spiking membrane and graded-synapse equations take the same forms as Equations (B1) and (B2) with the spiking threshold and accommodation terms removed, and the afferent resting-tone offsets are zeroed so that Ia and Ib signals are pure activity measures; the rhythm-generator layer has been verified to oscillate at a human-like period with the two legs in antiphase under tonic drive (Section 7.2.4). An explicit parameter-mapping and validation study between the two implementations remains future work; the tonic stimulation tests of Section 7.2.4 correspond to constant I_{app} injections into Equation (B1) at the motoneuron,

pattern-formation, and rhythm-generator layers respectively.

B.4 Spiking-Mirror Implementation, Calibration, and Verification

This appendix documents the numeric details behind the spiking-neuron mirrors of Section D.4, in the same file-traceable form as Appendix A: the cell and synapse equations as implemented, the calibration factors, the full verification-gate table of the MuJoCo mirror, and the variant-by-variant AnimatLab conversion table. The artifacts live under `Code/MuJoCo_SNS/spinal/reports_spiking_20261002/` (reports, gate scripts, logs, and metric files) and `Code/MuJoCo_SNS/spinal/` (the mirror builder `build_network_spiking.py` and its calibration script `calibrate_spiking.py`); the Simulink blocks are built by `Code/Matlab/SNS_Sims` and tested by `sns_units_test_spiking.m` in the same folder.

B.4.1 Adapting Integrate-and-Fire Cells and Hybrid Synapses

Between spikes, every spiking cell of the MuJoCo mirror obeys the membrane equation

$$C_m \dot{V} \;=\; -\,g_m\,(V - V_{\text{rest}}) \;+\; I_{\text{syn}} \;+\; I_{\text{bias}}, \tag{B3}$$

with a spike emitted when V reaches the firing threshold θ, after which V resets to V_{reset} and θ increases by the increment $\Delta\theta$; between spikes θ decays back toward its initial value with time constant τ_θ, bounded by the adaptation floor. The tonic bias I_{bias} parks the equilibrium a few millivolts below threshold, reproducing the analog cells' near-threshold operating point. Table B2 lists the per-class values as coded (SNS-Toolbox's `SpikingNeuron` class;

the toolbox's membrane update is dimensionally consistent only at $g_m = 1\,\mu\text{S}$, which the mirror therefore uses throughout). The rhythm-generator adaptation time constant reuses the analog network's persistent-sodium inactivation value, so that tuned adaptation values transfer between builds; at the rhythm gate's operating point it is 0.35 s. The mechanism behind that value, recorded here rather than in the chapter body, is as follows. In the analog rhythm cells of the contact-driven cores, burst termination is carried by a persistent sodium current, a slowly inactivating inward current; in every port of this dissertation its inactivation time constant is held fixed rather than voltage-dependent, because the voltage-dependent form of the AnimatLab original quenches burst termination and leaves the network latched in a static state, the failure documented with the archived-model results of Chapter 6.

Table B2: Spiking cell parameters of the MuJoCo mirror (real millivolts; `build_network_spiking.py`, `SPIKE` table, with the rhythm-generator increment superseded by the calibration of record). All cells rest at $-70\,\text{mV}$ with membrane conductance 1 µS and the tonic bias of Table B3.

Cell class	θ **(mV)**	V_{reset} **(mV)**	$\Delta\theta$ **(mV/spike)**	τ_θ **(s)**
Rhythm generator E/F	-50	-60	4.0 (calibrated)	0.35 (analog value)
Pattern-formation	-50	-60	0.3	0.15
Interneurons, commissurals	-50	-60	—	—
Afferent encoders	-52	-58	—	—
Renshaw	-52.5	-58	—	—

Spiking chemical synapses, in both the toolbox class used in the MuJoCo mirror and the Simulink `SpikingSynapse` block, carry a conductance that decays between events and

jumps by a fixed increment per presynaptic spike, saturating at its maximum,

$$\dot{g} \;=\; -\,g/\tau_{\text{syn}} \;\text{ between spikes,} \qquad g \leftarrow \min(g_{\max},\, g + g_{\text{inc}}) \;\text{ per spike,} \tag{B4}$$

with $I_{\text{syn}} = g\,(E_{\text{syn}} - V_{\text{post}})$ as for the graded synapse of Equation (B2). The Simulink `HybridSpikingSynapse` implements the same dynamics with the spike detected from the presynaptic voltage crossing a detection threshold that must sit strictly between the presynaptic resting and firing potentials (a detection threshold equal to the firing potential fires nothing, because the presynaptic voltage resets at threshold). Reversal potentials are 0 mV (excitation, above rest) and -70 mV (inhibition, shunting at rest) in the spiking frame, and $+8/-5$ mV in the motor neurons' 0–5 mV frame. The graded-to-spiking command pathways (descending drive, posture, and bias surrogates) keep the graded synapse of Equation (B2).

B.4.2 Calibration of Record

Table B3 lists every calibration factor produced by `calibrate_spiking.py` (`spiking_calibration.json`), with the measured residual of each. The spike-to-motor-neuron factors are bisected so that a steady 40 Hz presynaptic train reproduces the non-spiking steady-state motor-neuron voltage; the spike-to-spiking factors are rate-calibrated (a 40 Hz source drives a relay cell at 40 Hz, and a calibrated inhibitory factor silences a 40 Hz-driven cell); the readout-tap factor maps a 50 Hz train to the full-scale analog level. The rhythm-generator threshold increment was set by an isolated-pair sweep followed by a full-network sweep over alternation quality and period target.

Table B3: Calibration factors for the spiking mirror (`spiking_calibration.json`). Residuals are the measured deviation at each factor's calibration target.

Factor	Value	Calibrated against	Residual
Graded saturation reference	0.9567	analog network measurement	—
Spike-to-MN, excitatory ($k_{sn,exc}$)	5.4087	MN steady voltage	1.0×10^{-4} n
Spike-to-MN, inhibitory ($k_{sn,inh}$)	1.2284	MN steady voltage	2.6×10^{-4} n
Graded-to-spiking drive (k_{ns})	0.35	relay firing rate	-0.67 Hz
Spike-to-spiking, excitation	1.8754	40 Hz rate transfer	0.0 Hz
Spike-to-spiking, inhibition	0.9996	silences 40 Hz cell	0.0 Hz
Readout tap (g_{inc})	1.7093	50 Hz → full scale	5.00 mV che
RG threshold increment	4.0 mV/spike	pair + full-net sweeps	period targ
Tonic bias current	16 nA	equilibrium -54 mV	—
Synaptic time constants	5 ms (exc), 20 ms (inh)	analog values	—

B.4.3 Verification Gates of the MuJoCo Mirror

Table B4 records the complete gate set of Section D.4 with verdicts and headline numbers. The two failures are reported as failures: the constant-drive period misses its tolerance with the tuning knobs identified, and the air-regime cadence clause misses while every structural clause passes. Gate scripts, logs, and machine-readable outputs are archived in the reports folder (`tools/`, `logs/`, and the `json` metric files respectively).

B.4.4 AnimatLab Conversion Variants

Table B5 catalogues every spiking copy built from the AnimatLab standalone exports. The conversion rules, applied by `tools/make_spiking.py` to copies (no original file touched), are: thresholds of neurons with unreachable (positive) thresholds move into the spiking regime, either -55 mV with -59 mV for motor-neuron-named cells in the v1 copies, or

Table B4: Complete verification-gate table of the spiking mirror. Baseline scores reproduce digit for digit; "split" denotes a gate whose clauses pass and fail in part.

Gate	Verdict	Headline numbers
Pipeline compile	PASS	runner, curriculum, parameters, and network builder all compile.
Defaults gate	PASS	(410, 376, 1186) populations, inputs, and connections after the selector edit; variant build (888, 382, 7430).
Non-spiking bit-identity	PASS	default baseline score -161.56754173676563 identical before and after the edit (all 17 digits); other baselines -144.46440311078322 and -200.37955834081458.
Topology mirror	PASS	directed edge multisets identical in three gain configurations (1,186 / 7,908 / 7,762); inputs and motor-neuron maps identical; readout taps 12 / 12 / 16.
Constant-drive rhythm	SPLIT	sustained bilateral antiphase (E 4.9 Hz, F 3.2 Hz per leg; swing 5.13 mV in every window) = PASS; period 0.621 s versus analog 0.399 s ($+55.6\,\% > 20\,\%$) = FAIL.
Perturbation basin	PASS	12/12 trials under 1 % jitter on every scalar; final swing minimum 4.58 mV (margin $0.89\times$); zero explosions.
Air-regime smoke	FAIL (cadence)	finite, 53 envelope rises, knee range 106.0° of twin's 117.9°, hip 50.4° of 71.1°; period 0.215 s versus 2.108 s.
Untuned ground evaluation	ran	score -236.35; bilateral (8 right / 9 left cycles); cycle knee minimum -79.8°; no non-finite states; stayed up. Single run, no statistics.
Simulink regression	PASS	all committed demonstrations and units tests reproduce exactly with the spiking blocks in the library (two-neuron test 6.07×10^{-4} mV).
AnimatLab baselines	PASS	five headless baseline runs, exit 0, periods and antiphase reproducing the records (0.443 s, 0.462 s, 1.305 s).

per-class calibrated values from the measured graded peaks in the v2 copies (rhythm-generator/pattern-formation/interneuron $-53.5\,\mathrm{mV}$, motor neurons $-44\,\mathrm{mV}$, Renshaw $-44.5\,\mathrm{mV}$); relative accommodation is zeroed; every graded chemical synapse type is re-emitted as a per-spike type with the same identifier and amplitude, 10 ms decay, and maximum conductance $\max(5,\, 2\times\text{amplitude})$; and each connection's own class tag is flipped to the spiking enum. That last rule is a documented format trap: the per-connection `<Type>` tag is stored independently of the synapse type, and converting only the type kills the run at load with no diagnostic detail. Structural validation confirms for every copy that neuron, connection, synapse-type, chart, and stimulus sets are identical to the original and that every class tag matches its synapse type.

Figure B1 shows the ablation result at the joint level for the two-layer family: with graded synapses kept, spiking-regime thresholds leave the joint oscillation essentially intact, while the full conversion pins the joints. This is the visual form of Table D5's conclusion that the synapse conversion is the breaking change.

The spiking copies are headless-run artifacts only: none has been opened in the AnimatLab graphical environment, and graphical verification (or an equivalent acceptance ruling) remains outstanding for every `.asim` listed in Table B5. The mirror work also scoped, but did not attempt, two further conversions: the full 410-population Simulink spinal network (whose every one of 1,392 gains is tuned for graded voltages, making a spiking conversion a full re-tune rather than a parameter port) and the spiking mirrors of the alternative MuJoCo network variants (separate architectures whose fresh tuning is future work).

Table B5: AnimatLab spiking-conversion variants of the spiking mirror. "Connections" counts flipped per-connection class tags; the ablation converts thresholds only. All copies validate structurally and run headless to their configured end; rhythm outcomes as in Table D5.

Variant	Thresholds	Neurons	Synapse types	Connections	Rhythm outcome
Two-layer v1	−55/−59 mV	80	17	150	dead (2 spikes / 5 s)
Two-layer v2	calibrated	80	17	150	dead
Two-layer v3	v1, 30 ms decay	80	17	150	dead
Bilateral v1	−55/−59 mV	96	20	208	latch, one side 442 Hz–450 Hz
Bilateral v2	calibrated	96	20	208	latch
Project export v1	−55/−59 mV	122	50	243	synchronous 11.8 Hz co-bursti
Project export v2	calibrated	122	50	243	dead (passive settle)
Ablation	spiking, graded synapses	80	0	0	oscillates, near-baseline ranges
2023 reference (baseline)	spiking from birth	56	3 (spiking)	84	1.305 s bursts, antiphase +0.49

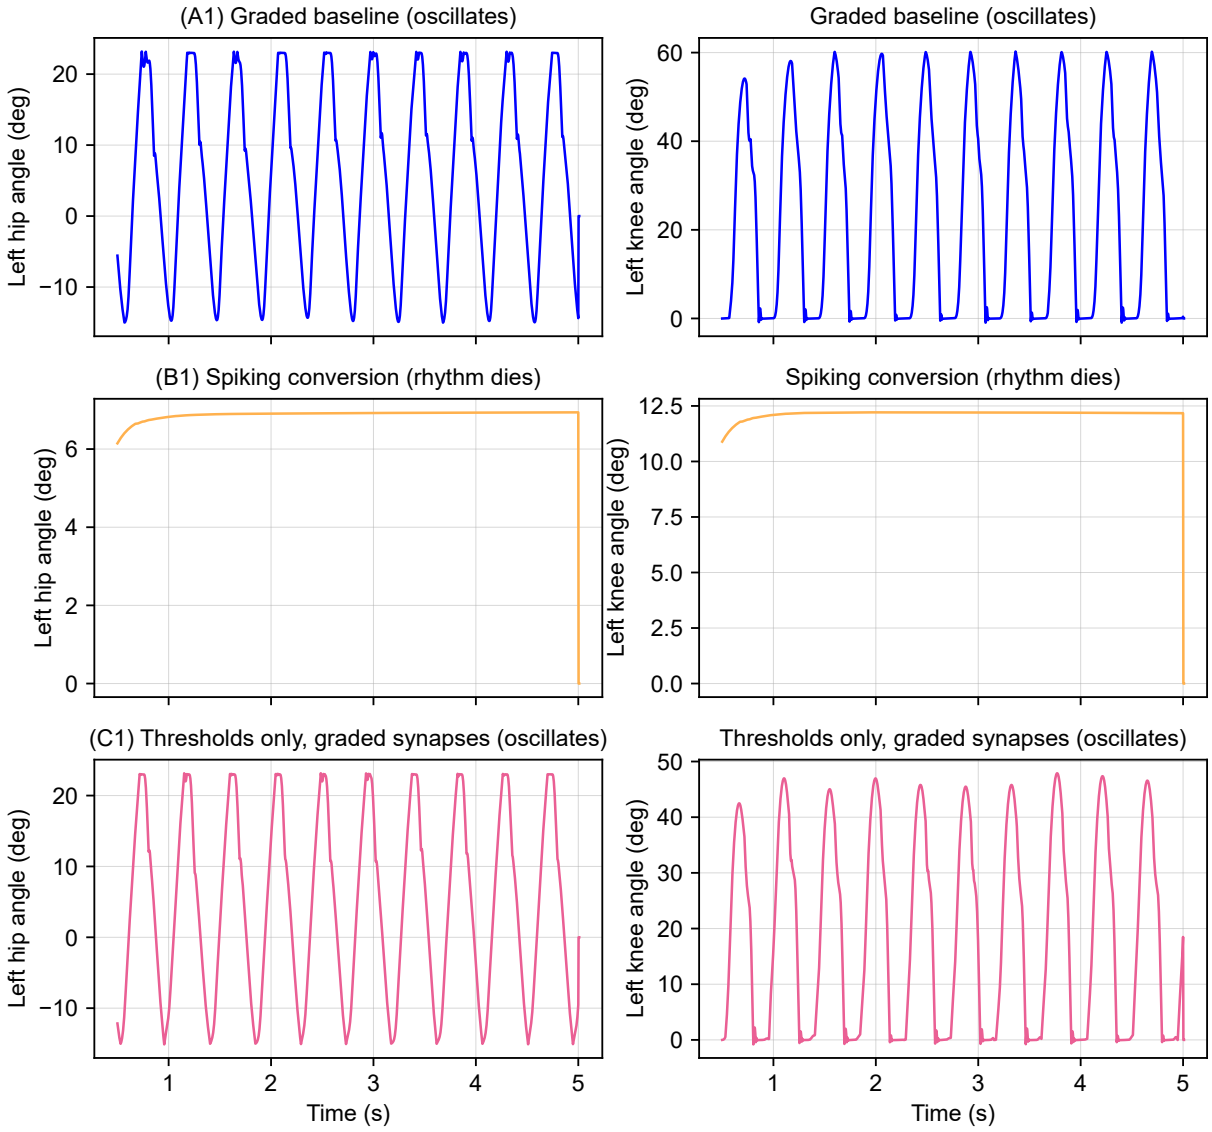


Figure B1: Left-leg joint angles of the AnimatLab two-layer walker family over the 0.5 s–5 s chart window: **(A, B)** graded baseline; **(C, D)** full spiking conversion (rhythm dies, joints pin); **(E, F)** thresholds-only ablation with graded synapses kept (oscillation preserved). Left column: hip angle. Right column: knee angle.

B.5 Connectomes of the Spinal-Network Models

The connectome of the production gait2392 spinal network (the 410-neuron SNS-Toolbox instantiation of Section 7.2.4) is documented in Figures B2 and B3. Both schematics are generated by `Code/MuJoCo_SNS/spinal/draw_circuit.py` from the compiled network at the fitted walking configuration (`fitted_walk_params.json`, the IK/NNLS back-solve refit, in which the pattern-formation and posture weights are refit against muscle activations back-solved from recorded walking kinematics, Section D.3; 410 neurons, 201 per leg plus 8 shared, and 376 inputs; the eight shared cells are the single `DRIVE` and `POSTURE` descending neurons and the six balance cells `BAL_PF`, `BAL_DF`, `BAL_TRK_EXT`, `BAL_TRK_FLX`, `BAL_LAT_R`, and `BAL_LAT_L`), with every motoneuron-pool membership pulled live from `muscle_map.py` rather than hardcoded, so the drawings cannot drift from the built network. The caption of Figure B2 and its alt-text sidecar repeat these population counts, so prose, caption, and alt text cannot drift apart either. Synapse markers follow the project convention throughout: open circles are neurons, white triangles mark excitatory synapses, and solid black dots mark inhibitory synapses.

Figure B2 draws both legs in full. A vertical dashed midline separates the sides and a pale central corridor gathers the interleg commissural inhibition. Each side is organized in labeled layer bands: descending DRIVE (MLR surrogate), POSTURE, and balance inputs; the rhythm-generator half-centers with their mutual inhibition and adaptation cells; the phase-window pattern-formation cells; the Ia, II, and Ib afferent populations with the stance-gated Ib load-sharing interneurons; and the ten motoneuron pools per side (HIP-E 7, HIP-F 4, KNEE-E 4, KNEE-F 4, ANK-PF 9, ANK-DF 3, HIP-AB 6, HIP-AD 6, TRK-E

1, and TRK-F 2, that is 46 motoneurons per leg) driving their muscles. Figure B3 shows the same network for one leg as a DRIVE/POSTURE to rhythm-generator to pattern-formation to motoneuron hierarchy, with the solved standing bias injected into every pool, the sensory column of Ia, II, and Ib afferents, and a boxed key listing every pool abbreviation with its member muscles. The spiking mirror of Section B.4 replicates this topology exactly (Table B4, topology-mirror gate), so these two figures also serve as the connectivity reference for that model.

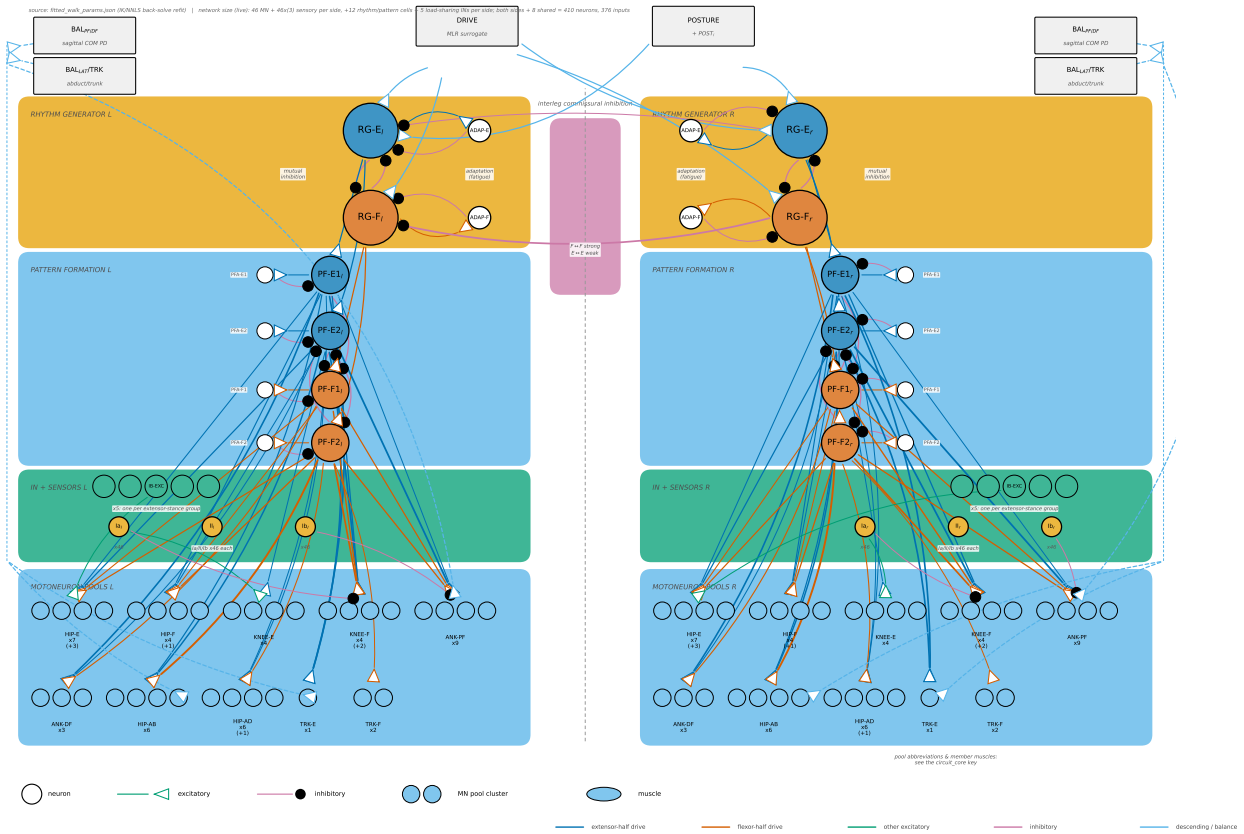


Figure B2: Bilateral connectome of the gait2392 spinal network rendered from the compiled network at the fitted walking configuration (410 neurons, 201 per side plus 8 shared; 376 inputs). Each side is organized in labeled layer bands: descending DRIVE (MLR surrogate), POSTURE, and balance inputs; rhythm-generator half-centers with mutual inhibition and adaptation cells; phase-window pattern-formation cells; Ia, II, and Ib afferents with the stance-gated Ib load-sharing interneurons; and the ten motoneuron pools per side driving their muscles. The central corridor gathers the interleg commissural inhibition across the dashed midline. Open circles are neurons, white triangles mark excitatory synapses, and solid black dots mark inhibitory synapses; the gray note names the source files.

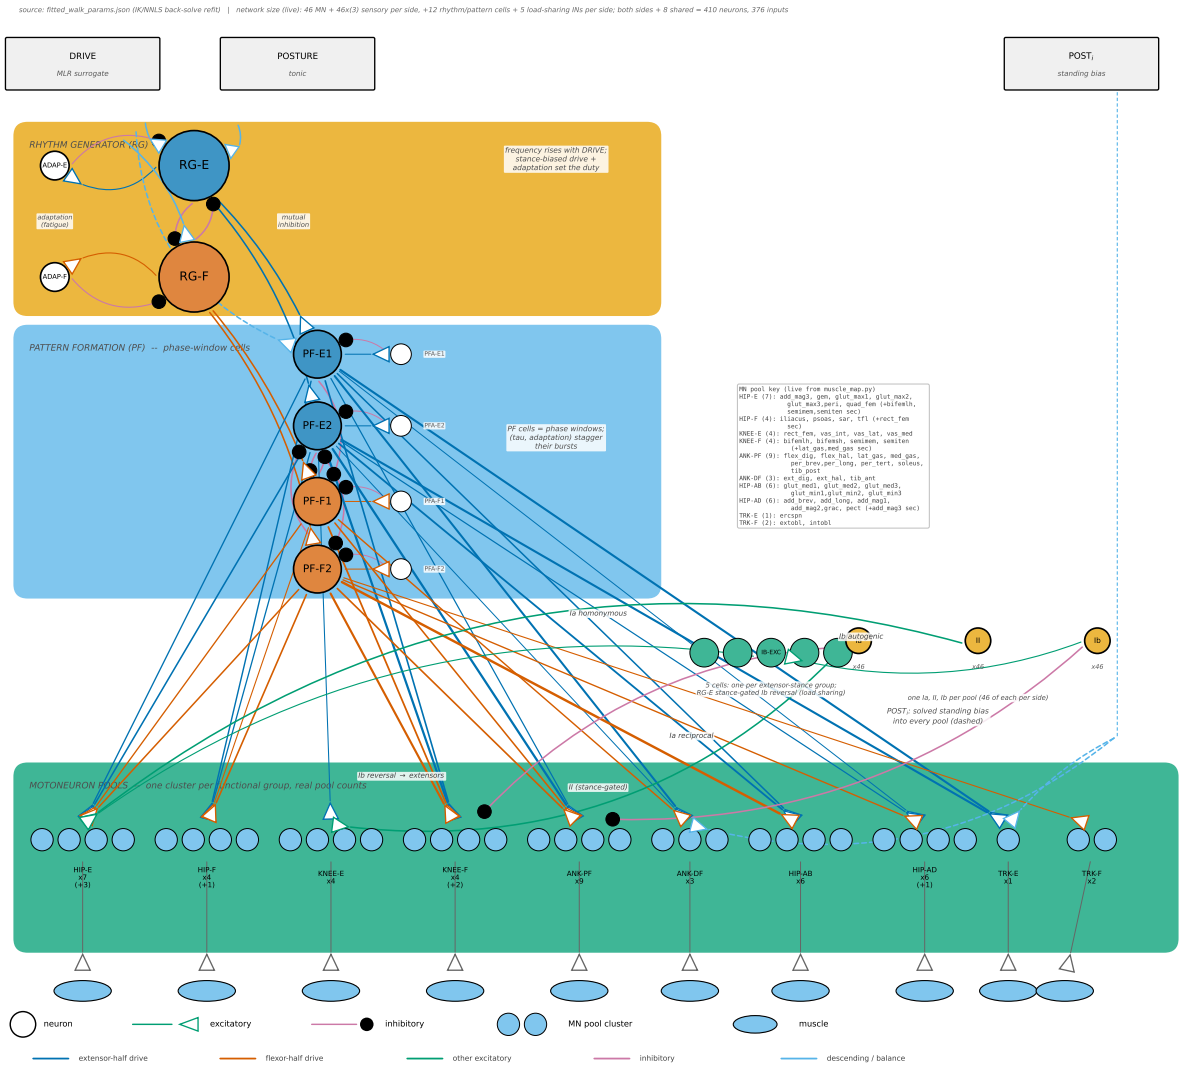


Figure B3: One-side hierarchy of the same gait2392 spinal network: descending DRIVE (MLR surrogate) and POSTURE inputs, the solved standing bias injected as dashed lines into every motoneuron pool, the persistent-sodium rhythm-generator half-centers, the phase-window pattern-formation cells, the Ia, II, and Ib afferent column, and the motoneuron pools (46 per leg) projecting to their muscles. A boxed key lists every pool abbreviation with its member muscles.

Appendix C: Model Geometry, Bracket Frames, and Compliance Parameters

This appendix documents the numeric geometry and compliance parameters behind the force and torque models of Chapters 3 and 4 as they exist in the project's MATLAB code, with file names given so that any value can be traced to the code that produced it. Route and bracket coordinates are coded in meters and reported here in millimeters; stiffnesses are coded and reported in newtons per meter. The bracket reference points are interpretation choices for the 3D-printed (Onyx FDM) attachment hardware, set from CAD by me; moving a reference point changes the identified stiffness values, so the points are recorded here alongside the values they produced.

C.1 Static BPA Pathway Data

The static BPA model routes each actuator as a straight-segment polyline through via points stored in a per-test `Location` array of dimensions $M \times 3 \times N$, where N is the number of joint frames in the motion sweep and M the number of via points; the first row in the distal link frame is marked by a cross-point index. There are no binormal or surface-frame quantities: at each frame the force direction is the normalized vector from the cross-point row to the previous row transformed into the crossing frame,

$$\hat{\mathbf{u}} \propto \mathbf{R}^{-1}\,\mathbf{p}_{C-1} \;-\; \mathbf{p}_{C}, \tag{C1}$$

where $\mathbf{p}_C$ is the cross-point row, $\mathbf{p}_{C-1}$ the previous row, and $\mathbf{R}$ the frame-crossing rotation [Provenance: `minimizeFlxPin.m` lines 160 and 524–537 (`UD`) and 201–235 (`computeForceVector`); `minimizeExtX3.m` lines 687–700; BPA struct fields, `minimizeFlxPin.m` lines 28–35]. The pinned-knee test routes are archived as `FlxPinBPASet.mat` (5 flexor tests, $M = 2$, knee angle -120° to $+10^\circ$) and `ExtPinBPASet.mat` (9 extensor tests, $M = 7$, knee angle -125° to $+35^\circ$).

Table C1 lists the five pinned-knee flexor tests; their routes are two points, femur-frame origin $\mathbf{p}_1 = [-75.0,\ 100.0,\ 32.8]$ mm and tibia-frame insertion $\mathbf{p}_2 = [-50.1,\ -45.0,\ 32.3]$ mm (test 1; `FlxPinBPASet.mat`, kf(1), frame 1). Table C3 lists the extensor pinned-knee route of test 1 at $\theta_k = -125^\circ$; rows 1–6 are femur-frame and row 7 is tibia-frame, and as the knee extends, rows 2–6 release onto $\mathbf{p}_1$ (wrap released at $+35^\circ$). Table C4 lists the per-test constants of the nine extensor tests; the correction terms of Section 4.1.4 used tests {1, 2, 5, 6, 7, 8}, of which the adopted fold trained on {5, 6, 7} and held out {1, 2, 8} for validation; tests 3, 4, and 9 were not used in the evaluation [Provenance: fold columns of `minimizeExt10mmX3_results.mat`; `Dig_crossPredict.m` line 46].

The biomimetic flexor route (`buildKneeFlexorContext20mm.m`) is a three-point path (femur-frame origin $\mathbf{p}_1$, a moving wrap point $\mathbf{p}_{\text{wrap}}$ on the tibia, and tibia-frame insertion $\mathbf{p}_{\text{end}}$), with the wrap point held on a standoff surface of radius `bpaRb` about the joint and its z placed on the projected $\mathbf{p}_1$–$\mathbf{p}_{\text{end}}$ line; the optimizer seed is $\mathbf{p}_{1,0} = [-50.0,\ 35.0,\ 50.0]$ mm and $\mathbf{p}_{2,0} = [-12.24,\ -8.87,\ 27.87]$ mm with $l_0 = 415$ mm and 25 mm tendon [Provenance: `buildKneeFlexorContext20mm.m` lines 209–218; `buildKneeFlexorRoute20mm.m` lines 98–186]. The biomimetic extensor seed route is given in Table C2, including the release schedule of the route-elimination rule.

Table C1: Pinned-knee flexor test BPA constants, read directly from `FlxPinBPASet.mat` (kf(1)–kf(5)). Labels are the laboratory shorthand used in the fitting logs (`minimizeFlxPin10mm_2brk.m` line 51); l_0, tendon, and pressure are the values used by the model. All tests use ∅10 mm BPAs; the 47 cm test's measured angles carry the $+5.3^\circ$ encoder correction at load time (`minimizeFlxPin.m` lines 77–80). The test was not used in the two-bracket identification, and the corrected-model fit reported for it in Table C9 is a post hoc evaluation; of the remaining four, the adopted cross-validation fold trained on kf(1) and kf(5) and held out kf(2) and kf(4) for validation.

Test	**Label**	$\boldsymbol{l_0}$ (m)	Tendon (m)	P (kPa)	F_m (N)	**Use**
kf(1)	48 cm	0.485	0	605.7	443.4	training
kf(2)	46 cm	0.457	0	614.2	441.4	validation
kf(3)	47 cm	0.479	0	623.2	443.0	not used
kf(4)	40 cm-tendon	0.406	0.035	623.4	436.9	validation
kf(5)	41 cm	0.410	0	615.4	437.3	training

C.2 Joint Transformation Matrices

The knee transform is planar: every joint frame is a rotation about z by the knee angle θ_k composed with a translation fitted to robot CAD tables of the migrating instantaneous center of rotation. With $\mathbf{R} = \mathbf{R}_z(\theta_k)$, the robot (femur→ICR) transform is $\mathbf{T}_{\text{Pam}} = [\mathbf{R},\ [f_3(\theta_k),\ f_4(\theta_k),\ 0]^\top;\ \mathbf{0}^\top\ 1]$ with f_3, f_4 cubic-spline fits (ICR x: 23.30 to 10.47 mm; y: -416.65 to -395.66 mm over $\theta_k = 0.17$ to -2.62 rad), and the ICR→tibia transform is the pure translation $[f_{13}(\theta_k),\ f_{14}(\theta_k),\ 0]^\top$ (x: 29.66 to -7.58 mm; y: 25.97 to 22.33 mm), composed as $\mathbf{T}_{t1\to f} = \mathbf{T}_{t1\to\text{ICR}}\ \mathbf{T}_{\text{Pam}}^{-1}$ over 100 frames spanning $\theta_k \in [-120^\circ, +10^\circ]$ [Provenance: `buildKneeFlexorContext20mm.m` lines 34–80; `buildKneeExtContext20mm.m` lines 30–78; `Functions/ModernRobotics/RpToTrans.m`]. In the archived test sets the per-frame

Table C2: Biomimetic knee, extensor 20 mm BPA route seed (femur frame: p1–p5; tibia frame: p6–p9; `buildKneeExtContext20mm.m` lines 431–440, my fcec-ray construction; the prior SolidWorks seed is kept commented at lines 421–430). Release schedule of the verified elimination, read from the redesign of record (`Vas_Pam_20mm_Result_20260925.mat`, `routeInfo.eliminatedAngleD`): p5 at -89.8°, p4 at -55.7°, p6 at -47.8°, p7 at -20.2°, p3 at $+6.1^\circ$, p8 at $+7.4^\circ$.

Point	x (mm)	y (mm)	z (mm)	**Frame**	**Role**
p1	40.00	35.00	0	femur	BPA origin (attachment)
p2	83.90	−274.76	0	femur	BPA contacts mounting base
p3	64.06	−426.80	0	femur	femoral contact
p4	52.83	−441.88	0	femur	femoral condyle contact
p5	37.77	−452.14	0	femur	femoral condyle contact
p6	69.85	34.33	0	tibia	tibial contact (upper)
p7	73.50	19.13	0	tibia	tibial contact (mid)
p8	71.98	−12.23	0	tibia	tibial contact (lower)
p9	34.06	−66.87	0	tibia	patellar ligament ring (distal, pEnd)

Table C3: Pinned-knee extensor BPA route of test 1 at $\theta_k = -125^\circ$ (`ExtPinBPASet.mat`, ke(1), frame 1; mm).

Point	x (mm)	y (mm)	z (mm)
p1	30.00	−50.00	0
p2	65.00	−350.00	0
p3	42.40	−408.50	0
p4	28.20	−421.60	0
p5	10.50	−425.20	0
p6	−8.60	−422.80	0
p7	42.50	−75.90	0

Table C4: Pinned-knee extensor test BPA constants, read directly from `ExtPinBPASet.mat` (ke(1)–ke(9)). Labels are the laboratory shorthand used in the fitting logs (`Dig_resultsTable.m` line 32); l_0, tendon, and pressure are the values used by the model; the shorthand rounds several lengths (e.g. "42 cm" is $l_0 = 0.415$ m and "43 cm-tendon" is $l_0 = 0.432$ m). All tests use ∅10 mm BPAs. The correction terms of Section 4.1.4 used tests $\{1, 2, 5, 6, 7, 8\}$: the adopted fold trained on $\{5, 6, 7\}$ and held out $\{1, 2, 8\}$ for validation; tests 3, 4, and 9 were not used in the evaluation.

Test	**Label**	$\boldsymbol{l_0}$ (m)	Tendon (m)	P (kPa)	F_m (N)	**Use**
ke(1)	40 cm	0.400	0	624.9	436.3	validation
ke(2)	40 cm-tendon	0.400	0.040	618.4	436.3	validation
ke(3)	42 cm	0.415	0	605.2	437.8	not used
ke(4)	42 cm-tendon	0.415	0.040	605.2	437.8	not used
ke(5)	43 cm	0.436	0	621.8	439.7	training
ke(6)	43 cm-tendon	0.432	0.022	622.9	439.3	training
ke(7)	46 cm	0.457	0	605.9	441.4	training
ke(8)	47 cm	0.465	0	605.9	442.0	validation
ke(9)	48 cm	0.480	0	605.9	443.1	not used

transform maps tibia to femur/ICR and takes the form

$$\mathbf{T}_k(\theta_k) \;=\; \begin{bmatrix} \cos\theta_k & -\sin\theta_k & 0 & d_x(\theta_k) \\ \sin\theta_k & \cos\theta_k & 0 & d_y(\theta_k) \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}, \tag{C2}$$

with rotation block $\mathbf{R}_z(\theta_k)$ and translation components d_x, d_y in meters. In the archived flexor test set the translation is constant across the sweep, $[d_x,\, d_y,\, 0] = [4.5,\, -395.8,\, 0]$ mm, and the rotation sweeps from $\mathbf{R}_z(-120^\circ)$ at full flexion to $\mathbf{R}_z(+10^\circ)$ at full extension (`FlxPinBPASet.mat`, kf(1)).

C.3 Bracket Reference Points and Frames

Table C5 lists the bracket reference points. Each compliant bracket carries an instantaneous body-fixed frame anchored at its reference point, constructed by two rotations (yaw about z, then pitch about the yawed y). With $\mathbf{p}_b$ the vector from the reference point to the bracket-to-BPA attachment target, the yaw and pitch angles are

$$\theta_{br} = \operatorname{atan2}(p_{by},\, p_{bx}), \qquad \phi_{br} = \operatorname{atan2}(p'_{bz},\, p'_{bx}) \quad \text{with} \quad \mathbf{p}'_b = \mathbf{R}_z(\theta_{br})^\top \mathbf{p}_b, \tag{C3}$$

and the bracket-frame pose is

$$\mathbf{R}_{kbr} = \mathbf{R}_z(\theta_{br})\, \mathbf{R}_y(\phi_{br})^\top, \qquad \mathbf{T}_{kbr} = [\mathbf{R}_{kbr},\; \mathbf{p}_{br};\; \mathbf{0}\;\; 1]\,, \tag{C4}$$

so the frame's x axis aims along the nominal load path. The angles are computed per configuration from the bracket geometry, not fixed [Provenance: `minimizeFlxPin2brk.m` lines 252–298; `minimizeExtX3.m` lines 461–483]. As executed, however, the flexor evaluator of record assembles *both* of its bracket frames from the yaw rotation alone (the pitch-only, one-transform pose $\mathbf{T}_{kbr} = [\mathbf{R}_z(\theta_{br}), \mathbf{p}_{br}; \mathbf{0}\ 1]$): insertion bracket at line 289 and origin bracket at line 305 of `minimizeFlxPin.m`, with the matching pitch-only point reconstructions at lines 318 and 331; the two-rotation assembly of Eq. (C4) is computed but left commented at lines 288, 304, and 330. For the pinned extensor the two-rotation assembly is live (transMode default `2trans`; `minimizeExtX3.m` lines 3–6 and 479–482). Force is transformed into the bracket frame by $\mathbf{RowVecTrans}(\mathbf{T}_{kbr}^{-1}, \mathbf{F}_k)$ and deformed points are mapped back by $\mathbf{RowVecTrans}(\mathbf{T}_{kbr}, \cdot)$. Because the insertion array $[X_1, X_2, X_1]$ is y-symmetric, the assembly choice is numerically immaterial there; for the origin bracket it is not (see Table C6). The pitch-only (yaw-only) variant that the earlier one-transform fit used survives in `minimizeFlxPinX3.m` lines 200–209.

Table C5: Bracket reference points (mm, as coded; divided by 1000 in code). All are author-set CAD interpretation points; frame denotes the frame the vector is expressed in.

Point	x	y	z	**Configuration / meaning**
Pbri	−48.11	−107.81	13.80	pinned-knee flexor, insertion bracket (knee-ICR frame)
Pbr2	−52.61	0	75.06	pinned-knee flexor, origin-side 2nd bracket (hip frame)
Pbr	−3.84	−46.44	62.50	pinned-knee extensor, rib midpoint, medial side (hip frame)
Pbr	8.38	20.75	25.10	biomimetic extensor 10 mm, cantilever centroid
Pbr	9.48	−33.38	33.90	biomimetic flexor 20 mm, bolt-pattern centroid
Pbr	−19.00	22.00	27.60	biomimetic flexor 10 mm, cantilever centroid

C.4 Stiffness Arrays and Compliance

Every configuration models bracket compliance as a diagonal Hooke-law spring in the bracket frame, $\mathbf{K}_{br} = \mathrm{diag}(k_1, k_2, k_3)$ newtons per meter. The scalar compliance along the force path and its series combinations are

$$c_b = \hat{\mathbf{u}}\,\mathbf{K}_{br}^{-1}\hat{\mathbf{u}}^{\top}, \qquad c_{\mathrm{eff}} = c_b + c_{\mathrm{ten}}\ (+\,c_{b2}\text{ for the second bracket}), \tag{C5}$$

exactly as in Equations (4.8)–(4.18); when an artificial tendon is present its axial spring ($\times 2$ cables, effective steel area 1.51×10^{-6} m^2, $E = 193$ GPa) adds in series, and the force balance $F_{\max} F^*(\varepsilon^*(\delta), P^*) = k_{\mathrm{eff}}\,\delta$ is solved per frame by `fzero` on the contraction interval [Provenance: stiffness arrays: `minimizeFlxPin.m` lines 387 and 400; `minimizeExtX3.m` lines 575–584; `minimizeExt.m` line 352; `minimizeFlx.m` line 314. `fzero` force balance: `minimizeFlxPin.m` line 438; `minimizeExtX3.m` line 611; `minimizeExt.m` line 389; `minimizeFlx.m` line 352]. The per-configuration stiffness arrays and their series combinations are given in Table C6.

The wrapping-loss term X_3 (Equations 4.24–4.26) is evaluated as $X_3\,R\,|\Delta\theta|\,\zeta^2$ per contact arc. On the pinned extensor the two physical contact arcs use $R_1 = 12$ mm (27° to −23°) and $R_2 = 40$ mm, with arc-length normalizers $27.75/(ang_1 - ang_2)$ and $92.46/(ang_2 + 120^\circ)$; on the biomimetic configurations the bend measure is $w_{\mathrm{rap}}|\theta_{\mathrm{turn}}|$ with $w_{\mathrm{rap}} = 25$ mm [Provenance: `minimizeExtX3.m` lines 333–359; `buildKneeFlexorRoute20mm.m` line 268; `MonoPamDataExplicit_balanceX3.m` lines 489–494]. Table C7 records the optimizer parameterization and bounds of the adopted identification: the two-bracket flexor fit solves all three terms, and the extensor refit locks (X_1, X_2) to the adopted flexor-front values (the

Table C6: Bracket stiffness arrays by configuration ($K = [k_1, k_2, k_3]$ in bracket-frame axes, N/m; optimizer solves X_1, X_2 assigned as shown), as executed by the evaluators of record. The flexor evaluator has no frame switch: both brackets use the pitch-only assembly (see the Bracket Reference Points and Frames section above). The pinned-extensor evaluator carries a `transMode` argument (default `2trans`, env `EXTX3_TRANS`); its axis assignment was revised by me, with the previous assignment kept commented in the file.

Configuration	$\boldsymbol{K}$ array	**Source**
Pinned flexor, single bracket (legacy)	$[X_1,\ X_2,\ X_1]$	`minimizeFlxPinX3.m` line 267
Pinned flexor, insertion bracket (pitch-only frame)	$[X_1,\ X_2,\ X_1]$	`minimizeFlxPin.m` line 387 (frame li
Pinned flexor, origin bracket (pitch-only frame)	$[X_1,\ X_1,\ X_2]$	`minimizeFlxPin.m` line 400 (frame li
Pinned extensor, 2trans (with X_3)	$[X_1,\ X_2,\ X_2]$	`minimizeExtX3.m` line 575
Biomimetic extensor	$[X_1,\ X_2,\ X_1]$	`minimizeExt.m` line 352
Biomimetic flexor	$[X_1,\ X_2,\ X_2]$	`minimizeFlx.m` line 314

extensor front was generated with the pair fixed to the adopted values of Table C8 and only (X_0, X_3) searched), enforcing a single stiffness pair across configurations.

Table C7: Optimizer parameterization and bounds of the adopted identification. Length offsets are carried in centimeters and stiffnesses on a base-10 logarithmic scale; the initial population of the flexor run was drawn from $X_0 \in [5, 15]$ mm, $X_1 \in [5 \times 10^4, 5 \times 10^5]$ N/m, $X_2 \in [7 \times 10^3, 1.5 \times 10^4]$ N/m.

Identification	**Solved terms**	**Bounds**
Two-bracket pinned flexor	X_0, X_1, X_2	$X_0 \in [0,\ 20]$ mm; $X_1 \in [3 \times 10^4,\ 10^6]$, $X_2 \in [5 \times 10^3,\ 2 \times 10^4]$
Pinned extensor refit	X_0, X_3	$X_0 \in [-20,\ +5]$ mm; $X_3 \in [0,\ 1]$

[Provenance: `minimizeFlxPin10mm_2brk.m` lines 82–83 (bounds) and 424–432 (parameterization); `minimizeExt10mmX3.m` lines 23–26 (lock load) and 67–68 (bounds); `Collect_ExtPinX3_sweep.m` lines 59–60.]

C.5 Adopted Correction-Term Values

Table C8 records the correction-term values adopted at the time of writing. The initial single-bracket identification reported with the results figures (Section 4.2) was superseded by the two-stage identification described in Section 4.3.2: a two-bracket flexor identification over the four tests of Table C1 other than the 47 cm test, cross-validated with two tests held out per fold (the adopted fold trained on the 48 cm and 41 cm tests and validated on the 46 cm and 40 cm-tendon tests), followed by an extensor refit that locks (X_1, X_2) to a member of the same flexor Pareto front and solves only the extensor-specific terms (X_0, X_3) on its three-test training set with three tests held out for validation. The adopted values and their front-row provenance are recorded below.

Table C8: Adopted correction-term values of record. The flexor fit solves (χ_0, χ_1, χ_2); the extensor refit solves (χ_0, χ_3) with the stiffness pair locked to the flexor values.

	χ_0 (mm)	χ_1 (N/m)	χ_2 (N/m)	χ_3
Flexor	$+3.94$	3.998×10^4	1.473×10^4	
Extensor	-6.38	3.998×10^4	1.473×10^4	0.158

The flexor values are pooled-front row 77 of `minimizeFlxPin10_results.mat` (212 members); the extensor values are row 32 of `minimizeExt10mmX3_results.mat`, a front generated with (χ_1, χ_2) locked to that same row-77 pair (mat fields `LOCKSRC`, `xi1lock`, `xi2lock`, `flxRow`) with only (χ_0, χ_3) searched. The redesign context builders load exactly these values: `buildKneeFlexorContext20mm.m` pick 77 at lines 171–182, with χ_3 from the extensor front's row 32 at lines 188–191, and `buildKneeExtContext20mm.m` pick 32 at lines 164–171.

Table C9 reports the per-test goodness of fit at the adopted values. The corrected model improves on the uncorrected rigid model in every extensor fit-set test, and the fraction of variance unexplained stays below unity on every fit-set test except the 43 cm-tendon test (FVU 1.33, against a rigid-baseline FVU of 4.73). The three excluded tests (3, 4, 9) are not part of the fit of record. On the biomimetic extensor test ($l_0 = 51.8$ cm) the same values give RMSE 1.17 N m, FVU 0.67, and maximum residual 2.21 N m, against a rigid baseline of RMSE 3.03 N m, FVU 4.49, and maximum residual 5.59 N m; the shared-pair correction now explains the bulk of that configuration's measured torque variance as well as its level [Provenance: `Dig_out/xi_audit_tmp/reeval_gof_20260926.m` and `reeval_bio_20260926.m`; extensor rigid baseline from the `baselineScores` field of the extensor results mat].

Table C9: Per-test torque-prediction goodness of fit at the adopted correction terms of Table C8: RMSE and FVU of the corrected model, with the uncorrected rigid-model RMSE for the extensor tests. The 47 cm flexor test was not used in the identification; its corrected-model fit improves over the zero-correction prediction but is a post hoc evaluation.

Configuration	**Test**	**RMSE** (N m)	**FVU**	**Rigid RMSE** (N m)
Pinned flexor	48 cm (training)	2.16	0.17	
	46 cm (validation)	1.37	0.05	
	47 cm (not used)	2.29	0.16	
	40 cm-tendon (validation)	1.61	0.04	
	41 cm (training)	1.08	0.03	
Pinned extensor identification	43 cm (training)	0.98	0.16	2.44
	43 cm-tendon (training)	1.52	1.33	2.88
	46 cm (training)	0.49	0.05	2.18
	40 cm (validation)	1.24	0.13	3.09
	40 cm-tendon (validation)	1.35	0.89	3.62
	47 cm (validation)	0.70	0.17	0.95

C.6 Redesigned Extensor Route Geometry

This section documents the direct redesigned extensor route, the first closed redesign of the extensor configuration. It remains the reference geometry for the route-elimination machinery and the seed from which the later work proceeded, but it is no longer the extensor design of record: when hardware testing of the direct route began, the printed and machined parts started breaking under the route's higher actuator forces, so the operating force had to be lowered, and the cam-profile redesign of Section 4.3.2, which converts the direct route's thin torque margin into headroom at reduced force, was developed in response and supersedes it as the current design. The geometry below is recorded as saved by the optimization results file (`Vas_Pam_20mm_Result_20260925.mat`, the run cited in Table C2). Figure C1 reproduces the optimizer's tiled route-geometry view from that results file: eight pose panels sweep the 100-frame solution, drawing full flexion ($\theta_k = -120.0^\circ$), the solver frame immediately before each of the six contact-elimination events, and full extension ($+10.0^\circ$). The panel angles are the last poses at which the released row is still active; the release angles of Table C2 are the first poses at which each row is off, one pose spacing of 1.31° later. Beyond the release schedule, the results file also carries a top-level `removedText` field reading '`-p8`'. This field is a plotting leftover rather than a route statement: the plotting block of `Opt_run_Ext.m` (lines 730–736) computes `removedText` only as part of each interior pose-panel title, naming the point released in the next frame, and the value left in the results file is the one from the last interior panel, the p8 release-event panel at $+6.06^\circ$ in Figure C1. The route of record does contain p8: like p3, released at $+6.1^\circ$, p8 stays active from full flexion through 97 of the 100 design-grid

frames until its release at $+7.4^\circ$ near full extension (Table C2), so the seed table and release schedule document the route as run. Each panel draws the active p1:p9 route in the femur frame (tibia-side rows transformed through the ICR) over the femur-cylinder, femur-wall, corrected-condyle, and tibia clearance envelopes of the route-elimination rule, with the local bend-radius checks at p2 and p8 drawn dashed. The optimized endpoint values of this design are $\mathbf{p}_1 = [40.0,\ 137.0,\ 0]$ mm (femur frame; the square marker of Figure C1) and $\mathbf{p}_{\mathrm{End}} = [34.4,\ -197.1,\ 0]$ mm (tibia frame; the diamond marker), with resting length 827.7 mm and 36.8 mm of tendon.

Figure C2 places the same optimized route on the OpenSim bone-surface point clouds (`Open_Sim_Bone_Geometry`) at exactly zero knee angle: the full skeleton (femur, tibia, pelvis and hip, sacrum, lumbar spine, and foot) is drawn behind the three human vastus reference paths in light orange, the optimized p1-to-pEnd route is drawn in magenta with its endpoints circled, and the optimized route carries a total BPA-plus-tendon path length of 0.784 m at this pose. Animated sweeps of both redesigned routes over the full flexion range are archived with the optimization results (`Knee_Extensor_20mm.gif`, 2.3 MB, and `Knee_Flexor_20mm.gif`, 4.7 MB, in `Code/Matlab/Mesh_Optimization/Results`).

C.7 Supporting Route Searches and Transmission Sweeps

The following diagnostics describe earlier optimizer runs and integration status. The five-site batch table concerns the original unity-transmission search; it is distinct from the subsequent size re-solves plotted below. Objective values from different formulations are not directly comparable.

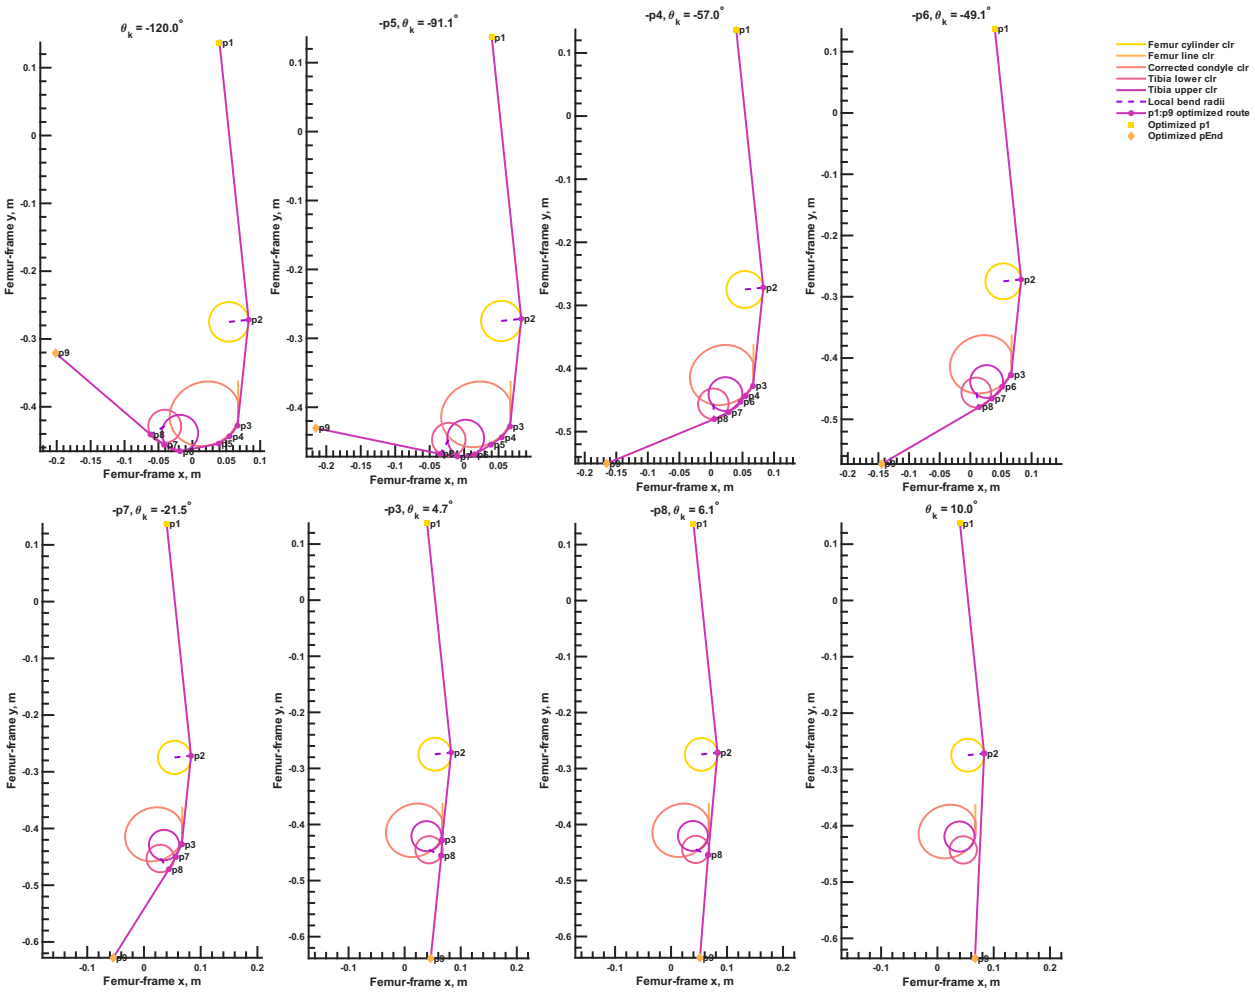


Figure C1: Eight-panel route-geometry view of the optimized 9-point extensor route of record (`Vas_Pam_20mm_Result_20260925.mat`), drawn in the femur frame. Regenerated by `Export_AppendixRouteGeo.m`, which ports the plotting block of `Opt_run_Ext.m` verbatim.

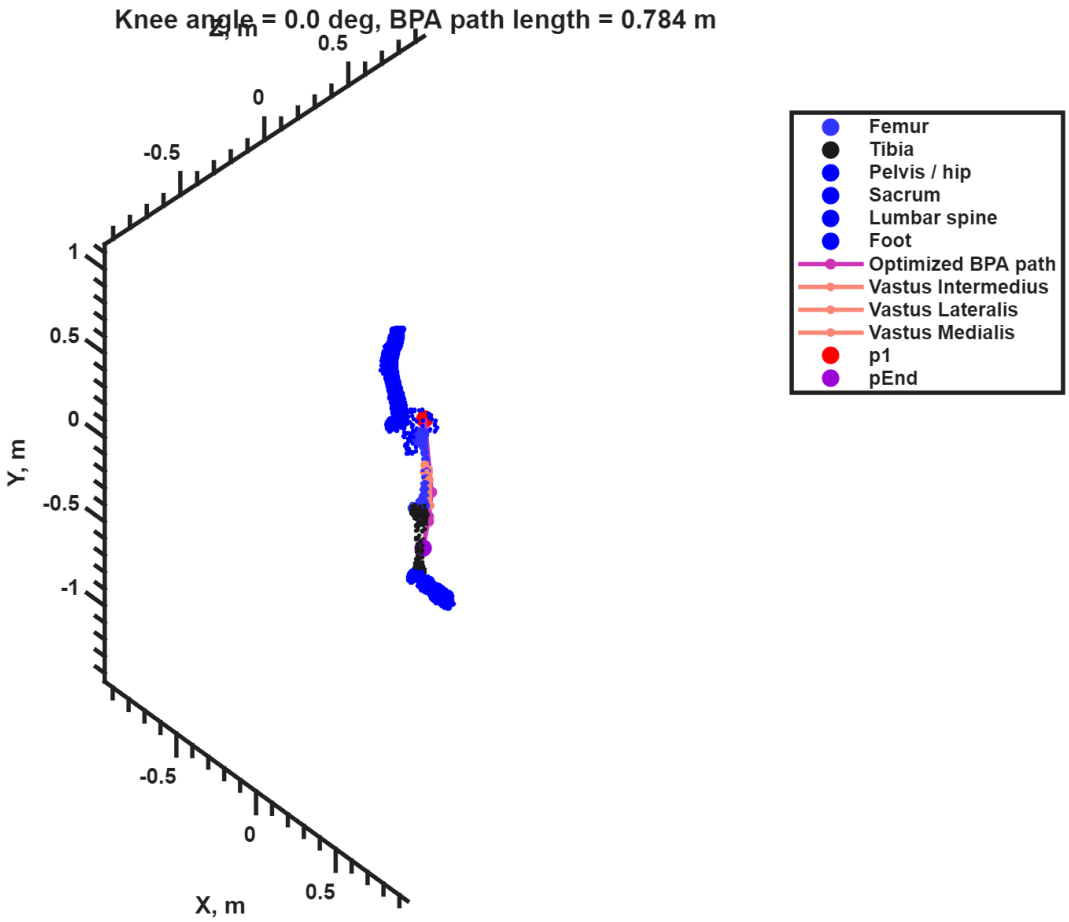


Figure C2: Static bone-mesh view of the optimized 20 mm extensor route at $\theta_k = 0^\circ$, rendered by `Export_AppendixBoneMesh.m`.

The route redesign of Section 4.3.2 placed both 20 mm actuators on direct attachment routes. The Mono/Bi-pulley algorithms extend that search space by allowing the artificial tendon to be reeved over pulleys before insertion, and the flexor pulley run quantifies what the added routing freedom is worth on the knee-flexor route. The pulley variant converged to a collision-feasible layout with zero bound violations (pattern-search exit flag 1) at an objective of 0.08465 under its own objective function, which differs from the direct route's objective of 74.47, so the two values are not commensurable. The best configuration on file reeves the two actuators over a single tackle-line part at a gain fixed at 1:1, so the pulley acts as a redirect rather than a force multiplier; even so, the redirect alone deepened the peak predicted flexion torque from 81.23 N · m to 85.74 N · m, a gain of 4.51 N · m (5.6 %), while holding the off-axis torque to at most 0.803 N · m, under 1 % of the peak on-axis value; the predicted peaks of the direct and pulley routes exceed the human flexor torque scale of 23.16 N · m by factors of 3.51 and 3.70. The Bi-pulley batch establishes where that freedom runs out on anatomical routes. Across the five OpenSim muscle sites attempted (soleus, tibialis anterior, rectus femoris, medial gastrocnemius, and biceps femoris short head), every run closed with a nonpositive torque margin (objectives of 0.929 for tibialis anterior, 0.997 for soleus, and 3.42 for rectus femoris with a residual worst-constraint value of 0.242), and the medial-gastrocnemius and biceps-femoris-short-head runs terminated at the infeasibility sentinel (objective 101, worst-constraint value 10, exit flag -2). For the four runs that report a finite route-to-route clearance, the binding constraint was geometry rather than torque: minimum pair clearances spanned 4.26 mm to 19.14 mm against the 25 mm separation required between adjacent routes (two 10 mm actuator radii plus a 5 mm bundle margin); the medial-gastrocnemius run records no finite clearance, so its failure

rests on the infeasibility sentinel alone. What the two algorithms buy is therefore reach and screening: the mono-pulley converts routing freedom into 4.51 N · m of additional predicted flexion torque without mechanical advantage and with off-axis torque below 1 % of the peak, and the bi-pulley sweep screens candidate anatomical sites, showing that its reeved routes do not yet fit these corridors and redirecting the design effort toward clearance-relaxed attachment sites. These pulley solutions are optimizer-stage results; they have not yet been propagated into the OpenSim or MuJoCo model trees used for the simulation work of Sections 6.3–6.4, and the hardware torque tests of the updated compliant system have begun with the extensor route.

The Bi-pulley batch applied the same reeving freedom to five muscle sites of the robot-body model's OpenSim leg, and all five runs failed the torque requirement, with the medial-gastrocnemius and biceps-femoris-short-head runs terminating at the infeasibility sentinel (Table C10). For the biceps femoris short head site, whose torque target is 40.2 N · m at the robot-body model's maximum isometric force of 804 N, the pattern search exited with flag -2 and returned the sentinel objective of 101 from an initial objective of 1000; Section 4.3.3 reports the per-run objectives and the clearance constraints behind these outcomes.

The robot-body model itself carries one update from this work. The adductor magnus attachment point P2 was restored to its full-precision location $(-0.1357, -0.0929, 0.0591)$ m, displacing a rounded copy that was disqualified for producing a zero-length via point that nullified the flexion torque; the regenerated mirrored model passed its validation with zero non-mirrored pairs. At the restored location, the muscle delivers 34.28 N · m of flexion torque, 1.12× the human target of 30.61 N · m, and 27.62 N · m of adduction torque,

Table C10: Bi-pulley batch on five muscle sites of the robot-body model's OpenSim leg. Objective 101 is the batch's infeasibility sentinel; all five runs closed with a nonpositive torque margin. The soleus and tibialis-anterior runs converged feasibly (pattern-search exit flag 1), and their objectives are the normalized torque shortfalls; the remaining three runs exited with flag -2.

Muscle site	**Objective**	**Outcome**
soleus	0.9968	converged (flag 1), nonpositive torque margin
tibialis anterior	0.9288	converged (flag 1), nonpositive torque margin
rectus femoris	3.423 (worst constraint 0.2423)	nonpositive torque margin
medial gastrocnemius	101, infeasibility sentinel (worst constraint 10)	pattern-search exit flag -2
biceps femoris short head	101, infeasibility sentinel (worst constraint 10)	pattern-search exit flag -2

$0.82\times$ the human target of $33.60\,\mathrm{N}\cdot\mathrm{m}$. These redesigned routes remain optimizer-stage results: the robot-body OpenSim model contains no pulley entities, and no via points or markers were added to it for these routes (its eight pre-existing moving-path-point vias, the rectus-femoris and vastii patella-tracking points, one per muscle per leg, belong to the original muscle paths). The MuJoCo conversion used for the simulation work predates the pulley batch and the adductor-magnus restoration, so it carries one stale path point, the pre-restoration rounded adductor-magnus P2. Physical torque tests of the redesigned routes have begun (Section 4.3.3).

The same redesign pipeline was then extended from the knee to all 27 muscle-actuators of the robot-body model, and this whole-body extension is reported as work in progress rather than a completed result. The first pass reached 19 of the 27 actuators before the run stopped, and its chosen solutions showed that the penalty-based formulation itself was

the binding limitation: all but one of the chosen routes violated the feasibility constraints, the sole exception being the biceps femoris long-head pair, whose worst-angle margin of $+0.03\,\%$ is indistinguishable from zero. The formulation was then restructured so that the task envelope, the joint-angle window each route must serve, acts as a hard feasible set rather than a penalty. In the restructured pass, twelve of the 27 actuators terminated on a concatenation error in the route builder (the three lumbar-wall actuators, tibialis posterior, the seven distal-foot actuators, and adductor magnus), leaving fifteen finalized routes. Eight of the fifteen were chosen from constraint-feasible configurations: semimembranosus, biceps femoris long head, tensor fasciae latae, soleus, lateral gastrocnemius, tibialis anterior, iliacus, and psoas. Of those eight, six also met the torque requirement at every angle, with worst-angle margins spanning $+0.25\,\%$ (biceps femoris long head) to $+250.7\,\%$ (tensor fasciae latae); the semimembranosus and soleus routes satisfy the geometry but remain short of their torque targets. The crash set and the two torque-short feasible routes define the remaining engineering: the route builder needs a fix for the concatenation failure, and the short routes need either larger actuators or relaxed corridors.

The unadjusted cam optimizer solution had minimum margins of $+31.30\,\%$ against the required envelope and $+42.34\,\%$ against the plain human target; the adjusted design used in Chapter 4 has margins of $+30.83\,\%$ and $+39.77\,\%$, respectively.

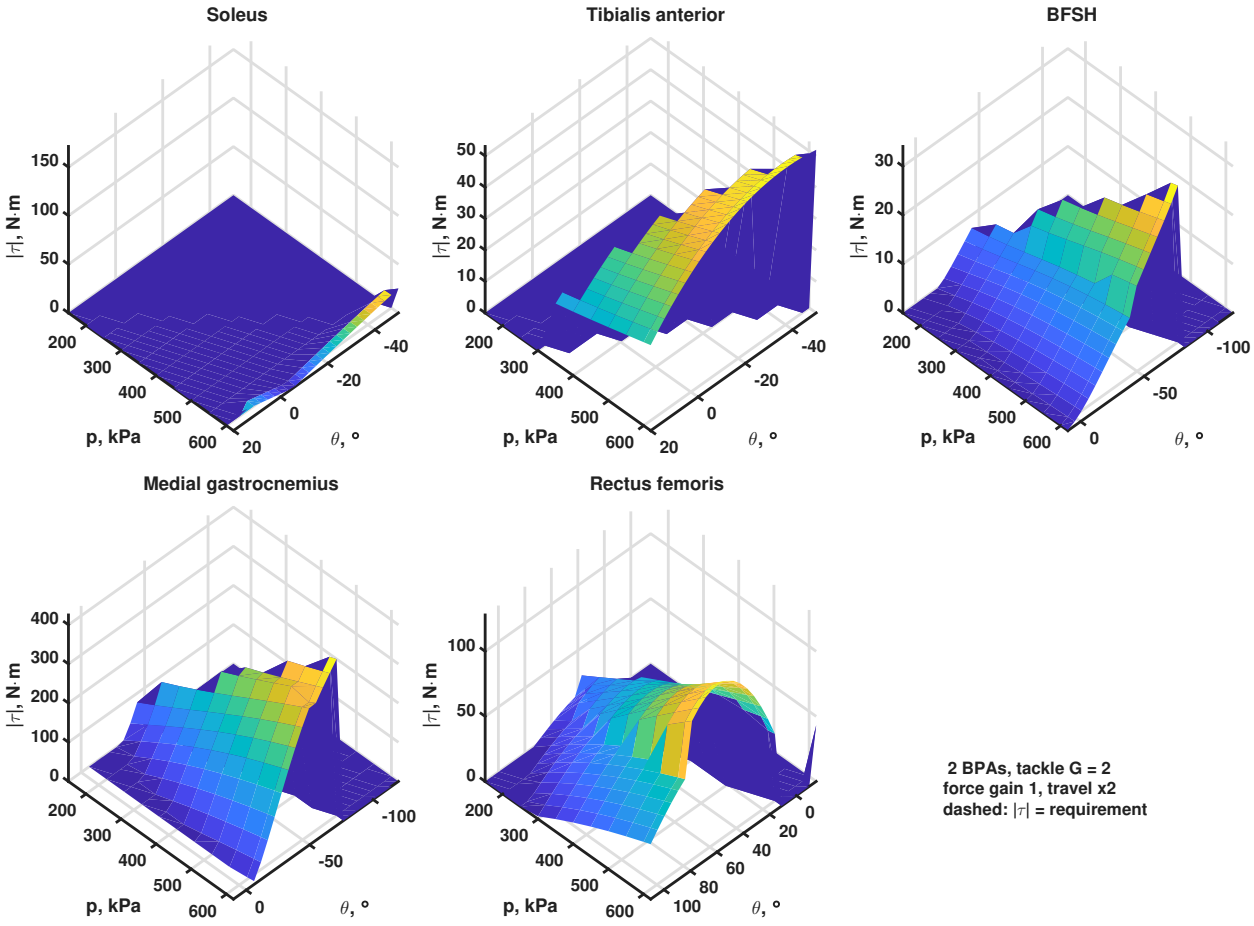


Figure C3: Predicted insertion-torque magnitude versus joint angle and supply pressure for two BPAs through a two-part tackle (force gain 1, travel gain 2), with actuator sizes re-solved. Panels show soleus, tibialis anterior, biceps femoris short head (BFSH), medial gastrocnemius with ankle fixed at zero, and rectus femoris evaluated at the hip. Dashed curves mark equality with the maximum-isometric OpenSim requirement; no equality curve lies within the soleus panel.

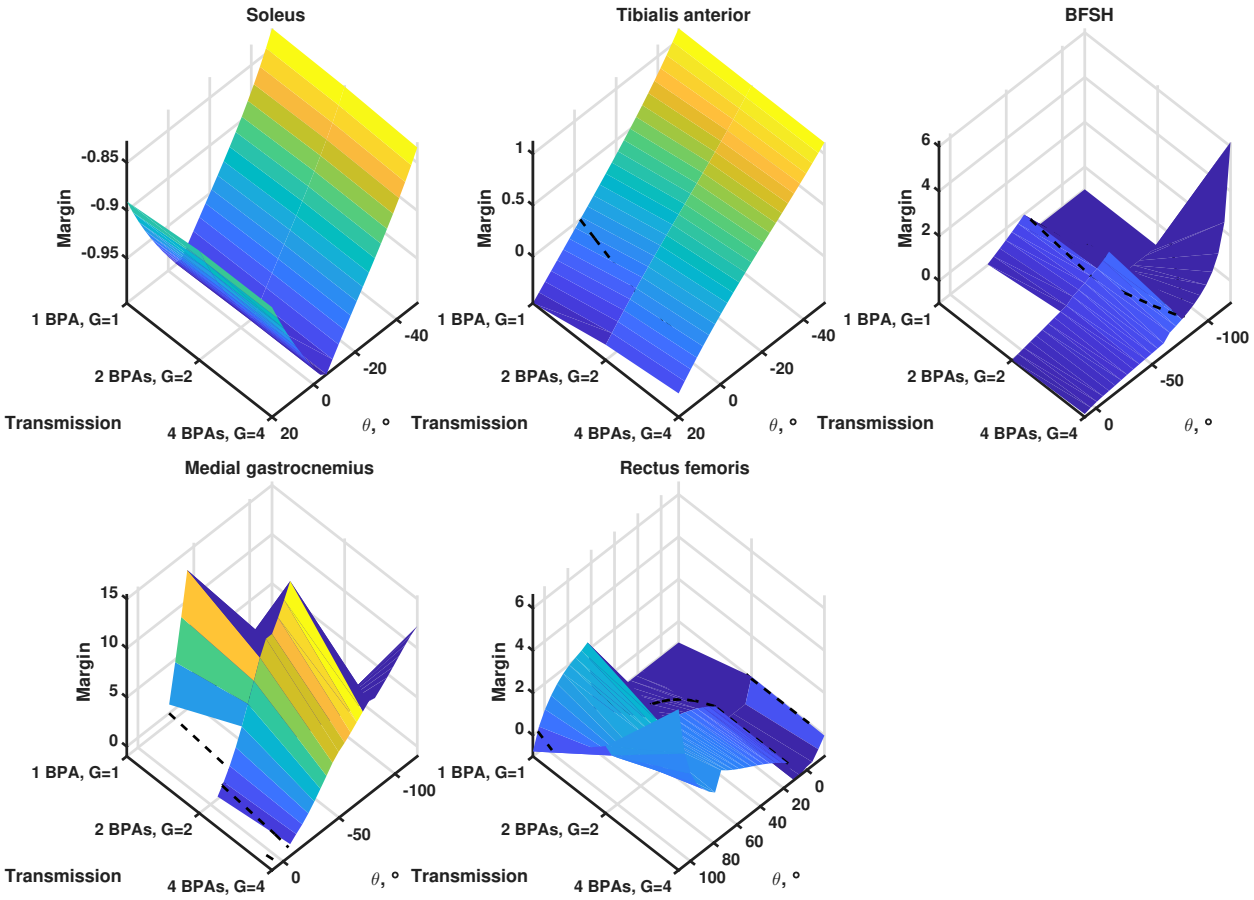


Figure C4: Torque margins at 620 kPa for the single-BPA baseline, two BPAs through a two-part tackle, and four BPAs through a four-part tackle. Actuator sizes are re-solved for each configuration. A positive fractional margin exceeds the maximum-isometric OpenSim requirement; the dashed line marks zero. Panels and evaluation planes follow Figure C3.

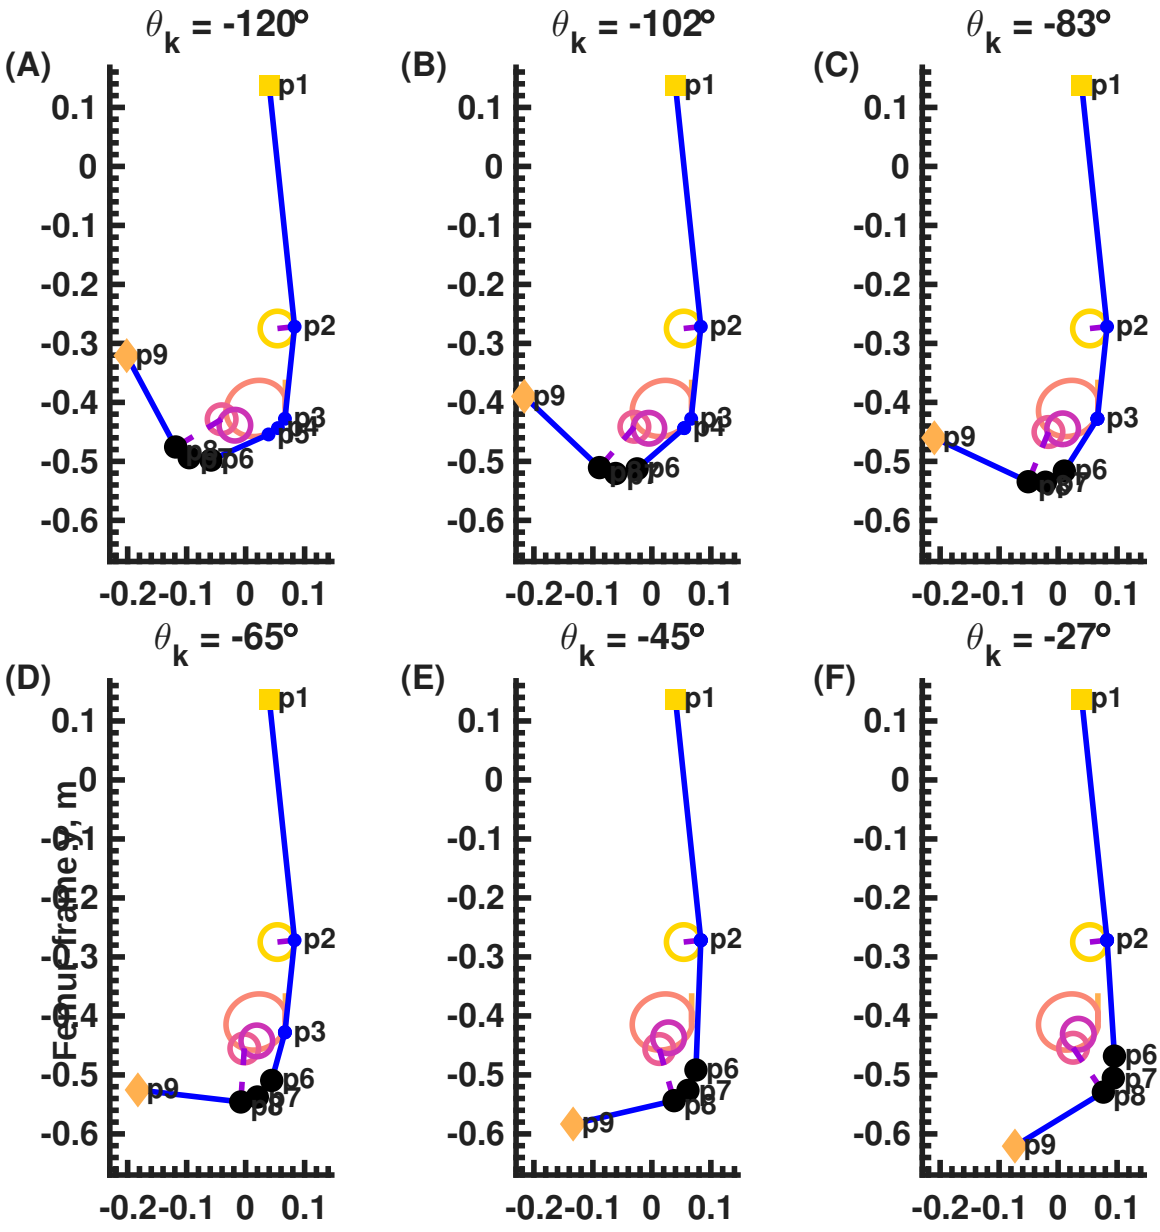


Figure C5: Cam-extensor poses A–F, spanning $-120°$ to $-27°$. Blue lines show the route; colored outlines show the clearance envelopes. The remaining poses and original legend are in Figure C6.

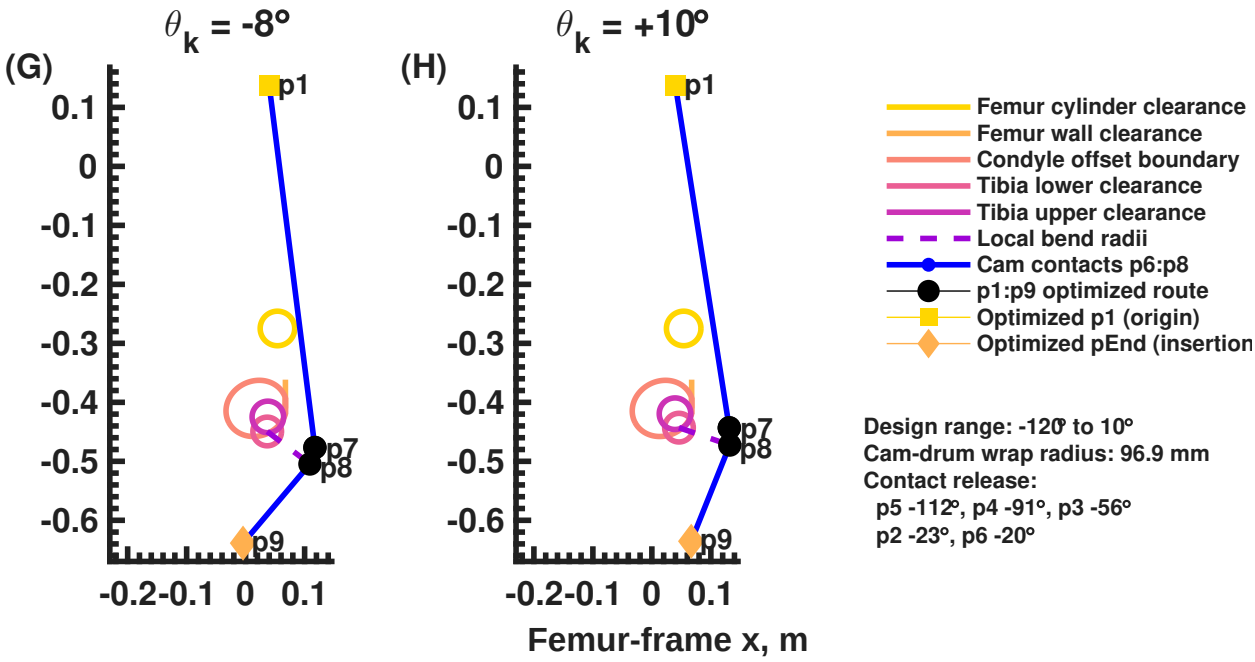


Figure C6: Cam-extensor poses G and H at -8° and 10°, with the original legend, 96.9 mm cam-drum radius, and contact-release schedule. Panel letters continue Figure C5.

C.8 Supporting Torque-Prediction Fit Statistics

The earlier cross-configuration fit below is distinct from the adopted two-stage identification. The pre-optimization comparison and the adopted pinned-flexor and pinned-extensor per-test fits appear in Chapter 4 (Tables 4.1, 4.2, and 4.3).

Table C11: Goodness of fit for the improved torque model across test configurations. RMSE and maximum absolute residual are in $\mathrm{N} \cdot \mathrm{m}$.

Configuration	**RMSE**	**FVU**	**Max. residual**
Flexor, pin joint, $l_0 = 48.5\,\mathrm{cm}$ (training)	1.616	0.0958	3.2527
Flexor, pin joint, $l_0 = 45.7\,\mathrm{cm}$ (validation)	2.1214	0.0878	5.1690
Extensor, biomimetic knee, $l_0 = 51.8\,\mathrm{cm}$	0.8920	0.3396	1.4188
Flexor, biomimetic knee, $\varnothing 10\,\mathrm{mm}$, $l_0 = 41.5\,\mathrm{cm}$	0.3007	0.0102	0.4647
Flexor, biomimetic knee, $\varnothing 20\,\mathrm{mm}$, $l_0 = 41.5\,\mathrm{cm}$ (620 kPa)	1.8024	0.0499	3.5678

Appendix D: Supplementary Neuromechanical Verification

This appendix supplies supporting model records and evaluation sets for Chapter 6. It separates implementation checks, gait comparisons, pathway-removal tests, and spiking conversions because they test different claims. Neuron equations and calibration tables are provided in Appendix B.

D.1 Model Construction and Recording

The archived AnimatLab projects provide reference circuits for implementation comparisons. Figure D1 shows their neural organization. A later reorganized bilateral project added commissural and biarticular pathways; it ran headless but latched under the tested drives. That assembly result is not evidence of stepping. The phase-1 recordings below concern an earlier instrumented configuration, not the reorganized network.

A 10 s instrumented run recorded joint, neural, sensory, and body-motion channels at a 0.2 ms interval. Contact-sensory recordings are membrane voltages, not contact forces. The initial external horizontal force and the limited window prevent the joint traces from establishing stable walking.

D.2 Controller Ports and Reference Checks

The matched W2L test compares implementations from a common circuit specification under identical drive and contact inputs, using burst period, bilateral phase, burst count,

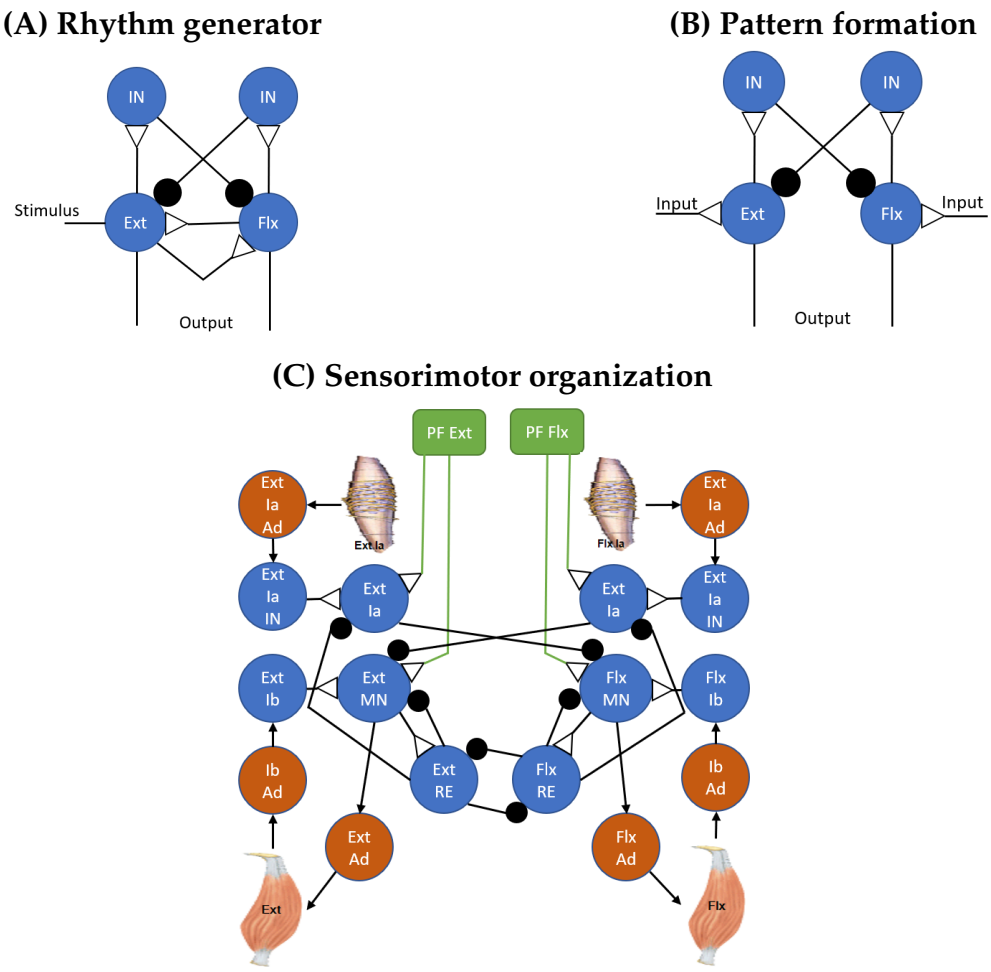


Figure D1: Archived AnimatLab motifs: rhythm generation (A), pattern formation (B), and sensorimotor organization (C). Ext and Flx denote antagonist pathways; IN denotes interneurons, MN motor neurons, and Ia and Ib muscle afferents. Triangles and circles distinguish excitatory and inhibitory connections.

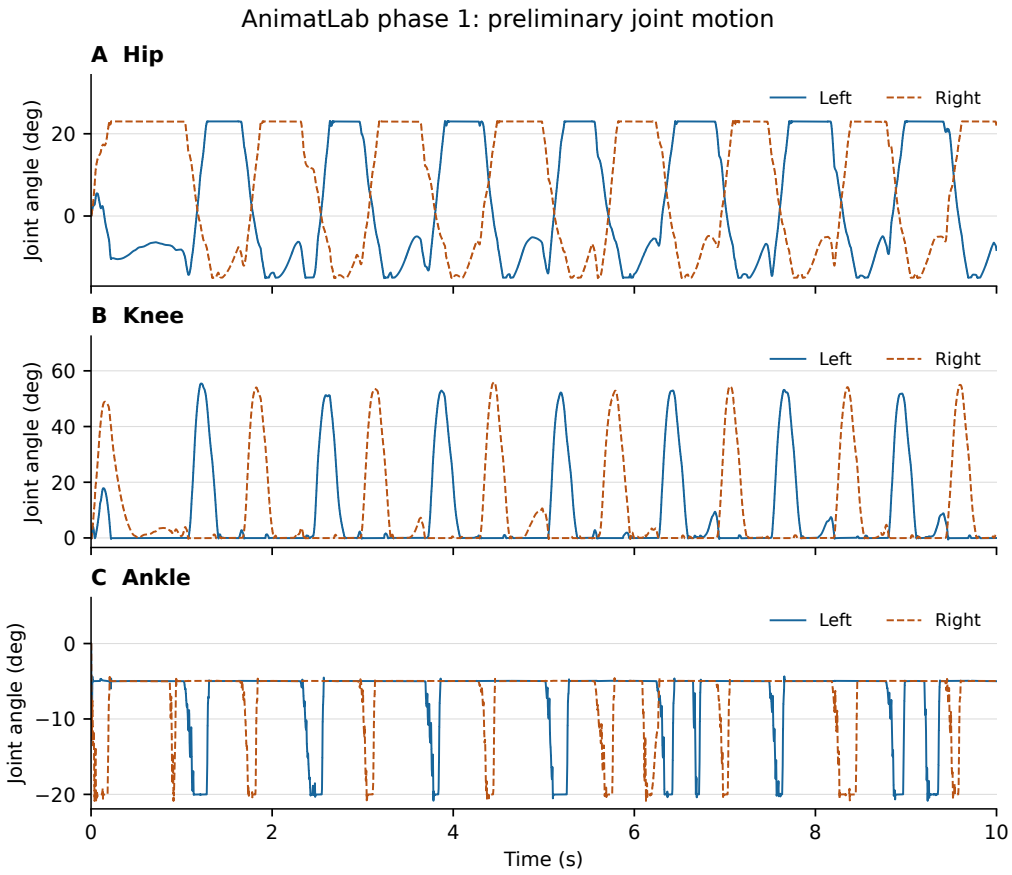
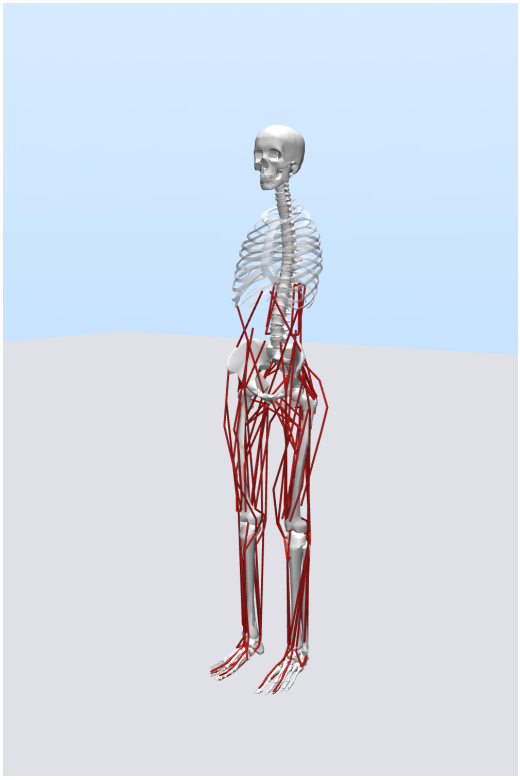


Figure D2: Bilateral hip (A), knee (B), and ankle (C) angles in the recorded joint conventions. All 50,001 samples in the 10 s window are retained. An external horizontal force acts during the first second. The recordings show motion, not validated walking or ground-contact gating.

(A) Converted robot



(B) Synthetic BPA knee

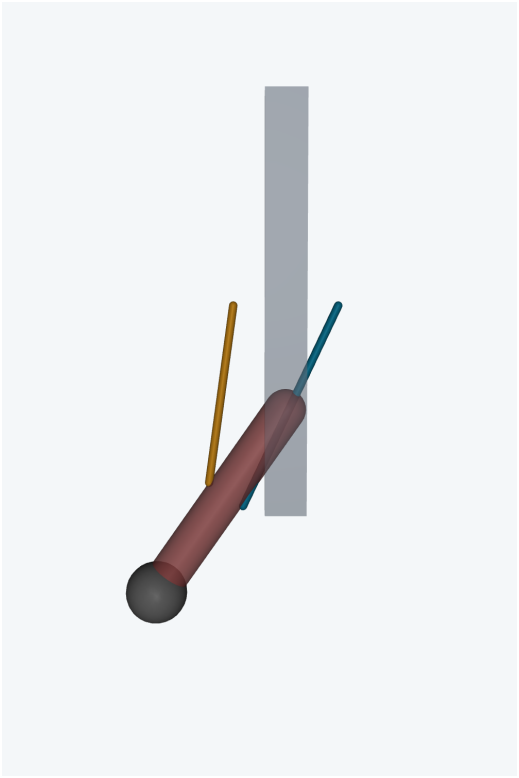


Figure D3: Converted Gait2392 robot at its reference pose (A) and independent synthetic antagonist-BPA knee at -35° (B). Cyan and gold identify flexor and extensor tendon routes. The images show construction, not controlled trajectories or validated torque predictions; scales differ.

and voltage correlation. The Deng controller port was checked separately: a current-pulse test produced a 1.938 s period against a 1.94 s NumPy reference, with antagonist correlation $r = -0.86$. Agreement on these tests does not establish long-term phase agreement in every network.

The full exported network agreed with its NumPy reference to 4.2×10^{-6} mV over the early transient across 410 neurons. Its subsequent divergence and the tonic-drive fixed-point discrepancy reported in Chapter 6 limit parameter transfer. The OpenSim-derived Simscape mechanism contains 1,446 blocks and 149 joints; mechanism compilation is distinct from completing its neural and muscle coupling.

D.2.1 Archived AnimatLab Models

Headless reruns tested whether archived models reproduced their reference behavior in the original simulator. Eleven targetable rhythm and gait metrics passed their acceptance bands. The bilateral contact-driven model's rhythm generators had periods of 0.450 and 0.467 s; the modern W2L model showed joint oscillations at 2.25 Hz. The Li-model rerun reached 0.637 m s^{-1} with step periods near 1.3 s [**?**]. These are reference-model checks, not results for the combined Deng–Li model or evidence of unsupported walking.

D.3 Tuning and Supported Evaluations

Free network gains were searched with a tree-structured Parzen estimator, with each candidate evaluated through simulation. The staged procedure separates suspended rhythm generation, sensory feedback, standing, and ground evaluations (Figure D4).

Finalists were checked under 1 % multiplicative parameter perturbations and changes in drive. Such checks probe the neighborhood of a selected configuration; they do not establish robustness to untested loads or removal of the harness.

The cycle-shape metric assigns zero to perfect agreement. A run without detectable cycles returns -25 as a failure sentinel. Because this value is numerically greater than the genuine-gait scores below, failed runs must be identified before ranking scores. The sentinel itself is not a valid measure of gait quality.

No.	Preparation	Capability	Score
1	**Deafferented air stepping**	Intrinsic rhythm; feedback off	Rhythm and swing (rhythm-gated)
2	**Afferented air stepping**	Heel, toe, and load feedback active	Rhythm and swing (rhythm-gated)
3	**Standing balance**	Posture under push perturbations	Posture maintenance under push
4	**Ground walking**	Ground contact and forward gait	Kinematics vs. human reference cycle
5	**Contact-driven ground walking**	Contact-driven pathway variant	Kinematics vs. human reference cycle

Each winner seeds the next stage. Finalists undergo perturbation and basin checks. Air stages are suspended; ground evaluations retain the support harness.

Figure D4: Staged tuning procedure. Air-stepping evaluations are suspended, and the reported ground evaluations retain mechanical support. Each stage introduces additional pathways or tasks. Scores require a separate rhythm/failure check because the no-cycle sentinel cannot be ranked as a gait score.

D.3.1 Reference-Library Comparison

The 43-cycle library contains 32 running trials from eight subjects at four speeds, ten walking references, and a predicted-walking cycle. The walking references span prescribed speeds of $0.5\,\mathrm{m\,s^{-1}}$ to $2.0\,\mathrm{m\,s^{-1}}$ and self-selected trials. The four configurations were scored without treating a good tuning-reference score as evidence of generalization. Library scores were substantially more negative (Table D1), and the evaluations showed unilateral or asymmetric gait. This comparison bounds transfer to the available references; it does not validate normal human walking.

Table D1: Gait-validation sweep of the four production walking configurations against the 43-cycle locomotion reference library. The score is the production pipeline's cycle-shape kinematic score (0 denotes perfect agreement); the tuning-reference score is each configuration's evaluation on the tuning pipeline's walking reference, and the library mean averages the same score over all 43 references (among genuine gait cycles, higher scores indicate closer agreement).

	Tuning-ref. score	**Library mean score (43 refs)**	**Most negative reference (score)**
`s3k`	−186.4	−8776	`Case_40_motion` (−8841)
`s3kpruned`	−176.1	−6924	`Case_40_motion` (−6986)
`syn6`	−200.4	−9961	`Case_40_motion` (−10040)
`w2lvar`	−144.5	−5808	`subject08_Run_20002` (−5896)

D.3.2 Pathway-Removal Tests

Individual pathway classes were removed from each selected configuration and the model re-evaluated on its tuning reference (Table D2). Removing Ia feedback from the production configuration eliminated detected cycles. Other removals changed the score without

necessarily eliminating movement. The different response of the six-synergy variant shows that these dependencies belong to the tested network configurations, rather than constituting a general claim about biological necessity. The systematic pattern-formation and motoneuron stimulation series was not completed.

Table D2: Pathway-removal sweep on the tuning-reference evaluation. The score is the production pipeline's cycle-shape kinematic score (0 denotes perfect agreement; −25 is the fixed no-cycle fallback, a failed evaluation that scores above every genuine gait in this table). The full-model rows repeat each configuration's unmodified production evaluation.

Configuration	**Removed pathway**	**Score**
`s3k` (full model)	none	−186.4
	Ia afferents	−25 (no cycles detected)
	all afferent classes	−267.8
	Renshaw cells	−267.3
	group-II afferents	−255.2
	Ib afferents	−180.3
	contact channels	−183.9
	interleg commissural pathways	−174.9
`syn6` (full model)	none	−200.4
	all afferent classes	−190.5
`w2lvar` (full model)	none	−144.5
	Ia afferents	−273.5
	weak rhythm-generator mutual excitation	−279.9
	Renshaw cells	−142.4

D.4 Spiking-Neuron Comparisons

The spiking variants replace afferent, rhythm-generator, and interneuron populations while retaining graded motor neurons. Synapses onto motor neurons are calibrated by mean conductance, preserving the continuous activation interface. This construction

does not require the network dynamics or resulting movement to remain unchanged. Calibration and full conversion tables are in Appendix B.

D.4.1 MuJoCo Network and Ground Evaluation

Under constant drive, the non-spiking and spiking rhythm periods were 0.399 and 0.621 s, respectively (Figure D5). The spiking period is 55.6 % longer. Timing agreement therefore requires more than copying the selected parameters. The untuned spiking variant produced bilateral cycling in a ground evaluation where the non-spiking winner cycled predominantly on one side (Table D3; Figure D6). This individual comparison motivates further controlled testing; it does not establish superiority of spiking control.

Table D3: Supported ground evaluation at the same selected parameters (16 s), with no spiking re-tuning. Knee minima and contact duty are computed over detected cycles. The kinematic score is a cycle-shape error metric: zero denotes perfect agreement, and less negative scores indicate closer agreement.

Metric	**Non-spiking winner**	**Spiking mirror (untuned)**
Kinematic score	-161.57	-236.35
Gait cycles, right / left	9 / 0 (left frozen)	8 / 9 (bilateral)
Cycle knee minimum	-64.3°	-79.8°
Contact duty, right / left	0.637 / 0.985	0.799 / 0.856
Maximum pelvis tilt	22.1°	19.2°
Non-finite states / fell	no / up	no / up

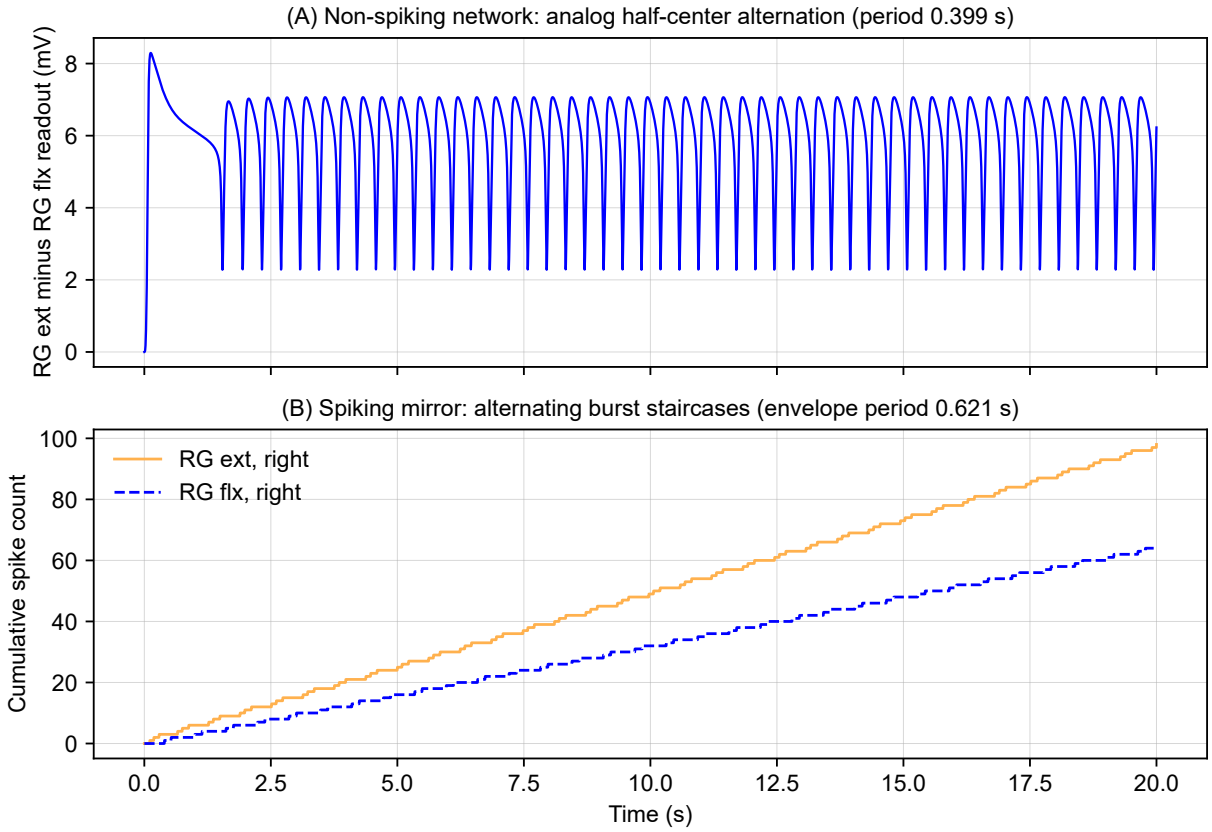


Figure D5: Constant-drive rhythm in the non-spiking network (A) and spiking variant (B). The observed periods are 0.399 and 0.621 s; the latter is 55.6 % longer. The panels use different neural readouts: graded population difference versus individual spiking cells.

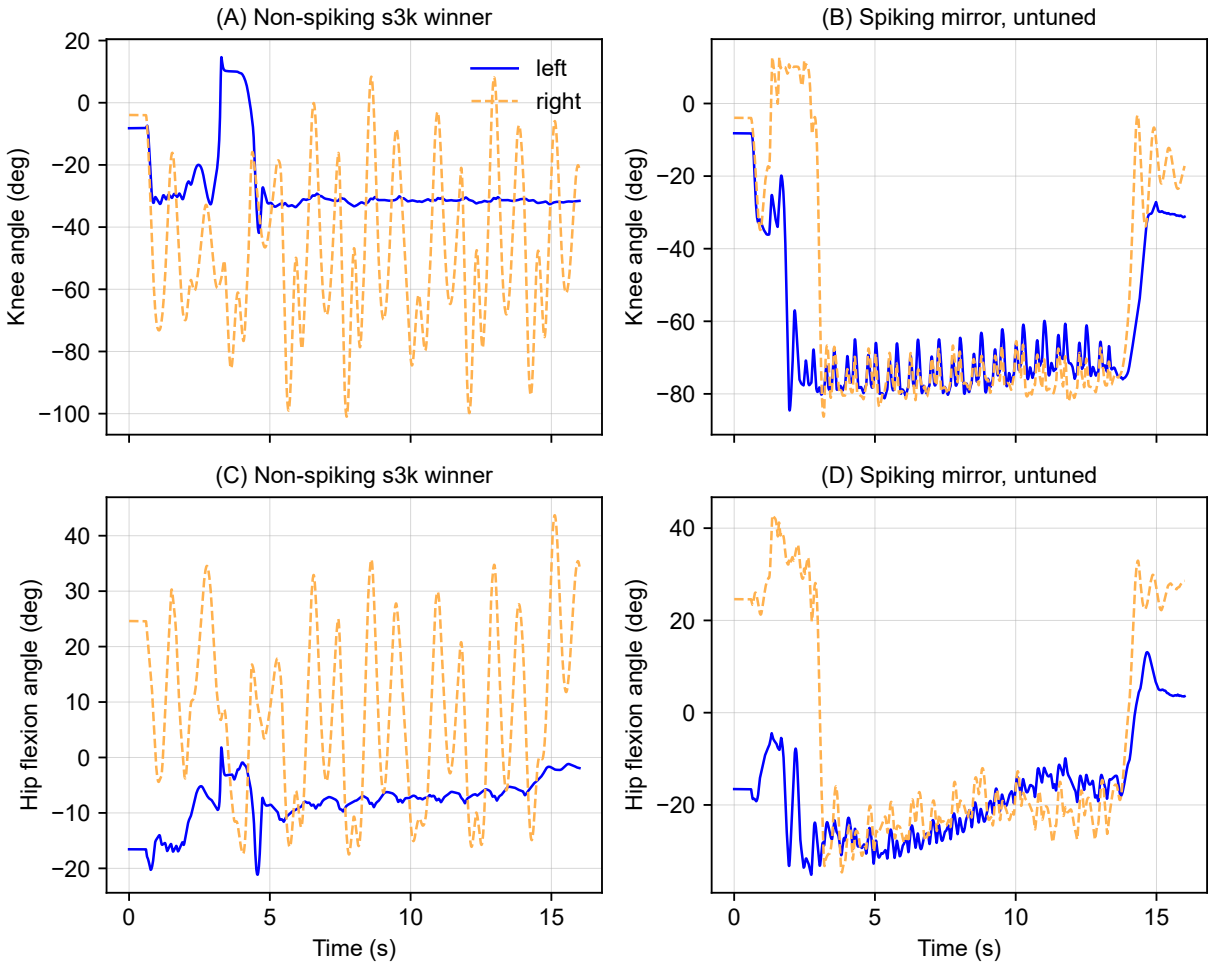


Figure D6: Ground-evaluation joint angles at the non-spiking winner parameters. The non-spiking model cycles predominantly on the right (A, C), while the untuned spiking variant cycles both legs (B, D). Solid indigo denotes left, and dashed orange right. Both use the supported evaluation protocol.

D.4.2 Simulink Demonstrations

The graded and spiking knee-reflex demonstrations were compared over matched 5 s runs (Table D4; Figure D7). The beverage example retained approximately the same sag, 2.02° versus 2.00°, compared with 7.85° without the reflex. These examples demonstrate transfer of particular reflex responses after calibration, not invariance of all controlled behavior to neural coding.

Table D4: Knee-reflex demonstration, graded baseline versus spiking twin at the matched operating point (5 s runs; settle statistics over the last 2 s).

Metric	Baseline	Spiking twin
Knee angle at 0.05 s	15.79°	15.79°
Settle mean (last 2 s)	43.5°	44.4°
Settle range	41.6°–44.9°	41.9°–46.7°
Extensor / flexor activation (window means)	0.53 / 0.53	0.58 / 0.56
Antagonist alternation	0.5 Hz–2 Hz	7.4 Hz (dominant peak)
Ia interneuron firing at settle	graded, subthreshold	95 Hz; Ib 0 Hz

D.4.3 AnimatLab Conversions

The tested spiking conversions did not preserve the graded models' rhythm (Table D5). Converting thresholds while retaining graded synapses preserved near-baseline oscillations, whereas converting synaptic amplitudes and decay changed or eliminated alternation. These results identify the synapse conversion as a limitation of this mapping. They do not demonstrate that a suitably calibrated spiking implementation cannot reproduce the reference behavior.

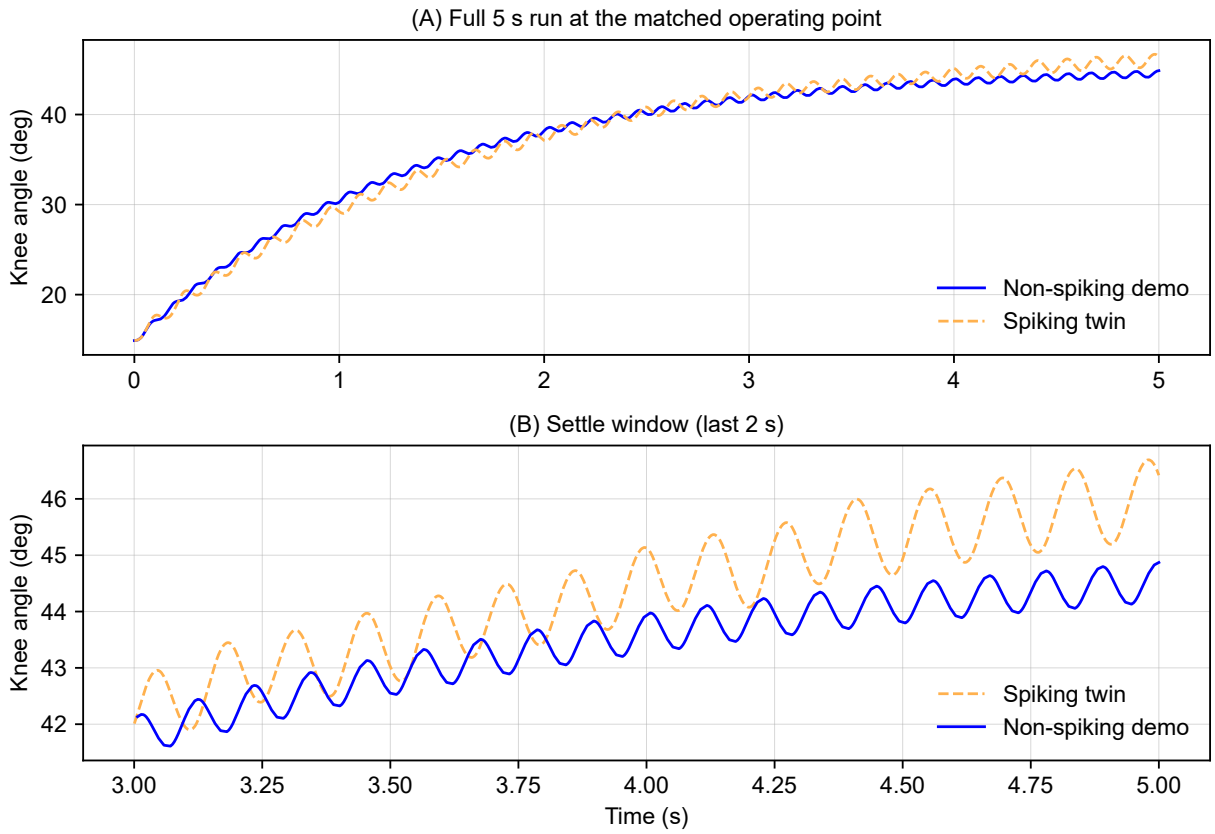


Figure D7: Knee-reflex demonstration at the matched operating point: graded baseline (solid indigo) and spiking variant (dashed orange). Full 5 s run (A) and final 2 s window (B). The spiking response has a slightly wider oscillation band.

Table D5: AnimatLab spiking conversions versus their graded baselines (5.1 s headless runs; joint ranges in degrees where the rhythm survives). The thresholds-only ablation converts the neurons but keeps the graded synapses.

Variant	Baseline (graded)	Spiking copy
Two-layer walker	0.443 s rhythm; hip -15 to $+23$, knee -1 to $+60$	dead (2 spikes; hip pinned near $+6$)
Bilateral variant	0.462 s; antiphase -0.42 cycle	latched; one side at 442 Hz–450 Hz
Reorganized project (standalone export)	latched (tonic drives)	synchronous 11.8 Hz co-bursting; same-leg envelopes $r = 1.000$ at lag 0
Thresholds only (graded synapses)	hip -15 to $+23$, knee -1 to $+60$	oscillates: hip -15.1 to $+23.2$, knee -0.8 to $+47.9$

D.5 Supported BPA-Robot Evaluations

The BPA-actuated walker exercised the chain from neural activation through the actuator force model to multi-joint motion. In the supported air-stepping run, the harness carried 70 % of the 411 N body weight. The period was 1.542 s, knee excursion was 63.5°, and peak forces were 880 and 1372 N, at a peak pressure of 372 kPa. This was a supported simulation, not a physical walking validation.

With split left/right drives and retuned feedback, the W2L robot progressed 1.72 m over 30 s under the support protocol. It retained progression at 60 % drive and passed 12 of 12 perturbed reruns, but did not satisfy the strict contact-duty walking criterion; below that drive the gait ceased. The mechanical support was not weaned.

In the separate human-model drive-weaning study, the selected score improved from -161.57 to -145.07, with ten contact cycles per leg and contact duties of 0.71 and 0.65. It passed eight of eight perturbed reruns and retained bilateral stepping down to zero hand drive. The motion remained in place under support. These two studies distinguish reduced

neural assistance from reduced mechanical assistance; neither demonstrates harness-free walking.

Appendix E: Supporting SADb Records

This appendix records the curation snapshot and pathway examples supporting Chapter 5. Counts refer to the stated audited or retained set; they are not estimates of the size of the neuroscience literature.

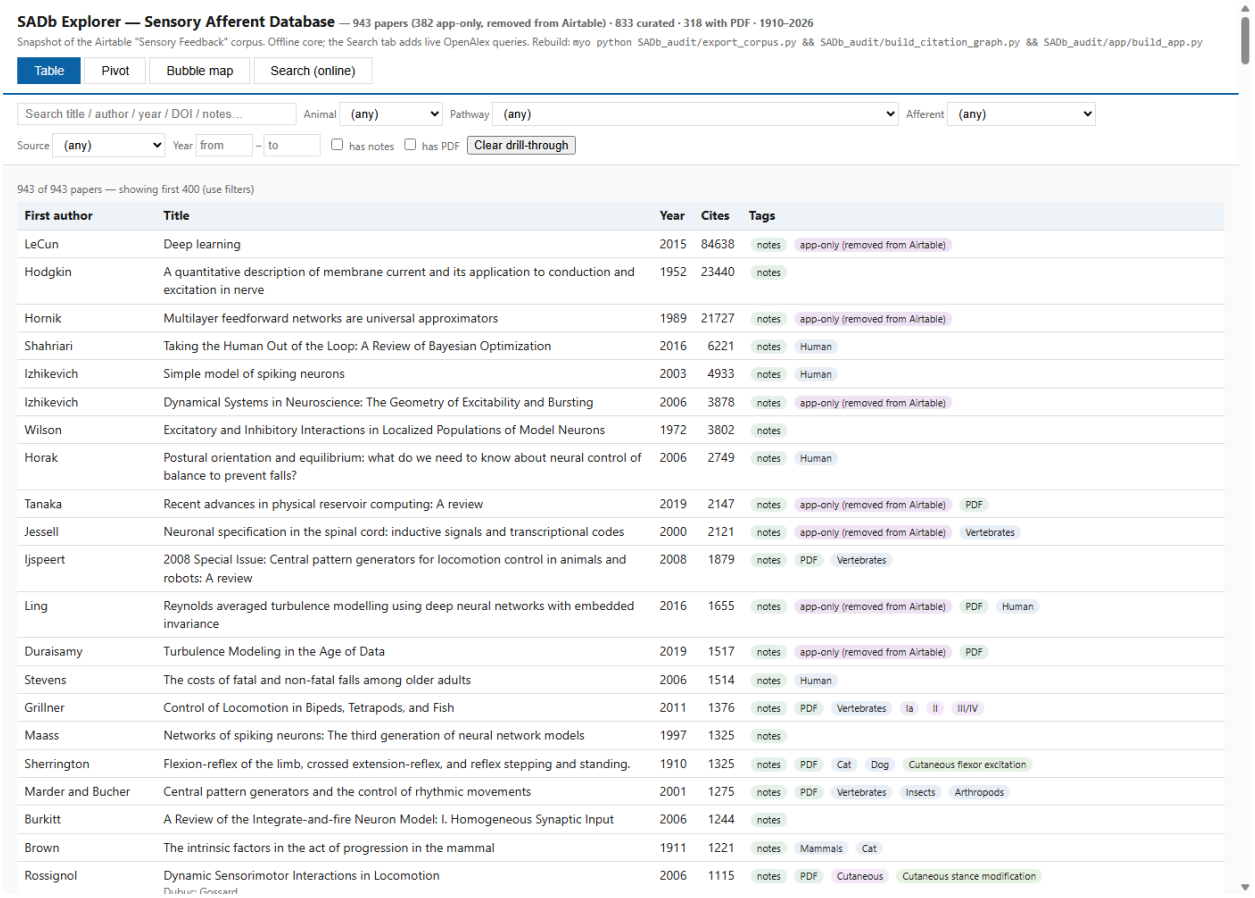
E.1 Record Sets and Coverage

The application contains 1,393 records from the five-table base: 943 Papers, 247 Models, 157 Review Papers, 42 Feedback pathways, and 4 Recorders. Of these, 436 are archived: 382 papers and 54 non-paper records. The retained working base contains 957 records: 561 Papers, 244 Models, 109 Review Papers, 39 Feedback pathways, and 4 Recorders. The application table includes all record types, while its pivot and citation views use the paper set.

Among the 561 retained papers, 523 carry curation notes, 426 preparation tags, 240 notes assessing animal-study potential, 118 links to named feedback pathways, 86 an afferent classification, and 258 robot/simulation translation notes. There are 550 DOIs; 38 records remain without grounding text. These fields overlap and should not be summed as disjoint categories.

The citation graph was constructed over the 943-paper audit: 730 papers were connected by 11,130 directed in-corpus citation pairs, and the Louvain partition produced 18 named groups. Citation proximity supplies a navigation aid, not evidence that two papers establish the same mechanism. The citation-neighborhood view is shown in Chapter 5

(Figure 5.1); Figure E1 provides the complementary paper-table view.



First author	Title	Year	Cites	Tags
LeCun	Deep learning	2015	84638	notes, app-only (removed from Airtable)
Hodgkin	A quantitative description of membrane current and its application to conduction and excitation in nerve	1952	23440	notes
Hornik	Multilayer feedforward networks are universal approximators	1989	21727	notes, app-only (removed from Airtable)
Shahriari	Taking the Human Out of the Loop: A Review of Bayesian Optimization	2016	6221	notes, Human
Izhikevich	Simple model of spiking neurons	2003	4933	notes, Human
Izhikevich	Dynamical Systems in Neuroscience: The Geometry of Excitability and Bursting	2006	3878	notes, app-only (removed from Airtable)
Wilson	Excitatory and Inhibitory Interactions in Localized Populations of Model Neurons	1972	3802	notes
Horak	Postural orientation and equilibrium: what do we need to know about neural control of balance to prevent falls?	2006	2749	notes, Human
Tanaka	Recent advances in physical reservoir computing: A review	2019	2147	notes, app-only (removed from Airtable), PDF
Jessell	Neuronal specification in the spinal cord: inductive signals and transcriptional codes	2000	2121	notes, app-only (removed from Airtable), Vertebrates
Ijspeert	2008 Special Issue: Central pattern generators for locomotion control in animals and robots: A review	2008	1879	notes, PDF, Vertebrates
Ling	Reynolds averaged turbulence modelling using deep neural networks with embedded invariance	2016	1655	notes, app-only (removed from Airtable), PDF, Human
Duraisamy	Turbulence Modeling in the Age of Data	2019	1517	notes, app-only (removed from Airtable), PDF
Stevens	The costs of fatal and non-fatal falls among older adults	2006	1514	notes, Human
Grillner	Control of Locomotion in Bipeds, Tetrapods, and Fish	2011	1376	notes, PDF, Vertebrates, Ia, II, III/IV
Maass	Networks of spiking neurons: The third generation of neural network models	1997	1325	notes
Sherrington	Flexion-reflex of the limb, crossed extension-reflex, and reflex stepping and standing.	1910	1325	notes, PDF, Cat, Dog, Cutaneous flexor excitation
Marder and Bucher	Central pattern generators and the control of rhythmic movements	2001	1275	notes, PDF, Vertebrates, Insects, Arthropods
Burkitt	A Review of the Integrate-and-fire Neuron Model: I. Homogeneous Synaptic Input	2006	1244	notes
Brown	The intrinsic factors in the act of progression in the mammal	1911	1221	notes, Mammals, Cat
Rossignol	Dynamic Sensorimotor Interactions in Locomotion	2006	1115	notes, PDF, Cutaneous, Cutaneous stance modification

Figure E1: SADb Explorer paper-table view. Searches and filters expose bibliographic information, curation notes, preparation tags, and afferent or pathway links. The embedded corpus is available offline.

E.2 Afferent Classes and Linked Pathways

The tables reproduce the curated signal summaries and representative linked pathways.

Their record names are database labels; pathway action depends on the experimental

preparation and locomotor state described in the cited sources. Translating these records into a controller requires an explicit encoding and verification step.

Table E1: Afferent classes represented in the Sensory Afferent Database, the signal each class carries, and the pathway roles the corpus records for it. Anchor references are corpus papers linked to the corresponding records; the signal summary for the Ia, II, and Ib rows follows the ensemble afferent recordings of **?**.

Class	**Receptor**	**Signal carried**	**Recorded pathway roles (anchors)**
Ia	muscle spindle primary	muscle length and velocity	homonymous stretch excitation [**?**]; reciprocal inhibition of the antagonist pool [**?**]
II	muscle spindle secondary	static muscle length	midlumbar disynaptic relay onto motor nuclei [**?**]; stance-phase support of extensor activity [**?**]
Ib	Golgi tendon organ	whole-muscle force	autogenic inhibition at rest, stance-phase reversal to excitation, and the load gate on swing onset [**??**]
III/IV	free nerve endings	metabolic and nociceptive signals, converging with cutaneous and group II–III afferents as flexor-reflex afferents	phase-dependent resetting of the locomotor rhythm [**?**]
Cutaneous	skin and foot-pad receptors	contact and perturbation events	flexor excitation and phase-modulated stance correction [**??**]

Table E2: Principal feedback-pathway records of the database's Feedback table, named by the fixed grammar of afferent, action, and phase. Linked papers are counted over the full audited corpus of 943 papers; for every pathway listed, the count is unchanged under the retained corpus of 561. Demonstrating references are the corpus papers linked to each record, as curated in the database.

Pathway record	Class	Recorded action in the step cycle	Linked papers	Demonstrating references
Ia monosynaptic excitation	Ia	stretch excitation of the homonymous motor pool, with the pathway gain modulated over the step cycle	7	[??]
Ia reciprocal inhibition	Ia	inhibition of the antagonist motor pool through the Ia inhibitory interneurons, rhythmically active in fictive locomotion	6	[??]
Ib disynaptic inhibition	Ib	autogenic non-reciprocal group I inhibition of the loaded muscle's own pool, reduced during stance	5	[??]
Ib disynaptic excitation	Ib	the locomotor reversal: group I input excites synergist extensor pools during the extensor phase, replacing the resting inhibition	6	[??]
Ib stance to swing	Ib	the load gate: Ib input from ankle extensors prolongs stance, so swing is not initiated while the leg is loaded	17	[??]
Ia or II stance to swing	Ia, II	flexor-length gate on stance termination: lengthening of flexor muscles in late stance ends extension and resets the cycle	10	[??]
type II excitatory	II	stance-phase support of ankle-extensor activity through the midlumbar group II relay	5	[??]
Cutaneous flexor excitation	Cutaneous	flexion withdrawal and the stumbling corrective reaction, phase-modulated over the cycle	16	[??]