Tools to Design a Biomimetic Humanoid Robot Actuated by Pneumatic Artificial Muscles and a Synthetic Neural Network

by

Ben Bolen

A dissertation submitted in partial fulfillment of the
requirements for the degree of

Doctor of Philosophy
in
Mechanical Engineering

Dissertation Committee:
Alexander Hunt, Chair
Bin Jiang
Radhika Reddy
David Turcic

Portland State University
2026

©Copyright by Ben Bolen
October, 2026
All Rights Reserved

Abstract

Biomimetic robots offer a platform for studying neuromechanical control through experiments that would be impractical or unethical with living subjects. This dissertation advances a long-running effort at Portland State University to build a humanoid robot, from the trunk down, actuated by pneumatic artificial muscles and controlled by a synthetic nervous system. The isometric force of Festo braided pneumatic actuators (BPAs) was characterized over a wide range of resting lengths, contractions, and pressures using a custom test jig. The analysis produced novel equations for the maximum isometric force as a function of resting length and a three-coefficient nondimensional force surface that predicts isometric force from relative strain and relative pressure. The actuator model was carried to the joint level through isometric knee torque experiments on two test stands: a simplified pinned knee and a biomimetic knee with a migrating instantaneous center of rotation. This approach isolated the sources of prediction error, revealing constant length offsets, series compliance in brackets and artificial tendons, and a loss of usable actuator length when BPAs wrap around the joint. The redesigned flexor and cam-extensor routes met or exceeded their selected human-muscle reference torques at all measured angles; these discrete isometric measurements do not establish dynamic performance or a continuous measured torque envelope. Tools were created for curating relevant neuromechanical research, as well as simulating neuromechanical models based on human

and robot biomechanical models.

Dedication

*

Darian

azīzam, jānam,

nūr-e cheshmam — the light of my eyes

qorbānat beram

Bābā kheyli dūset dāre

↘ ↘ ↘

"Los corazones obreros que laten

por nuestra causa, felices serán.

Si entusiasmados y unidos combaten,

de la victoria, la palma obtendrán."

Acknowledgments

This work was funded by the Department of Mechanical and Materials Engineering at Portland State University, the National Science Foundation (NSF) grant for NeuroNex: Communication, Coordination, and Control in Nervous Systems (C3NS) #2015317, NSF grant #1943483, and the NIH BUILD EXITO program (grant numbers 5RL5GM118963-06, 5TL4GM118965-06, 5RL5GM118963-07, and 5TL4GM118965-07).

I would like to thank my advisor, Dr. Alexander Hunt, for his guidance, patience, and the freedom to pursue this project, and my committee members — Dr. Bin Jiang, Dr. Radhika Reddy, and Dr. David Turcic — for their time and thoughtful feedback.

I am grateful to my collaborators and labmates: Mohammad Elzein and Lawrence Pang for their work on the actuator characterization experiments; Lindie Burgess and Connor Morrow for their work on the knee torque experiments; Cody Scharzenberger for building the Muscle Mutt platform and for his assistance proofreading; Alex Steele, who designed the initial biomimetic four-bar knee linkage used in the torque tests and generously shared solid models and prototype experience; and Jasmine Bradley, whose work reworking the figures improved their aesthetic quality and accessibility for readers with color vision deficiencies.

Finally, I thank my family for their support throughout this long process.

Contents

Abstract **ii**

Dedication **iv**

Acknowledgments **v**

1 Introduction **1**

1.1 Motivation 1

1.2 The Biomimetic Humanoid Robot Project 4

1.3 Research Objectives 5

1.4 Dissertation Organization 6

2 Background **9**

2.1 Biomimetic and Musculoskeletal Robots 9

2.2 Robot Actuation Technologies 10

2.2.1 Overview 10

2.2.2 Braided Pneumatic Actuators 14

2.2.3 BPA Force–Length–Pressure Models 16

2.3 Human Lower-Limb Biomechanics 18

2.4 Biomimetic Knee and Bipedal Robot Model 19
2.5 Neural Control of Locomotion: Central Pattern Generators 21
2.6 Sensory Feedback in Locomotion 24
2.7 Computational Neuroscience 28
2.8 Synthetic Nervous Systems 28
2.9 Neuromechanical Simulation 31
2.9.1 Simulation Tools 33
2.10 Testing Neuromechanical Hypotheses on Robots: The Ethical Case 34

3 Isometric Force Characterization of Braided Pneumatic Actuators 39
3.1 Methods 40
3.1.1 Actuator Force Model 40
3.1.2 BPA Isometric Force Characterization Experiment 41
3.2 Results 43
3.2.1 BPA Force Characterization 43
3.3 Discussion 49
3.3.1 Effects of Resting Length 49
3.3.2 Nondimensionalized Force Prediction 54
3.3.3 Comparison with Other Published BPA Models 55
3.3.4 Implications for Real-Time Control and Simulation 59
3.3.5 Biomimetic Artificial Muscle Analogy 60

4 Isometric Knee Torque 64
4.1 Methods 64

4.1.1 Robot Leg Model and Knee Test Stands 65
4.1.2 Knee Torque Data Collection and Equipment 69
4.1.3 Hybrid Torque Calculation Method 71
4.1.4 Improved Torque Model 72
4.1.5 Actuator Placement and Route Redesign 86
4.2 Results 94
4.2.1 Flexor on Pin Joint 94
4.2.2 Extensor on Pin Joint 99
4.2.3 Extensor on Biomimetic Knee 102
4.2.4 Flexor on Biomimetic Knee 104
4.3 Discussion 105
4.3.1 Reasons for the Correction Terms 105
4.3.2 Optimizer Updates 109
4.3.3 Mono- and Bi-Pulley Tendon Routing 116
4.3.4 Whole-Leg Transmission Capability 119
4.3.5 Physical Tests of the Redesigned Routes 120

5 The Sensory Afferent Database 124
5.1 The Resource 124
5.2 Curation Methods 125
5.3 Results: Application and Corpus Coverage 125
5.3.1 AI-Assisted Literature Triage 127
5.4 Discussion 129

5.4.1 Robot and Simulation Use . . . 129
5.4.2 Further Improvements . . . 131

6 Synthetic Nervous System Simulation 133

6.1 The AnimatLab Bipedal Walker . . . 133
6.2 Attempting W2L Integration in MATLAB . . . 134
6.3 Knowledge Gaps and the New Tools . . . 136
6.4 Methods: Software Tools . . . 139
6.4.1 Model Conversion and Actuator Coupling . . . 139
6.4.2 Block Programming Neural Pathways . . . 141
6.4.3 Connecting Simulink Controllers to MuJoCo . . . 142
6.4.4 A Curriculum for Tuning and Controlled Experiments . . . 142
6.5 Results: Verification of the Implemented Interfaces . . . 146
6.5.1 Coupled Block-Model Tests . . . 146
6.5.2 Matched Neural Implementations . . . 147
6.5.3 Layer-Specific Stimulation . . . 151
6.5.4 Conversion Checks and Muscle-Force Agreement . . . 151
6.5.5 Closed-Loop Model Coupling . . . 152
6.5.6 Testing Alternative Controllers on Gait2392 . . . 153
6.5.7 Feedback Contribution at Fixed Controller Settings . . . 155
6.5.8 Repeatability and Mechanical Support . . . 155
6.6 Discussion: What the Tools Establish . . . 157

7 Discussion 160

7.1 The Work as a Whole . . . 160
7.2 Future Work . . . 163
7.2.1 Expanding the BPA Force Characterization . . . 165
7.2.2 Expanding Attachment and Routing Design . . . 166
7.2.3 Expanding the Sensory Afferent Database . . . 168
7.2.4 Closing the Loop in Simulation . . . 168
7.2.5 Adding Sensory and Higher-Level Control . . . 170
7.2.6 Real-Time Hardware Realization . . . 173
7.2.7 Toward Dynamic Whole-Body Maneuvers . . . 174

8 Conclusion **175**

Appendices **177**
A Appendix A: Muscle Force and Joint Torque Calculations . . . 178
A.1 Human muscle force calculations . . . 178
A.1.1 OpenSim Thelen2003 and Millard2012 Formulations . . . 178
A.1.2 OpenSim Moment Arms and Human Torque Targets . . . 182
A.2 Converted MuJoCo Muscle Formulation . . . 183
A.2.1 MuJoCo Muscle Actuator . . . 183
A.3 BPA Force Model Implementation . . . 183
A.3.1 Strain Measures . . . 183
A.3.2 Normalized Force Surface . . . 184
A.3.3 Maximum Force . . . 185
A.4 Muscle Path Length and Moment Arm . . . 185

A.5 Compliance Model Implementation 186
A.6 Wrapping-Loss Implementation 188
B Appendix B: Neuron Model Equations 189
B.1 Spiking Neuron Model 189
B.2 Non-Spiking (Graded) Chemical Synapses 190
B.3 Relation to the SNS-Toolbox Implementation 191
B.4 Spiking-Mirror Implementation, Calibration, and Verification 192
B.4.1 Adapting Integrate-and-Fire Cells and Hybrid Synapses 192
B.4.2 Calibration of Record 194
B.4.3 Verification Gates of the MuJoCo Mirror 196
B.4.4 AnimatLab Conversion Variants 196
B.5 Connectome of the Verified W2L Model 199
C Appendix C: Model Geometry, Bracket Frames, and Compliance Parameters . 204
C.1 Static BPA Pathway Data 204
C.2 Joint Transformation Matrices 210
C.3 Bracket Reference Points and Frames 211
C.4 Stiffness Arrays and Compliance 212
C.5 Adopted Correction Terms and Complete Pinned-Knee Results 215
C.6 Redesigned Extensor Route Geometry 217
C.7 Contact and Release Angles of the Wrapped Route 218
C.8 Supporting Route Searches and Transmission Sweeps 224
C.9 Supporting Torque-Prediction Fit Statistics 231
C.9.1 Goodness-of-Fit Metrics 231

C.10 Configuration-Survey Coordinates and Actuator Sizes 231

C.10.1 Torque and Active-Force Coverage 232

D Appendix D: Supplementary Neuromechanical Verification 235

D.1 Model Construction and Recording 235

D.2 Controller Ports and Reference Checks 235

D.2.1 Transient Knee Sensory-Current Response 239

D.2.2 A Published Neural Mechanism as a Verification Target 242

D.2.3 RG, PF, and MN Stimulus Tests 242

D.2.4 Archived AnimatLab Models 245

D.3 Tuning and Supported Evaluations 246

D.3.1 Reference-Library Comparison 246

D.3.2 Direct Kinematics Overlay 247

D.3.3 Pathway-Removal Tests 248

D.4 Spiking-Neuron Comparisons 250

D.4.1 MuJoCo Network and Ground Evaluation 250

D.4.2 Simulink Demonstrations 250

D.4.3 AnimatLab Conversions 255

D.5 Supported BPA-Robot Evaluations 255

D.5.1 Independent Replays and Remaining Mechanical Assistance 256

D.5.2 Pathway Deletions and Stimulus Tests 257

D.5.3 Di Russo Controller Replication Attempt 257

D.5.4 CAD-Derived Robot Integration 259

D.6 OpenSim-to-MuJoCo Conversion Repairs 262

D.6.1 Cross-Platform Force Comparison 264
E Appendix E: Supporting SADb Records 267
E.1 Record Sets and Coverage 267
E.2 Afferent Classes and Linked Pathways 269

List of Tables

2.1 Coefficients of the 10 mm Festo BPA pressure model of Equation (2.1), corrected from **?**. 17

3.1 Maximum force equation coefficient values with confidence intervals (CI) and goodness-of-fit measures. Equation (3.5) is compared to data taken at 620 kPa. Equations (3.7) and (3.8) are compared against maximum force from using the Festo tool. It was not necessary to fit an adjusted R^2 in these cases. 48
3.2 Normalized isometric BPA force equation (Equation (3.4)) coefficient values and goodness-of-fit measures. 48
3.3 BPA force-model comparison. Diameter and resting length are in mm; RMSE and maximum absolute error are in N. FVU is dimensionless. 63

4.1 Terms of the knee-flexor mesh-placement cost. The cost is minimized; a surplus reduces it, while a direction mismatch increases it. 88

4.2 Torque penalties in the constrained knee-route search. FLEX and EXT identify configuration-specific terms. 91

4.3 Feasibility and geometric penalties in the constrained knee-route search. . . 92

4.4 Additional cam-profile penalties and rejection values. 92

4.5 Lengths l_0 (resting BPA) and l_{ten} (artificial tendon) are both in cm. Pre-optimization pinned-knee flexor torque-prediction goodness of fit per test, comparing the three pre-optimization torque predictions of Figure 4.7: the Hunt-equation rigid-body prediction, the unoptimized prediction of the improved BPA characterization, and the back-calculated (hybrid) torque. RMSE and maximum absolute residual (Max.) are in N m, and FVU denotes the fraction of variance unexplained. 96

4.6 Lengths l_0 (resting BPA) and l_{ten} (artificial tendon) are both in cm. Pinned-knee flexor goodness of fit before and after correction, using the adopted fold. RMSE and maximum absolute residual (Max.) are in N m; FVU is dimensionless. 98

4.7 Pinned-knee extensor goodness of fit before and after correction, with stiffness shared with the flexor. RMSE and maximum absolute residual (Max.) are in N m; FVU is dimensionless. 102

4.8 Biomimetic-knee goodness of fit before and after design-parameter optimization. RMSE and maximum absolute residual (Max.) are in N m; FVU is dimensionless. 105

4.9 Adopted transmission corrections. The stiffness pair is shared. The wrapping factor was identified on the extensor and is also used in the flexor redesign; it was not fitted to the pinned-flexor tests. 109
4.10 Mean torque and active-force coverage across 27 actuator sites at 620 kPa. . 120

6.1 Models and tools developed for neuromechanical experiments 138
6.2 Five-stage controller-tuning curriculum 145

List of Figures

1.1 Dissertation roadmap 7

2.1 Kengoro human-mimetic musculoskeletal humanoid 11
2.2 Representative robotic actuator technologies 13
2.3 Festo BPAs with and without inflation 15
2.4 Steele biomimetic knee mechanism 20
2.5 Robot skeleton and robot/human muscle paths 21
2.6 Equivalent geometric constructions of joint torque 22
2.7 Biceps femoris long-head paths and torque surfaces in OpenSim (A, C, E) and the robot model (B, D, F). Panels C–D show knee torque; E–F show hip torque over the hip–knee range [**?**, Fig. 8]. 23
2.8 Deng two-layer CPG and sensory–motor network 26

2.9 Antagonist-muscle spinal network with Ia, II, and Ib reflex pathways, Renshaw cells, and locomotor-pattern input. Reproduced from **?**, Fig. 5 under CC BY 4.0. 26

2.10 Shinohara bilateral neuromusculoskeletal model 27

2.11 Rybak two-level locomotor CPG concept 29

2.12 Shevtsova extended CPG model with two-joint-muscle circuits 29

2.13 Decerebrate cat treadmill experiment, restored footage frame 36

3.1 A picture of a BPA in the isometric force test stand with components labeled. 44

3.2 Results for finding the relationship between F_{620} and l_0. (**a**) $F_{620_{10}}$ as a function of l_0, at P_{620}. Dashed line is the $F_{620_{10}}$ data from Festo. Solid line is the fit from Equation (3.7). (**b**) $F_{620_{20}}$ as a function of l_0 at P_{620}. Dashed line is $F_{620_{20}}$ data from Festo. Solid line is the fit from Equation (3.8). 46

3.3 Results for finding the relationship between F_{620}, l_0, P_{620}, and $\epsilon_{620_{10}}$. (**a**) Isoclines of the surface fit for $F_{620_{10}}(l_0, P_{620})$. (**b**) $\epsilon_{620_{10}}$ versus l_0 at P_{620}. Although there is a general trend of longer resting lengths producing more contraction, no conclusive relationship between ϵ_{620} and l_0 could be deduced from this experiment. The error bars show the effect of being ± 1 mm in measuring l_0 and l_{620}. 50

3.4 Surface fit for $F^*(\epsilon^*, P^*)$. Solid lines are from our model. Dashed lines from Festo supplied data. In total, 537 data points were collected and used to fit the 3D surfaces (321 data points for ∅10 mm BPAs and 216 data points for ∅20 mm BPAs). In total, 80% of these data were used for training, and 20% were used for validation. For figure clarity, not all data is included and plotted circles represent collected data at ± 10 kPa of the stated pressure. (**a**) Fit data for ∅10 mm Festo BPAs. (**b**) Fit data for ∅20 mm Festo BPAs. 51

3.5 Comparison between measured and predicted BPA force at varying amounts of contraction and pressure. A $y = x$ perfect fit line is added for reference. The BPAs compared are: (**a**) ∅10 mm, $l_0 = 257$ mm, (**b**) ∅10 mm, $l_0 = 233$ mm, (**c**) ∅20 mm, $l_0 = 300$ mm, and (**d**) ∅20 mm, $l_0 = 450$ mm. 57

4.1 Pinned-knee test apparatus: CAD assembly (A), labeled mechanism at -30° (B), and extensor isometric test at -120° (C). 66

4.2 Biomimetic-knee instantaneous center of rotation relative to the tibial frame. White circles are measured points; knee outlines show 0° (solid), -60°, and -120° (dash-dot). 68

4.3 Bracket frame orientations for: **(A)** flexor configuration on the pinned knee, **(B)** extensor configuration on the pinned knee, **(C)** flexor and extensor configurations on the biomimetic knee with ∅10 mm BPAs, and **(D)** flexor configuration on the biomimetic knee with the ∅20 mm BPA. 74

4.4 Bracket-frame construction from attachment point $\mathbf{p}_A$ or $\mathbf{p}_B$: yaw about fixed $\hat{\mathbf{z}}_s$ aligns the sagittal projection (A); pitch about current $\hat{\mathbf{y}}_{br}$ aligns the full point vector (B). 76

4.5 Projection of bracket displacement onto the actuator force path. Each bracket has a local basis $\hat{\mathbf{e}}_1^{(j)}$, $\hat{\mathbf{e}}_2^{(j)}$, and $\hat{\mathbf{e}}_3^{(j)}$. The displacement $\boldsymbol{\eta}_{br,j}$ has an axial projection $\delta_{br,j}$ along $\hat{\mathbf{u}}_j$; the two bracket projections and tendon elongation contribute to the change in actuator length. Displacements are exaggerated. 81

4.6 BPA crossing a flexed knee 84

4.7 Pinned-knee flexor torque before design-parameter optimization: **(A)** $l_0 =$ 48.5 cm, **(B)** $l_0 = 41.0$ cm, **(C)** $l_0 = 45.7$ cm, and **(D)** $l_0 = 40.6$ cm with a 35 mm tendon. Panels **(A)** and **(B)** are training tests; **(C)** and **(D)** are validation tests. The legend distinguishes measurements, the Hunt model, the improved BPA model, and the hybrid calculation. 95

4.8 Optimized pinned-knee flexor torque: **(A)** $l_0 = 48.5$ cm, **(B)** $l_0 = 41.0$ cm, **(C)** $l_0 = 45.7$ cm, and **(D)** $l_0 = 40.6$ cm with a 35 mm tendon. Panels **(A)** and **(B)** are training tests; **(C)** and **(D)** are validation tests. Circles show measurements, dashed curves the improved BPA model, and solid magenta curves the optimized model. 97

4.9 Optimized pinned-knee extensor torque: **(A)** $l_0 = 43.6$ cm, **(B)** $l_0 = 43.2$ cm with a 22 mm tendon, **(C)** $l_0 = 45.7$ cm, **(D)** $l_0 = 40.0$ cm, **(E)** $l_0 = 40.0$ cm with a 40 mm tendon, and **(F)** $l_0 = 46.5$ cm. Panels **(A)**–**(C)** are training tests; **(D)**–**(F)** are validation tests. The legend applies to all panels. 100

4.10 Biomimetic-knee extensor test: **(A)** apparatus with the BPA passing between the femoral condyles; **(B)** measured torque, the improved BPA baseline, the optimized model at the adopted terms, and the approximate hybrid calculation for the 10 mm-diameter, 51.8 cm BPA. 103
4.11 Biomimetic-knee flexor torque: single 10 mm BPA (A), single 20 mm BPA (B), and redesigned two-BPA route (C). Measurements, predictions, and human targets follow the shared legend; B and C use the same axis ranges. . 106
4.12 Cam-extensor geometry at -120° (A) and 10° (H), with the assembled two-BPA hardware below. Panel letters match the complete sequence in Figures C5 and C6. 115
4.13 Reverse-pulley force and travel relations 118
4.14 Cam-extensor measurements, pressure-specific predictions, the optimized 620 kPa curve, and the vastus medialis reference. Marker area scales with pressure. Series follow the plot legend. 122
4.15 Physical tests of the redesigned flexor route 123

5.1 SADb Explorer citation-map functions 128
5.2 The SADb Explorer pivot view 130

6.1 Original AnimatLab walker (A) and shared-rhythm-generator controller (B), developed by Morrow and Bolen. The four-neuron RG drives hip and knee/ankle pattern-formation groups on both sides. 135
6.2 Model construction and neural-controller interfaces 143
6.3 Simulink SNS library blocks and neuron implementation 144

6.4 Transient knee response to a 2 nA, 30 ms sensory-current pulse after settling (A), antagonist-BPA leg schematic in an illustrative pose (B), and the six implemented feedback connections (C). Flexion is negative. The graded-response configuration uses broadened MN and Ia synaptic saturation ramps; its response is compared with a matched no-stimulus control. . 148

6.5 Loaded-arm demonstration: empty and loaded cup positions (A1, A2), feedback-on/off kinematics (B), and all six implemented feedback connections (C). 149

6.6 W2L rhythm-core traces across engines 150

6.7 Gait2392 model motion compared with inverse kinematics 154

6.8 Sensory-feedback comparison in the recovered W2L model 156

6.9 CAD-derived robot under supported hybrid control 158

7.1 The first-pass design workflow the enablers make possible 164

7.2 Proposed layered BPA finite-element model: contact and boundary conditions (A), and illustrative material cutaway (B). Geometry and layer spacing are schematic; the middle rubber layer is separated only in A. 167

7.3 Control stack with the planned supraspinal additions 171

7.4 Simulation-to-hardware workflow 172

1 Introduction

1.1 Motivation

Biomimetic robots provide a platform for investigating how mechanical structure and control interact to generate movement, offering complementary insights to simulation-based studies [**???**]. Wheeled robots remain the platform of choice in structured environments, where their speed, efficiency, and payload capacity are difficult to match; in the unstructured terrain of the natural and built world, legged machines hold the advantage, and biomimetic legs that reproduce human force capabilities have the potential to excel where wheels cannot [**??**]. Musculoskeletal humanoids and limb systems driven by tendon-like actuators have been shown to support biologically meaningful experiments on motion generation, internal force transmission, and neural control hypotheses [**?????**]. The dual perspective of using robots to learn about biology, and using biology to design more adaptive robots, has been emphasized in both biomimetic robotics and biomechanics [**??**]. With such robotic systems, experiments can be performed that would not be practical or ethical with human test subjects, such as removing a specific muscle from the system or blocking a sensory feedback pathway. Experiments of this kind allow direct testing of how the components of the neuromechanical system interact, and they offer new opportunities for testing neural control theories that are faster and cheaper than human observation studies [**???**].

Artificial muscles are a key enabling technology for this class of robot. Among the

available actuation technologies, McKibben-style braided pneumatic actuators (BPAs) combine low mass, high force-to-weight ratio, and intrinsic compliance with a force–length curve that qualitatively resembles that of skeletal muscle [**??**]. BPAs have been used in quadrupedal robots [**???**], musculoskeletal humanoids [**??**], and dexterous manipulation systems [**?**]. However, BPAs also differ from biological muscle in important ways: their maximum tensile force occurs at resting length rather than at an optimal fiber length, their maximum contraction ratio is substantially lower than that of human muscle, and their force output is a highly nonlinear function of internal pressure, contraction, and loading state [**???**]. My experience with 10 mm actuators made this gap concrete: manufacturer specifications overstate the force and contraction available at the short resting lengths a leg design requires, and designs based on those specifications delivered less range of motion and joint torque than intended. Nor could the published empirical models fill in the gaps, because each is fitted to measurements at a single resting length [**???**]; none accurately predicts how force changes when a designer chooses a different actuator length. Because muscle resting length is one of the primary degrees of freedom available to the robot designer, the absence of a force model valid at arbitrary resting length blocked accurate actuator selection. Designing a biomimetic robot around BPAs therefore requires quantitative, predictive models of the actuator and of the joint torques it can produce.

Actuator force alone, however, does not determine the torque a joint produces. The robot knee transmits BPA force to the skeleton through mounting brackets, artificial tendons, and paths that wrap over the joint, and each of these elements departs from the rigid-body idealization: brackets and fixtures deflect under load, tendons stretch, and a wrapping actuator loses usable length as the knee flexes. At the start of this work it

was unknown how much these transmission effects reduce the torque delivered at the knee relative to the rigid-body prediction, and no method existed for correcting that prediction in a form that would transfer between joint configurations. The earlier robot routes were selected from predicted muscle lengths, moment arms, and torque profiles [**?**]. The placement searches developed during this dissertation extended that design approach [**?**], but their rigid-body predictions still required experimental validation at the knee.

Researchers implementing neural control models face two related problems. Sensory-afferent findings are distributed across species, preparations, and experimental protocols, so selecting a feedback pathway requires locating its measured properties and the conditions under which they were observed [**???**]. I addressed that evidence-access problem with the Sensory Afferent Database. A second problem arises when a published neural circuit is coupled to a different mechanical model: its behavior depends on the actuator mechanics, sensory signals, support conditions, and numerical implementation as well as on the circuit itself. Published architectures and software supply essential components [**????**], but the behavior of their particular combination must be established. I developed coupled human and artificial-muscle models, graphical neural blocks, and controlled evaluation procedures to distinguish these contributions. Representative tests establish circuit agreement across implementations, corrective reflex behavior, and the dependence of a supported stepping configuration on selected feedback pathways.

The dissertation therefore addresses four connected needs: an actuator-force model that accounts for resting length, corrections for transmission losses at the knee, a curated sensory-afferent resource, and tools for developing and checking neuromechanical simulations. Together, they support the design and evaluation of the robot; achieving

unsupported walking remains a separate control and integration problem.

1.2 The Biomimetic Humanoid Robot Project

This dissertation is part of a long-running effort at Portland State University to build a biomimetic humanoid robot from the trunk down, actuated by pneumatic artificial muscles and controlled by a synthetic nervous system (SNS). The project is deliberately iterative. Knowledge from biomechanics and neuroscience guides the design of robots intended to perform complex tasks in unstructured environments with improved efficiency, balance, adaptability, and resilience. Those robots then become physical models with which biological hypotheses can be tested, validated, and refined through controlled changes that would be impractical or unethical in living subjects. The resulting biological insight informs the next generation of robot design, and lessons from both directions can be translated into improved prosthetic and orthotic devices.

The present work builds directly on a series of studies by me and my collaborators. In my master's project, the kinematics and dynamics of a biomimetic robot leg were matched against the Gait2392 OpenSim model, artificial muscle attachment and wrapping locations were determined for the robot legs, and MATLAB software was developed for evaluating muscle lengths, moment arms, and isometric joint torques over the full range of motion [**??**]. That work showed that simply transplanting human muscle attachment locations onto a BPA-actuated robot does not reproduce human torque profiles; the placement optimization developed during this dissertation addresses this limitation and is introduced in Section 7.1. The rest of the platform advanced in parallel, and this dissertation builds on its hardware

and demonstrations: the biomimetic four-bar knee linkage with a migrating instantaneous center of rotation, designed and validated by Steele [**???**], is the knee of the robot leg whose torque transmission Chapter 4 measures, and the canine-inspired quadruped platform Muscle Mutt, formerly called DoggyDeux, demonstrated antagonistic BPA pairs under synthetic nervous system control on a physical robot [**?**], extending an earlier AARL demonstration of central pattern generator control of a pneumatically actuated robot [**?**]. The humanoid leg of this dissertation carries that line forward.

1.3 Research Objectives

This dissertation investigates actuator force and knee torque experimentally, consolidates sensory-afferent evidence in the Sensory Afferent Database [**??**], and develops neural-model and simulation interfaces with representative verification checks. The objectives are:

1. **Characterize the isometric force of Festo BPAs** with uninflated diameters of 10 mm and 20 mm over a wide range of resting lengths, contractions, and pressures; derive improved maximum-force and normalized force–strain–pressure models; and compare them against existing models (Chapter 3).

2. **Validate and improve the prediction of isometric knee torque** produced by BPAs on pinned-joint and biomimetic four-bar knees; identify and model losses caused by constant length offsets, series compliance, and actuator wrapping; and compare the resulting torque envelopes with human capabilities (Chapter 4).

3. **Consolidate the sensory afferent literature into a searchable resource**, the Sensory Afferent Database, by auditing and curating the published record of afferent properties across animal species and their role in motor control, and by evaluating AI-assisted screening of candidate papers for expanding the collection (Chapter 5).

4. **Develop and demonstrate neuromechanical modeling tools** by implementing the Li–Deng walker, Gait2392-based models with alternative neural controllers and artificial-muscle actuation, SNS-Simscape blocks for neural pathways, and staged tuning and experiment procedures; verify representative circuit responses, implementation agreement, and neural–mechanical signal exchange (Chapters 6 and 7).

1.4 Dissertation Organization

Chapter 2 introduces the robot, BPA actuation, musculoskeletal design, locomotor control, and the scientific rationale for robotic models. Chapter 3 presents the isometric force characterization and actuator models, adapted from my publication [**?**]. Chapter 4 tests the joint-torque predictions, identifies compliance and wrapping corrections, and uses them to redesign actuator routes.

Chapter 5 presents the Sensory Afferent Database, its curation, and its use limits. Chapter 6 begins with the AnimatLab attempt to combine the Deng CPG and Li walker, identifies the development requirements exposed by that attempt, and presents the tools and representative checks built to address them. Chapter 7 discusses the contributions together and develops the future work; Chapter 8 states the conclusions. Figure 1.1 summarizes these connections.

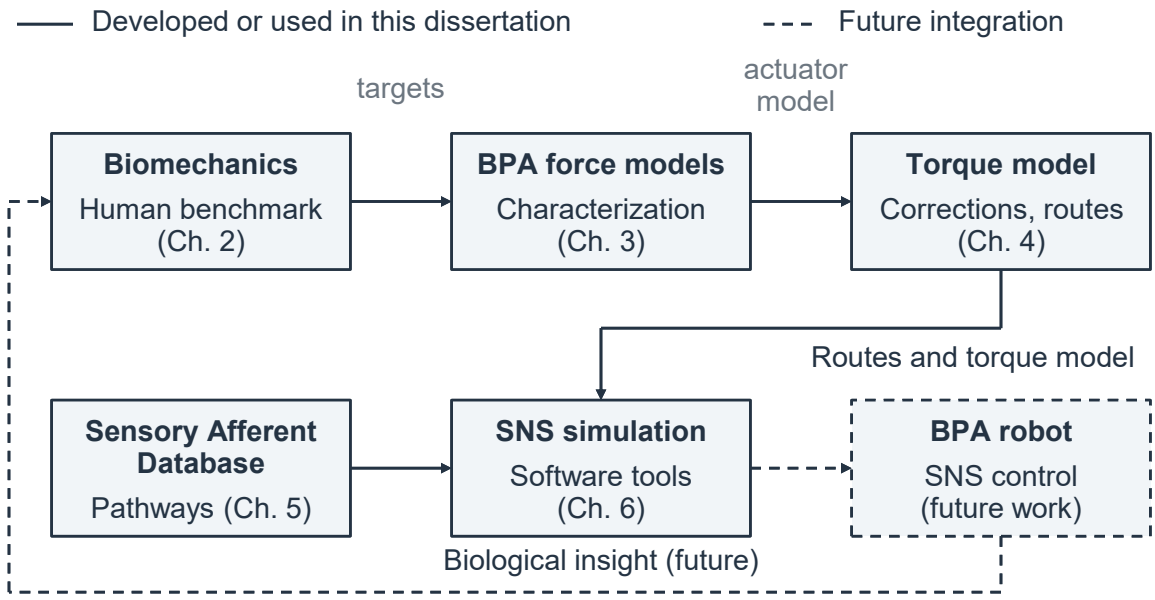


Figure 1.1: Dissertation roadmap. Biomechanical targets guide actuator characterization and route design; the sensory-afferent resource informs feedback selection; and simulation tools support controller evaluation. Dashed connections indicate future integration and the intended return of biological insight.

Appendices A–C provide the muscle calculations, neural equations, model geometry, and supporting mechanical results. Appendix D collects the larger simulation evaluations, and Appendix E provides the database coverage and pathway tables.

2 Background

This chapter reviews the established knowledge against which the dissertation's contributions are defined: biomimetic and musculoskeletal robotics, the actuation technologies available to them, the human lower-limb biomechanics they imitate, the neuroscience of locomotor control they implement, and the simulation and ethical context in which they are tested. Work from the author's own program that goes beyond this background, including the prior AnimatLab bipedal walker, the Sensory Afferent Database, and the neuromechanical simulation pipeline, is presented as dissertation work in Chapters 5 and 6 rather than repeated here.

2.1 Biomimetic and Musculoskeletal Robots

Biomimetic robots offer many potential applications, ranging from assisting in factories and environmental exploration to serving as physical models of human movement. By developing high-fidelity biomimetic robots, we can perform tests that help improve our understanding of both biomechanics and the underlying neuromechanical systems that control movement [**??????**]. A few such biomimetic robots exist, such as Kengoro and Kenshiro robots of **?** and **?**, which have proportionally human bodies, anatomically correct musculotendon arrangements, rolling-sliding joints, and joint ranges and output power comparable with human-measured values; their artificial muscles are implemented with brushless DC motors, wires as tendons, gears, and integrated tension measurement. Kengoro, the more recent of the two, stands 1670 mm tall, weighs 56.5 kg, and carries 174

degrees of freedom when the end effectors are included [**?**]. A dog-like robot developed by **?** and actuated with Festo BPAs was designed from a planar kinematic model scaled to an adult greyhound, with three sagittal degrees of freedom per leg driven by antagonistic BPA pairs; its front-leg pair alone lifted a 13.5 kg payload, almost double the complete robot's predicted mass. That platform was used to show how central pattern generators provide robust locomotion patterns that are adaptive to perturbations in ground speed [**?**].

Robots built with artificial muscles thus provide opportunities to investigate how muscles are controlled to produce desired motions; however, artificial muscles do not exactly replicate muscle properties, and existing robots do not capture all of the degrees of freedom and joint torques of the animal limbs they mimic. The biomimetic leg developed and tested in this dissertation builds on this approach: a human-proportioned skeleton, muscle-like actuators, and the goal of human-matched joint torques.

Figure 2.1 illustrates the human-scale skeletal structure and tendon routing of this approach.

2.2 Robot Actuation Technologies

2.2.1 Overview

A wide range of actuation technologies have been explored in robotics, including electric motors, hydraulic systems, smart-material actuators, and artificial muscles [**??**]. Electric motors are widely used due to their controllability and maturity, but they are limited by thermal constraints, gearing requirements, and reduced compliance compared to muscle-like actuators [**??**]. Hydraulic actuators offer exceptionally high power density and have

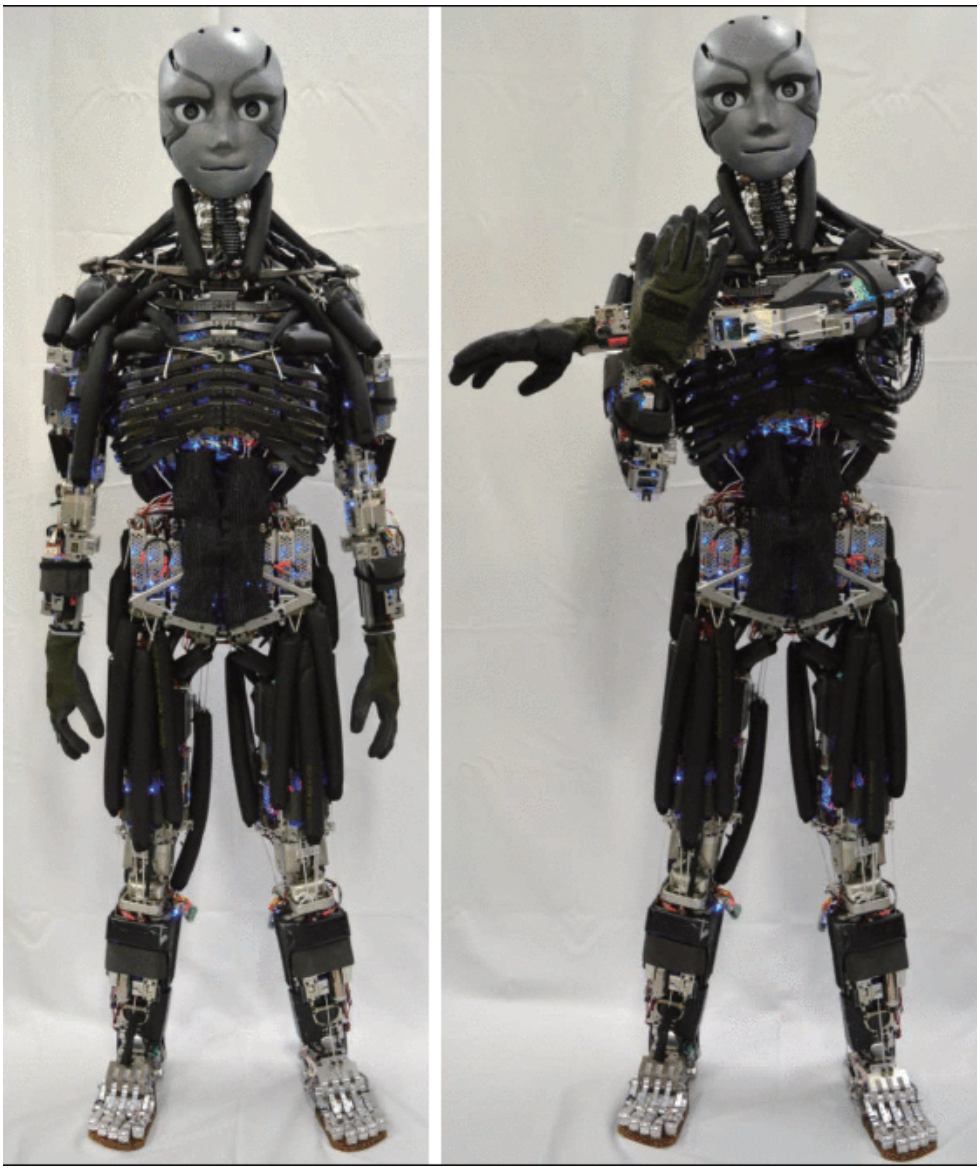


Figure 2.1: Kengoro human-mimetic musculoskeletal humanoid in neutral and upper-limb poses, illustrating its human-scale skeletal structure and densely routed tendon-driven musculature [**?**].

enabled highly dynamic legged systems, but they remain relatively heavy, less efficient than biological muscle, and require a dedicated fluid infrastructure [**?**].

Among smart-material actuators, dielectric elastomer actuators (DEAs), a class of electronic electroactive polymers, contract by a different mechanism: an applied voltage squeezes a soft polymer film between compliant electrodes through electrostatic stress, and the thinning and radial spreading of the film produces strains that can exceed biological muscle [**?**]. Comparative studies report high specific work and good bandwidth for DEAs [**?**], but kilovolt-range driving voltages, electromechanical instability, and the compliant-electrode requirement have kept them out of legged robots. Thermally activated artificial muscles reach the most extreme performance figures of any technology surveyed by **?**: twisted-fiber muscles support loads more than one hundred times those of a biological muscle of the same length and weight, with output stresses of 22 MPa. Their activation route, however, is shared with shape-memory alloys, whose heating is fast but whose cooling is limited by heat transfer away from the element, and that asymmetry sets the cycle time.

Closer to muscle, series-elastic and pneumatic actuators trade peak force and bandwidth for compliance, a trade that wearable and rehabilitation robotics exploit because force control and backdrivability matter most when a machine interacts with a human [**???**]. A recent review of pneumatic artificial muscle technology documents continued progress in actuator design and in assistive applications [**?**]. None of these technologies yet combines muscle-like force, strain, and compliance in one actuator, which is the combination a biomimetic humanoid leg demands and that braided pneumatic actuators approximate most closely.

Figure 2.2 compares the operating mechanisms of these actuator technologies.

Figure 2.2 shows representative operating mechanisms for robotic actuators. **(A)** An electric torque motor produces controlled shaft rotation. **(B)** A braided pneumatic actuator shortens as internal pressure expands its sleeve radially. **(C)** A hydraulic cylinder converts fluid pressure into linear piston force. **(D)** A dielectric elastomer thins and expands laterally under an applied electric field. The panels use a consistent schematic style to emphasize the distinct energy-conversion mechanisms and design tradeoffs summarized in the text [**??**].

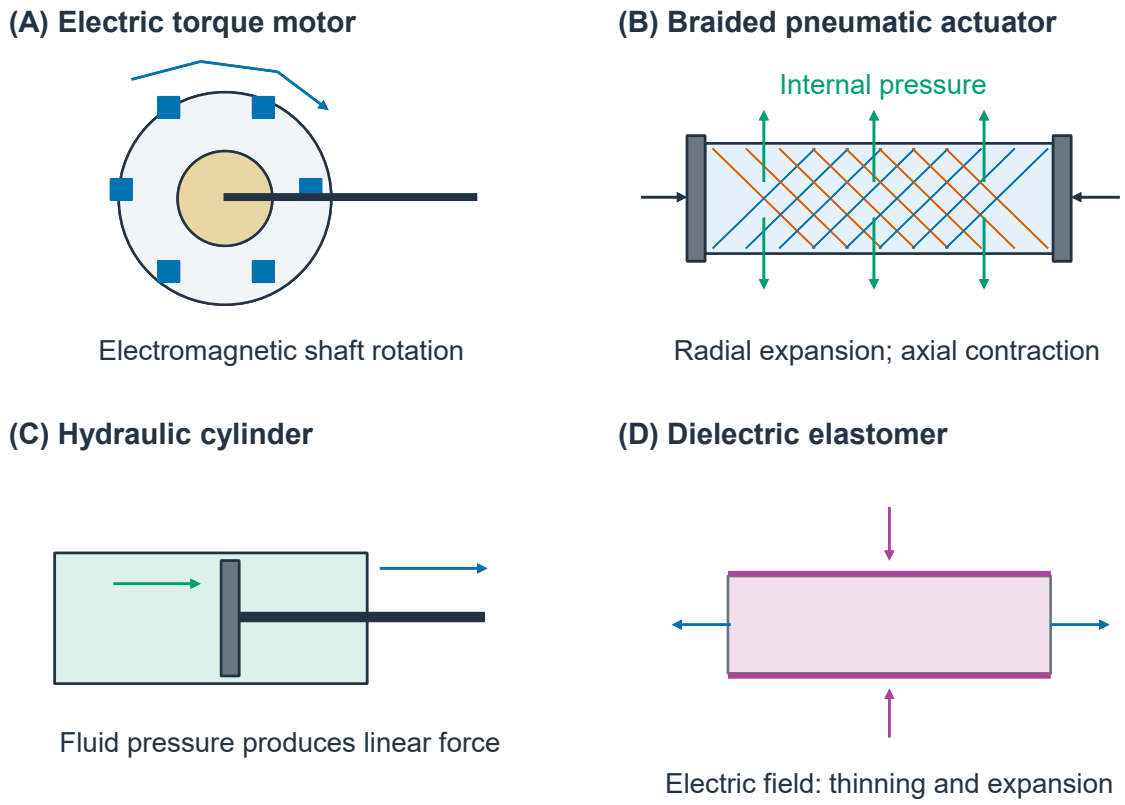


Figure 2.2: Representative actuator mechanisms: electric motor (A), braided pneumatic actuator (B), hydraulic cylinder (C), and dielectric elastomer (D) [**??**].

2.2.2 Braided Pneumatic Actuators

In contrast to the technologies above, McKibben-style braided pneumatic actuators (BPAs) combine low mass, high force-to-weight ratio, and intrinsic compliance with a force–length curve that qualitatively resembles that of skeletal muscle [**??**]. A BPA consists of an inflatable bladder inside a braided sleeve; pressurizing the bladder expands the bladder radially and contracts the sleeve axially, producing tensile force. As shown by Chou and Hannaford, the force–length relationship of McKibben actuators is governed by geometric constraints imposed by the braided shell, leading to a nonlinear dependence of force on actuator length through the braid angle [**?**]: actuator force is proportional to pressure and to a nonlinear geometric term $(3\cos^2\theta - 1)$, which governs both the maximum contraction limit and the curvature of the force–length relationship. Predicting their force as a function of length and pressure is therefore the first improvement of the torque tools developed in this dissertation, and the state of that prediction is the subject of the next subsection.

Figure 2.3 shows the change in actuator shape with inflation.

However, BPAs exhibit several important differences from biological muscle. First, BPA maximum tensile force occurs at resting (i.e., uninflated) length (l_0) rather than at an optimal fiber length [**?**]. Second, their maximum contraction ratio is substantially lower than that of human muscle, presenting challenges for achieving biomimetic joint ranges of motion [**???**]. Third, BPAs have a highly nonlinear dependence on internal pressure (P), contraction (ϵ), and loading state [**??**], which complicates control and motivates high-fidelity force–length–pressure modeling. Despite these differences, their muscle-like properties have made them attractive for biomimetic quadrupeds [**??**], humanoids [**??**],



Figure 2.3: Festo BPAs in the test fixture: the inflated actuator is wider and shorter than the deflated actuator beside it.

and musculoskeletal arms [**??**].

2.2.3 BPA Force–Length–Pressure Models

Previous studies have generated empirical models of the force–length–pressure relationship of BPAs. **?** discuss several high-fidelity BPA force models. They present a static force model with a 21-coefficient polynomial function for a Festo MAS-20-200N (i.e., ∅20 mm, $l_0 = 200$ mm) with an impressive $R^2 = 0.9994$, as well as Sárosi's static force model for a Festo DMSP-20-400N-RM-RM (i.e., ∅20 mm, $l_0 = 400$ mm), with six coefficients and a $R^2 = 0.9995$ [**??**]. Martens and Boblan present an even more accurate model, in terms of absolute error, using five coefficients for a DMSP-10-250 (i.e., ∅10 mm, $l_0 = 250$ mm) and a DMSP-20-300 (i.e., ∅20 mm, $l_0 = 300$ mm) [**?**]. However, as demonstrated in this work, different resting lengths produce different force–length curves, preventing the use of these single-length models for designing systems with different resting lengths.

? looked at six resting lengths of ∅10 mm Festo BPAs, accounting for differences in maximum contractile percentages, and elucidated the force-length-pressure relationship. In particular, for a given contraction percent and force F (in Newtons), the scalar pressure P (in kPa) required to produce this force for a ∅10 mm Festo artificial muscle can be determined by solving the equation:

$$P = P_0 + k_F F + k_S S + P_t \tan\left(c_a \left[\frac{\epsilon}{\epsilon_{\max} - k_E F} - c_b\right]\right) \tag{2.1}$$

where F is the required force, ϵ is the contraction, S encodes hysteresis ($S = 1$ shortening, $S = -1$ lengthening, $S = 0$ static), and the seven coefficients are defined in Table 2.1. The

coefficients have been updated with corrected values, as those reported in **?** contained typographical errors. The contraction ϵ and the maximum free-load contraction at 620 kPa, which appears as $\epsilon_{\max}$ in Equation (2.1), are defined in Section 3.1.1, together with the relative contraction $\epsilon^* = \epsilon/\epsilon_{620}$ on which the force surface developed in this dissertation is built. Because Equation (2.1) expresses pressure in closed form as a function of force and contraction, it can also be solved onto a grid of pressures and contractions to build a force–pressure–relative-contraction lookup table for design studies.

Table 2.1: Coefficients of the 10 mm Festo BPA pressure model of Equation (2.1), corrected from **?**.

Symbol	Value	Role
P_0	254 kPa	pressure offset
k_F	1.23 kPa N^{-1}	linear force-to-pressure gain
k_S	15.6 kPa	hysteresis offset (sign of S)
P_t	192 kPa	tangent-term amplitude
c_a	2.03	contraction-shape coefficient
c_b	0.46	contraction-shift coefficient
k_E	0.331×10^{-3} N^{-1}	force-dependent reduction of usable contraction

Festo provides a manufacturer tool, MuscleSIM, for predicting force as a function of pressure and contraction [**?**], but experimental comparisons indicate that it does not adequately account for variation in l_0 or actuator-specific contraction limits. Comparisons with experimental data indicate that none of the existing models effectively take into account differences in maximum force due to initial actuator length, and the Hunt model was taken with low external forces ($\leqslant$106 N), so it is unclear how well it captures actuator behavior at higher forces. Chapter 3 compares these models with data collected for this dissertation and develops the improved characterization on which the torque and

controller work of the later chapters depends.

2.3 Human Lower-Limb Biomechanics

Designing a biomimetic leg requires a quantitative benchmark of human muscle–joint behavior. This dissertation uses OpenSim [**???**] and the Gait2392 model, named for its 23 degrees of freedom and 92 musculotendon actuators. Gait2392 descends from the interactive lower-extremity model of **?** and includes musculoskeletal geometry from **?** and planar knee kinematics from **?**; related lower-limb models of the same family are characterized and validated by **?**. Muscle moment arms, force-generating properties, and functional roles in such models have been characterized extensively [**??????**], and standard reporting conventions for kinematics and joint mechanics are established by the International Society of Biomechanics [**????**].

Moment arms relate muscle force to joint torque. For a fixed-axis joint, the signed moment arm is $r_{\mathrm{m}} = \hat{\mathbf{a}} \cdot (\mathbf{r} \times \hat{\mathbf{u}})$, where $\hat{\mathbf{a}}$ is the joint-axis direction, $\mathbf{r}$ locates a point on the muscle's line of action relative to the joint, and $\hat{\mathbf{u}}$ is the force direction [**?**]. The tendon-excursion method instead evaluates the sensitivity of musculotendon length to the joint coordinate, with sign determined by the coordinate and force conventions [**?**]. For multiarticular muscles, the result also depends on adjacent-joint configurations [**??**]. OpenSim uses length sensitivities, so the reference moment arms in this dissertation are properties of the Gait2392-derived geometry [**??**]. Chapter 4 compares the robot's resulting torque predictions with test-stand measurements.

Knee modeling requires particular care because the instantaneous center of rotation

changes with flexion. Four-bar descriptions reproduce this behavior through linkage geometry [**??**]. Gait2392 instead prescribes planar translations as functions of knee flexion [**?**]. These different kinematics produce different moment arms for the same nominal attachment locations. Joint-coordinate conventions must therefore be stated consistently when comparing models and measurements [**?**]. Gait2392 provides the human reference in this dissertation, while Steele's four-bar mechanism supplies the physical robot knee.

2.4 Biomimetic Knee and Bipedal Robot Model

Steele's biomimetic knee joint supplies the physical knee used in this dissertation [**???**]. It is a one-degree-of-freedom four-bar linkage whose crossed links reproduce the posterior migration of the human knee's instantaneous center of rotation (ICR) during flexion. The rolling-and-sliding behavior of the human joint therefore emerges from the link geometry itself rather than from cams or gear profiles, and the design was calculated to reproduce an average sliding-to-rolling ratio of 0.079 over a 151° range of motion [**?**]. Because the ICR migrates, muscle moment arms about the knee change with flexion angle in a human-like way, and a torque prediction tool for this knee must capture that variation rather than assume a constant arm. The biomimetic knee tests of Chapter 4 therefore exercise the corrected torque model against this geometry, and Figure 2.4 summarizes the mechanism.

In my master's work, muscle lengths, moment arms, and isometric joint torques were computed over the range of motion for both the Gait2392 model and a PAM-actuated robot leg model, demonstrating that several robot muscle torque profiles could closely match the human model while many fell far short for part of the RoM [**??**]. Hand placement of

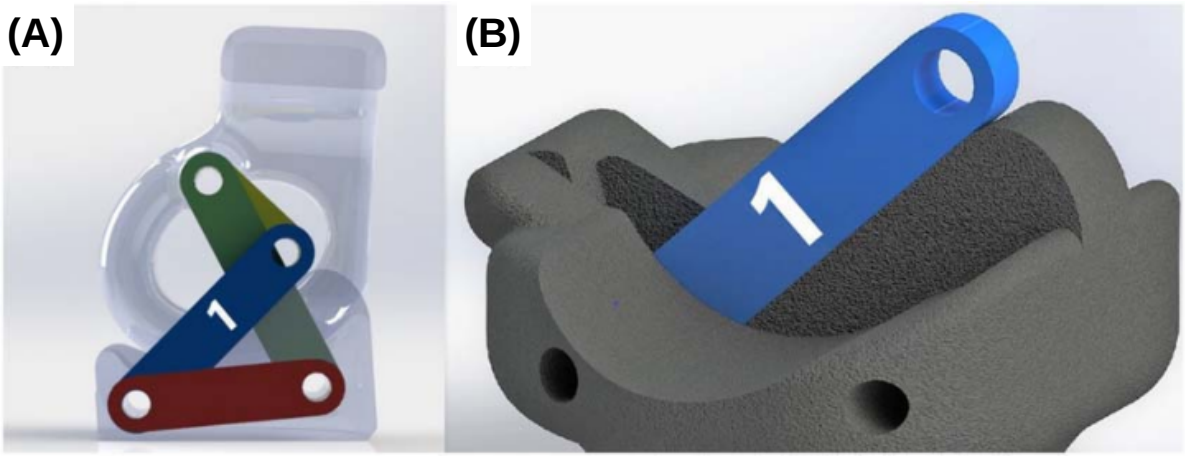


Figure 2.4: Biomimetic knee mechanism. **(A)** Sagittal cutaway at full extension showing the crossed four-bar linkage. **(B)** Exploded view of the mechanical stop. Reproduced from **?**, Fig. 6 under CC BY 4.0.

muscle origins and insertions was a painstaking and iterative task. This resulted in the creation of a placement optimization algorithm [**?**], discussed in Chapter 4. Figure 2.5 places the laboratory skeleton beside the robot PAM routes and human muscle paths used in that comparison.

The comparison evaluated muscle-path lengths, moment arms, and isometric torques across the human and robot models [**?**]. I adjusted BPA origins, insertions, and via-point locations to respect usable contraction and avoid interference. The biceps femoris long head illustrates the remaining mismatch: the robot path and its knee- and hip-torque surfaces differ substantially from the human reference. These discrepancies motivated placement optimization [**?**]. Figures 2.6 and 2.7 show the moment-arm construction and the torque comparison.

(A) **(B)** **(C)**

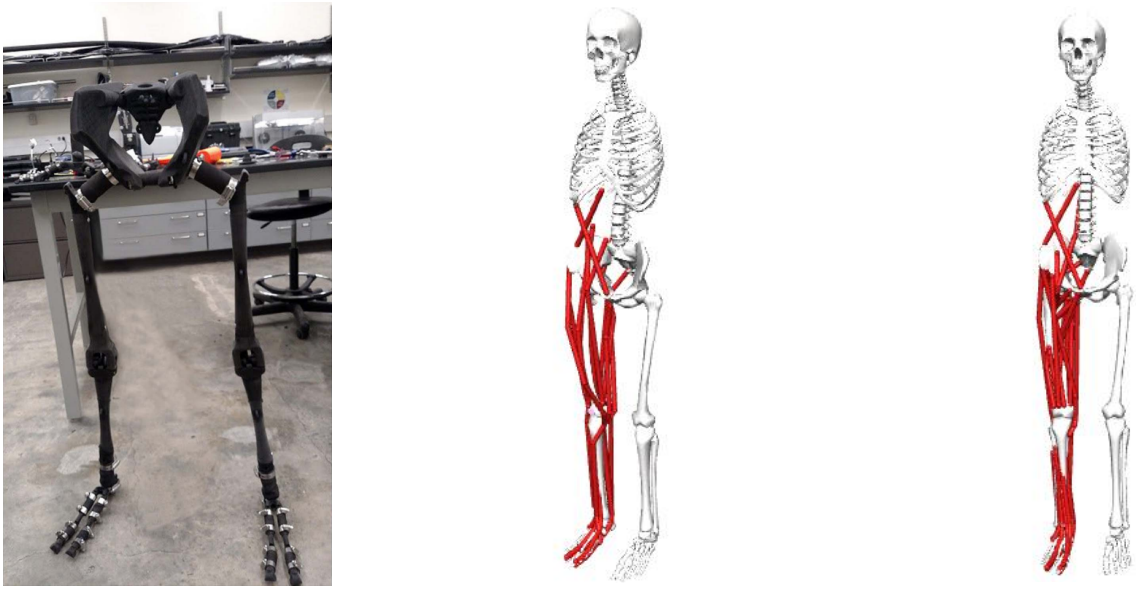


Figure 2.5: **(A)** Laboratory robot skeleton. **(B)** Robot PAM routes. **(C)** Human Gait2392 muscle paths used as design references [**??**].

2.5 Neural Control of Locomotion: Central Pattern Generators

Locomotion in mammals is driven by spinal central pattern generators (CPGs), neural circuits that can produce alternating rhythmic activity of flexor and extensor motoneurons in the absence of rhythmic input and proprioceptive feedback [**??**]. The founding hypothesis is Brown's half-center: alternating flexor and extensor activity arises from two reciprocally inhibitory central populations with opposing bias [**?**]. Modern instantiations of the half-center endow each neuron with persistent sodium channels, whose slow inactivation terminates each burst so that reciprocal inhibition switches the populations without any phasic input [**?**]; rhythmogenesis, in this tradition, is a property of the intrinsic membrane dynamics, not of a clock elsewhere in the network. The modern architectural

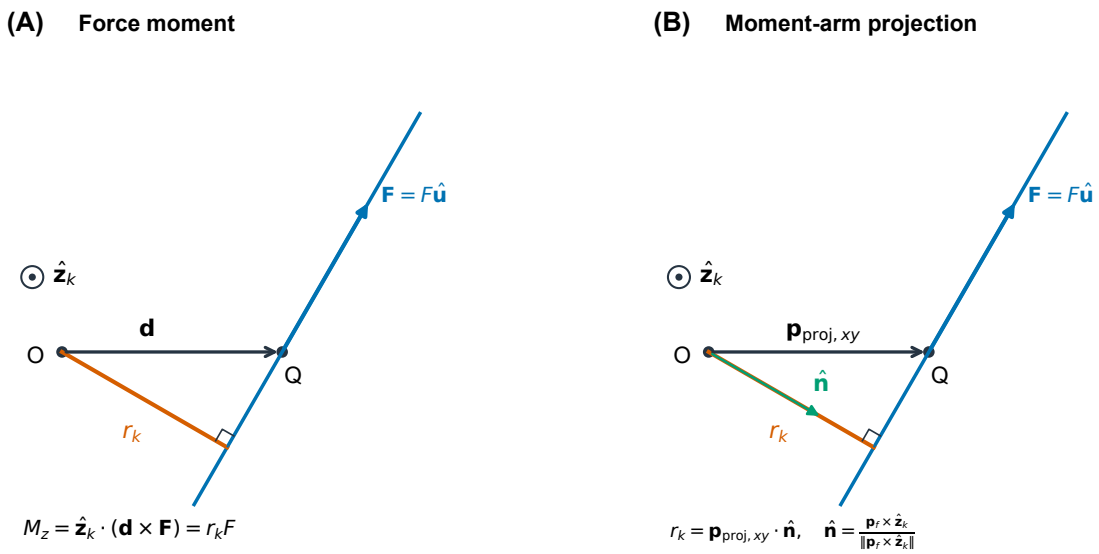


Figure 2.6: Equivalent fixed-axis constructions. **(A)** The cross product gives the moment of a force about the joint. **(B)** Projecting a joint-to-path vector onto the force-path normal gives the signed moment arm used in Equation (4.2). The joint axis points out of the page.

understanding of mammalian locomotor CPGs comes largely from the modeling work of the Rybak group and collaborators. **?** proposed that the locomotor CPG consists of two levels: a half-center rhythm generator (RG) and a pattern formation (PF) network, with reciprocal inhibitory interactions between antagonist neural populations at each level. This two-level organization explains, among other phenomena, spontaneous deletions of motoneuron activity during fictive locomotion in which the locomotor period and phase are maintained, and it permits independent control of the level of motoneuron activity and of step-cycle timing. **?** reviewed the electrophysiological support for this organization, and **?** related the model architecture to genetically identified spinal interneuron classes. Left–right coordination is provided by commissural pathways between the two half-centers [**??**], and the rhythmogenic core of the CPG continues to be refined experimentally [**?**].

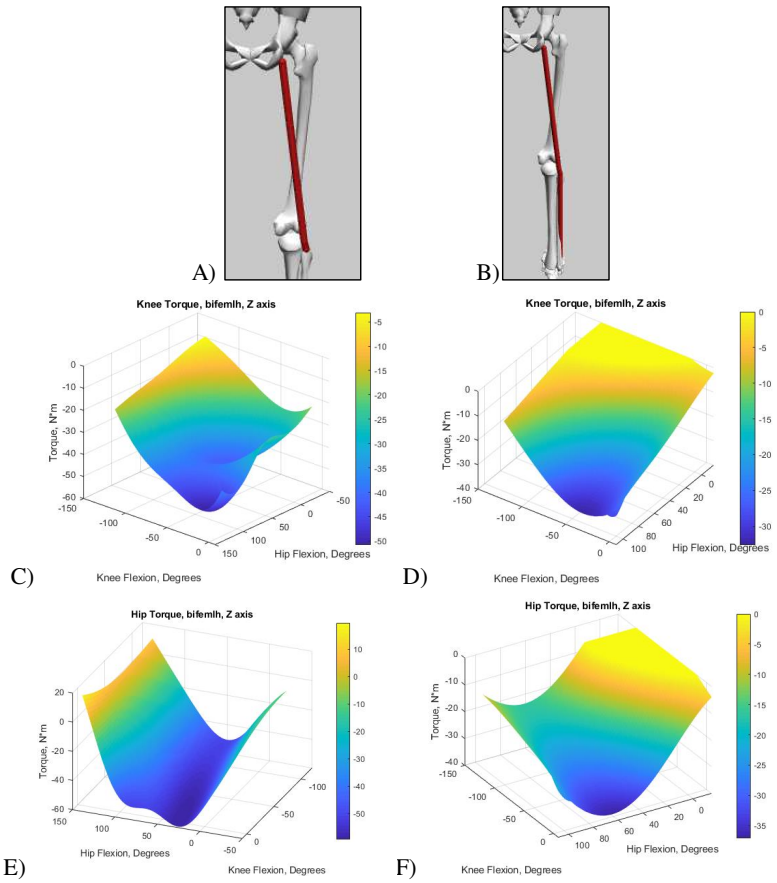


Figure 2.7: Biceps femoris long-head paths and torque surfaces in OpenSim (A, C, E) and the robot model (B, D, F). Panels C–D show knee torque; E–F show hip torque over the hip–knee range [**?**, Fig. 8].

The two-level CPG organization is the current working architecture, and its models account for specific observations beyond the deletions that motivated the split: the genetically identified interneuron classes mapped onto the model [**?**] and the organization of circuits controlling two-joint muscles [**?**], a convergence problem a robot muscle routed across two joints also presents. Coupling such circuits to musculoskeletal models shows how afferent feedback shapes CPG output [**?**] and how biomechanics and sensory feedback regularize the behavior of different CPG architectures [**?**]. Physiologically plausible models of human locomotion use the same machinery to dissect the relative contributions of reflexes and CPGs, with locomotor primitives, Ia, II, and Ib pathways, and Renshaw cells represented explicitly [**?**], and recent models address interlimb coordination under abrupt perturbations such as the sudden loss of ground support [**?**]. How this architecture is expressed as simulable circuit models is the subject of Section 2.7; the locomotor controller developed in Chapter 6 is built directly on this two-level organization.

2.6 Sensory Feedback in Locomotion

Sensory feedback is not a passive correction layer but an integral part of locomotor pattern generation. Reviews of afferent control of the CPG [**?**] and of muscle proprioceptive feedback and spinal networks [**??**] summarize the classical picture: group Ia afferents from muscle spindles signal muscle length and velocity, group Ib afferents from Golgi tendon organs signal force, and group II afferents signal length, with load-related feedback playing a role distinct from length-related feedback during locomotion [**?**]. Removing proprioceptive feedback degrades locomotor patterns in mice [**?**]; proprioceptive input

regulates and adapts locomotor activity [**?**]; and the relative contributions of proprioceptive and vestibular systems to locomotion are the subject of ongoing study [**??**]. Tonic central and sensory stimuli have been shown to facilitate involuntary air-stepping in humans [**?**], demonstrating that simple graded stimulation of motoneuronal and rhythmic circuits can elicit stepping-like output, an experiment readily reproduced in simulation by tonic stimulation of pattern-formation and rhythm-generator populations. The proprioceptive and contact feedback channels designed in Chapter 6 follow this afferent organization.

Three published circuits, reproduced in Figures 2.8, 2.9, and 2.10, place these afferent channels in their circuit context. The rat hindlimb controller of **?** is organized exactly as Section 2.5 describes: a rhythm-generator half-center and a pattern-formation layer sit above a sensory–motor network in which extensor and flexor muscle receptors feed back through interneurons and Renshaw cells onto the same motoneurons the CPG drives. The human-scale model of **?** carries the idea to a full lower limb: each agonist–antagonist muscle pair carries Ia, II, and Ib reflex populations with Renshaw cells, and pattern-formation nodes from the locomotor controller terminate on the same motoneuron pools. The cat interlimb model of **?** adds the left–right half of the problem, coupling the two rhythm generators through V3 and commissural interneuron populations while force, length, and velocity feedback returns from every muscle. The schematics differ in animal, scale, and emphasis, but they wire the same afferent classes of this section into the circuit at the level each hypothesis concerns, whether that is the rhythm generator, the pattern-formation layer, or the motor pool.

The two-level organization these circuits share descends from the modeling program of the Rybak group, and its foundational schematics are reproduced in Figures 2.11 and 2.12.

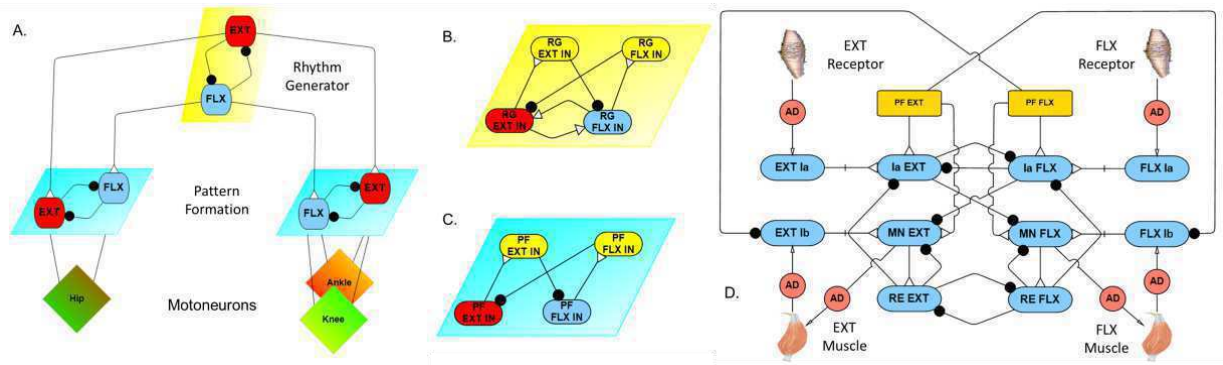


Figure 2.8: Synthetic nervous system of the rat hindlimb model: **(A)** single limb control hierarchy, **(B)** rhythm-generator configuration, **(C)** pattern-formation network, and **(D)** sensory–motor network. Reproduced from **?**, Fig. 4 under CC BY 4.0.

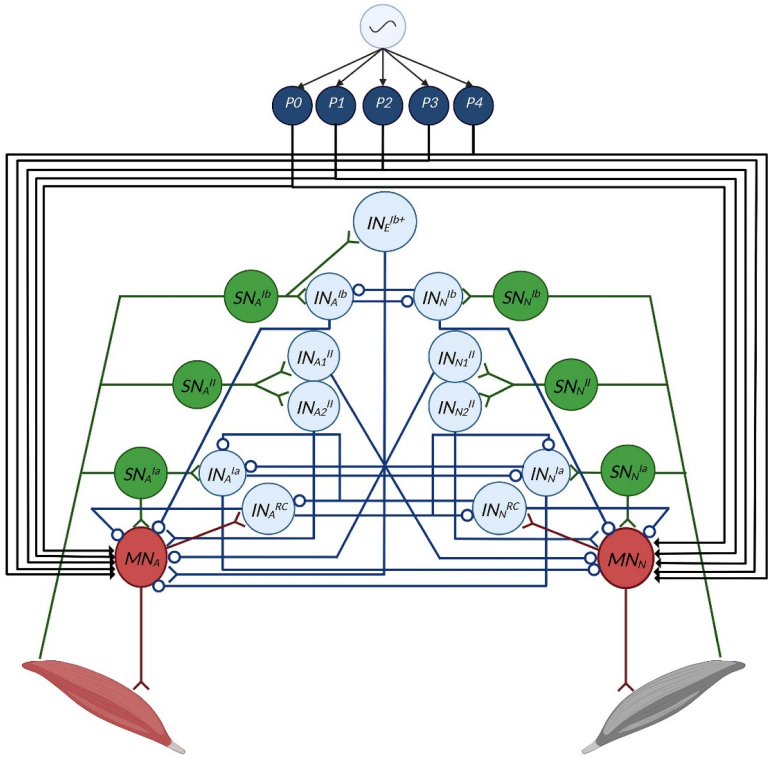


Figure 2.9: Antagonist-muscle spinal network with Ia, II, and Ib reflex pathways, Renshaw cells, and locomotor-pattern input. Reproduced from **?**, Fig. 5 under CC BY 4.0.

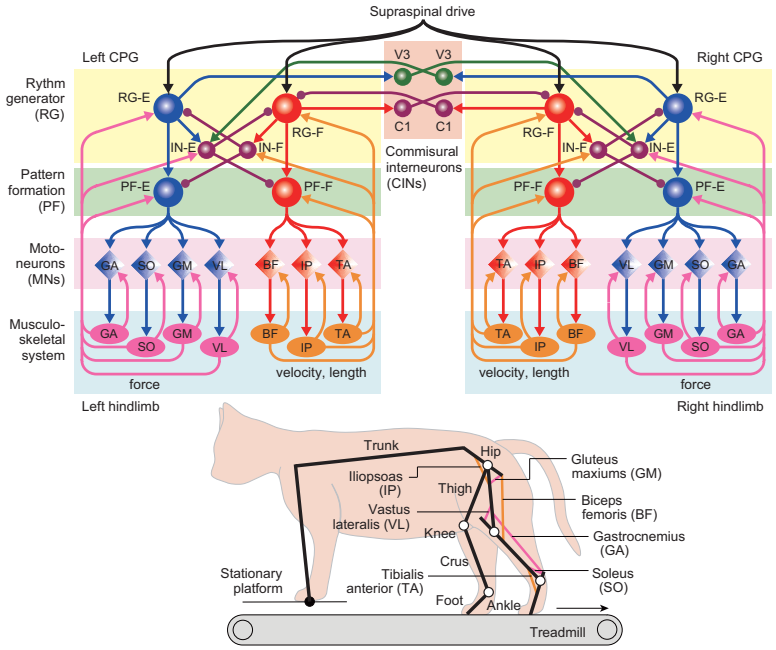


Figure 2.10: Neuromusculoskeletal cat model of **?**: bilateral rhythm-generator and pattern-formation layers coupled by V3 and commissural interneurons, motoneuron pools for seven muscles per hindlimb, and force, length, and velocity afferent feedback. Reproduced from **?**, Fig. 1 under CC BY 4.0.

In the concept model of **?**, afferent feedback ascends to both the rhythm generator and the pattern-formation network, so a sensory perturbation can reset the rhythm at the generator or delete motoneuron activity at the pattern level without necessarily doing both, precisely the deletion phenomenology that motivated the two-level split. The extended model of **?** carries the same organization to the two-joint muscles, embedding dedicated pattern-formation populations for the bifunctional posterior biceps–semitendinosus and rectus femoris pools, and that biarticular convergence problem is the one a robot muscle routed across two joints presents to the controllers developed in this dissertation.

2.7 Computational Neuroscience

Computational neuroscience expresses hypotheses about neural circuitry as executable models. In conductance-based models, membrane voltages evolve through capacitance, leak, synaptic, and applied currents, and connections represent identifiable pathways [**?**]. Model predictions can be compared with recordings or with responses to stimulation and pathway removal. This provides a way to test a proposed mechanism rather than infer it from a wiring diagram alone. Synthetic nervous systems use this approach to control physical and simulated bodies.

2.8 Synthetic Nervous Systems

For robots, these biological architectures can be implemented as synthetic nervous systems (SNSs): networks of conductance-based spiking and non-spiking neurons that are small enough to run in real time yet biologically grounded enough to be compared against

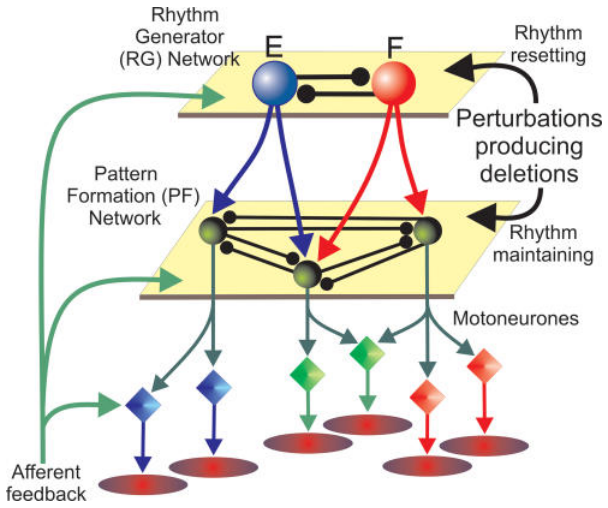


Figure 2.11: Two-level CPG: rhythm generation (RG), pattern formation (PF), motor output, and afferent feedback. Reproduced from **?**, Fig. 1.

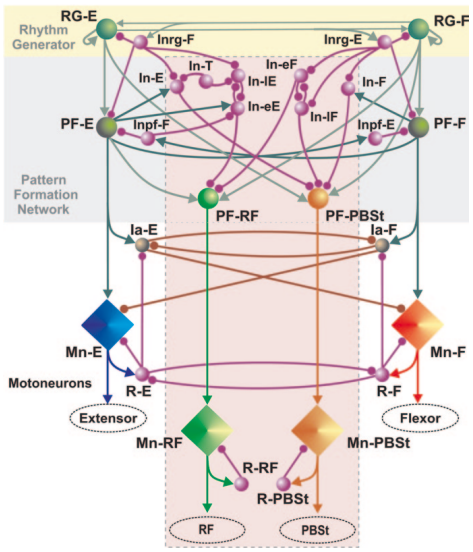


Figure 2.12: Extended CPG with pattern-formation and motor populations for the two-joint posterior biceps–semitendinosus (PBSt) and rectus femoris (RF). Reproduced from **?**, Fig. 5.8.

animal data. The approach has been pursued most persistently by the insect-biorobotics lineage of the Quinn laboratory at Case Western Reserve University and by laboratories closely connected to it, including my own. Works in this lineage span several platforms: CPG-based neural control of a pneumatically actuated dog robot [**?**], analytically designed dynamic neural networks for human balance control [**?**], the full-muscle rat hindlimb model and its moment-arm analysis [**?**], the two-layer CPG neuromechanical walker built on it [**?**], which supplies this dissertation's controller architecture, and the Muscle Mutt quadruped, a canine-inspired BPA robot designed at my laboratory, the Agile and Adaptive Robotics Laboratory (AARL) at Portland State University, for synthetic nervous system control [**?**]. Related controllers from the wider neuromechanical community include sensory-entrained CPG control of bipedal walking [**?**], inter-leg coordination analyses in stick-insect-inspired controllers [**?**], reinforcement-reflex hybrids in insects [**?**], and adaptive hindlimb walking controllers [**?**].

The SNS-Toolbox [**?**] makes such networks practical: it is an open-source Python package that simulates hundreds to thousands of spiking and non-spiking neurons in real time or faster, supports multiple software and hardware backends including GPUs and embedded platforms, and couples directly to physics simulators such as MuJoCo for musculoskeletal control. The field it serves is active and growing: the perspective review of **?** traces neuromorphic legged-locomotion control from its insect-robot beginnings through current platforms to open problems, and physical testbeds in this lineage now extend from insect-scale robots to a robotic cat hindlimb [**?**].

The boundary between background and contribution runs through this section as well. The prior AnimatLab walker with which the author's program first applied these

architectures to a biped [**?**] is presented as prior work of this dissertation in Section 6.1, the Sensory Afferent Database that organizes the afferent data such networks require [**?**] is the subject of Chapter 5, and the simulation pipeline that couples the controller to the robot model is built in Chapter 6. None of it is repeated as background here.

2.9 Neuromechanical Simulation

Neuromechanical simulation, the coupling of a musculoskeletal model to a model of the neural controller that drives it, has advanced along three partially overlapping lines. The first line embeds reflex physiology directly into the controller. The muscle-reflex walking model of **?** demonstrated that positive force, length, and velocity feedback around each muscle, organized with simple interlimb coordination, is sufficient to produce self-stabilizing human-like walking dynamics and realistic muscle activations without any central pattern generator. That model became the reference point for a generation of reflex-based neuromechanical controllers, including the cat hindlimb models of Section 2.7 [**??**] and their extensions. Hybrid work makes the complementarity of the two strategies explicit: **?** combined feed-forward muscle activations with neuronal reflex feedback in simulated human walking and found that the robustness benefit of each depends on which muscle carries it.

The second line replaces hand-designed circuitry with machine learning. **?** trained deep reinforcement learning policies on a physiologically realistic musculoskeletal model in OpenSim, reproducing human kinematics and muscle activations across locomotion tasks, with the learned policy re-usable for experiments that would be difficult in the

laboratory. This approach inherits the fidelity of the musculoskeletal model but encodes control in opaque function approximations: the trained policy can reproduce behavior, but the individual "neurons" of the controller do not correspond to identifiable physiological circuits, so the class of scientific questions it can answer, such as lesions, deletions, and pathway stimulation, is narrower than its behavioral repertoire. Recent perspectives on predictive musculoskeletal simulation [**?**] place both learned and optimal-control approaches in the same frame: the models are mature, and the open problems are the cost of solving them and the biological interpretability of the control they find.

The third line, in which this dissertation sits, uses synthetic nervous systems whose neurons and synapses explicitly represent candidate physiological circuits. SNS-Toolbox supports conductance-based networks at speeds suitable for selected real-time robotic applications [**?**]. Related work includes motor-module bipedal walkers [**?**], two-layer rhythm-generator and pattern-formation control of a musculoskeletal macaque model [**?**], and CPG control of muscle-driven robotic hind legs [**?**]. These studies establish the value of coupled neural and mechanical models. For the humanoid design studied here, the remaining methodological question is how to implement and verify candidate spinal circuits with human musculotendon mechanics and custom BPA actuation, while distinguishing neural effects from conversion or coupling errors. Chapter 6 addresses that question through specific model implementations, graphical neural blocks, and controlled verification and feedback tests.

2.9.1 Simulation Tools

All three lines depend on a software substrate that is itself part of the method: elaborate musculoskeletal geometry and complex neural models. The tools used in this dissertation are the following.

OpenSim [**??**], distributed through the SimTK (Simulation Toolkit) repository, provides the biomechanical modeling side: musculoskeletal models with detailed bone geometry, muscle routing, and force-generating properties, along with tools for analyzing moment arms and isometric torque. These models are approximations with known limits rather than exact anatomies: muscle paths are interpolations over discrete attachment and wrapping data [**?**]. Their geometry is accurate for a purpose, namely moment arms and torque predictions over the functional range of motion, and that is the standard against which the robot model of this program is compared.

The same SimTK infrastructure hosts SCONE, an open-source package for predictive simulation and optimization of neuromuscular control of biological motion [**?**].

AnimatLab [**?**] provides a graphical environment in which physical bodies and synthetic neural systems are constructed and co-simulated; the AARL's prior walker (Section 6.1) was built in it.

MuJoCo [**?**] provides the physics-engine side: fast, accurate simulation of rigid-body dynamics with contacts, along with muscle-like actuators that include activation dynamics.

MyoConverter bridges the biomechanical and physics-engine worlds by converting OpenSim musculoskeletal models to MuJoCo format with optimized muscle kinematics and kinetics [**??**].

SNS-Toolbox supplies the neuron and synapse models used to simulate the controller [**?**].

I used these tools to develop a robotic OpenSim leg model, convert it to MuJoCo, and couple it to a synthetic nervous system. Chapter 6 describes the assembled interfaces and their evaluation. The coupled simulation supports repeated parameter, stimulation, and pathway-removal tests under controlled conditions. Whether those tests support a biological hypothesis depends on the model's validation and the experimental comparison, as well as simulation speed.

2.10 Testing Neuromechanical Hypotheses on Robots: The Ethical Case

The experiments this research program is designed to enable, including deleting a muscle from a limb, blocking a sensory pathway, lesioning part of a rhythm generator, and applying tonic stimulation to spinal circuits, cannot be performed on healthy humans. Invasive intervention on the human neuromuscular system is limited to clinical populations, is tightly constrained by ethics review, and permits neither the within-subject control nor the reversibility that neuromechanical hypotheses demand. The same experiments in animals also carry real moral cost, and their justification is governed by the Three Rs framework: replacement of animal subjects with non-animal methods, reduction of the numbers required, and refinement of the remaining procedures [**??**]. Robotic and computational models allow proposed neural interventions to be tested under controlled conditions before corresponding animal experiments are considered. The same model and initial conditions can be reused across interventions to isolate the effect of a changed pathway.

Repeating a model does not create independent biological subjects, and its relevance to animal experiments depends on validation against biological observations.

Schematics of decerebrate-cat locomotion emphasize the neural circuits under study, but often leave the surgical preparation out of view. A simplified drawing of a walking cat can make the animal appear intact, alert, or even contented. That appearance should not be confused with the condition of the experimental animal: a brainstem transection has disconnected the forebrain, while neural circuits below the transection can still generate treadmill stepping [**??**]. Figure 2.13, taken from archival footage, shows the physical preparation behind the circuit diagrams [**?**]. This distinction makes the ethical meaning of replacement concrete: the robotic experiments discussed here are intended, in part, to reproduce selected neural mechanisms without repeating these invasive preparations in living animals.

The procedures behind that record are documented in films that remain publicly available. A 50 min classroom recording posted by the UTHealth Houston School of Dentistry first shows an intact cat playing with its handler, righting itself when lowered and placing its paws accurately on the edge of a table, and then documents the decerebration performed for the class [**?**]: the animal is anesthetized with open-drop ether, more than the horse in the same demonstration series received, the carotid arteries are ligated and a tracheal cannula installed, and the skull is opened with a trephine before the brainstem is transected between the superior and inferior colliculi. The treadmill preparations that established locomotor neurophysiology are likewise preserved in the historical literature: **?** reproduce the arrangement of the high-decerebrate cat treadmill experiments, with electromyographic recording from the hindlimb muscles, from the original Russian-language

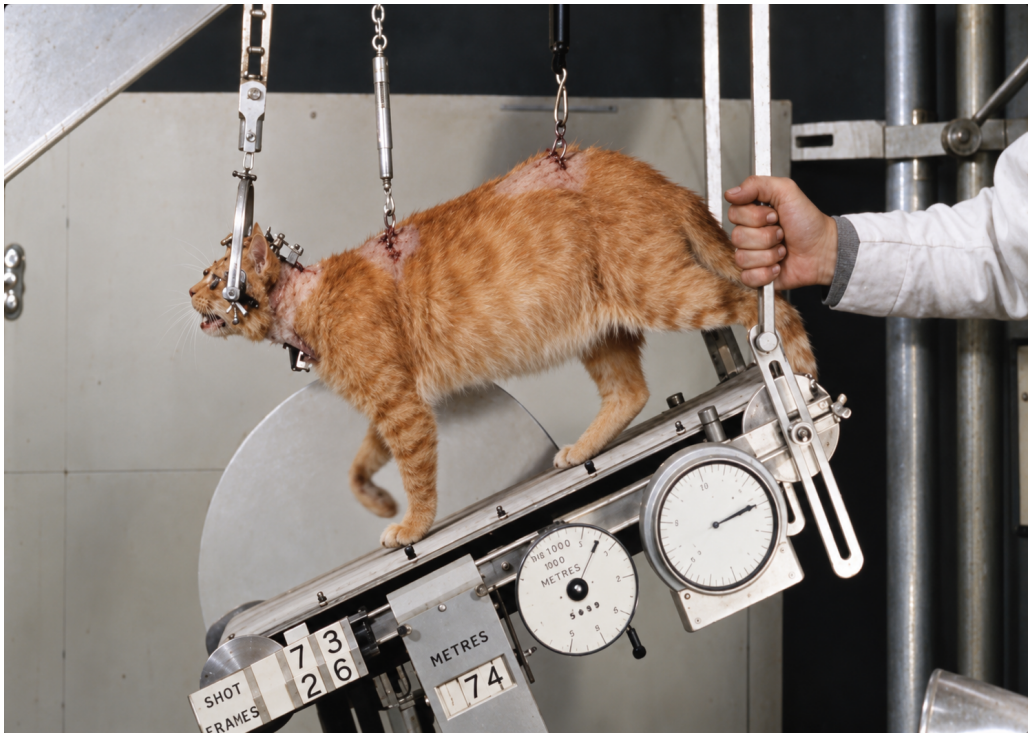


Figure 2.13: Archival decerebrate-cat treadmill experiment [**?**]. The frame is restored and machine colorized; color and fine detail are inferred. It provides documentary context, not gait evidence.

reports [**?**].

The honest accounting is murkier than the Three Rs framing alone suggests, because the models a robot implements are themselves products of animal experimentation. The half-center organization, the two-level rhythm and pattern architecture, and the length, force, and load signaling of the afferents were each established in preparations like the ones described above [**???**], and the databases of afferent properties that parameterize sensory feedback are compiled from animal recordings [**?**]. A synthetic nervous system therefore replaces future experiments; it does not replace, and cannot repay, the animal work on which its architecture is built. Recognizing this inheritance sharpens rather than undermines the case for robotic subjects: the knowledge purchased at that cost is precisely what makes a robotic replacement scientifically meaningful, and the robot is the instrument that keeps that knowledge productive without adding to the ledger.

For a robot to serve as a replacement subject, however, it must first earn standing as a scientific model. Webb's influential analyses of bio-robotic methodology [**??**] argue that a robot is a good model of biological behavior when the mechanism it implements is tested at the same level of organization as the target phenomenon, matching dynamics, control structure, and sensory ecology rather than surface appearance. The validation chain of this dissertation is constructed to that standard: the actuator force model is validated against bench data (Section 3.2.1), the joint torque model against isometric measurements and human torque envelopes (Section 4.2), and selected controller responses are compared with published tonic stimulation and deletion phenomena [**??**], the verification step that precedes locomotion tuning in this program. Only then can a robot experiment be read as evidence about the hypothesized biological circuit rather than about the robot.

The argument is ultimately bidirectional, completing the loop described in Section 1.2. Experiments that are unethical, impractical, or statistically unreachable in humans and animals become routine on the robot; the resulting insights into neuromechanical control feed back into biology, into the design of the next robot iteration, and directly into prosthetic and orthotic devices that face the same problem of controlling muscle-like actuators for human users, whose central pattern generator controllers are an active topic in lower-limb exoskeleton research [**??**]. The ethical case and the scientific case are the same case: the robot is the instrument that makes these experiments possible.

3 Isometric Force Characterization of Braided Pneumatic Actuators

Publication note: Sections 3.1.1–3.1.2, Section 3.2.1, and Sections 3.3.1–3.3.5, with their associated figures and tables, appeared in **?** (Bolen, Elzein, Pang, and Hunt, *Actuators*, 2026). For that work, B.B. performed the conceptualization, methodology, software, validation, formal analysis, investigation, data curation, visualization, and original draft preparation; M.E. and L.P. contributed methodology, software, investigation, and data curation; and A.H. supervised the work, administered the project, and acquired funding. All co-authors contributed to writing, review, and editing.

This chapter develops and validates the actuator force models on which the rest of the dissertation builds, and it reads as one self-contained study, from methods through results to discussion. Sections 3.1.1 and 3.1.2 present the methods: the force equations used throughout this work, and the test jig and measurement procedure used to characterize Festo BPAs with 10 mm and 20 mm uninflated diameters over a wide range of resting lengths, contractions, and pressures. Section 3.2.1 reports the resulting maximum-force and normalized force–strain–pressure models and compares them with the manufacturer's tool. The remaining sections discuss the effects of resting length, the nondimensionalized force prediction, the comparison with other published BPA models, the implications for real-time control and simulation, and the analogy between these artificial muscles and biological muscle. The validated force model is the one used for the knee-torque predictions of Chapter 4 and for the BPA actuation of the neuromechanical simulations in

Chapter 6.

3.1 Methods

3.1.1 Actuator Force Model

The actuator force model used throughout this dissertation is described here. First, the amount of contraction ε in a BPA is calculated as

$$\varepsilon = \frac{(l_0 - l_{\mathrm{m}})}{l_0} \tag{3.1}$$

where l_0 is uninflated resting length of the BPA measured from hose clamp to hose clamp. Current muscle length is l_{m}.

Similarly, ε_{620} is the amount of contraction in a BPA without external load when inflated at 620 kPa, calculated as

$$\varepsilon_{620} = \frac{(l_0 - l_{620})}{l_0} \tag{3.2}$$

where l_{620} is the muscle length measured at 620 kPa and free from external loads.

The amount of relative contraction of a BPA ε^* can then be calculated by dividing equation (3.1) by (3.2), such that

$$\varepsilon^* = \frac{\varepsilon}{\varepsilon_{620}} \tag{3.3}$$

describes the BPA's current contraction as a normalized percentage of its total.

The amount of normalized isometric force F^* in a Festo BPA is then calculated as

$$F^*(\varepsilon^*, P^*) = \begin{cases} c_0 \cdot \left(\exp(-c_1 \cdot \varepsilon^*) - 1\right) + P^* \cdot \exp\left(-c_2(\varepsilon^*)^2\right) & \text{for } F^* > 0 \\ 0 & \text{for } F^* \leqslant 0 \end{cases} \tag{3.4}$$

where P^* is pressure P normalized by 620 kPa. BPA diameter dependent coefficients are determined by the actuator characterization experiments reported in Section 3.2.1 (Table 3.2).

The maximum force F_{max} in a BPA is calculated as

$$F_{max}(l_0, P_{\max}) = a_1 \cdot P \cdot \arctan\left(a_2 \cdot P \cdot (l_0 - 0.0075)\right) \tag{3.5}$$

with coefficients determined by the actuator characterization experiments reported in Section 3.2.1 (Table 3.1). When maximum system pressure is set to $P = 620\,\text{kPa}$, $F_{\max}(l_0, 620) = F_{620}(l_0)$. Multiplying (3.4) by (3.5) yields

$$F(\varepsilon^*, P^*, l_0) = F^*(\varepsilon^*, P^*) \cdot F_{620}(l_0) \tag{3.6}$$

as the equation for isometric force F in a Festo BPA.

3.1.2 BPA Isometric Force Characterization Experiment

To measure artificial muscle force as a function of length and pressure, we built a test jig from extruded aluminum and 3D printed parts. Festo BPAs with 10 mm and 20 mm diameter were tested with various resting lengths and different amounts of contraction.

These data are used to develop an improved model that more accurately predicts maximum muscle force based on l_0, current muscle length l_m, and pressure P. The results are compared with existing models and the new data are used to create an updated model.

A test jig was made of 80/20 brand 1515 series extruded aluminum (Figure 3.1). Artificial muscles were placed vertically in the jig one at a time. The upper end was attached to a load cell. Force data was collected using a CALT DYLY-103 100 kg S-shaped load cell for the ∅10 mm BPAs and a Loadstar RAS1-2HKS-S for the ∅20 mm BPAs. The load cell was used in conjunction with a HX711 Load Cell Amplifier. Pressure data was obtained from a Freescale MPX5700 GP 5 V pressure sensor. Building air supply pressure was controlled with two pressure regulators in series. The first was a Parker model 20R113GC $0 - 120$ psi pressure regulator. The second was a Husky 3/8 in. High Performance Air Regulator HDA72200. A Festo VTUG-10-MRCR-S1T-26V20-T516LA-UL-532S-8K valve manifold VTUG-G was used to deliver air from the pressure regulator to the actuator. This manifold is comprised of eight two-in-one bidirectional normally closed Festo VUVG-S10-T32C-AZT-M5-1T1L valves.

Pressure and load cell amplifier data were sent to MATLAB 2021b via an Arduino Uno style Sparkfun BlackBoard C microcontroller. The computers that used MATLAB were running either Windows 10 or Windows 11. During phases when the Arduino was collecting force and pressure data to send to MATLAB, the Arduino would also trigger (via MATLAB) an Onsemi 2N4401 NPN transistor to make the valve manifold open or close the valve. For other data collection, the valve was manually opened and closed.

The inner distance between the hose clamps on each BPA was measured with a FANUC tape measure to determine the muscle's length l and resting length (l_0). This is how **?**

defines l_0, although in **?** it was measured to also include end cap length. We then inflated each BPA to $P_{620} = 620\,\text{kPa}$, with one end allowed to move freely in the axial Degree of Freedom (DoF), and measured the length l_{620} to calculate maximum contraction at 620 kPa using Equation (3.2). The distance between the crossmembers was then controlled to obtain different amounts of contraction, the muscles were inflated to various pressures, and the contractile force was recorded at the pressure–contraction pairs. This was done for ∅10 mm BPA resting lengths (l_0) of 112 mm–518 mm. For ∅20 mm BPAs, resting lengths (l_0) of 300 mm–509 mm were used.

3.2 Results

3.2.1 BPA Force Characterization

This section reports the results of the isometric force characterization experiments described in Section 3.1.2. First, maximum isometric force at 620 kPa is modeled as a function of resting length. The remaining data are then normalized to build a predictive force surface, and the resulting model is compared with the manufacturer's tool and with other published BPA models.

Maximum Force at 620 kPa

Data of the maximum force from BPA characterization tests of the 10 mm and 20 mm diameters at 620 kPa show a dependency on resting length (Figure 3.2). This is a previously unreported characteristic of these artificial muscles. Detailed analysis of the Festo tool [**?**]

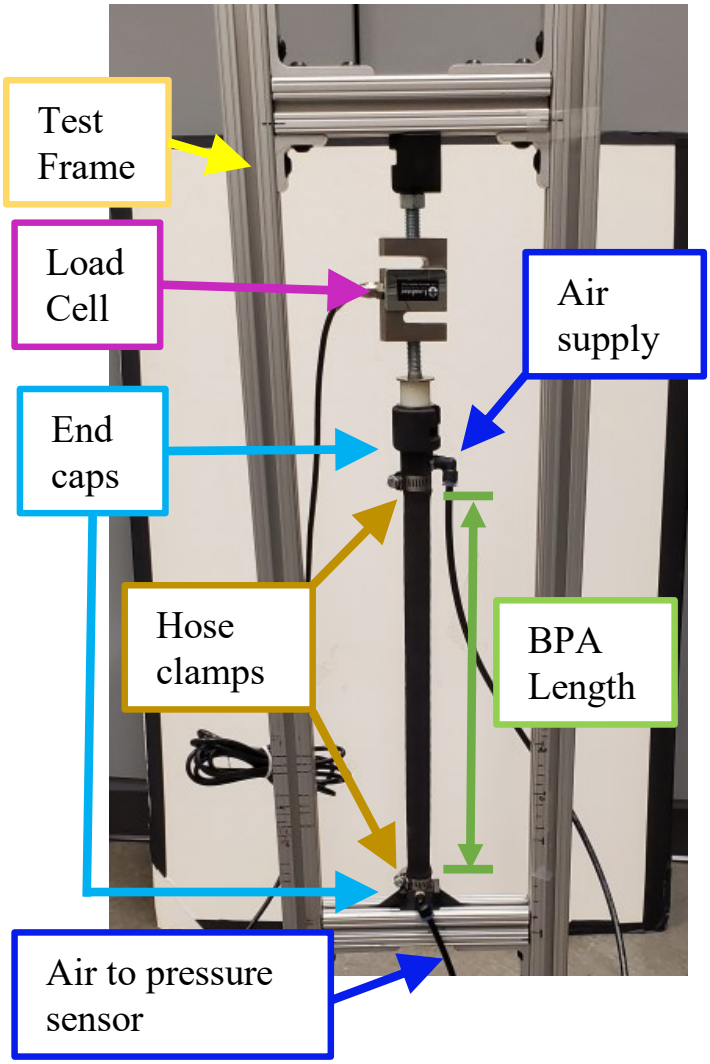


Figure 3.1: A picture of a BPA in the isometric force test stand with components labeled.

does in fact predict a change in maximum force with the resting length; however, the Festo tool predicts increased force with shorter lengths, while our collected data indicate decreasing force with shorter lengths. The data show a force response resembling an arctan curve along the l_0 dimension. Using the Nonlinear Least Squares method and a Least Absolute Residual robustness, we fit an arctan curve to the data to get the maximum force at 620 kPa given a resting length as:

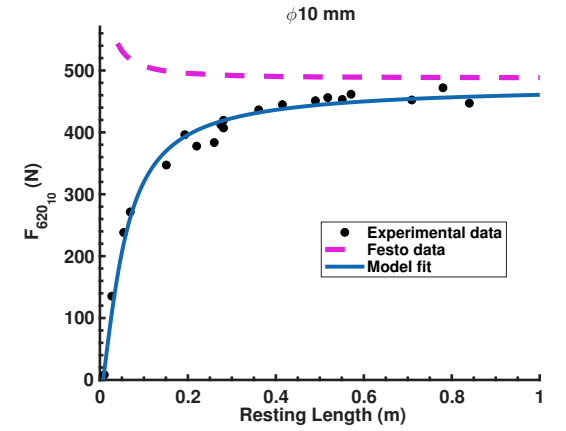
$$F_{620_{10}} = 303.5\,\mathrm{N} \cdot \arctan\left(19.03\,\mathrm{m}^{-1} \cdot (l_0 - 0.0075)\right) \tag{3.7}$$

$$F_{620_{20}} = 922.4\,\mathrm{N} \cdot \arctan\left(15.37\,\mathrm{m}^{-1} \cdot (l_0 - 0.013)\right) \tag{3.8}$$

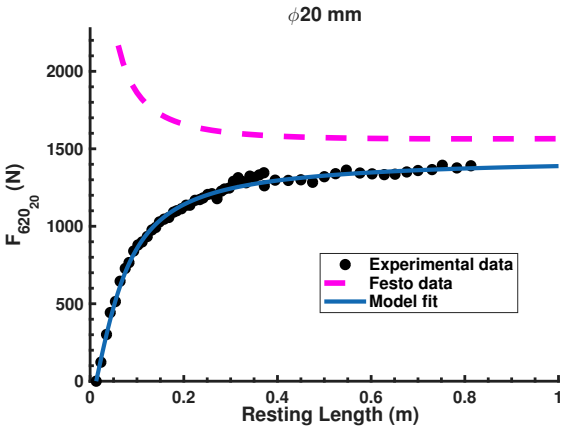
The length is offset by 0.0075 m and 0.013 m because solid modeling showed that the end caps contact each other at these lengths. At these lengths, the actuator would not be able to contract to produce force.

Force as a Function of Pressure and Resting Length

Data from BPA characterization tests of the 10 mm muscles resulted in force–pressure pairings for different muscle resting lengths (l_0) (Figure 3.3a). Similar to the maximum force data, these data show a force response resembling an arctan curve along the l_0 dimension with a more linear response to changes in pressure. Using the Nonlinear Least Squares method and a Least Absolute Residual robustness, we fit an arctan curve to the data to get the maximum force given a resting length and pressure as described by the model form given in Section 3.1.1, where $a_1 = 0.4895\,\mathrm{N\,kPa}^{-1}$ and $a_2 = 0.030\,68\,\mathrm{kPa}^{-1}\,\mathrm{m}^{-1}$



(**a**)



(**b**)

Figure 3.2: Results for finding the relationship between F_{620} and l_0. (**a**) $F_{620_{10}}$ as a function of l_0, at P_{620}. Dashed line is the $F_{620_{10}}$ data from Festo. Solid line is the fit from Equation (3.7). (**b**) $F_{620_{20}}$ as a function of l_0 at P_{620}. Dashed line is $F_{620_{20}}$ data from Festo. Solid line is the fit from Equation (3.8).

for the 10 mm actuator, and $a_1 = 1.49\,\mathrm{N\,kPa^{-1}}$, $a_2 = 0.0248\,\mathrm{kPa^{-1}\,m^{-1}}$ for the 20 mm actuator. Goodness-of-fit measures are given in Table 3.1; throughout this dissertation, goodness of fit is reported as the root-mean-square error (RMSE), the adjusted fraction of variance unexplained (FVU, for which $R^2 = 1 - \mathrm{FVU}$), and the maximum absolute residual, as defined alongside the reference implementation in Appendix A.

Maximum Contraction at 620 kPa

Figure 3.3(b) shows an attempt at a linear fit for maximum contraction at 620 kPa (ϵ_{620}) as a function of resting length (l_0). Error bars have been included to show the effect of ± 1 mm accuracy with measurements of l_0 and l_{620}. There was a large amount of variance in the data, with the linear fit giving an adjusted $\mathrm{R}^2 = 0.4124$ and an RMSE $= 0.0083$. Since there is no direct, predicable relationship between maximum contraction and l_0, this value should be recorded in each muscle used on a robot to best predict the force it may produce at different pressures and contraction.

Force as a Function of Pressure and Contraction

In addition to being a function of the pressure and resting length, the force produced by the actuator is also a function of the amount of actuator contraction, with less force being applied as the actuator contracts. To build this model, the collected data of force, pressure, and contraction are normalized by dividing by F_{620}, P_{620}, and ϵ_{620}, respectively. This had the effect of compressing the data into a 3D surface ranging from 0 to 1 on all axes. Normalized data are compared with pressure isoclines of force predicted by the

Table 3.1: Maximum force equation coefficient values with confidence intervals (CI) and goodness-of-fit measures. Equation (3.5) is compared to data taken at 620 kPa. Equations (3.7) and (3.8) are compared against maximum force from using the Festo tool. It was not necessary to fit an adjusted R^2 in these cases.

Equation	Coefficient	CI (95%)	Our Model			Festo Tool		
			Adj. R^2	RMSE	Max. Error	Adj. R^2	RMSE	Max. Error
(3.5)	$a_1 = 0.4895\,\mathrm{N\,kPa^{-1}}$ $a_2 = 0.030\,68\,\mathrm{kPa^{-1}\,m^{-1}}$	(0.4822, 0.4968) (0.0282, 0.03317)	0.9945	11.61 N	55.3 N	0.9854	14.7 N	30.9 N
(3.7)	$b_1 = 303.5\,\mathrm{N}$ $b_2 = 19.03\,\mathrm{m^{-1}}$	(300, 308) (17.48, 20.57)	0.9854	14.72 N	30.9 N	–	189.9 N	375.6 N
(3.8)	$b_1 = 922.4\,\mathrm{N}$ $b_2 = 15.37\,\mathrm{m^{-1}}$	(914.2, 930.7) (14.75, 15.98)	0.9945	23.83 N	62.1 N	–	668.4 N	1590.1 N

Table 3.2: Normalized isometric BPA force equation (Equation (3.4)) coefficient values and goodness-of-fit measures.

BPA	Coefficient	CI (95%)	Model			Validation		
			Adj. R^2	RMSE	Max. Error	Adj. R^2	RMSE	Max. Error
∅10 mm	$c_0 = 0.5682$	(0.5584, 0.578)	0.9998	0.005118	10.3%	0.9994	0.0245409	10.6%
	$c_1 = 4.254$	(4.126, 4.383)						
	$c_2 = 0.5597$	(0.5429, 0.5766)						
∅20 mm	$c_0 = 0.2579$	(0.2401, 0.2756)	0.992	0.02294	6.8%	0.9943	0.0231303	5.7%
	$c_1 = 6.477$	(5.558, 7.396)						
	$c_2 = 1.321$	(1.239, 1.403)						

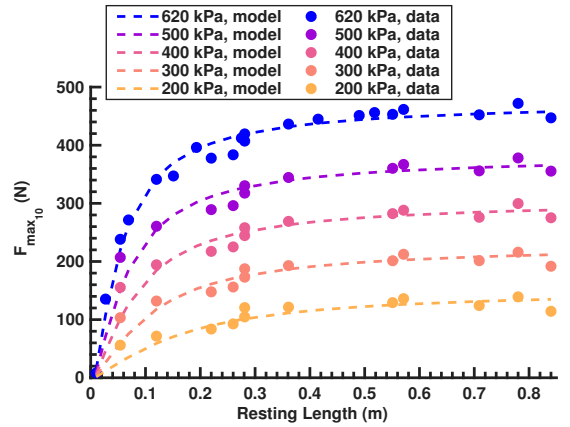
Festo tool divided by the maximum force equation described in the previous section in Figure 3.4. With this comparison, it is clear that the Festo tool over-predicts the expected force, especially at lower pressures, contraction, and resting lengths.

We therefore derived our own equation for isometric force in the BPA as a function of pressure and contraction. Visual analysis of the experimental data shows an exponential relationship between ϵ^* and F^*, and a linear relationship between P^* and F. We fit a surface to the data using Nonlinear Least Squares and a Least Absolute Residual robustness such that it is described by Equation (3.4). For the ∅10 mm BPAs, the result of the improved fit can be seen in Figure 3.4a. The adjusted $R^2 = 0.9998$, a RMSE $= 0.004537$, and a maximum absolute residual of 10.6%. Solving (3.4) for the ∅20 mm BPA, yields different coefficients, and results are seen in Figure 3.4b. Coefficient values and goodness-of-fit statistics are found in Table 3.2.

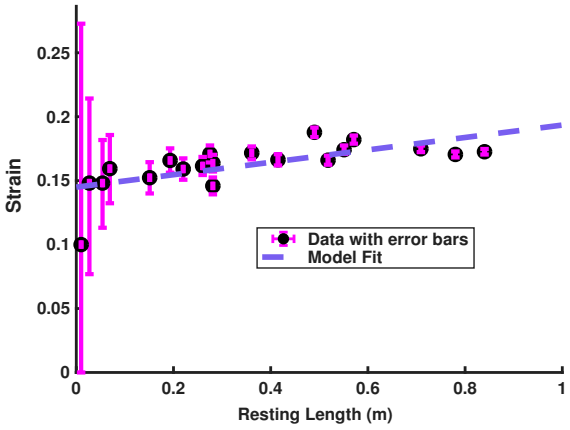
3.3 Discussion

3.3.1 Effects of Resting Length

We have developed new F_{620} equations for 10 mm and 20 mm diameter Festo BPAs that capture the change in maximum force produced by the actuators at 620 kPa as a function of their resting length. When examining Equation (3.7) and measured data in Figure 3.2b, it can be seen that as l_0 goes to infinity, $F_{620_{10}}$ goes to 470 N. When the Festo tool [**?**] was queried at $l_0 = 1$ m as an approximation of infinity, it predicts the ∅10 mm BPA can produce $F = 490$ N at P_{620}. This is within the 10% variability that Festo states may occur in manufacturing tolerances. Similarly, the Festo data predict a ∅20 mm BPA will

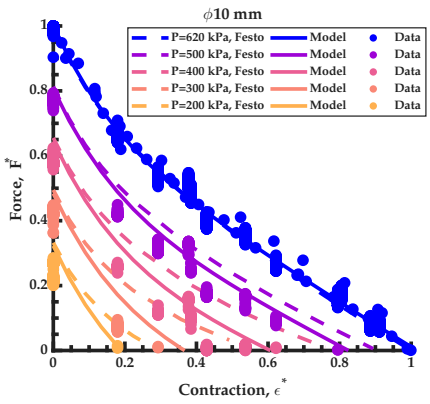


(**a**)

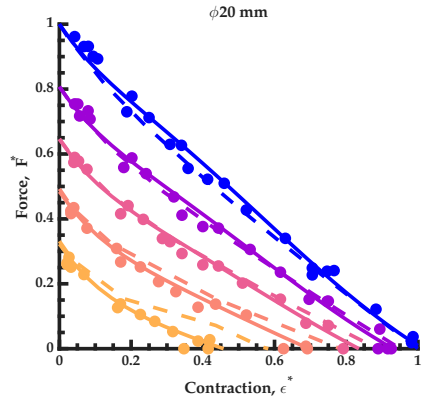


(**b**)

Figure 3.3: Results for finding the relationship between F_{620}, l_0, P_{620}, and $\epsilon_{620_{10}}$. (**a**) Isoclines of the surface fit for $F_{620_{10}}(l_0, P_{620})$. (**b**) $\epsilon_{620_{10}}$ versus l_0 at P_{620}. Although there is a general trend of longer resting lengths producing more contraction, no conclusive relationship between ϵ_{620} and l_0 could be deduced from this experiment. The error bars show the effect of being ± 1 mm in measuring l_0 and l_{620}.



(**a**)



(**b**)

Figure 3.4: Surface fit for $F^*(\epsilon^*, P^*)$. Solid lines are from our model. Dashed lines from Festo supplied data. In total, 537 data points were collected and used to fit the 3D surfaces (321 data points for ∅10 mm BPAs and 216 data points for ∅20 mm BPAs). In total, 80% of these data were used for training, and 20% were used for validation. For figure clarity, not all data is included and plotted circles represent collected data at ± 10 kPa of the stated pressure. (**a**) Fit data for ∅10 mm Festo BPAs. (**b**) Fit data for ∅20 mm Festo BPAs.

produce $F = 1570\,\text{N}$. Our results for Equation (3.8) show that $F_{620_{20}}$ goes to 1460 N as l_0 approaches infinity. Although the difference is much larger, it is also within the 10% manufacturing tolerance.

Our measured data differ significantly from the Festo model as l_0 goes to zero. For the measured data, as the length gets smaller, the force does as well, whereas the Festo tool predicts exponential force increase (see Figure 3.2). As a specific example, when specifying $l_0 = 112\,\text{mm}$ for the 10 mm diameter BPAs, the Festo predicted F_{620} is 498.6 N. However, using Equation (3.7), F_{620} is calculated to be 335.3 N. Actual force in the $l_0 = 112\,\text{mm}$ was measured at 350.9 N. Therefore, the error in predicted F_{620} for the ∅10 mm BPA is 4.5% for our model and 42.1% using the Festo tool. Researchers using BPA resting lengths under 300 mm should take note of these results.

? (p. 12) states that force will be reduced by approximately 10% for a minimum nominal length. In **?** (p. 14) they say the force can deviate by 10% based on factors such as production variation, material variation, and deviation from nominal length (l_0). **?** (p. 13) notes that testing was done with BPAs that are ten times the uninflated diameter, and that theoretical force can increase by up to 10%. In each datasheet, the characteristic curves appear to be the same. As demonstrated above, the 100 mm ∅10 mm BPA should have much less force than what is shown in the curve. This is our experience and the reason for the testing. It is unclear why the MuscleSIM tool [**?**] then predicts higher force as resting length decreases. It could be a reversed sign coefficient in their internal curve-fit model. We have shown here that as the nominal length goes to zero, so does force.

BPA models commonly predict higher forces than are measured experimentally because they do not account for end effects in regions of the pressure vessel where those effects

dominate. Nominally, the clamped ends fix the inner and outer diameter while allowing axial and radial displacement. In the isometric case, there is only radial displacement. The boundary conditions force steep gradients in stress and strain near the ends. Near the boundary conditions, axial, radial, and circumferential strains deviate from analytic solutions for a cylindrical model. For a constant pressure, the end effect region appears to have the same length. Therefore, in shorter BPAs, the end effects dominate. The axial length scale of the boundary section is comparable to l_0, and the active force differs substantially from that of an assumed perfect cylinder.

Manufacturing variations seem to play an important role in the behavior of these actuators. It remains to be seen how ϵ_{620} can be known a priori. We still suspect that it is a function of l_0, but it might also be a function of other factors such as product batch number or UV light exposure [**?**] (p. 15); [**?**] (p. 6); [**?**] (p. 13). We found there to be larger variations in the 10 mm BPAs than in the 20 mm BPAs. Modeling work on ideal McKibben actuators describes how the relationship between contraction and initial length is a function of the total number of twists and the angle of twist in the usable muscle area [**?**]. It is unclear if there are variations in the angle and number of twists in the Festo BPAs due to manufacturing. Uncovering the relationship between ϵ_{620} and l_0 for the Festo DMSP artificial muscles, if it exists in a meaningful way, will require additional controlled tests.

Similarly, it is not clear how the maximum force might change with manufacturing variations. This was reasonably corrected for by normalizing the equations to each individual's maximum contraction, but more nuanced details might emerge with more data collection. Not knowing these relationships a priori means there is still a need for

robots and other systems to go through an iterative design stage if it is discovered that the system does not produce the torque or have the RoM that the design team expects. However, the data and models provided here should enable faster redesigns with fewer additional measurements.

3.3.2 Nondimensionalized Force Prediction

We have also introduced the concept of nondimensionalized isometric force that is a function of relative strain and relative pressure (Equation (3.4)). This elegant equation has only three coefficients with low error (see Table 3.2). In the work presented here, we introduce the relative pressure term, $P^* = P/P_{620}$. Our data show that by normalizing force, contraction, and pressure, we are able to create a simplified force equation as a function of contraction and pressure that scales well with initial actuator length, as expressed by Equation (3.6) in Section 3.1.1.

For example, given $l_0 = 54\,\text{mm}$ and $P = 300\,\text{kPa}$, the measured force was $F_{actual} = 103\,\text{N}$. Force prediction using only the Festo tool predicts the force to be $F_{predict} = 270\,\text{N}$. This is 161% greater than the measured force. Using Equations (3.7) and (3.4) with Equation (3.6) gives $F_{predict} = F^*(0, 0.48) \cdot F_{620}(0.054) = 0.49 \cdot 230\,\text{N} = 112\,\text{N}$. This is an error of 8%, which is much more accurate than using the Festo tool only.

The presented work is capable of predicting force output for 10 mm and 20 mm muscles with different initial muscle lengths, internal pressure, and contraction amount with minimal initial data collection. However, it remains to be seen if these relationships are maintained over the full lifecycle of the muscles. There are not yet any peer reviewed

studies investigating the long-term behavior of these artificial muscles, though the Festo datasheet provides some indication of what can be expected [**?**] (p.19). They report that the 40 mm BPA can complete 1 million cycles at 6000 N (maximum rated load), or 10 million cycles at 4000 N. The service life of the other size BPAs is not reported; however, it is mentioned that thermal stressing may reduce the lifespan and they should be inspected every 500,000 cycles for signs of aging [**?**] (pp. 10–11). Future research should look at the behavior of these muscles over time.

In our work, $P_{620} = 620\,\text{kPa}$ is the maximum supply pressure for our system, and other users of this actuator may use a different supply pressure. However, this supply pressure is not required, and normalizing by this value was chosen as a method for improving the optimization techniques by reducing the sensitivity of the term associated with pressure. This equation will work for any pressure supplies below 620 kPa. We also anticipate the equation will work for pressures up to 800 kPa (the maximum rated pressure for Festo ∅10 mm BPAs); however, we do not have the data to confirm this.

3.3.3 Comparison with Other Published BPA Models

The proposed model was compared with the formulations by **?**, **?**, and **?**, to evaluate their predictive performance across actuator sizes and operating conditions (see Table 3.3 and Figure 3.5). For the ∅20 mm BPA, parameters from [**?**] were used, corresponding to actuators with resting lengths of 200 mm and 400 mm. These were compared against the experimental data at $l_0 = 300\,\text{mm}$ and $l_0 = 450\,\text{mm}$, respectively. For the ∅10 mm BPA with resting lengths of 257 mm and 233 mm, parameters from [**?**] were applied.

In each case, the maximum contraction parameter (k_{max}) was determined numerically by solving $F(6.2\,\text{bar}, k_{max}) = 0$, ensuring consistency with the experimental maximum pressure condition.

While the **?** and **?** models capture general trends in force–contraction behavior, discrepancies were observed when applying parameters derived from different actuator lengths. In particular, the model tends to over-predict force at higher pressures and low amounts of contraction, suggesting sensitivity to geometric scaling effects not explicitly accounted for in the formulation.

The **?** model is derived from a geometric and material-based framework and provides a more physically grounded representation of actuator behavior. However, its accuracy depends strongly on precise geometric parameters, including braid angle, wall thickness, and fiber length. Small deviations in these parameters can lead to significant differences in predicted force, particularly at higher pressures.

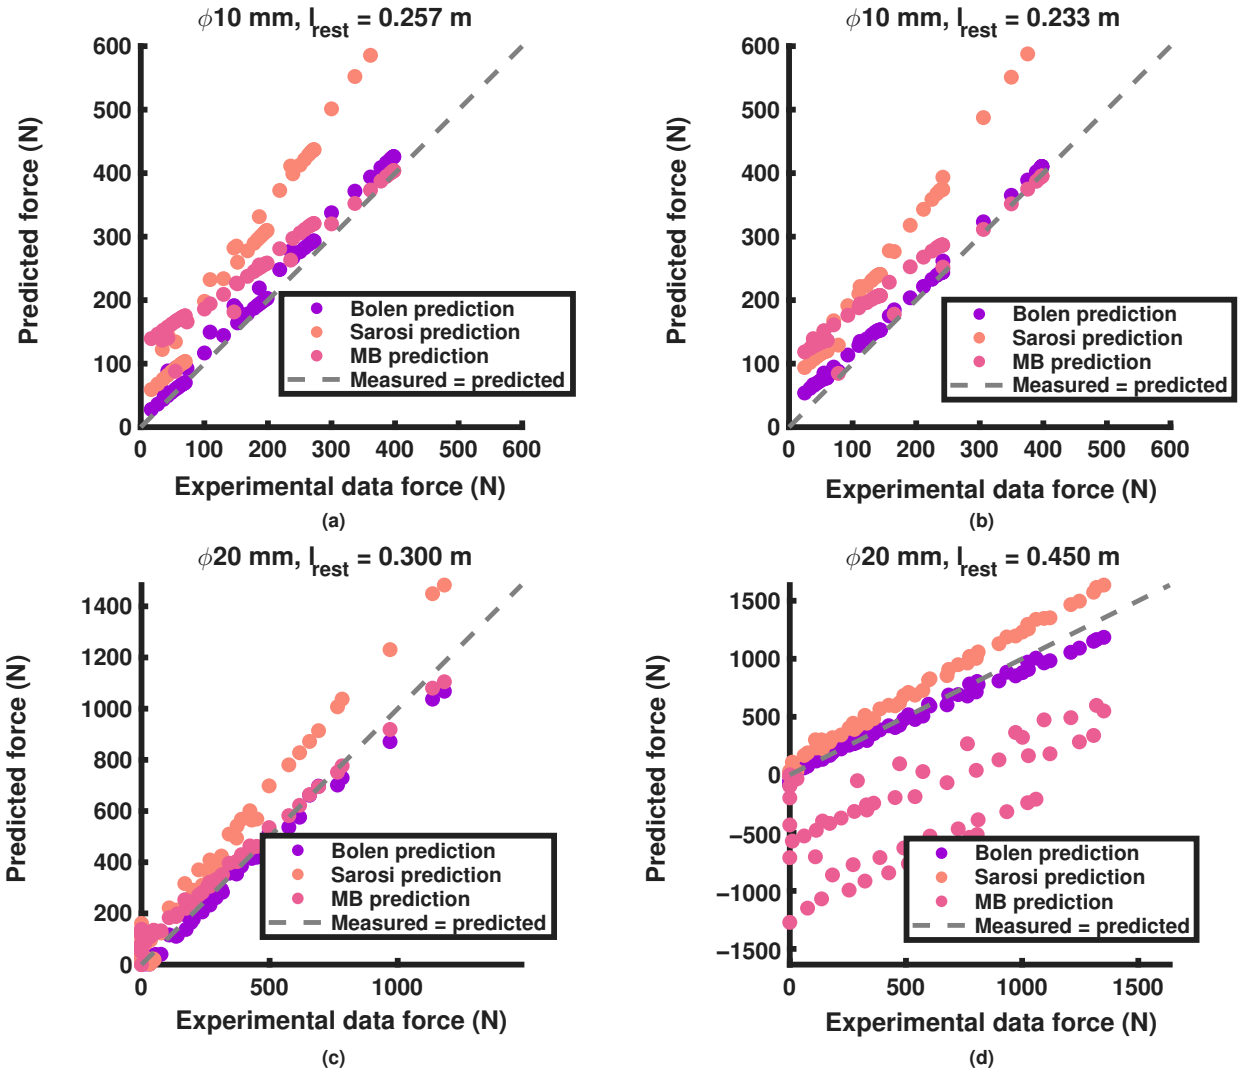


Figure 3.5: Comparison between measured and predicted BPA force at varying amounts of contraction and pressure. A $y = x$ perfect fit line is added for reference. The BPAs compared are: (**a**) ∅10 mm, $l_0 = 257$ mm, (**b**) ∅10 mm, $l_0 = 233$ mm, (**c**) ∅20 mm, $l_0 = 300$ mm, and (**d**) ∅20 mm, $l_0 = 450$ mm.

The Martens and Boblan model was first evaluated at the same pressure–contraction points as in their experimental data [**?**] (Tables A1 and A2), indicating that it reproduces

the general published force map shape but is sensitive to the assumed initial geometry. Examining their ∅20 mm BPA, for example, using the explicitly stated $l_0 = 300\,\text{mm}$ and initial diameter $D_0 = 20\,\text{mm}$ significantly under-predicted the results compared to using the implied $l_0 = 296\,\text{mm}$ ([**?**] (Table A1)) and $D_0 + H_0 = 21.8\,\text{mm}$, where $H_0 = 1.8\,\text{mm}$ is the material thickness. The diameter adjustment can be inferred from [**?**]. This adjusted diameter was also included for the ∅10 mm BPA to give a better fit. As can be seen in Figure 3.5 and Table 3.3, the highest accuracy is observed when ∅20 mm and $l_0 = 300\,\text{mm}$, which is at or very near their experimental measurements. For the ∅10 mm BPAs the accuracy decreases even though their experimentally measured $l_0 = 250\,\text{mm}$ is close to our $l_0 = 257\,\text{mm}$ and $l_0 = 233\,\text{mm}$. The largest error occurs at ∅20 mm and $l_0 = 450\,\text{mm}$ and the model starts to predict negative force. We have calculated that the $F_{PE} \cdot \frac{dD}{DL}$ term becomes very large compared to the others [**?**] (p. 5). This suggests that the model is sensitive to geometric parameter selection and interpretation, particularly with respect to initial length and diameter definitions.

In comparison, the present normalized model demonstrates improved consistency across actuator lengths by incorporating experimentally derived scaling through F_{620} and ϵ_{620}. This normalization reduces sensitivity to geometric variation and enables a unified representation of actuator behavior, while maintaining predictive accuracy over the tested range of pressures and contractions.

3.3.4 Implications for Real-Time Control and Simulation

Although pneumatic muscle actuators exhibit hysteresis, thermodynamic effects, and other dynamic behavior, accurate and high-performance controller designs can still be achieved using a static force map approach when the underlying static force model is accurate. Martens and Boblan explicitly note that BPA dynamics are often dominated by the air dynamics inside the actuator and the mass flow behavior of the valve, such that precise modeling of the static force characteristic remains crucial for torque control accuracy [**?**]. This view is consistent with more recent BPA control work, including a sensor-less torque control interface for BPA-driven joints [**?**], stable force/position/stiffness control of antagonistic pneumatic-muscle joints [**??**], and sensor-less angle stiffness control of antagonistic BPA systems [**?**].

In this context, the present model is well suited for real-time use because it reduces the static isometric force prediction to a low-order normalized surface, $F^*(\epsilon^*, P^*)$, scaled by the resting-length-dependent term $F_{620}(l_0)$. This structure is computationally inexpensive to evaluate online and is also suitable for lookup-table implementation. While the present work does not model the full actuator–valve dynamics, it provides the static force relationship needed as a feedforward or inner-model component in real-time torque control, simulation, and preliminary design workflows. It is anticipated that improved real-time control for high-speed and high-precision systems will benefit more from improved modeling of airflow dynamics than from modeling actuator hysteresis, damping, and other dynamic behavior. Characterizing the actuators under dynamic operating conditions, such as isokinetic (constant-velocity), isobaric, and isotonic (constant-load) contractions,

is a natural next step beyond the isometric characterization reported here, for which the pressure–force dynamics model of **?** provides the modeling foundation.

3.3.5 Biomimetic Artificial Muscle Analogy

BPA artificial muscles are often said to be analogous to biological muscles because they have force–length curves and can produce force only in tension. This analogy is worth closer inspection. For example, the passive force–length curve has been described as an exponential function arising from connective tissue elements, while the active force–length relationship is typically modeled as a bell-shaped curve reflecting actin–myosin overlap [**?**]. These formulations emphasize that muscle force generation is governed by nonlinear geometric and material interactions rather than simple linear elasticity. McKibben-type pneumatic artificial muscles (of which Festo is an example) exhibit analogous nonlinear behavior, though arising from fundamentally different physical mechanisms. As shown by Chou and Hannaford, the force–length relationship of McKibben actuators is derived from geometric constraints imposed by the braided shell, leading to a nonlinear dependence of force on actuator length through the braid angle [**?**]. In their formulation, actuator force is proportional to pressure and a nonlinear geometric term $(3\cos^2\theta - 1)$, which governs both the maximum contraction limit and the curvature of the force–length relationship. The work presented here demonstrates that this nonlinear curve can be well approximated by an exponential as well.

However, the behavior of these artificial muscles differs distinctly from that of biological muscles. For example, the optimal fiber length of a biological muscle, l_{OFL}, is not equivalent

to l_0 in an artificial muscle. Artificial muscles produce their maximum active force at their resting length, which is also their maximum length. In contrast, the maximum active force that a biological muscle can produce is somewhere in the middle of the total contractile length of the muscle.

Furthermore, biological muscles can adapt the optimal operating length through structural remodeling (e.g., sarcomere addition or tendon compliance) while BPAs have a fixed geometric configuration. This allows biological systems to self-optimize for the particular tasks the animal is subjected to. In order to re-optimize a robot for new tasks, new actuators would have to be swapped into the system. This distinction highlights a key limitation of BPAs in biomimetic applications, while also reinforcing the importance of accurate nonlinear modeling for control and design.

The analogy also suggests two design observations. First, the force–length penalty is shared by both families: force falls as the actuator operates away from its optimum, so a limb design must route each actuator such that its highest-force operating range coincides with the joint angles where torque is needed most, because the force lost at the extremes of the range of motion cannot be recovered by pressure alone. Second, the two families gain moment arm in different ways. Skeletal muscle attaches to the skeleton itself, through tendons or directly over broad areas of bone, so the moment arms of large animals follow from skeletal geometry. The robot's BPAs instead act through brackets that protrude from the limb segments; that hardware is mechanically exposed, and its series compliance is one of the loss mechanisms quantified in Chapter 4.

The analogy finally raises a question that this dissertation can pose but not settle. Every muscle-like actuator loses force as it operates away from its optimal length, and a soft

actuator working against soft structure has no rigid skeleton from which to gain a moment arm. Given those two facts, it may be no accident that nearly all large land animals are built on rigid skeletons and that large soft-bodied land animals are essentially absent. Whether the force–length penalty contributes to that absence, or whether it reflects other constraints such as support and protection, is a question for comparative biomechanics rather than for a single actuator study. For the designer, the lesson holds regardless of the answer: some moment arm must come from somewhere, and on this robot it comes from protruding brackets whose behavior must be modeled rather than ignored.

Table 3.3: BPA force-model comparison. Diameter and resting length are in mm; RMSE and maximum absolute error are in N. FVU is dimensionless.

Diameter	Length	Model	RMSE	FVU	Max. error
10 mm	257	Bolen	22.9	0.0353	46.0
10 mm	257	S'arosi	143.8	1.3903	235.8
10 mm	257	Martens and Boblan	73.8	0.3668	122.9
10 mm	233	Bolen	16.8	0.0187	30.3
10 mm	233	S'arosi	132.7	1.1655	222.6
10 mm	233	Martens and Boblan	62.6	0.2590	101.2
20 mm	300	Bolen	35.6	0.0144	114.0
20 mm	300	S'arosi	132.1	0.1992	313.4
20 mm	300	Martens and Boblan	72.8	0.0605	137.8
20 mm	450	Bolen	67.6	0.0258	166.6
20 mm	450	S'arosi	177.7	0.1785	294.8
20 mm	450	Martens and Boblan	862.1	4.2022	1321.1

4 Isometric Knee Torque

Publication note: This chapter is the subject of an unpublished manuscript in preparation by Bolen, Burgess, Morrow, Elzein, Pang, Scharzenberger, and Hunt [**?**]. For that work, B.B. performed the conceptualization, methodology, software, validation, formal analysis, investigation, data curation, visualization, and original draft preparation; M.E., L.P., L.B., and C.M. contributed methodology, software, investigation, and data curation; and A.H. supervised the work, administered the project, and acquired funding. All co-authors contributed to writing, review, and editing.

This chapter connects the actuator force characterization of Chapter 3 to torque about the robot knee. The earlier torque model assumed inextensible muscle paths on rigid links. Isometric tests on pinned and biomimetic knees show where those assumptions fail. The methods develop corrections for length offsets, attachment compliance, and wrapping loss, then explain how actuator-placement searches use the corrected prediction. Held-out tests evaluate the correction terms, and physical tests assess whether the redesigned routes meet selected human-muscle torque targets.

4.1 Methods

I built a modular test stand to measure flexor and extensor torque with either a pinned knee or a biomimetic four-bar knee. The tests varied actuator diameter and resting length, joint contact, and tendon attachment. I first compared measured torque with predictions based on the characterized BPA force model. I then set the length, compliance, and wrapping

corrections using separate training and validation tests, and used the corrected model to redesign routes against human-muscle torque targets. Isometric tests evaluated the resulting hardware.

4.1.1 Robot Leg Model and Knee Test Stands

The primary components of the robot leg are the knee joint, femur, tibia, BPA assemblies, and artificial tendons (Fig. 4.1**A–C**). The artificial bone components are 3D printed using a combination of Onyx, carbon fiber, and high-strength high-temperature fiberglass on Markforged Onyx One and Mark Two printers, as well as machined links for the knee mechanism. Onyx is a proprietary Markforged material that consists of chopped carbon fiber in nylon. Artificial tendons are made with Shimano bicycle brake cable.

The modular test stand accommodates either knee joint and both flexor and extensor configurations (Figure 4.1). This allows measured joint torque to be compared with predictions while changing the joint geometry and actuator route. The extensor test at high flexion also exposes the actuator deformation caused by contact with the knee.

The leg assembly geometry and dimensions are based on bones scans of a 6 ft (1.83 m) tall person [**????**]. Two knee joints were tested, one pure rotational joint, and one biomimetic joint with rotation and translation. Each joint is driven by two antagonistic Festo actuators. Both robot joints have the same center of rotation in the femur frame at $\theta_{\mathbf{k}} = 0\,\text{deg}$.

Details on the biomimetic robot knee used in this study can be found in [**???**]. In brief, this 1-DoF joint (Fig. 4.2) uses a four-bar linkage to change the effective moment arm during rotation. The intersection point of the internal crossed links define the instantaneous

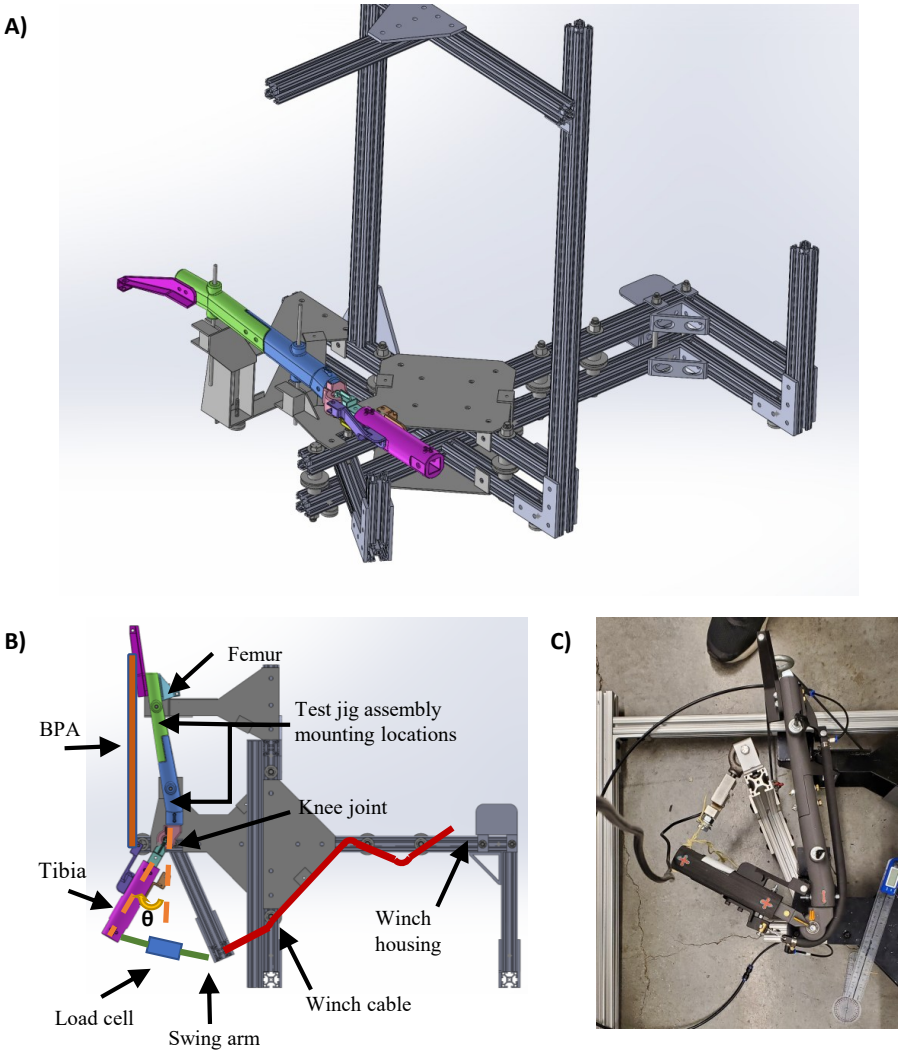


Figure 4.1: Pinned-knee test apparatus: CAD assembly (A), labeled mechanism at -30° (B), and extensor isometric test at -120° (C).

center of rotation (ICR). This causes the joint ICR location in the femur frame to translate posteriorly in the X direction and distally in the Y direction except for a brief proximal move near full flexion. Reference International Society of Biomechanics (ISB) coordinate frames are given by **??**. This is similar to how the human knee behaves [**?????**]. The correct location of the ICR is important for both theoretical and experimental torque calculations.

Both knee types use uniarticular flexor and extensor actuators. Each knee joint type was tested with a single ∅10 mm Festo BPA per muscle. The biomimetic knee was additionally tested with ∅20 mm Festo BPAs in a paired configuration: the extensor with two BPAs in parallel, and the flexor with two BPAs whose pathways are mirrored about the Z axis. Each artificial muscle is pinned at the muscle origin location on the femur while the other end is either pinned or attached to the muscle insertion location on the tibia via an artificial tendon.

The modeled geometry and torque calculation follow **??**. Torque about the knee is generated by the moment of the BPA force about the joint's instantaneous center of rotation (ICR). Two vectors enter that moment: the position vector $\mathbf{d}$, drawn from the ICR to the point where the actuator force is applied, and the force vector $\mathbf{F} = \|\mathbf{F}\|\hat{\mathbf{u}}$, whose direction $\hat{\mathbf{u}}$ is the line of action of the actuator. The moment of the force about the joint is $\mathbf{M} = \mathbf{d} \times \mathbf{F}$, and the signed knee-axis torque is $M_{\mathbf{k}} = \hat{\mathbf{z}}_{\mathbf{k}} \cdot (\mathbf{d} \times \mathbf{F})$. For force acting in the plane normal to the knee axis, the moment-arm magnitude is

$$r = \|\mathbf{d} \times \hat{\mathbf{u}}\| \tag{4.1}$$

This planar moment arm is a scalar distance and can be computed from the vectors or

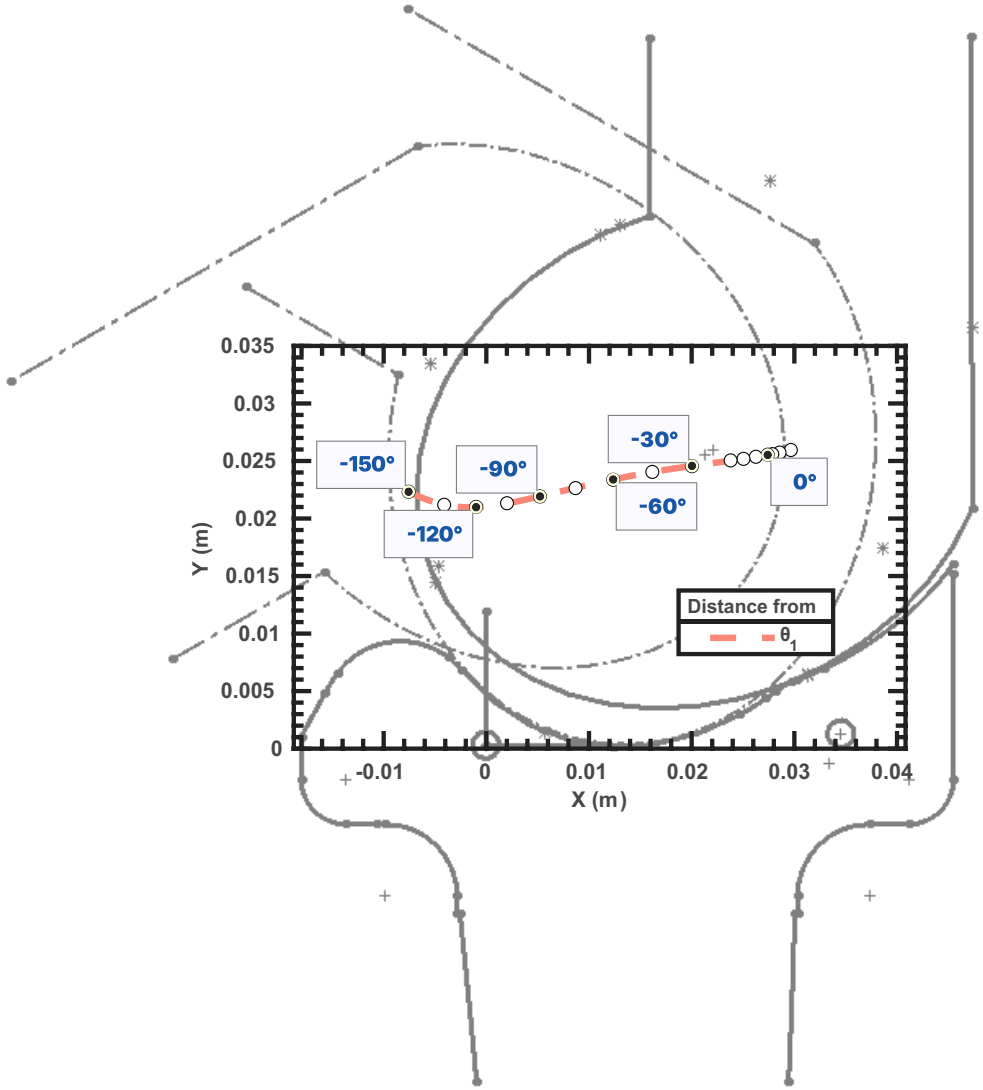


Figure 4.2: Biomimetic-knee instantaneous center of rotation relative to the tibial frame. White circles are measured points; knee outlines show 0° (solid), -60°, and -120° (dash-dot).

measured directly. For an out-of-plane force, the knee-axis moment must instead use the axis projection above; the magnitude of the full three-dimensional moment is not generally the knee torque. In this study it is calculated using the method specified by **?**. The scalar moment arm $r_{\mathbf{k}}$ about the joint axis $\hat{\mathbf{z}}_{\mathbf{k}}$ (i.e., the Z axis in the knee frame) is calculated as

$$r_{\mathbf{k}} = \mathbf{p}_{proj,xy} \cdot \frac{\mathbf{p}_f \times \hat{\mathbf{z}}_{\mathbf{k}}}{\|\mathbf{p}_f \times \hat{\mathbf{z}}_{\mathbf{k}}\|} \tag{4.2}$$

A muscle path may be modeled as a single segment running from origin to insertion, or as a chain of segments joined at via points. The segment that crosses from one reference frame into the other, here from the femur frame to the tibia frame, is the free muscle segment: because its endpoints ride on different bodies, its length and its direction respond directly to joint rotation, whereas a segment contained within one body's frame rotates rigidly with that body and changes neither. Here $\mathbf{p}_f$ is the force-path segment vector, and $\mathbf{p}_{proj,xy}$ locates a point on that line relative to the joint, projected onto the plane normal to the joint axis. The normalized cross product in Equation (4.2) is the in-plane normal to the force path. Its dot product with the projected position gives the signed moment arm of the projected force path (Figure 2.6). Multiplying this arm by the in-plane force magnitude gives the knee-axis torque; the full wrench calculation retains all three force components. Moment arm is calculated in the tibial body frame.

4.1.2 Knee Torque Data Collection and Equipment

A test jig was built to collect knee torque measurements at different knee angles over its range of motion. The test stand frame is made predominantly out of 80/20 1515 series

extruded aluminum components. The knee joint is free to rotate the tibia while the femur is fixed to the frame. A force sensor has one end connected to the tibia and the other is connected to a swing arm. The swing arm is tied to a winch with 3/16 inch Kevlar rope from Quality Nylon Rope routed over intermediate idler pulleys (see Fig. 4.1**B** and 4.1**C**). At each measured position of $\theta_{\mathbf{k}}$ the winch would be locked and the BPA would be inflated and then deflated. This test procedure kept the measured results as isometric as possible. The knee joint is co-axial with gravity to eliminate gravitational torque (see Fig. 4.1**A** and Fig. 4.1**C**).

Force data was collected using one of two different sensors. The first is a MODERN STEP 300 kg digital crane scale. The second force sensor is a CALT DYLY-103 100 kg S shaped load cell. The load cell is used in conjunction with a HX711 Load Cell Amplifier. Pressure is measured with a Freescale MPX5700 GP 5 V pressure sensor. Building air supply pressure is controlled with two pressure regulators in series. The first is a Parker model 20R113GC $0 - 120$ psi pressure regulator. The second is a Husky 3/8 inch High Performance Air Regulator HDA72200. A Festo VTUG-10-MRCR-S1T-26V20-T516LA-UL-T532S-8K valve manifold was used to supply air to the actuator. Pressure and load cell amplifier data are sent to MATLAB via an Arduino Uno style Sparkfun BlackBoard C microcontroller.

The length measurements were made using a tape measure. When there was sufficient curvature in the BPA during knee extensor torque tests, a flexible piece of string was used to mark the axial length and then transferred to the tape measure. Angle measurements were taken with a Medigauge digital electronic goniometer or with MALENOO analog goniometers of 6, 8, or 12 inch lengths.

4.1.3 Hybrid Torque Calculation Method

A Hybrid Torque calculation method was developed to help uncover inconsistencies between the model and recorded torque values. This method measures the actual muscle length and angles of the robot to calculate force and joint torque. With this method, it is easier to determine which parts of the torque equation did not achieve their predicted value and provide a correction to that part of the equation. The classical mechanics way to calculate the torque about a joint is to take the cross product of the position vector **d** and the force vector **F**, both defined in Section 4.1.1,

$$\mathbf{M} = \mathbf{d} \times \mathbf{F} \tag{4.3}$$

Knee torque is the scalar component of this moment about the joint axis, written $M_{\mathbf{k}}$ to make its axis explicit. For the planar hybrid calculation, its magnitude is the measured perpendicular arm multiplied by the actuator-force magnitude,

$$|M_{\mathbf{k}}| = r\,F \tag{4.4}$$

where r is the scalar moment arm of Eq. (4.1), here measured directly with a tape measure as the perpendicular distance from the knee axis to the actuator's line of action, and F is the scalar force magnitude, calculated by measuring the BPA's P, l_0, l_{620} and l_{m} values and using the model fit equations determined in the previous chapter. The hybrid method therefore replaces the model-predicted moment arm $r_{\mathbf{k}}$ of Eq. (4.2) and the model-predicted muscle length with their measured counterparts, leaving the force model

as the only predicted quantity. The flexor and extensor signs are then assigned using the adopted knee convention.

4.1.4 Improved Torque Model

The classical torque calculation treated links and attachments as rigid: the routed path depended on joint geometry but did not change under actuator load. Data collected during the first flexor experiments (Section 4.2) showed that this assumption does not hold for a pneumatic system: measured muscle lengths were consistently shorter than predicted, and measured torque deviated from prediction in a load-dependent manner. This section describes the corrections that were developed to account for a constant length offset, series compliance in the attachment hardware and artificial tendons, and the loss of usable muscle length when a BPA wraps around the joint.

The correction terms represent the combined effects of the attachment brackets, fixtures, artificial tendons, and test-stand loading path. They are identified from measured torque because the torque prediction requires their aggregate effect on actuator length, rather than a separate structural model of every component. The constant term χ_0 corrects the unloaded length discrepancy, χ_1 and χ_2 describe effective stiffness in the local bracket frame, and χ_3 scales the wrapping-length loss. The compliance calculation couples force and displacement: force deflects the attachments, and the resulting change in muscle length changes force. The multiobjective procedure of Section 4.1.4 evaluates candidate terms on fitting and validation tests.

Length Offset and System Stiffness

Analysis of measured muscle length data shows that it was consistently shorter than the predicted muscle length at all joint angles. This shorter length results in less muscle force, a shifted RoM, and less torque in the experimental setup. That this shorter length occurs when $\varepsilon = 1$ and $F = 0$ indicates that a constant length offset should be used. Therefore, to produce an improved torque prediction model, l_{m} was recalculated as

$$l_{\mathrm{m}} = l_{LMT} - 2\, l_{fitting} - \chi_0 \tag{4.5}$$

where l_{LMT} is the path length of the musculotendon from origin to insertion point, $l_{fitting}$ is the length of the end caps, and χ_0 is a constant length offset.

Actuator force deforms the attachments, changing effective BPA length and the predicted force. I approximate this load-dependent deformation with Hooke's law along the bracket-frame principal directions. The fitted stiffnesses χ_1 and χ_2 represent the combined attachment system, not direct measurements of individual hardware components. Because fixtures and loading hardware contribute to that compliance, beam theory alone cannot determine the fitted values.

A local bracket frame was established at the attachment bracket. The $\hat{\mathbf{x}}_{br}$ axis was aligned with the vector from the bracket reference point to the BPA attachment point and represents the stiffest direction of the bracket. The $\hat{\mathbf{y}}_{br}$ axis was chosen to point in the posterior direction with respect to the tibia frame, corresponding to the primary cantilever bending direction (see Fig. 4.3), while the $\hat{\mathbf{z}}_{br}$ axis completes the right-handed frame. In this frame the bracket is modeled as three independent linear springs along the principal

axes, with stiffnesses k_x, k_y, and k_z in the x_{br}, y_{br}, and z_{br} directions.

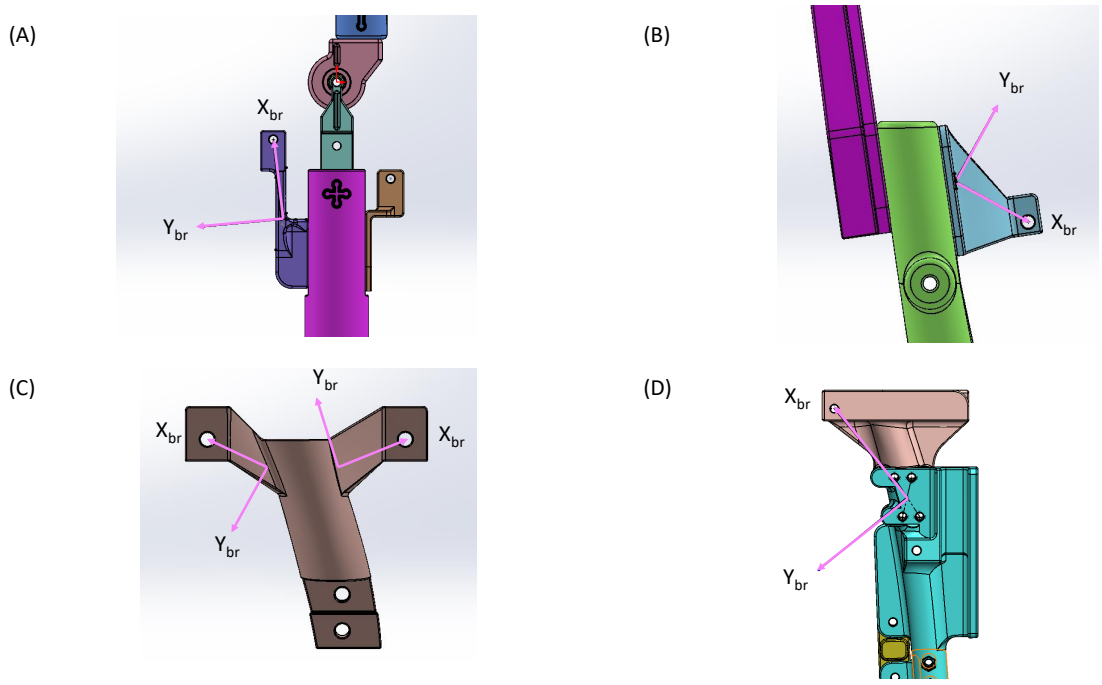


Figure 4.3: Bracket frame orientations for: **(A)** flexor configuration on the pinned knee, **(B)** extensor configuration on the pinned knee, **(C)** flexor and extensor configurations on the biomimetic knee with ∅10 mm BPAs, and **(D)** flexor configuration on the biomimetic knee with the ∅20 mm BPA.

For the two-rotation evaluator, the bracket frame was constructed from the vector joining the bracket origin to attachment point $\mathbf{p}_A$ or $\mathbf{p}_B$. The first rotation, through θ_z about the fixed space-frame axis $\hat{\mathbf{z}}_s$, points the intermediate axis $\hat{\mathbf{x}}_1$ toward the sagittal-plane projection, $\mathbf{p}_A^{\text{sag}}$ for the origin point $\mathbf{p}_A$ and $\mathbf{p}_B^{\text{sag}}$ for the insertion point $\mathbf{p}_B$. The second rotation, through $-\theta_y$ about the current axis $\hat{\mathbf{y}}_1 = \hat{\mathbf{y}}_{br}$, brings the final bracket axis $\hat{\mathbf{x}}_{br}$ into alignment with the full three-dimensional point vector. Figure 4.4 shows these geometric targets. The resulting orientation and homogeneous transform are

$$\mathbf{R}_{sb} = \mathbf{R}_z(\theta_z)\mathbf{R}_y(\theta_y)^T. \tag{4.6}$$

$$\mathbf{T}_{sb} = \begin{bmatrix} \mathbf{R}_{sb} & \mathbf{p}_{br} \\ \mathbf{0}^T & 1 \end{bmatrix}. \tag{4.7}$$

$$\mathbf{K}_{br} = \mathrm{diag}\big(k_x,\ k_y,\ k_z\big) \tag{4.8}$$

The fitted values χ_1 and χ_2 define the principal stiffnesses, with configuration-specific axis assignments in Table C9. The adopted pinned-flexor fit uses two brackets with two-rotation frames. For the insertion bracket, $k_x = k_z = \chi_1$ and $k_y = \chi_2$; this symmetric assignment makes rotation about $\hat{\mathbf{y}}_{br}$ immaterial to its compliance. The origin bracket uses $[\chi_1, \chi_1, \chi_2]$, for which that invariance does not hold. Appendix C records the remaining assignments. For the insertion bracket,

$$\mathbf{K}_{br} = \begin{bmatrix} \chi_1 & 0 & 0 \\ 0 & \chi_2 & 0 \\ 0 & 0 & \chi_1 \end{bmatrix} \tag{4.9}$$

The corresponding local compliance matrix is

$$\mathbf{C}_{br} = \mathbf{K}_{br}^{-1} \tag{4.10}$$

The force applied to the bracket was expressed in the bracket frame as $\mathbf{F}_{br}$. This force vector can be written in terms of its magnitude and unit direction as

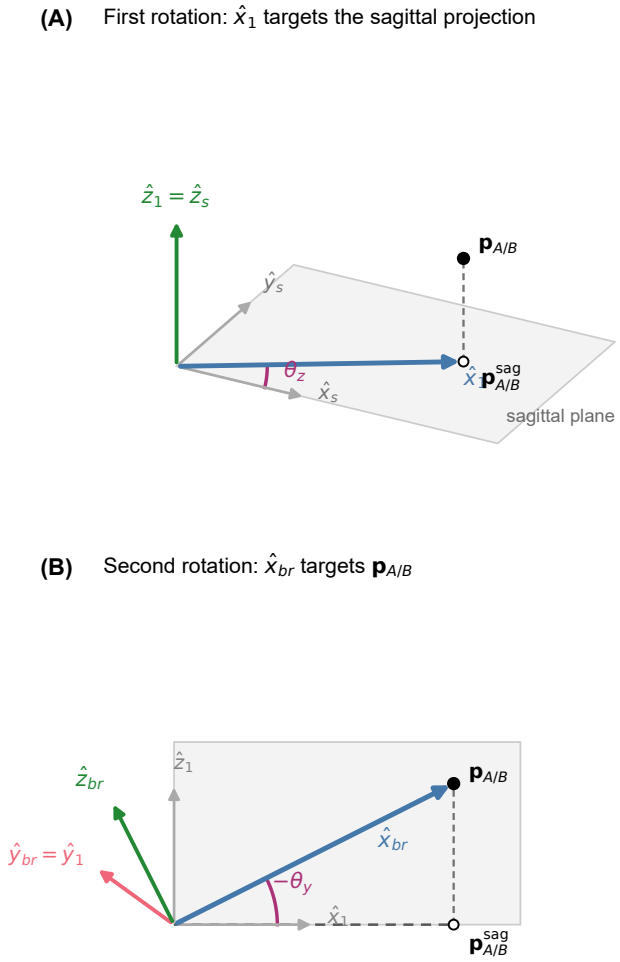


Figure 4.4: Bracket-frame construction from attachment point $\mathbf{p}_A$ or $\mathbf{p}_B$: yaw about fixed $\hat{\mathbf{z}}_s$ aligns the sagittal projection (A); pitch about current $\hat{\mathbf{y}}_{br}$ aligns the full point vector (B).

$$\mathbf{F}_{br} = F\,\hat{\mathbf{u}} \tag{4.11}$$

where $F = \|\mathbf{F}\|$ is the scalar force magnitude of Eq. (4.4) and $\hat{\mathbf{u}} = [u_x, u_y, u_z]^T$ is the BPA force direction expressed in the bracket frame. The bracket deflection vector, $\boldsymbol{\epsilon}_{br}$, was calculated from Hooke's law as

$$\mathbf{F}_{br} = \mathbf{K}_{br}\,\boldsymbol{\epsilon}_{br} \tag{4.12}$$

Solving Eq. 4.12 for the bracket deflection gives

$$\boldsymbol{\epsilon}_{br} = \mathbf{K}_{br}^{-1}\,\mathbf{F}_{br} = \mathbf{C}_{br}\,\mathbf{F}_{br} \tag{4.13}$$

Substituting Eq. 4.11 into Eq. 4.13 gives

$$\boldsymbol{\epsilon}_{br} = F\,\mathbf{K}_{br}^{-1}\,\hat{\mathbf{u}} \tag{4.14}$$

For the diagonal stiffness matrix in Eq. 4.9, Eq. 4.13 is equivalent to the component-wise expression

$$\boldsymbol{\epsilon}_{br} = \begin{bmatrix} F\,u_x/\chi_1 \\ F\,u_y/\chi_2 \\ F\,u_z/\chi_1 \end{bmatrix} \tag{4.15}$$

Although $\boldsymbol{\epsilon}_{br}$ is a three-dimensional deflection vector, the BPA force model depends on the scalar change in effective BPA length (δ). Therefore, only the component of $\boldsymbol{\epsilon}_{br}$

along the BPA line of action directly changes the effective BPA length. This scalar bracket deflection is the projection of $\boldsymbol{\epsilon}_{br}$ onto $\hat{\mathbf{u}}$:

$$\delta_{br} = \hat{\mathbf{u}}^T \boldsymbol{\epsilon}_{br} \tag{4.16}$$

Substituting Eq. 4.13 into Eq. 4.16 gives

$$\delta_{br} = \hat{\mathbf{u}}^T \mathbf{K}_{br}^{-1} \mathbf{F}_{br} \tag{4.17}$$

Using Eq. 4.11, Eq. 4.17 becomes

$$\delta_{br} = F\, \hat{\mathbf{u}}^T \mathbf{K}_{br}^{-1} \hat{\mathbf{u}} \tag{4.18}$$

With $\mathbf{C}_{br} = \mathbf{K}_{br}^{-1}$, the scalar compliance along the BPA line of action is

$$c_{\hat{u},br} = \frac{\delta_{br}}{F} = \hat{\mathbf{u}}^T \mathbf{C}_{br} \hat{\mathbf{u}} \tag{4.19}$$

This projected-compliance form is used because the force is applied along $\hat{\mathbf{u}}$, while the bracket is allowed to deflect independently along its local $\hat{\mathbf{x}}_{br}$, $\hat{\mathbf{y}}_{br}$, and $\hat{\mathbf{z}}_{br}$ directions. The reduced-order Hooke's law expression for the bracket deflection along the BPA line of action is therefore

$$F = k_{eq}\, \delta \tag{4.20}$$

Here k_{eq} is the equivalent stiffness of every compliant element in the actuator path, and the bracket-only case is the special case with no tendon. When the BPA attaches directly

to the bracket, k_{eq} excludes the tendon and equals $c_{\hat{u},br}^{-1}$, where $c_{\hat{u},br} = \hat{\mathbf{u}}^T \mathbf{K}_{br}^{-1} \hat{\mathbf{u}}$. When an artificial tendon was used, it was modeled as an axial spring with stiffness k_{ten}: bracket deflection and tendon stretch both shorten the effective BPA length along the actuator line of action, the two elements carry the same force, and their scalar compliances add in series.

The projected bracket compliance is $c_{\hat{u},br} = k_{\hat{u},br}^{-1}$, tendon compliance is $c_{ten} = k_{ten}^{-1}$, and the equivalent stiffness is $k_{eq} = c_{eq}^{-1}$.

$$c_{eq} = c_{\hat{u},br} + c_{ten} \tag{4.21}$$

The series model assumes quasi-static tensile loading, linear compliance, additive axial displacements, and a tendon without preload. Rotational compliance and force–moment coupling are neglected. Within each force-balance calculation the line of action is fixed; the attachment positions, force direction, and moment arm are then recalculated at the deflected geometry. Antagonistic loading that lengthens the actuator lies outside this treatment.

Since δ is the scalar compliance displacement along the BPA line of action, then the compliance-corrected active BPA length used in the force balance is

$$l_m(\delta) = l_m(0) - \delta \tag{4.22}$$

where $l_m(0)$ is the initial muscle length calculated from (4.5).

For a specified set of correction parameters, the following calculation solves the attachment-compliance displacement at each knee angle; it is distinct from the optimization that identifies those parameters from experimental torque data (Section 4.1.4). The

corresponding relative contraction is $\varepsilon^*(\delta)$. Since the BPA force depends on contraction, and contraction depends on the compliance displacement, δ is the root of the nonlinear force-balance equation

$$F_{\max} \cdot F_{\mathrm{BPA}}(\varepsilon^*(\delta), P^*) - k_{eq}\,\delta = 0 \tag{4.23}$$

MATLAB `fzero` solves Equation (4.23) between zero displacement and the maximum usable contraction, using its default tolerance. If the actuator is already at that contraction limit or the residual has no sign change, the evaluator returns $\delta = 0$. Solver tolerance is distinct from physical-model accuracy.

After solving Equation (4.23), I used the force to calculate vector bracket displacement, transformed the displaced attachment into the limb frame, and recalculated routed length, force direction, moment arm, and torque. The fitted χ_1 and χ_2 represent effective attachment-system stiffnesses, including fixtures and loading hardware, rather than isolated bracket properties.

The actuator path can carry compliant hardware at both ends: the insertion-side bracket of the preceding development together with a second, origin-side bracket. Figure 4.5 distinguishes the three-dimensional displacement of each bracket from the scalar displacement along the BPA force path. For bracket j, the displacement resolves into components $\eta_{1,j}$, $\eta_{2,j}$, and $\eta_{3,j}$ along the local bracket-frame basis, whereas only its projection $\delta_{br,j}$ along the corresponding force direction $\hat{\mathbf{u}}_j$ changes the actuator length. The two bracket projections act in series with the tendon elongation δ_{tendon}:

$$
\begin{aligned}
\boldsymbol{\eta}_{br,j} &= \sum_{i=1}^{3} \eta_{i,j} \hat{\mathbf{e}}_i^{(j)}, \\
\delta_{br,j} &= \hat{\mathbf{u}}_j^T \boldsymbol{\eta}_{br,j}, \qquad j \in \{1,2\}, \\
\delta_{\text{series}} &= \delta_{br,1} + \delta_{\text{tendon}} + \delta_{br,2}.
\end{aligned}
\tag{4.24}
$$

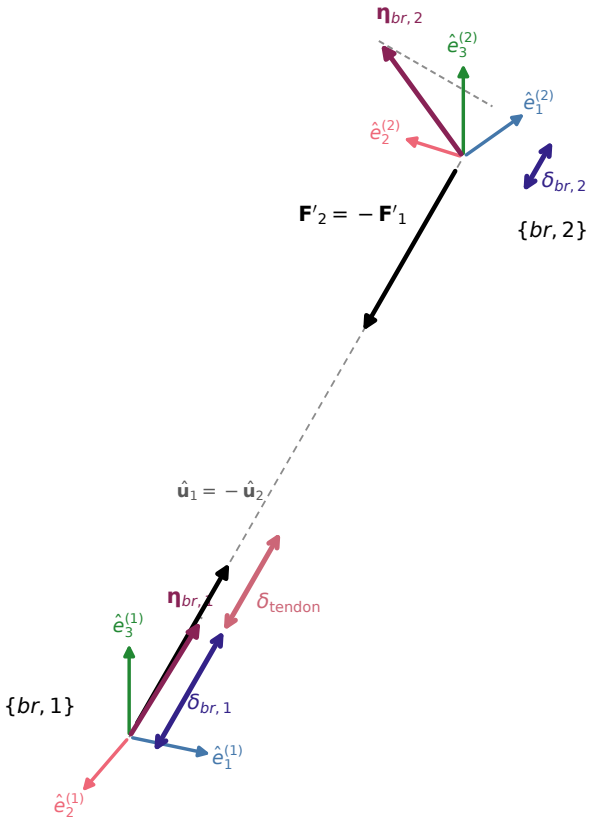


Figure 4.5: Projection of bracket displacement onto the actuator force path. Each bracket has a local basis $\hat{\mathbf{e}}_1^{(j)}$, $\hat{\mathbf{e}}_2^{(j)}$, and $\hat{\mathbf{e}}_3^{(j)}$. The displacement $\boldsymbol{\eta}_{br,j}$ has an axial projection $\delta_{br,j}$ along $\hat{\mathbf{u}}_j$; the two bracket projections and tendon elongation contribute to the change in actuator length. Displacements are exaggerated.

BPA Fictive Contraction

As the knee flexes, the extensor BPA contacts the anterior joint surface and bends around it (Figure 4.6). The rigid routed-length calculation does not account for the resulting change in usable contraction. I represent this effect with fictive strain: an empirical increase in the contraction seen by the force model, beyond the actuator's geometric contraction. The correction is phenomenological; it does not resolve wall stresses or the deformation of the braided actuator. A wrapping-loss term χ_3 scales the associated reduction in usable length:

$$\Delta l = \begin{cases} \chi_3 \; R \; |\theta_{\text{wrap}} - \theta_{\mathbf{k}}| \; \zeta^2 & \theta_{\mathbf{k}} \leqslant \theta_{\text{wrap}} \\ 0 & \theta_{\mathbf{k}} \geqslant \theta_{\text{wrap}} \end{cases} \tag{4.25}$$

where Δl is the change in usable length of the BPA, R is the wrapping radius, $|\theta_{\text{wrap}} - \theta_{\mathbf{k}}|$ is the angle through which the BPA is bent once contact begins, and θ_{wrap} is the knee flexion angle at which the BPA starts to bend over the knee joint, determined from solid modeling. ζ is the additive complement to relative strain, such that

$$\zeta = 1 - \varepsilon^* \tag{4.26}$$

so the fitted wrapping loss grows quadratically toward resting length and vanishes at the normalized free-contraction limit. Combining (4.26) into (4.25), and subtracting Δl from (4.5) yields

$$l_{\text{m}} = l_{LMT} - 2 \, l_{fitting} - \chi_0 - \Delta l \tag{4.27}$$

Figure 4.6 shows a sagittal schematic of an extensor BPA crossing a flexed robot knee, with its braided body following an anterior contact arc between the proximal attachment and distal tendon. Dashed radii bound the contact arc; R is the BPA-path wrapping radius, and $\alpha = |\theta_{\text{wrap}} - \theta_{\mathbf{k}}|$ is the bend angle after contact begins, expressed in radians in Equation (4.25). The drawing illustrates the local circular-wrap approximation, not the full multi-contact route. The fitted usable-length loss Δl is not an arc-minus-chord length.

Identification of the Torque-Model Correction Terms

I set the correction terms in two stages using `gamultiobj` in MATLAB. The flexor stage fits (χ_0, χ_1, χ_2) with compliance at both attachments. The selected training/validation split trains on the $l_0 = 48.5$ cm and 41.0 cm tests and validates on the 45.7 cm and 40.6 cm tendon tests. The extensor stage locks (χ_1, χ_2) to that same flexor-front member and fits only (χ_0, χ_3). Its selected training/validation split trains on the 43.6 cm, 43.2 cm-tendon, and 45.7 cm tests and validates on the 40.0 cm plain, 40.0 cm-tendon, and 46.5 cm tests. The extensor length offset is allowed to be negative. Appendix C records the test constants, bounds, stiffness-axis assignments, and adopted values in Tables C1, C4, C10, C9, and 4.9.

Each candidate is evaluated using RMSE, FVU, and maximum absolute residual. These metrics jointly measure torque error; none isolates a unique source such as level offset, waveform shape, or angular misalignment.

The objective vector of the identification is defined explicitly. For a candidate correction vector $\boldsymbol{\chi}$, (χ_0, χ_1, χ_2) in the flexor stage and (χ_0, χ_3) on the locked stiffness pair in the extensor refit, the predicted torque curve of test i is interpolated to its n_i measured knee

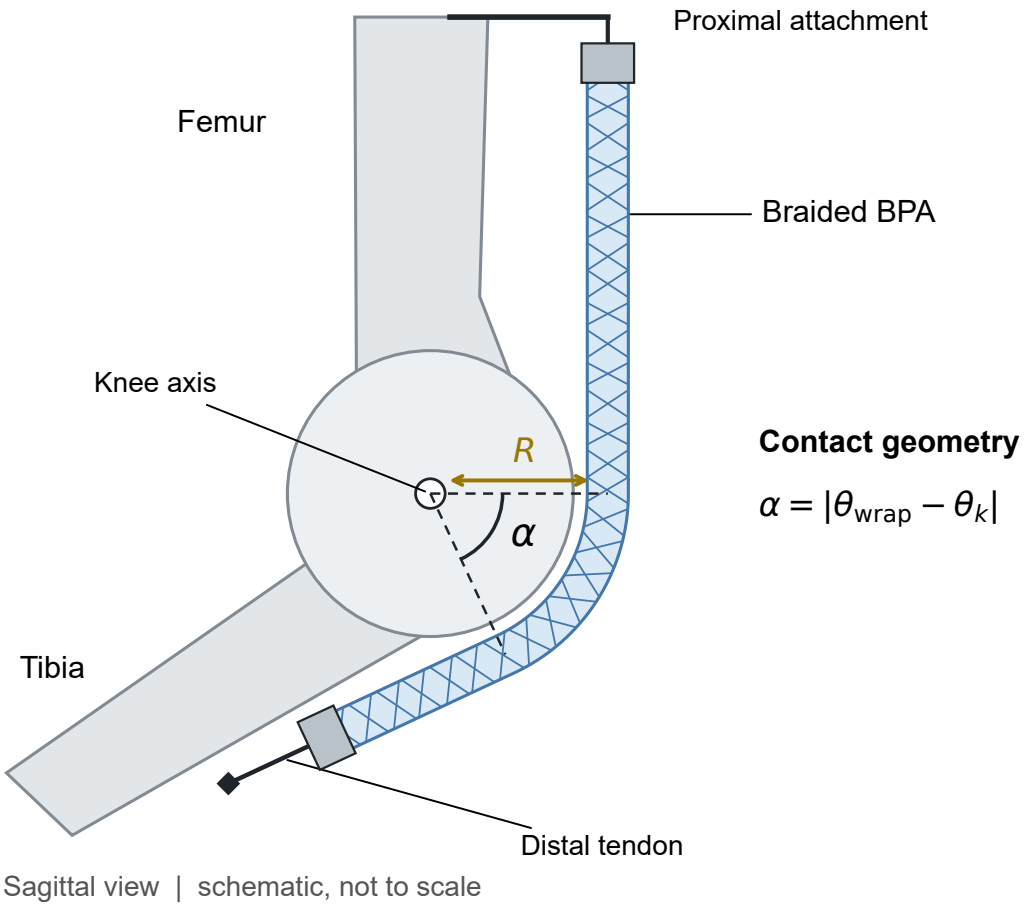


Figure 4.6: Extensor BPA wrapping around a flexed knee. Dashed radii bound the local contact arc, with wrapping radius R and bend angle α.

angles θ_{ij}, giving residuals $e_{ij} = M_{ij} - \hat{M}(\theta_{ij}; \boldsymbol{\chi})$ against the measured torque M_{ij}, and the three goodness-of-fit components of that test are

$$\mathrm{RMSE}_i = \sqrt{\frac{1}{n_i} \sum_j e_{ij}^2} \tag{4.28}$$

$$\mathrm{FVU}_i = \frac{n_i - 1}{n_i} \cdot \frac{\sum_j e_{ij}^2}{\sum_j (M_{ij} - \bar{M}_i)^2} \tag{4.29}$$

$$e_i^{\max} = \max_j |e_{ij}|. \tag{4.30}$$

RMSE summarizes the overall error magnitude. FVU expresses the residual sum of squares relative to the sum of squared deviations of the measured torque from its mean $\bar{M}_i$, with the factor $(n_i - 1)/n_i$ carrying the small-sample adjustment of the evaluators. The maximum absolute residual records the largest pointwise error. RMSE and the maximum residual have torque units, whereas FVU is dimensionless. For a given test, FVU is proportional to squared RMSE, so the two metrics are not independent measures of shape and level. Across tests, their different scaling changes the relative weighting of the errors. The solver minimizes the objective vector

$$\mathbf{g}(\boldsymbol{\chi}) = \frac{1}{|\mathcal{T}|} \sum_{i \in \mathcal{T}} \begin{bmatrix} g_{1,i}(\boldsymbol{\chi}) / g_{1,i}^{(0)} \\ g_{2,i}(\boldsymbol{\chi}) / g_{2,i}^{(0)} \\ g_{3,i}(\boldsymbol{\chi}) / g_{3,i}^{(0)} \end{bmatrix}, \tag{4.31}$$

where $g_{1,i}$, $g_{2,i}$, and $g_{3,i}$ are the RMSE, FVU, and maximum absolute residual of Equations (4.28)–(4.30), the superscript (0) marks the zero-correction baseline value of the same

component, the same kinematics and force model with every correction term disabled, and $\mathcal{T}$ is the set of tests used for training. The mean of the baseline-normalized components makes the three objectives dimensionless and the training tests commensurate. The search carried no constraints beyond the parameter bounds of each stage, and the returned Pareto fronts were filtered to members scoring at or below the zero-correction baseline on every component of every test in that stage's identification pool.

MATLAB `gamultiobj` uses NSGA-II to return a Pareto front: candidate parameter sets for which improving one objective requires worsening another [**??**]. Candidate predictions are checked on the training and validation tests. The search uses bounded, normalized parameters, with logarithmic stiffness coordinates; Appendix C gives the bounds and adopted values. Population, mutation, crossover, and stopping settings are retained with the archived run records.

4.1.5 Actuator Placement and Route Redesign

The placement search varies actuator attachments to control moment arm and muscle excursion over the joint range of motion [**??**]. For the erector spinae, internal oblique, and external oblique at the lumbosacral joint, the targets are Gait2392 torque surfaces sampled at 100 positions along each of two joint coordinates.

At each search scale, offsets of $\pm\epsilon$ along the Cartesian coordinates define eight candidates per attachment and 64 candidate pairs. The lowest-cost pair becomes the next search center. When no further improvement is found, ϵ is halved. The search uses an initial $\epsilon = 0.1\,\mathrm{m}$ and stops below $1\,\mathrm{mm}$ [**?**].

The original cost combined the summed absolute torque mismatch over the sampled poses, actuator length, displacement of the attachments from their starting locations, and a diameter-dependent penalty. The published weights for the first three terms were 10^{-4}, 10^{-2}, and 1, respectively; the penalties for 10 mm, 20 mm and 40 mm BPAs were 1, 5, and 15. These were numerical weights chosen to balance differently scaled cost terms. The geometric terms discouraged impractical placements but did not impose collision avoidance: attachments could lie away from the modeled bones, and routes could intersect them. This study established the placement-search approach; it did not establish physical torque accuracy or satisfaction of the human torque requirement at every pose.

A subsequent knee implementation restricted the initial candidates to points on the femur and tibia meshes. For the biceps femoris short-head replacement, every candidate origin–insertion pair was evaluated over 100 knee poses from approximately -120° to 10°, including neutral. The BPA calculation used a fixed 20 mm diameter, determined actuator and tendon lengths from the route geometry, and rejected candidates outside its allowable contraction range. It retained the lowest-cost usable pair. This discrete mesh search was distinct from the earlier cube-perturbation search. Both used rigid attachment geometry without the experimentally identified compliance corrections.

The knee-flexor mesh search also used a different objective from the original torque-mismatch cost. It rewarded torque surplus in the human torque's direction while penalizing disagreement between the human and robot torque-vector directions:

$$C_{\text{mesh}} = -\frac{1}{3N}\sum_{j=1}^{N}\left[w_{\mathbf{k}}\,D_{\mathbf{k},j} + \sum_{a\in\{x,y\}} w_a\,D_{a,j}\right] + \frac{w_u}{N}\sum_{j=1}^{N}\left(1 - \hat{\mathbf{M}}_j^{\mathrm{H}}\cdot\hat{\mathbf{M}}_j^{\mathrm{R}}\right). \tag{4.32}$$

Here N is the number of evaluated poses, and $D_{a,j} = M^R_{a,j} - M^H_{a,j}$ when $M^H_{a,j} \geqslant 0$, with the subtraction reversed otherwise. The implementation assigns weights $w_{\mathbf{k}} = 5$ to the knee-axis component and $w_x = w_y = 0.25$ to the two off-axis components. It averages these weighted component surpluses over $3N$ entries. The direction term averages over N torque vectors, with weight $w_u = 10$; it is zero for aligned vectors and contributes 20 for opposite directions. A zero robot-torque vector is assigned the opposite human direction in this term. The equation assumes a nonzero human torque vector at the evaluated pose. Torques are supplied in newton meters; this is a numerically weighted score, rather than a sum of normalized physical errors. Minimizing it favors torque capacity in the required direction, but an average surplus does not guarantee that the requirement is met at every pose. Table 4.1 summarizes this knee-specific objective.

Table 4.1: Terms of the knee-flexor mesh-placement cost. The cost is minimized; a surplus reduces it, while a direction mismatch increases it.

Term	**Weight**	**Meaning**
Knee torque surplus	5	Signed surplus $D_{\mathbf{k},j}$, summed with the off-axis terms and divided by $3N$; enters with a negative sign.
Off-axis torque surplus	0.25 each	Signed surplus $D_{x,j}$ and $D_{y,j}$, using the human torque's sign on each axis; enters with a negative sign.
Direction mismatch	10	Mean $1 - \hat{\mathbf{M}}^H_j \cdot \hat{\mathbf{M}}^R_j$; zero for aligned vectors and two for opposing vectors.

Constrained Knee-Route Redesign

The dissertation's route redesign extends the placement search as a constrained continuous optimization over the design vector $\mathbf{x} = [\mathbf{p}_1,\ \mathbf{p}_{\text{end}},\ l_0,\ l_{\text{tendon}}]$, the BPA origin $\mathbf{p}_1$, insertion $\mathbf{p}_{\text{end}}$, resting length, and tendon length (Appendix C), so that the flexor and extensor configurations meet or exceed human knee torque capability over the range of motion. Because the route prediction (contact activation, wrapping, and compliance correction) is expensive to evaluate and only moderately smooth, the redesign minimizes a single scalar cost assembled from weighted penalty terms. The cost has the form

$$J = w_w\, J_{\text{worst}} + w_t\, J_{\text{torque}} + \textstyle\sum_a J_a + J_{\text{geom}} + J_{\text{len}}, \tag{4.33}$$

where J_{worst} is the largest fraction by which robot torque falls short of the human target at any angle across the range of motion, weighted heavily ($w_w = 10^5$); J_{torque} is the mean-squared shortfall fraction ($w_t = 10^3$), weighted two orders of magnitude below the worst-angle term so that no shortfall at a single angle can be traded away for a better average; the J_a are configuration-specific auxiliary penalties; and J_{geom} and J_{len} are regularizers that keep the endpoints and actuator lengths near their seed values. The two leading terms are built from the normalized shortfall

$$s_j = \frac{\max\Big(0,\ (1+m_r)|M_j^{\text{H}}| - |M_j^{\text{R}}|\Big)}{\tau}, \qquad J_{\text{worst}} = \max_j s_j, \qquad J_{\text{torque}} = \frac{1}{N}\sum_j s_j^2, \tag{4.34}$$

where $m_r = 0.05$ is the required torque margin and $\tau = \max(1, \max_j |M_j^{\mathrm{H}}|)$ is the human torque scale. The flexor and extensor configurations carry the remaining terms of Equation (4.33) as

$$J_{\mathrm{flx}} = 10^5 J_{\mathrm{worst}} + 10^3 J_{\mathrm{torque}} + 10^{-2} J_{\mathrm{shape}} + J_{\mathrm{off}} + J_{\mathrm{rest}} + J_{\mathrm{strain,hi}} + J_{\mathrm{strain,lo}} + J_{\mathrm{wrap}} + J_{\mathrm{geom}} + J_{\mathrm{len}} + J_{\mathrm{anchor}}, \tag{4.35}$$

$$J_{\mathrm{ext}} = 10^5 J_{\mathrm{worst}} + 10^3 J_{\mathrm{torque}} + J_{\mathrm{over}} + J_{\mathrm{rest}} + J_{\mathrm{strain,hi}} + J_{\mathrm{strain,lo}} + J_{\mathrm{contr}} + J_{\mathrm{geom}} + J_{\mathrm{len}}. \tag{4.36}$$

Tables 4.2–4.4 define the normalized penalties and their weights. A candidate is rejected if its route cannot be constructed or its torque or strain calculation is not finite.

Human torque targets are computed from the Gait2392 lower-extremity model: biceps femoris short head for the flexor and vastus medialis for the extensor [**??**]. At each knee angle, the OpenSim moment arm is multiplied by the muscle's maximum isometric force [**???**]. Appendix A gives the model parameters and exported values. The search also enforces bone clearance, feasible wrapping, actuator contraction limits, and tendon geometry; a candidate that violates these requirements cannot be accepted.

Both configurations use a global search followed by local refinement. MATLAB's `surrogateopt` approximates the expensive route calculation from sampled designs and uses that approximation to choose further candidates [**?**]. The evaluation budgets are 3,000 for the direct extensor and 1,000 for the flexor. The best feasible design is refined with `patternsearch`, which compares nearby candidates without requiring derivatives, for up to 15,000 further evaluations [**??**]. The cam variant uses 1,000 surrogate evaluations and

Table 4.2: Torque penalties in the constrained knee-route search. FLEX and EXT identify configuration-specific terms.

Term	**Weight**	**Meaning**
J_{worst} worst-angle shortfall	10^5	Largest fraction by which predicted torque magnitude falls short of the required envelope (human curve plus margin m_r) at any angle; dominates so that no single-angle shortfall can be traded for a better average.
J_{torque} mean-square shortfall	10^3	Mean squared shortfall fraction across all angles.
J_{shape} shape agreement (FLEX)	10^{-2}	Small mean-square tracking error of the full torque curve beyond the requirement.
J_{over} overshoot (EXT)	10^{-3}	Mean-square fraction by which torque exceeds the plain human curve, discouraging over-designed moment arms.
J_{off} off-axis excess (FLEX)	10^3	Worst and mean squared excess of the off-axis (x, y) torques beyond the original single-BPA route's off-axis torques, each axis normalized separately.

Table 4.3: Feasibility and geometric penalties in the constrained knee-route search.

Term	Weight	Meaning
J_{rest} rest-length feasibility	10^6	Squared violation of resting length plus tendon plus fittings covering the longest path over the range of motion.
$J_{\text{strain,hi}}$ upper strain limit	10^7 FLEX, 10^4 EXT	Squared violation of the maximum admissible relative strain and contraction.
$J_{\text{strain,lo}}$ stretch limit	10^7 FLEX, 10^4 EXT	Squared violation of the minimum strain, forbidding a stretched (slack-free) actuator.
J_{contr} contraction cap (EXT)	10^6	Squared violation of the measured maximum usable contraction; also enforced as a hard constraint.
J_{wrap} wrap release (FLEX)	10^3	Penalty when the moving wrap contact fails to release during the extension sweep, scaled by the remaining turn of the wrap.
J_{geom} endpoint regularizer	10^{-2}	Keeps origin and insertion near their seed values; the flexor scales each axis separately, allowing large proximal motion of the origin.
J_{len} length regularizer	10^{-3}	Keeps the resting length near its seed value (40 mm norm).
J_{anchor} second-BPA anchor (FLEX)	10	Keeps the second actuator's derived endpoints near the build line of the loaded design (10 mm norm); a score term only, never a constraint.

Table 4.4: Additional cam-profile penalties and rejection values.

Term	Weight	Meaning
J_{arm} moment-arm reward (CAM)	-10^2, capped	Negative reward on the worst-angle margin against the required envelope, capped at $+100\%$ of the torque scale, so the solver keeps buying moment arm instead of stopping at the requirement.
J_{armStep} arm-step penalty (CAM)	10^3	Penalizes adjacent-pose moment-arm steps beyond the accepted release discontinuity of 2.2 mm.
J_{reg} cam regularizer (CAM)	10^{-2}	Tie-break toward the first-cut cam neighborhood (25.4 mm norm).
Infeasible sentinels	10^{11}, 10^{12}	Non-finite predictions, or a route the predictor cannot build, are rejected outright.

1,500 refinement evaluations. The budgets limit computation; local refinement can also stop when its search-step size falls below the specified tolerance and the constraints are satisfied.

Although both configurations share this pipeline, their cost functions differ in ways that reflect the mechanics. The flexor BPA routes around a single moving wrap point, so clearance from the tibia and femoral condyles is enforced by continuous signed-distance constraints, and its cost adds an off-axis torque penalty (torques about the joint axes other than the knee axis, relative to the original single-BPA route) and a penalty if the wrap fails to release cleanly as the knee extends. The extensor BPA wraps over the front of the joint through as many as seven sequential contacts, so clearance is enforced instead by the contact-elimination rule that governs the route, and its cost penalizes torque overshoot above the human curve and excessive actuator contraction, which risks bottoming out at high flexion. The flexor and extensor targets include a 5 % torque margin as a design factor for additional effects not represented by the model. Thus, the required prediction is 1.05 times the human target at every design angle. This prescribed margin is a design allowance, rather than an experimentally established error bound. With two BPAs per joint, the flexor route is solved as two separately evaluated paths, while the extensor model scales the force of a single shared route by the actuator count.

The redesign retains the objective of obtaining adequate torque from anatomically constrained routes, but changes both the prediction and the search formulation. Actuator and tendon lengths become independent design variables, geometry is constrained continuously, and pointwise torque shortfall is penalized using the corrected force model. Section 4.3.2 discusses the consequences for hardware design.

4.2 Results

The pinned-flexor results first compare the Hunt actuator model with the improved BPA force model. Subsequent comparisons use the improved model before and after design-parameter optimization. Hybrid calculations substitute measured muscle length and moment arm to help distinguish force-model error from geometric error on the pinned knee. The biomimetic-extensor hybrid curve is an approximate diagnostic because the curved actuator and its distance from the ICR are difficult to measure within the femoral condyles.

4.2.1 Flexor on Pin Joint

The pre-optimization comparisons distinguish force-model error from errors in the predicted muscle geometry (Figure 4.7). The hybrid calculation uses measured length and moment arm, whereas the other predictions use the routed geometry. Flexion angles and torques are negative; the caption identifies the individual tests.

Both pre-optimization force models overpredicted measured flexor-torque magnitude. On the common domain where both predictions are defined, the improved BPA characterization gives the smaller errors (Table 4.5). Those scores use a different domain from the full-test corrected-model scores in Table 4.6. The hybrid calculation, using measured length and moment arm, follows the measurements more closely, identifying the predicted geometry as an important source of torque error.

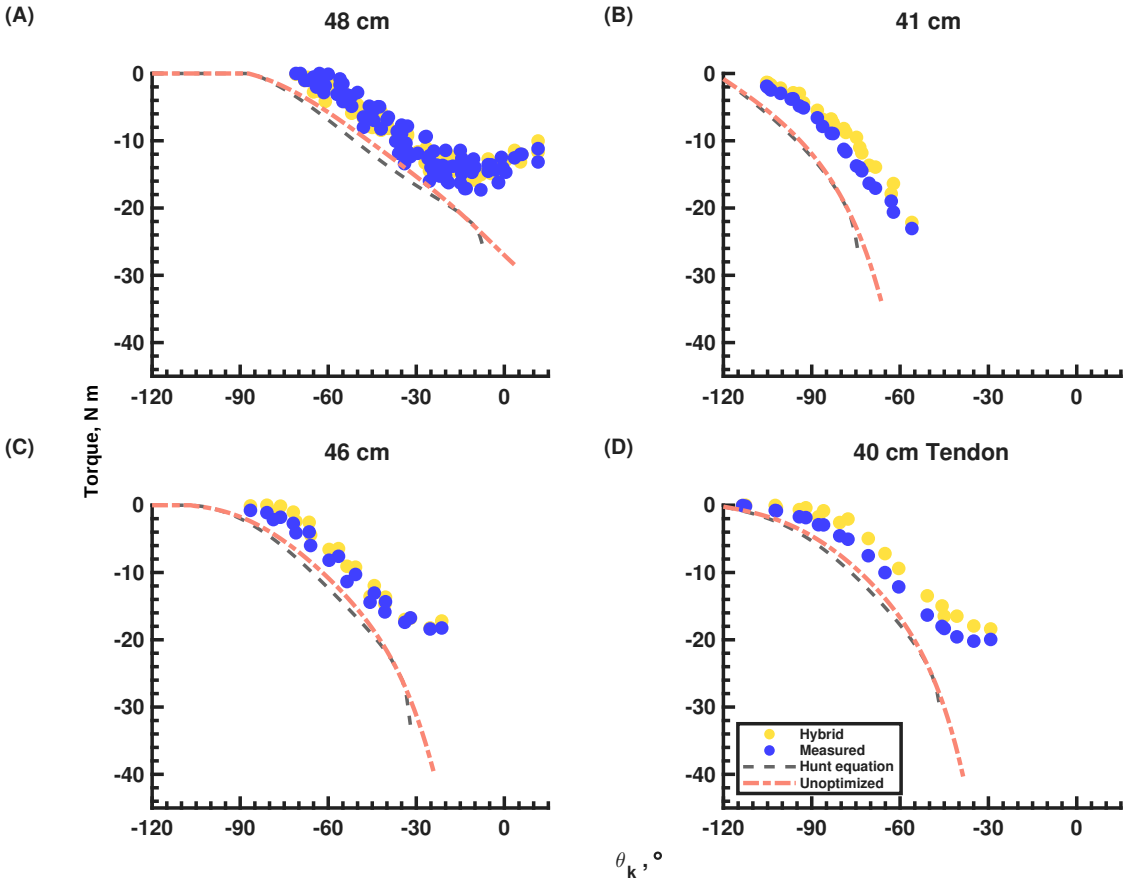


Figure 4.7: Pinned-knee flexor torque before design-parameter optimization: **(A)** $l_0 = 48.5\,\text{cm}$, **(B)** $l_0 = 41.0\,\text{cm}$, **(C)** $l_0 = 45.7\,\text{cm}$, and **(D)** $l_0 = 40.6\,\text{cm}$ with a 35 mm tendon. Panels **(A)** and **(B)** are training tests; **(C)** and **(D)** are validation tests. The legend distinguishes measurements, the Hunt model, the improved BPA model, and the hybrid calculation.

Table 4.5: Lengths l_0 (resting BPA) and l_{ten} (artificial tendon) are both in cm. Pre-optimization pinned-knee flexor torque-prediction goodness of fit per test, comparing the three pre-optimization torque predictions of Figure 4.7: the Hunt-equation rigid-body prediction, the unoptimized prediction of the improved BPA characterization, and the back-calculated (hybrid) torque. RMSE and maximum absolute residual (Max.) are in N m, and FVU denotes the fraction of variance unexplained.

		Hunt equation			**Improved BPA**			**Back-calculated (hybrid)**		
l_0 (cm)	l_{ten} (cm)	**RMSE**	**FVU**	**Max.**	**RMSE**	**FVU**	**Max.**	**RMSE**	**FVU**	**Max.**
48.5	–	5.779	1.1761	9.798	4.871	0.8354	9.865	1.635	0.0991	3.842
41.0	–	6.633	3.2313	11.409	6.074	2.7090	9.262	2.187	0.1184	4.243
45.7	–	5.315	0.9408	9.991	4.465	0.6640	9.536	1.246	0.0409	2.275
40.6	3.5	3.945	0.6750	7.527	3.226	0.4512	7.376	2.085	0.0755	3.073

The flexor adopted identification is the two-bracket fit of Section 4.1.4 on the four cross-validation tests: the $l_0 = 48.5$ cm and 41.0 cm tests in training and the $l_0 = 45.7$ cm and 40.6 cm-tendon tests held out. Its adopted terms are a length offset $\chi_0 = 3.94$ mm and effective stiffnesses $\chi_1 = 39.98\,\text{kN m}^{-1}$ and $\chi_2 = 14.73\,\text{kN m}^{-1}$ (Table 4.9). At these terms the optimized model follows the measured torque across the range of motion, including the loss of torque near full extension that the rigid-body prediction does not produce (Figure 4.8). The training-set RMSE is 2.165 N m for the $l_0 = 48.5$ cm test and 1.084 N m for the $l_0 = 41.0$ cm test, and the held-out validation tests give 1.368 N m at 45.7 cm and 1.611 N m at 40.6 cm with tendon (Table 4.6). The complete five-test comparison before and after correction is provided in Table C11.

The optimized prediction also captures the loss of torque near full extension (Figure 4.8). Agreement on the held-out tests checks the correction beyond the measurements used to

set its parameters; the training and validation errors are reported separately in Table 4.6.

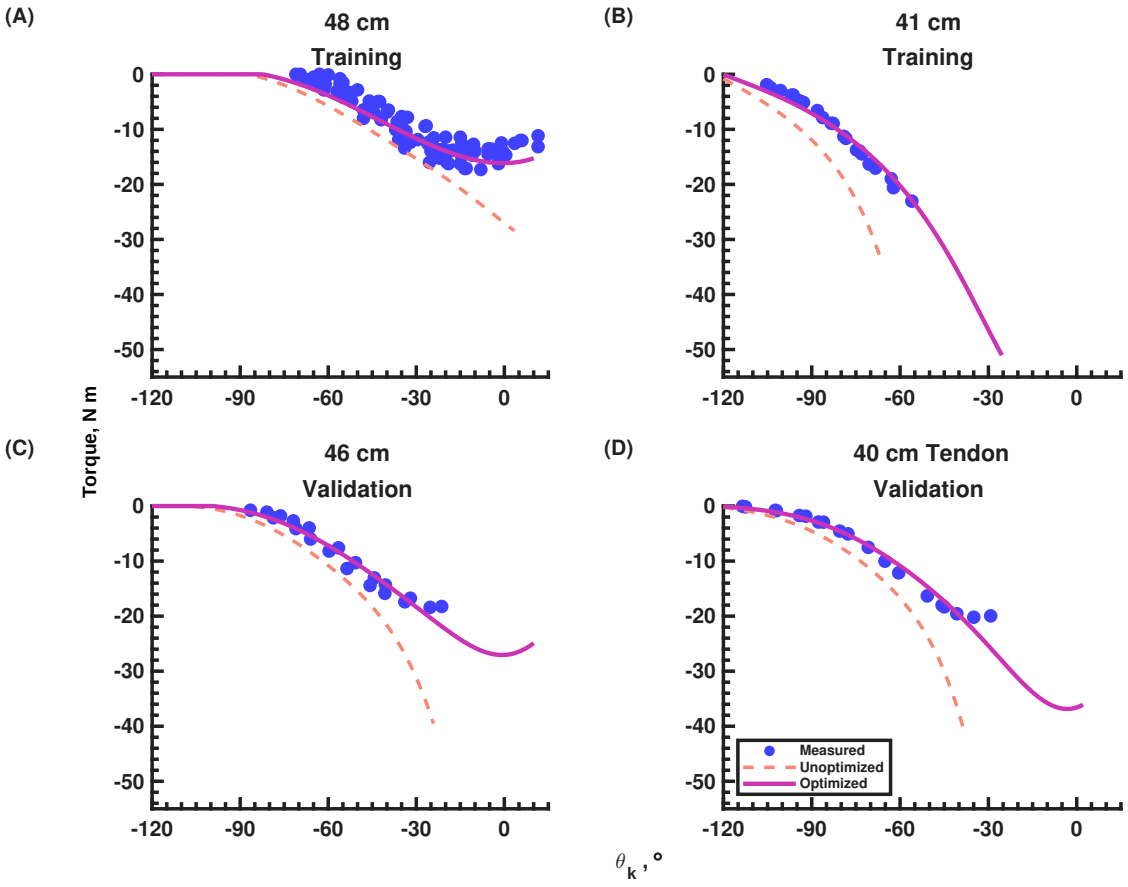


Figure 4.8: Optimized pinned-knee flexor torque: **(A)** $l_0 = 48.5\,\text{cm}$, **(B)** $l_0 = 41.0\,\text{cm}$, **(C)** $l_0 = 45.7\,\text{cm}$, and **(D)** $l_0 = 40.6\,\text{cm}$ with a 35 mm tendon. Panels **(A)** and **(B)** are training tests; **(C)** and **(D)** are validation tests. Circles show measurements, dashed curves the improved BPA model, and solid magenta curves the optimized model.

Table 4.6: Lengths l_0 (resting BPA) and l_{ten} (artificial tendon) are both in cm. Pinned-knee flexor goodness of fit before and after correction, using the adopted fold. RMSE and maximum absolute residual (Max.) are in N m; FVU is dimensionless.

		Improved BPA			**Optimized model**		
l_0 (cm)	l_{ten} (cm)	**RMSE**	**FVU**	**Max.**	**RMSE**	**FVU**	**Max.**
Training							
48.5	–	7.617	2.1278	25.531	2.165	0.1718	5.467
41.0	–	9.694	2.1149	21.076	1.084	0.0277	1.766
Mean		8.656	2.1213	23.303	1.624	0.0998	3.617
Validation							
45.7	–	8.803	1.9380	26.264	1.368	0.0468	3.614
40.6	3.5	6.373	0.5935	16.859	1.611	0.0427	5.891
Mean		7.588	1.2657	21.562	1.490	0.0447	4.752

The corrected model reduces mean torque-prediction error in both the training and validation tests (Table 4.6). For the $l_0 = 48.5$ cm and 41.0 cm training tests, mean RMSE decreases from 8.66 N m to 1.62 N m. For the held-out 45.7 cm and 40.6 cm tendon tests, it decreases from 7.59 N m to 1.49 N m. The baseline uses the same kinematics and force model with the correction terms disabled; Figure 4.7 separately compares the original rigid-body prediction.

4.2.2 Extensor on Pin Joint

When the new stiffness and length offset model was used to predict torque for the knee extensor, measured torque was slightly higher than the original rigid-body model predicts for the 48.0 cm resting length BPA, slightly lower than expected for the 45.7 cm BPA, and much lower for the 41.5 cm BPA. Only the χ_1 and χ_2 terms from the previous optimization results were used with the extensor configuration. By inspection, the zero point of the measured torque was shifted to the right of the original model, indicating an offset opposite in sign to the χ_0 term in the previous section. The BPA with $l_0 = 45.7$ cm illustrates the pattern: the hybrid-method calculated torque correlates well with the originally predicted torque, while the actual measured torque diverges.

Figure 4.9 shows pinned-knee torque versus knee angle for the extensor BPAs: the six identification tests at the adopted correction terms, with the extensor-specific terms (χ_0, χ_3) on the shared stiffness pair (Table 4.9). Knee angle is negative in flexion, and extension torque is positive. Gold circles are the hybrid calculation; indigo filled circles are experimentally measured torque; the dashed light orange line is the unoptimized prediction of the improved BPA characterization; the magenta line is the optimized prediction of the optimized model. Panel titles give the nominal resting length; the actual resting lengths are $l_0 = 43.6$ cm **(A)**, $l_0 = 43.2$ cm with tendon **(B)**, $l_0 = 45.7$ cm **(C)**, $l_0 = 40.0$ cm **(D)**, $l_0 = 40.0$ cm with tendon **(E)**, and $l_0 = 46.5$ cm **(F)**. Panels **(A)–(C)** are the training tests and panels **(D)–(F)** the held-out validation tests of the adopted cross-validation fold; the legend applies to all panels.

The adopted two-stage extensor identification uses the shared flexor stiffness pair and

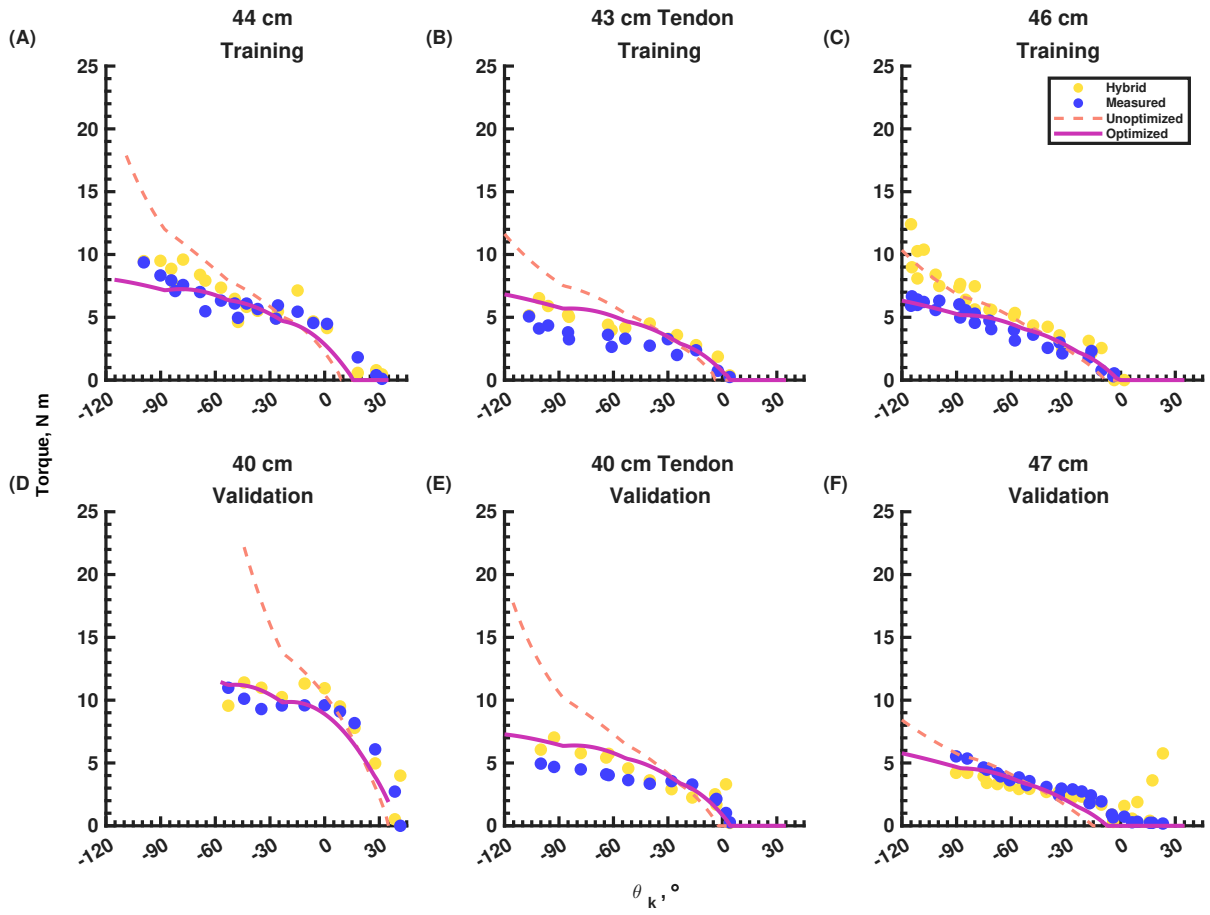


Figure 4.9: Optimized pinned-knee extensor torque: **(A)** $l_0 = 43.6\,\text{cm}$, **(B)** $l_0 = 43.2\,\text{cm}$ with a 22 mm tendon, **(C)** $l_0 = 45.7\,\text{cm}$, **(D)** $l_0 = 40.0\,\text{cm}$, **(E)** $l_0 = 40.0\,\text{cm}$ with a 40 mm tendon, and **(F)** $l_0 = 46.5\,\text{cm}$. Panels **(A)**–**(C)** are training tests; **(D)**–**(F)** are validation tests. The legend applies to all panels.

the six-test fold defined in Section 4.1.4. Figure 4.9 and Table 4.7 report that adopted fit.

At the adopted correction terms of the full re-identification (Section 4.1.4), Table 4.7 reports the per-test goodness of fit on the six extensor tests used by the identification, three for training and three for validation: the optimized model improves the RMSE of every used test over the zero-correction prediction and holds FVU below unity on all but the $l_0 = 43.2\,\mathrm{cm}$ tendon test. The original rigid-body model on these tests is shown graphically in Fig. 4.9.

Table 4.7 gives the per-test extensor fit, including training and validation means. The full nine-test before/after comparison is in Table C12; the separate biomimetic-knee fits are in Table 4.8.

Table 4.7: Pinned-knee extensor goodness of fit before and after correction, with stiffness shared with the flexor. RMSE and maximum absolute residual (Max.) are in N m; FVU is dimensionless.

		Improved BPA			**Optimized model**		
l_0 (cm)	l_{ten} (mm)	**RMSE**	**FVU**	**Max.**	**RMSE**	**FVU**	**Max.**
Training							
43.6	–	2.437	0.9618	4.284	0.978	0.1634	1.818
43.2	22	2.878	4.7304	4.817	1.525	1.3270	2.481
45.7	–	2.178	1.0586	4.974	0.490	0.0535	1.206
Mean		2.498	2.2503	4.692	0.997	0.5146	1.835
Validation							
40.0	–	3.092	0.5638	4.842	1.244	0.1303	2.196
40.0	40	3.622	6.3919	7.845	1.350	0.8878	1.890
46.5	–	0.949	0.3152	2.272	0.699	0.1714	1.559
Mean		2.554	2.4236	4.986	1.098	0.3965	1.882

4.2.3 Extensor on Biomimetic Knee

The biomimetic knee was tested with a ∅10 mm BPA in the extensor configuration. The bracket frame origin is the centroid of the cantilevered segment that holds the muscle origin bracket to the proximal femur (see Fig. 4.3**C**, right side).

I tested the extensor configuration on the biomimetic knee with a ∅10 mm BPA at $l_0 = 51.8$ cm. The measured torque fell well below both the original prediction and the hybrid calculation (Figure 4.10). Applying the extensor terms identified on the pinned knee improved the prediction: at the adopted two-stage terms, the optimized model reached an

RMSE of 1.1702 N m, an FVU of 0.6681, and a maximum absolute residual of 2.2088 N m, against 3.0322 N m, 4.4856, and 5.5870 N m for the original rigid-body prediction. On this configuration the actuator path passes between the femoral condyles close to the instantaneous center of rotation, a geometry that is difficult to represent with the modeled path (Figure 4.10**(A)**). The hybrid points are an approximate diagnostic because the curved BPA length and its distance from the ICR are difficult to measure within the femoral condyles.

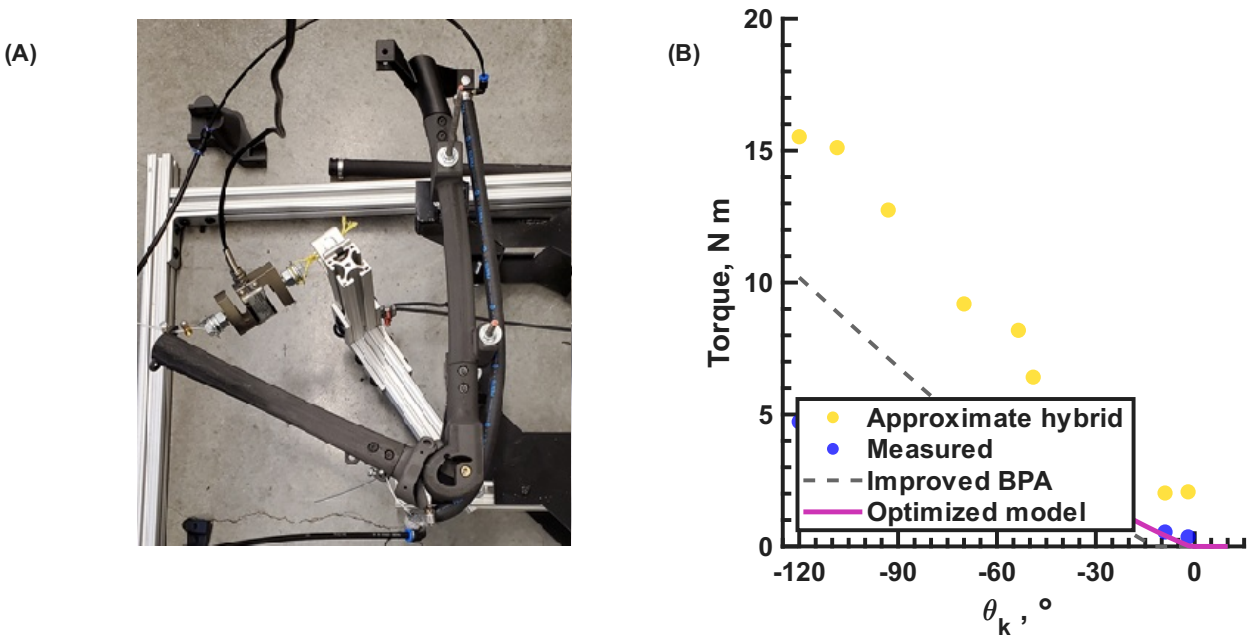


Figure 4.10: Biomimetic-knee extensor test: **(A)** apparatus with the BPA passing between the femoral condyles; **(B)** measured torque, the improved BPA baseline, the optimized model at the adopted terms, and the approximate hybrid calculation for the 10 mm-diameter, 51.8 cm BPA.

4.2.4 Flexor on Biomimetic Knee

Fig. 4.11 shows the biomimetic-knee flexor measurements and the corrected-model predictions using the adopted shared terms (Table 4.8); the bracket frame for the ∅10 mm BPA is shown in Fig. 4.3**C**, and for the ∅20 mm BPA in Fig. 4.3**D**. Panel **A** presents the ∅10 mm actuator at $l_0 = 41.5$ cm with a 21 mm artificial tendon (∅1.5 mm Shimano bicycle brake cable wrapped four times), pressurized to 604 kPa; the optimized model reached an RMSE of 2.8002 N m, an FVU of 0.8835, and a maximum absolute residual of 4.1201 N m. Panel **B** presents the ∅20 mm actuator at $l_0 = 41.5$ cm with a 14 mm artificial tendon made out of the same brake cable, wrapped three times so that six cable segments share the load, each wrap contributing two tendon lengths and cross-sectional areas. Panel **B** was tested at 620 kPa, for which the optimized model reached an RMSE of 2.2070 N m, an FVU of 0.0748, and a maximum absolute residual of 3.7320 N m.

The torque of the single ∅20 mm flexor actuator of panel **B** did not exceed the magnitude of the maximum human isometric flexion torque between 60° and 100° of knee flexion, and within a ±20% band about the human mean it still fell short of human-level torque between 78° and 103° of flexion. Because operating at 620 kPa across the full desired range of motion would load the printed femur beyond its ultimate strength, single points were also recorded at the reduced pressures 560 kPa, 421 kPa, 325 kPa and 281 kPa, which show that the desired range of motion is achievable at reduced pressure. This shortfall in flexion torque defines the target that the attachment-route redesign of Section 4.3.2 addresses. Panel **C** shows the outcome for the redesigned two-BPA flexor route of Section 4.3.5: across its fifteen tested measurements from −119° to 5°, the measured torque magnitude met or

exceeded the human biceps-femoris short-head maximum at every point, with margins from 0.42 N m at -105° to 17.0 N m at -14°.

Table 4.8: Biomimetic-knee goodness of fit before and after design-parameter optimization. RMSE and maximum absolute residual (Max.) are in N m; FVU is dimensionless.

	Improved BPA			Optimized model		
Configuration	**RMSE**	**FVU**	**Max.**	**RMSE**	**FVU**	**Max.**
Extensor, ∅10 mm	3.0322	4.4856	5.5870	1.1702	0.6681	2.2088
Flexor, ∅10 mm	7.4513	5.5612	13.3163	2.8002	0.8835	4.1201
Flexor, ∅20 mm	14.8064	3.3672	23.7218	2.2070	0.0748	3.7320

4.3 Discussion

4.3.1 Reasons for the Correction Terms

The correction terms of Section 4.1.4 exist because the physical transmission between the actuator and the joint is neither rigid nor lossless, and each term names one family of physical reasons. A constant length offset χ_0 exists because of manufacturing and positioning tolerances: the printed bone segments, the machined knee links, and the end-cap fittings each carry dimensional tolerance, the actuator is pinned at a measured location on that stack of tolerances, and the accumulated discrepancy appears as a length that is consistently shorter than the predicted musculotendon path at every joint angle, including at zero load where no compliance can act. The stiffness terms χ_1 and χ_2 exist because loaded components displace: the isometric system is not rigid, and there is compliance in the winch cable of the test stand, the artificial tendons, and the cantilevered attachment brackets that bend and stretch under actuator load. The wrapping-loss term χ_3 exists

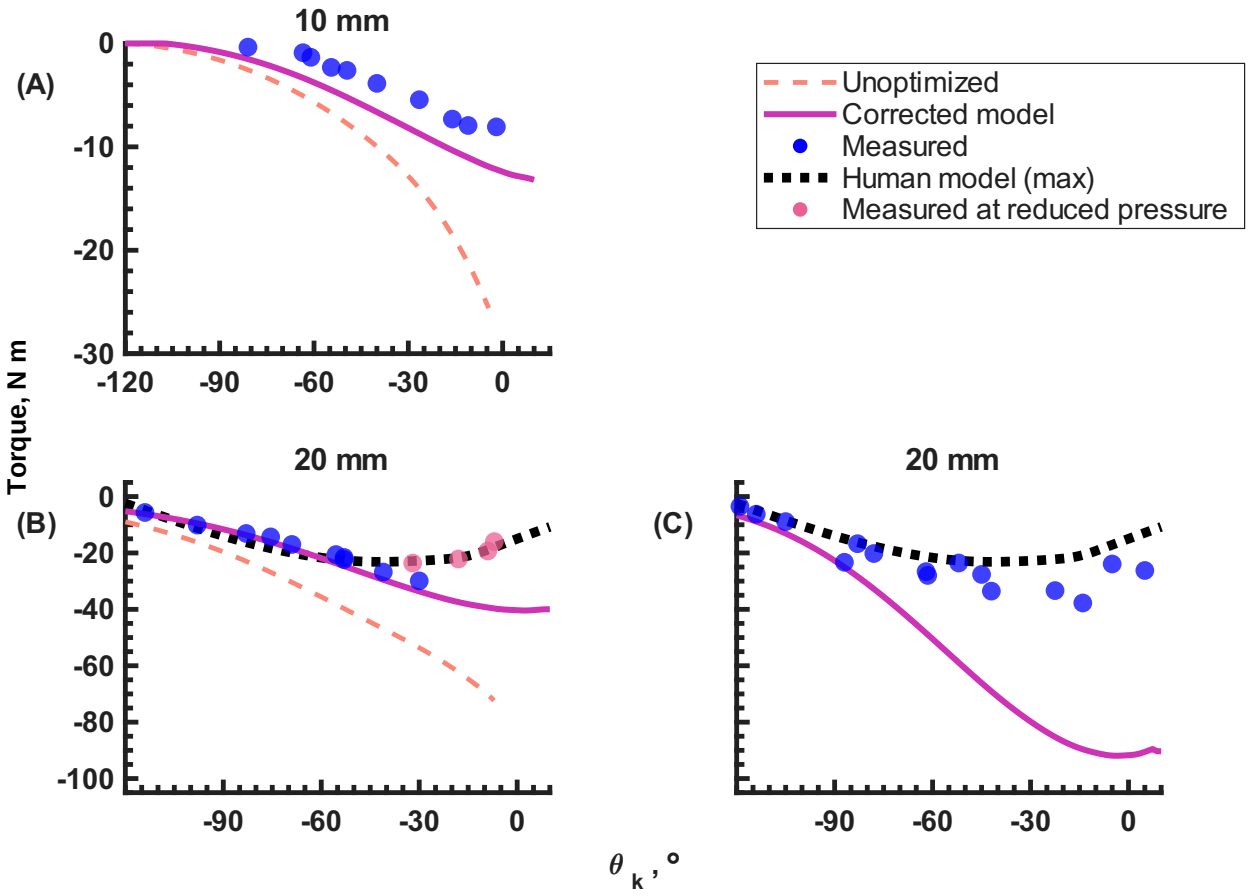


Figure 4.11: Biomimetic-knee flexor torque: single 10 mm BPA (A), single 20 mm BPA (B), and redesigned two-BPA route (C). Measurements, predictions, and human targets follow the shared legend; B and C use the same axis ranges.

because a BPA that bends over the joint loses usable contraction length that no straight-line model accounts for, and because the same loading could displace the actuator off its intended seat, as the observed slip off the anterior head of the femur bolt shows.

Our improved BPA characterization predicts torque more accurately than the original BPA model (see Fig. 4.7). This only accounts for part of the error. Many factors affect the torque results. Compliance in the testing system, positioning tolerances, flexibility of brackets, imperfect path modeling, and variation in individual BPAs are other factors.

The adopted identification uses four pinned-flexor tests and six pinned-extensor tests, with separate fitting and validation subsets. Sharing the effective stiffness pair between the two configurations tests whether the same representation of attachment compliance can support both torque predictions.

The modeling of extensor BPA paths that wrap around the knee joint is an approximation with inherent error (Fig. 4.10). High actuator force and a short moment arm make the torque sensitive to attachment deflection and small geometric errors. Small changes in moment arm can therefore produce substantial torque errors [**?**]. It was assumed that wrapping the BPA around the joint would be easier to account for than adding a patella, deterministically designing its location as a function of knee angle, and accounting for friction. At greater knee flexion angles, shorter l_0 BPAs were observed to stretch and compress (see Fig. 4.1**C**). The same loading sometimes caused the BPA to slip off the anterior head of the bolt that holds the femur to the knee. This produced a $\pm 20\,\text{mm}$ displacement in the Z direction, changing both the force vector and the torque. Fig. 4.9 shows a deviation between measured and originally predicted results that widens as expected torque increases. This indicates system compliance. However, the hybrid-calculated torque matches

or exceeds the predicted torque, except at greater flexion angles for the $l_0 = 41.5\,\text{cm}$ BPA, indicating a loss of functional l_m; this is why the χ_3 term was created in (4.25). Fig. 4.9 shows the improved fit from running the optimization that includes this term.

The biomimetic knee joint with an extensor BPA was used to validate the extensor optimization results. The corrected prediction more closely matches the measured torque values, although it remains higher than the measurements.

The pinned-knee tests distinguish actuator-force error from errors introduced by attachment compliance and routed geometry. Applying the identified terms to the biomimetic knee then tests their usefulness when the ICR and moment arm change with knee angle. This sequence links the correction terms to measured mechanisms and identifies the remaining limits of torque prediction.

The isometric force characterization supplies a static force–length relation for later dynamic actuator models, including the pressure–force formulation of **?**. It does not identify velocity dependence or damping. Isometric testing measures torque at fixed joint positions and therefore establishes a static capacity envelope. Isokinetic testing controls joint speed, allowing torque to be compared across angles at a defined velocity and repeated at several velocities. For this robot, those tests would show how contraction rate, pressure regulation, and changing moment arm affect the torque available during movement, and whether the static design margin remains adequate. Isotonic testing maintains a prescribed external load while the actuator shortens; it would assess shortening speed, achievable excursion, and the ability to support a load through the movement. Loading and unloading sequences would quantify hysteresis, while pressure-step tests would characterize pneumatic response. A torque-controlled motor and a less compliant

loading connection would reduce winch slack and permit controlled-speed and controlled-load tests on the same rig. Together, these measurements would identify the additional dynamics needed to extend the static model to movement and load-bearing control.

4.3.2 Optimizer Updates

The results figures use the adopted two-stage identification. The procedure of Section 4.1.4 yields $\chi_0 = 3.94\,\mathrm{mm}$, $\chi_1 = 3.998 \times 10^4\,\mathrm{N\,m^{-1}}$, and $\chi_2 = 1.473 \times 10^4\,\mathrm{N\,m^{-1}}$ for the flexor, and $\chi_0 = -6.38\,\mathrm{mm}$ with $\chi_3 = 0.158$ for the extensor on that same locked stiffness pair (Table 4.9); the search bounds and the adopted values are recorded in Appendix C (Tables C10 and 4.9).

Table 4.9: Adopted transmission corrections. The stiffness pair is shared. The wrapping factor was identified on the extensor and is also used in the flexor redesign; it was not fitted to the pinned-flexor tests.

Configuration	χ_0 (m)	χ_1 (N/m)	χ_2 (N/m)	χ_3
Flexor	0.00394	3.998×10^4	1.473×10^4	0.158
Extensor	−0.00638	3.998×10^4	1.473×10^4	0.158

What the correction terms do to the prediction is quantified in Tables 4.6 and 4.7. The corrected flexor model holds FVU at or below 0.18 on the four tests of the identification, the worst case being the 48.5 cm training test at 0.172. On the extensor identification the optimized model improves the RMSE on every used test, cutting the training-triple mean from 2.50 N m to 1.00 N m and the validation-triple mean from 2.55 N m to 1.10 N m, and holds FVU below unity on every used test except the 43.2 cm tendon test. On the biomimetic extensor test ($l_0 = 51.8\,\mathrm{cm}$) the corrected prediction improves RMSE from

3.03 N m to 1.17 N m and FVU from 4.49 to 0.67, so the shared-pair model explains the bulk of that configuration's measured torque variance as well as its level. The genetic multiobjective solver is retained for the identification, where the Pareto front is required for holdout validation, and the surrogate-plus-pattern-search pipeline serves the route redesign (Section 4.1.5).

The original placement methods selected routes by their predicted torque profiles under rigid attachment assumptions. The isometric tests show why those predictions require compliance corrections before they can guide hardware selection. Measured muscle lengths were shorter than predicted, and torque error depended on load. The redesigned routes therefore use the corrected force balance when evaluating candidate geometries.

The biomimetic-knee tests assess the corrected torque model on a four-bar joint and on different actuator sizes. On the biomimetic extensor, evaluating the shared adopted terms reduces FVU from 4.49 to 0.67. The biomimetic flexor comparisons use the same adopted stiffness pair and give FVU values of 0.8835 for the 10 mm actuator and 0.0748 for the 20 mm actuator (Table 4.8). The larger error for the smaller actuator limits the accuracy of transfer across configurations. The redesigned two-BPA routes provide a further physical test: their measured torque meets or exceeds the human-reference magnitude at every tested angle (Section 4.3.5).

The final mesh optimizers close the loop between the two families of Section 4.1.5, and their cost functions record the evolution explicitly. The original mesh optimization scored rigid-body torque surplus with the reward of Equation (4.32). The route redesign replaced the reward with the penalized requirement of Equations (4.35) and (4.36), shortfall against a required envelope, evaluated through the stiffness-aware force balance. The extensor's

final form, the cam-profile redesign, adds three terms to the family objective,

$$J_{\text{cam}} = J_{\text{ext}}(\mathbf{x}_{\text{full}}) + J_{\text{arm}} + J_{\text{armStep}} + J_{\text{reg}}, \qquad J_{\text{arm}} = -10^2 \, \min\left(1, \, \min_j m_j^{\text{req}}\right), \quad (4.37)$$

where $m_j^{\text{req}} = (|M_j^{\text{R}}| - (1 + m_r)|M_j^{\text{H}}|)/\tau$ is the margin against the required envelope of Equation (4.34), and the moment-arm step and cam regularizer carry the weights of Table 4.2. The reward becomes a capped payment for margin beyond the requirement, so the solver keeps buying moment arm after the envelope is met, and the step term holds the release schedule smooth while it does.

The initial route searches bounded the BPA origins and insertions near the corresponding human muscle attachments to preserve biomimetic routing. Those anatomical bounds had to be relaxed where the available BPA contraction could not cover the knee range of motion. The resulting routes then required a second tradeoff: reducing the actuator forces carried by the attachment hardware after parts failed during testing. The cam redesign changes the moment arm and contact geometry to retain the torque requirement at lower actuator force.

The final flexor optimizer applies the wrapping-loss term χ_3, identified on the extensor, because the flexor route can also bring the BPA into contact with the tibia and the femoral condyles at high flexion, and an actuator that a possible contact could bend must be scored with the length that bending removes. Its search box was deliberately expanded as an acknowledged tradeoff with biomimicry: the muscle origin was freed to move more than 100 mm proximal of the femoral shaft, far from the human biceps femoris origin, purely for reserve contraction, the extra actuator length, and therefore force and torque, available at

high flexion where the single-actuator tests fell short. Two 20 mm BPAs are carried, both to reach the required torque and to limit the off-axis torque that the single-mirrored-actuator arrangement was found to load parts past their breaking point. A mirrored two-actuator solution, the two routes exact reflections of one another about the mid-sagittal plane, was found and rejected: it is not bio-realistic if a pelvis is ever introduced, since the medial origin would sit where pelvic structure belongs. The adopted solution is a staggered-offset pair, with a separately solved route for each actuator; the second insertion mirrors the first across the sagittal plane, while its origin preserves the first route's lateral endpoint offset, so both origins sit off the same side of the knee and the two routes do not cross. Each route is given its own wrapping calculation and its own collision checks.

The final extensor optimizer faces a requirement of a different order. The extensor target is set by vastus medialis, the quadriceps muscle with the smallest maximum isometric force in the reference set, and even that weakest-quad envelope demands roughly twice the peak torque of the flexor requirement, 50.32 N m against 23.16 N m. The two 20 mm extensor BPAs nominally share one path, a single route whose force is scaled by the actuator count, because the extensor route wraps over the front of the joint through as many as seven sequential contacts and duplicating that chain would double the potential for interference. That shared-path economy has a programming cost: the large solution space of the wrapped route makes the contact and release angles, the knee angles at which each wrap contact activates and releases, difficult to schedule by inspection, and the calculation that produces them is formalized in Appendix C.7. The routing was additionally modified to reduce the force the actuators exert, because the direct route's higher actuator forces broke the printed and machined parts during testing; reducing moment-arm demand and

contact severity at lower force is what led into the cam algorithm of Equation (4.37).

The route redesign built on these values addresses the single-muscle shortfall at high flexion angles documented in Section 4.2: each configuration now carries two 20 mm BPAs per route. The flexor solution meets the flexor torque target with a minimum margin of $+5.01\,\%$ at -107°, with a separately solved route for each actuator. The second insertion mirrors the first across the sagittal plane, while its origin preserves the first route's lateral endpoint offset; its wrapping and clearance are evaluated independently. The direct extensor route meets or exceeds its human-scaled target with the required $5\,\%$ overshoot as the thinnest margin at every angle, $+5.02\,\%$ at worst, while respecting the wrap-release schedule and contraction limits of the route-elimination rule (Appendix C); it was later superseded by the cam redesign of Section 4.3.2. These are mechanical design predictions. Their complete incorporation into the converted neuromechanical models remains future integration work.

Redesigned-Route Results on the Gait2392 Robot-Body Model

This subsection records the quantitative outcomes of the redesigned routes on the Gait2392-derived robot-body model, whose muscle force and torque targets define the redesign objectives. All three result families were produced by the surrogate-plus-pattern-search pipeline of Section 4.1.5 under the adopted correction terms (Appendix C, Table 4.9) with one shared actuator plan: two 20 mm BPAs per route, a 620 kPa pressure ceiling, and a maximum usable contraction of 0.255 in the force model. Each family minimizes its own objective function, a unitless weighted sum of normalized penalty terms in which

zero corresponds to no violated penalty, led by the worst-angle and mean-square torque shortfall against the required 5 %-overshoot human envelope, followed by overshoot, rest-length, strain, contraction, and mild design-change penalties. The three functions differ in penalty families and normalizations (the flexor and pulley objectives add off-axis-excess, wrap-release, and second-BPA anchor terms, and the pulley objective evaluates its torque requirement on the insertion-side torque after the tackle's force loss), so their numerical values are not commensurable with one another. The human-scaled requirements are the OpenSim isometric torque profiles of the reference human muscles over the same design-grid angles, the vastus medialis for the extensor and the biceps femoris short head for the flexor.

The direct extensor route exceeds the human envelope in magnitude at every angle, with zero shortfall points: across the design range the predicted two-BPA extensor torque spans 35.45 N m to 109.4 N m against the human-scaled requirement of 20.25 N m to 50.32 N m, and the required 5 % overshoot is consumed as the thinnest margin exactly as designed, at a minimum margin of $+5.02$ % at the worst-case angle of -20.2°.

Allowing the tibial-side route points to trace a cam profile produced the current extensor design (Figure 4.12). The final cam geometry includes the optimized route and the verified adjustment of the distal route point and tendon length. Its minimum torque margin is $+30.83$ % with respect to the required envelope, the human target plus the 5 % overshoot, and $+39.77$ % with respect to the plain human target. The pressure required to meet the envelope spans 269 kPa–507 kPa across the range of motion, comfortably above the 150 kPa regulated floor of the VTUG valve manifold that supplies the hardware. The geometry confirms that the headroom is not bought with routing risk: the tibial wrap follows a

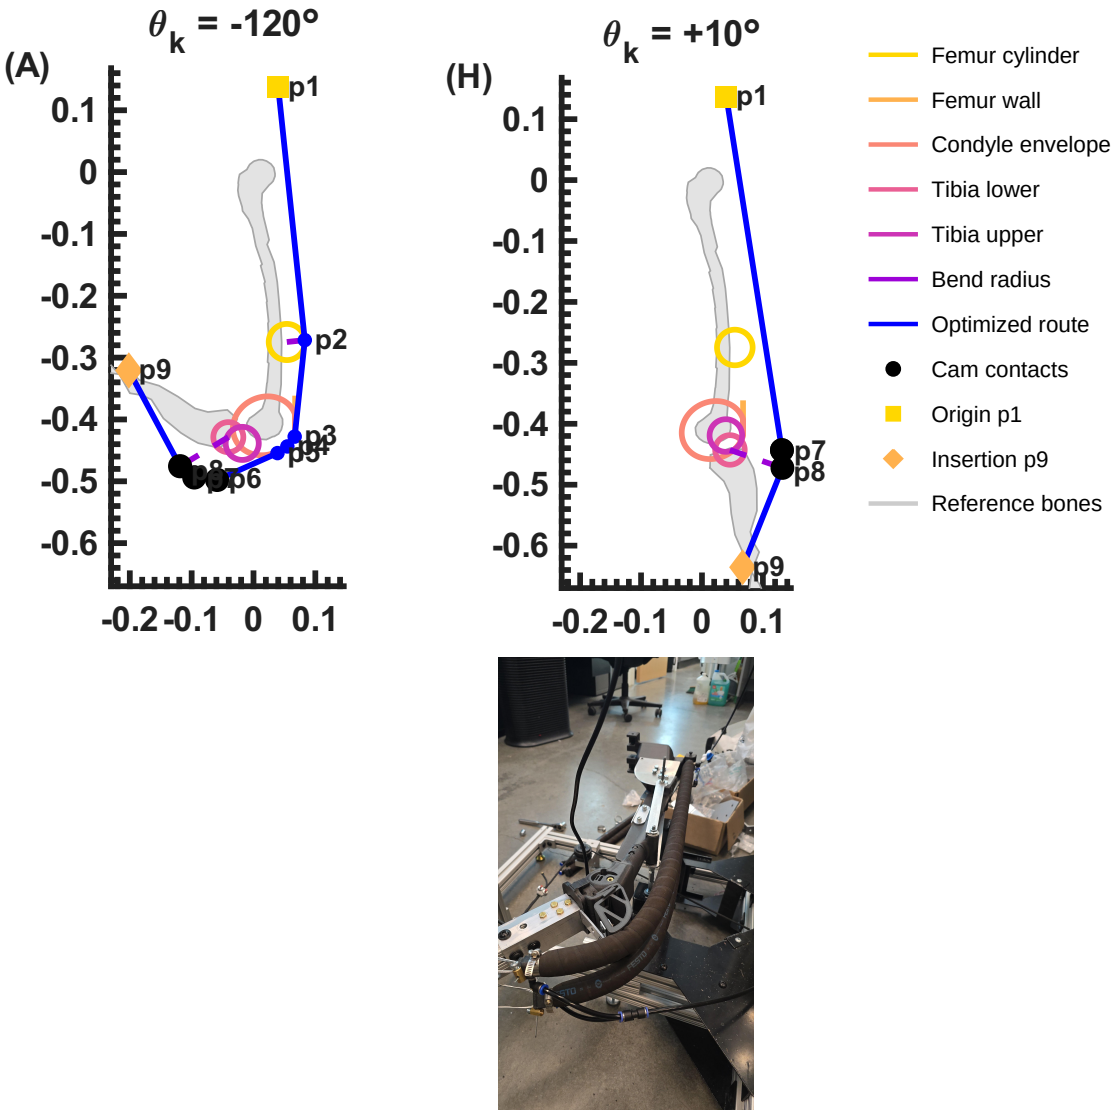


Figure 4.12: Cam-extensor geometry at -120° (A) and 10° (H), with the assembled two-BPA hardware below. Panel letters match the complete sequence in Figures C5 and C6.

bend radius of 96.9 mm, the guiding route point clears the femoral-condyle envelope by 43.1 mm at 10° of knee extension, and the moment arm spans 59.4 mm–102.8 mm (mean 72.4 mm) over the range of motion.

The flexor route, built from two separately evaluated BPA paths, meets the flexor torque target with a minimum margin fraction of 0.0501 against the required 0.05. Its predicted knee-axis torque spans $-81.23\,\mathrm{N\,m}$ to $-3.864\,\mathrm{N\,m}$ against the human-scaled requirement of $-23.16\,\mathrm{N\,m}$ to $-2.461\,\mathrm{N\,m}$ (peak magnitude $23.16\,\mathrm{N\,m}$), again exceeding the human envelope at every angle. The flexor redesign applies the adopted correction terms of Table 4.9, including the wrapping factor identified on the extensor configuration.

The Mono/Bi-pulley algorithms extend the route search space by allowing the artificial tendon to be reeved over pulleys before insertion. On the knee-flexor route, the pulley variant converged to a collision-feasible layout with zero bound violations and a peak predicted flexion torque of $85.56\,\mathrm{N\,m}$ on the insertion-side measure, the torque applied at the insertion after the tackle's force loss; Section 4.3.3 compares the optimized configurations using knee-axis and off-axis torque.

4.3.3 Mono- and Bi-Pulley Tendon Routing

Every optimizer of the preceding sections scores its routes with the MonoPamDataExplicit-type calculation: a fixed chain of via points carrying one segment across from the parent body into the child, its contacts wrapped and released by the elimination rule, its length and moment arm read from that chain. The Mono- and Bi-pulley algorithms are the generalization of exactly that calculation. The artificial tendon is reeved over pulleys

before insertion, forming a reverse pulley, a block-and-tackle transmission: with G tackle line parts supporting the moving block, the tendon travel at the insertion is multiplied by G while the force delivered there is divided by G, and at $G = 1$ the reeved route reduces to the straight-tendon route of the direct redesign. The tendon exit is either body-fixed on the proximal body so that its direction rotates with the joint, or anchored through a housing to the distal body, and the tackle's increased internal force and the pulley-block bracket deflections enter the same stiffness-aware force balance as the direct route's brackets. Read this way the pulley variants are not a new actuator concept but the general routing case of the mesh-optimization machinery, of which every direct route is the $G = 1$ special case. Figure 4.13 distinguishes the travel multiplier from the force gain and shows the separate mounting loads.

At unity transmission, the knee-flexor pulley route increased peak predicted knee-axis torque from 81.23 N m to 85.74 N m, while limiting off-axis torque to 0.803 N m. This comparison measures the effect of the two optimized routing configurations; their different objective functions do not isolate the pulley alone. The pulley serves as a redirect in this configuration rather than a force multiplier. Appendix C.8 records the optimizer diagnostics and the broader route searches.

A subsequent transmission study re-solved actuator sizes for each tackle configuration at five anatomical muscle sites. A two-BPA pair through a two-part tackle preserves net actuator force while doubling insertion travel; increasing travel can extend the angle range over which a route meets its torque requirement. The corresponding torque and margin surfaces are shown in Figures C3 and C4. These calculations use the moving-force-point routing, with the adjacent joint fixed for the biarticular evaluation planes; they do not

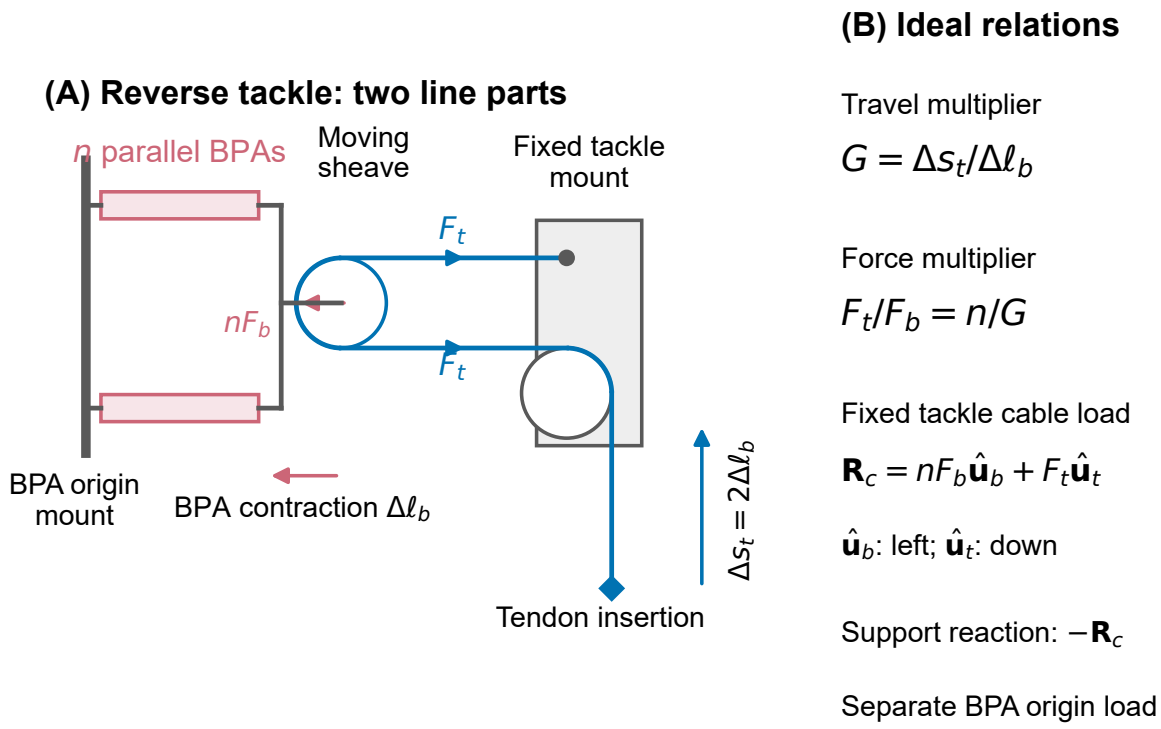


Figure 4.13: Reverse-pulley principle for two moving-block cable parts (A), and ideal force, travel, and mounting-load relations (B). Each of the n parallel BPAs supplies force F_b. The travel multiplier is G; force gain relative to one BPA is n/G. The cable load on the fixed tackle mount combines the forces along the moving-block and insertion-side cable directions. Its support reaction is equal and opposite. The BPA origin mount carries a separate load. Compliance and pulley losses are omitted from this schematic.

establish simultaneous performance across both joints. Physical verification of the pulley configurations remains future work.

4.3.4 Whole-Leg Transmission Capability

With improved torque prediction and route-design methods in place, I revisited the robot-leg configurations of **?**. The expanded design search varies parallel actuator count and transmission ratio and re-solves actuator and tendon lengths on the inherited routing. Table 4.10 summarizes the improvement in coverage of the human torque requirements.

The configuration survey grades each actuator against the same-named single muscle of the Gait2392 model [**??**] on a 5×5 grid of two joint coordinates. The hip adduction domain is clipped by its flexion-dependent envelope. A cell meets the requirement when the BPA is taut and its torque has the human torque's sign and at least its magnitude. The coverage metric is the fraction of included grid cells meeting that condition; the aggregate is the unweighted mean of the per-actuator fractions, not a continuous-area integral or a whole-leg torque sum.

The BPA force model is evaluated at 620 kPa [**?**]. The comparison includes all 27 actuator sites (Table C14). The baseline uses the master's-era attachment routing at unity transmission; its eight nominal 40 mm sites are evaluated with four parallel 20 mm BPAs, while the other sites use the survey's base configuration. The updated side selects the saved transmission-ladder winner at each site, with actuator length and tendon length re-solved per configuration. It retains the same underlying attachment routing and the original diameter specification, including 40 mm actuators at those eight sites. The comparison

therefore combines transmission, actuator sizing, and the baseline diameter substitution; it does not isolate an effect of attachment placement or compare two all-20 mm assemblies.

Active-force coverage is the fraction of included grid cells with positive insertion tension, excluding slack or infeasible configurations. It does not require the torque to match the human reference. Both coverage measures are calculated for each actuator and then averaged equally across the 27 sites. Table 4.10 reports the resulting means.

The mean fraction meeting the human torque requirement increases from 37.28 % for the baseline to 65.19 % for the saved transmission configurations, improving at 23 sites and tying at four. Mean active-force coverage increases from 92.44 % to 98.81 %. The higher active-force coverage shows that the transmission configurations maintain tension over more of the graded domain, while the lower torque coverage shows that positive tension alone is insufficient to meet the human torque requirement.

Table C15 gives the coverage at each actuator site.

Table 4.10: Mean torque and active-force coverage across 27 actuator sites at 620 kPa.

Configuration	**Human torque met (%)**	**Active force (%)**
Baseline	37.28	92.44
Transmission configurations	65.19	98.81

4.3.5 Physical Tests of the Redesigned Routes

The redesigned cam-assisted extensor route was tested at nine isometric measurements, and the two-BPA flexor route at fifteen. Both test sets span knee angles from -119° to 5°. The analysis transforms the measured load-cell wrench to the moving knee center and

evaluates the model at each measured angle and pressure. Table C7 in Appendix C records the measured per-actuator dimensions and distinguishes them from the shared prediction inputs. Figures 4.14 and 4.15 compare these results with the selected human-muscle targets and the optimized 620 kPa predictions.

All nine extensor measurements meet or exceed the vastus medialis reference torque at the measured angles. Measured extension torque ranges from 29.02 N m to 51.17 N m, with margins of 0.24 N m to 3.52 N m, at pressures of 256 kPa to 488 kPa. The measured torque exceeds the pressure-specific prediction at eight angles, by up to 13.75 N m; at -86.5°, the measurement exceeds the human reference and is 2.06 N m below the pressure-specific prediction. This prediction uses the recorded angle and pressure with the modeled route. It is not the hybrid calculation, which additionally requires measured muscle length and moment arm.

All fifteen flexor measurements meet or exceed the biceps femoris short-head reference magnitude, with margins from 0.42 N m at -105° to 17.0 N m at -14°. Agreement with the pressure-specific prediction remains uneven. These measurements support torque capacity at the sampled configurations; variable test pressure and discrete sampling do not establish a measured full-pressure envelope between every pair of angles.

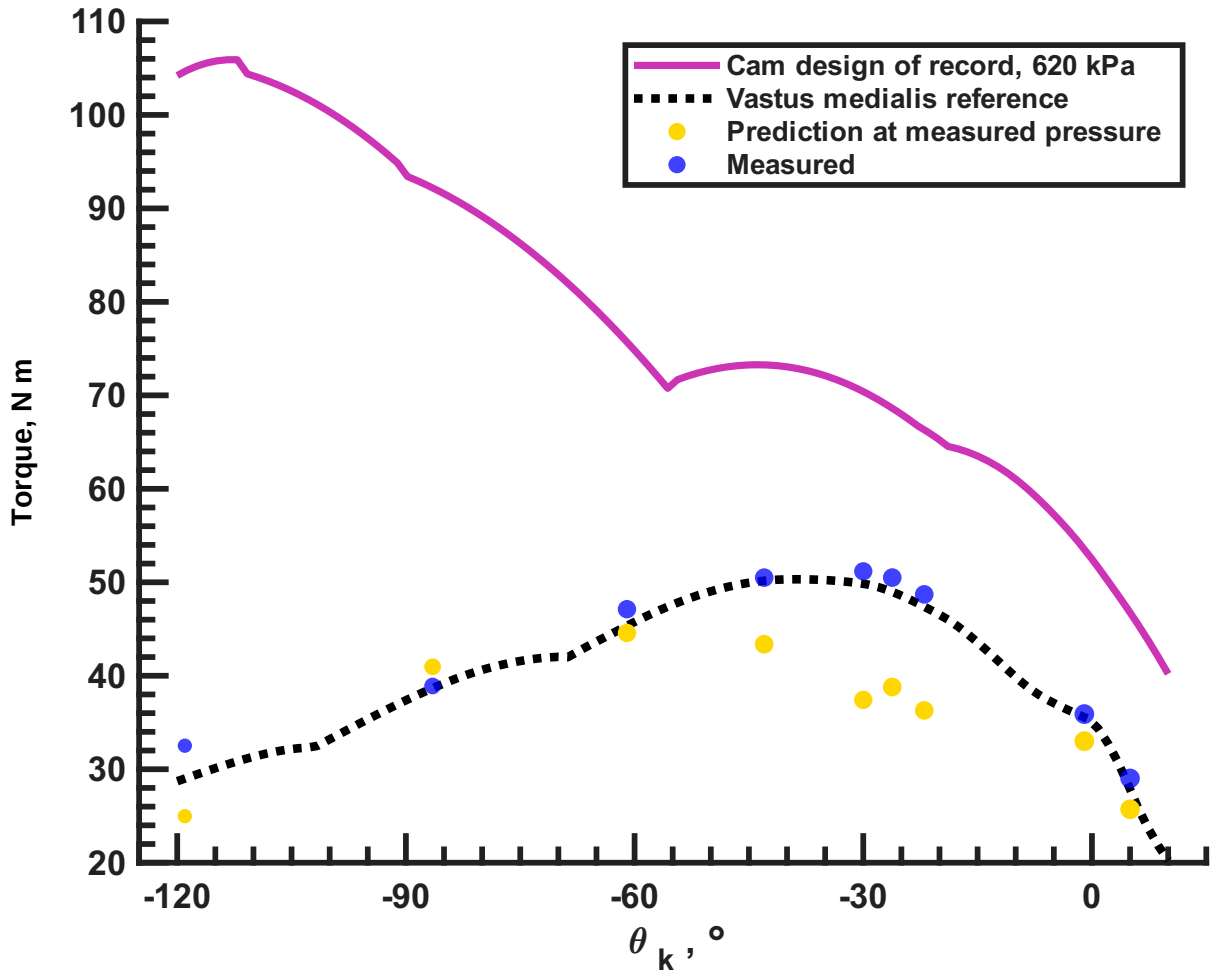


Figure 4.14: Cam-extensor measurements, pressure-specific predictions, the optimized 620 kPa curve, and the vastus medialis reference. Marker area scales with pressure. Series follow the plot legend.

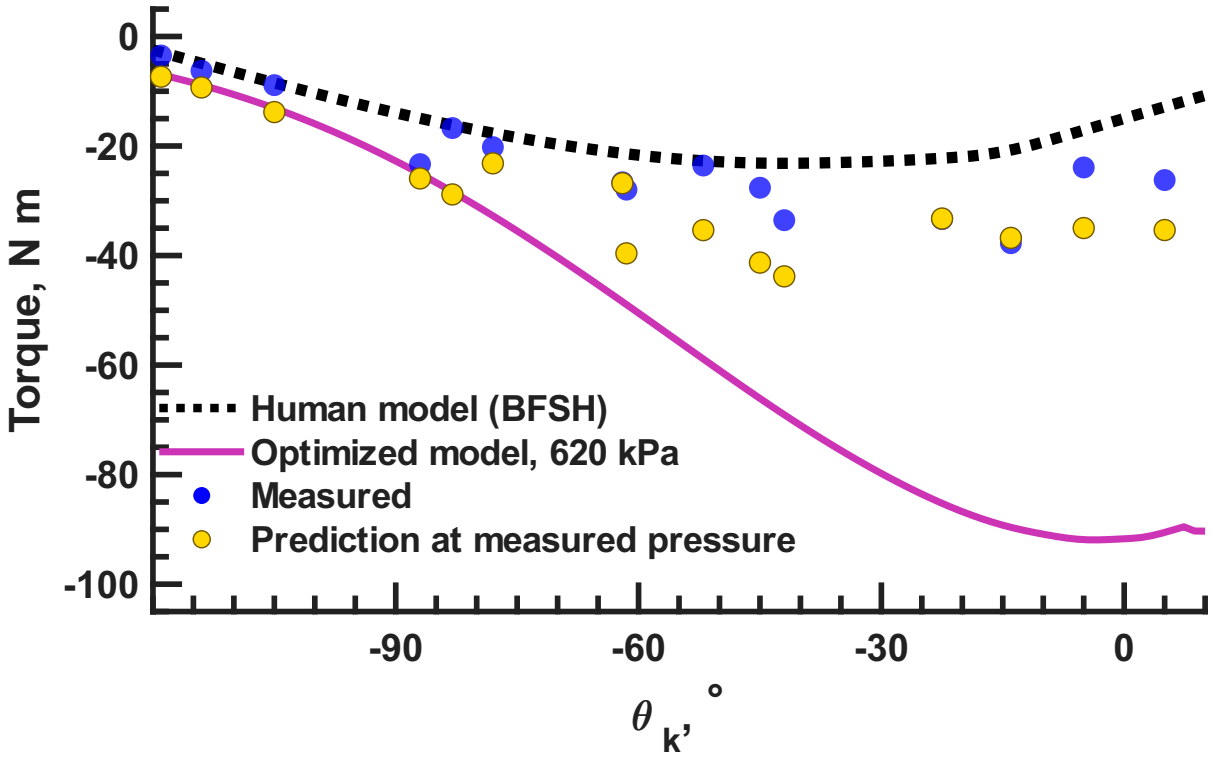


Figure 4.15: Two-BPA flexor tests (nominal 562 mm plot label); the installed resting lengths were 561 mm and 560 mm. Blue points are measured torques; gold points are model predictions at the measured angles and pressures. The solid magenta curve is the optimized 620 kPa prediction, and the black dotted curve is the human biceps femoris short-head target.

5 The Sensory Afferent Database

5.1 The Resource

The Sensory Afferent Database (SADb) is a curated literature resource for studying sensory feedback in locomotion and posture. It connects papers to afferent classes, experimental preparations, feedback pathways, and computational models. I developed it within the AARL's NeuroNex Communication, Coordination, and Control in Nervous Systems (C3NS) effort and presented the resource at Neuroscience 2023 [**?**].

SADb addresses the need to recover measured afferent properties, feedback pathways, and experimental conditions from a dispersed literature. Existing reviews organize parts of this evidence [**???**]; the application links individual findings to their sources and candidate controller roles, while distinguishing observations from modeling assumptions.

The primary deliverable is SADb Explorer, a browser application for searching literature records, comparing their classifications, and following groups of papers connected by citations (Figure 5.1). Its embedded snapshot contains 1,393 records, and the table, pivot, and citation views can be used offline. A separate online-search tab identifies candidate additions. The app preserves the collected material independently of the size of the editable curation base, so a reader can inspect both current and archived records.

I used Airtable for editable curation. Linked Papers, Feedback, Models, Review Papers, and Recorders tables supply bibliographic information, source-based notes, classifications, and pathway links. DOI and normalized-title matching reconcile records. SADb Explorer

brings these categories together in the distributed application.

5.2 Curation Methods

The initial collection drew on the literature libraries used for previous AARL papers on locomotion and postural balance, together with records already curated in Airtable. Records were reconciled by DOI and title, then expanded through citation searches. Reviews of spinal and proprioceptive feedback supplied an initial organization of the literature [**???**]. This was a curated research collection, rather than a registered systematic review or an exhaustive search of every relevant publication.

Records were retained when they reported sensory-afferent measurements, investigated feedback in locomotion or posture, or modeled a relevant sensory pathway. Reviews were retained as sources for organizing and locating that evidence. Duplicate records and material outside this feedback scope were removed from the working base. Classification used controlled vocabularies for preparations and afferent classes. Each mechanism summary was grounded in available source text, using the abstract when full text was unavailable; missing text was flagged rather than supplied from recollection.

5.3 Results: Application and Corpus Coverage

The application contains 1,393 entries in five record categories: 943 Papers, 247 Models, 157 Review Papers, 42 Feedback pathways, and 4 Recorders. These categories describe how the application organizes the literature and its linked material; they are not a count of unique publications across categories. The table view searches the complete 1,393-entry collection.

The pivot and citation views use the 943 entries in the Papers category. Archived records remain searchable in the application. Appendix E records the smaller maintenance subset separately.

Within the application's Papers category, containing 943 entries, 833 carry curation notes, 606 carry animal-preparation tags, 316 carry animal-study assessments, 119 link to feedback pathways, 95 carry an afferent classification, and 363 carry robot-or-simulation translation notes. DOI identifiers are present for 892 papers, and 318 have attached-PDF indicators. These fields overlap and measure the available curation, not independent experimental support. Within that category, 730 papers are connected by 11,130 directed citation pairs, and the Louvain partition resolves 18 named topic clusters. The application distributes the records, citation graph, and cluster assignments together. Figure 5.1 illustrates the map; Appendix E shows the table view.

The table view combines text search with filters for preparation, pathway, afferent class, record type, archive status, year, and note or PDF availability (Figure 5.1A). Selecting a record opens its source information, curation notes, and proposed robot-or-simulation translation. The table, pivot, and citation views use the embedded snapshot offline; the separate online-search tab queries OpenAlex for candidate additions and marks DOIs already represented.

The bubble map supports both publication-year versus citation-count axes and a topic landscape based on the citation-network layout. Its controls select the layout, display style, record grouping, citation-neighborhood depth, and click action (Figure 5.1B). Selecting a paper as the focus identifies it in the left panel (C), which exposes its curation note and proposed robot-or-simulation translation (D). The plotted neighborhood can extend

through one, two, or three citation links in either direction (E). Direct connections retain color, while more distant papers are gray. In the illustrated Grillner (1985) neighborhood, 57 in-corpus papers cite the focus, the focus cites 13 in-corpus papers, and 555 papers lie within two links. Green triangles identify incoming citations and red circles identify the focus paper's references. These synapse-like symbols encode bibliographic direction; they do not assert excitatory or inhibitory physiological effects.

The connection list separates papers citing the focus from its references (Figure 5.1F). Selecting a listed paper makes it the new focus, allowing a reader to follow the development of an idea without repeatedly searching for individual titles. The reference-only and cited-by views isolate either direction, and the detail action opens the underlying record. A researcher can therefore start from a known experiment, inspect earlier studies supporting it, follow subsequent models or experiments citing it, and read the associated translation notes before choosing a candidate pathway for simulation.

5.3.1 AI-Assisted Literature Triage

I evaluated AI-assisted literature identification, summarization, and categorization in work presented at Neuroscience 2024 [**?**]. The comparison included ResearchRabbit, Litmaps, txyz, SciSpace, and MirrorThink. Some tools helped identify related papers or summarize straightforward material, but categorization remained inadequate: none received a good rating, and only txyz was rated adequate. Complex models and nuanced physiological findings required substantial human correction. The tools therefore offered limited help with the central curation task of accurately assigning papers to database categories. I

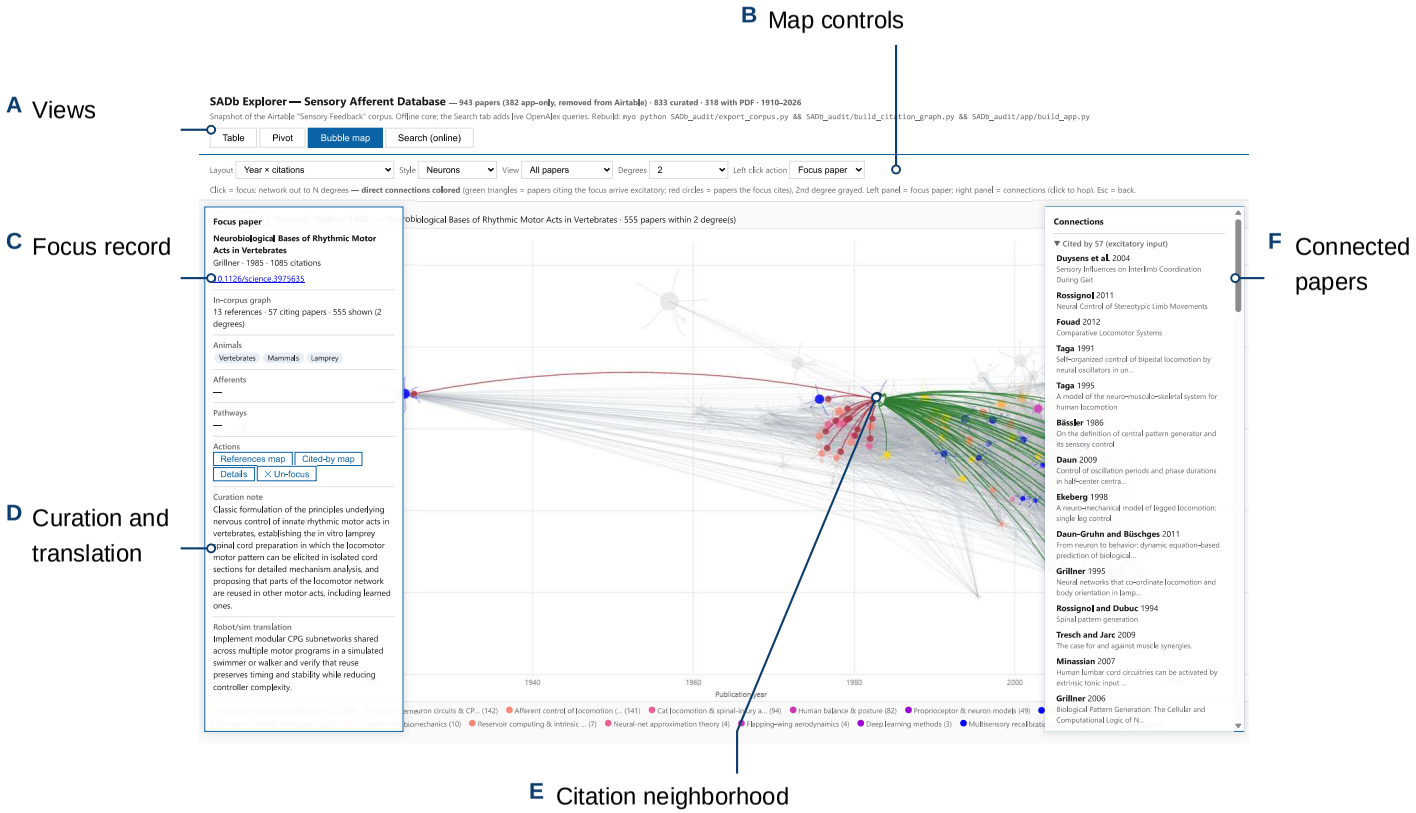


Figure 5.1: SADb Explorer focused on Grillner (1985): views (A), map controls (B), focus record (C), curation and translation notes (D), citation neighborhood (E), and connected papers (F).

retained responsibility for inclusion, classification, and deletion decisions after checking the source.

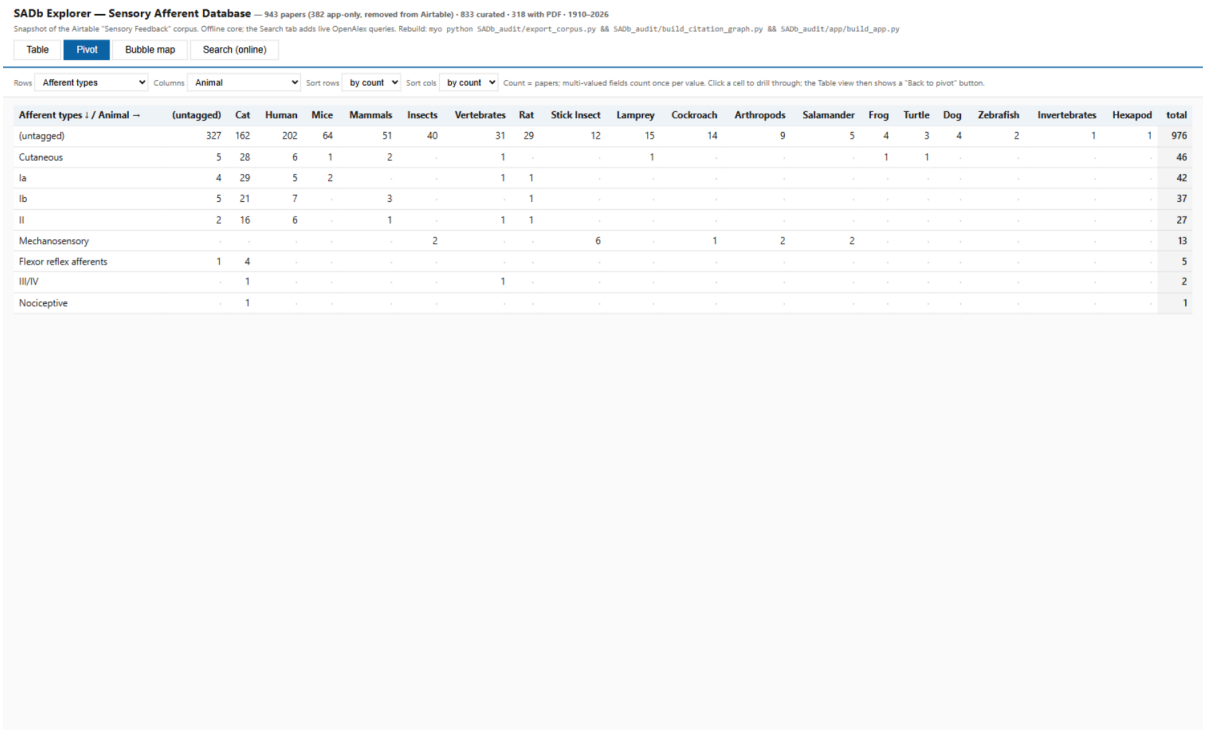
5.4 Discussion

5.4.1 Robot and Simulation Use

SADb supports sensory-pathway selection for Chapter 6. Its afferent and pathway tables connect candidate signals, connection signs, and phase dependence to the cited evidence (Tables E1 and E2). Receptor dynamics, gains, delays, and species-specific assumptions still require explicit selection and testing. The model records connect proposed implementations to their literature, including two-level rhythm generation and pattern formation [**?**], muscle-reflex walking [**?**], and the laboratory walker lineage [**?**].

The pivot view counts papers by two selected dimensions, including topic cluster, decade, primary author, animal preparation, feedback pathway, and afferent classification. Figure 5.2 crosses afferent classes with animal preparations: each cell counts papers carrying both tags. Rows and columns can be ordered alphabetically or by count. Selecting a cell opens the corresponding records in the table view, and a return control restores the pivot. The pivot is computed over the 943 entries in the Papers category, including archived entries, rather than all 1,393 records. Because a paper can carry several animal or afferent tags, it can contribute to several cells. The counts describe the distribution of curated records, not independent experimental observations or the strength of evidence for a pathway.

The feedback-pathway records extend beyond the mammalian core of Table E2: paired



SADb Explorer — Sensory Afferent Database — 943 papers (382 app-only, removed from Airtable) · 833 curated · 318 with PDF · 1910–2026

Snapshot of the Airtable "Sensory Feedback" corpus. Offline core; the Search tab adds live OpenAlex queries. Rebuild: `myo python SADb_audit/export_corpus.py && SADb_audit/build_citation_graph.py && SADb_audit/app/build_app.py`

Table | Pivot | Bubble map | Search (online)

Rows: Afferent types | Columns: Animal | Sort rows: by count | Sort cols: by count | Count = papers; multi-valued fields count once per value. Click a cell to drill through; the Table view then shows a "Back to pivot" button.

Afferent types ↓ / Animal →	(untagged)	Cat	Human	Mice	Mammals	Insects	Vertebrates	Rat	Stick Insect	Lamprey	Cockroach	Arthropods	Salamander	Frog	Turtle	Dog	Zebrafish	Invertebrates	Hexapod	total
(untagged)	327	162	202	64	51	40	31	29	12	15	14	9	5	4	3	4	2	1	1	976
Cutaneous	5	28	6	1	2	·	1	·	·	1	·	·	·	1	1	·	·	·	·	46
Ia	4	29	5	2	·	·	1	1	·	·	·	·	·	·	·	·	·	·	·	42
Ib	5	21	7	·	3	·	·	1	·	·	·	·	·	·	·	·	·	·	·	37
II	2	16	6	·	1	·	1	1	·	·	·	·	·	·	·	·	·	·	·	27
Mechanosensory	·	·	·	·	·	2	·	·	6	·	1	2	2	·	·	·	·	·	·	13
Flexor reflex afferents	1	4	·	·	·	·	·	·	·	·	·	·	·	·	·	·	·	·	·	5
III/IV	·	1	·	·	·	·	1	·	·	·	·	·	·	·	·	·	·	·	·	2
Nociceptive	·	1	·	·	·	·	·	·	·	·	·	·	·	·	·	·	·	·	·	1

Figure 5.2: SADb Explorer pivot table by afferent classification and animal preparation. Each cell counts papers carrying both tags and links to those records.

excitation and inhibition records cover the chordotonal organs and trochanteral hair plates of insect legs, a load-feedback record covers the campaniform sensilla, and further records hold fictive locomotion without sensory feedback (12 linked papers), group III/IV fatigue, and one paper each for somatosensory, vestibular, and visual contributions to postural control. The controller of Chapter 6 implements the mammalian subset.

The balance, posture, and multisensory-recalibration clusters provide entry points for the vestibular and visual extensions proposed in Section 7.2.5. Preparation-based searches also let new users approach the literature through the experiment they intend to perform.

5.4.2 Further Improvements

Classification remains incomplete within the application's 943-entry Papers category: 95 carry an afferent classification, and 119 link to named pathways. Completing these fields would improve comparisons across preparations and help expose where a proposed pathway depends on a small or indirect body of evidence. Missing source text should remain explicit until the paper can be checked; model and review records should link to the primary experiments supporting their assumptions.

Further curation should connect pathway records to the papers that establish their signals, connection signs, and task dependence, and connect model records to the pathways they implement. Vestibular and visual contributions to posture require expanded classifications and source coverage before the database can support the additional sensory channels proposed in Section 7.2.5. New records should include a source identifier and a substantive curation note, with incomplete classifications flagged for review. These improvements

would make the resource more useful for selecting experiments and constructing testable neural models.

The application makes physiological evidence accessible for controller design, but selecting a pathway is only the beginning of a neuromechanical experiment. Its encoding, connection sign, gain, and interaction with the body must still be implemented and tested. Chapter 6 develops the models and simulation tools needed to make that next step explicit.

6 Synthetic Nervous System Simulation

The actuator and torque models of Chapters 3 and 4 describe the robot's mechanics. The next task is to investigate how those mechanics interact with neural control. My initial AnimatLab walker and subsequent W2L integration work in MATLAB exposed the difficulty of distinguishing a controller limitation from an implementation or model-conversion error. I developed neuromechanical models, graphical neural blocks, simulation interfaces, and a staged evaluation procedure to make those distinctions testable. This chapter follows that progression, then presents tests of the circuits, interfaces, and coupled models.

6.1 The AnimatLab Bipedal Walker

For a course project in ME 610: Neuromechanical Modeling, Morrow and I combined the bipedal mechanics of **?** with the two-level central pattern generator (CPG) of **?** in AnimatLab [**?**]. The mechanical model had a pelvis and bilateral thigh, shank, foot, and toe segments, with sagittal-plane joints and antagonist Hill-type muscles at the hip, knee, and ankle. I added a passive toe joint. The initial controller used one rhythm generator (RG) to control both sides, with hip and knee/ankle pattern-formation (PF) groups for each leg. We added afferent feedback at the motoneuron (MN) level. Figure 6.1 shows the model and controller organization.

We adjusted applied currents, synaptic weights, and muscle parameters by hand. The suspended model developed alternating muscle activity and stepping-like motion, presented at the 2021 NIH BRAIN Initiative Investigators Meeting [**?**]. It did not achieve

weight-bearing walking.

I subsequently replaced the shared rhythm generator with one rhythm generator on each side and linked the two through commissural pathways. The bilateral pattern-formation and motor organization was already present in the original model. I added afferent feedback to the RG and PF layers, and heel- and toe-contact signals during the subsequent development. This later model is referred to as the Walker Two-Layer (W2L) model. Further circuit schematics and implementation records are provided in Appendix D.

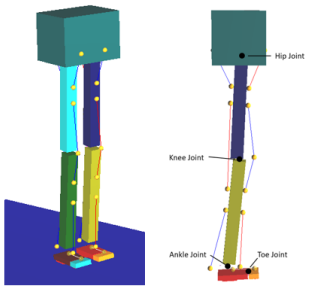
The AnimatLab model established that the combined circuit could generate alternating activity and move the suspended limbs. To investigate the contribution of sensory feedback and actuator mechanics, I needed to vary those components while retaining a specified neural organization and mechanical support condition. Extending that approach to the human-scale Gait2392 model and its BPA-routed robot counterpart also required checking joint directions, muscle paths, actuator forces, and the signals exchanged with the controller.

Figure 6.1 precedes the later split into one rhythm generator per side and the addition of afferent and foot-contact feedback to the RG and PF layers.

6.2 Attempting W2L Integration in MATLAB

I next pursued W2L integration in MATLAB, using Simulink for the neural controller and Simscape Multibody for the mechanical model. Bringing these parts into one environment required explicit interfaces for neural activation, muscle force, joint motion, and sensory

(A)



(B)

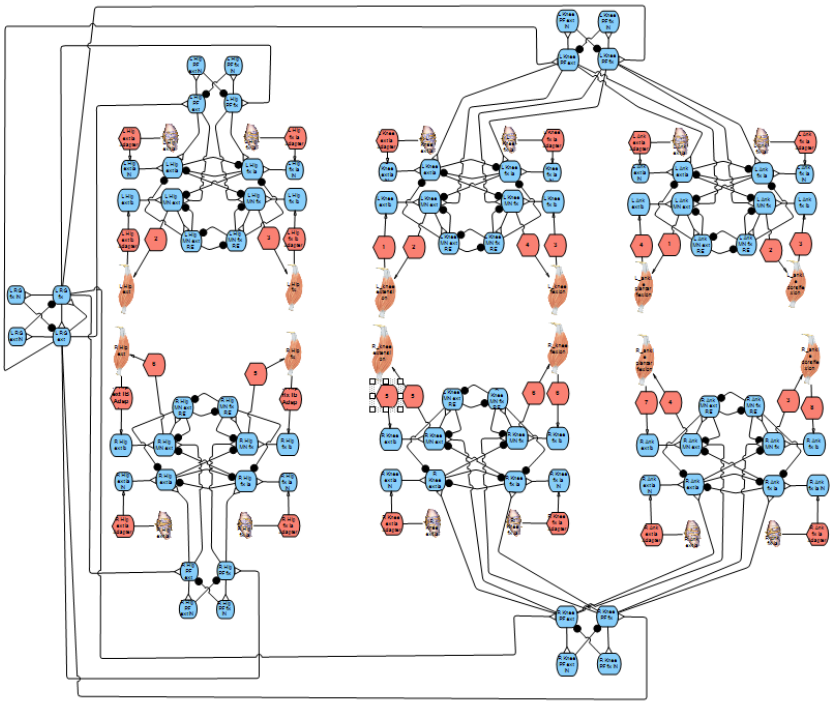


Figure 6.1: Original AnimatLab walker (A) and shared-rhythm-generator controller (B), developed by Morrow and Bolen. The four-neuron RG drives hip and knee/ankle pattern-formation groups on both sides.

feedback. A mechanism that imported successfully and a neural circuit that oscillated separately were necessary components, but neither established that the coupled model retained the intended behavior.

The MATLAB work therefore became a sequence of component and coupling tests. Selected neural implementations reproduced their reference responses, and mechanical import provided models that could be assembled and simulated. More extensive tests also exposed limits: a full-network export agreed with its reference over an early transient but later diverged, and one Simulink realization reached a steady, nonoscillating state under constant neural drive while its NumPy reference oscillated (Appendix D.2). Those results did not establish a complete W2L walking implementation in MATLAB. They showed why importing a model or reproducing a short trace could not substitute for testing the neural dynamics, actuator action, and feedback loop separately.

This experience motivated two complementary development routes. SNS-Simscape makes neural and mechanical connections inspectable in block diagrams, while the scripted MuJoCo/SNS-Toolbox models support repeated trials, parameter searches, and controlled pathway changes. The MuJoCo–Simulink bridge connects the two approaches without requiring the mechanical model to be rebuilt for every controller experiment.

6.3 Knowledge Gaps and the New Tools

The knowledge gap concerns how to attribute movement in a coupled neuromechanical model. When a proposed circuit fails to produce the expected behavior, the cause may lie in its neural organization, its sensory encoding, the actuator or mechanical model, or the

numerical implementation. For the humanoid models studied here, I needed to determine whether selected neural mechanisms survived transfer between implementations and how their effects changed when coupled to human musculotendon mechanics or artificial muscles. That requires common inputs, observable intermediate signals, and interventions that hold the other parts of the model fixed.

I addressed this gap by constructing models in which the neural pathways, actuator characteristics, sensory inputs, and support conditions can be specified and varied, together with tests that distinguish circuit behavior from implementation and coupling errors. The tools include the Li–Deng/W2L and Gait2392 model implementations, the SNS-Simscape library and GUI, neural and BPA-actuator coupling to a SolidWorks-derived Simscape assembly, the MuJoCo–Simulink interface, and experiment drivers for staged tuning and controlled interventions. Literature-derived and spiking-neuron implementations support comparisons of circuit mechanisms and coding assumptions. Table 6.1 summarizes the models and tools developed in this work and what each is used for. The following sections describe their implementation and demonstrations. The feedback experiments also provide a limited mechanistic result: removing sensory pathways can disrupt bilateral contact coordination in a fixed configuration, identifying a dependence that would not be apparent from its intact motion alone.

OpenSim supplied musculoskeletal geometry, MyoConverter translated models to MuJoCo, SNS-Toolbox supplied neural formulations, and Simulink/Simscape supplied graphical neural and mechanical modeling [**????**]. My contributions are the model implementations, custom blocks, actuator integrations, and evaluation procedures built with those packages. Matched-input tests check selected circuits before their outputs are

Table 6.1: Models and tools developed for neuromechanical experiments.

Model or tool	**What it is used for**
Li–Deng/W2L walker and controller	Combine walker mechanics and neural control in AnimatLab and MuJoCo with SNS-Toolbox; compare matched implementations of the contact-driven neural core in NumPy and Simulink.
Gait2392-based models	Compare alternative neural controllers on the human reference model and a robot counterpart with BPA routing and the force model from Chapter 3.
SNS-Simscape block library and editor	Build neural pathways graphically from sensory, neuron, synapse, activation, and muscle-force blocks, including custom BPA blocks.
SOLIDWORKS-derived robot assembly	Couple a CAD-derived Simscape Multibody mechanism to a neural controller and BPA actuators using the force model.
MuJoCo–Simulink integration	Run an SNS-Simscape controller with a MuJoCo mechanical model through the MathWorks bridge, exchanging actuator commands and sensory measurements.
Curriculum and experiment drivers	Tune controllers in stages, vary support and sensory pathways, apply neural stimuli or mechanical perturbations, and record neural and mechanical responses.
Literature circuits and spiking variants	Compare selected published neural mechanisms and the behavior of graded and spiking implementations.

interpreted as control of a mechanical model.

6.4 Methods: Software Tools

The tools provide three complementary capabilities: constructing a neuromechanical model, testing its component behavior and signal exchange, and evaluating controlled changes in neural or mechanical conditions. Figure 6.2 shows the model-construction routes and neural-controller interfaces used in this chapter. The broader simulation-to-hardware workflow is discussed in Chapter 7 (Figure 7.4).

6.4.1 Model Conversion and Actuator Coupling

My first conversion attempt used the robot model, which combined BPA routing on the Gait2392 skeleton with unchanged human muscle parameters. MyoConverter's muscle-equilibrium checks rejected infeasible combinations of routed length and muscle parameters. I therefore separated conversion of the human reference from reconstruction of the robot routing. This distinction matters because an error caused by an artificial-actuator route should not be interpreted as a property of the original human model.

The human-model conversion also required specific repairs. Three small hip rotators per side, the quadratus femoris, gemelli, and piriformis, were disabled at model load. The converted model retains 92 actuator entries, but these six contribute no force. The converted knee retains the OpenSim flexion-negative sign convention. Spatial frame transformations and individual-actuator tests check joint motion and torque directions; the CAD-derived model uses the anatomical frame specified in Appendix D.5.4.

The human-model repairs addressed coupled moving muscle-path points, patella tracking, joint limits, and unintended contacts. Conflicting limits were removed from coupled path-point joints while anatomical driver limits were retained; rectus femoris required a patella-tracking via point, and contact was restricted to intended foot–ground pairs. These choices restore specified geometry and contact behavior rather than establish general model equivalence. Appendix D.6 records the repairs and force comparison.

Comparisons with the OpenSim attachment paths and single-actuator tests checked the transferred geometry and force directions. They do not establish that every muscle reproduces the source force throughout its range or that the controller produces walking. The neural and mechanical models exchange signals at a 2 ms solver time step. Appendix D.6 records the individual repairs and their scope.

For artificial-muscle simulations, I implemented an actuator-force calculation that evaluates the quasi-static BPA model from Chapter 3 at the current simulated actuator length and commanded pressure. The calculated tensile force acts along the model's specified tendon route. Length, velocity, force, and foot-contact measurements are then returned to the neural controller through sensory channels. This closes the neural-command–actuator-force–sensory-feedback loop. The calculation represents the measured static force dependence; it does not add valve-flow or pressure-transient dynamics. Routing geometry and dynamic force transmission therefore require their own checks beyond the isometric force fit.

The SADb of Chapter 5 supplies the physiological context for selecting Ia, II, Ib, and contact pathways. A database link identifies the supporting observation; it does not by itself determine a numerical gain or guarantee the same function in another model.

The implemented encoding, connection sign, phase gating, and tested response therefore remain explicit parts of the controller specification.

The resulting model set separates the mechanical representation from the neural controller. The human Gait2392 reference retains biological-muscle actuators; the robot counterpart uses custom artificial-muscle paths and BPA forces. The Li–Deng/W2L walker and the SolidWorks-derived robot are separate mechanical models, rather than interchangeable versions of Gait2392. Within the Gait2392-based pipeline, the controller comparisons include the full graded network, a version with selected connections removed, a version organized around six grouped muscle-activation patterns, and a version adapting the W2L controller organization; selected spiking counterparts change neural dynamics while retaining a continuous motor-activation interface. These alternatives make controller and actuator assumptions available for comparison. Appendix D distinguishes their configurations and reported support conditions.

6.4.2 Block Programming Neural Pathways

The SNS-Simscape library implements neuron, synapse, afferent-encoder, activation, and muscle-force elements as editable Simulink blocks. The library retains the conductance-based neuron and synapse formulation supplied by SNS-Toolbox; it does not introduce a new neuron model. Verification uses selected circuits under matched inputs. Figure 6.3 shows the core blocks and graded-neuron signal flow.

I imported OpenSim-derived and SolidWorks-derived mechanisms into Simscape Multibody and coupled the SolidWorks robot assembly to a neural controller and BPA

actuators using my force model. The block library also supports reduced-order demonstrations in which the sensory pathway and mechanical response can be inspected directly. These smaller models provide interpretable checks of the components used in larger assemblies.

6.4.3 Connecting Simulink Controllers to MuJoCo

The MathWorks MuJoCo–Simulink bridge provides an alternative to importing the mechanical model into Simscape. I connected SNS-Simscape neural outputs to the bridge's actuator-command interface and returned muscle length, velocity, force, and joint measurements to the controller. This arrangement allows a block-programmed neural controller to act on an existing MuJoCo model. The integration tests compare activation and sensor readback against a Python reference before interpreting the coupled motion. The bridge supplies the simulation interface; the neural and actuator integrations and their verification are the work described here.

6.4.4 A Curriculum for Tuning and Controlled Experiments

The larger models require a tuning procedure as well as a working interface. I organized the evaluations as a staged curriculum, beginning with suspended rhythm generation, adding sensory feedback and standing requirements, and progressing to supported ground evaluations (Table 6.2). This ordering separates a failure to generate rhythmic commands from a failure to support or coordinate the mechanical model. Free network gains, which scale connection strengths and sensory contributions, were searched with a tree-structured

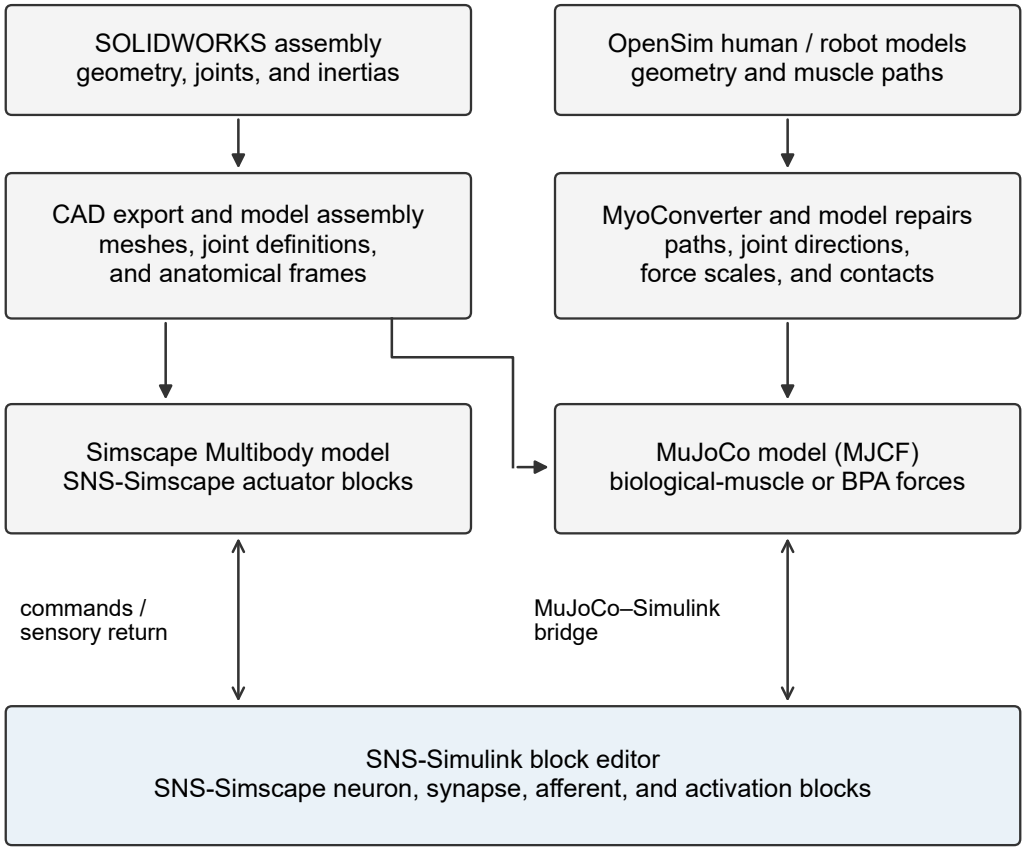


Figure 6.2: Model construction and neural-controller interfaces. SOLIDWORKS-derived geometry supplies Simscape Multibody and separately assembled MuJoCo models; OpenSim models follow the MyoConverter and repair route. The SNS-Simulink block editor assembles controllers from SNS-Simscape library blocks. Commands and sensory measurements couple these controllers to Simscape or, through the MuJoCo–Simulink bridge, to MuJoCo.

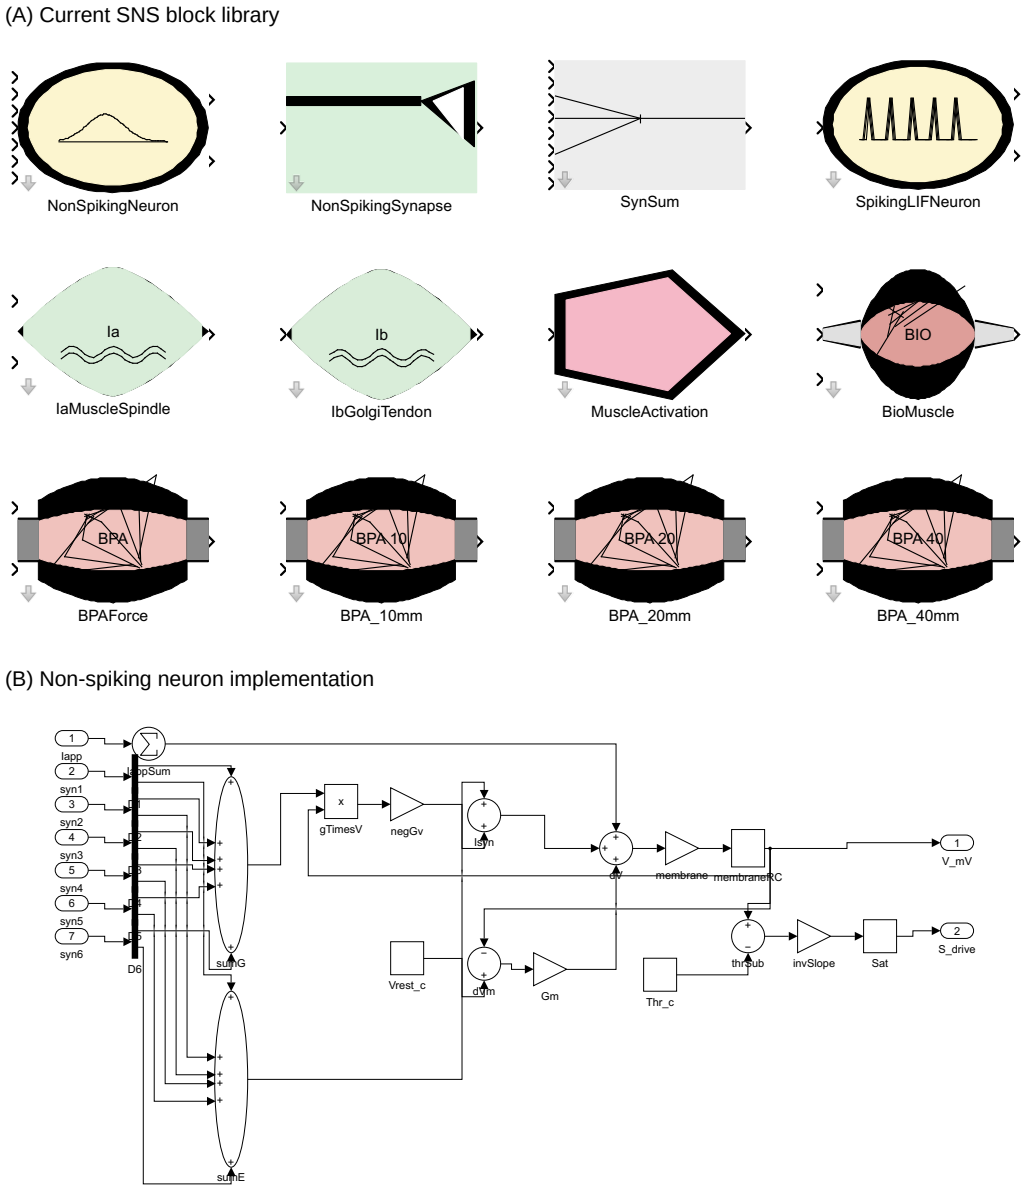


Figure 6.3: Current SNS-Simscape library (A) and native non-spiking neuron implementation (B). The twelve library blocks include biological-muscle and diameter-specific BPA elements. Summed currents drive membrane integration; the threshold and saturation bound neural output.

Parzen estimator. This sequential search uses earlier simulation results to propose new parameter combinations; each candidate is then evaluated through simulation. Saved configurations and fixed evaluation protocols allow the candidate to be replayed before a new condition is introduced.

The experiment drivers then hold a selected controller fixed while changing sensory pathways, drive, mechanical support, or external perturbations. Finalists are also replayed under small parameter perturbations. Gait scoring requires detected cycles before a score is interpreted as movement quality; unsuccessful runs cannot be promoted by a reserved numerical failure value. The reference-library comparison tests transfer beyond the tuning reference. These procedures provide a repeatable way to identify where a controller works and which changes disrupt it, rather than making a single favorable trajectory the endpoint of tuning. Detailed scoring and configuration records are in Appendix D.3.

Table 6.2: Five-stage controller-tuning curriculum. Each selected controller seeds the next stage. Air-stepping stages are suspended; ground stages retain mechanical support. The table describes the training tasks and evaluation measures.

Stage	Training task	Condition introduced	What is evaluated
1	Deafferented air stepping	Legs suspended; sensory feedback disabled	Rhythm generation and swing depth
2	Afferented air stepping	Heel, toe, and load feedback enabled	Rhythm and swing with sensory feedback
3	Standing balance	Standing posture and push perturbations	Ability to maintain posture during a push
4	Ground walking	Ground contact and forward-progression task	Joint kinematics against a human reference cycle
5	Contact-driven ground walking	Contact-driven pathway variant	Joint kinematics against a human reference cycle

6.5 Results: Verification of the Implemented Interfaces

The tests address three distinct questions: whether a circuit produces its expected response, whether implementations agree under matched inputs, and whether the neural and mechanical models exchange signals in a closed loop. The larger evaluation sets and their limitations are collected in Appendix D; neuron equations and calibration details remain in Appendix B.

The demonstrations below test the models and tools summarized in Table 6.1. Matched-input comparisons check neural implementations, the loaded-arm experiment tests corrective feedback, and bridge tests check command and measurement exchange. Saved-controller replays and pathway removals then test whether the coupled experiments can be repeated and compared.

6.5.1 Coupled Block-Model Tests

The library assembles Ia stretch excitation, reciprocal inhibition, and Ib autogenic inhibition in a reduced-order antagonist knee model (Figure 6.4C). After the model reaches a steady operating point, a 2 nA, 30 ms current pulse at the extensor Ia sensory-neuron input evokes 1.09° of extension, peaking 111 ms after pulse onset. The knee then returns toward the matched no-stimulus trajectory, meeting the specified recovery thresholds after 6.355 s. The displayed configuration broadens the MN and Ia synaptic saturation ramps; the original narrow-ramp settings maintained oscillations. This is a transient test of the implemented feedback and actuator equations, rather than a physiologically calibrated tendon-tap model. Appendix D.2.1 gives the configuration and numerical checks.

A second demonstration tests feedback in an arm holding a beverage container as it is loaded (Figure 6.5). With the reflex enabled, arm sag was 2.00°, compared with 7.85° when it was disabled. As the forearm sags, biceps stretch feedback excites the biceps motoneuron and reciprocally inhibits the triceps motoneuron. The separate Ib pathways provide autogenic inhibition of each muscle. The reduced sag demonstrates a corrective effect of the implemented sensory pathways in the coupled neural and actuator blocks.

6.5.2 Matched Neural Implementations

The contact-driven W2L rhythm core was implemented in NumPy and native Simulink blocks from a common connectivity specification. Both implementations received the same input protocol. Burst periods matched at 1.000 s, all ten expected bursts were reproduced, and voltage-trace correlation was at least 0.99996 (Figure 6.6). The agreement is specific to this circuit and input protocol.

Removing the phasic heel-contact inputs caused the circuit to latch in both implementations. The alternating activity therefore depends on the specified sensory drive.

A complementary neural-only test distinguishes mechanisms that can produce a missing motor burst. In the reduced Rybak implementation, pattern-formation stimulation suppresses flexor motor output while the rhythm generator continues; rhythm-generator stimulation shifts the subsequent phase. Reproducing that distinction demonstrates the use of controlled interventions to test circuit organization before adding mechanical coupling (Appendix D.2.2).

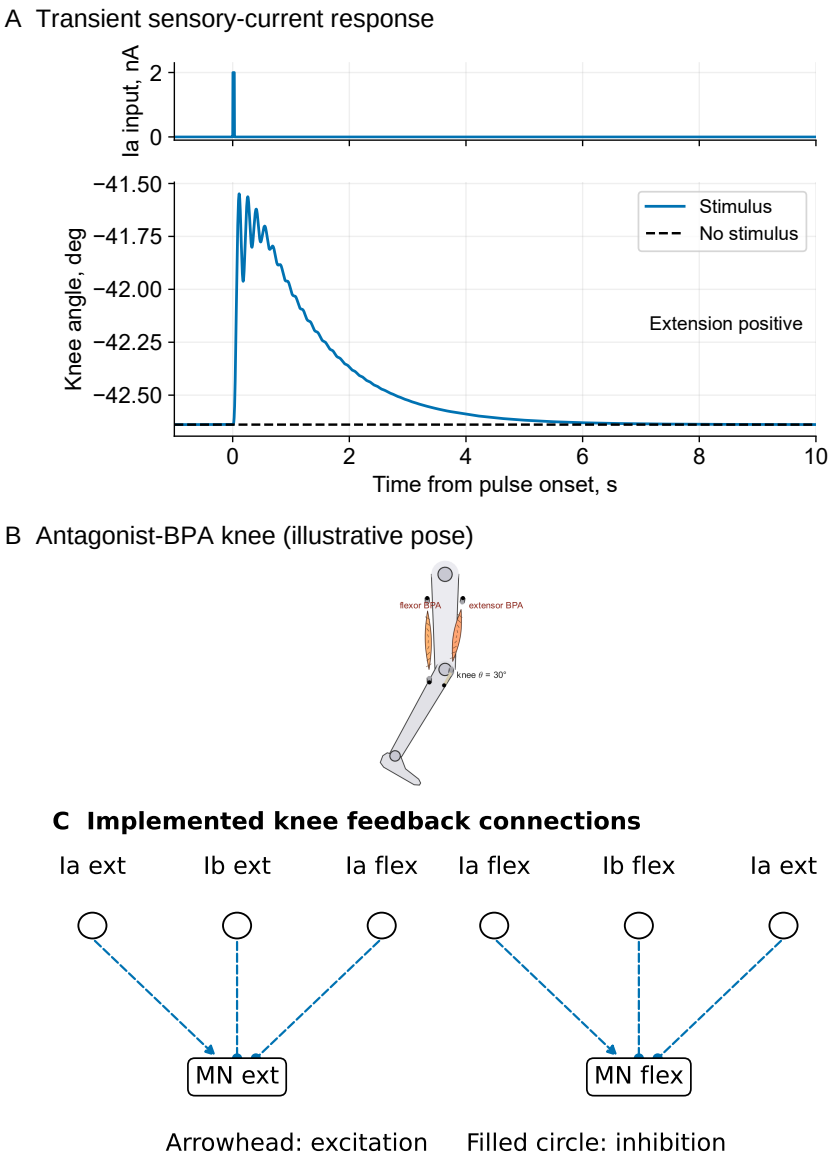


Figure 6.4: Transient knee response to a 2 nA, 30 ms sensory-current pulse after settling (A), antagonist-BPA leg schematic in an illustrative pose (B), and the six implemented feedback connections (C). Flexion is negative. The graded-response configuration uses broadened MN and Ia synaptic saturation ramps; its response is compared with a matched no-stimulus control.

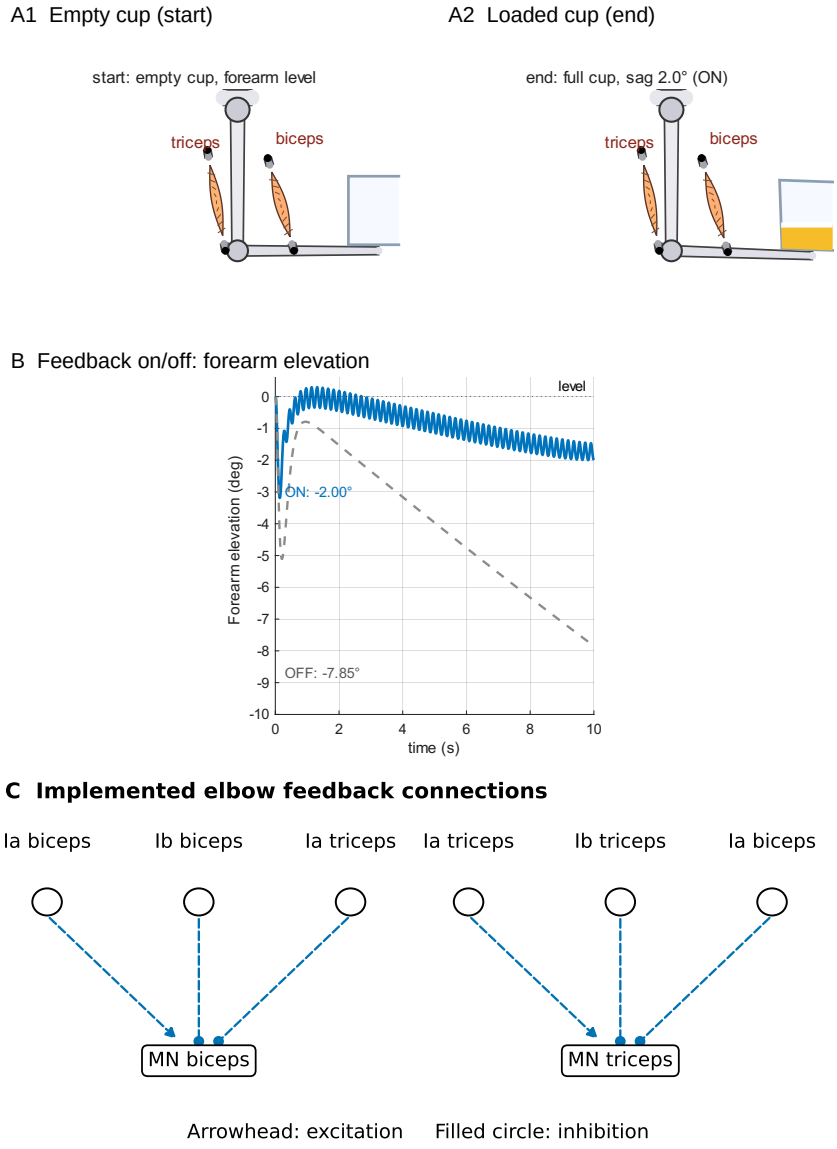


Figure 6.5: Loaded-arm demonstration: empty and loaded cup positions (A1, A2), feedback-on/off kinematics (B), and all six implemented feedback connections (C).

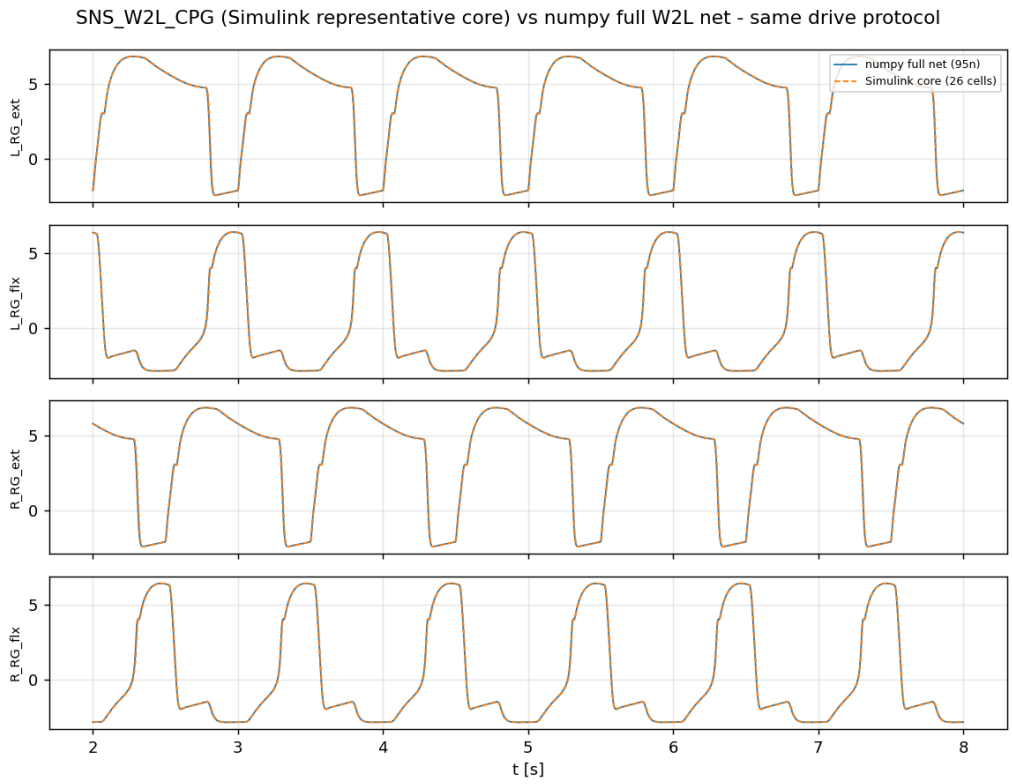


Figure 6.6: W2L rhythm-generator voltages under matched inputs in NumPy and Simulink. Burst periods and counts agree; full-trace correlation is at least 0.99996. The residual 2 ms lag is one discrete solver step. The test concerns the driven rhythm core, not walking mechanics.

6.5.3 Layer-Specific Stimulation

A missing motor burst can arise from a disturbance of the rhythm generator or from suppression farther downstream. I applied stimuli separately to the RG, PF, and MN layers of a reduced Deng-derived circuit [**??**]. A resetting deletion suppresses expected motor activity and shifts the subsequent rhythm; a non-resetting deletion suppresses activity while preserving the original schedule. Matched baseline and stimulus runs distinguish the rhythm-generator clock from the motor output.

Stimuli lasting two baseline periods suppressed two expected flexor motor bursts at each layer. RG stimulation shifted both RG and motor timing by 0.211 cycles, whereas PF stimulation and direct MN inhibition preserved the original schedule. A separate one-second PF stimulus shifted motor timing by 0.208 cycles while leaving the RG clock unchanged. Thus, a motor phase shift alone did not establish resetting of the rhythm generator. The smaller-current sweep produced no observable deletions. Appendix D.2.3 gives the protocols, all stronger-current controls, and representative neural responses. These are controlled tests of a translated circuit coupled to a suspended model, rather than a replication of the original rat's locomotion.

6.5.4 Conversion Checks and Muscle-Force Agreement

The conversion checks test the quantities needed to interpret a neural-control experiment: the direction of each muscle's joint action, the coupled motion of path points, the force scale, and the intended ground contacts. They exposed reversed hip actions, inconsistent path-point limits, unintended contact pairs, and incorrect force scales in individual ac-

tuator entries. Correcting those faults made the converted human model usable for the subsequent controller tests. Appendix D.6 gives the individual repairs and the full force comparison.

At matched poses and OpenSim static-optimization activations, eight of twelve sampled muscles had force-prediction R^2 between 0.87 and 0.98 and mean forces within 10 % of the reference. Rectus femoris and biceps femoris short head had negative R^2, indicating errors greater than a constant mean-force prediction. These results support selected force predictions while identifying remaining muscle-specific differences. Because the comparison sets pose and activation frame by frame without reproducing trajectory velocities or tendon dynamics, it does not establish dynamic equivalence between the muscle formulations. That distinction matters when interpreting feedback based on muscle force.

6.5.5 Closed-Loop Model Coupling

The MathWorks MuJoCo–Simulink bridge exchanges activation commands and mechanical measurements between a Simulink controller and an existing MuJoCo mechanical model. I paired this bridge with SNS-Simscape and tested single-muscle activation, different control rates, and sensor readback against a Python integration. Readback was bit-exact over the scored window. An exported neural network also drove all 92 actuator channels during a 5 s run without afferent return; a separate closed-loop reflex experiment exercised both commands and sensory measurements. These tests verify signal exchange through the interface. Under constant neural drive, one Simulink circuit realization reached a steady, nonoscillating state while its NumPy reference oscillated, showing why successful

signal exchange must be followed by circuit-specific dynamical checks before transferring tuned parameters.

6.5.6 Testing Alternative Controllers on Gait2392

The Gait2392-based evaluations use the same reference-library and pathway-removal procedures to compare alternative neural organizations. The full graded, pruned, six-synergy, and W2L-layout configurations show different responses to feedback removal and different agreement with the tuning and library references. Removing Ia feedback eliminated detected cycles in the tested production configuration, whereas other removals and other network organizations had different effects. The library comparisons also exposed asymmetric activity and limited transfer beyond the tuning reference (Appendix D.3.1). This demonstrates that the models and experiment drivers can discriminate controller dependence and transfer limits; the comparison does not establish normal human gait.

A replay of the saved human-model controller produced ten detected cycles per leg. Figure 6.7 compares the joint trajectories directly with the Gait2392 inverse-kinematics reference and its associated ground-reaction forces. The model's mean stance fractions were 0.71 and 0.65 on the right and left, compared with 0.61 in the reference. Substantial hip and knee timing differences remain visible despite the detected cycles. The overlay therefore complements the scalar gait score by showing which parts of the movement disagree. Appendix D.3.2 reports the angle errors and alignment method.

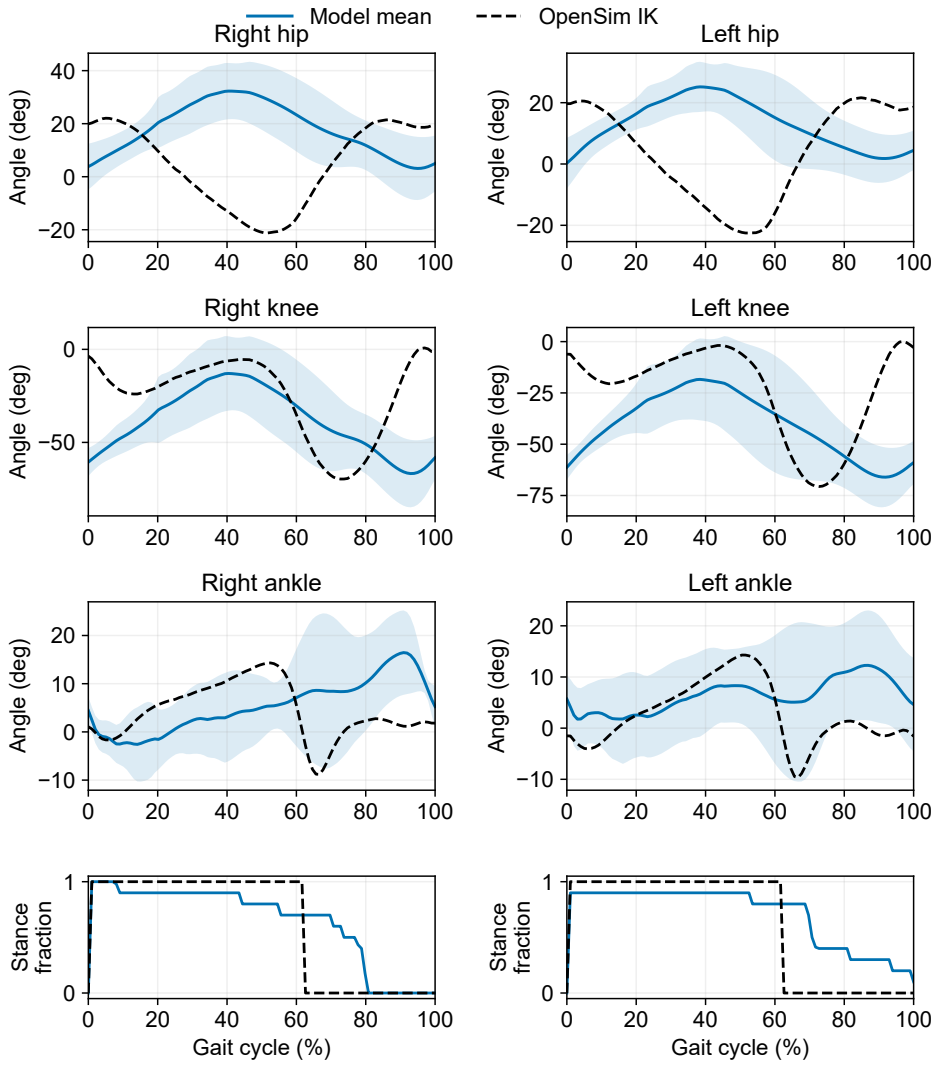


Figure 6.7: Supported human-model joint motion compared with the Gait2392 inverse-kinematics reference. Model curves show the mean and one standard deviation across ten cycles per leg. Each leg is aligned to its own foot-loading onset; no angle offset, amplitude adjustment, or fitted phase shift is applied. Bottom strips compare model contact probability across cycles with reference stance and swing from ground-reaction forces, using a 50 N threshold. Separate leg alignment does not establish bilateral phase agreement.

6.5.7 Feedback Contribution at Fixed Controller Settings

The recovered W2L controller was tested with all saved effective gains delivered to the compiled network. With controller, drive, and support settings fixed, disabling all sensory feedback shifted median bilateral contact phase from 0.498 to 0.223 cycles. A phase of 0.5 cycles indicates alternating foot loading; the supported-stepping criterion accepts 0.35–0.65 cycles. Forward progression and joint cycling persisted, but the feedback-disabled condition failed this criterion for supported stepping (Figure 6.8). Over the 4.5 s to 30 s scored window, forward displacement was 2.239 m with feedback and 1.996 m without it. Thus, the demonstrated feedback contribution is to bilateral contact coordination, rather than the existence of oscillation or forward motion. Additional pathway deletions and stimulus tests are summarized in Appendix D.5.2.

6.5.8 Repeatability and Mechanical Support

A separate W2L configuration reproduced its supported-stepping trajectory in four independent replays, and all twelve reruns with 1 % parameter perturbations retained supported stepping. This demonstrates that a saved controller can be replayed and tested in its local parameter neighborhood. The repeatability tests and the pathway-removal panels use separately specified support settings; their event counts are therefore reported separately. Appendix D.5.1 retains the event counts, bilateral phase, displacement, speed, and support accounting.

Mechanical support remains part of these results. Removing a constant weight-support term did not remove the position-control harness, and contact duty remained below the

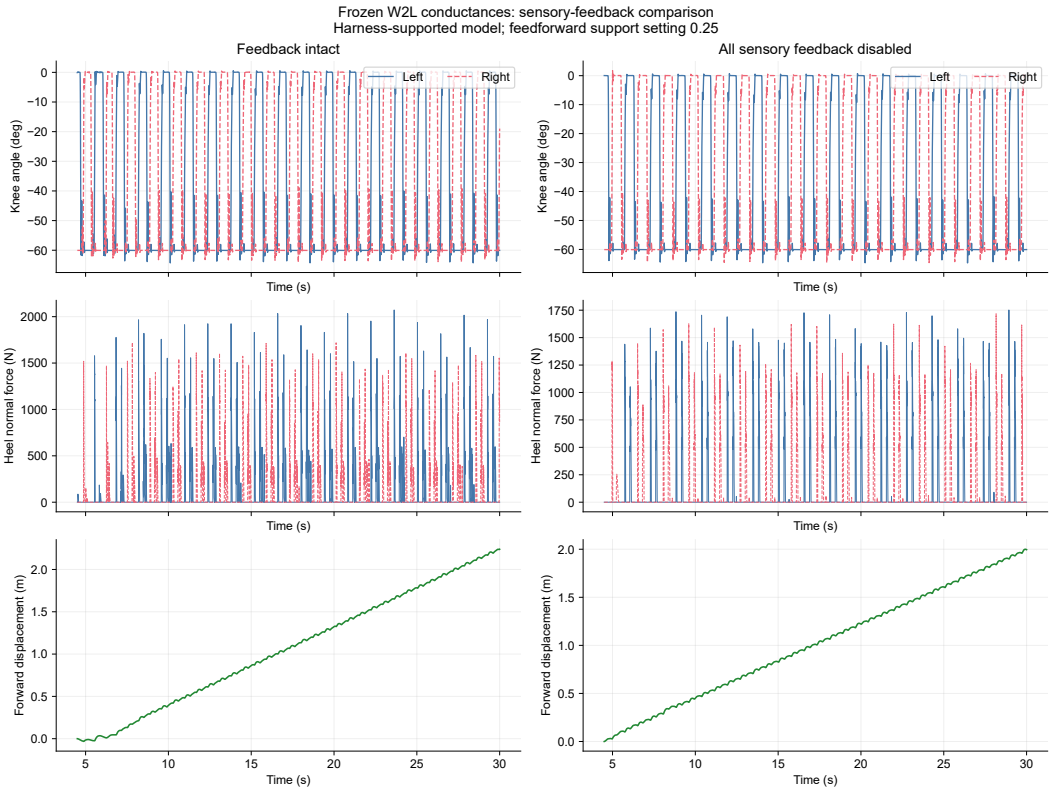


Figure 6.8: Sensory-feedback comparison with fixed W2L conductances, drive, and support settings. Knee angles, heel contact forces, and displacement are shown with feedback intact (left) and all sensory feedback disabled (right). Cycling persists in both conditions; the feedback-disabled condition fails the bilateral contact-phase criterion.

strict walking criterion. A separate CAD-derived model completed a coupled stance demonstration under substantial rig support and reference actuator tensions. The model-specific support and actuator limitations are documented in Appendix D.5.4. No harness-free walking was demonstrated.

A newer CAD-model comparison tests whether additional abductor control prevents the legs from crossing (Figure 6.9). Under the same assistance settings, muscle-mediated separation feedback changed reversed ankle ordering from 67 % of sampled states to zero.

The model uses a layered SNS for timing together with an engineering force allocator. Its MN target-current compensation cancels incoming PF and reflex currents, so this comparison demonstrates supported model integration and a separation-control effect, rather than full neural pattern control. Appendix D.5.4 gives the matched measures and foot-mount sensitivity.

These distinctions are part of the experimental method: circuit agreement, successful signal exchange, supported stepping, and balance are different capabilities. A result in one cannot be used as an acceptance test for another.

6.6 Discussion: What the Tools Establish

The principal contribution is a set of executable neuromechanical models and experimental tools that make the roles of neural pathways and actuator mechanics accessible to controlled tests. The initial AnimatLab model established the combined walker and CPG organization. The Gait2392 implementations extend controller development to a human-scale model with many joints and actuator channels, while the BPA implementations connect the measured actuator-force model to neural control. SNS-Simscape provides a graphical representation of the pathways, and the curriculum and experiment drivers provide repeatable changes in drive, support, feedback, and perturbation conditions.

The representative results answer different parts of the knowledge gap. Agreement between the NumPy and Simulink rhythm cores establishes that the specified circuit response can be preserved across those implementations under matched inputs. The loaded-arm comparison establishes a corrective effect of feedback in the coupled block

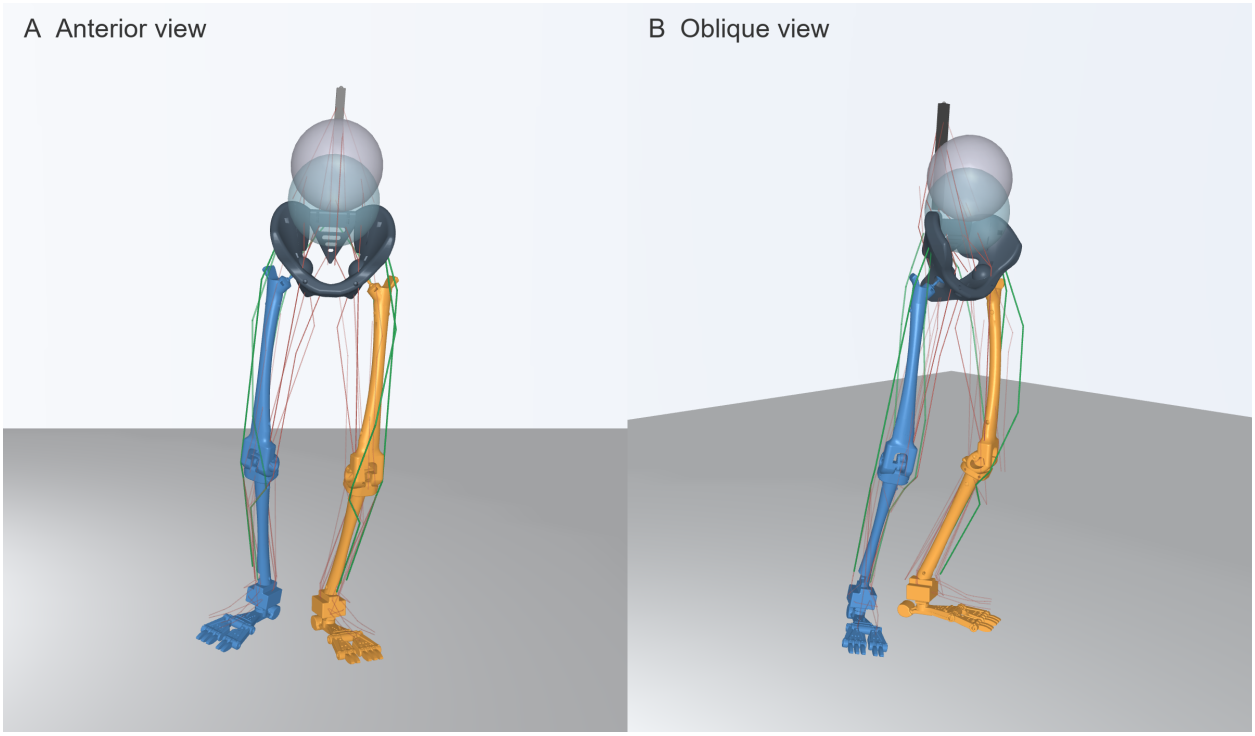


Figure 6.9: Anterior (A) and oblique (B) views of the same recorded CAD-model state. Right and left legs are blue and gold; selected hip-abducting routes are green. The 75 kg model uses 54 reference-tension actuators with hybrid neural and postural control. Vertical assistance supplies 80% of body weight in this interval; lateral and angular assistance remain active. Attachments, the welded foot mount, and passive knee stops are provisional.

model. The fixed-controller W2L comparison identifies a contribution of sensory feedback to bilateral contact coordination. The bridge tests establish the command and measurement exchange needed to carry out such experiments across software environments. Together, these results provide an operational method for separating implementation agreement, sensory contribution, and mechanical response.

Spiking-neuron comparisons preserved a continuous motor-activation interface while changing rhythm timing and, in some runs, limb coordination (Appendix D).

The gap identified at the start of the chapter is addressed through a working experimental workflow: construct a defined neural–mechanical model, check its implementation and interfaces, tune it in stages, and compare controlled changes. The demonstrations establish selected capabilities of that workflow and identify where additional validation is needed. The next chapter brings these tools together with the actuator characterization, torque models, and SADb application, then separates the uses already supported by the dissertation from the studies required for broader robot behavior and hardware deployment.

7 Discussion

This chapter considers how the actuator models, torque corrections, sensory-afferent resource, and simulation tools support biomimetic robot design, and identifies the experiments and integration work still required.

7.1 The Work as a Whole

The outcome that organizes this discussion is that improved first-pass design of a BPA-actuated biomimetic leg is now achievable: a designer can predict actuator force, joint torque, and range of motion accurately enough to commit to an actuator set, a resting length, and a route before the hardware exists (Figure 7.1). At the center of that outcome are the identified design factors that remove the rigid-body assumptions from the torque model (Chapter 4): a constant length offset χ_0, effective stiffnesses χ_1 and χ_2 that lump the series compliance of the attachment brackets, fixtures, and artificial tendons into the bracket frame's principal directions, and a wrapping-loss factor χ_3 that removes usable actuator length as an extensor bends over the joint. The terms were identified on the pinned knee. Applied to the biomimetic knee, they reduce prediction error, although the remaining error depends on the configuration. The biomimetic extensor's fraction of variance unexplained falls from 4.49 to 0.67 with the shared terms (Section 4.3.2). Whether the set is complete is itself now an experimental question: a fourth term χ_4, a stiffness acting along the force path independently of force orientation and presently absorbed by the stiffness pair (χ_1 and χ_2), is arguably the next candidate, and the identification

machinery delivered here is the instrument for settling it.

Improved BPA characterization is the first enabler (Chapter 3). The maximum isometric force of Festo BPAs was shown to depend strongly on resting length, a previously unreported characteristic that corrects the manufacturer's tool at exactly the short lengths a leg design requires, and a nondimensional force surface predicts isometric force from relative strain and relative pressure with only three diameter-dependent coefficients. Because resting length is a primary degree of freedom in actuator selection, this model is what makes that selection rational. The quasi-static surface is complemented at the dynamic end by the state-space, optimization-based model of pressure and force dynamics in pulse-modulated BPAs [**?**], which supplies the actuator behavior that the isometric maps cannot.

Attachment and routing optimization is the second enabler. Morrow, Bolen, and Hunt developed a mesh search for three lumbosacral muscle replacements, balancing torque matching against actuator length, attachment displacement, and diameter [**?**]. This provided a systematic response to the finding that human attachment locations do not transfer directly to a BPA-actuated robot [**??**]. The later knee-route optimization uses the corrected torque model and a continuous search over attachment, routing, and actuator variables, with explicit mechanical constraints (Section 4.3.2). In the route-design predictions, with two 20 mm BPAs per route and routes chosen with the actuators' bending in mind, the redesigned flexor and extensor each meet or exceed the magnitude of their human-scaled torque targets across the full knee range of motion, and the cam-profile extensor converts the direct route's thin margin into genuine design headroom (Section 4.3.2). Stated as a design program, the generalization is one or more BPAs replacing one or more human muscles across one or more joints, serving the two degrees of freedom that matter most

at each joint, with routes grown outward from positions close to the bone and held to the spatial constraints of the candidate actuators and of the BPAs already mounted; the broader route-search diagnostics are collected in Appendix C.8. Where the attachments are fixed and only the route and the tendon length can change, the cam-shape solver applies: the tibial-side route points trace a designed cam profile, trading biomimicry for torque-profile optimization and reduced component forces.

The Sensory Afferent Database (SADb) is the third enabler (Chapter 5). The delivered application makes 1,393 literature and linked records searchable offline. The Papers category contains 943 entries; among those, 833 carry curation notes, 95 carry an afferent classification, and 363 carry robot-or-simulation translation notes. Its table, pivot, and citation-neighborhood views connect a broad search to the source observations and assumptions behind a proposed sensory pathway. Citation symbols encode bibliographic direction, rather than physiological excitation or inhibition. The contribution is both the curated evidence and a practical means of navigating it. Continued classification and source review remain necessary; the AI-assisted triage evaluation did not establish reliable automated categorization [**?**].

The fourth contribution is a functioning neuromechanical experiment workflow (Chapter 6). The AnimatLab and MATLAB integration attempts showed why a moving model, an imported mechanism, or an isolated neural rhythm was insufficient to establish the intended coupled behavior. The resulting Gait2392 and W2L implementations, SNS-Simscape blocks, MuJoCo–Simulink integration, and curriculum make those components and their interactions testable. Matched-input checks establish selected implementation agreement, reflex demonstrations show corrective feedback, and fixed-controller pathway

removal exposes sensory dependence in supported stepping. The curriculum and replay procedures connect these demonstrations into a reusable method for constructing, tuning, and interrogating a model. Their limits, including retained mechanical support and incomplete transfer between models, also define the next experiments.

The coupled-model results must be interpreted under their stated support conditions. Supported stepping and forward progression establish functioning neural–mechanical loops, while removal of tonic drive tests a different dependence from removal of mechanical support. The balance-platform studies provide a separate hardware demonstration of a BPA-actuated limb under pseudo-vestibular inertial feedback [**???**]. Together, the simulation tools and hardware precedents support further study of balance and weight bearing without treating those capabilities as already demonstrated by the stepping tests.

The proposed workflow connects these contributions through explicit checks (Figure 7.1). Force and route predictions must meet the human torque and motion targets before a candidate proceeds to closed-loop simulation and hardware. Failure at either check returns the design to actuator selection or route revision.

7.2 Future Work

The completed contributions support actuator selection, corrected route evaluation, source-based sensory-pathway selection, and controlled simulation experiments. The work below extends those capabilities in distinct directions: dynamic actuator characterization, routing and packaging, continued literature curation, weight-bearing control, and embedded hardware. These are further investigations built on the delivered tools, rather than prerequisites

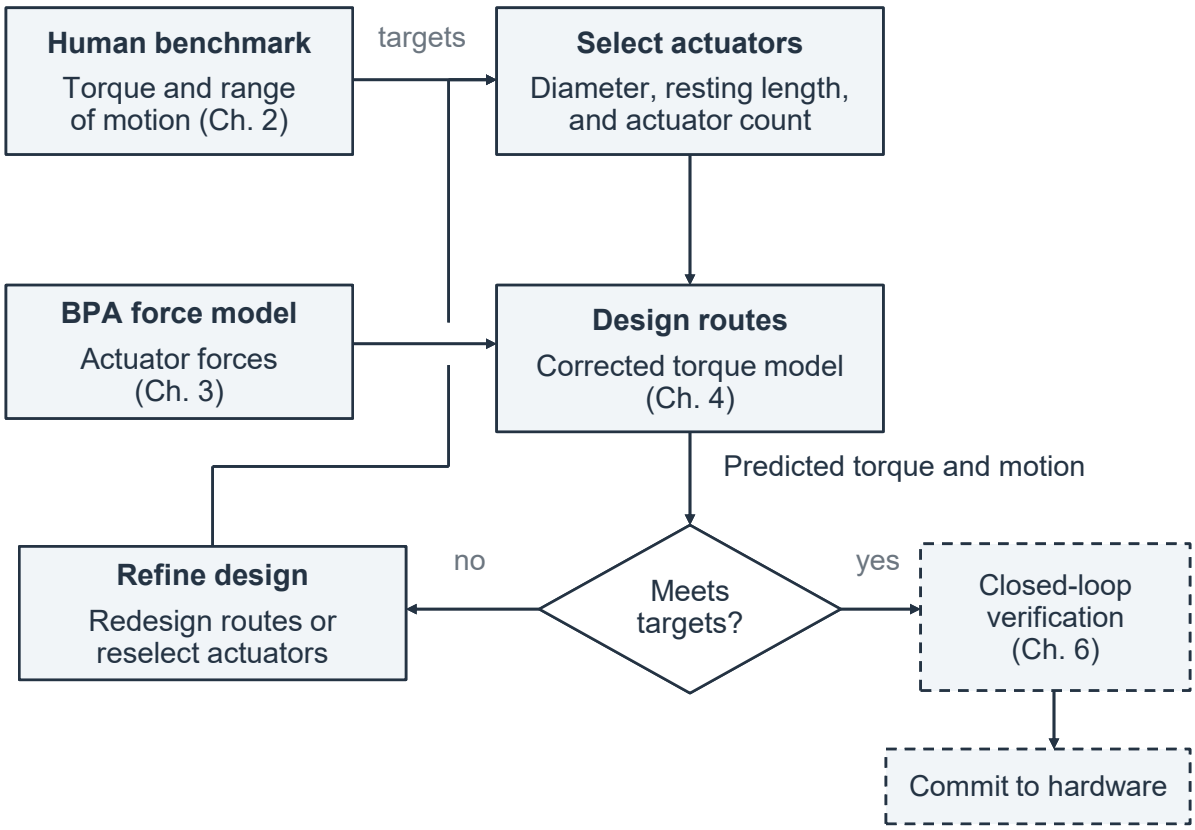


Figure 7.1: Proposed workflow connecting human torque targets, BPA characterization, route design, neuromechanical simulation, and hardware development.

for using the results already established.

7.2.1 Expanding the BPA Force Characterization

The isometric characterization of Chapter 3 leaves one actuator curve unidentified and one regime unmeasured, and both are tractable next experiments. The unidentified curve is the maximum-contraction function: ε_{620} could not be pinned to a predictive relationship with resting length, so each actuator's free contraction is recorded individually before use. Controlled tests designed for that relationship, across multiple product batches of the 10 mm, 20 mm and 40 mm diameters, would either produce general fit surfaces or establish batch variation as the binding uncertainty. The 20 mm actuators characterized here came from a single product batch, so batch-to-batch variation is unquantified and those fits in particular carry that limitation.

The unmeasured regime is everything off the isometric plateau. Biological muscle models separate active and passive force and include activation, contraction, and tendon dynamics [**??**]. Appendix A presents these relations and explains their parameters. A corresponding BPA model would require separate characterization of active pressurized-braid force, passive resistance of the unpressurized actuator and sleeve, and contraction dynamics. Isokinetic and isotonic or even isobaric characterizations could then be assimilated into the same framework rather than fitted separately, with the pressure–force dynamics model of **?** supplying the actuator-side dynamics.

Tensile testing with the Onyx and aluminum end caps, under pressurized and unpressurized conditions, would test the proposed end-effect explanation for resting-length-

dependent force. A complementary finite-element model could represent the rubber layers, crossed aramid braid, end constraints, and contact interfaces described in the Festo patent literature [**?**]. Figure 7.2 shows the proposed geometry and boundary conditions. Comparing predicted deformation with those measurements would test the empirical force surfaces against a mechanical explanation.

In Figure 7.2, longitudinal symmetry permits in-plane expansion. At the half-length boundary, axial displacement can be fixed or axial motion allowed under an optional load. Internal pressure loads the contact interfaces; cap details are schematic.

7.2.2 Expanding Attachment and Routing Design

Each routing tier of Section 7.1 leaves a named next step. The knee's adopted flexor redesign evaluates two staggered paths separately. Physical verification of the two paths' clearances, load sharing, and off-axis loading remains necessary before extending that configuration to a complete leg. The bi-pulley sweep showed that reeved routes do not yet fit the anatomical corridors, where geometry rather than torque was the binding constraint, so clearance-relaxed attachment sites are where that freedom should pay. The whole-body pass needs its route-builder failure fixed and either larger actuators or relaxed corridors on the two routes that satisfy geometry but fall short of torque.

An artificial patella is a further routing option. Its position could be optimized to preserve extensor moment arm and shield the BPA from joint contact, while explicitly accounting for friction and contact transitions. The cam-profile result establishes that a moving routing surface can be optimized; it does not establish the performance of a

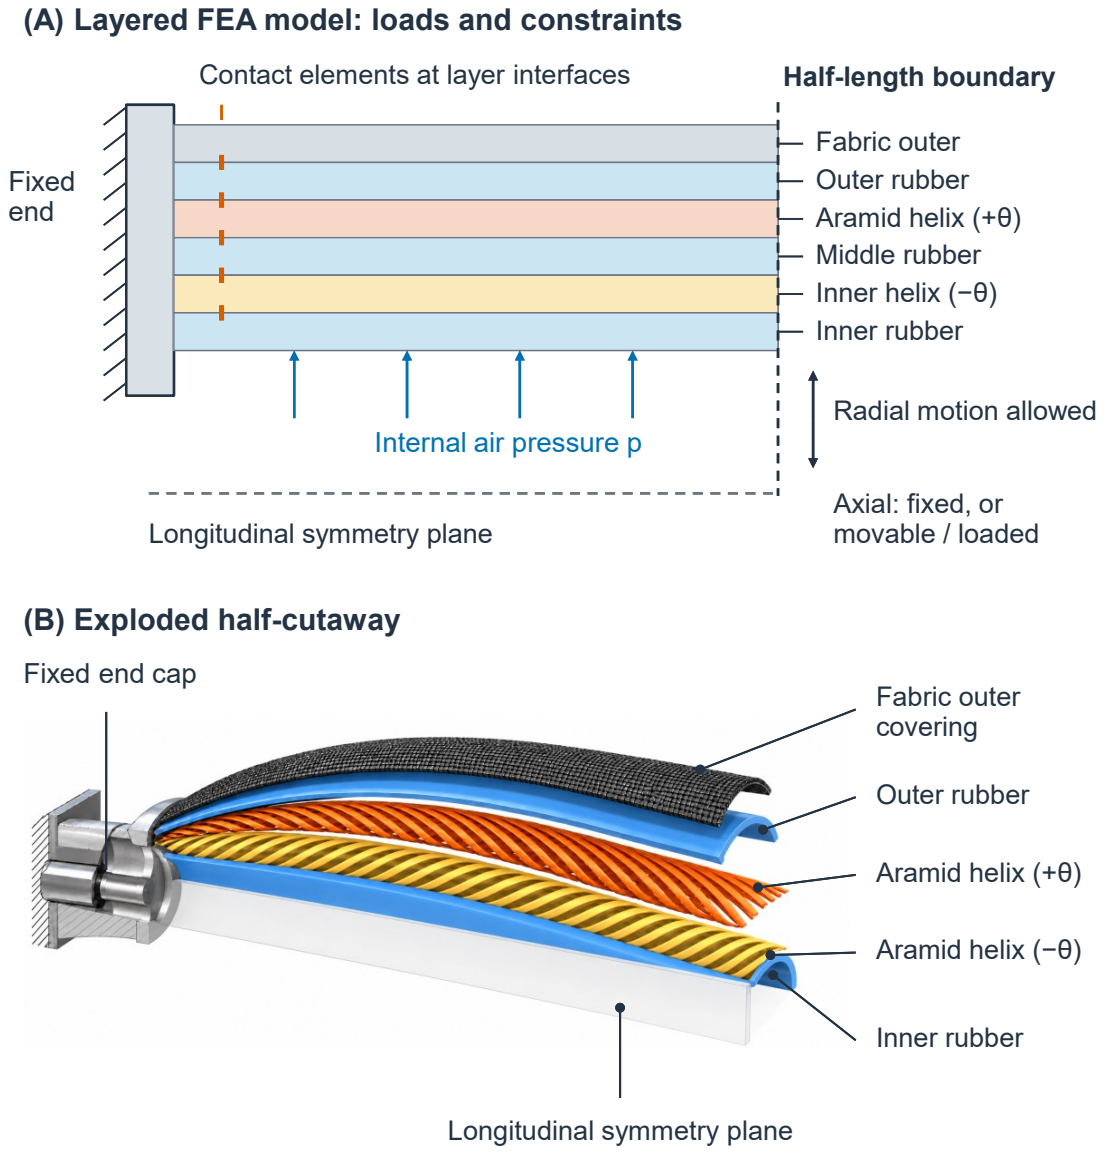


Figure 7.2: Proposed layered BPA finite-element model: contact and boundary conditions (A), and illustrative material cutaway (B). Geometry and layer spacing are schematic; the middle rubber layer is separated only in A.

patella-based design (Chapter 4).

7.2.3 Expanding the Sensory Afferent Database

A monthly literature-scraping extension is being implemented to identify candidate additions to the database. Its retrieval and classification must be evaluated against human decisions before it can support routine updates. Further curation can also target missing afferent classifications, pathway links, and preparation metadata in the existing records. The AI-assisted triage study establishes the relevant division of responsibility: automated tools draft and retrieve, while I retain the inclusion and interpretation decisions [**?**].

7.2.4 Closing the Loop in Simulation

The simulation laboratory of Chapter 6 is assembled, audited, and in active use. The converted models support individual checks of actuator geometry, force direction, and signal exchange, and the BPA force interface evaluates the isometric force model inside the loop; the implemented interfaces and their limits are described in Section 6.4. The 410-neuron SNS-Toolbox instantiation of the spinal network, with per-leg rhythm-generator half-centers, phase-shifted pattern-formation groups, motoneuron pools for every actuator, and Ia, Ib, and II afferent populations, couples to the converted model at a 2 ms neural step, and supported stepping configurations have been tuned against a walking reference and evaluated on the ground (Section D.3.1). A drive-weaning study has since produced a feedback-sustained bilateral configuration whose stepping persists with the externally supplied neural drive removed, still in place under the support rig, and the robot-side

model progresses under its harness with reduced neural drive (Section D.5). These evaluations identify the tasks researchers can now take: the balance and weight-bearing layers that unsupported, bilateral walking requires; transfer from a single tuning reference to the diversity of the locomotion reference library; and the systematic verification series that follows.

Further verification should apply specified stimuli to the rhythm-generator, pattern-formation, and motoneuron layers, while recording both neural timing and mechanical responses. The reduced-circuit tests in Section D.2.3 distinguish RG resetting from suppression of downstream motor output. Extending these tests to the full Gait2392 networks will require matched stimulus protocols and separate assessment of the RG clock and motor pattern. Physiological and modeling studies provide reference phenomena for these comparisons [**????**].

The next experiments should stimulate, remove, or add specified pathways while retaining the other controller and mechanical settings. Existing removal tests establish a starting point (Section D.3.3); adding a candidate sensory or commissural pathway would test whether it improves a defined response. Matched comparisons of spiking and non-spiking controllers should use the same tuning and evaluation conditions. The observed differences in their limb coordination motivate that comparison (Section D.4).

The existing integrations provide a starting point for extending the model set. A full-body OpenSim model with arms and trunk would require its own geometry and actuator checks after conversion. Its controller could then be assembled or inspected with the neural block tools and evaluated with the relevant simulation interface. The matched-circuit tests and pathway-intervention procedures provide a reusable verification method, but each

new model and controller combination must establish its own behavior.

7.2.5 Adding Sensory and Higher-Level Control

Four planned extensions address sensory estimation, coordination, and action selection (Figure 7.3). First, vestibular feedback will be simulated with IMU-like sensing of body orientation and angular velocity, extending the balance-platform work, in which an inverted pendulum with pseudo-vestibular IMU feedback maintained balance on an oscillating platform [**???**] and analytic neural networks reproduced human balance dynamics [**?**]. Second, ocular feedback will be modeled to investigate how visual stabilization cues can help reset IMU drift, a well-known practical problem for inertial sensing on robots, by treating visual measurements of the horizon and surroundings as a low-frequency correction to the vestibular estimate. Third, a cerebellar layer will be explored to adapt commands and coordination between higher-level control and the CPG: a structure that can gate between gaits or movement patterns (e.g., walking, running, and turning) and tune the phase relationships among them, in line with the cerebellum's established role in adaptation and coordination. Multilevel spinal-plus-cerebellar synthetic nervous systems have already been demonstrated for insect-like legs [**?**], and hierarchical cortico-cerebellar-spinal models have been proposed for adaptive step-length control in bipedal locomotion [**?**]. The fourth addition, basal-ganglia circuits realizing the inhibitory and excitatory gating that selects among actions and locomotor modes, would give the controller a substrate for choosing between behaviors rather than only executing one. The database already holds entry points for this program: the human balance and posture cluster and the multisensory

recalibration cluster are curated and connected to the locomotor core through the citation graph (Chapter 5), and the proprioceptive–vestibular interaction literature will guide how the signals are weighted [**?**].

The hierarchy distinguishes implemented spinal and body-model components from proposed supraspinal and sensory additions (Figure 7.3). The balance-platform work provides a vestibular-feedback precedent; integrating that channel with action selection and adaptation remains future work.

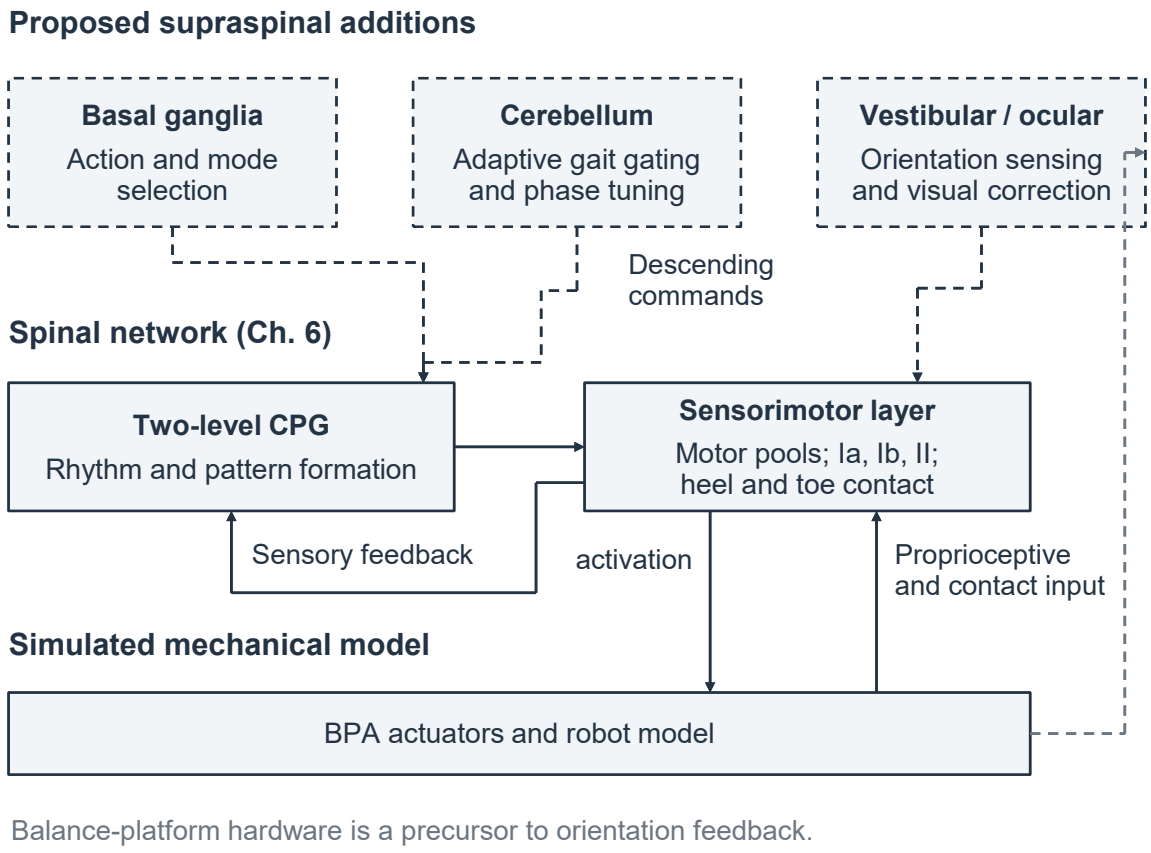


Figure 7.3: Control hierarchy connecting spinal pathways and the BPA-actuated model to proposed supraspinal and sensory extensions. Dashed elements indicate future work.

Figure 7.4 places the implemented simulation interfaces within the proposed hardware workflow. The model-construction and controller-editing routes are detailed separately in Figure 6.2; the embedded controller, valve-control interfaces, and treadmill deployment below remain the next integration steps.

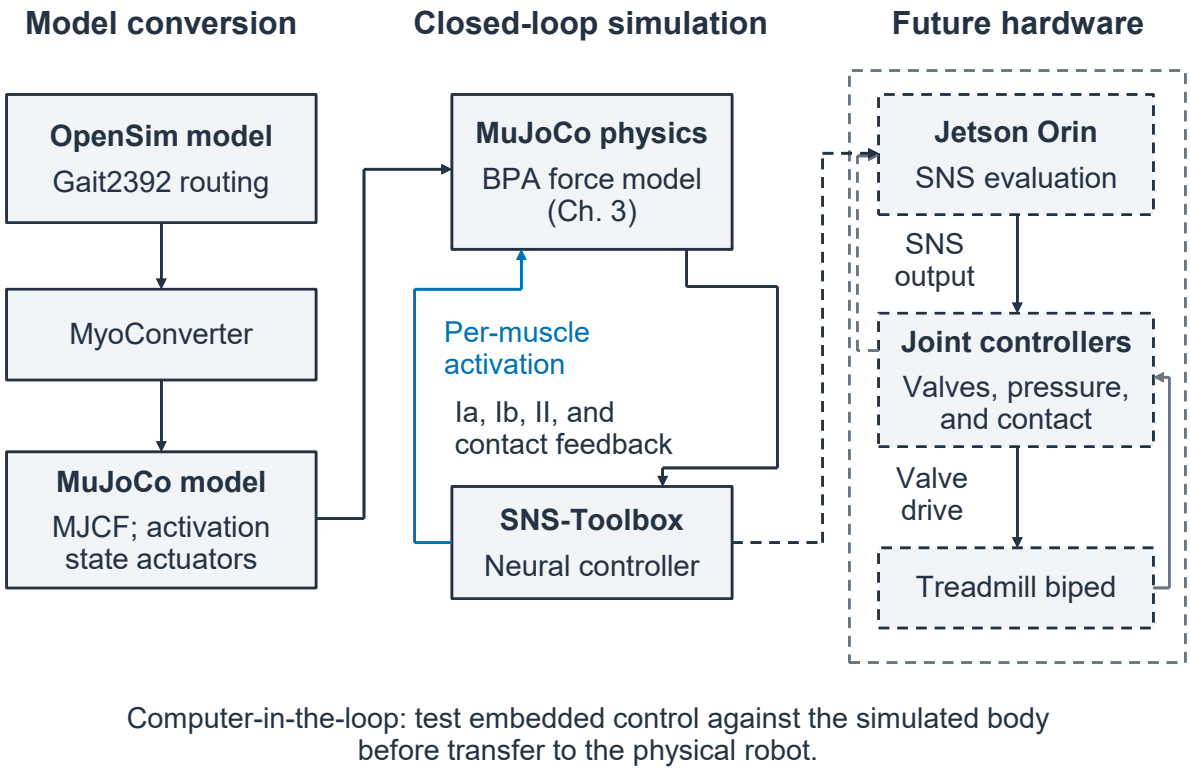


Figure 7.4: Model conversion, closed-loop simulation, and proposed hardware deployment. Simulation interfaces exchange actuator commands and afferent signals. The embedded valve-control and network-evaluation system and treadmill deployment are future work.

7.2.6 Real-Time Hardware Realization

The hardware objective is an embedded SNS controller with low-latency joint loops on Teensy-class microcontrollers and network evaluation on an NVIDIA Jetson Orin. Encoders, strain gauges, pressure sensors, contact sensors, and IMUs would supply the sensory channels. Ia and II signals require length and velocity estimates; Ib force estimates must be checked against actuator length and loading direction. Computer-in-the-loop testing should precede physical deployment, with serial communication reserved for logging and high-level commands. Appendix B supplies the existing neural-pathway schematics; wiring and commissioning remain integration work.

Prior hardware studies support the feasibility of further controller transfer. **?** and colleagues demonstrated SLIP-based control of a physical PAM-actuated biped; **?** coupled the Muscle Mutt MuJoCo model to SNS-Toolbox and implemented rough closed-loop stepping on the physical quadruped using potentiometer and pressure feedback. These are precedents, not validation of the present biped. Reverse-pulley transmissions offer a separate mechanical option: the balance-platform demonstrations traded transmitted force for tendon travel, with increased structural reaction loads [**???**]. Their use at other joints requires route-specific force, excursion, and mounting-load checks.

Because the static force and torque models of Chapters 3 and 4 are computationally inexpensive lookup-table candidates, they will serve as the feedforward model inside the real-time controller, while the dynamic pressure–force behavior under pulse-modulated control will draw on the state-space model of **?**. Embedded execution of the complete CPG-plus-reflex network is a future hardware objective. The simulation framework provides

a place to test the controller before that transfer, including its responses to stimulation, pathway removal, and changed mechanical conditions.

7.2.7 Toward Dynamic Whole-Body Maneuvers

Beyond steady walking, the long-term goal of this line of research is dynamic, whole-body maneuvering: transitions between gaits, reactive balancing, and ultimately ballistic maneuvers that require coordinating the full actuator set through brief, precisely timed activation bursts. My aspiration for this program is a robot that can perform a kickflip, a maneuver that, like its skateboard namesake, demands explosive coordinated extension, precise timing, aerial body orientation, and controlled landing. It is exactly the kind of behavior that exercises every layer of the architecture described above, from contact receptors and spinal pattern generation to vestibular orientation sensing and cerebellar coordination, and it would demonstrate that biomimetic actuation and neuromechanical control scale from steady locomotion to dynamic whole-body movement.

The next experiments can therefore be selected for the question they answer: whether a dynamic actuator model predicts movement, whether a sensory pathway changes coordination, or whether a controller transfers to less-supported mechanics and hardware. Each extends an experimental foundation that is already usable. The common purpose is to make the connection between a design assumption and an observed response traceable.

8 Conclusion

This dissertation moved the biomimetic humanoid robot project from an unvalidated design concept to a quantitatively characterized actuator set and experimentally tested knee-torque models, and it defined the path from those experimental foundations to neuromechanical control.

Chapter 3 produced equations for isometric force in Festo 10 mm and 20 mm-diameter BPAs. The resting-length dependence $F_{620}(l_0)$ corrects the manufacturer tool's opposing trend at 620 kPa, and the nondimensional surface $F^*(\epsilon^*, P^*)$ uses three diameter-dependent coefficients. Together, these fits support actuator selection and route evaluation across resting lengths before prototyping, within the tested isometric conditions.

At the joint level (Chapter 4), the BPA model was combined with transmission corrections for effective length, series compliance, and wrapping. Identification on the pinned knee improved torque prediction, and the shared correction terms also improved predictions on the biomimetic four-bar knee. The corrected model supports route redesign: the optimized two-BPA flexor and cam-extensor configurations are predicted to meet their human-scaled torque targets across the design range. The isometric tests demonstrate that the redesigned configurations meet or exceed their selected human-muscle reference torques at all fifteen flexor and nine extensor measurement points. This establishes capacity at the sampled angles and pressures, while agreement with pressure-specific predictions remains uneven.

The SADb application in Chapter 5 makes physiological evidence searchable by pathway, preparation, and modeling relevance. The tools in Chapter 6 turn selected pathways into testable neural–mechanical models: the Li–Deng/W2L and Gait2392 implementations, SNS-Simscape blocks, MuJoCo–Simulink integration, and staged tuning and intervention procedures. Demonstrations establish matched-circuit agreement, corrective feedback, command and measurement exchange, and pathway-dependent stepping under specified support. These tests provide a repeatable framework for comparing neural and mechanical hypotheses. General controller transfer and independent balance remain open.

These actuator and joint models, the sensory-afferent application, and the neuromechanical tools make design assumptions testable: force against measurement, routing against torque requirements, and sensory pathways against controlled changes in movement. The future studies in Chapter 7 extend these tests toward dynamic actuation, independent balance, and hardware. This work provides an experimental foundation for biomimetic robot design and for studying how neural control interacts with mechanical structure.

APPENDICES

Appendix A: Muscle Force and Joint Torque Calculations

This appendix documents the reference implementation of the force, torque, and compliance equations used in Chapters 3 and 4. When present in MATLAB code, file names are given so that any number in this dissertation can be traced to the code that produced it.

A.1 Human muscle force calculations

A.1.1 OpenSim Thelen2003 and Millard2012 Formulations

The Thelen2003 and Millard2012 equilibrium muscle models describe active and passive fiber force, pennation, and a series tendon [**??**]. Their common force balance is

$$F_T(l_T) = F_M(l_M, \dot{l}_M, a)\cos\alpha, \qquad l_{MT} = l_T + l_M\cos\alpha, \tag{A1}$$

where a is activation, α is pennation, and l_M, l_T, and l_{MT} are fiber, tendon, and musculotendon lengths. The cosine projects fiber force onto the tendon direction. When tendon compliance is enabled, fiber length is a state determined by the equilibrium calculation rather than the full musculotendon path length.

In the damped Millard formulation, fiber force has the form

$$F_M = F_0\left[a f_L(\tilde{l}_M) f_V(\tilde{v}_M) + f_{PE}(\tilde{l}_M) + \beta\tilde{v}_M\right], \tag{A2}$$

with normalized fiber length $\tilde{l}_M$, normalized velocity $\tilde{v}_M$, maximum isometric force F_0, and damping coefficient β. The active, passive, and tendon curves and the treatment of singular configurations differ between the Thelen and Millard models. The OpenSim implementation can also disable tendon compliance or activation dynamics, so the model class alone does not establish which dynamics were exercised in an analysis [**?**].

Thelen activation and contraction dynamics

The following seven relations correspond, in order, to Equations (1)–(7) of **?**. They specify the activation state, passive and active force–length curves, tendon force, and fiber velocity. Bars denote normalization by maximum isometric fiber force F_0^M or optimal fiber length L_0^M, as appropriate.

Activation follows

$$\frac{da}{dt} = \frac{u - a}{\tau_a(a, u)}, \tag{A3}$$

with an activation-dependent time constant

$$\tau_a(a, u) = \begin{cases} \tau_{\text{act}}(0.5 + 1.5a), & u > a, \\ \tau_{\text{deact}}/(0.5 + 1.5a), & u \leqslant a. \end{cases} \tag{A4}$$

Here $u \in [0, 1]$ is neural excitation, a is muscle activation, and τ_{act} and τ_{deact} set the activation and deactivation time scales.

The normalized passive fiber force and active force–length multiplier are

$$\bar{F}^{PE} = \frac{\exp\left[k^{PE}(\bar{L}^M - 1)/\varepsilon_0^M\right] - 1}{\exp(k^{PE}) - 1}, \tag{A5}$$

$$f_l = \exp\left[-\frac{(\bar{L}^M - 1)^2}{\gamma}\right]. \tag{A6}$$

The tendon relation is

$$\bar{F}^T = \begin{cases} \dfrac{\bar{F}^T_{\text{toe}}}{\exp(k_{\text{toe}}) - 1}\left[\exp\left(\dfrac{k_{\text{toe}}\varepsilon^T}{\varepsilon^T_{\text{toe}}}\right) - 1\right], & \varepsilon^T \leqslant \varepsilon^T_{\text{toe}}, \\ k_{\text{lin}}(\varepsilon^T - \varepsilon^T_{\text{toe}}) + \bar{F}^T_{\text{toe}}, & \varepsilon^T > \varepsilon^T_{\text{toe}}. \end{cases} \tag{A7}$$

The fiber-velocity relation and its shortening/lengthening denominator are

$$V^M = (0.25 + 0.75a)V^M_{\max}\frac{\bar{F}^M - af_l}{b}, \tag{A8}$$

$$b = \begin{cases} af_l + \bar{F}^M/A_f, & \bar{F}^M \leqslant af_l, \\ \dfrac{(2 + 2/A_f)(af_l\bar{F}^M_{\text{len}} - \bar{F}^M)}{\bar{F}^M_{\text{len}} - 1}, & \bar{F}^M > af_l. \end{cases} \tag{A9}$$

Here $\bar{F}^M$ is the normalized active fiber force entering the force–velocity relation, $V^M = \dot{L}^M$ is negative during shortening, and $V^M_{\max}$ is the maximum shortening-speed scale. Table A1 defines the curve parameters. A shape factor is a dimensionless parameter controlling a curve's curvature or width; it is not a stiffness in N/m. For example, increasing γ broadens the active force–length curve while preserving its peak at $\bar{L}^M = 1$. Changing k^{PE} changes the passive curve's shape while retaining unit force at strain ε^M_0.

Table A1: Length, strain, and shape parameters in the Thelen formulation.

Symbol	Definition
$\bar{L}^M = L^M / L_0^M$	Fiber length relative to optimal fiber length.
ε_0^M	Passive fiber strain at force F_0^M.
k^{PE}	Passive force–length exponential shape factor (5 in the paper).
γ	Active force–length width parameter (0.45 in the paper).
ε^T	Tendon strain relative to tendon slack length.
ε_0^T	Tendon strain at force F_0^M.
$\varepsilon_{\text{toe}}^T$	Strain at the end of the nonlinear tendon region; $0.609\varepsilon_0^T$.
$\bar{F}_{\text{toe}}^T$	Normalized force at that transition; 0.33.
k_{toe}, k_{lin}	Tendon exponential shape factor (3) and linear slope ($1.712/\varepsilon_0^T$).
A_f	Force–velocity shape factor (0.25 in the paper).
$\bar{F}_{\text{len}}^M$	Maximum normalized lengthening-force factor.

The passive expressions describe tensile loading; they do not imply that a slack muscle or tendon transmits compression. These are the original paper's constitutive relations. Numerical implementations must also handle slack states, minimum activation, pennation limits, and the inversion of the force–velocity relation. Consequently, the original equations alone do not specify every numerical safeguard in an OpenSim implementation.

Millard equilibrium and damping

The equilibrium relation of **?** can be written in normalized form as

$$\left[a f_L(\tilde{l}_M) f_V(\tilde{v}_M) + f_{PE}(\tilde{l}_M)\right] \cos\alpha - f_T(\tilde{l}_T) = 0. \tag{A10}$$

The damped formulation adds $\beta\tilde{v}_M$ inside the brackets, as in Equation A2. The functions f_L, f_V, f_{PE}, and f_T describe active force–length, force–velocity, passive fiber force, and

tendon force. Millard's smooth characteristic curves and damping address numerical behavior while retaining the series force balance. Figure 3 of that paper illustrates the equilibrium, damped-equilibrium, and rigid-tendon formulations. A rigid-tendon model removes tendon stretch as a state; it does not simply set the passive fiber force to zero.

A.1.2 OpenSim Moment Arms and Human Torque Targets

The human torque targets of Section 4.1.5 are OpenSim tendon-excursion moment arms: for each muscle, the moment arm about the knee is the partial derivative of muscle-tendon length with respect to the knee coordinate, with positive tensile force and torque positive in the coordinate direction, $r(\theta_k) = -\partial l_{\mathrm{MT}} / \partial \theta_k$, evaluated across the range of motion. The exported analysis tables archived with the test data (`OpenSim_Vasti_MomentArm.txt`, `OpenSim_Vasti_Results.txt`, `OpenSim_Bifem_MomentArm.txt`, `OpenSim_Bifem_Results.txt`) carry the per-muscle moment arms and the corresponding maximum-isometric torque curves (moment arm multiplied by each muscle's maximum isometric force) for vastus intermedius, lateralis, and medialis, and for the biceps femoris short head; the optimization loads the vastus medialis and biceps femoris short-head columns as its targets (`buildKneeExtContext20mm.m` lines 119–140; `Knee_Flexor_Data_20mm.m` line 546). For an unconstrained planar rotation, this derivative definition and the geometric perpendicular-foot construction of Equation (A14) below are equivalent for a straight muscle path crossing a planar revolute joint; the constrained and moving-path-point geometry is evaluated by OpenSim rather than by substituting the robot-side straight-path construction.

A.2 Converted MuJoCo Muscle Formulation

A.2.1 MuJoCo Muscle Actuator

The converted model uses MuJoCo muscle actuators, which combine activation-dependent active force with passive tensile force. Let L denote the routed musculotendon length, L_T the constant biological-tendon length, and L_0 the optimal fiber length inferred from the actual and normalized operating ranges. The normalized muscle length is $\bar{L} = (L - L_T)/L_0$. The signed actuator force is

$$F_{\text{act}} = -F_0 \left[a f_L(\bar{L}) f_V(\dot{L}/(L_0 v_{\max})) + f_P(\bar{L}) \right]. \tag{A11}$$

Here $v_{\max}$ is the maximum shortening speed in optimal fiber lengths per second. The negative sign denotes a pulling actuator force; both active and passive terms contribute to tension. The rigid tendon has a constant length, but no elastic state. MuJoCo also omits pennation, so its differences from the OpenSim equilibrium formulations are not limited to tendon compliance [**?**].

A.3 BPA Force Model Implementation

A.3.1 Strain Measures

The musculotendon path length l_{LMT} is the sum of straight segments between route via points, transformed into a common frame at each frame crossing. The contractile strain used by the force model subtracts the artificial tendon and both end-cap fittings from the

path length before normalization by the resting length,

$$\varepsilon \;=\; \frac{l_0 - \left(l_{LMT} - l_{ten} - 2\,l_{fitting} - \chi_0\right)}{l_0}, \tag{A12}$$

where χ_0 is the constant length offset of Section 4.1.4 (zero in the uncorrected model). [Provenance: `MonoPamDataExplicit.m` get.Contraction, lines 210–217; `minimizeFlxPin.m` lines 415–426. End-cap fitting length $l_{fitting} = 25.4\,\text{mm}$ is a constructor default, `MonoPamDataExplicit.m` line 69.]

The maximum free-load contraction at 620 kPa is $\varepsilon_{620} = (l_0 - l_{620})/l_0$, where l_{620} is the BPA length at 620 kPa with no external load (code variable `KMAX`; `MonoPamDataExplicit.m` line 278). The relative strain is $\varepsilon^* = \varepsilon/\varepsilon_{620}$ (line 281).

A.3.2 Normalized Force Surface

The normalized isometric force surface of Equation (3.4) is implemented in closed form in `Functions/f_festo.m` (lines 1–26),

$$F^* = c_0\left(e^{-c_1\varepsilon^*} - 1\right) + P^*\,e^{-c_2(\varepsilon^*)^2}, \qquad F^* \leqslant 0 \;\Rightarrow\; 0, \tag{A13}$$

with $P^* = P/620\,\text{kPa}$. The zero-clamp $F^* \leqslant 0 \Rightarrow 0$ is applied by the callers (e.g. `Functions/festo4.m` line 35), not inside `f_festo.m` itself. The coefficients were fitted with `fittype` nonlinear least squares and least-absolute-residual robustness (`Testing_Data/bpaFits.m`, lines 45–60 and 115–140, lower bounds $[0, 0, 0]$) and are reported in Table 3.2. The code additionally holds a 40 mm fit not used in this dissertation ($c_0 = 0.1224$, $c_1 = 10.47$,

$c_2 = 2.023$; `f_festo.m` lines 17–20).

At runtime the torque model may alternatively evaluate a spline interpolation of the manufacturer's datasheet curves (`Functions/festo4.m`, backed by `FestoLookup.mat`), with force clamped to zero for $\varepsilon^* > 1$ or $F^* < 0$. Equation (A13) replaced this lookup in the corrected model.

A.3.3 Maximum Force

The resting-length-dependent maximum force of Equation (3.5) is implemented in `Functions/maxBPAforce.m` (10 mm: line 46; 20 mm: line 57; duplicate for 20 mm in `Functions/Fmax20.m`) with the fixed 7.5 mm end-contact offset. The fitted 20 mm coefficients carry confidence intervals $a_1 \in [1.4745, 1.5011]\,\mathrm{N\,kPa^{-1}}$ and $a_2 \in [0.0238, 0.0258]\,\mathrm{kPa^{-1}\,m^{-1}}$ (`Fmax20.m` lines 26–27). Total force is the product $F = F^*(\varepsilon^*, P^*)\, F_{620}(l_0)$, resolved into a vector along the force direction (`MonoPamDataExplicit.m` lines 283–287).

A.4 Muscle Path Length and Moment Arm

Path length accumulates the Euclidean norm of each segment between via points, applying the body-frame transformation whenever the path crosses between the femur and tibia frames (`MonoPamDataExplicit.m` lines 78–95, total at lines 122–133).

The vector moment arm is implemented as the perpendicular foot of the joint origin onto the line of action,

$$\mathbf{r} \;=\; \mathbf{p} \;-\; \hat{\mathbf{u}}\,(\hat{\mathbf{u}} \cdot \mathbf{p})\,, \tag{A14}$$

where $\mathbf{p}$ is the vector from the joint origin to the muscle-path crossing point and $\hat{\mathbf{u}}$ is the unit force direction (`MonoPamDataExplicit.m` lines 156–168; identical implementations in `MonoMuscleData.m` lines 111–123, the `Mom` function of `minimizeFlxPin.m` at line 539, and `MonoPam_mult.m` lines 186–201). This is the same quantity as the cross-product projection of Equation (4.2); the planar scalar moment arm about the joint axis is the magnitude of the in-plane components, computed as `hypot(mA(:,1), mA(:,2))` (`predictKneeFlexor20mm.m` line 217). Torque about the joint z-axis is $\mathbf{m} = \mathbf{r} \times \mathbf{f}$, evaluated only where the BPA is in tension (values are discarded for $\varepsilon < -0.02$, the excessive stretching region; `MonoPamDataExplicit.m` line 308).

A.5 Compliance Model Implementation

The compliance-corrected muscle length entering Equation (4.5) is computed as $l_m = l_{LMT} - \chi_0 - l_{ten} - 2\,l_{fitting}$ (`minimizeFlxPin.m` lines 415–426; identically `minimizeExtX3.m` line 345). In the two-bracket fit of record, the insertion-bracket stiffness matrix and the projected scalar stiffness along the force direction are

$$\mathbf{K}_{br} = \mathrm{diag}(\chi_1,\, \chi_2,\, \chi_1), \qquad k_{\hat{u},br} = \left(\hat{u}^{\,T}\mathbf{K}_{br}^{-1}\hat{u}\right)^{-1}, \tag{A15}$$

matching Equation (4.9), and the origin-side second bracket contributes an additional scalar compliance c_{b2} in series (Appendix C, Table C9). As executed by the evaluator of record, both bracket frames are assembled pitch-only, at `minimizeFlxPin.m` line 289 (insertion) and line 306 (origin), with the two-rotation assemblies carried as commented alternatives; the insertion array is $[X_1, X_2, X_1]$ (line 387) and the origin array $[X_1, X_1, X_2]$

(line 400). The two-rotation-capable twin of the same block is `minimizeFlxPin2brk.m` lines 387–418. An earlier arm-1 array $[X_1, X_2, X_2]$ is superseded and survives only in the provenance comment at line 387, and the legacy single-bracket evaluator uses $[X_1, X_2, X_1]$ (`minimizeFlxPinX3.m` line 267).

The artificial tendon is modeled as an axial spring with stiffness

$$k_{ten} \;=\; \frac{n\,A_c\,E}{L_{ten}}, \tag{A16}$$

with $n = 2$ parallel cable strands ($\times 6$ for the 20 mm assembly), $A_c = 1.51 \times 10^{-6}\,\mathrm{m}^2$ (19-strand cable), and $E = 193 \times 10^9$ Pa (`MonoPam_mult.m`, function `Spr()`, lines 627–637; the evaluators obtain the spring rate through `Spr(bpa_i)`, `minimizeFlxPin.m` line 156). Series combination with the projected bracket compliance gives $k_{eq} \;=\; \left(k_{\hat{u},br}^{-1} + k_{ten}^{-1} + c_{b2}\right)^{-1}$ (`minimizeFlxPin.m` lines 394–409).

The scalar compliance displacement δ of Equation (4.23) is found numerically with `fzero` on the bracket $[0,\, l_m - l_{620}]$,

$$F_{620}(l_0)\,F^*(\varepsilon^*(\delta), P^*) - k_{eq}\,\delta = 0, \tag{A17}$$

with $\delta \equiv 0$ when the uncorrected relative strain already exceeds unity (`minimizeFlxPin.m` lines 434–441). The vector bracket deflection is then $\boldsymbol{\epsilon}_{br} = \mathbf{K}_{br}^{-1}\,(\|\mathbf{f}\|\,\hat{\mathbf{u}})$ (line 461), the cable stretch is $\|\mathbf{f}\|/k_{ten}$ (line 478), and the attachment point, path, moment arm, and torque are recomputed. For the two-BPA flexor grouping, the coupled balance shares one common force between both muscles and is solved analogously on $[0, F_{620}]$ (`MonoPam_mult.m` lines 556–612).

A.6 Wrapping-Loss Implementation

For the pinned-knee extensor, Equation (4.25) is implemented segment-wise over the wrap arc with radii $R_1 = 12\,\text{mm}$ and $R_2 = 40\,\text{mm}$ and segment boundaries $\theta_1 = 27^\circ$ and $\theta_2 = -23^\circ$ (`minimizeExtX3.m`, nested `Contraction`, lines 327–370). For the biomimetic-knee 20 mm cell, the arc-length factor is the routed bend measure $b(\theta_k)$ (arc length swept by the wrap, built with the `wRap` bend radius in `buildKneeFlexorRoute20mm.m` lines 265–268), giving $\Delta l = \chi_3\, b(\theta_k)\,(1-\varepsilon^*)^2$ (`MonoPam_mult.m` lines 398–441), with a hardcoded-geometry fallback for degenerate routes (lines 428–434). The loss enters only the force-producing strain; the kinematic (prediction) strain excludes it (`minimizeExtX3.m` lines 236–240).

Appendix B: Neuron Model Equations

This appendix defines the neuron and synapse equations and records the parameters used in the paired graded/spiking comparisons. The archived AnimatLab and 410-neuron Gait2392 comparisons are distinct from the independently replayed W2L configuration in Chapter 6. The final section documents that W2L configuration directly from its compiled network.

B.1 Spiking Neuron Model

Each spiking neuron is a conductance-based leaky integrator with threshold and adaptation,

$$C_m \dot{V} \;=\; -\,g_{\text{leak}}\,(V - E_{\text{rest}}) \;+\; I_{\text{syn}} \;+\; I_{\text{app}}, \tag{B1}$$

where V is the membrane potential, C_m the membrane capacitance ($g_{\text{leak}} = C_m/\tau_m$ with membrane time constant τ_m), E_{rest} the resting potential, I_{syn} the summed synaptic current, and I_{app} any applied (tonic stimulation) current. A spike is emitted when V exceeds the firing threshold, after which V resets and an after-hyperpolarization (AHP) of amplitude A_{AHP} with time constant τ_{AHP} is applied; the threshold itself accommodates with the firing history (relative accommodation r_{acc}, time constant τ_{acc}). Optional calcium channels are defined but disabled in the archived AnimatLab model ($G_{\text{max,Ca}} = 0$).

The parameter values shared by the neurons in the walker model are given in Table B1.

Table B1: Spiking neuron parameters in the AnimatLab bipedal walker model (`Biped_2xCPG_wSubs`). All neurons share these values except the motor neuron threshold noted below.

Parameter	Symbol	Value
Resting potential	E_{rest}	$-60\,\text{mV}$
Membrane time constant	τ_m	$5\,\text{ms}$
Initial firing threshold	V_{th}	$-55\,\text{mV}$ (motor neurons: $-59\,\text{mV}$)
Relative accommodation	r_{acc}	0.3
Accommodation time const	τ_{acc}	$10\,\text{ms}$
AHP amplitude	A_{AHP}	$1\,\text{mV}$
AHP time constant	τ_{AHP}	$3\,\text{ms}$

B.2 Non-Spiking (Graded) Chemical Synapses

Pattern-formation and motoneuron-drive connections in the walker use graded chemical synapses, whose conductance is a clipped linear ramp of the presynaptic potential,

$$g_{\text{syn}}(V_{\text{pre}}) \;=\; g_{\max}\;\text{clamp}\!\left(\frac{V_{\text{pre}}-V_{\text{th}}}{V_{\text{sat}}-V_{\text{th}}},\,0,\,1\right), \qquad I_{\text{syn}} = g_{\text{syn}}\left(E_{\text{syn}}-V_{\text{post}}\right), \tag{B2}$$

Here inward current is positive, consistently with Equation (B1); an excitatory reversal potential above the postsynaptic voltage therefore produces depolarization. The quantities V_{th} and V_{sat} are the synaptic threshold and saturation potentials and E_{syn} the reversal potential. The walker model instantiates three non-spiking synapse types:

- **Nicotinic ACh (excitatory drive):** $E_{\text{syn}} = -10\,\text{mV}$, $g_{\max} = 0.5\,\mu\text{S}$, $V_{\text{th}} \to V_{\text{sat}} = -65\,\text{mV} \to -20\,\text{mV}$.
- **Hyperpolarizing IPSP (inhibition):** $E_{\text{syn}} = -70\,\text{mV}$, $V_{\text{th}} \to V_{\text{sat}} = -55\,\text{mV} \to -40\,\text{mV}$.

- **Depolarizing IPSP:** $E_{\text{syn}} = -55\,\text{mV}$, $V_{\text{th}} \rightarrow V_{\text{sat}} = -60\,\text{mV} \rightarrow -40\,\text{mV}$.

Spiking chemical synapses and electrical synapses (rectifying coupling weight 0.1, non-rectifying 0.2) are also defined in the model and are available for the sensory afferent pathways of the MuJoCo/SNS implementation.

B.3 Relation to the SNS-Toolbox Implementation

These are the standard AnimatLab firing-rate/conductance models. SNS-Toolbox [**?**] implements corresponding classes of non-spiking and spiking neurons with conductance-based synapses, and the AnimatLab parameters above served as the physically motivated initialization for the first SNS-Toolbox instantiation of the spinal network: a 410-neuron network with per-leg rhythm-generator half-centers, phase-shifted pattern-formation groups, motoneuron pools for all 92 muscle-tendon actuators of the converted model, and Ia, Ib, and II afferent populations. In that implementation the non-spiking membrane and graded-synapse equations take the same forms as Equations (B1) and (B2) with the spiking threshold and accommodation terms removed, and the afferent resting-tone offsets are zeroed so that Ia and Ib signals are pure activity measures; the rhythm-generator layer has been verified to oscillate at a human-like period with the two legs in antiphase under tonic drive (Section 7.2.4). The matched circuit checks in Chapter 6 verify selected ports, but do not establish full equivalence of these larger networks; the tonic stimulation tests of Section 7.2.4 correspond to constant I_{app} injections into Equation (B1) at the motoneuron, pattern-formation, and rhythm-generator layers respectively.

B.4 Spiking-Mirror Implementation, Calibration, and Verification

This appendix documents the numeric details behind the spiking-neuron mirrors of Section D.4, in the same file-traceable form as Appendix A: the cell and synapse equations as implemented, the calibration factors, the full verification-gate table of the MuJoCo mirror, and the variant-by-variant AnimatLab conversion table. The artifacts live under `Code/MuJoCo_SNS/spinal/reports_spiking_20261002/` (reports, gate scripts, logs, and metric files) and `Code/MuJoCo_SNS/spinal/` (the mirror builder `build_network_spiking.py` and its calibration script `calibrate_spiking.py`); the Simulink blocks are built by `Code/Matlab/SNS_Simscape/sns_build_spiking.m` and tested by `sns_units_test_spiking.m` in the same folder.

B.4.1 Adapting Integrate-and-Fire Cells and Hybrid Synapses

Between spikes, every spiking cell of the MuJoCo mirror obeys the membrane equation

$$C_m \dot{V} \;=\; -\,g_m\,(V - V_{\text{rest}}) \;+\; I_{\text{syn}} \;+\; I_{\text{bias}}, \tag{B3}$$

with a spike emitted when V reaches the firing threshold θ, after which V resets to V_{reset} and θ increases by the increment $\Delta\theta$; between spikes θ decays back toward its initial value with time constant τ_θ, bounded by the adaptation floor. The tonic bias I_{bias} parks the equilibrium a few millivolts below threshold, reproducing the analog cells' near-threshold operating point. Table B2 lists the per-class values as coded (SNS-Toolbox's `SpikingNeuron` class; the numeric membrane update used by this archived implementation is consistent

with its specified current units at $g_m = 1\,\mu S$, which the mirror therefore uses throughout). The rhythm-generator adaptation time constant reuses the analog network's persistent-sodium inactivation value, so that tuned adaptation values transfer between builds; at the rhythm gate's operating point it is 0.35 s. The mechanism behind that value, recorded here rather than in the chapter body, is as follows. In the analog rhythm cells of the contact-driven cores, burst termination is carried by a persistent sodium current, a slowly inactivating inward current; in the contact-driven implementations described here its inactivation time constant is held fixed rather than voltage-dependent, because the tested voltage-dependent form quenched burst termination and leaves the network latched in a static state, the failure documented with the archived-model results of Chapter 6.

Table B2: Spiking cell parameters of the MuJoCo mirror (real millivolts; `build_network_spiking.py`, `SPIKE` table, with the rhythm-generator increment superseded by the calibration of record). All cells rest at -70 mV with membrane conductance 1 µS and the tonic bias of Table B3.

Cell class	θ (mV)	V_{reset} (mV)	$\Delta\theta$ (mV/spike)	τ_θ (s)
Rhythm generator E/F	-50	-60	4.0 (calibrated)	0.35 (analog value)
Pattern-formation	-50	-60	0.3	0.15
Interneurons, commissurals	-50	-60	—	—
Afferent encoders	-52	-58	—	—
Renshaw	-52.5	-58	—	—

Spiking chemical synapses, in both the toolbox class used in the MuJoCo mirror and the Simulink `SpikingSynapse` block, carry a conductance that decays between events and

jumps by a fixed increment per presynaptic spike, saturating at its maximum,

$$\dot{g} \;=\; -\,g/\tau_{\text{syn}} \;\text{ between spikes,} \qquad g \leftarrow \min(g_{\max},\, g + g_{\text{inc}}) \;\text{ per spike,} \tag{B4}$$

with $I_{\text{syn}} = g\,(E_{\text{syn}} - V_{\text{post}})$ as for the graded synapse of Equation (B2). The Simulink `HybridSpikingSynapse` implements the same dynamics with the spike detected from the presynaptic voltage crossing a detection threshold that must sit strictly between the presynaptic resting and firing potentials (a detection threshold equal to the firing potential fires nothing, because the presynaptic voltage resets at threshold). Reversal potentials are 0 mV (excitation, above rest) and -70 mV (inhibition, shunting at rest) in the spiking frame, and $+8/-5$ mV in the motor neurons' 0–5 mV frame. The graded-to-spiking command pathways (descending drive, posture, and bias surrogates) keep the graded synapse of Equation (B2).

B.4.2 Calibration of Record

Table B3 lists every calibration factor produced by `calibrate_spiking.py` (`spiking_calibration.json`), with the measured residual of each. The spike-to-motor-neuron factors are bisected so that a steady 40 Hz presynaptic train reproduces the non-spiking steady-state motor-neuron voltage; the spike-to-spiking factors are rate-calibrated (a 40 Hz source drives a relay cell at 40 Hz, and a calibrated inhibitory factor silences a 40 Hz-driven cell); the readout-tap factor maps a 50 Hz train to the full-scale analog level. The rhythm-generator threshold increment was set by an isolated-pair sweep followed by a full-network sweep over alternation quality and period target.

Table B3: Calibration factors for the spiking mirror (`spiking_calibration.json`). Residuals are the measured deviation at each factor's calibration target.

Factor	Value	Calibrated against	Residual
Graded saturation reference	0.9567	analog network measurement	—
Spike-to-MN, excitatory ($k_{sn,exc}$)	5.4087	MN steady voltage	1.0×10^{-4} mV
Spike-to-MN, inhibitory ($k_{sn,inh}$)	1.2284	MN steady voltage	2.6×10^{-4} mV
Graded-to-spiking drive (k_{ns})	0.35	relay firing rate	-0.67 Hz
Spike-to-spiking, excitation	1.8754	40 Hz rate transfer	0.0 Hz
Spike-to-spiking, inhibition	0.9996	silences 40 Hz cell	0.0 Hz
Readout tap (g_{inc})	1.7093	50 Hz → full scale	5.00 mV check
RG threshold increment	4.0 mV/spike	pair + full-net sweeps	period target
Tonic bias current	16 nA	equilibrium -54 mV	—
Synaptic time constants	5 ms (exc), 20 ms (inh)	analog values	—

B.4.3 Verification Gates of the MuJoCo Mirror

Table B4 records the complete gate set of Section D.4 with verdicts and headline numbers. The two failures are reported as failures: the constant-drive period misses its tolerance with the tuning knobs identified, and the air-regime cadence clause misses while every structural clause passes. Gate scripts, logs, and machine-readable outputs are archived in the reports folder (`tools/`, `logs/`, and the `json` metric files respectively).

B.4.4 AnimatLab Conversion Variants

Table B5 catalogues every spiking copy built from the AnimatLab standalone exports. The conversion rules, applied by `tools/make_spiking.py` to copies (no original file touched), are: thresholds of neurons with unreachable (positive) thresholds move into the spiking regime, either $-55\,\mathrm{mV}$ with $-59\,\mathrm{mV}$ for motor-neuron-named cells in the v1 copies, or per-class calibrated values from the measured graded peaks in the v2 copies (rhythm-generator/pattern-formation/interneuron $-53.5\,\mathrm{mV}$, motor neurons $-44\,\mathrm{mV}$, Renshaw $-44.5\,\mathrm{mV}$); relative accommodation is zeroed; every graded chemical synapse type is re-emitted as a per-spike type with the same identifier and amplitude, 10 ms decay, and maximum conductance $\max(5,\ 2\times\text{amplitude})$; and each connection's own class tag is flipped to the spiking enum. That last rule is a documented format trap: the per-connection `<Type>` tag is stored independently of the synapse type, and converting only the type kills the run at load with no diagnostic detail. Structural validation confirms for every copy that neuron, connection, synapse-type, chart, and stimulus sets are identical to the original and that every class tag matches its synapse type.

Table B4: Complete verification-gate table of the spiking mirror. Baseline scores reproduce digit for digit; "split" denotes a gate whose clauses pass and fail in part.

Gate	Verdict	Headline numbers
Pipeline compile	PASS	runner, curriculum, parameters, and network builder all compile.
Defaults gate	PASS	(410, 376, 1186) populations, inputs, and connections after the selector edit; variant build (888, 382, 7430).
Non-spiking bit-identity	PASS	default baseline score -161.56754173676563 identical before and after the edit (all 17 digits); other baselines -144.46440311078322 and -200.37955834081458.
Topology mirror	PASS	directed edge multisets identical in three gain configurations (1,186 / 7,908 / 7,762); inputs and motor-neuron maps identical; readout taps 12 / 12 / 16.
Constant-drive rhythm	SPLIT	sustained bilateral antiphase (E 4.9 Hz, F 3.2 Hz per leg; swing 5.13 mV in every window) = PASS; period 0.621 s versus analog 0.399 s ($+55.6\,\% > 20\,\%$) = FAIL.
Perturbation basin	PASS	12/12 trials under 1 % jitter on every scalar; final swing minimum 4.58 mV (margin $0.89\times$); zero explosions.
Air-regime smoke	FAIL (cadence)	finite, 53 envelope rises, knee range 106.0° of twin's 117.9°, hip 50.4° of 71.1°; period 0.215 s versus 2.108 s.
Untuned ground evaluation	ran	score -236.35; bilateral (8 right / 9 left cycles); cycle knee minimum -79.8°; no non-finite states; stayed up. Single run, no statistics.
Simulink regression	PASS	all committed demonstrations and units tests reproduce exactly with the spiking blocks in the library (two-neuron test 6.07×10^{-4} mV).
AnimatLab baselines	PASS	five headless baseline runs, exit 0, periods and antiphase reproducing the records (0.443 s, 0.462 s, 1.305 s).

Table B5: AnimatLab spiking-conversion variants of the spiking mirror. "Connections" counts flipped per-connection class tags; the ablation converts thresholds only. All copies validate structurally and run headless to their configured end; rhythm outcomes as in Table D7.

Variant	**Thresholds**	**Neurons**	**Synapse types**	**Connections**	**Rhythm outcome**
Two-layer v1	−55/−59 mV	80	17	150	dead (2 spikes / 5 s)
Two-layer v2	calibrated	80	17	150	dead
Two-layer v3	v1, 30 ms decay	80	17	150	dead
Bilateral v1	−55/−59 mV	96	20	208	latch, one side 442 Hz–450 Hz
Bilateral v2	calibrated	96	20	208	latch
Project export v1	−55/−59 mV	122	50	243	synchronous 11.8 Hz co-bursting
Project export v2	calibrated	122	50	243	dead (passive settle)
Ablation	spiking, graded synapses	80	0	0	oscillates, near-baseline ranges
2023 reference (baseline)	spiking from birth	56	3 (spiking)	84	1.305 s bursts, antiphase +0.494 cycle

Figure B1 shows the ablation result at the joint level for the two-layer family: with graded synapses kept, spiking-regime thresholds leave the joint oscillation essentially intact, while the full conversion pins the joints. This is the visual form of Table D7's conclusion that the synapse conversion is the breaking change.

The spiking copies are headless-run artifacts only: none has been opened in the AnimatLab graphical environment, and graphical verification (or an equivalent acceptance ruling) remains outstanding for every `.asim` listed in Table B5. The mirror work also scoped, but did not attempt, two further conversions: the full 410-population Simulink spinal network (whose every one of 1,392 gains is tuned for graded voltages, making a spiking conversion a full re-tune rather than a parameter port) and the spiking mirrors of the alternative MuJoCo network variants (separate architectures whose fresh tuning is future work).

B.5 Connectome of the Verified W2L Model

Figures B2 and B3 document the W2L configuration used for the independent replays in Section D.5.1. The corrected bilateral template compiles to 96 neurons, 204 chemical synapses, and 22 input channels. The node and connection lists were read from the compiled SNS-Toolbox network using the saved configuration `w2l_verified_config.json`. This is a different network from the 410-neuron Gait2392 model used in the earlier spiking comparisons; their neuron counts and results are not interchangeable.

The two sides each contain 48 neurons. RG denotes rhythm-generator cells; RG-IN, their local interneurons; PF, pattern-formation cells; PF-IN, their local interneurons; MN, motor

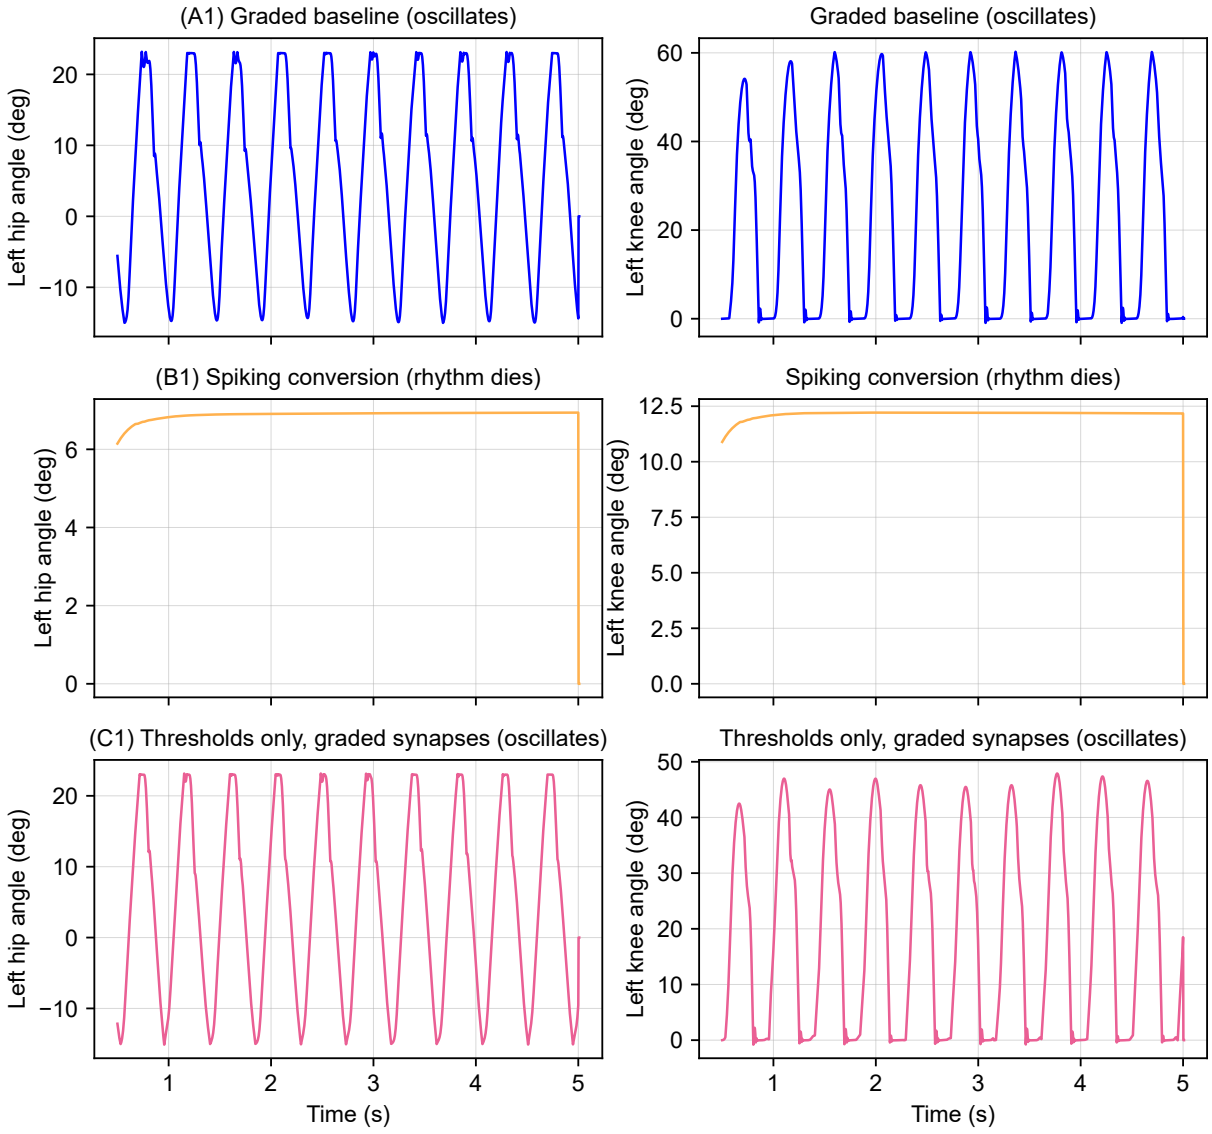


Figure B1: AnimatLab left hip (left column) and knee (right) angles over 0.5 s to 5 s: graded baseline (A, B), full spiking conversion (C, D), and spiking thresholds with graded synapses retained (E, F).

neurons; RC, Renshaw cells; Ia, Ib, and II, muscle-afferent encoders; and Ia-IN, reciprocal-inhibition interneurons. The c1 and V3 populations provide commissural pathways, and Contact comprises heel and toe encoders. The motor outputs drive six actuator channels per side. External inputs and actuator outputs are interfaces, not additional neurons.

Figure B2 groups individual connections by their presynaptic and postsynaptic populations. A symbol indicates at least one connection of that sign; a population pair can contain both excitatory and inhibitory connections. The diagonal quadrants contain within-side pathways, and the other quadrants show interleg pathways. The matrix exposes the sensory projections to both RG and PF populations as well as to motor circuitry.

Figure B3 shows the actual individual maximum conductance values, grouped by postsynaptic population. Equal values overlap. This representation retains the spread of synaptic strengths without implying that every connection in a pathway family shares a single gain. Conductance values alone do not predict the closed-loop contribution of a pathway; that requires the controlled interventions described in Chapter 6.

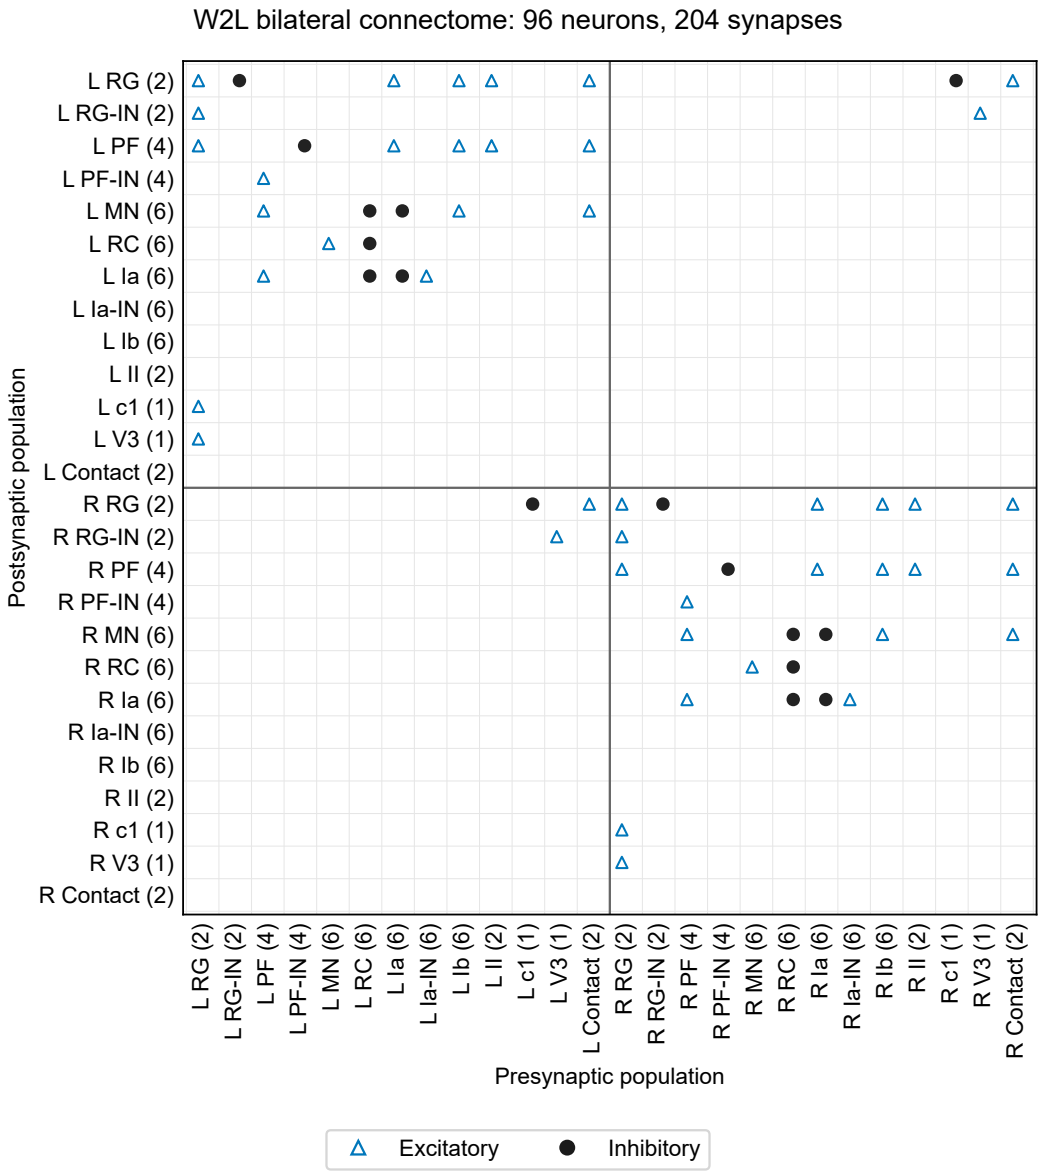


Figure B2: Population connectivity of the verified W2L network. Parentheses give neuron counts. Open triangles denote excitation, and filled circles inhibition.

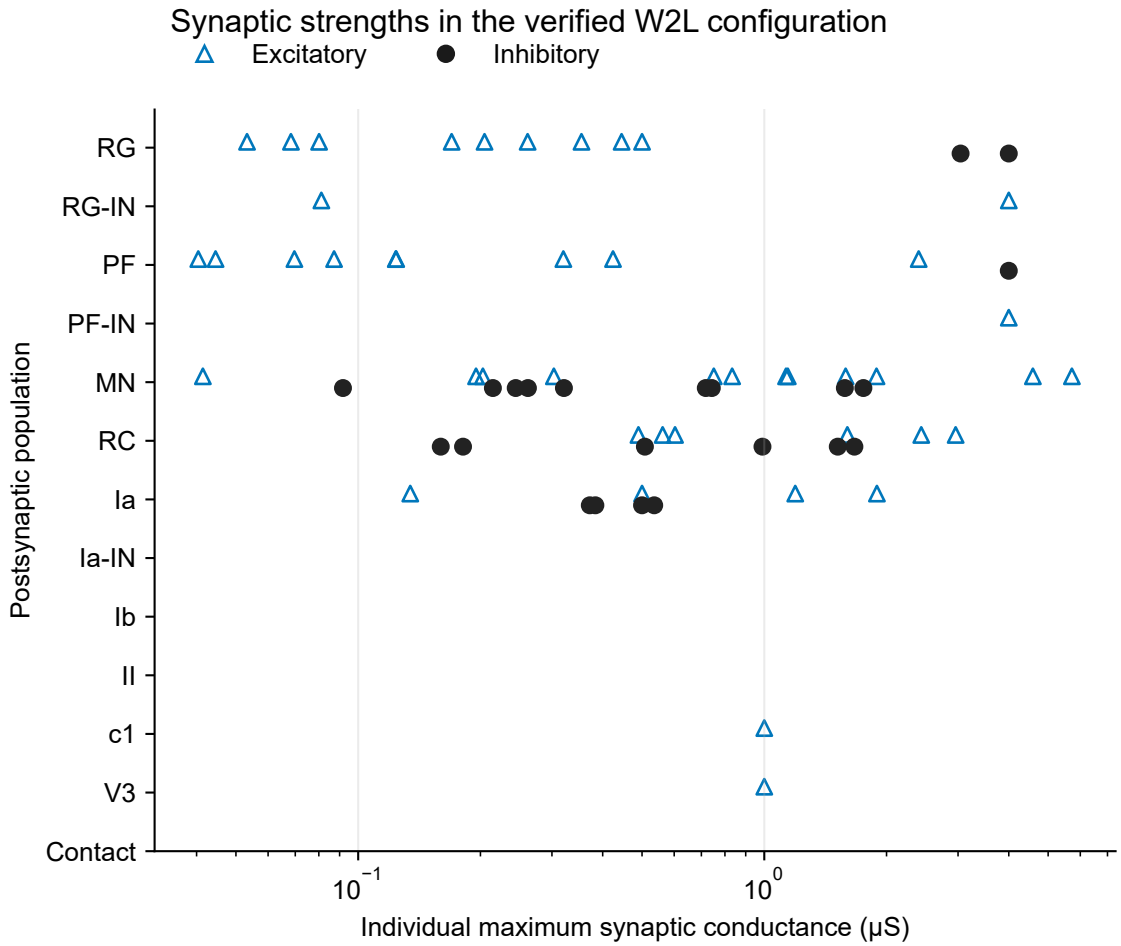


Figure B3: Synaptic conductances in the verified W2L configuration, grouped by target population. The horizontal axis is logarithmic; coincident values overlap.

Appendix C: Model Geometry, Bracket Frames, and Compliance Parameters

This appendix documents the numeric geometry and compliance parameters behind the force and torque models of Chapters 3 and 4 as they exist in the project's MATLAB code, with file names given so that any value can be traced to the code that produced it. Route and bracket coordinates are coded in meters and reported here in millimeters; stiffnesses are coded and reported in newtons per meter. The bracket reference points are interpretation choices for the 3D-printed (Onyx FDM) attachment hardware, set from CAD by me; moving a reference point changes the identified stiffness values, so the points are recorded here alongside the values they produced.

C.1 Static BPA Pathway Data

The static BPA model routes each actuator as a straight-segment polyline through via points stored in a per-test `Location` array of dimensions $M \times 3 \times N$, where N is the number of joint frames in the motion sweep and M the number of via points; the first row in the distal link frame is marked by a cross-point index. There are no binormal or surface-frame quantities: at each frame the force direction is the normalized vector from the cross-point row to the previous row transformed into the crossing frame,

$$\mathbf{d} = \mathbf{R}_{ft}^{\mathsf{T}}(\mathbf{p}_{C-1} - \mathbf{t}_{ft}) - \mathbf{p}_C, \qquad \hat{\mathbf{u}} = \frac{\mathbf{d}}{\|\mathbf{d}\|}, \tag{C1}$$

where $\mathbf{p}_C$ is expressed in the tibia frame, $\mathbf{p}_{C-1}$ in the femur frame, and $\mathbf{T}_{ft} = [\mathbf{R}_{ft}, \mathbf{t}_{ft}; \mathbf{0}^\mathsf{T}, 1]$ maps tibia coordinates to femur coordinates. The inverse transform includes both rotation and translation, as in the implemented point transformation [Provenance: `minimizeFlxPin.m` lines 160 and 524–537 (`UD`) and 201–235 (`computeForceVector`); `minimizeExtX3.m` lines 687–700; BPA struct fields, `minimizeFlxPin.m` lines 28–35]. The pinned-knee test routes are archived as `FlxPinBPASet.mat` (5 flexor tests, $M = 2$, knee angle -120° to $+10^\circ$) and `ExtPinBPASet.mat` (9 extensor tests, $M = 7$, knee angle -125° to $+35^\circ$).

Table C1 lists the five pinned-knee flexor tests; their routes are two points, femur-frame origin $\mathbf{p}_1 = [-75.0,\ 100.0,\ 32.8]$ mm and tibia-frame insertion $\mathbf{p}_2 = [-50.1,\ -45.0,\ 32.3]$ mm (test 1; `FlxPinBPASet.mat`, kf(1), frame 1).

Table C3 lists the extensor pinned-knee route of test 1 at $\theta_\mathbf{k} = -125^\circ$; rows 1–6 are femur-frame and row 7 is tibia-frame, and as the knee extends, rows 2–6 release onto $\mathbf{p}_1$ (wrap released at $+35^\circ$).

Table C4 lists the per-test constants of the nine extensor tests; the correction terms of Section 4.1.4 used tests {1, 2, 5, 6, 7, 8}, of which the adopted fold trained on {5, 6, 7} and held out tests 1, 2, and 8 for validation; tests 3, 4, and 9 were not used in identification; their retrospective predictions are reported below [Provenance: fold columns of `minimizeExt10mmX3_results.mat`; `Dig_crossPredict.m` line 46].

Table C1 lists pinned-knee flexor test BPA constants, read directly from `FlxPinBPASet.mat` (kf(1)–kf(5)). Labels are the laboratory shorthand used in the fitting logs (`minimizeFlxPin10mm_2brk.m` line 51); l_0, tendon, and pressure are the values used by the model. All tests use ∅10 mm BPAs; the 47 cm test's measured angles carry the $+5.3^\circ$ encoder correction at load time (`minimizeFlxPin.m` lines 77–80). The test was not used in the

two-bracket identification, and the corrected-model fit reported for it in Table C11 is a post hoc evaluation; of the remaining four, the adopted cross-validation fold trained on kf(1) and kf(5) and held out kf(2) and kf(4) for validation.

Table C1: Pinned-knee flexor test constants from `FlxPinBPASet.mat`. All actuators have 10 mm diameter; test names use laboratory shorthand.

Test	**Label**	l_0 (m)	Tendon (m)	P (kPa)	F_m (N)	**Use**
kf(1)	48 cm	0.485	0	605.7	443.4	training
kf(2)	46 cm	0.457	0	614.2	441.4	validation
kf(3)	47 cm	0.479	0	623.2	443.0	not used
kf(4)	40 cm-tendon	0.406	0.035	623.4	436.9	validation
kf(5)	41 cm	0.410	0	615.4	437.3	training

The seed in Table C2 follows my femoral-condyle ray construction. The saved direct-route elimination schedule releases p5 at -89.8°, p4 at -55.7°, p6 at -47.8°, p7 at -20.2°, p3 at $+6.1^\circ$, and p8 at $+7.4^\circ$. These values come from `routeInfo.eliminatedAngleD` in `Vas_Pam_20mm_Result_20260925.mat`; the seed construction is in `buildKneeExtContext20mm.m`, lines 431–440.

Table C4 lists pinned-knee extensor test BPA constants, read directly from `ExtPinBPASet.mat` (ke(1)–ke(9)). Labels are the laboratory shorthand used in the fitting logs (`Dig_resultsTable.m` line 32); l_0, tendon, and pressure are the values used by the model; the shorthand rounds several lengths (e.g. "42 cm" is $l_0 = 0.415$ m and "43 cm-tendon" is $l_0 = 0.432$ m). All tests use ∅10 mm BPAs. The correction terms of Section 4.1.4 used tests $\{1, 2, 5, 6, 7, 8\}$: the adopted fold trained on $\{5, 6, 7\}$ and held out tests 1, 2, and 8 for validation; tests 3, 4, and 9 were not used in identification; their retrospective predictions are included in Table C12.

Table C2: Direct-route seed for the 20 mm biomimetic-knee extensor. Points p1–p5 use the femur frame; p6–p9 use the tibia frame.

Point	x (mm)	y (mm)	z (mm)	**Frame**	**Role**
p1	40.00	35.00	0	femur	BPA origin (attachment)
p2	83.90	−274.76	0	femur	BPA contacts mounting base
p3	64.06	−426.80	0	femur	femoral contact
p4	52.83	−441.88	0	femur	femoral condyle contact
p5	37.77	−452.14	0	femur	femoral condyle contact
p6	69.85	34.33	0	tibia	tibial contact (upper)
p7	73.50	19.13	0	tibia	tibial contact (mid)
p8	71.98	−12.23	0	tibia	tibial contact (lower)
p9	34.06	−66.87	0	tibia	patellar ligament ring (distal, pEnd)

Table C3: Pinned-knee extensor BPA route of test 1 at $\theta_{\mathbf{k}} = -125^{\circ}$ (`ExtPinBPASet.mat`, ke(1), frame 1; mm).

Point	x (mm)	y (mm)	z (mm)
p1	30.00	−50.00	0
p2	65.00	−350.00	0
p3	42.40	−408.50	0
p4	28.20	−421.60	0
p5	10.50	−425.20	0
p6	−8.60	−422.80	0
p7	42.50	−75.90	0

Table C4: Pinned-knee extensor test constants from `ExtPinBPASet.mat`. All actuators have 10 mm diameter; test names use laboratory shorthand.

Test	**Label**	$\boldsymbol{l_0}$ (m)	Tendon (m)	P (kPa)	F_m (N)	**Use**
ke(1)	40 cm	0.400	0	624.9	436.3	validation
ke(2)	40 cm-tendon	0.400	0.040	618.4	436.3	validation
ke(3)	42 cm	0.415	0	605.2	437.8	not used
ke(4)	42 cm-tendon	0.415	0.040	605.2	437.8	not used
ke(5)	43 cm	0.436	0	621.8	439.7	training
ke(6)	43 cm-tendon	0.432	0.022	622.9	439.3	training
ke(7)	46 cm	0.457	0	605.9	441.4	training
ke(8)	47 cm	0.465	0	605.9	442.0	validation
ke(9)	48 cm	0.480	0	605.9	443.1	not used

The biomimetic flexor route (`buildKneeFlexorContext20mm.m`) is a three-point path (femur-frame origin $\mathbf{p}_1$, a moving wrap point $\mathbf{p}_{\text{wrap}}$ on the tibia, and tibia-frame insertion $\mathbf{p}_{\text{end}}$), with the wrap point held on a standoff surface of radius `bpaRb` about the joint and its z placed on the projected $\mathbf{p}_1$–$\mathbf{p}_{\text{end}}$ line; the optimizer seed is $\mathbf{p}_{1,0} = [-50.0,\ 35.0,\ 50.0]$ mm and $\mathbf{p}_{2,0} = [-12.24,\ -8.87,\ 27.87]$ mm with $l_0 = 415$ mm and 25 mm tendon [Provenance: `buildKneeFlexorContext20mm.m` lines 209–218; `buildKneeFlexorRoute20mm.m` lines 98–186]. The biomimetic extensor seed route is given in Table C2, including the release schedule of the route-elimination rule.

The biomimetic-knee evaluator inputs are listed in Tables C5 and C6. These tables report the saved model dimensions loaded by `minimizeFlx` and `minimizeExt`; the laboratory test names are nominal labels. The reduced-pressure flexor row identifies the reference model used with the separate pressure measurements.

Table C5: Biomimetic-knee flexor inputs loaded by `minimizeFlx`. Tendon length is the effective value stored in the saved actuator model.

Model	Diameter (mm)	l_0 (mm)	Tendon (mm)	P (kPa)	F_m (N)
10 mm	10	415.0	11.0	603.5	437.8
20 mm, full pressure	20	415.0	14.0	620.0	1302.9
20 mm, reduced pressure	20	415.0	14.0	325.0	1302.9

Table C6: Biomimetic-knee extensor inputs loaded by `minimizeExt` from the saved 10 mm actuator model.

Model	Diameter (mm)	l_0 (mm)	Tendon (mm)	P (kPa)	F_m (N)
10 mm extensor	10	520.0	38.0	596.4	445.7

Table C7: Measured dimensions of the four BPAs used in the final physical tests. Resting lengths and tendon lengths are recorded separately for each actuator in the extensor and flexor test worksheets. All four BPAs have 20 mm diameter.

Actuator	Resting length (mm)	Tendon length (mm)
Flexor 1	561	28
Flexor 2	560	30
Extensor 1	828	92
Extensor 2	826	85

The prediction inputs are distinct from the measured dimensions in Table C7. The flexor predictions retain a shared optimized resting length of 561.49 mm and an effective tendon length of 5.00 mm, obtained from the 25 mm route parameter less a 20 mm fitting adjustment; the physical-test plot uses the nominal 562 mm resting-length label. The extensor predictions retain a shared resting length of 827.74 mm and the approximate 92.49 mm tendon parameter. The reported physical-test comparisons use these original prediction inputs; the measured dimensions above do not constitute a new model evaluation.

C.2 Joint Transformation Matrices

The knee transform is planar: every joint frame is a rotation about z by the knee angle $\theta_{\mathbf{k}}$ composed with a translation fitted to robot CAD tables of the migrating instantaneous center of rotation. With $\mathbf{R} = \mathbf{R}_z(\theta_{\mathbf{k}})$, the robot (femur→ICR) transform is $\mathbf{T}_{\text{Pam}} = [\mathbf{R},\ [f_3(\theta_{\mathbf{k}}),\ f_4(\theta_{\mathbf{k}}),\ 0]^\top;\ \mathbf{0}^\top\ 1]$ with f_3, f_4 cubic-spline fits (ICR x: 23.30 to 10.47 mm; y: -416.65 to -395.66 mm over $\theta_{\mathbf{k}} = 0.17$ to -2.62 rad), and the ICR→tibia transform is the pure translation $[f_{13}(\theta_{\mathbf{k}}),\ f_{14}(\theta_{\mathbf{k}}),\ 0]^\top$ (x: 29.66 to -7.58 mm; y: 25.97 to 22.33 mm), composed as $\mathbf{T}_{t1\to f} = \mathbf{T}_{t1\to\text{ICR}}\ \mathbf{T}_{\text{Pam}}^{-1}$ over 100 frames spanning $\theta_{\mathbf{k}} \in [-120^\circ, +10^\circ]$ [Prove-

nance: `buildKneeFlexorContext20mm.m` lines 34–80; `buildKneeExtContext20mm.m` lines 30–78; `Functions/ModernRobotics/RpToTrans.m`]. In the archived test sets the per-frame transform maps tibia to femur/ICR and takes the form

$$\mathbf{T}_{\mathbf{k}}(\theta_{\mathbf{k}}) \;=\; \begin{bmatrix} \cos\theta_{\mathbf{k}} & -\sin\theta_{\mathbf{k}} & 0 & d_x(\theta_{\mathbf{k}}) \\ \sin\theta_{\mathbf{k}} & \cos\theta_{\mathbf{k}} & 0 & d_y(\theta_{\mathbf{k}}) \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}, \tag{C2}$$

with rotation block $\mathbf{R}_z(\theta_{\mathbf{k}})$ and translation components d_x, d_y in meters. In the archived flexor test set the translation is constant across the sweep, $[d_x,\, d_y,\, 0] = [4.5,\, -395.8,\, 0]$ mm, and the rotation sweeps from $\mathbf{R}_z(-120^\circ)$ at full flexion to $\mathbf{R}_z(+10^\circ)$ at full extension (`FlxPinBPASet.mat`, kf(1)).

C.3 Bracket Reference Points and Frames

Table C8 lists the bracket reference points. Each compliant bracket carries an instantaneous body-fixed frame anchored at its reference point, constructed by two rotations (yaw about z, then pitch about the yawed y). With $\mathbf{p}_b$ the vector from the reference point to the bracket-to-BPA attachment target, the yaw and pitch angles are

$$\theta_{br} = \operatorname{atan2}(p_{by},\, p_{bx}), \qquad \phi_{br} = \operatorname{atan2}(p'_{bz},\, p'_{bx}) \quad \text{with} \quad \mathbf{p}'_b = \mathbf{R}_z(\theta_{br})^\top \mathbf{p}_b, \tag{C3}$$

and the bracket-frame pose is

$$\mathbf{R}_{kbr} = \mathbf{R}_z(\theta_{br})\,\mathbf{R}_y(\phi_{br})^\top, \qquad \mathbf{T}_{kbr} = [\mathbf{R}_{kbr},\ \mathbf{p}_{br};\ \mathbf{0}\ 1]\,, \tag{C4}$$

The frame's x axis follows the nominal load path, with angles calculated from the bracket geometry. The adopted pinned-flexor result is row 77 of `minimizeFlxPin10_results_20260908_2brkt_2trans_noT3.mat`: both brackets use the two-rotation construction of Equation (C4). The pinned extensor also uses this construction. Forces are transformed into each bracket frame, and displaced attachment points are transformed back into the limb frame. The symmetric insertion stiffness array makes its compliance invariant to the additional rotation; the origin array does not have that invariance (Table C9).

Table C8: Bracket reference points (mm, as coded; divided by 1000 in code). All are author-set CAD interpretation points; frame denotes the frame the vector is expressed in.

Point	x	y	z	**Configuration / meaning**
Pbri	−48.11	−107.81	13.80	pinned-knee flexor, insertion bracket (knee-ICR frame)
Pbr2	−52.61	0	75.06	pinned-knee flexor, origin-side 2nd bracket (hip frame)
Pbr	−3.84	−46.44	62.50	pinned-knee extensor, rib midpoint, medial side (hip frame)
Pbr	8.38	20.75	25.10	biomimetic extensor 10 mm, cantilever centroid
Pbr	9.48	−33.38	33.90	biomimetic flexor 20 mm, bolt-pattern centroid
Pbr	−19.00	22.00	27.60	biomimetic flexor 10 mm, cantilever centroid

C.4 Stiffness Arrays and Compliance

Every configuration models bracket compliance as a diagonal Hooke-law spring in the bracket frame, $\mathbf{K}_{br} = \mathrm{diag}(k_1, k_2, k_3)$ newtons per meter. The scalar compliance along the

force path and its series combinations are

$$c_b = \hat{\mathbf{u}}^\top \mathbf{K}_{br}^{-1} \hat{\mathbf{u}}, \qquad c_{\text{eff}} = c_b + c_{\text{ten}} \ (+\, c_{b2} \text{ for the second bracket}), \tag{C5}$$

exactly as in Equations (4.9)–(4.19); when an artificial tendon is present its axial spring ($\times 2$ cables, effective steel area 1.51×10^{-6} m^2, $E = 193$ GPa) adds in series, and the force balance of Equation (4.23) is solved per frame by `fzero` on the contraction interval [Provenance: stiffness arrays: `minimizeFlxPin.m` lines 387 and 400; `minimizeExtX3.m` lines 575–584; `minimizeExt.m` line 352; `minimizeFlx.m` line 314. `fzero` force balance: `minimizeFlxPin.m` line 438; `minimizeExtX3.m` line 611; `minimizeExt.m` line 389; `minimizeFlx.m` line 352]. The per-configuration stiffness arrays and their series combinations are given in Table C9.

Table C9 lists bracket stiffness arrays by configuration ($K = [k_1, k_2, k_3]$ in bracket-frame axes, N/m; optimizer solves X_1, X_2 assigned as shown), as executed by the evaluators of record. The flexor evaluator has no frame switch: both brackets use the pitch-only assembly (see the Bracket Reference Points and Frames section above). The pinned-extensor evaluator carries a `transMode` argument (default `2trans`, env `EXTX3_TRANS`); its axis assignment was revised by the author, with the previous assignment kept commented in the file.

The wrapping-loss term X_3 (Equations 4.25–4.27) is evaluated as $X_3\, R\, |\Delta\theta|\, \zeta^2$ per contact arc. On the pinned extensor the two physical contact arcs use $R_1 = 12$ mm (27° to -23°) and $R_2 = 40$ mm, with arc-length normalizers $27.75/(ang_1 - ang_2)$ and $92.46/(ang_2 + 120^\circ)$; on the biomimetic configurations the bend measure is $w_{\text{rap}}|\theta_{\text{turn}}|$ with $w_{\text{rap}} = 25$ mm [Provenance: `minimizeExtX3.m` lines 333–359; `buildKneeFlexorRoute20mm.m` line 268;

Table C9: Bracket-frame stiffness arrays by configuration, in N/m. The fitted variables X_1 and X_2 are assigned to the axes shown.

Configuration	$\boldsymbol{K}$ array	**Source**
Pinned flexor, single bracket (legacy)	$[X_1,\ X_2,\ X_1]$	`minimizeFlxPinX3.m` line 267
Pinned flexor, insertion bracket (pitch-only frame)	$[X_1,\ X_2,\ X_1]$	`minimizeFlxPin.m` line 387 (frame line 289)
Pinned flexor, origin bracket (pitch-only frame)	$[X_1,\ X_1,\ X_2]$	`minimizeFlxPin.m` line 400 (frame line 305)
Pinned extensor, 2trans (with X_3)	$[X_1,\ X_2,\ X_2]$	`minimizeExtX3.m` line 575
Biomimetic extensor	$[X_1,\ X_2,\ X_1]$	`minimizeExt.m` line 352
Biomimetic flexor	$[X_1,\ X_2,\ X_2]$	`minimizeFlx.m` line 314

`MonoPamDataExplicit_balanceX3.m` lines 489–494]. Table C10 records the optimizer parameterization and bounds of the adopted identification: the two-bracket flexor fit solves all three terms, and the extensor refit locks (X_1, X_2) to the adopted flexor-front values (the extensor front was generated with the pair fixed to the adopted values of Table 4.9 and only (X_0, X_3) searched), enforcing a single stiffness pair across configurations.

The physical correction terms X_i used elsewhere in this appendix correspond to χ_i; the lowercase x_i below are the optimizer coordinates. The initial flexor population in Table C10 used $X_0 \in [5, 15]$ mm, $X_1 \in [5 \times 10^4, 5 \times 10^5]$ N/m, and $X_2 \in [7 \times 10^3, 1.5 \times 10^4]$ N/m.

[Provenance: `minimizeFlxPin10mm_2brk.m` lines 82–83 (bounds) and 424–432 (parameterization); `minimizeExt10mmX3.m` lines 23–26 (lock load) and 67–68 (bounds); `Collect_ExtPinX3_sweep.m` lines 59–60.]

Table C10: Optimizer coordinates and equivalent physical correction-term bounds. The numerical optimizer variables are x_i; length offsets χ_0 are reported in millimeters, stiffnesses in N/m, and χ_3 is dimensionless. The extensor refit fixes the stiffness pair at the adopted flexor values.

Identification	Variable	Optimizer bounds	Physical bounds / mapping
Pinned flexor	x_0	$[0, 2]$	$\chi_0 = 10x_0 \in [0, 20]$ mm
	x_1	$[\log_{10}(3 \times 10^4), 6]$	$\chi_1 = 10^{x_1} \in [3 \times 10^4, 10^6]$ N/m
	x_2	$[\log_{10}(5 \times 10^3), \log_{10}(2 \times 10^4)]$	$\chi_2 = 10^{x_2} \in [5 \times 10^3, 2 \times 10^4]$ N/m
Extensor refit	x_0	$[-2, 0.5]$	$\chi_0 = 10x_0 \in [-20, 5]$ mm
	x_3	$[0, 1]$	$\chi_3 = x_3 \in [0, 1]$

C.5 Adopted Correction Terms and Complete Pinned-Knee Results

The adopted correction terms are listed in Table 4.9. The flexor uses row 77 of `filtered_results` in `minimizeFlxPin10_results.mat`; the extensor uses row 32 in `minimizeExt10mmX3_results.mat`. The extensor identification holds (χ_1, χ_2) at the flexor values and solves (χ_0, χ_3), preserving one effective stiffness pair across both configurations.

Tables C11 and C12 report every recorded pinned-knee test before and after applying the adopted corrections. "Before" uses the characterized BPA force model with zero length and wrapping corrections and rigid attachments; "after" uses the adopted correction terms. Both predictions use the same measured angles and test inputs. RMSE and maximum absolute residual have units of N m, and FVU is dimensionless. The sample-size convention for FVU is given in Section C.9.

The adopted flexor fold uses tests 1 and 5 for training and tests 2 and 4 for validation. Test 3 (47 cm) was not used in identification. The extensor fold uses tests 5, 6, and 7 for training and tests 1, 2, and 8 for validation. Tests 3, 4, and 9 (42 cm, 42 cm with tendon,

and 48 cm) were not used because the attachment bracket shifted during those early tests. Results for unused tests are retrospective predictions, not held-out validation results.

Table C11: Complete pinned-knee flexor test results before and after correction. Test names follow the laboratory shorthand.

		Before			After		
Test	Use	RMSE	FVU	Max.	RMSE	FVU	Max.
1: 48 cm	Training	7.617	2.1278	25.531	2.165	0.1718	5.467
2: 46 cm	Validation	8.803	1.9380	26.264	1.368	0.0468	3.614
3: 47 cm	Unused	7.670	1.8463	21.650	2.287	0.1641	7.818
4: 40 cm+tendon	Validation	6.373	0.5935	16.859	1.611	0.0427	5.891
5: 41 cm	Training	9.694	2.1149	21.076	1.084	0.0277	1.766

Table C12: Complete pinned-knee extensor test results before and after correction. Test names follow the laboratory shorthand.

		Before			After		
Test	Use	RMSE	FVU	Max.	RMSE	FVU	Max.
1: 40 cm	Validation	3.092	0.5638	4.842	1.244	0.1303	2.196
2: 40 cm+tendon	Validation	3.622	6.3919	7.845	1.350	0.8878	1.890
3: 42 cm	Unused	3.528	2.2190	7.665	1.299	0.3418	2.700
4: 42 cm+tendon	Unused	2.910	6.2250	6.367	1.289	1.2209	2.003
5: 43 cm	Training	2.437	0.9618	4.284	0.978	0.1634	1.818
6: 43 cm+tendon	Training	2.878	4.7304	4.817	1.525	1.3270	2.481
7: 46 cm	Training	2.178	1.0586	4.974	0.490	0.0535	1.206
8: 47 cm	Validation	0.949	0.3152	2.272	0.699	0.1714	1.559
9: 48 cm	Unused	0.578	0.1085	1.569	0.784	0.1996	2.282

The adopted corrections reduce all three errors on every test used in identification. The unused 48 cm extensor test has larger errors after correction, as shown in Table C12. The 43 cm-tendon extensor test retains an FVU above unity despite improvement over its baseline. These outcomes bound the accuracy of the shared correction model.

The values were evaluated with `minimizeFlxPin.m` and `minimizeExtX3.m`. On the biomimetic extensor test ($l_0 = 51.8\,\text{cm}$), the shared adopted terms give RMSE 1.17 N m, FVU 0.67, and maximum residual 2.21 N m, compared with the rigid-model values 3.03 N m, 4.49, and 5.59 N m. Table C13 collects the biomimetic comparisons at the adopted terms.

C.6 Redesigned Extensor Route Geometry

The direct redesigned extensor route is preserved in `Vas_Pam_20mm_Result_20260925.mat`; the cam-assisted route supersedes it as the design of record. Figure C1 shows the direct route at full flexion, immediately before each contact release, and at full extension. The plotted event pose precedes the first inactive grid pose used to report each release angle in Table C2. All route points, including p8, are present in the saved geometry; plotting-title metadata do not define the route.

The optimized origin is $\mathbf{p}_1 = [40.0, 137.0, 0]$ mm in the femur frame, and the insertion is $\mathbf{p}_{\text{End}} = [34.4, -197.1, 0]$ mm in the tibia frame. Resting length is 827.7 mm, and tendon length is 36.8 mm. The displayed envelopes and bend-radius checks explain the route's geometric constraints.

Figure C2 places the same optimized route on the OpenSim bone-surface point clouds (`Open_Sim_Bone_Geometry`) at exactly zero knee angle: the full lower skeleton (femur, tibia, pelvis and hip, sacrum, lumbar spine, and foot) is drawn behind the three human vastus reference paths in light orange, the optimized p1-to-pEnd route is drawn in magenta with its endpoints circled, and the optimized route carries a total BPA-plus-tendon path length

of 0.784 m at this pose. Animated sweeps of both redesigned routes over the full flexion range are archived with the optimization results (`Knee_Extensor_20mm.gif`, 2.3 MB, and `Knee_Flexor_20mm.gif`, 4.7 MB, in `Code/Matlab/Mesh_Optimization/Results`).

C.7 Contact and Release Angles of the Wrapped Route

The route-elimination rule that schedules the wrapped extensor route is a route-topology state machine rather than a shortest-path solver: the fully flexed route carries every candidate contact, and each contact is released at the first more-extended knee angle at which the straight chord between its surviving neighbors can replace it without entering the exclusion envelopes. This section states the calculation in generic terms, so that it applies to any wrapped actuator route on a rotating joint pair; the extensor route of Table C2 is the worked example, and the release schedule the calculation produces is the one quoted there.

The route is stored as a fixed number of rows $p_1, \ldots, p_9$: rows p_1 through p_5 are femur-frame points, rows p_6 through p_9 are tibia-frame points, and the segment from p_5 to p_6 crosses from one body frame into the other. The endpoints p_1 and p_9 are the actuator origin and insertion, and they are always active. Each interior row is either active, a genuine wrap contact, or eliminated; an eliminated femur-side row repeats the preceding active femur row, and an eliminated tibia-side row repeats the neighboring active tibia row, so the route keeps a fixed size and its segment chain crosses the joint exactly once. Let $\mathcal{A}(\theta_{\mathbf{k}}) \subseteq \{2, \ldots, 8\}$ denote the active interior rows at knee angle $\theta_{\mathbf{k}}$. The sweep over the design grid, here 100 poses from -120° to 10° at a spacing of 1.31°, proceeds from full

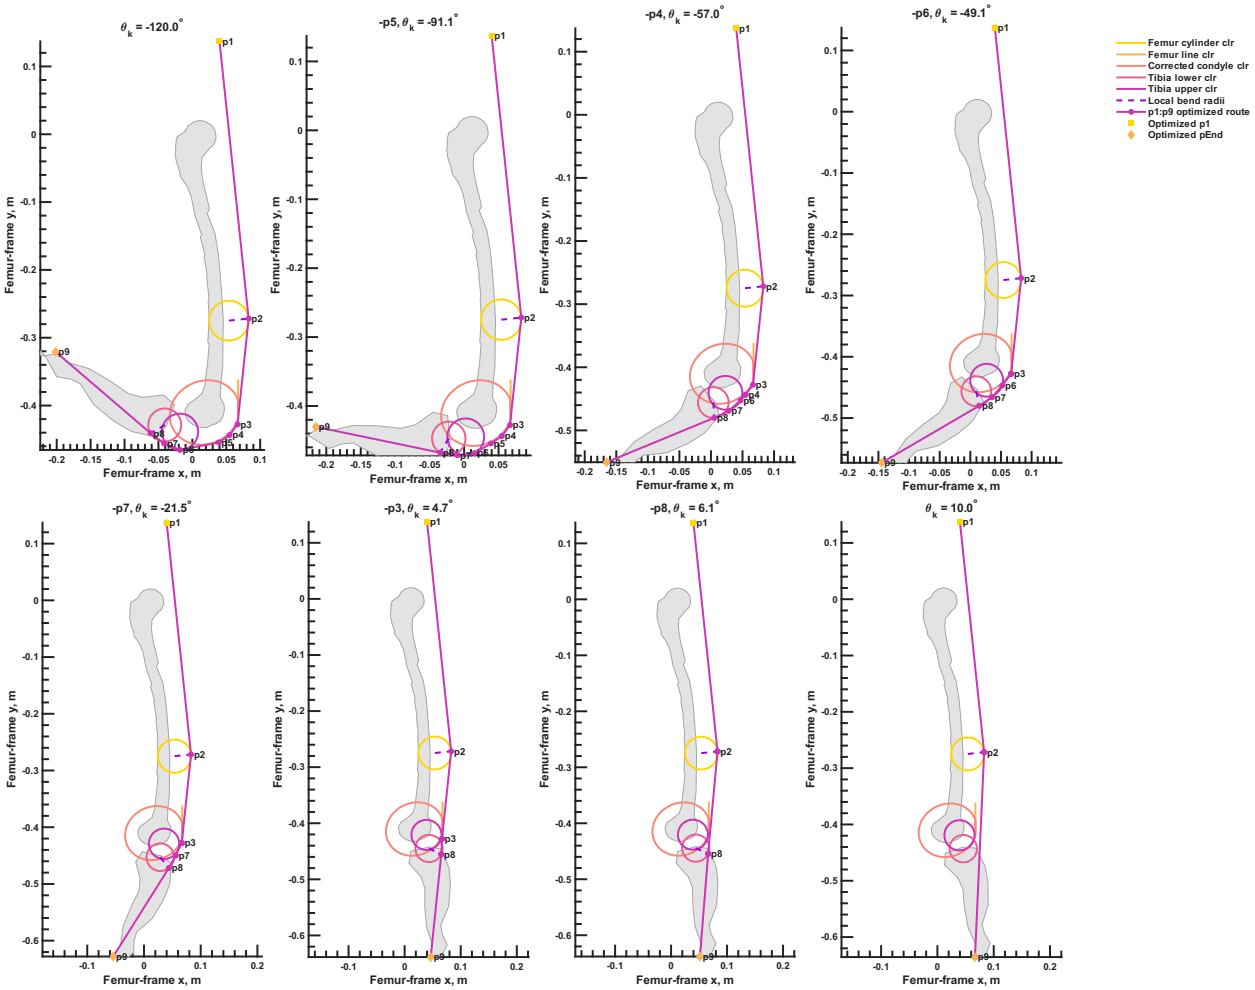


Figure C1: Eight-panel route-geometry view of the optimized 9-point extensor route of record (`Vas_Pam_20mm_Result_20260925.mat`), drawn in the femur frame. Gray silhouettes show registered OpenSim femur and proximal tibia as anatomical context; the colored curves are the robot route and clearance envelopes.

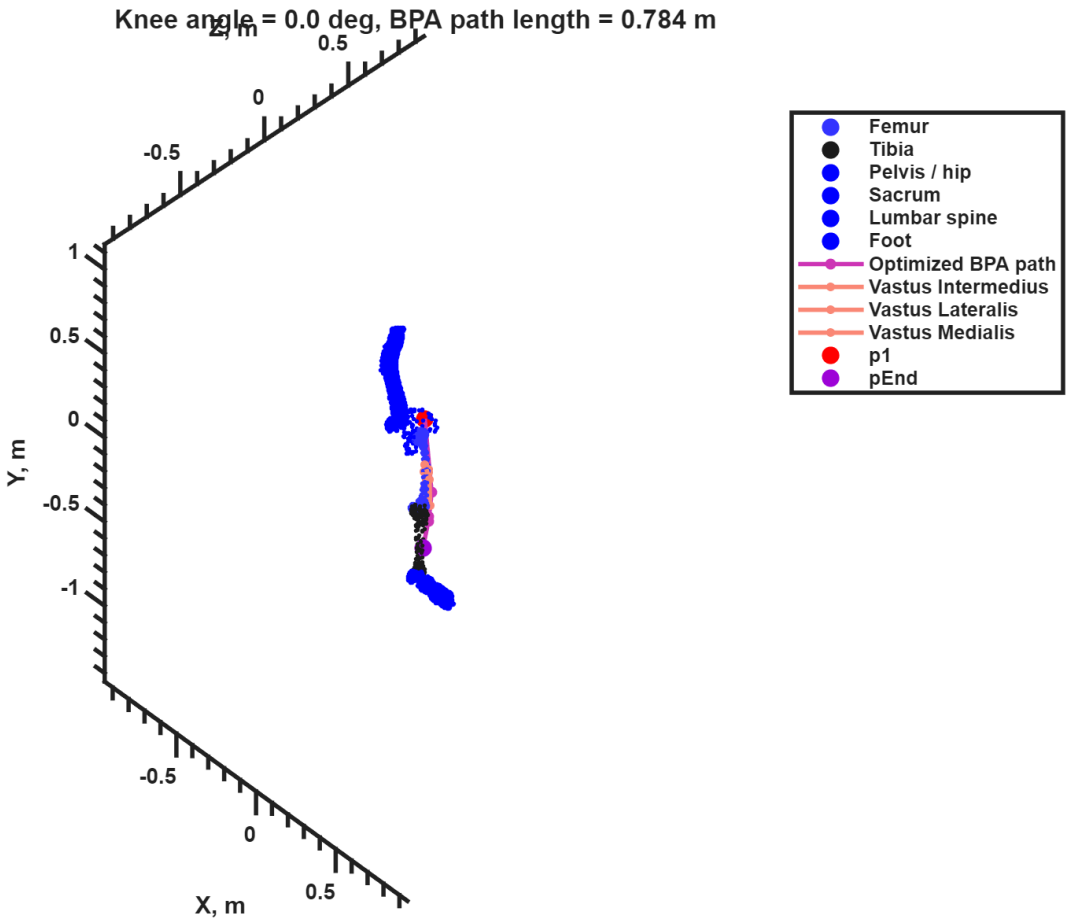


Figure C2: Static bone-mesh view of the optimized 20 mm extensor route at $\theta_k = 0^\circ$, rendered by `Export_AppendixBoneMesh.m`.

flexion toward extension, each pose inherits the active set of the previous pose, and rows are only ever removed, so $\mathcal{A}$ shrinks monotonically as the knee extends.

At every pose, every active interior row is tested between its nearest active anchors,

$$i_j^- = \max\{i < j : i \in \mathcal{A} \cup \{1\}\}, \qquad i_j^+ = \min\{i > j : i \in \mathcal{A} \cup \{9\}\}, \tag{C6}$$

with the anchors expressed in the candidate's native frame through the per-angle joint transforms. Write A, B, and C for the planar, sagittal positions of $p_{i_j^-}$, $p_{i_j^+}$, and p_j, and let $s(x, y)$ denote the signed angle from vector x to vector y, positive counterclockwise, evaluated as a single atan2 of the planar cross product and the dot product so that no principal-value seam can corrupt the comparison. The elimination margin of row j is

$$\mu_j(\theta_{\mathbf{k}}) = \begin{cases} s(C - A,\ B - A)\,, & j \leqslant 5, \\ \\ s(A - B,\ C - B)\,, & j \geqslant 6, \end{cases} \tag{C7}$$

a signed angle between the through-contact leg and the bypass chord, measured in the native wrap sense of the candidate's side of the crossing segment; the two cases differ only in that sense of rotation, because the femur-side contacts turn the route one way around the distal femur and the tibia-side contacts turn it the opposite way around the proximal tibia. The rotation rule releases the contact once the margin turns positive, the orientation at which the chord, rather than the contact, provides the outer boundary of the route on that side. The gate is opened slightly early,

$$\mu_j(\theta_{\mathbf{k}}) > -\Delta, \qquad \Delta = 1.5^\circ, \tag{C8}$$

as a release hysteresis: in the last poses before the contact becomes collinear with its chord, the computed moment arm of the wrapping segment collapses, and releasing the contact slightly before that orientation, with the collision gate below still controlling validity, removes the artifact without changing the route family.

Opening the rotation gate is necessary but not sufficient. A removal additionally requires the straight bypass chord $\overline{AB}$ to clear the exclusion envelopes of the candidate's side,

$$\operatorname{dist}\big(\overline{AB}_{\text{trim}},\, \mathcal{E}_j\big) \;\geqslant\; g, \qquad g = 0.5\,\text{mm}, \tag{C9}$$

where dist is the planar clearance between the trimmed chord and the envelope and tangency is allowed. The chord is trimmed by 2 mm at each end, capped at 45 % of its length, because a legitimate contact departs from the clearance boundary itself. On the femur side the envelopes $\mathcal{E}_j$ are the femoral-cylinder clearance circle, the condyle clearance polygon, and the femoral-wall span; on the tibia side they are the two tibial-cylinder clearance circles and the wall band between them. Every envelope is the normal offset of a bone profile by the fully inflated BPA radius R_{BPA},

$$\mathbf{q}(\vartheta) = \mathbf{e}(\vartheta) + R_{\text{BPA}}\,\hat{\mathbf{n}}(\vartheta), \tag{C10}$$

where $\mathbf{e}(\vartheta)$ traces the physical profile and $\hat{\mathbf{n}}(\vartheta)$ is its outward normal; the offset curve shares the tangent of the physical profile, its radius of curvature is $\rho(\vartheta) = \rho_{\text{phys}}(\vartheta) + R_{\text{BPA}}$, and the candidate contact positions themselves are seeded at tangent points of this curve. Because the envelopes are drawn at the fully inflated radius while the routed BPA is contracted at the release angles, each envelope radius may be relaxed by $\rho = 3\,\text{mm}$ inside

the gate, floored at the bone profile plus 0.5 mm; a rotation-gate opening that still fails the relaxed collision gate indicates that the seed placement, not the gate, must move.

At the first pose of the sweep, full flexion, the seed route is settled to a fixpoint: every contact whose two gates already open there is removed, and removals repeat until no further gate opens, so that the sweep starts from a valid topology; a colinearity guard may additionally remove a seed contact that starts the sweep already aligned with its neighbors, and its tolerance is zero in the design of record, so every seed row of Table C2 starts active. At every later pose at most one contact is released: if several rows open both gates in the same pose, the row with the largest elimination margin, the one whose rotation rule is most decisively open, is released first, and the remaining rows are retested at the next pose. The release angle of row j is recorded as the first grid pose at which the row is inactive. The activation angle is read from the same record with the sweep reversed, as the last pose at which the row is active, which is the pose at which the contact re-engages when the route is driven from extension toward flexion.

The same active-set record drives the geometric bend measure that scales the fictive-strain wrapping loss of Equations (4.25) and (4.27) inside the stiffness-aware force balance. Each wrap family, the femoral-cylinder contact, the femoral-condyle curve, the tibial contact curve, and the crossing-point curve at the ICR, contributes an arc length over its effective radius, and the total

$$b(\theta_{\mathbf{k}}) = \sum_{f} R_f(\theta_{\mathbf{k}})\, \alpha_f(\theta_{\mathbf{k}}) \tag{C11}$$

is the bend signal handed to the force model. Each family's turn angle α_f ramps with

knee angle between the two route-state events that bound it, holding zero before the first event and its plateau after the second, so the topology state and the bending signal are two readings of one record. The formulation is joint-agnostic: it requires only a fixed route representation with one crossing segment, per-angle transforms between the two body frames, and clearance envelopes on each side, so it carries unchanged to the other wrapped routes of this program.

C.8 Supporting Route Searches and Transmission Sweeps

The transmission studies extend the route search by allowing an artificial tendon to pass through a tackle before insertion. A redirect changes the force direction; a multi-part tackle also changes the relation between actuator travel and insertion force. The comparisons below re-solve actuator sizes for each transmission configuration and evaluate insertion-side torque against the same human reference. They remain model predictions; physical tests of the redesigned flexor and cam extensor are reported in Section 4.3.5.

The robot-body model itself carries one update from this work. The adductor magnus attachment point P2 was restored to its full-precision location $(-0.1357, -0.0929, 0.0591)$ m, displacing a rounded copy that was disqualified for producing a zero-length via point that nullified the flexion torque; the regenerated mirrored model passed its validation with zero non-mirrored pairs. At the restored location, the muscle delivers 34.28 N m of flexion torque, $1.12\times$ the human target of 30.61 N m, and 27.62 N m of adduction torque, $0.82\times$ the human target of 33.60 N m. These redesigned routes remain optimizer-stage results: the robot-body OpenSim model contains no pulley entities, and no via points or

markers were added to it for these routes (its eight pre-existing moving-path-point vias, the rectus-femoris and vastii patella-tracking points, one per muscle per leg, belong to the original muscle paths). The MuJoCo conversion used for the simulation work predates the pulley batch and the adductor-magnus restoration, so it carries one stale path point, the pre-restoration rounded adductor-magnus P2. Physical torque tests of the redesigned routes are reported in Section 4.3.5.

The same redesign pipeline was then extended from the knee to all 27 muscle-actuators of the robot-body model, and this whole-body extension is reported as work in progress rather than a completed result. The first pass reached 19 of the 27 actuators before the run stopped, and its chosen solutions showed that the penalty-based formulation itself was the binding limitation: all but one of the chosen routes violated the feasibility constraints, the sole exception being the biceps femoris long-head pair, whose worst-angle margin of $+0.03\,\%$ is indistinguishable from zero. The formulation was then restructured so that the task envelope, the joint-angle window each route must serve, acts as a hard feasible set rather than a penalty. In the restructured pass, twelve of the 27 actuators terminated on a concatenation error in the route builder (the three lumbar-wall actuators, tibialis posterior, the seven distal-foot actuators, and adductor magnus), leaving fifteen finalized routes. Eight of the fifteen were chosen from constraint-feasible configurations: semimembranosus, biceps femoris long head, tensor fasciae latae, soleus, lateral gastrocnemius, tibialis anterior, iliacus, and psoas. Of those eight, six also met the torque requirement at every angle, with worst-angle margins spanning $+0.25\,\%$ (biceps femoris long head) to $+250.7\,\%$ (tensor fasciae latae); the semimembranosus and soleus routes satisfy the geometry but remain short of their torque targets. The crash set and the two torque-short feasible routes define

the remaining engineering: the route builder needs a fix for the concatenation failure, and the short routes need either larger actuators or relaxed corridors.

The unadjusted cam optimizer solution had minimum margins of $+31.30\,\%$ against the required envelope and $+42.34\,\%$ against the plain human target; the adjusted design used in Chapter 4 has margins of $+30.83\,\%$ and $+39.77\,\%$, respectively.

Figure C3 shows predicted insertion-torque magnitude versus joint angle and supply pressure for two BPAs through a two-part tackle (force gain 1, travel gain 2), with actuator sizes re-solved. Panels show soleus, tibialis anterior, biceps femoris short head (BFSH), medial gastrocnemius with ankle fixed at zero, and rectus femoris evaluated at the hip. Dashed curves mark equality with the maximum-isometric OpenSim requirement; no equality curve lies within the soleus panel.

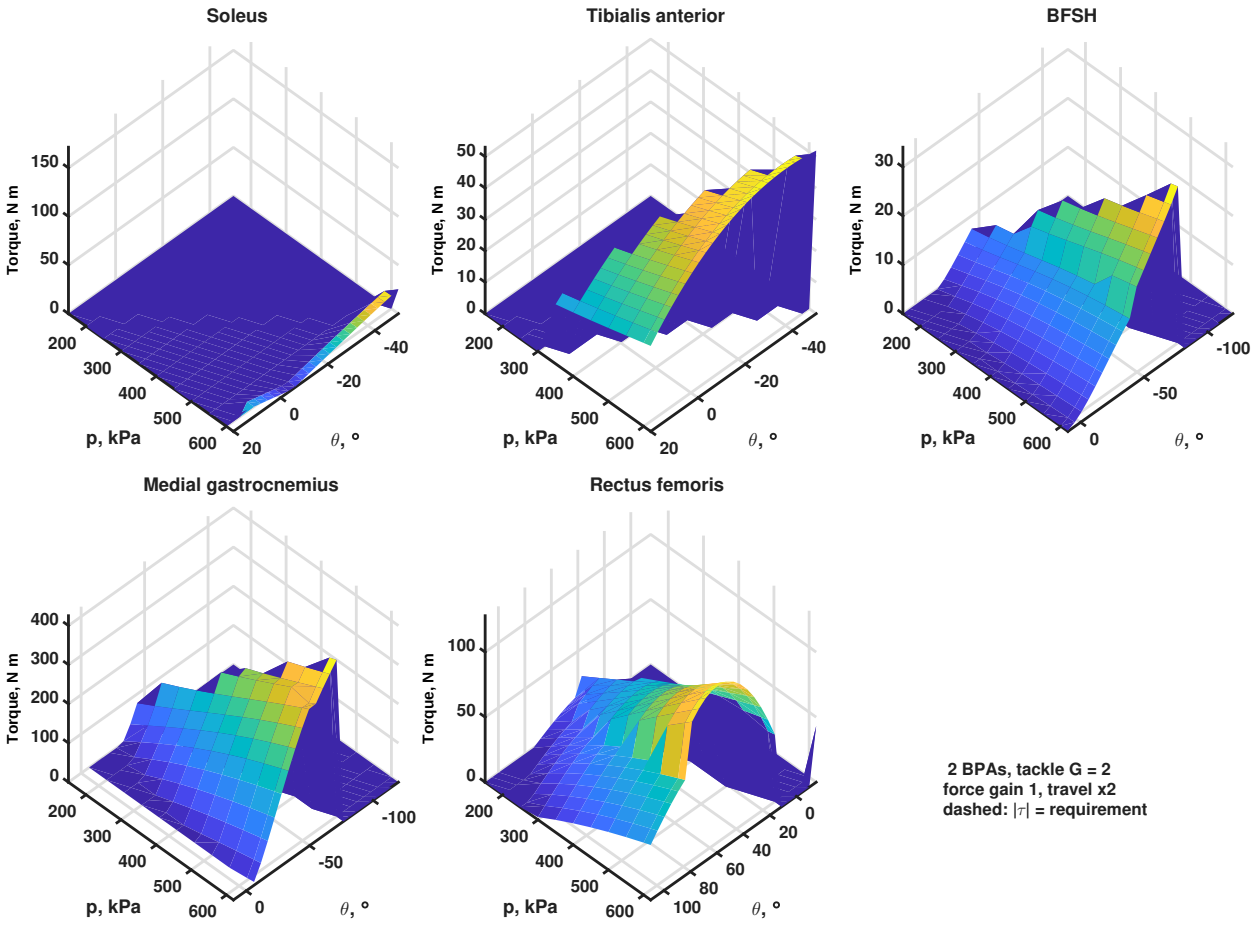


Figure C3: Predicted insertion-torque magnitude for two BPAs through a two-part tackle, with sizes re-solved. Dashed curves mark equality with the maximum-isometric human target.

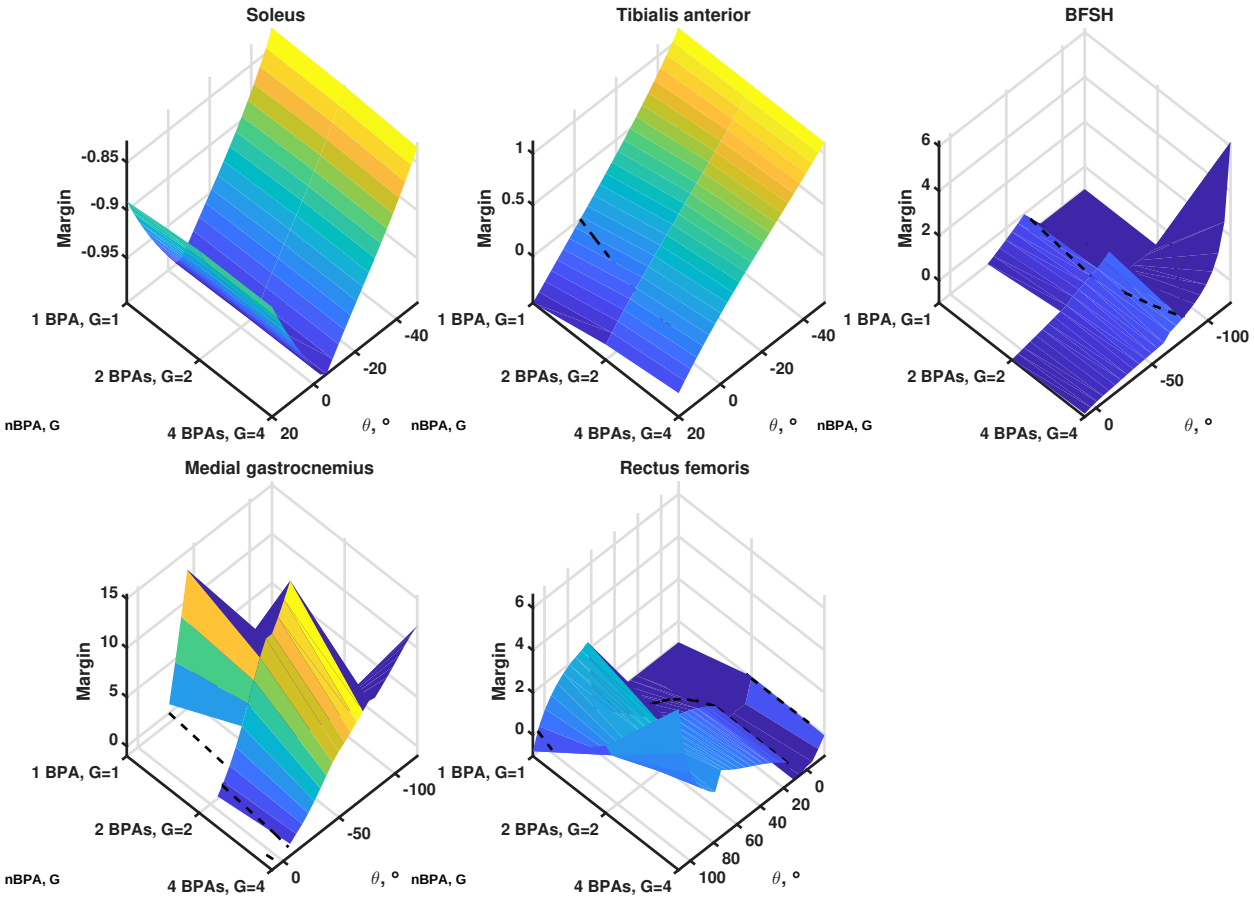


Figure C4: Torque margins at 620 kPa for the single-BPA baseline, two BPAs with a two-part tackle, and four BPAs with a four-part tackle. The categorical axis lists BPA count and travel multiplier G; G is not force gain. Ideal force gain relative to one BPA is n/G, equal to one for each displayed configuration. Sizes are re-solved. Positive values exceed the reference; panels follow Figure C3.

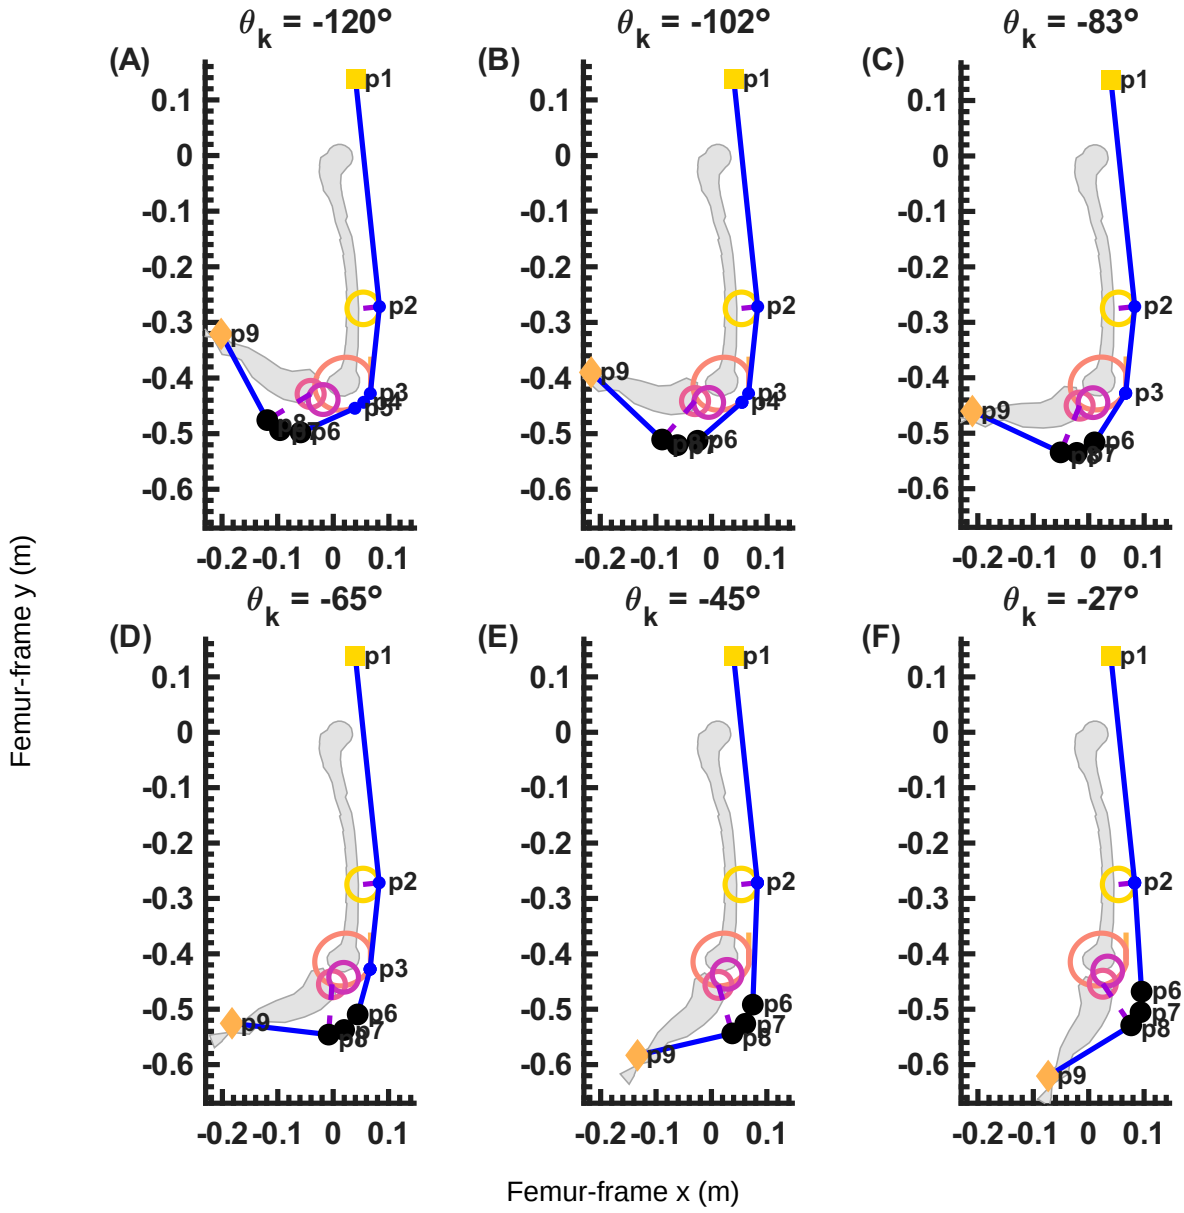
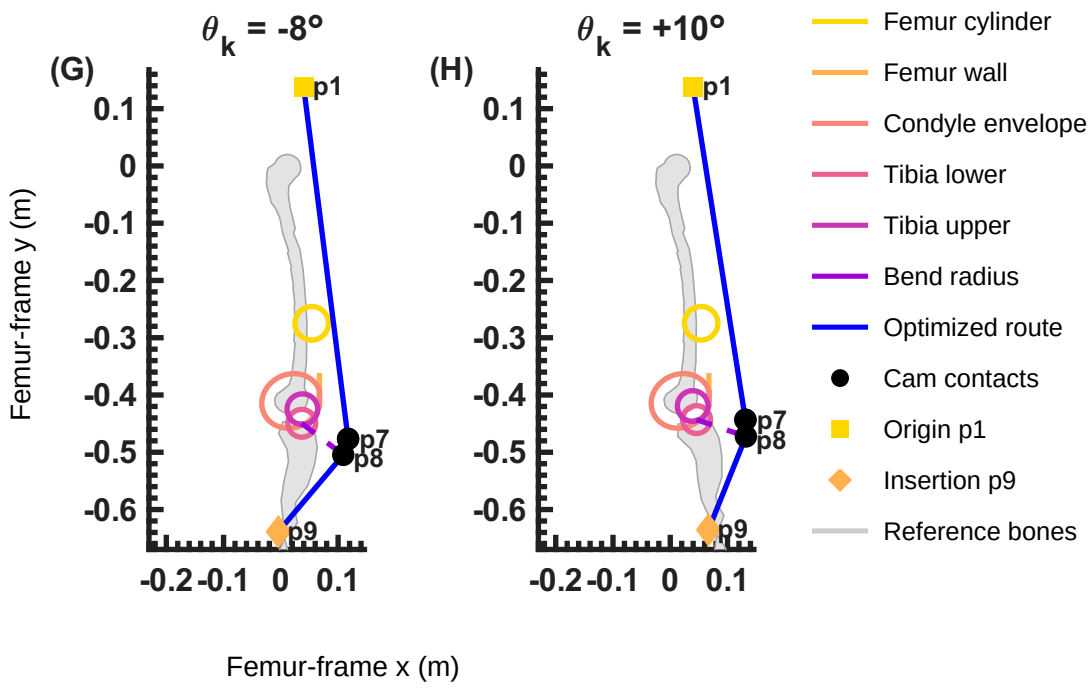


Figure C5: Cam-extensor poses A–F, spanning −120° to −27°. Blue lines show the route; colored outlines show the clearance envelopes. Gray silhouettes show the registered OpenSim anatomical reference; no foot is drawn. The remaining poses and legend are in Figure C6.



Cam-drum radius: 96.9 mm

Release angles: p5 -110.8 deg; p4 -89.8 deg; p3 -54.3 deg;
p2 -21.5 deg; p6 -18.9 deg. Gray bones are anatomical references.

Figure C6: Cam-extensor poses G and H at -8° and 10°, with the legend, 96.9 mm cam-drum radius, and contact-release schedule. Panel letters continue Figure C5.

C.9 Supporting Torque-Prediction Fit Statistics

C.9.1 Goodness-of-Fit Metrics

The MATLAB torque-fit comparisons use the shared routine `Functions/Go_OfF.m`: the root-mean-square error RMSE $= \sqrt{\frac{1}{n}\sum_i (y_i - \hat{y}_i)^2}$, the adjusted fraction of variance unexplained FVU $= \frac{n-1}{n}\ \sum_i (y_i - \hat{y}_i)^2 \ /\ \sum_i (y_i - \bar{y})^2$ (a sample-size-scaled residual ratio, distinct from the conventional $1 - R^2 = \mathrm{SSE}/\mathrm{SST}$), and the maximum absolute residual $\max_i |y_i - \hat{y}_i|$ (`Go_OfF.m` lines 45–52). Lower is better for all three, which makes the triplet directly usable inside optimization objectives.

Table C13 reports the biomimetic-knee comparisons at the adopted shared stiffness pair, with the configuration-specific length and wrapping terms defined in Chapter 4. The complete pinned-knee results are given in Tables C11 and C12.

Table C13: Biomimetic-knee goodness of fit before and after design-parameter optimization. RMSE and maximum absolute residual (Max.) are in N m; FVU is dimensionless.

	Improved BPA			**Optimized model**		
Configuration	**RMSE**	**FVU**	**Max.**	**RMSE**	**FVU**	**Max.**
Extensor, ∅10 mm	3.0322	4.4856	5.5870	1.1702	0.6681	2.2088
Flexor, ∅10 mm	7.4513	5.5612	13.3163	2.8002	0.8835	4.1201
Flexor, ∅20 mm	14.8064	3.3672	23.7218	2.2070	0.0748	3.7320

C.10 Configuration-Survey Coordinates and Actuator Sizes

The two coordinate columns of Table C14 define the graded surface, rather than every joint crossed by a multiarticular route. Other baseline sites use one actuator of the listed

diameter.

C.10.1 Torque and Active-Force Coverage

Table C15 expands the means reported in Table 4.10. Torque coverage requires the reference torque direction and magnitude; active-force coverage requires positive insertion tension. Both exclude cells outside the graded domain, and slack or infeasible configurations do not count as active.

Table C14: Graded coordinates and nominal BPA sizes for the 27-site survey. At nominal 40 mm sites, the baseline uses four parallel 20 mm BPAs; the updated configurations retain the inherited diameter.

Site	Primary DoF	Secondary DoF	Dia. (mm)	Baseline (mm)
ercspn_r	lumbar ext.	lumbar bend.	40	4×20
intobl_r	lumbar ext.	lumbar bend.	20	1×20
extobl_r	lumbar ext.	lumbar bend.	20	1×20
bifemlh_r	knee	hip flex.	20	1×20
semimem_r	knee	hip flex.	40	4×20
sar_r	knee	hip flex.	10	1×10
tfl_r	knee	hip flex.	10	1×10
rect_fem_r	knee	hip flex.	20	1×20
grac_r	knee	hip flex.	10	1×10
bifemsh_r	knee	hip flex.	20	1×20
vas_int_r	knee	hip flex.	40	4×20
soleus_r	ankle	knee	40	4×20
med_gas_r	ankle	knee	40	4×20
lat_gas_r	ankle	knee	20	1×20
tib_post_r	subtalar	ankle	40	4×20
tib_ant_r	ankle	subtalar	20	1×20
per_brev_r	subtalar	ankle	10	1×10
per_long_r	subtalar	ankle	20	1×20
per_tert_r	subtalar	ankle	10	1×10
flex_dig_r	MTP	ankle	10	1×10
flex_hal_r	MTP	ankle	10	1×10
ext_dig_r	MTP	ankle	10	1×10
ext_hal_r	MTP	ankle	10	1×10
glut_max_r	hip flex.	hip add.	40	4×20
add_mag_r	hip add.	hip flex.	40	4×20
iliacus_r	hip flex.	hip add.	20	1×20
psoas_r	hip flex.	hip add.	20	1×20

Table C15: Torque and active-force coverage at each actuator site. Percentages use the included grid cells for that site.

	Baseline (%)		Transmission (%)	
Site	**Torque met**	**Active force**	**Torque met**	**Active force**
ercspn_r	100.00	100.00	100.00	100.00
intobl_r	8.00	96.00	72.00	96.00
extobl_r	0.00	100.00	76.00	92.00
bifemlh_r	64.00	96.00	64.00	100.00
semimem_r	24.00	100.00	32.00	80.00
sar_r	40.00	100.00	60.00	100.00
tfl_r	16.00	64.00	40.00	100.00
rect_fem_r	12.00	68.00	24.00	100.00
grac_r	20.00	100.00	36.00	100.00
bifemsh_r	80.00	80.00	100.00	100.00
vas_int_r	20.00	100.00	20.00	100.00
soleus_r	40.00	100.00	100.00	100.00
med_gas_r	0.00	100.00	20.00	100.00
lat_gas_r	0.00	100.00	20.00	100.00
tib_post_r	96.00	100.00	96.00	100.00
tib_ant_r	56.00	100.00	100.00	100.00
per_brev_r	8.00	24.00	88.00	100.00
per_long_r	4.00	100.00	80.00	100.00
per_tert_r	80.00	96.00	96.00	100.00
flex_dig_r	56.00	80.00	96.00	100.00
flex_hal_r	48.00	100.00	60.00	100.00
ext_dig_r	8.00	92.00	12.00	100.00
ext_hal_r	80.00	100.00	88.00	100.00
glut_max_r	46.67	100.00	53.33	100.00
add_mag_r	26.67	100.00	60.00	100.00
iliacus_r	33.33	100.00	80.00	100.00
psoas_r	40.00	100.00	86.67	100.00
Mean	37.28	92.44	65.19	98.81

Appendix D: Supplementary Neuromechanical Verification

This appendix supplies supporting model records and evaluation sets for Chapter 6. It separates implementation checks, gait comparisons, pathway-removal tests, and spiking conversions because they test different claims. Neuron equations and calibration tables are provided in Appendix B.

D.1 Model Construction and Recording

The archived AnimatLab projects provide reference circuits for implementation comparisons. Figure D1 shows their neural organization. A later reorganized bilateral project added commissural and biarticular pathways; it ran headless but latched under the tested drives. That assembly result is not evidence of stepping. The phase-1 recordings below concern an earlier instrumented configuration, not the reorganized network.

A 10 s instrumented run recorded joint, neural, sensory, and body-motion channels at a 0.2 ms interval. Contact-sensory recordings are membrane voltages, not contact forces. The initial external horizontal force and the limited window prevent the joint traces from establishing stable walking.

D.2 Controller Ports and Reference Checks

The matched W2L test compares implementations from a common circuit specification under identical drive and contact inputs, using burst period, bilateral phase, burst count,

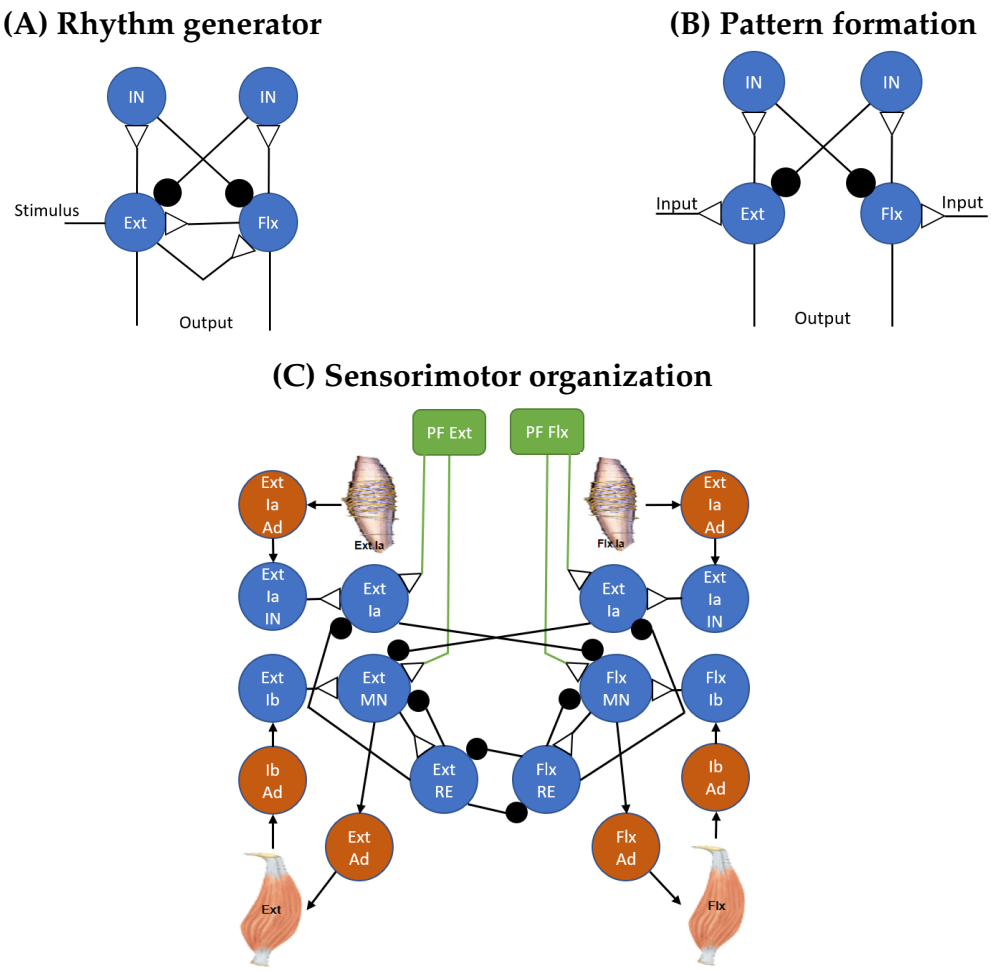


Figure D1: Archived AnimatLab motifs: rhythm generation (A), pattern formation (B), and sensorimotor organization (C). Ext and Flx denote antagonist pathways; IN denotes interneurons, MN motor neurons, and Ia and Ib muscle afferents. Triangles and circles distinguish excitatory and inhibitory connections.

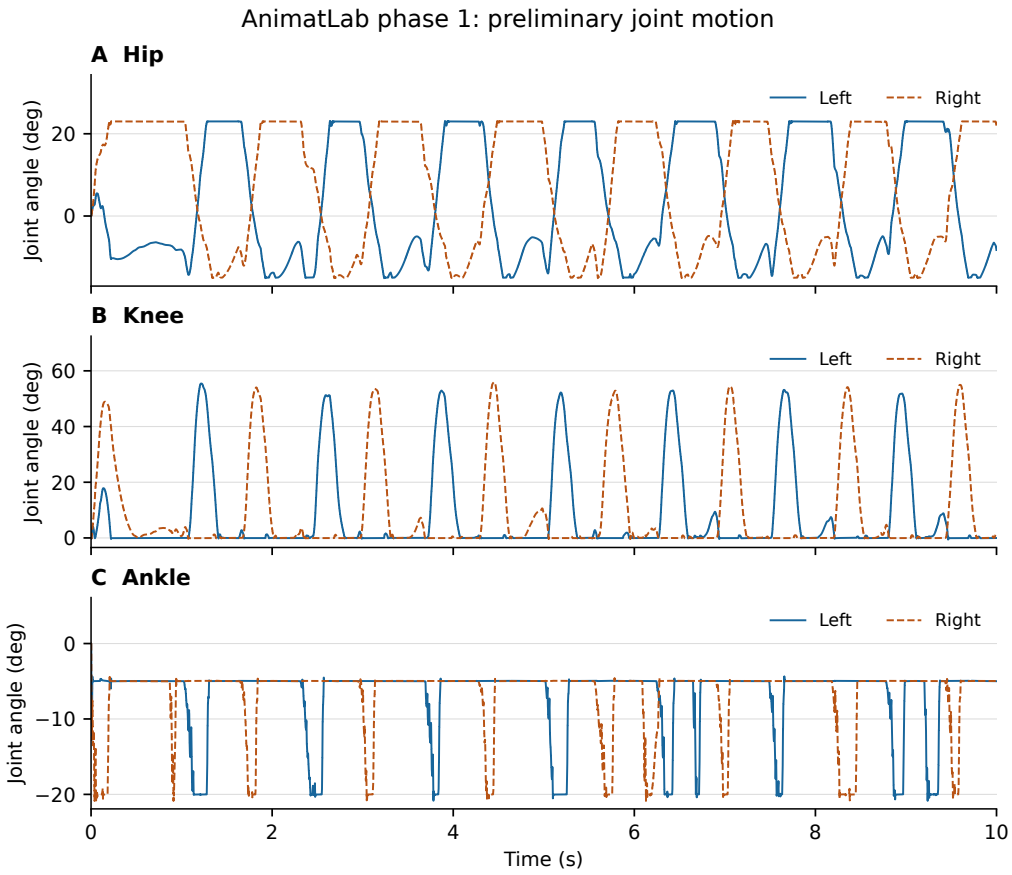
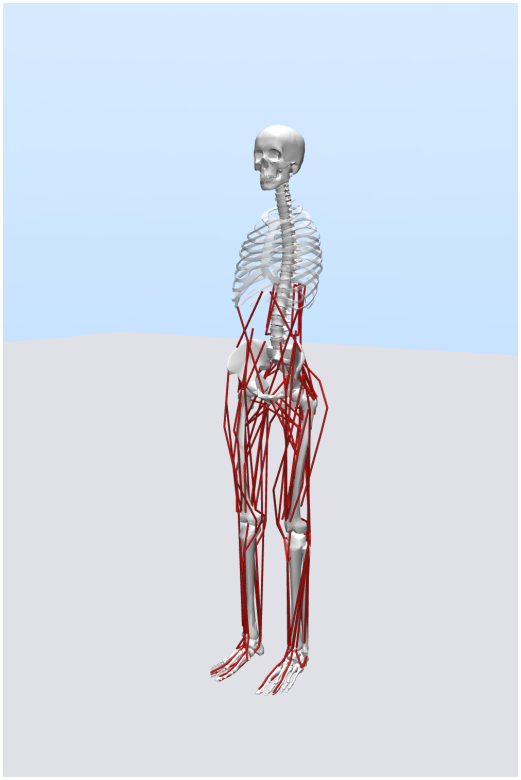


Figure D2: Bilateral hip (A), knee (B), and ankle (C) angles in the recorded joint conventions. All 50,001 samples in the 10 s window are retained. An external horizontal force acts during the first second. The recordings show motion, not validated walking or ground-contact gating.

(A) Converted robot



(B) SNS-Simscape knee model

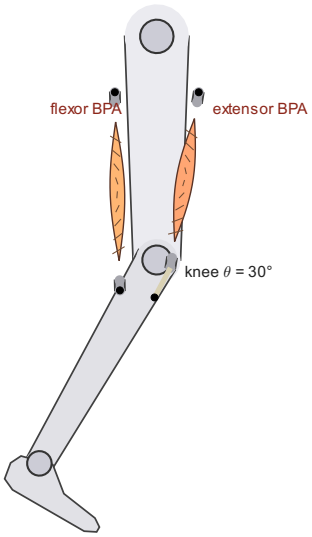


Figure D3: Converted Gait2392 robot (A) and the antagonist-BPA knee illustration used in the current SNS-Simscape demonstration (B). The two panels represent different models; scales differ.

and voltage correlation. The Deng controller port was checked separately: a current-pulse test produced a 1.938 s period against a 1.94 s NumPy reference, with antagonist correlation $r = -0.86$. Agreement on these tests does not establish long-term phase agreement in every network.

The archived original AnimatLab model produced a period of approximately 0.221 s, compared with 1.938 s in the translated test. Both alternated antagonists, but the original bilateral rhythm was in phase. Different parameterizations prevent quantitative equivalence. In both implementations an RG perturbation disturbed activity without sustained suppression; comparing perturbed and baseline sequences, rather than counting peaks in a short window, establishes that qualitative response.

The full exported network agreed with its NumPy reference to 4.2×10^{-6} mV over the early transient across 410 neurons. Its subsequent divergence and the tonic-drive fixed-point discrepancy reported in Chapter 6 limit parameter transfer. The OpenSim-derived Simscape mechanism contains 1,446 blocks and 149 joints; mechanism compilation is distinct from completing its neural and muscle coupling.

D.2.1 Transient Knee Sensory-Current Response

The knee block model was tested from a settled operating point using a rectangular current pulse added after the spindle transducer and before the extensor Ia sensory-neuron membrane input. Stimulus and no-stimulus cases start from the same initial conditions and run continuously for 50 s; the 2 nA pulse acts from 30.000 to 30.030 s. The original configuration did not settle: its knee angle spanned 0.712° during the 25–30 s baseline

window. Its former startup comparison is therefore not used to demonstrate transient reflex recovery.

The graded-response comparison changes exactly six parameters: the two MN output ramp widths from 1 to 10 mV, and the four Ia synaptic ramp widths from 0.5 to 10 mV. Equations, thresholds, conductances, wiring, mechanics, tonic drives, and activation dynamics are retained. This reduces loop gain as an explicit engineering comparison, rather than fitting physiological parameters. The simple actuator-force law is distinct from the characterized BPA model of Chapter 3. The flexor activation is zero at baseline, so the experiment does not establish suppression of a tonically active antagonist.

The modified configuration passes baseline settling tests for angle, velocity, all six membrane states, and both activations. Its paper-coordinate baseline is -42.6392°; positive excursions indicate extension. Extension peaks at 1.0908° after 111 ms. Recovery requires displacement within 0.01° of baseline and absolute angular velocity below $1 \times 10^{-4}\,\mathrm{rad\,s^{-1}}$; both remain within those bounds from 6.355 s after stimulus onset. The late residual is approximately $9.3 \times 10^{-7\circ}$. A fresh-process replay reproduces the resampled pulse arrays exactly. Halving the ODE45 maximum step from 1 to 0.5 ms, at unchanged tolerances, changes the angle trace by at most $5.21 \times 10^{-5\circ}$. Comparisons use linear interpolation onto a 1 ms grid, without smoothing. These checks establish numerical recovery in the model, not physiological reflex fidelity.

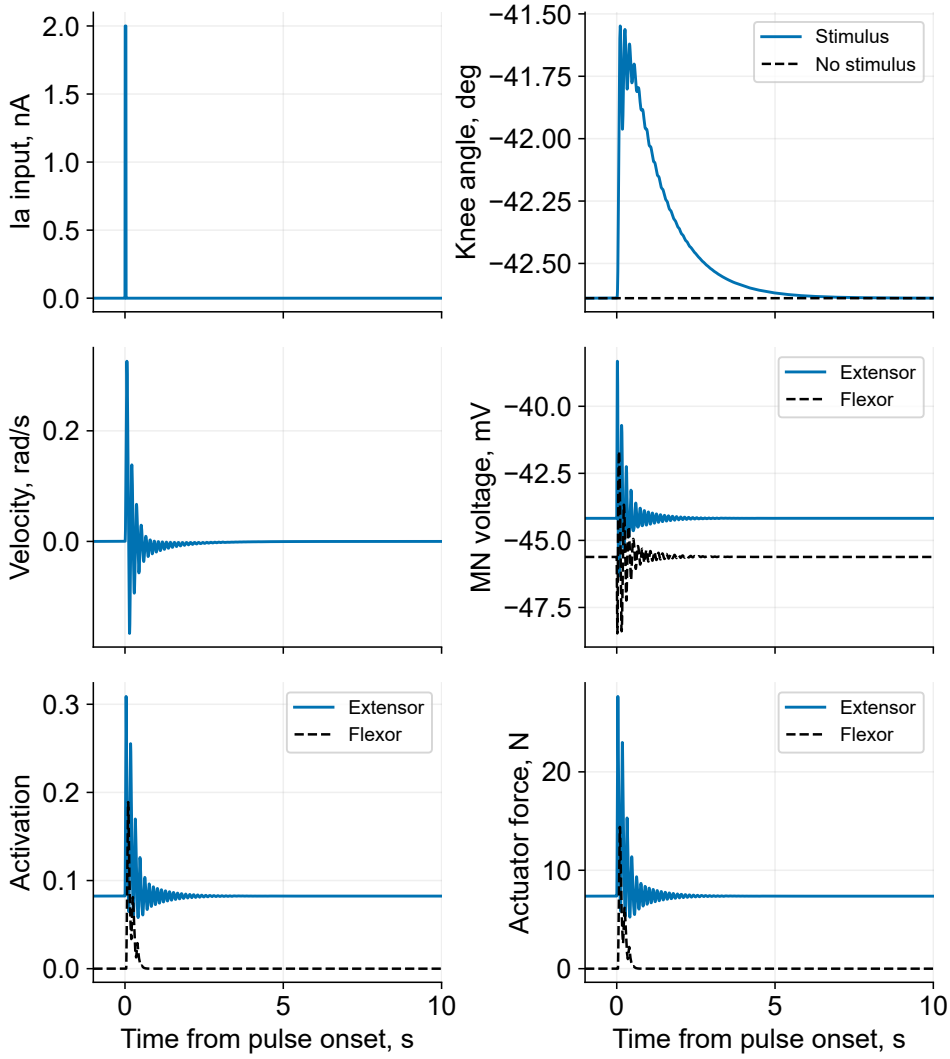


Figure D4: Stimulus, knee angle, angular velocity, MN voltages, activations, and actuator forces in the graded-response knee model. The pulse enters the extensor Ia sensory neuron after settling. Knee angle is compared with the no-stimulus control; extensor and flexor signals are distinguished in the other paired panels. The display applies only the sign conversion to flexion-negative coordinates.

D.2.2 A Published Neural Mechanism as a Verification Target

The reduced Python implementation of the two-level circuit of **?** passed 13 checks covering drive-dependent activity, repeatability, and the distinction between resetting and non-resetting deletions. Stimulating the pattern-formation extensor population suppressed flexor motor output while the rhythm generator continued, so flexor activity resumed at its expected phase. Stimulating the rhythm-generator extensor population also suppressed flexor output but shifted the subsequent rhythm phase. This comparison tests the functional separation between rhythm generation and pattern formation: similar missing motor bursts can arise through different neural mechanisms. It is a neural-only verification of the translated circuit, independent of a mechanical model and separate from the Simulink demonstrations.

D.2.3 RG, PF, and MN Stimulus Tests

The fourteen-neuron Deng-derived implementation follows the circuit used in SNS-Toolbox demonstrations [**??**]. A 0.1 ms Euler step reproduced the existing neural integrator's complete 24-second baseline voltage trace exactly in the tested process. The baseline period was 1.938 s. Neural output drove the existing suspended, unilateral W2L mechanical model without afferent return. Hip commands came from the MN outputs; knee and ankle commands used the existing PF-to-activation mapping. This configuration tests transmitted circuit effects, not the original rat mechanics or bilateral walking.

The initial sweep applied $+2$ and $+10$ nA to each layer for 0.1 and 0.5 seconds, together with two direct MN inhibitory controls. None met the deletion criterion. Subsequent

stronger-current control series, including stimuli lasting two baseline periods (3.876 s), retained the same circuit and mechanical parameters (Table D1). RG and PF stimuli began one-quarter cycle after a baseline RG onset; MN inhibition began 0.05 seconds into a baseline motor burst. Positive currents target the extensor RG or hip-PF population; negative currents target the hip-flexor MN. These currents are mechanistic controls, not experimentally calibrated physiological stimuli.

A deletion required at least 20 ms of baseline motor activity in the stimulus window and suppression of at least half of that activity, using an activation threshold of 0.5. Phase residuals compare the final four recovered onsets with the nearest baseline onset, normalized by the baseline period. Residuals at or below 0.05 cycles were classified as non-resetting, and those at or above 0.15 as shifted; intermediate values remained indeterminate. RG and motor residuals were assessed separately. The two-cycle stimuli suppressed both expected flexor burst onsets at all three layers. Baseline and stimulated neural and mechanical trajectories were identical before stimulus onset. The one-second PF pulse shifted the motor pattern without resetting the RG clock, whereas the two-cycle PF pulse preserved both schedules. The duration dependence cautions against treating every motor phase shift as an RG reset.

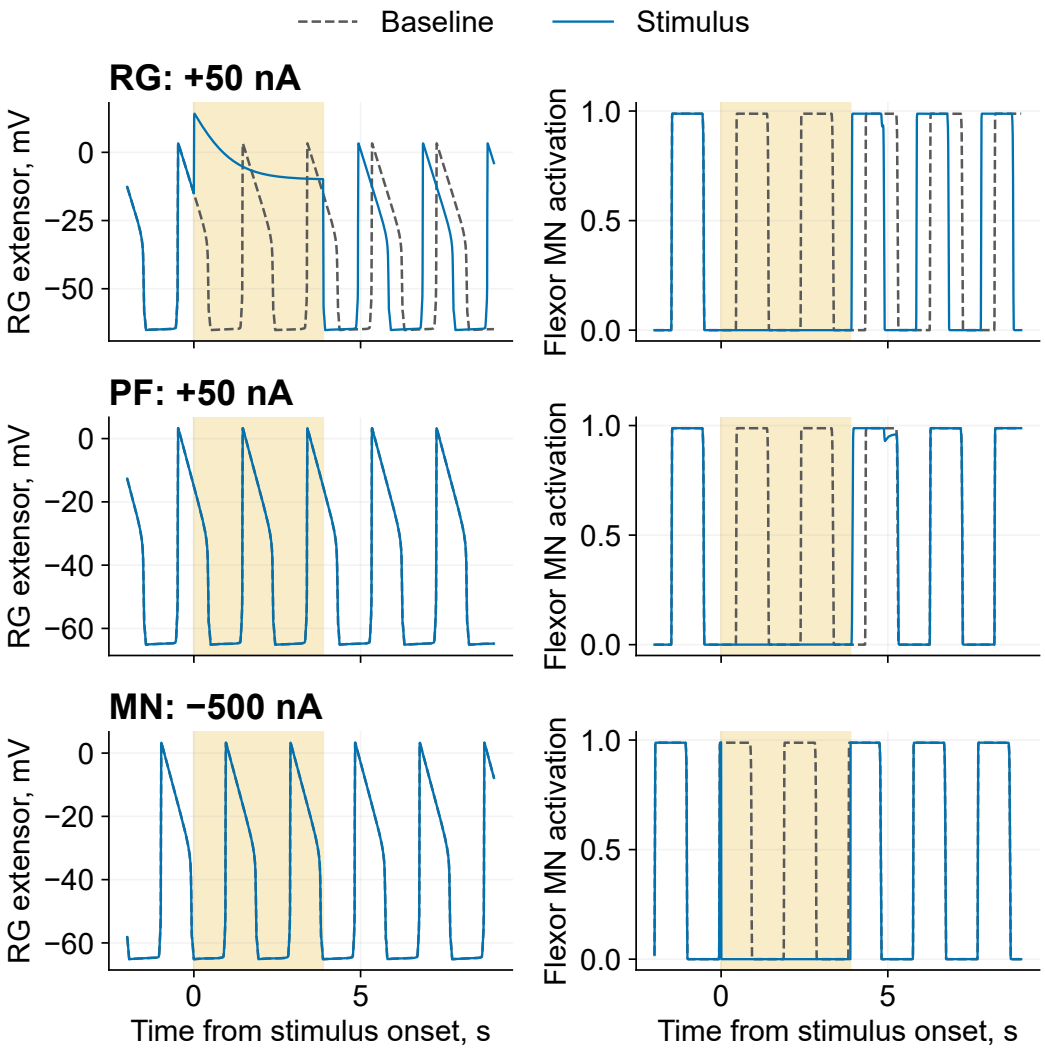


Figure D5: Matched baseline and stimulus responses in the reduced Deng-derived circuit coupled to a suspended model. All stimuli last two baseline periods. Rows show RG stimulation, PF stimulation, and direct flexor-MN inhibition; columns show RG voltage and flexor-MN activation. Shading marks the applied stimulus. The RG intervention shifts subsequent timing, whereas the displayed PF and MN interventions preserve the RG schedule.

Table D1: Stronger-current controls in the reduced Deng-derived circuit. Phase residuals are fractions of the baseline cycle. "No deletion" means the specified suppression criterion was not met.

Target	nA	s	RG phase	MN phase	Outcome
RG	+50	0.5	0.052	0.052	Brief suppression; phase indeterminate
RG	+50	1.0	0.307	0.307	Suppression; RG and MN shifted
PF	+50	0.5	0.000	0.000	Brief suppression; schedule preserved
PF	+50	1.0	0.000	0.208	Motor pattern shifted; RG preserved
MN	−100	0.5	0.000	0.000	No deletion
MN	−500	0.5	0.000	0.000	Brief suppression; schedule preserved
RG	+50	3.876	0.211	0.211	Resetting deletion
PF	+50	3.876	0.000	0.000	Non-resetting deletion
MN	−500	3.876	0.000	0.000	Non-resetting deletion

D.2.4 Archived AnimatLab Models

Headless reruns tested whether archived models reproduced their reference behavior in the original simulator. Eleven targetable rhythm and gait metrics passed their acceptance bands. The bilateral contact-driven model's rhythm generators had periods of 0.450 and 0.467 s; the modern W2L model showed joint oscillations at 2.25 Hz. The Li-model rerun reached 0.637 m s^{-1} with step periods near 1.3 s [**?**]. These are reference-model checks, not results for the combined Deng–Li model or evidence of unsupported walking.

A separate MuJoCo translation of the Li controller and body did not reproduce that original-simulator walking result. In a 20-second ground test, mean forward speed was 0.040 m s^{-1}, with eleven detected right-foot stance episodes and none on the left. The right-side period of 1.294 s alone therefore did not establish bilateral stepping. This failure distinguishes reproduction of the source model from successful transfer to another simulator.

D.3 Tuning and Supported Evaluations

Free network gains were searched with a tree-structured Parzen estimator, with each candidate evaluated through simulation. The staged procedure separates suspended rhythm generation, sensory feedback, standing, and ground evaluations (Table 6.2). Finalists were checked under 1 % multiplicative parameter perturbations and changes in drive. Such checks probe the neighborhood of a selected configuration; they do not establish robustness to untested loads or removal of the harness.

The cycle-shape metric assigns zero to perfect agreement. A run without detectable cycles returns -25 as a reserved failure value. Because this value is numerically greater than the genuine-gait scores below, failed runs must be identified before ranking scores. That value itself is not a valid measure of gait quality.

The curriculum is shown in Table 6.2 in Chapter 6.

D.3.1 Reference-Library Comparison

The 43-cycle library contains 32 running trials from eight subjects at four speeds, ten walking references, and a predicted-walking cycle. The walking references span prescribed speeds of $0.5\,\mathrm{m\,s^{-1}}$ to $2.0\,\mathrm{m\,s^{-1}}$ and self-selected trials. The production network (`s3k`), reduced-connection network (`s3kpruned`), six-group activation network (`syn6`), and W2L-organization network (`w2lvar`) were scored without treating a good tuning-reference score as evidence of generalization. Library scores were substantially more negative (Table D2), and the evaluations showed unilateral or asymmetric gait. This comparison bounds transfer to the available references; it does not validate normal human walking.

Table D2 uses the production pipeline's cycle-shape kinematic score. The tuning-reference column reports each configuration on its tuning reference; the library mean averages that score over all 43 reference cycles.

Table D2: Four controller configurations evaluated on their tuning reference and the 43-cycle reference library. For genuine gait cycles, larger kinematic scores indicate closer agreement; zero is perfect agreement.

	Tuning-ref. score	Library mean score (43 refs)	Most negative reference (score)
`s3k`	−186.4	−8776	`Case_40_motion` (−8841)
`s3kpruned`	−176.1	−6924	`Case_40_motion` (−6986)
`syn6`	−200.4	−9961	`Case_40_motion` (−10040)
`w2lvar`	−144.5	−5808	`subject08_Run_20002` (−5896)

D.3.2 Direct Kinematics Overlay

The saved human-model controller reproduced its cycle-shape score of −145.07 in an independent process. The direct overlay in Figure 6.7 uses the Gait2392 walking inverse-kinematics trajectory and its paired ground-reaction-force file. Model cycles begin when vertical foot loading crosses 50 N; reference stance is defined using the same force threshold. Ten cycles per side, with durations between 0.35 and 2.5 seconds, were retained using the existing cycle detector. Each leg was aligned to its own onset and resampled over the gait cycle, without fitting phase, subtracting the mean angle, or scaling amplitude.

Mean cycle periods were 1.285 s on the right and 1.254 s on the left, compared with 1.233 s and 1.213 s in the reference. Model stance fractions were 0.710 and 0.649, compared with 0.610 in the reference. Table D3 reports absolute-angle errors between mean cycles.

These errors retain posture offsets and are distinct from the mean-removed cycle-shape errors used in tuning. The hip and knee mismatches show why cycle counts and a scalar score must be accompanied by the trajectories and contact timing.

Table D3: Unshifted absolute-angle RMSE between mean model and reference cycles, in degrees.

Joint	Right	Left
Hip	28.03	24.65
Knee	27.44	28.45
Ankle	8.71	7.38

D.3.3 Pathway-Removal Tests

Individual pathway classes were removed from each selected configuration and the model re-evaluated on its tuning reference (Table D4). Removing Ia feedback from the production configuration eliminated detected cycles. Other removals changed the score without necessarily eliminating movement. The different response of the six-synergy variant shows that these dependencies belong to the tested network configurations, rather than constituting a general claim about biological necessity. The systematic pattern-formation and motoneuron stimulation series for these full Gait2392 networks was not completed. The reduced-circuit tests in Section D.2.3 address a separate, smaller model.

Table D4: Pathway-removal tests on the tuning reference. Full-model rows give unchanged baselines. A score of −25 is the fixed no-cycle fallback and denotes failure, despite exceeding the genuine-gait scores.

Configuration	**Removed pathway**	**Score**
`s3k` (full model)	none	−186.4
	Ia afferents	−25 (no cycles detected)
	all afferent classes	−267.8
	Renshaw cells	−267.3
	group-II afferents	−255.2
	Ib afferents	−180.3
	contact channels	−183.9
	interleg commissural pathways	−174.9
`syn6` (full model)	none	−200.4
	all afferent classes	−190.5
`w2lvar` (full model)	none	−144.5
	Ia afferents	−273.5
	weak rhythm-generator mutual excitation	−279.9
	Renshaw cells	−142.4

D.4 Spiking-Neuron Comparisons

The spiking variants replace afferent, rhythm-generator, and interneuron populations while retaining graded motor neurons. Synapses onto motor neurons are calibrated by mean conductance, preserving the continuous activation interface. This construction does not require the network dynamics or resulting movement to remain unchanged. Calibration and full conversion tables are in Appendix B.

D.4.1 MuJoCo Network and Ground Evaluation

Under constant drive, the non-spiking and spiking rhythm periods were 0.399 and 0.621 s, respectively (Figure D6). The spiking period is 55.6 % longer. Timing agreement therefore requires more than copying the selected parameters. The untuned spiking variant produced bilateral cycling in a ground evaluation where the non-spiking winner cycled predominantly on one side (Table D5; Figure D7). This individual comparison motivates further controlled testing; it does not establish superiority of spiking control.

D.4.2 Simulink Demonstrations

The graded and spiking knee-reflex demonstrations were compared over matched 5 s runs (Table D6; Figure D8). The beverage example retained approximately the same sag, 2.02° versus 2.00°, compared with 7.85° without the reflex. These examples demonstrate transfer of particular reflex responses after calibration, not invariance of all controlled behavior to neural coding.

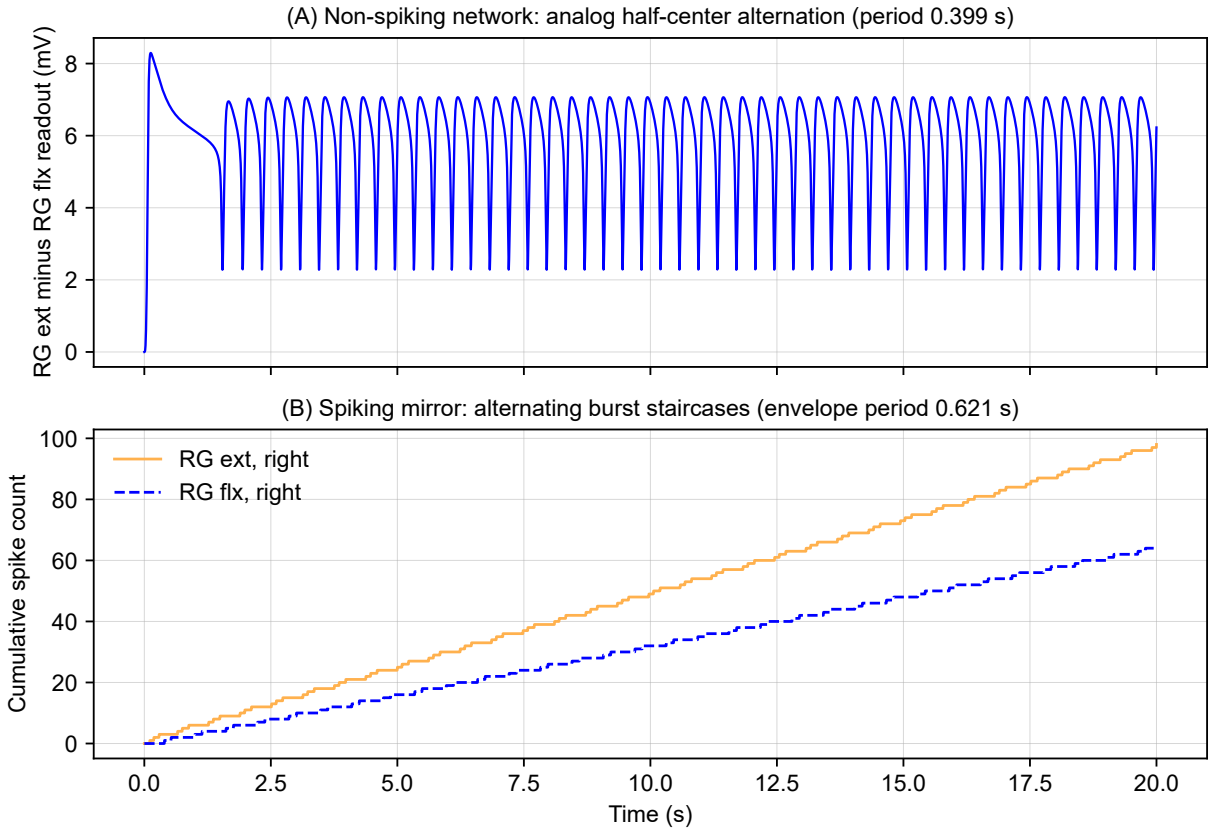


Figure D6: Constant-drive rhythm in the non-spiking network (A) and spiking variant (B). The observed periods are 0.399 and 0.621 s; the latter is 55.6 % longer. The panels use different neural readouts: graded population difference versus individual spiking cells.

Table D5: Supported ground evaluation at the same selected parameters (16 s), with no spiking re-tuning. Knee minima and contact duty are computed over detected cycles. The kinematic score is a cycle-shape error metric: zero denotes perfect agreement, and less negative scores indicate closer agreement.

Metric	**Non-spiking winner**	**Spiking mirror (untuned)**
Kinematic score	−161.57	−236.35
Gait cycles, right / left	9 / 0 (left frozen)	8 / 9 (bilateral)
Cycle knee minimum	−64.3°	−79.8°
Contact duty, right / left	0.637 / 0.985	0.799 / 0.856
Maximum pelvis tilt	22.1°	19.2°
Non-finite states / fell	no / up	no / up

Table D6: Knee-reflex demonstration, graded baseline versus spiking twin at the matched operating point (5 s runs; settle statistics over the last 2 s).

Metric	**Baseline**	**Spiking twin**
Knee angle at 0.05 s	15.79°	15.79°
Settle mean (last 2 s)	43.5°	44.4°
Settle range	41.6°–44.9°	41.9°–46.7°
Extensor / flexor activation (window means)	0.53 / 0.53	0.58 / 0.56
Antagonist alternation	0.5 Hz–2 Hz	7.4 Hz (dominant peak)
Ia interneuron firing at settle	graded, subthreshold	95 Hz; Ib 0 Hz

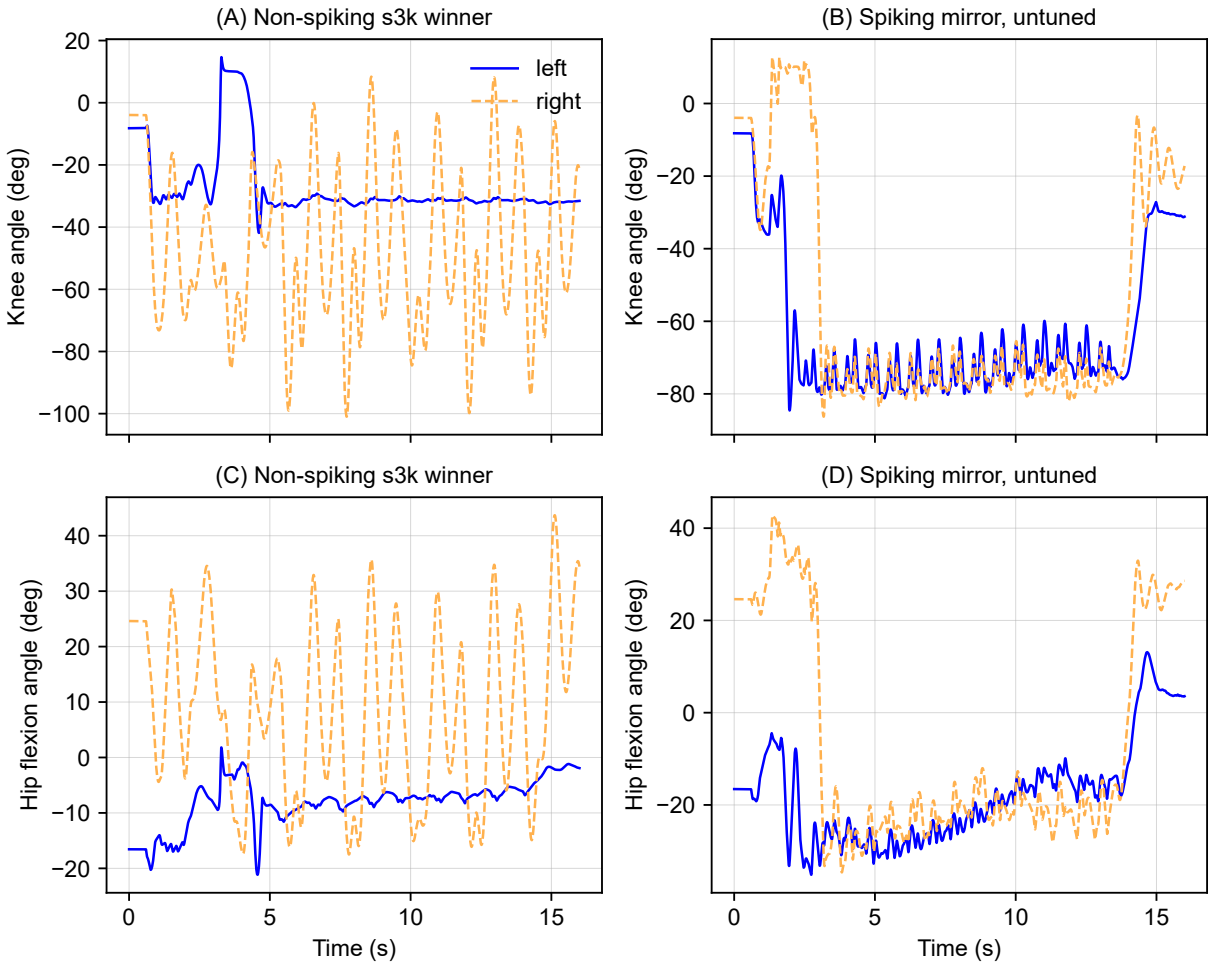


Figure D7: Ground-evaluation joint angles at the non-spiking winner parameters. The non-spiking model cycles predominantly on the right (A, C), while the untuned spiking variant cycles both legs (B, D). Solid indigo denotes left, and dashed orange right. Both use the supported evaluation protocol.

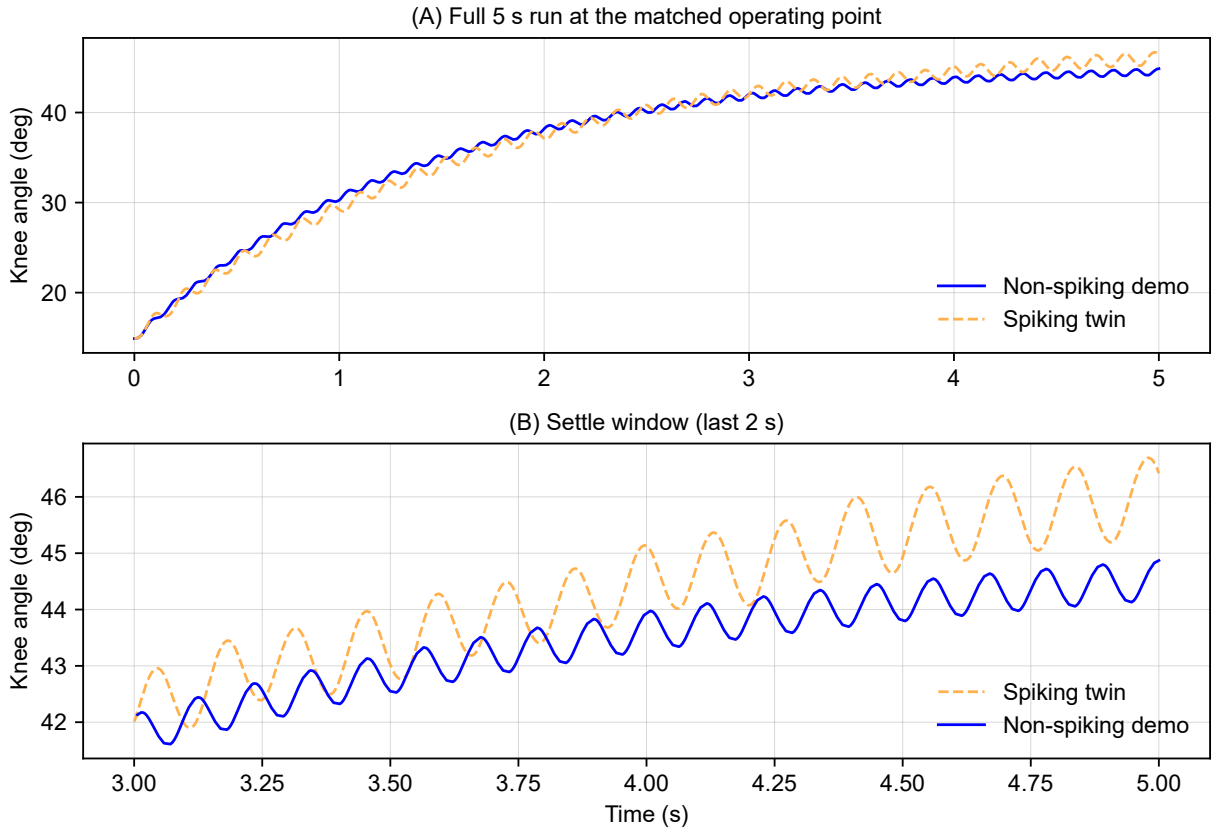


Figure D8: Knee-reflex demonstration at the matched operating point: graded baseline (solid indigo) and spiking variant (dashed orange). Full 5 s run (A) and final 2 s window (B). The spiking response has a slightly wider oscillation band.

D.4.3 AnimatLab Conversions

The tested spiking conversions did not preserve the graded models' rhythm (Table D7). Converting thresholds while retaining graded synapses preserved near-baseline oscillations, whereas converting synaptic amplitudes and decay changed or eliminated alternation. These results identify the synapse conversion as a limitation of this mapping. They do not demonstrate that a suitably calibrated spiking implementation cannot reproduce the reference behavior.

Table D7: AnimatLab spiking conversions versus their graded baselines (5.1 s headless runs; joint ranges in degrees where the rhythm survives). The thresholds-only ablation converts the neurons but keeps the graded synapses.

Variant	Baseline (graded)	Spiking copy
Two-layer walker	0.443 s rhythm; hip -15 to $+23$, knee -1 to $+60$	dead (2 spikes; hip pinned near $+6$)
Bilateral variant	0.462 s; antiphase -0.42 cycle latched (tonic drives)	latched; one side at 442 Hz–450 Hz
Reorganized project (standalone export)		synchronous 11.8 Hz co-bursting; same-leg envelopes $r = 1.000$ at lag 0
Thresholds only (graded synapses)	hip -15 to $+23$, knee -1 to $+60$	oscillates: hip -15.1 to $+23.2$, knee -0.8 to $+47.9$

D.5 Supported BPA-Robot Evaluations

The BPA-actuated walker exercised the chain from neural activation through the actuator force model to multi-joint motion. In the supported air-stepping run, the virtual harness carried 70 % of the 411 N body weight. The period was 1.542 s, knee RoM was 63.5°, and peak forces were 880 and 1372 N, at a peak pressure of 372 kPa.

With split left/right drives and retuned feedback, the W2L robot progressed 1.72 m

over 30 s under the support protocol. It retained progression at 60 % drive and passed 12 of 12 perturbed reruns, but did not satisfy the strict contact-duty walking criterion; below that drive the gait ceased. The mechanical support was not weaned.

In the separate human-model drive-weaning study, the selected score improved from -161.57 to -145.07, with ten contact cycles per leg and contact duties of 0.71 and 0.65. It passed eight of eight perturbed reruns and retained bilateral stepping down to zero hand drive. The motion remained in place under support. These two studies distinguish reduced neural assistance from reduced mechanical assistance.

D.5.1 Independent Replays and Remaining Mechanical Assistance

A separate W2L configuration was evaluated after correcting the delivery of saved per-target gains to the evaluator. Four independent 30 s replays produced identical trajectories. Over the 4.5 s to 30 s analysis window, the run had 28 left and 21 right contact events, a median bilateral phase of 0.496 cycles, forward displacement of 2.246 m, and fitted speed of $0.0918\,\mathrm{m\,s^{-1}}$. All twelve independent reruns with 1 % parameter perturbations passed the supported-stepping criterion. These are results for this configuration, distinct from the drive-weaning study above.

Removing nominal constant weight support did not remove the position-control harness: its mean vertical force still supplied approximately 16.76 % of body weight. Contact duty remained below the strict walking criterion. Thus, repeatable supported stepping was demonstrated, while harness-free walking was not. An accelerated neural backend preserved the complete reference trajectories and reduced the isolated neural-step cost by

a factor of 1.425; this does not establish the same speed improvement for the complete neuromechanical simulation.

D.5.2 Pathway Deletions and Stimulus Tests

The recovered W2L configuration was evaluated with nominal weight support set to 0.25 and the remaining controller settings fixed. Mild and strong drive reductions, forward and backward 40 N pushes, and separate Ia, Ib, Renshaw, crossed-heel, and heel-cascade deletions retained the supported-stepping criterion. Their effects differed: Renshaw deletion reduced forward displacement to 0.081 m, while commissural deletion failed bilateral stepping. Disabling all sensory feedback retained joint cycling and 1.996 m progression but shifted median contact phase to 0.223 cycles, outside the accepted alternation range. The intact comparison produced 2.239 m and a phase of 0.498 cycles. These experiments distinguish rhythm generation, progression, and bilateral contact coordination, rather than treating them as interchangeable outcomes.

D.5.3 Di Russo Controller Replication Attempt

The original Di Russo controller uses a reduced planar Delp/Thelen model with nine active coordinates and 18 Thelen muscles, rather than the full Gait2392 model. The source mechanical model loaded and its OpenSim system initialized successfully; this checks model construction, rather than gait. Its optimized parameter file contains 256 parameters, including individual reflex connections, muscle delays, and balance terms. Recovering that source clarified differences from the exploratory adaptation used here. Corrections

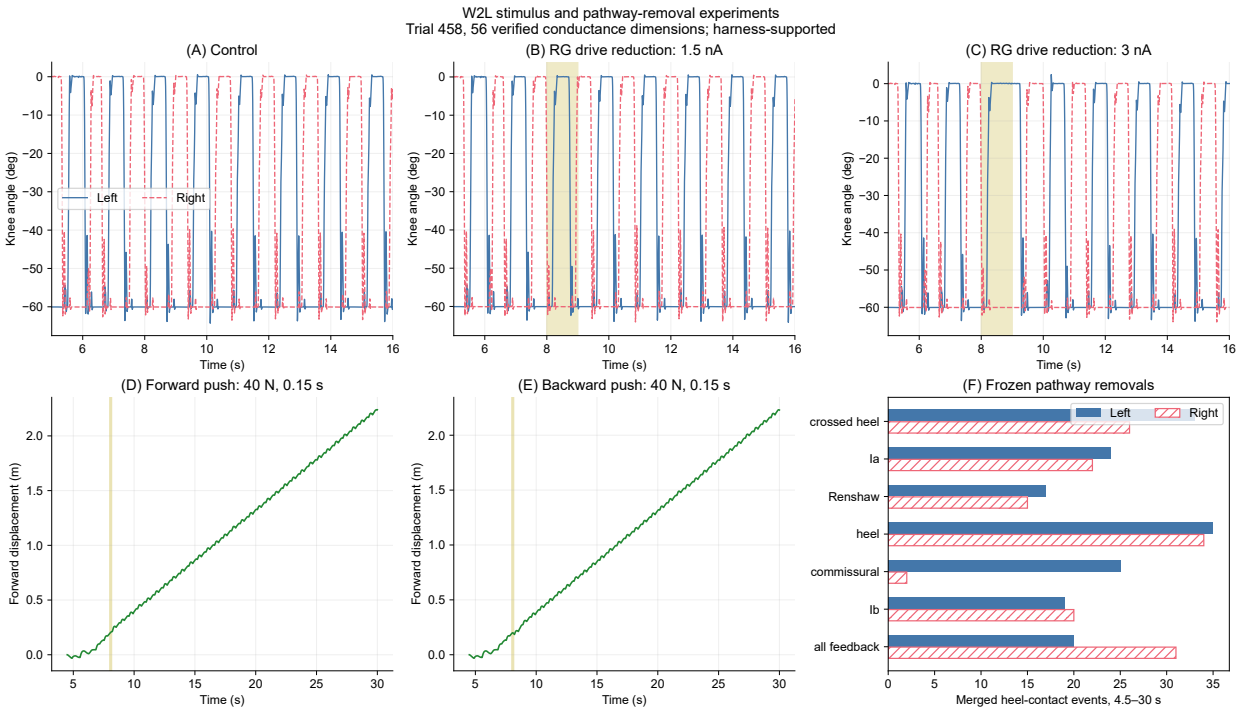


Figure D9: Recovered W2L controller responses to drive changes, pushes, and pathway deletions. The same effective controller gains are used across conditions, with the specified intervention applied in each panel.

to the adaptation addressed phase units, synchronous updates, output saturation, sensor ordering, contact-force direction, and initialization. Its compiled neural equations agreed with the corrected Python equations within 3.34×10^{-16} over 2000 steps, establishing agreement between those two implementations.

A 960-trial exploratory duration-and-speed search nevertheless selected a tumbling trajectory, with pelvis rotation reaching approximately 223°. That candidate was rejected as walking, and further optimization of the incomplete adaptation was stopped. The original controller's custom runtime was unavailable for a native baseline replay. The useful outcome is therefore the recovered source and the identified implementation requirements; replication of the published walking behavior remains incomplete. Source inventories, corrected-equation checks, and the rejected trajectory are retained in `Code/MuJoCo_SNS/spinal/reports_20261027_codex/`.

D.5.4 CAD-Derived Robot Integration

A separate CAD-derived robot model retained the four-bar knees and spherical hips, with 27 selected actuator routes per side. Its anatomical frame is x anterior, y superior, and z right; explicit frame transformations distinguish it from older converted models. The controller drives simulated reference tensions rather than calibrated BPA forces, so this model does not yet validate the complete actuator-to-controller integration proposed in Chapter 7.

The model completed a 20 s supported crouched-stance test. During the final five seconds, knee excursions were 1.35° right and 0.95° left, while the vertical rig carried

94.68 % of body weight and angular stabilization remained active. Disabling muscle forces under the same support settings caused lumbar failure after 0.202 s. This comparison establishes a contribution of the muscle controller under strong assistance, not independent balance. A separate run with rhythmic drive completed its horizon, but alternating gait and progression after settling have not been verified.

A subsequent matched 40-second comparison retained the same geometry, initial pose, reference-tension capacities, and external assistance in both conditions. During 15–40 s, vertical assistance supplied 80% of body weight. The additional control combines muscle-mediated separation feedback (160 N/m, capped at 20 N) with a 6% activation floor on currently abducting routes, compared with the 5% global floor (Table D8). Full recorded trajectories reproduce bit for bit between fresh reference and accelerated runs. This is numerical replay within the specified model, rather than a robustness test.

Table D8: Matched CAD-model comparison over 15–40 s. Vertical assistance remains 80% of body weight in both conditions.

Measure	Control off	Control on
Reversed ankle ordering (%)	67	0
Minimum ankle separation (mm)	−38.3	142.4
Median ankle separation (mm)	−6.2	159.0
Mean right TFL activation (%)	22.3	81.8
Mean left TFL activation (%)	99.6	99.8
Maximum sampled knee-loop error (mm)	4.485	0.387
Maximum sampled interleg penetration (mm)	2.700	0

The 262-neuron controller receives 230 inputs and drives 54 reference-tension actuators. Its MN target-current compensation cancels incoming PF and reflex synaptic currents to track the force allocator's targets. Consequently, the effective comparison concerns RG

timing, engineered force allocation, and mechanical response, rather than normal PF/reflex authority over activation. Toe-out and small post-ramp knee excursions remain, and the stepping acceptance tests fail.

The foot attachment is an unverified assembly datum. A separate hypothesis translates each whole foot 60 mm toward the proximal shank, retaining tibial geometry and mass. With otherwise unchanged hybrid control, median ankle separation changes from 159.0 to 57.3 mm, and sampled interleg penetration reaches 1.316 mm. A separate W2L-organization timing probe has 544 neurons and 230 inputs; four named ankle indices alias knee-PF cells and are not additional neurons. Its shifted-foot condition passes bilateral unloading and lifting counts but fails repeated alternating foot apices and forward progression. The inherited MN compensation remains active, so this is not a full PF-driven controller comparison.

Route sites can lie outside CAD surfaces and transmit force through their named parent bodies, equivalent to ideal rigid offset brackets or guides. Bracket mass, compliance, packaging, and collision-aware wrapping remain unmodeled. These attachment assumptions are separate from verification of the source moving-point functions and complete path derivatives. The older 20-second stance test and newer 80%-assistance comparison are different configurations, not a matched support-reduction experiment.

The W2L replay metrics and support accounting are archived in `Code/MuJoCo_SNS/spinal/reports_20261027_codex/verification.json` and the accompanying support-ladder records. CAD-model stance metrics are in `Code/MuJoCo_SNS/robot_torso/controller_assessment.json`. These records retain the configurations and evaluation conditions needed to distinguish the two models.

D.6 OpenSim-to-MuJoCo Conversion Repairs

The conversion work in Section 6.4.1 involved two distinct problems: importing the robot's artificial-muscle routing with human muscle parameters, and repairing the converted human reference used for controller development. The first combination failed muscle-equilibrium checks because the routed lengths were incompatible with the retained human parameters. The subsequent human-model repairs described below belong to that reference-model loading path; they are not evidence that the BPA-routed model reproduces every human muscle force.

Joint directions and coordinate conventions. Single-actuator tests exposed reversed hip-flexion and hip-adduction actions. The corresponding hinge-axis directions were negated so that the muscle actions agreed with the adopted OpenSim conventions. The knee required a different treatment: its converted flexion-negative convention was retained. Rotating the knee without updating its coupled translational coordinates gives a misleading geometry comparison, so those coordinates must follow the prescribed coupling during a pose or actuator test. These are specific joint and path checks, rather than a change from a left-handed to a right-handed model.

Small muscles and force scales. The quadratus femoris, gemelli, and piriformis actuators were disabled on both sides. Their six entries remain in the 92-actuator model, with no force contribution. A separate check found eight muscle-actuator entries with an active maximum-force scale of only 1 N. Their active-force scales were restored from the source-model values. Table D9 lists the corrected values; each applies separately to the left and

right actuator.

Table D9: Active maximum-force scales restored in the converted human model.

Muscle	Converted scale (N)	Restored scale (N)
Erector spinae	1	2500
Internal oblique	1	900
External oblique	1	900
Extensor hallucis longus	1	162

These are source-model maximum-force parameters, not measured BPA capacities or the realized force at every pose. The passive-force scales already contained the source values. The converted length ranges for these entries remain a limitation, so restoring the active scale alone does not establish full force–length agreement.

Patella-related paths and coupled joints. The converted moving muscle-path points use joints coupled to knee angle. Their sampled limits conflicted with the coupling constraints away from the sampled pose, producing large constraint forces and restricted motion. The coupled path-point joints were therefore left without those conflicting limits while retaining their prescribed trajectories. The anatomical driver joints were treated separately: limits were re-enabled on 14 hinges using the converted ranges. The rectus femoris route also lacked the patella tracking needed for knee-extension action; its intermediate path point was replaced by the vastus medialis patella-tracking point. The force comparison below quantifies the remaining differences at matched poses and activations.

Contact and virtual geometry. The converted model contained 19 explicit ground–body contact pairs, independent of geometry contact flags. Suspended tests remove these

pairs; ground-contact tests retain only the six foot–ground pairs. Path-marker geometry is noncolliding. These settings isolate intended foot loading from unintended contact with other model components.

These checks verify specific joint actions, coupled paths, and intended contacts. They retain the rigid-tendon approximation and do not establish full OpenSim equivalence or walking; the matched-pose force comparison below quantifies the remaining discrepancies.

D.6.1 Cross-Platform Force Comparison

The comparison script reads OpenSim static-optimization forces and activations and assigns the corresponding poses and activations to the converted MuJoCo model. It calls the forward calculation at each pose; it does not assign the trajectory velocities or integrate muscle activation and fiber dynamics through time. OpenSim static optimization likewise solves frame-by-frame force allocation rather than a compliant-tendon forward simulation [**?**]. The comparison therefore tests force predictions along sampled poses at matched activation, not dynamic equivalence or tendon-energy storage.

Table D10 retains the saved per-muscle statistics: $R^2 = 1 - \mathrm{SSE}/\mathrm{SST}$, squared Pearson correlation, and the best-aligning lag in analysis frames. The calculation and saved values are in `Code/MuJoCo_SNS/spinal/_muscle_force_compare.py` and `muscle_force_compare.csv`.

Eight of the twelve representative muscles have R^2 between 0.87 and 0.98 and mean-force ratios within 10 %. Soleus and semimembranosus have lower agreement, and the negative R^2 values for rectus femoris and biceps femoris short head indicate errors larger

Table D10: MuJoCo vs. OpenSim muscle force along the subject01 walk at matched activations. Lag in analysis frames (≈ 17 ms each).

Muscle	R^2	shape corr^2	lag [frames]
soleus	0.742	0.800	1
med. gastrocnemius	0.888	0.894	2
lat. gastrocnemius	0.884	0.887	2
tibialis anterior	0.945	0.945	0
vastus lateralis	0.941	0.945	1
rectus femoris	-1.15	0.384	0
semimembranosus	0.565	0.621	0
biceps femoris short	-0.26	0.868	0
psoas	0.927	0.934	1
gluteus maximus (2)	0.926	0.943	1
tensor fasciae latae	0.933	0.935	0
peroneus brevis	0.977	0.989	0

than predicting the corresponding mean reference force. Shape correlation alone does not remove those errors. The rectus femoris path was altered during conversion, while the biceps femoris short-head result has a substantial force-scale difference. These observations identify muscles requiring closer inspection; they do not isolate routing, force-law, and velocity effects from one another.

The frame-lag calculation describes alignment of these saved sequences. It cannot establish a physical delay caused by tendon compliance because the comparison does not simulate that mechanism. The results support selected pose-wise force checks, with the reported exceptions, rather than a general claim of equivalence for neural-control studies. A dynamic comparison would require matched velocities, activation histories, and explicitly specified tendon dynamics.

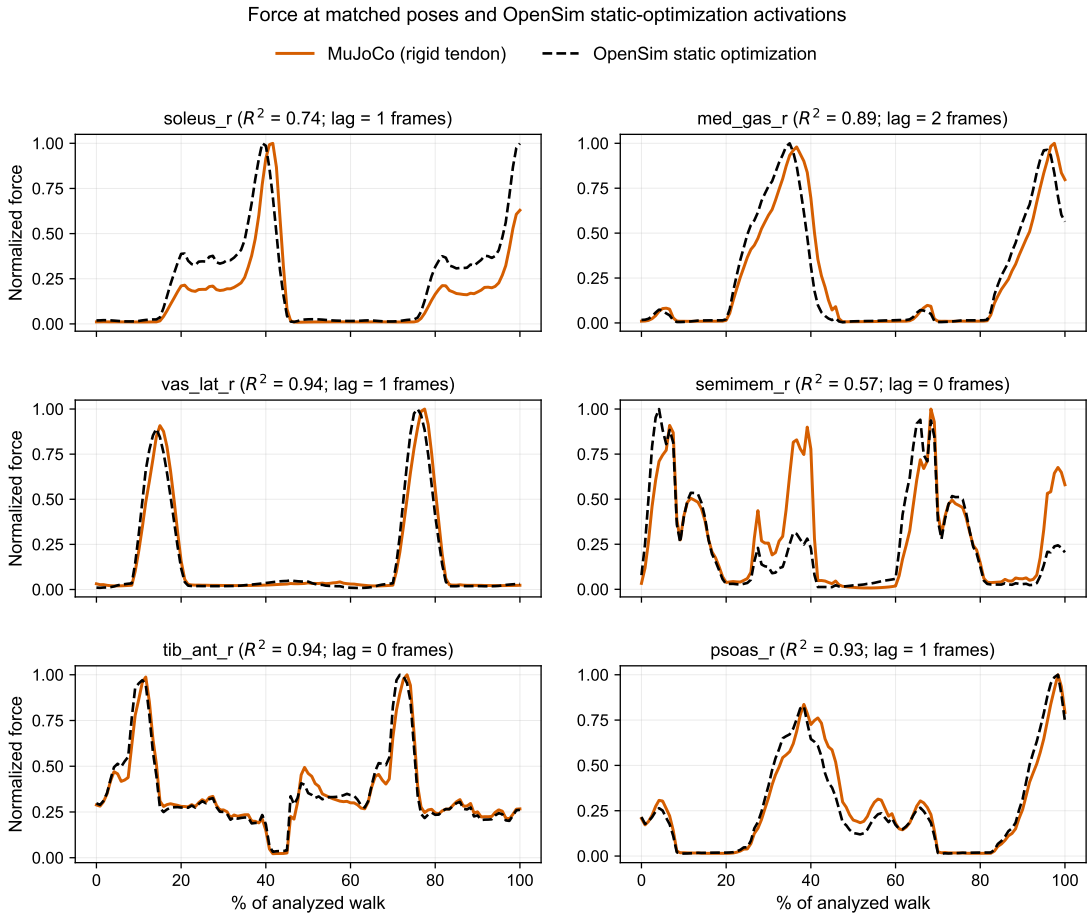


Figure D10: Force sequences at matched poses and activations, normalized to each curve's peak. Absolute-force agreement is reported in Table D10.

Appendix E: Supporting SADb Records

This appendix records the curation snapshot and pathway examples supporting Chapter 5. Counts refer to the stated audited or retained set; they are not estimates of the size of the neuroscience literature.

E.1 Record Sets and Coverage

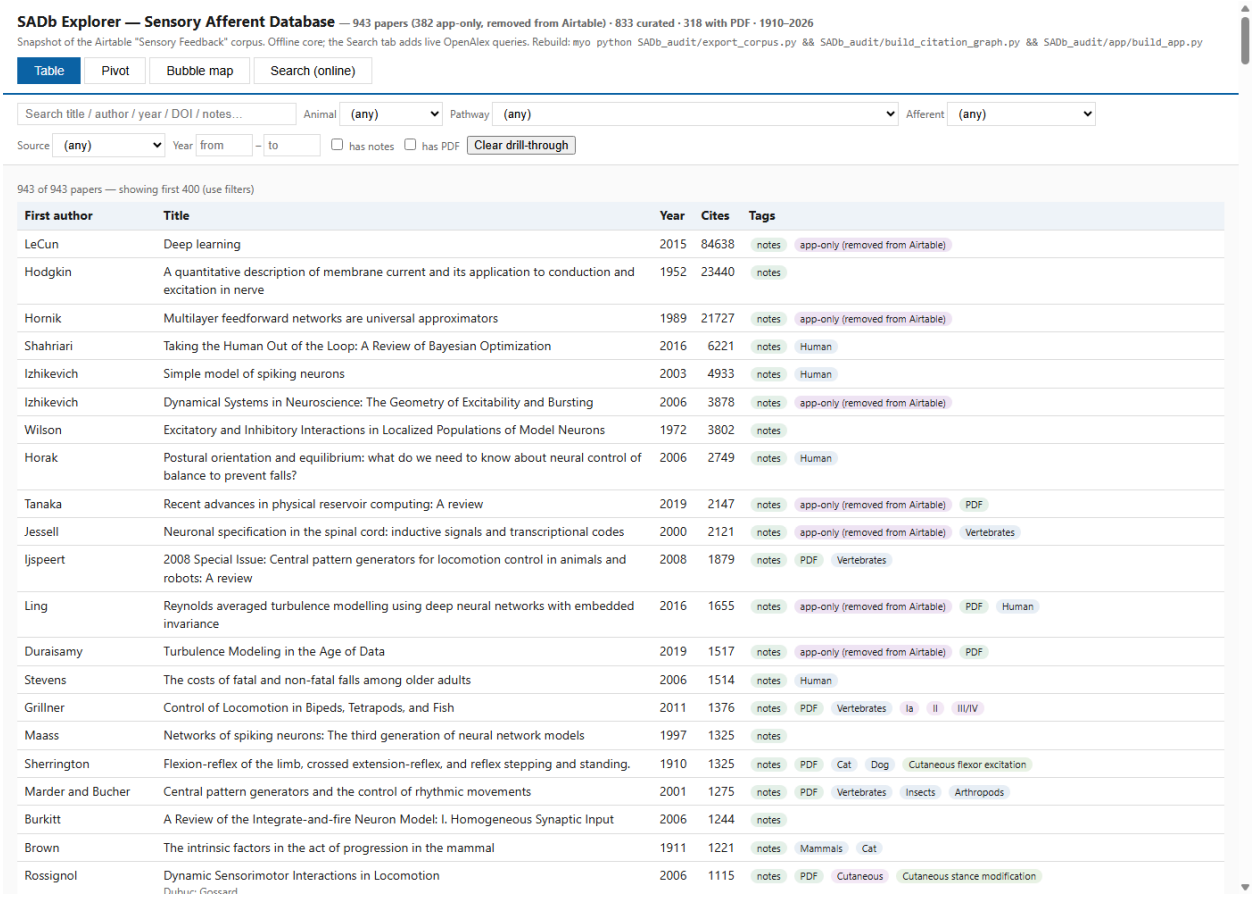
The application contains 1,393 records from the five-table base: 943 Papers, 247 Models, 157 Review Papers, 42 Feedback pathways, and 4 Recorders. Of these, 436 are in the app but not in Airtable due to their 1000 entry limitation. The retained working base in Airtable contains 957 records: 561 Papers, 244 Models, 109 Review Papers, 39 Feedback pathways, and 4 Recorders. The application table includes all record types, while its pivot and citation views use the paper set.

Coverage of the application's Papers category is reported in Chapter 5: among its 943 entries, 833 have curation notes, 606 preparation tags, 316 animal-study assessments, 119 pathway links, 95 afferent classifications, 363 robot/simulation translation notes, 892 DOIs, and 318 attached-PDF indicators. The following retained-set counts describe the maintenance subset only and use a different denominator.

Among the 561 retained papers, 523 carry curation notes, 426 preparation tags, 240 notes assessing animal-study potential, 118 links to named feedback pathways, 86 an afferent classification, and 258 robot/simulation translation notes. There are 550 DOIs; 38 records remain without grounding text. These fields overlap and should not be summed

as disjoint categories.

The citation graph covers the 943-paper audit: 730 connected papers, 11,130 directed in-corpus citation pairs, and 18 Louvain groups. Citation proximity is a navigation aid, not mechanistic evidence. Chapter 5 shows the citation view (Figure 5.1); Figure E1 shows the paper table.



First author	Title	Year	Cites	Tags
LeCun	Deep learning	2015	84638	notes app-only (removed from Airtable)
Hodgkin	A quantitative description of membrane current and its application to conduction and excitation in nerve	1952	23440	notes
Hornik	Multilayer feedforward networks are universal approximators	1989	21727	notes app-only (removed from Airtable)
Shahriari	Taking the Human Out of the Loop: A Review of Bayesian Optimization	2016	6221	notes Human
Izhikevich	Simple model of spiking neurons	2003	4933	notes Human
Izhikevich	Dynamical Systems in Neuroscience: The Geometry of Excitability and Bursting	2006	3878	notes app-only (removed from Airtable)
Wilson	Excitatory and Inhibitory Interactions in Localized Populations of Model Neurons	1972	3802	notes
Horak	Postural orientation and equilibrium: what do we need to know about neural control of balance to prevent falls?	2006	2749	notes Human
Tanaka	Recent advances in physical reservoir computing: A review	2019	2147	notes app-only (removed from Airtable) PDF
Jessell	Neuronal specification in the spinal cord: inductive signals and transcriptional codes	2000	2121	notes app-only (removed from Airtable) Vertebrates
Ijspeert	2008 Special Issue: Central pattern generators for locomotion control in animals and robots: A review	2008	1879	notes PDF Vertebrates
Ling	Reynolds averaged turbulence modelling using deep neural networks with embedded invariance	2016	1655	notes app-only (removed from Airtable) PDF Human
Duraisamy	Turbulence Modeling in the Age of Data	2019	1517	notes app-only (removed from Airtable) PDF
Stevens	The costs of fatal and non-fatal falls among older adults	2006	1514	notes Human
Grillner	Control of Locomotion in Bipeds, Tetrapods, and Fish	2011	1376	notes PDF Vertebrates Ia II III/IV
Maass	Networks of spiking neurons: The third generation of neural network models	1997	1325	notes
Sherrington	Flexion-reflex of the limb, crossed extension-reflex, and reflex stepping and standing.	1910	1325	notes PDF Cat Dog Cutaneous flexor excitation
Marder and Bucher	Central pattern generators and the control of rhythmic movements	2001	1275	notes PDF Vertebrates Insects Arthropods
Burkitt	A Review of the Integrate-and-fire Neuron Model: I. Homogeneous Synaptic Input	2006	1244	notes
Brown	The intrinsic factors in the act of progression in the mammal	1911	1221	notes Mammals Cat
Rossignol	Dynamic Sensorimotor Interactions in Locomotion	2006	1115	notes PDF Cutaneous Cutaneous stance modification

Figure E1: SADb Explorer paper-table view. Searches and filters expose bibliographic information, curation notes, preparation tags, and afferent or pathway links. The embedded corpus is available offline.

E.2 Afferent Classes and Linked Pathways

The tables reproduce the curated signal summaries and representative linked pathways. Their record names are database labels; pathway action depends on the experimental preparation and locomotor state described in the cited sources. Translating these records into a controller requires an explicit encoding and verification step.

Table E1: Afferent classes, encoded signals, and pathway roles represented in the Sensory Afferent Database. References link to corpus records; Ia, II, and Ib signal summaries follow **?**.

Class	Receptor	Signal carried	Recorded pathway roles (anchors)
Ia	muscle spindle primary	muscle length and velocity	homonymous stretch excitation [**?**]; reciprocal inhibition of the antagonist pool [**?**]
II	muscle spindle secondary	static muscle length	midlumbar disynaptic relay onto motor nuclei [**?**]; stance-phase support of extensor activity [**?**]
Ib	Golgi tendon organ	whole-muscle force	autogenic inhibition at rest, stance-phase reversal to excitation, and the load gate on swing onset [**??**]
III/IV	free nerve endings	metabolic and nociceptive signals, converging with cutaneous and group II–III afferents as flexor-reflex afferents	phase-dependent resetting of the locomotor rhythm [**?**]
Cutaneous	skin and foot-pad receptors	contact and perturbation events	flexor excitation and phase-modulated stance correction [**??**]

Table E2: Principal feedback pathways, named by afferent, action, and phase, with linked demonstrating papers. Counts are identical in the audited 943-paper corpus and retained 561-paper corpus.

Pathway record	Class	Recorded action in the step cycle	Linked papers	Demonstrating references
Ia monosynaptic excitation	Ia	stretch excitation of the homonymous motor pool, with the pathway gain modulated over the step cycle	7	[??]
Ia reciprocal inhibition	Ia	inhibition of the antagonist motor pool through the Ia inhibitory interneurons, rhythmically active in fictive locomotion	6	[??]
Ib disynaptic inhibition	Ib	autogenic non-reciprocal group I inhibition of the loaded muscle's own pool, reduced during stance	5	[??]
Ib disynaptic excitation	Ib	the locomotor reversal: group I input excites synergist extensor pools during the extensor phase, replacing the resting inhibition	6	[??]
Ib stance to swing	Ib	the load gate: Ib input from ankle extensors prolongs stance, so swing is not initiated while the leg is loaded	17	[??]
Ia or II stance to swing	Ia, II	flexor-length gate on stance termination: lengthening of flexor muscles in late stance ends extension and resets the cycle	10	[??]
type II excitatory	II	stance-phase support of ankle-extensor activity through the midlumbar group II relay	5	[??]
Cutaneous flexor excitation	Cutaneous	flexion withdrawal and the stumbling corrective reaction, phase-modulated over the cycle	16	[??]